

GNUPRO TOOLKIT™

GNUPro Compiler Tools

98r2

*Using GNU CC
The C Preprocessor*

CYGNUS

Copyright © 1988-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the sections entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Published by: **Free Software Foundation**
59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

Part #: 300-400-10100042-98r2

GNUPro warranty

The GNUPro Toolkit is free software, covered by the GNU General Public License, and you are welcome to change it and/or distribute copies of it under certain conditions. This version of GNUPro Toolkit is supported for customers of Cygnus.

For non-customers, GNUPro Toolkit software has NO WARRANTY.

Because this software is licensed free of charge, there are no warranties for it, to the extent permitted by applicable law. Except when otherwise stated in writing, the copyright holders and/or other parties provide the software “as is” without warranty of any kind, either expressed or implied, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose. The entire risk as to the quality and performance of the software is with you. Should the software prove defective, you assume the cost of all necessary servicing, repair or correction.

In no event, unless required by applicable law or agreed to in writing, will any copyright holder, or any other party who may modify and/or redistribute the program as permitted above, be liable to you for damages, including any general, special, incidental or consequential damages arising out of the use or inability to use the program (including but not limited to loss of data or data being rendered inaccurate or losses sustained by you or third parties or a failure of the program to operate with any other programs), even if such holder or other party has been advised of the possibility of such damages.

How to contact Cygnus

Use the following means to contact Cygnus.

Cygnus Headquarters

1325 Chesapeake Terrace
Sunnyvale, CA 94089 USA
Telephone (toll free): +1 800 CYGNUS-1
Telephone (main line): +1 408 542 9600
Telephone (hotline): +1 408 542 9601
FAX: +1-408 542 9699
(Faxes are answered 8 a.m.–5 p.m., Monday through Friday.)
email: info@cygnus.com
Website: www.cygnus.com.

Cygnus United Kingdom

36 Cambridge Place
Cambridge CB2 1NS
United Kingdom
Telephone: +44 1223 728728
FAX: +44 1223 728728
email: info@cygnus.co.uk/

Cygnus Japan

Nihon Cygnus Solutions
Madre Matsuda Building
4-13 Kioi-cho Chiyoda-ku
Tokyo 102-0094
Telephone: +81 3 3234 3896
FAX: +81 3 3239 3300
email: info@cygnus.co.jp
Website: <http://www.cygnus.co.jp/>

Use the hotline (+1 408 542 9601) to get help, although the most reliable and most expedient means to resolve problems with GNUPro Toolkit is by using the Cygnus Web Support site:

<http://support.cygnus.com>

Contents

GNUPro warranty	iii
How to contact Cygnus	iv
Using GNU CC	
<i>GNU General Public License</i>	3
Preamble.....	3
Terms and conditions for copying, distribution and modification	4
How to apply these terms to your new programs.....	9
<i>Contributors to GNU CC</i>	11
<i>Funding free software</i>	15
<i>Protect your freedom; fight “Look and Feel”</i>	17
<i>Introduction to the compiler for C, C++, or Objective C</i>	21
<i>Installing GCC</i>	23
Installing GCC on Unix systems.....	24
Configurations supported by GCC.....	33
Compilation in a separate directory	47
Building and installing a cross-compiler.....	48
Steps of cross-compilation	48
Configuring a cross-compiler.....	49
Tools and libraries for a cross-compiler	49
libgcc.a and cross-compilers	50

<i>Cross-compilers and header files</i>	51
Standard header file directories	52
Actually building the cross-compiler	53
collect2 and cross-compiling	54
Installing GCC on the Sun	55
Installing GCC on VMS	55
<i>Using GCC on VMS</i>	59
GNU CC command options	65
Option summary for GCC	67
<i>Overall options</i>	67
<i>C language options</i>	67
<i>C++ language options</i>	67
<i>Warning options</i>	68
<i>Debugging options</i>	68
<i>Optimization options</i>	68
<i>Preprocessor options</i>	68
<i>Assembler option</i>	69
<i>Linker options</i>	69
<i>Directory options</i>	69
<i>Target options</i>	69
<i>Machine dependent options</i>	69
<i>Code generation options</i>	74
Options controlling the kind of output	75
Options controlling C dialect	79
Options controlling C++ dialect	85
<i>Compiling C++ programs</i>	91
Options to request or suppress warnings	93
Options for debugging	103
Options that control optimization	111
Options controlling the preprocessor	117
Passing options to the assembler	121
Options for linking	123
Options for directory search	127
Specifying target machine and compiler version	129
Hardware models and configurations	131
AMD29K options	133
ARC options	135
ARM options	136

Clipper options	140
Convex options.....	141
D10V options	143
DEC Alpha options	144
Hitachi H8/300 options	147
Hitachi SH options	148
HPPA options	149
IBM RS/6000 and PowerPC options.....	151
IBM RT options.....	159
Intel x86 options.....	160
Intel 960 options.....	163
M32R/D options.....	165
MIPS options.....	166
MN10300 options.....	170
Motorola 68K options	171
Motorola 88K options	174
SPARC options.....	178
System V options	182
Thumb options	183
Vax options.....	185
<i>Options for code generation conventions.....</i>	<i>187</i>
<i>The offset-info option.....</i>	<i>193</i>
<i>Environment variables affecting GCC.....</i>	<i>195</i>
<i>Running the protoize program</i>	<i>197</i>
<i>Extensions to the C language family.....</i>	<i>201</i>
Statements and declarations in expressions	203
Locally declared labels.....	204
Labels as values.....	205
Nested functions.....	206
Constructing function calls	208
Naming an expression's type	209
Referring to a type with the typeof keyword.....	210
Generalized lvalues	211
Conditionals with omitted operands.....	213
Double-word integers.....	214
Complex numbers	215
Arrays of length zero.....	216
Arrays of variable length.....	217
Macros with variable numbers of arguments	219

Non-lvalue arrays may have subscripts.....	220
Arithmetic on <code>void</code> - and function-pointers	221
Non-constant initializers	222
Constructor expressions	223
Labeled elements in initializers.....	224
<i>Case ranges</i>	225
<i>Cast to a union type</i>	225
Declaring attributes of functions.....	226
Prototypes and old-style function definitions	232
Compiling functions for interrupt calls	233
<i>C++ style comments</i>	233
<i>Dollar signs in identifier names</i>	233
<i>The character <code>ESC</code> in constants</i>	233
Inquiring on alignment of types or variables	234
Specifying attributes of variables.....	235
Specifying attributes of types.....	239
An inline function is as fast as a macro.....	243
Assembler instructions with C expression operands.....	245
Constraints for <code>asm</code> operands.....	249
<i>Simple constraints</i>	249
<i>Multiple alternative constraints</i>	251
<i>Constraint modifier characters</i>	252
<i>Constraints for particular machines</i>	253
Controlling names used in assembler code	262
Variables in specified registers.....	263
<i>Defining global register variables</i>	263
<i>Specifying registers for local variables</i>	265
Alternate keywords	266
Incomplete <code>enum</code> types	267
Function names as strings	268
<i>Getting the return or frame address of a function</i>	268
<i>Extensions to the C++ language</i>	271
Named return values in C++	272
Minimum and maximum operators in C++.....	274
The <code>goto</code> and destructors in GNU C++	275
Declarations and definitions in one header	276
Where's the template?.....	278
Type abstraction using signatures	282
<i>gcov: a test coverage program</i>.....	285
Introduction to <code>gcov</code> test coverage	286

Invoking the <code>gcov</code> program	287
Using <code>gcov</code> with GCC optimization	290
Brief description of <code>gcov</code> data files	291
<i>Known causes of trouble with GCC</i>	293
Actual bugs we haven't fixed yet	294
Installation problems	294
Cross-compiler problems	299
Interoperation	300
Problems compiling certain programs	305
Incompatibilities of GCC	306
Fixed header files	309
Standard libraries	310
Disappointments and misunderstandings	310
Common misunderstandings with GNU C++	312
<i>Declare and define static members</i>	312
<i>Temporaries may vanish before you expect</i>	312
<code>protoize</code> and <code>unprotoize</code> warnings	313
Certain changes we don't want to make	315
Warning messages and error messages	318
<i>Reporting bugs</i>	319
Have you found a bug?	320
Where to report bugs	321
How to report bugs	322
Sending patches for GCC	327
<i>How to get help with GCC</i>	329

The C Preprocessor

<i>Overview of the C preprocessor</i>	333
What the C preprocessor provides	334
<i>Transformations made globally</i>	335
<i>Preprocessing directives</i>	337
<i>Header files</i>	339
Uses of header files	340
The <code>#include</code> directive	341
How <code>#include</code> works	343
Once-only include files	344
Inheritance and header files	345
<i>Macros</i>	347
Simple macros	348

Macros with arguments	350
Predefined macros	353
<i>Standard predefined macros</i>	353
<i>Non-standard predefined macros</i>	356
Stringification.....	358
Concatenation.....	360
Undefining macros	362
Redefining macros	363
Pitfalls and subtleties of macros.....	364
<i>Improperly nested constructs</i>	364
<i>Unintended grouping of arithmetic</i>	364
<i>Swallowing the semicolon</i>	365
<i>Duplication of side effects</i>	366
<i>Self-referential macros</i>	367
<i>Separate expansion of macro arguments</i>	367
<i>Cascaded use of macros</i>	370
<i>Newlines in macro arguments</i>	370
Conditionals	371
Why conditionals are used	372
Syntax of conditionals.....	373
<i>The #if directive</i>	373
<i>The #else directive</i>	374
<i>The #elif directive</i>	374
<i>Keeping deleted code for future reference</i>	375
Conditionals and macros	376
Assertions.....	378
The #error and #warning directives.....	380
Combining source files	381
Other preprocessing directives	383
C preprocessor output	385
Invoking the C preprocessor.....	387
Index	393

Using GNU CC

Copyright © 1988-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the following documentation entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Published by: **Free Software Foundation**
59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

GNU General Public License

Version 2, June 1991

Copyright © 1989, 1991 Free Software Foundation, Inc.
59 Temple Place / Suite 330, Boston, MA 02111-1307 USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software—to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you

these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps: (1) copyright the software, and (2) offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

Terms and conditions for copying, distribution and modification

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The "Program", below, refers to any such program or work, and a "work based on the Program" means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term "modification".) Each licensee is addressed as "you".

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty;

keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:

You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.

You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.

If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3. You may copy and distribute the Program (or a work based on it, under Section 2)

in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:

- a. Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- b. Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- c. Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.
5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and

all its terms and conditions for copying, distributing or modifying the Program or works based on it.

6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.
7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.
9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in

spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and “any later version”, you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.
12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

How to apply these terms to your new programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

one line: the program's name and a brief idea of what it does.
Copyright © 19yy name of author

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like the following example when it starts in an interactive mode:

Gnomovision version 69, Copyright © 19yy name of author Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type 'show w'. This is free software, and you are welcome to redistribute it under certain conditions; type 'show c' for details.

The hypothetical commands `show w` and `show c` should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than `show w` and `show c`; they can be mouse-clicks or menu items—whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a “copyright disclaimer” for the program, if necessary. The following is a sample (when copying, *alter the names*).

Yoyodyne, Inc., hereby disclaims all copyright interest in the program

How to apply these terms to your new programs

`'Gnomovision' (which makes passes at compilers) written by James
Hacker.`

`signature of Ty Coon, 1 April 1989`

`Ty Coon, President of Vice`

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.

Contributors to GNU CC

In addition to Richard Stallman, several people have written parts of the GNU compiler, GNU CC.

- The idea of using RTL and some of the optimization ideas came from the program PO written at the University of Arizona by Jack Davidson and Christopher Fraser. See “Register Allocation and Exhaustive Peephole Optimization”, *Software Practice and Experience* 14 (9), Sept. 1984, pages 857-866.
- Paul Rubin wrote most of the preprocessor.
- Leonard Tower wrote parts of the parser, RTL generator, and RTL definitions, and the Vax machine description.
- Ted Lemon wrote parts of the RTL reader and printer.
- Jim Wilson implemented loop strength reduction and some other loop optimizations.
- Nobuyuki Hikichi of Software Research Associates, Tokyo, contributed the support for the Sony NEWS machine.
- Charles LaBrec contributed the support for the Integrated Solutions 68020 system.
- Michael Tiemann of Cygnus wrote the front end for C++, as well as the support for inline functions and instruction scheduling. Also the descriptions of the National Semiconductor 32000 series CPU, the SPARC CPU and part of the

Motorola 88000 CPU.

- Gerald Baumgartner added the signature extension to the C++ front-end.
- Jan Stein of the Chalmers Computer Society provided support for Genix, as well as part of the 32000 machine description.
- Randy Smith finished the Sun FPA support.
- Robert Brown implemented the support for Encore 32000 systems.
- David Kashtan of SRI adapted GCC to VMS.
- Alex Crain provided changes for the 3b1.
- Greg Satz and Chris Hanson assisted in making GCC work on HP-UX for the 9000 series 300.
- William Schelter did most of the work on the Intel 80386 support.
- Christopher Smith did the port for Convex machines.
- Paul Petersen wrote the machine description for the Alliant FX/8.
- Dario Dariol contributed the four varieties of sample programs that print a copy of their source.
- Alain Lichnewsky ported GCC to the MIPS CPU.
- Devon Bowen, Dale Wiles and Kevin Zachmann ported GCC to the Tahoe.
- Jonathan Stone wrote the machine description for the Pyramid computer.
- Gary Miller ported GCC to Charles River Data Systems machines.
- Richard Kenner of the New York University Ultracomputer Research Laboratory wrote the machine descriptions for the AMD 29000, the DEC Alpha, the IBM RT PC, and the IBM RS/6000 as well as the support for instruction attributes. He also made changes to better support RISC processors including changes to common subexpression elimination, strength reduction, function calling sequence handling, and condition code support, in addition to generalizing the code for frame pointer elimination.
- Richard Kenner and Michael Tiemann jointly developed `reorg.c`, the delay slot scheduler.
- Mike Meissner and Tom Wood of Data General finished the port to the Motorola 88000.
- Masanobu Yuhara of Fujitsu Laboratories implemented the machine description for the Tron architecture (specifically, the Gmicro).
- NeXT, Inc. donated the front end that supports the Objective C language.
- James van Artsdalen wrote the code that makes efficient use of the Intel 80387 register stack.

- Mike Meissner at the Open Software Foundation finished the port to the MIPS CPU, including adding ECOFF debug support, and worked on the Intel port for the Intel 80386 CPU.
- Ron Guilmette implemented the `protoize` and `unprotoize` tools, the support for Dwarf symbolic debugging information, and much of the support for System V Release 4. He has also worked heavily on the Intel 386 and 860 support.
- Torbjorn Granlund implemented multiply- and divide-by-constant optimization, improved long long support, and improved leaf function register allocation.
- Mike Stump implemented the support for Elxsi 64 bit CPU.
- John Wehle added the machine description for the Western Electric 32000 processor used in several 3b series machines (no relation to the National Semiconductor 32000 processor).
- Holger Teutsch provided the support for the Clipper CPU.
- Kresten Krab Thorup wrote the run time support for the Objective C language.
- Stephen Moshier contributed the floating point emulator that assists in cross-compilation and permits support for floating point numbers wider than 64 bits.
- David Edelsohn contributed the changes to RS/6000 port to make it support the PowerPC and POWER2 architectures.
- Steve Chamberlain wrote the support for the Hitachi SH processor.
- Peter Schauer wrote the code to allow debugging to work on the Alpha.
- Oliver M. Kellogg of Deutsche Aerospace contributed the port to the MIL-STD-1750A.
- Michael K. Gschwind contributed the port to the PDP-11.

Funding free software

If you want to have more free software a few years from now, it makes sense for you to help encourage people to contribute funds for its development. The most effective approach known is to encourage commercial redistributors to donate.

Users of free software systems can boost the pace of development by encouraging for-a-fee distributors to donate part of their selling price to free software developers—the Free Software Foundation, and others.

The way to convince distributors to do this is to demand it and expect it from them. So when you compare distributors, judge them partly by how much they give to free software development. Show distributors they must compete to be the one who gives the most.

To make this approach work, you must insist on numbers that you can compare, such as, “We will donate ten dollars to the Frobnitz project for each disk sold.” Don’t be satisfied with a vague promise, such as “A portion of the profits are donated,” since it doesn’t give a basis for comparison.

Even a precise fraction “of the profits from this disk” is not very meaningful, since creative accounting and unrelated business decisions can greatly alter what fraction of the sales price counts as profit. If the price you pay is \$50, ten percent of the profit is probably less than a dollar; it might be a few cents, or nothing at all.

Some redistributors do development work themselves. This is useful too; but to keep everyone honest, you need to inquire how much they do, and what kind. Some kinds of development make much more long-term difference than others. For example, maintaining a separate version of a program contributes very little; maintaining the standard version of a program for the whole community contributes much. Easy new ports contribute little, since someone else would surely do them; difficult ports such as adding a new CPU to the GNU C compiler contribute more; major new features or packages contribute the most.

By establishing the idea that supporting further development is “the proper thing to do” when distributing free software for a fee, we can assure a steady flow of resources into making more free software.

Copyright © 1994 Free Software Foundation, Inc.

Verbatim copying and redistribution of this section is permitted without royalty; alteration is not permitted.

Protect your freedom; fight “Look and Feel”

This section is a political message from the League for Programming Freedom to the users of GCC. We have included it here because the issue of interface copyright is important to the GNU project.

Apple, Lotus, and now CDC have tried to create a new form of legal monopoly: a copyright on a user interface.

An interface is a kind of language—a set of conventions for communication between two entities, human or machine. Until a few years ago, the law seemed clear: interfaces were outside the domain of copyright, so programmers could program freely and implement whatever interface the users demanded. Imitating de-facto standard interfaces, sometimes with improvements, was standard practice in the computer field. These improvements, if accepted by the users, caught on and became the norm; in this way, much progress took place.

Computer users, and most software developers, were happy with this state of affairs. However, large companies such as Apple and Lotus would prefer a different system—one in which they can own interfaces and thereby rid themselves of all serious competitors. They hope that interface copyright will give them, in effect, monopolies on major classes of software.

Other large companies such as IBM and Digital also favor interface monopolies, for the same reason: if languages become property, they expect to own many de-facto standard languages. But Apple and Lotus are the ones who have actually sued.

Apple’s lawsuit was defeated, for reasons only partly related to the general issue of interface copyright.

Lotus won lawsuits against two small companies, which were thus put out of business. Then they sued Borland; they won in the trial court (no surprise, since it was the same court that had ruled for Lotus twice before), but the decision was reversed by the court of appeals, with help from the League for Programming Freedom in the form of a friend-of-the- court brief. We are now waiting to see if the Supreme Court will hear the case. If it does, the League for Programming Freedom will again submit a brief.

The battle is not over. Just this summer a company that produced a simulator for a CDC computer was shut down by a copyright lawsuit by CDC, which charged that the simulator infringed the copyright on the manuals for the computer.

If the monopolists get their way, they will hobble the software field:

- Gratuitous incompatibilities will burden users. Imagine if each car manufacturer had to design a different way to start, stop, and steer a car.
- Users will be “locked in” to whichever interface they learn; then they will be prisoners of one supplier, who will charge a monopolistic price.
- Large companies have an unfair advantage wherever lawsuits become commonplace. Since they can afford to sue, they can intimidate smaller developers with threats even when they don’t really have a case.
- Interface improvements will come slower, since incremental evolution through creative partial imitation will no longer occur.

If interface monopolies are accepted, other large companies are waiting to grab theirs:

- Adobe is expected to claim a monopoly on the interfaces of various popular application programs, if Lotus ultimately wins the case against Borland.
- Open Computing magazine reported a Microsoft vice president as threatening to sue people who imitate the interface of Windows.

Users invest a great deal of time and money in learning to use computer interfaces. Far more, in fact, than software developers invest in developing and even implementing the interfaces. Whoever can own an interface, has made its users into captives, and misappropriated their investment.

To protect our freedom from monopolies like these, a group of programmers and users have formed a grass-roots political organization, the League for Programming Freedom.

The purpose of the League is to oppose monopolistic practices such as interface copyright and software patents. The League calls for a return to the legal policies of the recent past, in which programmers could program freely. The League is not

concerned with free software as an issue, and is not affiliated with the Free Software Foundation.

The League’s activities include publicizing the issues, as is being done here, and filing friend-of-the-court briefs on behalf of defendants sued by monopolists.

The League’s membership rolls include Donald Knuth, the foremost authority on algorithms, John McCarthy, inventor of Lisp, Marvin Minsky, founder of the MIT Artificial Intelligence lab, Guy L. Steele, Jr., author of well-known books on Lisp and C, as well as Richard Stallman, the developer of GCC. Please join and add your name to the list. Membership dues in the League are \$42 per year for programmers, managers and professionals; \$10.50 for students; \$21 for others.

Activist members are especially important, but members who have no time to give are also important. Surveys at major ACM conferences have indicated a vast majority of attendees agree with the League on both issues (interface copyrights and software patents). If just ten percent of the programmers who agree with the League join the League, we will probably triumph.

To join, or for more information, phone (617) 243-4091 or write to:

League for Programming Freedom
1 Kendall Square #143
P.O. Box 9171
Cambridge, MA 02139

You can also send electronic mail to lpf@uunet.uu.net.

In addition to joining the League, here are some suggestions from the League for other things you can do to protect your freedom to write programs:

- Tell your friends and colleagues about this issue and how it threatens to ruin the computer industry.
- Mention that you are a League member in your ‘.signature,’ and mention the League’s email address for inquiries.
- Ask the companies you consider working for or working with to make statements against software monopolies, and give preference to those that do.
- When employers ask you to sign contracts giving them copyright on your work, insist on a clause saying they will not claim the copyright covers imitating the interface.
- When employers ask you to sign contracts giving them patent rights, insist on clauses saying they can use these rights only defensively. Don’t rely on “company policy,” since policies can change at any time; don’t rely on an individual executive’s private word, since that person may be replaced. Get a commitment just as binding as the commitment they get from you.

- Write to Congress to explain the importance of these issues.

House Subcommittee on Intellectual Property

2137 Rayburn Building

Washington, DC 20515

Senate Subcommittee on Patents, Trademarks and Copyrights

United States Senate

Washington, DC 20510

(These committees have received lots of mail already; let's give them even more.)

Democracy means nothing if you don't use it. Stand up and be counted!

Introduction to the compiler for C, C++, or Objective C

The GNU compiler is for compiling programs written in C, C++, or Objective C, which are all integrated. “GCC” is a common shorthand term for the GNU C compiler and the name used when the emphasis is on compiling C programs.

IMPORTANT: With the GNUPro tools, the documentation for “Installing GCC” on page 23 serves as an introduction to the installation of GCC, only necessary for reconfiguring.

The following documentation discusses some of the essentials within GCC.

- “GNU CC command options” on page 65
- “Extensions to the C language family” on page 201
- “Extensions to the C++ language” on page 271
- “gcov: a test coverage program” on page 285
- “Known causes of trouble with GCC” on page 293
- “Reporting bugs” on page 319

The following documentation discusses the Free Software Foundation, the community behind the essentials of the GNU tools (the compiler, the debugger, the assembler, the linker, as well as many other GNU tools).

- “GNU General Public License” on page 3
- “Contributors to GNU CC” on page 11
- “Funding free software” on page 15

- “Protect your freedom; fight “Look and Feel”” on page 17

We use the name *GCC* to refer to the compilation system as a whole, and more specifically to the language-independent part of the compiler. For example, we refer to the optimization options as affecting the behavior of GCC, or sometimes as affecting just the *compiler*.

When referring to C++ compilation, it is customary to call the compiler G++. Since there is only one compiler, it is also accurate to call it `gcc`, no matter what the language context; however, the term, G++, is more useful when the emphasis is on compiling C++ programs.

Front ends for other languages, such as Ada 9X, Fortran, Modula-3, and Pascal, are or have been in development. These front-ends, like that for C++, are built in subdirectories of GCC and link to it. The result is an integrated compiler that can compile programs written in C, C++, Objective C, or any of the languages for which you have installed front-ends.

In all the GNUPro documentation, we only discuss the options for the C, Objective C, and C++ compilers, and those of the GCC core. Consult the documentation of the other front ends for the options to use when compiling programs or when using tools that are written in other languages.

G++ is a *compiler*, not merely a preprocessor. G++ builds object code directly from your C++ program source. There is no intermediate C version of the program. (By contrast, for example, some other implementations use a program that generates a C program from your C++ source.) Avoiding an intermediate C representation of the program means that you get better object code, and better debugging information. The GNU debugger, GDB, works with this information in the object code to give you comprehensive C++ source-level editing capabilities (see “Using GDB with different languages” on page 109 and “C and C++” on page 116 in *Debugging with GDB in GNUPro Debugging Tools*).

Installing GCC

The following documentation discusses installation and what you need to get the GNU compiler, GCC, to work for you.

- “Installing GCC on Unix systems” on page 24
- “Configurations supported by GCC” on page 33
- “Compilation in a separate directory” on page 47
- “Building and installing a cross-compiler” on page 48
- “Steps of cross-compilation” on page 48
- “Standard header file directories” on page 52
- “Actually building the cross-compiler” on page 53
- “collect2 and cross-compiling” on page 54

See also “Installing GCC on the Sun” on page 55 and “Installing GCC on VMS” on page 55.

Installing GCC on Unix systems

What follows is the procedure for installing GCC on a UNIX system. In this documentation, we assume you compile in the same directory that contains the source files; see “Compilation in a separate directory” on page 47 to find out how to compile in a separate directory on UNIX systems. For VMS systems, see “Installing GCC on VMS” on page 55. For Windows, you need to get the complete compilation package, DJGPP, which includes binaries as well as sources along with all the necessary compilation tools and libraries.

1. If you have built GCC previously in the same directory for a different target machine, do `make distclean` to delete all files that might be invalid. One of the files `make distclean` deletes is `Makefile`; if `make distclean` complains that `Makefile` does not exist, it probably means that the directory is already suitably clean.
2. On a System V release 4 system, make sure `/usr/bin` precedes `/usr/ucb` in `PATH`. The `cc` command in `/usr/ucb` uses libraries which have bugs.
3. Specify the host, build and target machine configurations. You do this by running the file, `configure`.

The *build* machine is the system which you are using, the *host* machine is the system where you want to run the resulting compiler (normally the build machine), and the *target* machine is the system for which you want the compiler to generate code.

If you are building a compiler to produce code for the machine it runs on (a native compiler), you normally do not need to specify any operands to `configure`; it will try to guess the type of machine you are on and use that as the build, host and target machines. So you don't need to specify a configuration when building a native compiler unless `configure` cannot figure out what your configuration is or guesses wrong. In those cases, specify the build machine's *configuration name* with the option, `--build`; the host and target will default to be the same as the build machine. (If you are building a cross-compiler, see “Building and installing a cross-compiler” on page 48.) The following is an example.

```
./configure--build=sparc-sun-sunos4.1
```

A configuration name may be canonical or it may be more or less abbreviated. A canonical configuration name has three parts, separated by dashes, like the following example where *cpu* designates the processor, *company* designates the company who makes the processor, and *system* designates the actual system for which the processor is configured.

cpu-company-system

The three parts may themselves contain dashes; `configure` can figure out which dashes serve which purpose. For example, `m68k-sun-sunos4.1` specifies a Sun 3.

You can also replace parts of the configuration by nicknames or aliases. For

example, `sun3` stands for `m68k-sun`, so `sun3-sunos4.1` is another way to specify a Sun 3. You can also use simply `sun3-sunos`, since the version of SunOS is assumed by default to be version 4. You can specify a version number after any of the system types, and some of the CPU types. In most cases, the version is irrelevant, and will be ignored. So you might as well specify the version if you know it.

WARNING: See “Configurations supported by GCC” for the supported configuration names and notes on many of the configurations. See also “Version numbers for programs” on page 52, “Embedded cross-configuration support” on page 50 and “Naming hosts and targets” on page 53 in *Introduction* within **Getting Started with GNUPro Toolkit**.

There are four additional options you can specify independently to describe variant hardware and software configurations.

These are: `--with-gnu-as`, `--with-gnu-ld`, `--with-stabs` and `--nfp`.

`--with-gnu-as`

If you will use GCC with the GNU assembler (GAS), you should declare this by using the `--with-gnu-as` option when you run `configure`.

Using this option does not install GAS. It only modifies the output of GCC to work with GAS. Building and installing GAS is up to you.

Conversely, if you do *not* wish to use GAS *and* do *not* specify `--with-gnu-as` when building GCC, it is up to you to make sure that GAS is *not* installed.

GCC searches for a program named `as` in various directories; if the program it finds is GAS, then it runs GAS. If you are not sure where GCC finds the assembler it is using, try specifying `-v` when you run it.

The systems where it makes a difference whether you use GAS are

`hppa1.0-any-any`, `hppa1.1-any-any`, `i386-any-sysv`, `i386-any-isc`, `i860-any-bsd`, `m68k-bull-sysv`, `m68k-hp-hpux`, `m68k-sony-bsd`, `m68k-altos-sysv`, `m68000-hp-hpux`, `m68000-att-sysv`, `any-lynx-lynxos` and `mips-any`. (*any* in these designations refers to any version of the specified configuration.) On any other system, `--with-gnu-as` has no effect.

On the previously listed systems (except for the HPPA, for ISC on the 386, and for `mips-sgi-irix5.*`), if you use GAS, you should also use the GNU linker (specifying with the option, `--with-gnu-ld`).

`--with-gnu-ld`

Specify the option, `--with-gnu-ld`, if you plan to use the GNU linker with GCC.

This option does not cause the GNU linker to be installed; it just modifies the behavior of GCC to work with the GNU linker. Specifically, it inhibits the installation of `collect2`, a program which otherwise serves as a front-end for

the system's linker on most configurations.

`--with-stabs`

On MIPS based systems and on Alphas, you must specify whether you want GCC to create the normal ECOFF debugging format, or to use BSD-style stabs passed through the ECOFF symbol table. The normal ECOFF debug format cannot fully handle languages other than C. BSD stabs format can handle other languages, but it only works with the GNU debugger GDB.

Normally, GCC uses the ECOFF debugging format by default; if you prefer BSD stabs, specify `--with-stabs` when you configure GCC.

No matter which default you choose when you configure GCC, you can use the `-gcoff` and `-gstabs+` options to specify explicitly the debug format for a particular compilation.

`--with-stabs` is meaningful on the ISC system on the 386 as well as if using `--with-gas`. It selects use of stabs debugging information embedded in COFF output. This kind of debugging information supports C++ well; ordinary COFF debugging information does not.

`--with-stabs` is also meaningful on 386 systems running SVR4. It selects use of stabs debugging information embedded in ELF output. The C++ compiler currently (2.6.0) does not support the DWARF debugging information normally used on 386 SVR4 platforms; stabs provide a workable alternative. This requires `gas` and `gdb`, as the normal SVR4 tools can not generate or interpret stabs.

`--nfp`

On certain systems, you must specify whether the machine has a floating point unit. These systems include `m68k-sun-sunosn` and `m68k-isi-bsd`. On any other system, `--nfp` currently has no effect, though perhaps there are other systems where it could usefully make a difference.

`--enable-objthreads=type`

Certain systems, notably Linux, can't be relied on to supply a threads facility for the Objective C runtime and so will default to single-threaded runtime. They may, however, have a library threads implementation available, in which case threads can be enabled with this option by supplying a suitable *type*, probably `'posix'`. The possibilities for *type* are `'single'`, `'posix'`, `'win32'`, `'solaris'`, `'irix'` and `'mach'`.

`--enable-leading-underscore`

`--disable-leading-underscore`

`--enable-leading-underscore` forces generation of underscores for assembly output using `a.out` object file format to override the default user label prefix used by the compiler and recognized by the binary utilities and the debugger. `--disable-leading-underscore` will do the opposite.

The `configure` script searches subdirectories of the source directory for other compilers that are to be integrated into GCC.

The GNU compiler for C++, called `g++`, is in a subdirectory named `cp`. `configure` inserts rules into `Makefile` to build all of those compilers.

In the following, we clarify which files will be set up by `configure`. Normally you need not be concerned with these files.

- A file named `config.h` is created that contains a `#include` of the top-level configuration file for the machine you will run the compiler on (for a specific CPU and system, see “Configurations supported by GCC” on page 33).

This file is responsible for defining information about the host machine. It includes `tm.h`.

The top-level configuration file is located in the subdirectory, `config`. Its name is always `xm- something.h`; usually, `xm- machine.h`, but there are some exceptions.

If your system does not support symbolic links, you might want to set up `config.h` to contain a `#include` command which refers to the appropriate file.

- A file named `tconfig.h` is created which includes the top-level configuration file for your target machine. This is used for compiling certain programs to run on that machine.
- A file named `tm.h` is created which includes the machine-description macro file for your target machine. It should be in the subdirectory `config` with the specific name of the system followed by `.h` (for a specific CPU and system, see “Configurations supported by GCC” on page 33).
- The command file, `configure`, also constructs the file, `Makefile`, by adding some text to the template file, `Makefile.in`. The additional text comes from files in the `config` directory, named `t-target` and `x-host`. *If these files do not exist, it means nothing needs to be added for a given target or host.*

4. The standard directory for installing GCC is `/usr/local/lib`.

If you want to install its files somewhere else, specify `--prefix=dir` when you run `configure`. `dir` is a directory name to use instead of `/usr/local` for all purposes with one exception: the directory `/usr/local/include` is searched for header files no matter where you install the compiler. To override this name, use the `--local-prefix` option in the following documentation, Step 5.

5. Specify `--local-prefix=dir` if you want the compiler to search directory, `dir/include`, for locally installed header files instead of `/usr/local/include`.

You should specify `--local-prefix` only if your site has a different convention

(not `/usr/local`) for where to put site-specific files.

The default value for `--local-prefix` is `/usr/local` regardless of the value of `--prefix`. Specifying `--prefix` has no effect on which directory GCC searches for local header files. This may seem counterintuitive, but actually it is logical. The purpose of `--prefix` is to specify where to install GCC. The local header files in `/usr/local/include`—if you put any in that directory—are not part of GCC. They are part of other programs—perhaps many others. (GCC installs its own header files in another directory which is based on the `--prefix` value.)

WARNING: Do not specify `/usr` as the `--local-prefix`!

The directory you use for `--local-prefix` *must not* contain any of the system's standard header files. If it did contain them, certain programs would be miscompiled (including GNU Emacs, on certain targets), because this would override and nullify the header file corrections made by the `fixincludes` script.

Indications are that people use this option use it based on mistaken ideas of its purpose. They use it as if it specified where to install GCC, perhaps on the assumption that installing GCC creates this directory.

6. Make sure the Bison parser generator is installed. (This is unnecessary if the Bison output files `c-parse.c` and `cexp.c` are more recent than `c-parse.y` and `cexp.y` and you do not plan to change the `.y` files.)
Bison versions older than Sept 8, 1988 will produce incorrect output for `c-parse.c`.
7. If you have chosen a configuration for GCC which requires other GNU tools (such as GAS or the GNU linker) instead of the standard system tools, install the required tools in the build directory under the names `as`, `ld`, or whatever is appropriate. This will enable the compiler to find the proper tools for compilation of the program `enquire`.

Alternatively, you can do subsequent compilation using a value of the `PATH` environment variable such that the necessary GNU tools come before the standard system tools.

8. Build the compiler. Just type `make LANGUAGES=c` in the compiler directory. `LANGUAGES=c` specifies that only the C compiler should be compiled. The Makefile normally builds compilers for all the supported languages; currently, C, C++ and Objective C. However, C is the only language that is sure to work when you build with other non-GNU C compilers. In addition, building anything but C, at this stage, is a waste of time.

In general, you can specify the languages to build by typing the argument `LANGUAGES="list"` where *list* is one or more words from the list `c`, `c++`, and

`objective-c`. If you have any additional GNU compilers as subdirectories of the GCC source directory, you may also specify their names in this list.

Ignore any warnings you may see about “`statement not reached`” in ‘`insn-emit.c`’; they are normal.

Also, warnings about “`unknown escape sequence`” are normal in ‘`genopinit.c`’ and perhaps some other files. Likewise, you should ignore warnings about “`constant is so large that it is unsigned`” in ‘`insn-emit.c`’ and ‘`insn-recog.c`’ and a warning about a comparison always being zero in ‘`enquire.o`’. Any other compilation errors may represent bugs in the port to your machine or operating system, and should be investigated and reported (see “General licenses and terms for using GNUPro Toolkit” on page 15 in *Installation* within ***Getting Started with GNUPro Toolkit***). Some commercial compilers fail to compile GCC because they have bugs or limitations. For example, the Microsoft compiler is said to run out of macro space. Some Ultrix compilers run out of expression space; then you need to break up the statement where the problem happens.

9. If you are building a cross-compiler, stop here. See “Building and installing a cross-compiler” on page 48.
10. Move the first-stage object files and executables into a subdirectory with the following command:

```
make stage1
```

The files are moved into a subdirectory named `stage1`. Once installation is complete, you may wish to delete these files with `rm -r stage1`.

11. If you have chosen a configuration for GCC which requires other GNU tools (such as GAS or the GNU linker) instead of the standard system tools, install the required tools in the `stage1` subdirectory under the names `as`, `ld` or whatever is appropriate. This will enable the stage 1 compiler to find the proper tools in the following stage.

Alternatively, you can do subsequent compilation using a value of the `PATH` environment variable such that the necessary GNU tools come before the standard system tools.

12. Recompile the compiler with itself, with this command:

```
make CC="stage1/xgcc -Bstage1/" CFLAGS="-g -O2"
```

This is called making the stage 2 compiler. The command shown in the previous example builds compilers for all the supported languages. If you don’t want them all, you can specify the languages to build by typing the argument, `LANGUAGES="list"`, where `list` should contain one or more of the words, `c`, `c++`, `objective-c` or `proto`. Separate the words with spaces. `proto` stands for the programs `protoize` and `unprotoize`; they are not a separate language, but you use `LANGUAGES` to enable or disable their installation. If you are going to build the stage 3 compiler, then you might want to build only the C language in stage 2.

Once you have built the stage 2 compiler, if you are short of disk space, you can delete the subdirectory `stage1`. On a 68000 or 68020 system lacking floating point hardware, unless you have selected a `'tm.h'` file that expects by default to find no hardware; instead, use the following statement.

```
make CC="stage1/xgcc -Bstage1/" CFLAGS="-g -O2 -msoft-float"
```

- 13.** If you wish to test the compiler by compiling it with itself one more time, install any other necessary GNU tools (such as GAS or the GNU linker) in the `stage2` subdirectory as you did in the `stage1` subdirectory; then, use the following. This is called making the stage 3 compiler.

```
make stage2
make CC="stage2/xgcc -Bstage2/" CFLAGS="-g -O2"
```

Aside from the `'-B'` option, the compiler options should be the same as when you made the stage 2 compiler. But the `LANGUAGES` option need not be the same. The command in the previous example builds compilers for all the supported languages; if you don't want them all, you can specify the languages to build by typing the argument `'LANGUAGES=" list "'` as described in Step 8. If you don't have to install any additional GNU tools, you may use the following command instead of making `stage1`, `stage2`, and performing the two compiler builds.

```
make bootstrap    LANGUAGES= language-list \
                  BOOT_CFLAGS= option-list
```

- 14.** Then compare the latest object files with the stage 2 object files—they ought to be identical, aside from time stamps (if any). On some systems, meaningful comparison of object files is impossible; they always appear *different*. This is currently true on Solaris and some systems that use ELF object file format. On some versions of Irix on SGI machines and DEC UNIX (OSF/1) on Alpha systems, you will not be able to compare the files without specifying `-save-temps`; see the description of individual systems in the previous discussions to see if you get comparison failures. You may have similar problems on other systems. Use the following command to compare the files.

```
make compare
```

This will mention any object files that differ between stage 2 and stage 3. Any difference, no matter how innocuous, indicates that the stage 2 compiler has compiled GCC incorrectly, and is therefore a potentially serious bug which you should investigate and report (see “General licenses and terms for using GNUPro Toolkit” on page 15 in *Installation* within ***Getting Started with GNUPro Toolkit***). If your system does not put time stamps in the object files, then use the following as a faster way to compare them (using the Bourne shell).

```
for file in *.o; do
  cmp $file stage2/$file
done
```

NOTE: If you built the compiler on MIPS machines with the option,

-mno-mips-tfile, you can't compare files.

15. Install the compiler driver, the compiler's passes and run-time support with 'make install'. Use the same value for CC, CFLAGS and LANGUAGES that you used when compiling the files that are being installed. One reason this is necessary is that some versions of Make have bugs and recompile files gratuitously when you use this step. If you use the same variable values, those files will be properly recompiled.

For example, if you have built the stage 2 compiler, you can use the following command.

```
make install CC="stage2/xgcc -Bstage2/" CFLAGS="-g -O" \
LANGUAGES=" list"
```

This copies the files 'cc1', 'cpp' and 'libgcc.a' to files 'cc1', 'cpp' and 'libgcc.a' in the directory /usr/local/lib/gcc-lib/target/version, which is where the compiler driver program looks for them. Here, target is the target machine type specified when you ran configure, and version is the version number of GCC.

This naming scheme permits various versions and/or cross-compilers to coexist. It also copies the executables for compilers for other languages (e.g., 'cc1plus' for C++) to the same directory.

This also copies the driver program 'xgcc' into /usr/local/bin/gcc, so that it appears in typical execution search paths. It also copies 'gcc.1' into /usr/local/man/man1 and info pages into /usr/local/info.

On some systems, this command causes recompilation of some files. This is usually due to bugs in make. You should either ignore this problem, or use GNU Make.

WARNING: There is a bug in alloca in the Sun library. To avoid this bug, be sure to *install the executables of GCC that were compiled by GCC*. (The executables from stage 2 or 3, *not* stage 1.) They use alloca as a built-in function, *never* the one in the library.

(It is usually better to install GCC executables from stage 2 or 3, as they usually run faster than ones compiled with some other compiler.)

16. If you're going to use C++, it's likely that you need to also install the libg++ distribution. It should be available from the same place where you got the GNU C distribution.

Just as GNU C does not distribute a C runtime library, it also does not include a C++ run-time library. All I/O functionality, special class libraries, etc., are available in the libg++ distribution.

17. GCC includes a runtime library for Objective-C because it is an integral part of the language. You can find the files associated with the library in the subdirectory

`'objc'`. The GNU Objective-C Runtime Library requires header files for the target's C library in order to be compiled, and also requires the header files for the target's thread library if you want thread support. See "Cross-compilers and header files" on page 51 for discussion about header files issues for cross-compilation.

When you run `'configure'`, it picks the appropriate Objective-C thread implementation file for the target platform. In some cases, you may wish to choose a different back-end as some platforms support multiple thread implementations or you may wish to disable thread support completely. To do this, specify a value for the `OBJC_THREAD_FILE` makefile variable on the command line when you run `make` with something like the following input.

```
make CC="stage2/xgcc -Bstage2/" CFLAGS="-g -O2"
                                OBJC_THREAD_FILE=thr-single
```

The following list shows the currently available back-ends.

- `thr-single`
Disable thread support, should work for all platforms.
- `thr-decosf1`
DEC OSF/1 thread support.
- `thr-irix`
SGI IRIX thread support.
- `thr-mach`
Generic MACH thread support, known to work on NEXTSTEP.
- `thr-os2`
IBM OS/2 thread support.
- `thr-posix`
Generix POSIX thread support.
- `thr-pthreads`
PThreads on Linux-based GNU systems.
- `thr-solaris`
SUN Solaris thread support.
- `thr-win32`
Microsoft Win32 API thread support.

Configurations supported by GCC

The following are the possible CPU types:

```
1750a, a29k, alpha, arm, cn, clipper, dsp16xx, elxsi, h8300, hppa1.0,
hppa1.1, i370, i386, i486, i586, i860, i960, m68000, m68k, m88k, mips, mipsel,
mips64, mips64el, ns32k, powerpc, powerpcle, pyramid, romp, rs6000, sh,
sparc, sparclite, sparc64, vax, we32k.
```

The following are the recognized company names. As you can see, customary abbreviations are used rather than the longer official names.

```
acorn, alliant, altos, apollo, apple, att, bull, cbm, convergent, convex,
crds, dec, dg, dolphin, elxsi, encore, harris, hitachi, hp, ibm, intergraph,
isi, mips, motorola, ncr, next, ns, omron, plexus, sequent, sgi, sony, sun,
tti, unicom, wrs.
```

The company name is meant only to clarify when the rest of the information supplied is insufficient. You can omit it, substituting ‘*cpu-system*’ (where *cpu* stands for your processor and *system* for your operating system), if the company name is not needed. For example, *vax-ultrix4.2* is equivalent to *vax-dec-ultrix4.2*.

The following is a list of system types:

```
386bsd, aix, acis, amigados, aos, aout, aux, bosx, bsd, clix, coff, ctix, cxux,
dgux, dynix, ebmon, ecoff, elf, esix, freebsd, hms, genix, gnu, gnu/linux,
hiux, hpux, iris, irix, isc, luna, lynxos, mach, minix, msdos, mvs, netbsd,
newsos, nindy, ns, osf, osfrose, ptx, riscix, riscos, rtu, sco, sim, solaris,
sunos, sym, sysv, udi, ultrix, unicos, uniplus, unos, vms, vsta, vxworks,
winnt, xenix.
```

You can omit the system type; then `configure` guesses the operating system from the CPU and company.

You can add a version number to the system type; this may or may not make a difference. For example, you can write ‘*bsd4.3*’ or ‘*bsd4.4*’ to distinguish versions of BSD. In practice, the version number is most needed for *sysv3* and *sysv4*, which are often treated differently.

If you specify an impossible combination such as *i860-dg-vms*, then you may get an error message from `configure`, or it may ignore part of the information and do the best it can with the rest. `configure` always prints the canonical name for the alternative that it used. GCC does not support all possible alternatives.

Often a particular model of machine has a name. Many machine names are recognized as aliases for CPU/company combinations. Thus, the machine name *sun3*, mentioned previously, is an alias for *m68k-sun*.

Sometimes we accept a company name as a machine name, when the name is popularly used for a particular machine. The following are the known machine

names:

```
3300, 3b1, 3bn, 7300, altos3068, altos, apollo68, att-7300, balance,
convex-cn, crds, decstation-3100, decstation, delta, encore, fx2800,
gmicro, hp7nn, hp8nn, hp9k2nn, hp9k3nn, hp9k7nn, hp9k8nn, iris4d, iris,
isi68, m3230, magnum, merlin, miniframe, mmax, news-3600, news800, news,
next, pbd, pc532, pmax, powerpc, powerpcle, ps2, risc-news, rtpc, sun2,
sun386i, sun386, sun3, sun4, symmetry, tower-32, tower.
```

Remember that a machine name specifies both the CPU type *and* the company name.

If you want to install your own homemade configuration files, you can use 'local' as the company name to access them.

If you use configuration `cpu-local`, the configuration name without the `cpu` prefix is used to form the configuration filenames. Thus, if you specify 'm68k-local', configuration uses files `m68k.md`, `local.h`, `m68k.c`, `xm-local.h`, `t-local`, and `x-local`, all in the directory, `config/m68k`.

What follows is a list of configurations that have special treatment or special things you must know.

1750a-**-*

MIL-STD-1750A processors.

The MIL-STD-1750A cross configuration produces output for `as1750`, an assembler/linker available under the GNU Public License for the 1750A.

Download from `ftp://ftp.fta-berlin.de/pub/crossgcc/1750gals/` to get `as1750`. A similarly licensed simulator for the 1750A is available from the same address.

You should ignore a fatal error during the building of `libgcc` (`libgcc` is not yet implemented for the 1750A.)

The `as1750` assembler requires the file, `ms1750.inc`, which is found in the directory `config/1750a`.

GCC produced the same sections as the Fairchild F9450 C Compiler; namely, they are the following sections.

- Normal
The program code section.
- Static
The read/write (RAM) data section.
- Konst
The read-only (ROM) constants section.
- Init
Initialization section (code to copy KREL to SREL).

The smallest addressable unit is 16 bits (BITS PER UNIT is 16). This means that type `char` is represented with a 16-bit word per character. The 1750A's

“Load/Store Upper/Lower Byte” instructions are not used by GCC.

`alpha-*-osf1`

Systems using processors that implement the DEC Alpha architecture and are running the DEC UNIX (OSF/1) operating system, for example the DEC Alpha AXP systems. (VMS on the Alpha is not currently supported by GCC.)

GCC writes a `.verstamp` directive to the assembler output file unless it is built as a cross-compiler. It gets the version to use from the system header file, `/usr/include/stamp.h`. If you install a new version of DEC UNIX, you should rebuild GCC to pick up the new version stamp.

NOTE: Since the Alpha is a 64-bit architecture, cross-compilers from 32-bit machines will not generate code as efficient as that generated when the compiler is running on a 64-bit machine. That is because many optimizations that depend on being able to represent a word on the target in an integral value on the host cannot be performed.

Building cross-compilers on the Alpha for 32-bit machines has only been tested in a few cases and may not work properly.

`make compare` may fail on old versions of DEC UNIX unless you add `savetemps` to `CFLAGS`. On these systems, the name of the assembler input file is stored in the object file, and that makes comparison fail if it differs between the `stage1` and `stage2` compilations. The option, `-save-temps`, forces a fixed name to be used for the assembler input file, instead of a randomly chosen name in `/tmp`.

Do not add `-save-temps` unless the comparisons fail without that option. If you add `-save-temps`, you will have to manually delete the `.i` and `.s` files after each series of compilations.

GCC now supports both the native (ECOFF) debugging format used by DBX and GDB and an encapsulated STABS format for use only with GDB. See Step 3 and its discussion of the `--with-stabs` option for `configure` on page 26 for more information on these formats and how to select them.

There is a bug in DEC’s assembler that produces incorrect line numbers for ECOFF format when the `.align` directive is used. To work around this problem, GCC will not emit such alignment directives while writing ECOFF format debugging information even if optimization is being performed. Unfortunately, this has the very undesirable side-effect that code addresses when `-o` is specified are different depending on whether or not `-g` is also specified.

To avoid this behavior, specify `-gstabs+` and use GDB instead of DBX. DEC is now aware of this problem with the assembler and hopes to provide a fix shortly. See “Options for debugging” on page 103.

`arm-*-aout`

Advanced RISC Machines ARM-family processors. These are often used in embedded applications. There are no standard UNIX configurations. This

configuration corresponds to the basic instruction sequences and will produce `a.out` format object modules.

You may need to make a variant of the file, `arm.h`, for your particular configuration.

`arm-*-linuxaout`

Any of the ARM family processors running the Linux-based GNU system with the ‘`a.out`’ binary format (ELF is not yet supported). You must use version 2.8.1.0.7 or later of the Linux binutils; download it from ‘`sunsite.unc.edu:/pub/Linux/GCC`’ and other mirror sites for Linux-based GNU systems.

`arm-*-riscix`

The ARM2 or ARM3 processor running RISC iX, Acorn’s port of BSD UNIX. If you are running a version of RISC iX prior to 1.2, then you must specify the version number during configuration.

NOTE: The assembler shipped with RISC iX does not support `stabs` debugging information; a new version of the assembler, with `stabs` support included, is now available from Acorn.

`a29k`

AMD Am29k-family processors. These are normally used in embedded applications. There are no standard UNIX configurations. This configuration corresponds to AMD’s standard calling sequence and binary interface and is compatible with other 29k tools.

You may need to make a variant of the file `a29k.h` for your particular configuration.

`a29k-*-bsd`

AMD Am29050 used in a system running a variant of BSD UNIX.

`decstation-*`

DECstations can support three different personalities: Ultrix, DEC OSF/1, and OSF/rose. To configure GCC for these platforms use the following configurations:

- `decstation-ultrix`
Ultrix configuration.
- `decstation-osf1`
Dec’s version of OSF/1.
- `decstation-osfrose`
Open Software Foundation reference port of OSF/1 which uses the OSF/rose object file format instead of ECOFF. Normally, you would not select this configuration.

The MIPS C compiler needs to be told to increase its table size for switch statements with the `-wf, -XNg1500` option in order to compile `cp/parse.c`. If you use the ‘`-O2`’ optimization option, you also need to use `-Olimit 3000`. Both of these options are automatically generated in the ‘`Makefile`’ that the shell script

‘configure’ builds. If you override the `CC` make variable and use the MIPS compilers, you may need to add `-Wf,-XNg1500 -Olimit 3000`.

`elxsi-elxsi-bsd`

The Elxsi’s C compiler has known limitations that prevent it from compiling GNU C. Please contact mrs@cygnus.com for more details.

`dsp16xx`

A port to the AT&T DSP1610 family of processors.

`h8300-*-*`

The calling convention and structure layout has changed in release 2.6. All code must be recompiled. The calling convention now passes the first three arguments in function calls in registers. Structures are no longer a multiple of 2 bytes.

`hppa*-*-*`

There are several variants of the HP-PA processor which run a variety of operating systems. GCC must be configured to use the correct processor type and operating system, or GCC will not function correctly. The easiest way to handle this problem is to avoid specifying a target when configuring GCC. The ‘configure’ script will try to automatically determine the right processor type and operating system.

`-g` does not work on HP/UX, since that system uses a peculiar debugging format about which GCC does not know. However, `-g` will work if you also use GAS and GDB in conjunction with GCC. We highly recommend using GAS for all HPPA configurations.

You should be using GAS-2.6 (or later) along with GDB-4.16 (or later). These can be retrieved from all the traditional GNU ftp archive sites. Install GAS into a directory before `/bin`, `/usr/bin`, and `/usr/ccs/bin` in your search path.

To enable debugging, configure GCC with the `--with-gnu-as` option before building.

`i370-*-*`

This port is very preliminary and has many known bugs. We hope to have a higher-quality port for this machine soon.

`i386-*-linuxoldld`

Use this configuration to generate `a.out` binaries on Linux if you do not have `gas/binutils` version 2.5.2 or later installed. This is an obsolete configuration.

`i386-*-linuxaout`

Use this configuration to generate ‘`a.out`’ binaries on Linux. This configuration is being superseded. You must use `gas/binutils` version 2.5.2 or later.

`i386-*-linux-gnu`

Use this configuration to generate ELF binaries on Linux. You must use `gas/binutils` version 2.5.2 or later.

`i386-*-sco`

Compilation with RCC is recommended. Also, it may be a good idea to link with GNU `malloc` instead of the `malloc` that comes with the system.

i386-*-sco3.2v4

Use this configuration for SCO release 3.2 version 4.0.

i386-*-sco3.2v5

Use this for SCO Open Server release 3.2 version 5.0. GCC can generate ELF binaries (if you specify `-melf`) or COFF binaries (the default). If you are going to build your compiler in ELF mode (once you have bootstrapped the first stage compiler) you must specify `-melf` as part of `CC`, not `CFLAGS`.

You should use some variant of the following example code statement.

```
CC="stage1/xgcc -melf" CFLAGS="-Bstage1/"
```

If you do not do this, the bootstrap will generate completely bogus versions of `libgcc.a` generated. You must have TLS597 (from <ftp.sco.com/TLS>) installed for ELF binaries to work correctly.

note: Open Server 5.0.2 does need TLS597 installed.

i386-*-isc

It may be a good idea to link with GNU `malloc` instead of the `malloc` that comes with the system. In ISC version 4.1, `sed` core dumps when building `'deduced.h'`. Use the version of `sed` from version 4.0.

i386-*-esix

It may be good idea to link with GNU `malloc` instead of the `malloc` that comes with the system.

i386-ibm-aix

You need to use GAS version 2.1 or later, and LD from GNU binutils version 2.2 or later.

i386-sequent-bsd

Go to the Berkeley universe before compiling. In addition, you probably need to create a file named `'string.h'` containing the following line:

```
#include <strings.h>
```

i386-sequent-ptx1*

Sequent DYNIX/ptx 1.x.

i386-sequent-ptx2*

Sequent DYNIX/ptx 2.x.

i386-sun-sunos4

You may find that you need another version of GCC to begin bootstrapping with, since the current version when built with the system's own compiler seems to get an infinite loop compiling part of `'libgcc2.c'`. GCC version 2 compiled with GCC (any version) seems not to have this problem. See "Installing GCC on the Sun" on page 55 for information on installing GCC on Sun systems.

i[345]86-*-winnt3.5

This version requires a GAS that has not yet been released. Until it is, you can get a pre-built binary version via anonymous ftp from `'cs.washington.edu:pub/gnat'` or `'cs.nyu.edu:pub/gnat'`. You must also use the Microsoft header files from the Windows NT 3.5 SDK. Find these on the CDROM in the `/mstools/h` directory dated September 4, 1994. You must use a

fixed version of Microsoft linker made especially for NT 3.5, which is also is available on the NT 3.5 SDK CDROM. If you do not have this linker, can you also use the linker from Visual C/C++ 1.0 or 2.0.

Installing GCC for NT builds a wrapper linker, called 'ld.exe', which mimics the behavior of UNIX 'ld' in the specification of libraries ('-L' and '-l'). 'ld.exe' looks for both UNIX and Microsoft named libraries. For example, if you specify '-lfoo', 'ld.exe' will look first for 'libfoo.a' and then for 'foo.lib'. You may install GCC for Windows NT in one of two ways, depending on whether or not you have a UNIX-like shell and various UNIX-like utilities.

- If you do not have a UNIX-like shell and few UNIX-like utilities, you will use a DOS style batch script called 'configure.bat'.

Invoke it as 'configure winnt' from an MSDOS console window or from the program manager dialog box. 'configure.bat' assumes you have already installed and have in your path a UNIX-like `sed` pro-gram which is used to create a working Makefile from 'Makefile.in'. Makefile uses the Microsoft Nmake program maintenance utility and the Visual C/C++ V8.00 compiler to build GCC.

You only need the utilities, `sed` and `touch`, to use this installation method, which only automatically builds the compiler itself. You must then examine what 'fixinc.winnt' does, edit the header files by hand and build 'libgcc.a' manually.

- The second type of installation assumes you are running a UNIX-like shell, have a complete suite of UNIX-like utilities in your path, and have a previous version of GCC already installed, either through building it via the previous installation method or acquiring a pre-built binary. In this case, use the `configure` script in the normal fashion.

i860-intel-osfl

This is the Paragon. If you have version 1.0 of the operating system, see "Installation problems" on page 294 for special things you need to do to compensate for peculiarities in the system.

*-lynx-lynxos

LynxOS 2.2 and earlier comes with GCC 1.x already installed as `/bin/gcc`. You should compile with `/bin/gcc` instead of `/bin/cc`. You can tell GCC to use the GNU assembler and linker, by specifying the following declaration when configuring.

```
--with-gnu-as --with-gnu-ld
```

These will produce COFF format object files and executables; otherwise GCC will use the installed tools, which produce 'a.out' format executables.

m32r-*-elf

Embedded M32R system.

m68000-hp-bsd

HP 9000 series 200 running BSD.

NOTE: The C compiler that comes with this system cannot compile GCC; contact 'law@cs.utah.edu' to get binaries of GCC for bootstrapping.

m68k-altos

Altos 3068. You must use the GNU assembler, linker and debugger. Also, you must fix a kernel bug. Details in the file, 'README.ALTOS'.

m68k-apple-aux

Apple Macintosh running A/UX. You may configure GCC to use either the system assembler and linker or the GNU assembler and linker.

You should use the GNU configuration if you can, especially if you also want to use G++. You enabled that configuration with the options, `--with-gnu-as` and `--with-gnu-ld`, to configure.

NOTE: The C compiler that comes with this system cannot compile GCC. You can find binaries of GCC for bootstrapping on `jagubox.gsfc.nasa.gov`. You will also a patched version of `/bin/ld` there that raises some of the arbitrary limits found in the original.

m68k-att-sysv

AT&T 3b1, a.k.a. 7300 PC. Special procedures are needed to compile GCC with this machine's standard C compiler, due to bugs in that compiler. You can bootstrap it more easily with previous versions of GCC if you have them.

Installing GCC on the 3b1 is difficult if you do not already have GCC running, due to bugs in the installed C compiler. However, the following procedure *might* work. (We were unable to test it.)

1. Comment out `'#include "config.h"'` line near the start of `'cccp.c'` and do `'make cpp'`. This makes a preliminary version of GNU `cpp`.
2. Save the old `'/lib/cpp'` and copy the preliminary GNU `cpp` to that filename.
3. Undo your change in `'cccp.c'`, or reinstall the original version, and do `'make cpp'` again.
4. Copy this final version of GNU `cpp` into `'/lib/cpp'`.
5. Replace every occurrence of `obstack_free` in the file, `'tree.c'`, with `_obstack_free`.
6. Run `make` to get the first-stage GCC.
7. Reinstall the original version of `/lib/cpp`.
8. Now you can compile GCC with itself and install it in the normal fashion.

m68k-bull-sysv

Bull DPX/2 series 200 and 300 with BOS2.00.45 up to BOS-2.01. GCC works either with native assembler or GNU assembler.

You can use GNU assembler with native COFF generation by providing the ‘`--with-gnu-as`’ option to the configure script or use GNU assembler with the ‘`dbx-in-coff`’ encapsulation by providing the ‘`--with-gnu-as-stabs`’ option. For any problem with native assembler or for availability of the DPX/2 port of GAS, contact: F.Pierresteguy@frcl.bull.fr.

m68k-crds-unox

Use ‘`configure unos`’ for building on Unos. The Unos assembler is named `casm` instead of `as`.

For some strange reason, linking `/bin/as` to `/bin/casm` changes the behavior, and does not work.

So, when installing GCC, you should install the following script as `as` in the subdirectory where the passes of GCC are installed:

```
#!/bin/sh
casm $*
```

The default Unos library is named ‘`libunos.a`’ instead of ‘`libc.a`’. To allow GCC to function, either change all references to ‘`-lc`’ in ‘`gcc.c`’ to ‘`-lunos`’ or link ‘`/lib/libc.a`’ to ‘`/lib/libunos.a`’.

When compiling GCC with the standard compiler, to overcome bugs in the support of `alloca`, do not use ‘`-o`’ when making stage 2. Then use the stage 2 compiler with ‘`-o`’ to make the stage 3 compiler. This compiler will have the same characteristics as the usual stage 2 compiler on other systems. Use it to make a stage 4 compiler and compare that with stage 3 to verify proper compilation.

Unos uses memory segmentation instead of demand paging, so you will need a lot of memory. 5 Mb is barely enough if no other tasks are running. If linking `cc1` fails, try putting the object files into a library and linking from that library.

m68k-hp-hpux

HP 9000 series 300 or 400 running HP/UX. HP/UX version 8.0 has a bug in the assembler that prevents compilation of GCC. To fix it, get patch PHCO 4484 from HP.

In addition, if you wish to use the `gas` function ‘`--with-gnu-as`’, you must use `gas`, version 2.1 or later, and you must use the GNU linker version 2.1 or later. Earlier versions of `gas` relied upon a program which converted the `gas` output into the native HP/UX format, but that program has not been kept up to date. `gdb` does not understand that native HP/UX format, so you must use `gas` if you wish to use `gdb`.

m68k-sun

Sun 3. We do not provide a configuration file to use the Sun FPA by default, because programs that establish signal handlers for floating point traps inherently

cannot work with the FPA. See “Installing GCC on the Sun” on page 55 for information on installing GCC on Sun systems.

m88k-*-svr3

Motorola m88k running the AT&T/Unisoft/Motorola V.3 reference port. These systems tend to use the Green Hills C, revision 1.8.5, as the standard C compiler. There are apparently bugs in this compiler that result in object files differences between stage 2 and stage 3. If this happens, make the stage 4 compiler and compare it to the stage 3 compiler. If the stage 3 and stage 4 object files are identical, this suggests you encountered a problem with the standard C compiler; the stage 3 and 4 compilers may be usable. It is best, however, to use an older version of GCC for bootstrapping if you have one.

m88k-*-dgux

Motorola m88k running DG/UX. To build 88open BCS native or cross compilers on DG/UX, specify the configuration name as ‘m88k-*-dguxbcs’ and build in the 88open BCS software development environment. To build ELF native or cross compilers on DG/UX, specify ‘m88k-*-dgux’ and build in the DG/UX ELF development environment. You set the software development environment by issuing ‘sde-target’ command and specifying either ‘m88kbcs’ or ‘m88kdguxelf’ as the operand. If you do not specify a configuration name, configure guesses the configuration based on the current software development environment.

m88k-tektronix-sysv3

Tektronix XD88 running UTekV 3.2e. Do not turn on optimization while building stage1 if you bootstrap with the buggy Green Hills compiler. Also, The bundled LAI System V NFS is buggy so if you build in an NFS mounted directory, start from a fresh reboot, or avoid NFS all together. Otherwise you may have trouble getting clean comparisons between stages.

mips-mips-bsd

MIPS machines running the MIPS operating system in BSD mode. It’s possible that some old versions of the system lack the functions memcpy, memcmp, and memset. If your system lacks these, you must remove or undo the definition of TARGET_MEM_FUNCTIONS in ‘mips-bsd.h’.

The MIPS C compiler needs to be told to increase its table size for switch statements with the ‘-Wf,-XNg1500’ option in order to compile ‘cp/parse.c’. If you use the ‘-O2’ optimization option, you also need to use ‘-Olimit 3000’. Both of these options are automatically generated in the Makefile that the shell script configure builds. If you override the CC make variable and use the MIPS compilers, you may need to add ‘-Wf,-XNg1500 -Olimit 3000’.

mips-mips-riscos*

The MIPS C compiler needs to be told to increase its table size for switch statements with the ‘-Wf,-XNg1500’ option in order to compile ‘cp/parse.c’.

If you use the ‘-O2’ optimization option, you also need to use ‘-Olimit 3000’. Both of these options are automatically generated in the ‘Makefile’ that the shell

script configure builds.

If you override the `CC` make variable and use the MIPS compilers, you may need to add `-Wf,-XNg1500 -Olimit 3000`. MIPS computers running RISC-OS can support four different personalities: default, BSD 4.3, System V.3, and System V.4 (older versions of RISC-OS don't support V.4). To configure GCC for these platforms use the following configurations:

- `mips-mips-riscosrev`
Default configuration for RISC-OS, revision, *rev*.
- `mips-mips-riscosrevbsd`
BSD 4.3 configuration for RISC-OS, revision, *rev*.
- `mips-mips-riscosrevsysv4`
System V.4 configuration for RISC-OS, revision, *rev*.
- `mips-mips-riscosrevsysv`
System V.3 configuration for RISC-OS, revision, *rev*.

The revision, *rev* (mentioned in the previous paragraphs for the option, `mips-mips-riscos*`), is the version of RISC-OS to use. You must reconfigure GCC when going from a RISC-OS revision 4 to RISC-OS revision 5. This has the effect of avoiding a linker bug (see “Installation problems” on page 294).

`mips-sgi-*`

In order to compile GCC on an SGI running IRIX 4, the `c.hdr.lib` option must be installed from the CD-ROM supplied from Silicon Graphics. This is found on the second CD in release 4.0.1.

In order to compile GCC on an SGI running IRIX 5, the `compiler dev.hdr` subsystem must be installed by the IDO CD-ROM, supplied by Silicon Graphics.

`make compare` may fail on version 5 IRIX unless you add `-save-temps` to `CFLAGS`. On these systems, the name of the assembler input file is stored in the object file, and that makes comparison fail if it differs between the `stage1` and `stage2` compilations. The `-save-temps` option forces a fixed name to be used for the assembler input file, instead of a randomly chosen name in `/tmp`.

Do not add `-save-temps` unless the comparisons fail without that option. If you do you `-save-temps`, you will have to manually delete the `.i` and `.s` files after each series of compilations. The MIPS C compiler needs to be told to increase its table size for switch statements with the `-Wf,-XNg1500` option in order to compile `cp/parse.c`. If you use the `-O2` optimization option, you also need to use `-Olimit 3000`. Both of these options are automatically generated in the `'Makefile'` that the shell script `configure` builds.

If you override the `CC` make variable and use the MIPS compilers, you may need to add `-Wf,-XNg1500 -Olimit 3000`. On Irix version 4.0.5F, and perhaps on some other versions as well, there is an assembler bug that reorders instructions incorrectly. To work around it, specify `mips-sgi-irix4loser` as the target configuration. This configuration inhibits assembler optimization. In a compiler

configured with target, `mips-sgi-irix4`, you can turn off assembler optimization by using the `'-noasmopt'` option.

This compiler option passes the option, `-O0`, to the assembler, to inhibit reordering. The `-noasmopt` option can be useful for testing whether a problem is due to erroneous assembler reordering.

Even if a problem does not go away with `-noasmopt`, it may still be due to assembler reordering—perhaps GCC itself was miscompiled as a result. To enable debugging under Irix 5, you must use GNU `as` 2.5 or later, and use the `--with-gnu-as` configure option when configuring `gcc`. GNU `as` is distributed as part of the `binutils` package.

`mips-sony-sysv`

Sony MIPS NEWS. This works in NEWSOS 5.0.1, but not in 5.0.2 (which uses ELF instead of COFF). Support for 5.0.2 will probably be provided soon by volunteers. In particular, the linker does not like the code generated by GCC when shared libraries are linked in.

`ns32k-encore`

Encore NS32000 system. Encore systems are supported only under BSD.

`ns32k-*-genix`

National Semiconductor NS32000 system. Genix has bugs in `alloca` and `malloc`; you must get the compiled versions of these from GNU Emacs.

`ns32k-sequent`

Go to the Berkeley universe before compiling. In addition, you probably need to create a file named `'string.h'` containing just one line:

```
#include <strings.h>
```

`ns32k-utek`

UTEK NS32000 system (“merlin”). The C compiler that comes with this system cannot compile GCC; contact `'tektronix!reed!mason'` to get binaries of GCC for bootstrapping.

`romp-*-aos`

`romp-*-mach`

The only operating systems supported for the IBM RT PC are AOS and MACH. GCC does not support AIX running on the RT.

We recommend you compile GCC with an earlier version of itself; if you compile GCC with `hc`, the Metaware compiler, it will work, but you will get mismatches between the stage 2 and stage 3 compilers in various files. These errors are minor differences in some floating-point constants and can be safely ignored; the stage 3 compiler is correct.

`rs6000-*-aix`

`powerpc-*-aix`

Various early versions of each release of the IBM XLC compiler will not bootstrap GCC. Symptoms include differences between the stage2 and stage3 object files, and errors when compiling `'libgcc.a'` or `'enquire'`. Known problematic releases include: `xlc-1.2.1.8`, `xlc-1.3.0.0` (distributed with AIX 3.2.5),

and xlc-1.3.0.19. Both xlc-1.2.1.28 and xlc-1.3.0.24 (PTF 432238) are known to produce working versions of GCC, but most other recent releases correctly bootstrap GCC. Also, releases of AIX prior to AIX 3.2.4 include a version of the IBM assembler which does not accept debugging directives: assembler updates are available as PTFs. Also, if you are using AIX 3.2.5 or greater and the GNU assembler, you must have a version modified after October 16, 1995 in order for the GNU C compiler to build.

See the file ‘README.RS6000’ for more details on of these problems.

GCC does not yet support the 64-bit PowerPC instructions.

Objective C does not work on this architecture because it makes assumptions that are incompatible with the calling conventions. AIX on the RS/6000 provides support (NLS) for environments outside of the United States. Compilers and assemblers use NLS to support locale-specific representations of various objects including floating-point numbers (“.” vs “,” for separating decimal fractions). There have been problems reported where the library linked with GCC does not produce the same floating-point formats that the assembler accepts. If you have this problem, set the `LANG` environment variable to “C” or “En_US”. Due to changes in the way that GCC invokes the binder (`linker`) for AIX 4.1, you may now receive warnings of duplicate symbols from the link step that were not reported before. The assembly files generated by GCC for AIX have always included multiple symbol definitions for certain global variable and function declarations in the original program. The warnings should not prevent the linker from producing a correct library or runnable executable. By default, AIX 4.1 produces code that can be used on either Power or PowerPC processors. You can specify a default version for the `-mcpu=cpu_type` switch by using the configure option, `--with-cpu-cpu_type`.

```
powerpc-*-elf
powerpc-*-sysv4
```

PowerPC system in big endian mode, running System V.4. You can specify a default version for the `-mcpu=cpu_type` switch using the option,

```
--with-cpu-cpu_type.
```

```
powerpc-*-linux-gnu
```

PowerPC system in big endian mode, running Linux. You can specify a default version for the `-mcpu=cpu_type` switch using the option,

```
--with-cpu-cpu_type.
```

```
powerpc-*-eabiaix
```

Embedded PowerPC system in big endian mode with ‘`-mcall-aix`’ selected as the default. You can specify a default version for the `-mcpu=cpu_type` switch by using the configure option, `--with-cpu-cpu_type`.

```
powerpc-*-eabisim
```

Embedded PowerPC system in big endian mode for use in running under the

PSIM simulator. You can specify a default version for the `-mcpu=cpu_type` switch by using the configure option, `--with-cpu-cpu_type`.

`powerpc-*eabi`

Embedded PowerPC system in big endian mode. You can specify a default version for the `-mcpu=cpu_type` switch by using the configure option, `--with-cpu-cpu_type`.

`powerpcle-*-elf`

`powerpcle-*-sysv4`

PowerPC system in little endian mode, running System V.4. You can specify a default version for the `-mcpu=cpu_type` switch using the option, `--with-cpu-cpu_type`.

`powerpcle-*-solaris2*`

PowerPC system in little endian mode for use in running Solaris 2.5.1 or higher.

You can specify a default version for the `-mcpu=cpu_type` switch by using the configure option, `--with-cpu-cpu_type`. Beta versions of the Sun 4.0 compiler do not seem to be able to build GCC correctly. There are also problems with the host assembler and linker that are fixed by using the GNU versions of these tools.

`powerpcle-*-eabisim`

Embedded PowerPC system in little endian mode for use in running under the PSIM simulator.

`powerpcle-*-eabi`

Embedded PowerPC system in little endian mode. You can specify a default version for the `-mcpu=cpu_type` switch by using the option, `--with-cpu-cpu_type`.

`powerpcle-*-winnt`

`powerpcle-*-pe`

PowerPC system in little endian mode running Windows NT. You can specify a default version for the `-mcpu=cpu_type` switch by using the configure option, `--with-cpu-cpu_type`.

`vax-dec-ultrix`

Don't try compiling with Vax C (`vcc`). It produces incorrect code in some cases (for example, when `alloca` is used).

Meanwhile, compiling `'cp/parse.c'` with `pcc` does not work because of an internal table size limitation in that compiler. To avoid this problem, compile just the GNU C compiler first, and use it to recompile building all the languages that you want to run.

`sparc-sun-*`

`sparc64-sun-*`

See "Installing GCC on the Sun" on page 55 for information on installing GCC on Sun systems.

`vax-dec-vms`

See Installing GCC on VMS for details on how to install GCC on VMS.

we32k--*

These computers are also known as the 3b2, 3b5, 3b20 and other similar names. (However, the 3b1 is actually a 68000; see “Configurations Supported by GCC”) Don’t use ‘-g’ when compiling with the system’s compiler. The system’s linker seems to be unable to handle such a large program with debugging information. The system’s compiler runs out of capacity when compiling ‘stmt.c’ in GCC. You can work around this by building ‘cpp’ in GCC first, then use that instead of the system’s preprocessor with the system’s C compiler to compile ‘stmt.c’. Use the following.

```
mv /lib/cpp /lib/cpp.att
cp cpp /lib/cpp.gnu
echo '/lib/cpp.gnu -traditional ${1+"$@"}' > /lib/cpp
chmod +x /lib/cpp
```

The system’s compiler produces bad code for some of the GCC optimization files. So you must build the stage 2 compiler without optimization. Then build a stage 3 compiler with optimization. That executable should work. Use the following.

```
make LANGUAGES=c CC=stage1/xgcc CFLAGS="-Bstage1/ -g"
make stage2
make CC=stage2/xgcc CFLAGS="-Bstage2/ -g -O"
```

You may need to raise the ULIMIT setting to build a C++ compiler, as the file ‘cclplus’ is larger than one megabyte.

Compilation in a separate directory

If you wish to build the object files and executables in a directory other than the one containing the source files, use the following.

1. Make sure you have a version of `make` that supports the `VPATH` feature. (GNU `make` supports it, as do `make` versions on most BSD systems.)
2. If you have ever run ‘`configure`’ in the source directory, you must undo the configuration. Do this by running `make distclean`.
3. Go to the directory in which you want to build the compiler before running `configure`.

```
mkdir gcc-sun3
cd gcc-sun3
```

On systems that do not support symbolic links, this directory must be on the same file system as the source code directory.

4. Specify where to find `configure` when you run it.
- ```
../gcc/configure...
```

This also tells `configure` where to find the compiler sources; `configure` takes the directory from the filename that was used to invoke it.

And, if you want to be sure, you can specify the source directory with the

`--srcdir` option like the following example demonstrates.

```
../gcc/configure--srcdir= ../gcc other options
```

The directory you specify with `--srcdir` need not be the same as the one in which `configure` is found.

Now, you can run `make` in that directory. You need not repeat the configuration steps shown previously when ordinary source files change. You must, however, run `configure` again when the configuration files change, if your system does not support symbolic links.

## Building and installing a cross-compiler

GCC can function as a cross-compiler for many machines, but not all.

- Cross-compilers for the MIPS as target using the MIPS assembler currently do not work, because the auxiliary programs `'mips-tdump.c'` and `'mips-tfile.c'` can't be compiled on anything but a MIPS. It does work to cross compile for a MIPS if you use the GNU assembler and linker.
- Cross-compilers between machines with different floating point formats have not all been made to work. GCC now has a floating point emulator with which these can work, but each target machine description needs to be updated to take advantage of it.
- Cross-compilation between machines of different word sizes is some what problematic and sometimes does not work.

Since GCC generates assembler code, you probably need a cross-assembler that GCC can run, in order to produce object files. If you want to link on other than the target machine, you need a cross-linker as well. You also need header files and libraries suitable for the target machine that you can install on the host machine.

## Steps of cross-compilation

To compile and run a program using a cross-compiler involves several steps:

1. Run the cross-compiler on the host machine to produce assembler files for the target machine. This requires header files for the target machine.
2. Assemble the files produced by the cross-compiler. You can do this either with an assembler on the target machine, or with a cross-assembler on the host machine.
3. Link those files to make an executable. You can do this either with a linker on the target machine, or with a cross-linker on the host machine. Whichever machine you use, you need libraries and certain startup files (typically, `crt.o`) for the target machine.

It is most convenient to do all of these steps on the same host machine, since then you can do it all with a single invocation of GCC. This requires a suitable cross-assembler

and cross-linker. For some targets, the GNU assembler and linker are available.

## Configuring a cross-compiler

To build GCC as a cross-compiler, you start out by running `configure`. Use the `--target=target` option to specify the target type. If `configure` was unable to correctly identify the system you are running on, specify the option, `--build=build`. For instance, the following example shows how to configure for a cross-compiler that produces code for an HP 68030 system running BSD on a system that `configure` can correctly identify.

```
./configure--target=m68k-hp-bsd4.3
```

## Tools and libraries for a cross-compiler

If you have a cross-assembler and cross-linker available, you should install them now. Put them in the directory, `/usr/local/target/bin`. What follows are the tools you should put in this directory.

- `as`  
This should be the cross-assembler.
- `ld`  
This should be the cross-linker.
- `ar`  
This should be the cross-archiver: a program which can manipulate archive files (linker libraries) in the target machine's format.
- `ranlib`  
This should be a program to construct a symbol table in an archive file.

The installation of GCC will find these programs in that directory, and copy or link them to the proper place to for the cross-compiler to find them when run later.

The easiest way to provide these files is to build the `binutils` package and `gas`.

Configure them with the same `--host` and `--target` options that you use for configuring GCC, then build and install them. They install their executables automatically into the proper directory. Alas, they do not support all the targets that GCC supports.

If you want to install libraries to use with the cross-compiler, such as a standard C library, put them in the directory `/usr/local/target/lib`; installation of GCC copies all the files in that subdirectory into the proper place for GCC to find them and link with them.

What follows is an example of copying some libraries from a target machine.

```
ftp target-machine
lcd /usr/local/target/lib
cd /lib
```

```
get libc.a
cd /usr/lib
get libg.a
get libm.a
quit
```

The precise set of libraries you'll need, and their locations on the target machine, vary depending on its operating system. Many targets require “*start files*” such as ‘crt0.o’ and ‘crtn.o’ which are linked into each executable. These too should be placed in ‘/usr/local/target/lib’.

There may be several alternatives for ‘crt0.o’, for use with profiling or other compilation options.

Check your target's definition of `STARTFILE_SPEC` to find out what start files it uses. The following is an example of copying these files from a target machine.

```
ftp target-machine
lcd /usr/local/target/lib
prompt
cd /lib
mget *crt*.o
cd /usr/lib
mget *crt*.o
quit
```

## libgcc.a and cross-compilers

Code compiled by GCC uses certain runtime support functions implicitly. Some of these functions can be compiled successfully with GCC itself, but a few cannot be. These problem functions are in the source file, ‘libgcc1.c’; the library made from them is called ‘libgcc1.a’.

When you build a native compiler, these functions are compiled with some other compiler—the one that you use for bootstrapping GCC. Presumably it knows how to open code these operations, or else knows how to call the run-time emulation facilities that the machine comes with. But this approach doesn't work for building a cross-compiler. The compiler that you use for building knows about the host system, not the target system.

So, when you build a cross-compiler you have to supply a suitable library ‘libgcc1.a’ that does the job it is expected to do.

To compile ‘libgcc1.c’ with the cross-compiler itself does not work. The functions in this file are supposed to implement arithmetic operations that GCC does not know how to open code for your target machine. If these functions are compiled with GCC itself, they will compile into infinite recursion.

On any given target, most of these functions are not needed. If GCC can open code an arithmetic operation, it will not call these functions to perform the operation. It is

possible that on your target machine, none of these functions is needed. If so, you can supply an empty library as 'libgcc1.a'.

Many targets need library support only for multiplication and division. If you are linking with a library that contains functions for multiplication and division, you can tell GCC to call them directly by defining the macros `MULSI3_LIBCALL`, and the like. These macros need to be defined in the target description macro file. For some targets, they are defined already. This may be sufficient to avoid the need for 'libgcc1.a'; if so, you can supply an empty library.

Some targets do not have floating point instructions; they need other functions in 'libgcc1.a', which do floating arithmetic. Recent versions of GCC have a file which emulates floating point. With a certain amount of work, you should be able to construct a floating point emulator that can be used as 'libgcc1.a'. Perhaps future versions will contain code to do this automatically and conveniently. That depends on whether someone wants to implement it.

Some embedded targets come with all the necessary 'libgcc1.a' routines written in C or assembler. These targets build 'libgcc1.a' automatically and you do not need to do anything special for them. Other embedded targets do not need any 'libgcc1.a' routines since all the necessary operations are supported by the hardware.

If your target system has another C compiler, you can configure GCC as a native compiler on that machine.

Build just 'libgcc1.a' with 'make libgcc1.a' on that machine, and use the resulting file with the cross-compiler. To do this, execute the following on the target machine.

```
cd target-build-dir
./configure-host=sparc-target=sun3
make libgcc1.a
```

And then, execute the following on the host machine:

```
ftp target-machine
binary
cd target-build-dir
get libgcc1.a
quit
```

Another way to provide the functions you need in 'libgcc1.a' is to define the appropriate `perform_...` macros for those functions. If these definitions do not use the C arithmetic operators that they are meant to implement, you should be able to compile them with the cross-compiler you are building. (If these definitions already exist for your target file, then you are all set.) To build 'libgcc1.a' using the `perform` macros, use `LIBGCC1=libgcc1.a OLDCC=./xgcc` when building the compiler.

Otherwise, you should place your replacement library under the name 'libgcc1.a' in the directory in which you will build the cross-compiler, before you run `make`.

## Cross-compilers and header files

If you are cross-compiling a stand-alone program or a program for an embedded system, then you may not need any header files except the few that are part of GCC (and those of your program). However, if you intend to link your program with a standard C library such as `libc.a`, then you probably need to compile with the header files that go with the library you use.

The GNU C compiler does not come with these files, because (1) they are system-specific, and (2) they belong in a C library, not in a compiler.

If the GNU C library supports your target machine, then you can get the header files from there (assuming you actually use the GNU library when you link your program).

If your target machine comes with a C compiler, it probably comes with suitable header files also. If you make these files accessible from the host machine, the cross-compiler can use them also. Otherwise, you're on your own in finding header files to use when cross-compiling.

When you have found suitable header files, put them in `/usr/local/target/include`, before building the cross compiler. Then installation will run `fixincludes` properly and install the corrected versions of the header files where the compiler will use them. Provide the header files before you build the cross-compiler, because the build stage actually runs the cross-compiler to produce parts of `libgcc.a`. (These are the parts that can be compiled with GCC.) Some of them need suitable header files.

To copy header files from a target machine, use the following example's input.

```
(cd /usr/include; tar cf - .) > tarfile
```

Then, on the host machine, use the following example's input where *target-machine* represents your intended target machine.

```
ftp target-machine
lcd /usr/local/target/include
get tarfile
quit
tar xf tarfile
```

## Standard header file directories

`GCC_INCLUDE_DIR` means the same thing for native and cross. It is where GCC stores its private include files, and also where GCC stores the fixed `include` files. A cross compiled GCC runs `fixincludes` on the header files in `$(tooldir)/include`. (If the cross compilation header files need to be fixed, they must be installed before GCC is built. If the cross compilation header files are already suitable for ANSI C and GCC, nothing special need be done).

`GPLUS_INCLUDE_DIR` means the same thing for native and cross. It is where `g++` looks first for header files. `libg++` installs only target independent header files in that directory.

LOCAL\_INCLUDE\_DIR is used only for a native compiler. It is normally `'/usr/local/include'`. GCC searches this directory so that users can install header files in `'/usr/local/include'`.

CROSS\_INCLUDE\_DIR is used only for a cross compiler. GCC doesn't install anything there.

TOOL\_INCLUDE\_DIR is used for both native and cross compilers. It is the place for other packages to install header files that GCC will use. For a cross-compiler, this is the equivalent of `'/usr/include'`. When you build a cross-compiler, `fixincludes` processes any header files in this directory.

```
(cd /usr/include; tar cf - .) > tarfile
```

Then, on the host machine, use the following example's input where *target-machine* represents your intended target machine.

```
ftp target-machine
lcd /usr/local/target/include
get tarfile
quit
tar xf tarfile
```

## Actually building the cross-compiler

Now you can proceed just as for compiling a single-machine compiler through the step of building stage 1. If you have not provided some sort of `'libgcc1.a'`, then compilation will give up at the point where it needs that file, printing a suitable error message. If you do provide `'libgcc1.a'`, then building the compiler will automatically compile and link a test program called `'libgcc1-test'`; if you get errors in the linking, it means that not all of the necessary routines in `'libgcc1.a'` are available.

You must provide the header file `float.h`. One way to do this is to compile `enquire` and run it on your target machine. The job of `enquire` is to run on the target machine and figure out by experiment the nature of its floating point representation. `enquire` records its findings in the header file `'float.h'`. If you can't produce this file by running `enquire` on the target machine, then you will need to come up with a suitable `'float.h'` in some other way (or else, avoid using it in your programs).

Do not try to build stage 2 for a cross-compiler. It doesn't work to rebuild GCC as a cross-compiler using the cross-compiler, because that would produce a program that runs on the target machine, not on the host. For example, if you compile a 386-to-68030 cross-compiler with itself, the result will not be right either for the 386 (because it was compiled into 68030 code) or for the 68030 (because it was configured for a 386 as the host). If you want to compile GCC into 68030 code, whether you compile it on a 68030 or with a cross-compiler on a 386, you must specify a 68030 as the host when you configure it.

To install the cross-compiler, use ‘make install’, as usual.

## collect2 and cross-compiling

Many target systems do not have support in the assembler and linker for “constructors”—initialization functions to be called before the official “start” of `main`. On such systems, GCC uses a utility called `collect2` to arrange to call these functions at start time.

The program `collect2` works by linking the program once and looking through the linker output file for symbols with particular names indicating they are constructor functions. If it finds any, it creates a new temporary ‘.c’ file containing a table of them, compiles it, and links the program a second time including that file.

The actual calls to the constructors are carried out by a subroutine called `__main`, which is called (automatically) at the beginning of the body of `main` (provided `main` was compiled with GCC).

Calling `__main` is necessary, even when compiling C code, to allow linking C and C++ object code together. (If you use ‘-nostdlib’, you get an unresolved reference to `__main`, since it’s defined in the standard GCC library. Include `-lgcc` at the end of your compiler command line to resolve this reference.)

The program, `collect2`, is installed as ‘ld’ in the directory where the passes of the compiler are installed. When `collect2` needs to find the *real* ‘ld’, it tries the following filenames.

- `real-ld` in the directories listed in the compiler’s search directories.
- `real-ld` in the directories listed in the environment variable, `PATH`.
- The file specified in the `REAL_LD_FILE_NAME` configuration macro, if specified.
- `ld` in the compiler’s search directories, except that `collect2` will not execute itself recursively.
- `ld` in `PATH`.

The “compiler’s search directories” means all the directories where `gcc` searches for compiler passes, including directories that you specify with ‘-B’.

Cross-compilers search a little differently than normal configurations, using the following filenames.

- `real-ld` in the compiler’s search directories.
- `target-real-ld` in `PATH`.
- The file specified in the `REAL_LD_FILE_NAME` configuration macro, if specified.
- `ld` in the compiler’s search directories.
- `target-ld` in `PATH`.



`collect2` explicitly avoids running `ld` using the filename under which `collect2` itself was invoked. In fact, it remembers up a list of such names—in case one copy of `collect2` finds another copy (or version) of `collect2` installed as `'ld'` in a second place in the search path.

`collect2` searches for the utilities `nm` and `strip` using the same algorithm as previous installation for `'ld'`.

## Installing GCC on the Sun

On Solaris (version 2.1), do not use the linker or other tools in `'/usr/ucb'` to build GCC. Use `'/usr/ccs/bin'`.

Make sure the environment variable `FLOAT_OPTION` is not set when you compile `'libgcc.a'`. If this option were set to `f68881` when `'libgcc.a'` is compiled, the resulting code would demand to be linked with a special startup file and would not link properly without special pains.

There is a bug in `alloca` in certain versions of the Sun library. To avoid this bug, install the binaries of GCC that were compiled by GCC. They use `alloca` as a built-in function and never the one in the library.

Some versions of the Sun compiler crash when compiling GCC. The problem is a segmentation fault in `cpp`. This problem seems to be due to the bulk of data in the environment variables. You may be able to avoid it by using the following command to compile GCC with Sun CC:

```
make CC="TERMCAP=x OBJDUMP=x LIBFUNCS=x STAGESTUFF=x cc"
```

SunOS 4.1.3 and 4.1.3 U1 have bugs that can cause intermittent core dumps when compiling GCC. A common symptom is an internal compiler error which does not recur if you run it again. To fix the problem, install Sun recommended patch 100726 (for SunOS 4.1.3) or 101508 (for SunOS 4.1.3 U1), or upgrade to a later SunOS release.

## Installing GCC on VMS

The VMS version of GCC is distributed in a backup “save set” containing both source code and precompiled binaries. To install the `gcc` command so you can use the compiler easily, in the same manner as you use the VMS C compiler, you must install the VMS CLD file for GCC as follows.

1. Define the VMS logical names `GNU_CC` and `GNU_CC_INCLUDE` to point to the directories where the GCC executables (`'gcc-cpp.exe'`, `'gcc-cc1.exe'`, etc.) and the C include files are kept respectively. This should be done with the following commands.

```
assign /system /translation=concealed -disk:[gcc.] gnu_cc
assign /system /translation=concealed -disk:[gcc.include.]
 gnu_cc_include
```

Include the appropriate disk and directory names. These commands can be placed in your system startup file so they will be executed whenever the machine is rebooted. You may, if you choose, do this using the `GCC_INSTALL.COM` script in the `[GCC]` directory.

2. Install the `gcc` command with the following declaration.

```
set command /table=sys$common:[syslib]dcltables \
- / output=sys$common:[syslib]dcltables \
gnu_cc:[000000]gcc
install replace sys$common:[syslib]dcltables
```

3. To install the help file, use the following declaration.

```
library/help sys$library:helpplib.hlb gcc.hlp
```

Now, invoke the compiler with a command like `'gcc /verbose file.c'`, which is equivalent to the command `'gcc -v -c file.c'` in UNIX.

If you wish to use G++, you must first install GCC, and then perform the following steps.

1. Define the VMS logical name `GNU_GXX_INCLUDE` to point to the directory where the preprocessor will search for the C++ header files. This can be done with the following command.

```
assign /system /translation=concealed -disk:[\
gcc.gxx_include.] gnu_gxx_include
```

Include the appropriate disk and directory name. If you are going to be using `libg++`, this is where the `libg++` install procedure will install the `libg++` header files.

2. Obtain the file `'gcc-cc1plus.exe'`, and place this in the same directory that `'gcc-cc1.exe'` is kept. The G++ compiler can be invoked with a command like `'gcc /plus /verbose file.cc'`, which is equivalent to the command `'g++ -v -c file.cc'` in UNIX.

We try to put corresponding binaries and sources on the VMS distribution tape. But sometimes the binaries will be from an older version than the sources, because we don't always have time to update them. (Use the `'/version'` option to determine the version number of the binaries and compare it with the source file `'version.c'` to tell whether this is so.) In this case, you should use the binaries you get to recompile the sources. If you must recompile, use the following steps.

1. Execute the command procedure, `'vmsconfig.com'` to set up the files, `'tm.h'`, `'config.h'`, `'aux-output.c'`, and `'md. '`, and to create files, `'tconfig.h'` and `'hconfig.h'`. This procedure also creates several linker option files used by `'make-cc1.com'` and a data file used by `'make-l2.com'`.

```
@vmsconfig.com
```

2. Setup the logical names and command tables as defined in Step 1. In addition, define the VMS logical name `GNU_BISON` to point to the directories where the Bison executable is kept. This should be done with the following command.

```
assign /system /translation=concealed -disk:[bison.] \
 gnu_bison
```

If you want, use the `INSTALL_BISON.COM` script in the `[BISON]` directory.

3. Install the `BISON` command with the following command.

```
set command /table=sys$common:[syslib]dcltables \
 -output=sys$common:[syslib]dcltables - \
 gnu_bison:[000000]bison
install replace sys$common:[syslib]dcltables
```

4. Type `@make-gcc` to recompile everything (alternatively, submit the file `'make-gcc.com'` to a batch queue). If you wish to build the G++ compiler as well as the GCC compiler, you must first edit `make-gcc.com` and follow the instructions that appear in the comments.
5. In order to use GCC, you need a library of functions which GCC compiled code will call to perform certain tasks, and these functions are defined in the file `'libgcc2.c'`.

To compile this you should use the command procedure, `'make-l2.com'`, which will generate the library `'libgcc2.olb'`. Build `'libgcc2.olb'` using the compiler built from the same distribution that `'libgcc2.c'` came from, and `make-gcc.com` will automatically do all of this for you. To install the library, use the following commands.

```
library gnu_cc:[000000]gcclib/delete=(new,eprintf)
library gnu_cc:[000000]gcclib/delete=L_*
library libgcc2/extract=*/output=libgcc2.obj
library gnu_cc:[000000]gcclib libgcc2.obj
```

The first command simply removes old modules that will be replaced with modules from `libgcc2` under different module names. The modules `new` and `eprintf` may not actually be present in your `gcclib.olb`—if the VMS librarian complains about those modules not being present, simply ignore the message and continue on with the next command. The second command removes the modules that came from the previous version of the library `libgcc2.c`. Whenever you update the compiler on your system, you should also update the library with the previous procedure.

6. You may wish to build GCC in such a way that no files are written to the directory where the source files reside. An example would be the when the source files are on a read-only disk. In these cases, execute the following DCL commands (substituting your actual path names) where the `'dual:[gcc.source_dir]'` directory contains the source code, and the `'dua0:[gcc.build_dir]'` directory is meant to contain all of the generated object files and executables.

```
assign dua0:[gcc.build_dir.]/translation=concealed, \
-dual:[gcc.source_dir.]/translation=concealed gcc_build
set default gcc_build:[000000]
```

Once you have done this, you can proceed building GCC as previously described. (Keep in mind that 'gcc\_build' is a rooted logical name, and thus the device names in each element of the search list must be an actual physical device name rather than another rooted logical name).

7. If you are building GCC with a previous version of GCC, you also should check to see that you have the newest version of the assembler. In particular, GCC version 2 treats global constant variables slightly differently from GCC version 1, and GAS version 1.38.1 does not have the patches required to work with GCC version 2. If you use GAS 1.38.1, then extern const variables will not have the read-only bit set, and the linker will generate warning messages about mismatched psect attributes for these variables. These warning messages are merely a nuisance, and can safely be ignored.

If you are compiling with a version of GCC older than 1.33, specify '/DEFINE=("inline=")' as an option in all the compilations. This requires editing all the gcc commands in 'make-cc1.com'. (The older versions had problems supporting inline.) Once you have a working 1.33 or newer GCC, you can change this file back.

8. If you want to build GCC with the VAX C compiler, you will need to make minor changes in 'make-cccp.com' and 'make-cc1.com' to choose alternate definitions of CC, CFLAGS, and LIBS. See comments in those files. However, you must also have a working version of the GNU assembler (GNU as, also known as GAS) as it is used as the back-end for GCC to produce binary object modules and is not included in the GCC sources. GAS is also needed to compile 'libgcc2' in order to build 'gcclib' (see Step 5); 'make-l2.com' expects to be able to find it operational in 'gnu\_cc:[000000]gnu-as.exe'.

To use GCC on VMS, you need the VMS driver programs 'gcc.exe', 'gcc.com', and 'gcc.cld'. They are distributed with the VMS binaries (gcc-vms) rather than the GCC sources. GAS is also included in 'gcc-vms', as is bison.

Once you have successfully built GCC with VAX C, you should use the resulting compiler to rebuild itself. Before doing this, be sure to restore the CC, CFLAGS, and LIBS definitions in 'make-cccp.com' and 'make-cc1.com'. The second generation compiler will be able to take advantage of many optimizations that must be suppressed when building with other compilers.

Under previous versions of GCC, the generated code would occasionally give strange results when linked with the sharable 'VAXCRT' library. Now this should work.

Even with this version, however, GCC itself should not be linked with the sharable 'VAXCRT'. The version of qsort in 'VAXCRT' has a bug (known to be present in VMS

versions V4.6 through V5.5) which causes the compiler to fail.

The executables are generated by `'make-cc1.com'` and `'make-cccp.com'` uses the object library version of `'VAXCRTL'` in order to make use of the `qsort` routine in `'gcclib.olb'`.

If you wish to link the compiler executables with the shareable image version of `'VAXCRTL'`, you should edit the file `'tm.h'` (created by `'vmsconfig.com'`) to define the macro `QSORT_WORKAROUND`.

`QSORT_WORKAROUND` is always defined when GCC is compiled with VAX C, to avoid a problem in case `'gcclib.olb'` is not yet available.

## Using GCC on VMS

See the following documentation for how to use GCC on VMS.

### Include files and VMS

Due to the differences between the file systems of UNIX and VMS, GCC attempts to translate filenames in `'#include'` into names that VMS will understand. The basic strategy is to prepend a prefix to the specification of the include file, convert the whole filename to a VMS filename, and then try to open the file. GCC tries various prefixes one by one until one of them succeeds:

- The first prefix is the `'GNU_CC_INCLUDE:'` logical name: this is where GNU C header files are traditionally stored. If you wish to store header files in non-standard locations, then you can assign the logical `'GNU_CC_INCLUDE'` to be a search list, where each element of the list is suitable for use with a rooted logical.
- The next prefix tried is `'SYS$SYSROOT:[SYSLIB.]'`. This is where VAX-C header files are traditionally stored.
- If the include file specification by itself is a valid VMS filename, the preprocessor then uses this name with no prefix in an attempt to open the include file.
- If the file specification is not a valid VMS filename (i.e., the specification does not contain a device or a directory specifier, and contains a `'/'` character), the preprocessor tries to convert it from UNIX syntax to VMS syntax. Conversion works like this: the first directory name becomes a device, and the rest of the directories are converted into VMS-format directory names. For example, the name `'x11/foobar.h'` is translated to `'x11:[000000]foobar.h'` or `'x11:foobar.h'`, whichever one can be opened. This strategy allows you to assign a logical name to point to the actual location of the header files.
- If none of these strategies succeeds, the `'#include'` fails. Include directives of the following form.

```
#include foobar
```

Such directives are a common source of incompatibility between VAX-C and GCC.

VAX-C treats them much like a standard `#include <foobar.h>` directive. That is incompatible with the ANSI C behavior implemented by GCC, to expand the name `'foobar'` as a macro.

Macro expansion should eventually yield one of the two standard formats for `#include`:

```
#include "file"
#include < file>
```

If you have this problem, the best solution is to modify the source to convert the `#include` directives to one of the two standard forms. That will work with either compiler. If you want a quick and dirty fix, define the filenames as macros with the proper expansion, like the following example.

```
#define stdio <stdio.h>
```

This will work, as long as the name doesn't conflict with anything else in the program. Another source of incompatibility is that VAX-C assumes the following.

```
#include "foobar"
```

The program is actually asking for the file `'foobar.h'`. GCC does not make this assumption, and instead takes what you ask for literally; it tries to read the file `'foobar'`. The best way to avoid this problem is to always specify the desired file extension in your include directives. GCC for VMS is distributed with a set of include files that is sufficient to compile most general purpose programs. Even though the GCC distribution does not contain header files to define constants and structures for some VMS system-specific functions, there is no reason why you cannot use GCC with any of these functions. You first may have to generate or create header files, either by using the public domain utility UNSDL (which can be found on a DECUS tape), or by extracting the relevant modules from one of the system macro libraries, and using an editor to construct a C header file. A `#include` filename cannot contain a DECNET node name. The preprocessor reports an I/O error if you attempt to use a node name, whether explicitly, or implicitly via a logical name.

## Global declarations and VMS

GCC does not provide the `globalref`, `globaldef` and `globalvalue` keywords of VAX-C. You can get the same effect with an obscure feature of GAS, the GNU assembler. (This requires GAS version 1.39 or later.) The following macros allow you to use this feature in a fairly natural way.

```
#ifdef __GNUC__
#define GLOBALREF(TYPE,NAME) \
 TYPE NAME \
 asm ("_$$PsectAttributes_GLOBSYMBOL$$" #NAME)
#define GLOBALDEF(TYPE,NAME,VALUE) \
 TYPE NAME \
 asm ("_$$PsectAttributes_GLOBSYMBOL$$" #NAME) \

```

```

 = VALUE
#define GLOBALVALUEREDEF(TYPE,NAME) \
 const TYPE NAME[1] \
 asm ("_$$PsectAttributes_GLOBALVALUE$$" #NAME)
#define GLOBALVALUEDEF(TYPE,NAME,VALUE) \
 const TYPE NAME[1] \
 asm ("_$$PsectAttributes_GLOBALVALUE$$" #NAME) \
 = {VALUE}
#else
#define GLOBALREF(TYPE,NAME) \
 globalref TYPE NAME
#define GLOBALDEF(TYPE,NAME,VALUE) \
 globaldef TYPE NAME = VALUE
#define GLOBALVALUEDEF(TYPE,NAME,VALUE) \
 globalvalue TYPE NAME = VALUE
#define GLOBALVALUEREDEF(TYPE,NAME) \
 globalvalue TYPE NAME
#endif

```

(The ‘\_\$\$PsectAttributes\_GLOBALSYMBOL’ prefix at the start of the name is removed by the assembler, after it has modified the attributes of the symbol). These macros are provided in the VMS binaries distribution in a header file ‘GNU\_HACKS.H’. An example of the usage is the following.

```

GLOBALREF (int, ijk);
GLOBALDEF (int, jkl, 0);

```

The macros, GLOBALREF and GLOBALDEF, cannot be used straightforwardly for arrays, since there is no way to insert the array dimension into the declaration at the right place. However, declare an array with these macros if you first define a typedef for the array type, like the following declaration.

```

typedef int intvector[10];
GLOBALREF (intvector, foo);

```

Array and structure initializers will also break the macros; you can define the initializer to be a macro of its own, or you can expand the GLOBALDEF macro by hand. You may find a case where you wish to use the GLOBALDEF macro with a large array, but you are not interested in explicitly initializing each element of the array. In such cases you can use an initializer like: ‘{0,}’, which will initialize the entire array to ‘0’.

A shortcoming of this implementation is that a variable declared with GLOBALVALUEREDEF or GLOBALVALUEDEF is always an array. For example, the following declaration gives the variable ‘ijk’ as an array of type ‘int [1]’.

```

GLOBALVALUEREDEF(int, ijk);

```

This is done because a globalvalue is actually a constant; its *value* is what the linker

would normally consider an address. That is not how an integer value works in C, but it is how an array works. So treating the symbol as an array name gives consistent results—with the exception that the value seems to have the wrong type. ***Don't try to access an element of the array.*** It doesn't have any elements. The array *address* may not be the address of actual storage.

The fact that the symbol is an array may lead to warnings where the variable is used. Insert type casts to avoid the warnings. Here is an example; it takes advantage of the ANSI C feature allowing macros that expand to use the same name as the macro itself.

```
GLOBALVALUEREDEF (int, ss$_normal);
GLOBALVALUEDEF (int, xyzzy, 123);
#ifdef __GNUC__
#define ss$_normal ((int) ss$_normal)
#define xyzzy ((int) xyzzy)
#endif
```

Don't use GLOBALDEF or GLOBALREF with a variable whose type is an enumeration type; this is not implemented. Instead, make the variable an integer, and use a GLOBALVALUEDEF for each of the enumeration values. An example of this would be the following declaration.

```
#ifdef __GNUC__
GLOBALDEF (int, color, 0);
GLOBALVALUEDEF (int, RED, 0);
GLOBALVALUEDEF (int, BLUE, 1);
GLOBALVALUEDEF (int, GREEN, 3);
#else
enum globaldef color {RED, BLUE, GREEN = 3};
#endif
```

## Other VMS issues

GCC automatically arranges for `main` to return 1 by default if you fail to specify an explicit return value. This will be interpreted by VMS as a status code indicating a normal successful completion. Version 1 of GCC did not provide this default. GCC on VMS works only with the GNU assembler, GAS. You need version 1.37 or later of GAS in order to produce value debugging information for the VMS debugger. Use the ordinary VMS linker with the object files produced by GAS. Under previous versions of GCC, the generated code would occasionally give strange results when linked to the sharable 'VAXCRT' library. Now this should work.

A caveat for use of `const` global variables: the `const` modifier must be specified in every external declaration of the variable in all of the source files that use that variable. Otherwise the linker will issue warnings about conflicting attributes for the variable. Your program will still work despite the warnings, but the variable will be placed in writable storage.

Although the VMS linker does distinguish between upper and lower case letters in



global symbols, most VMS compilers convert all such symbols into upper case and most run-time library routines also have upper case names. To be able to reliably call such routines, GCC (by means of the assembler, GAS) converts global symbols into upper case like other VMS compilers. However, since the usual practice in C is to distinguish case, GCC (using GAS) tries to preserve usual C behavior by augmenting each name that is not all lower case. This means truncating the name to at most 23 characters and then adding more characters at the end which encode the case pattern of those 23. Names which contain at least one dollar sign are an exception; they are converted directly into upper case without augmentation.

Name augmentation yields bad results for programs that use precompiled libraries (such as Xlib) which were generated by another compiler. You can use the compiler option `‘/NOCASE_HACK’` to inhibit augmentation; it makes external C functions and variables case-independent as is usual on VMS. Alternatively, you could write all references to the functions and variables in such libraries using lower case; this will work on VMS, but is not portable to other systems. The compiler option `‘/NAMES’` also provides control over global name handling.

Function and variable names are handled somewhat differently with G++. The G++ compiler performs *name mangling* on function names, which means that it adds information to the function name to describe the data types of the arguments that the function takes. One result of this is that the name of a function can become very long. Since the VMS linker only recognizes the first 31 characters in a name, special action is taken to ensure that each function and variable has a unique name that can be represented in 31 characters. If the name (plus a name augmentation, if required) is less than 32 characters in length, then no special action is performed. If the name is longer than 31 characters, the assembler (GAS) will generate a hash string based upon the function name, truncate the function name to 23 characters, and append the hash string to the truncated name. If the `‘/VERBOSE’` compiler option is used, the assembler will print both the full and truncated names of each symbol that is truncated.

The `‘/NOCASE_HACK’` compiler option should not be used when you are compiling programs that use `libg++`. `libg++` has several instances of objects (i.e., `Filebuf` and `filebuf`) which become indistinguishable in a case-insensitive environment. This leads to cases where you need to inhibit augmentation selectively (if you were using `libg++` and Xlib in the same program, for example). There is no special feature for doing this, but you can get the result by defining a macro for each mixed case symbol for which you wish to inhibit augmentation. The macro should expand into the lower case equivalent of itself, as in the following example (such macro definitions can be placed in a header file to minimize the number of changes to your source code).

```
#define StuDlyCapS studlycaps
```



# GNU CC command options

---

When you invoke the GNU compiler, GCC, it normally does preprocessing, compilation, assembly and linking. The *overall options* allow you to stop this process at intermediate stages. Some options control the preprocessor, others the compiler itself, while still other options control the assembler and linker. For example, the ‘-c’ option says not to run the linker. Then the output consists of object files’ output by the assembler. Other options are passed on to one stage of processing. Most of the command line options that you can use with GCC are useful for C programs; when an option is only useful with another language (usually C++), the explanation says so explicitly.

The following documentation summarizes the command options for GCC.

- “Overall options” on page 67
- “C language options” on page 67
- “C++ language options” on page 67
- “Warning options” on page 68
- “Debugging options” on page 68
- “Optimization options” on page 68
- “Preprocessor options” on page 68
- “Assembler option” on page 69
- “Linker options” on page 69
- “Directory options” on page 69

- “Target options” on page 69
- “Code generation options” on page 74

The following documentation summarizes the command options for GCC for specific machines.

- “Machine dependent options” on page 69
  - “AMD29K options” on page 69
  - “ARC options” on page 69
  - “ARM options” on page 70
  - “Clipper options” on page 70
  - “Convex options” on page 70
  - “D10V options” on page 70
  - “DEC Alpha options” on page 70
  - “Hitachi H8/300 options” on page 70
  - “Hitachi SH options” on page 71
  - “HPPA options” on page 71
  - “IBM RS/6000 and PowerPC options” on page 71
  - “IBM RT options” on page 71
  - “i386 options” on page 72
  - “i960 options” on page 72
  - “M32R/D options” on page 72
  - “MIPS options” on page 72
  - “MN10300 options” on page 72
  - “Motorola 68K options” on page 73
  - “Motorola 88K options” on page 73
  - “SPARC options” on page 73
  - “System V options” on page 73
  - “Thumb options” on page 73
  - “Vax options” on page 74

# Option summary for GCC

What follows is a summary of all the options, grouped by type.

If the description for a particular option does not mention a source language, you can use that option with all supported languages. The `gcc` program accepts options and filenames as operands. Many options have multi-letter names; therefore multiple single-letter options may *not* be grouped. You can mix options and other arguments. Order does matter when you use several options of the same kind; for example, if you specify `-L` more than once, the directories are searched in the order specified.

Many options have long names starting with `-f` or with `-w` (for example, `-fstrength-reduce` or `-Wformat` with both having positive and negative forms; the negative form of `-ffoo` would be `-fno-foo`). This documentation generally discusses only one of the forms, whichever one not being default.

For each option, refer to the corresponding documentation for more details.

## Overall options

See “Options controlling the kind of output” on page 75.

```
-c -S -E -o file -pipe -v -x language
```

## C language options

See “Options controlling C dialect” on page 79.

```
-ansi -fallow-single-precision -fcond-mismatch
-fno-asm -fno-builtin -ffreestanding -fhosted
-fsigned-bitfields -fsigned-char
-funsigned-bitfields -funsigned-char
-fwritable-strings -traditional -traditional-cpp
-trigraphs
```

## C++ language options

See “Options controlling C++ dialect” on page 85.

```
-fall-virtual -fdollars-in-identifiers
-felide-constructors -fenum-int-equiv
-fexternal-templates -ffor-scope -fno-for-scope
-fhandle-signatures -fmemoize-lookups
-fname-mangling-version-n -fno-default-inline
-fno-gnu-keywords -fnonnull-objects
-fguiding-decls -foperator-names -fstrict-prototype
-fthis-is-variable -ftemplate-depth-n
-fthis-is-variable -nostdinc++ -traditional +en
```

## Warning options

See “Options to request or suppress warnings” on page 93.

```
-fsyntax-only -pedantic -pedantic-errors -w -W -Wall
-Waggregate-return -Wbad-function-cast -Wcast-align
-Wcast-qual -Wchar-subscript -Wcomment -Wconversion
-Werror -Wformat -Wid-clash-len -Wimplicit -Wimplicit-int
-Wimplicit-function-declarations -Wimport -Winline
-Wlarger-than- len -Wmain -Wmissing-declarations
-Wmissing-prototypes -Wnested-externs -Wno-import
--Wold-style-cast -Woverloaded-virtual
-Wparentheses -Wpointer-arith -Wredundant-decls
-Wreorder -Wreturn-type -Wshadow -Wsign-compare
-Wstrict-prototypes -Wswitch -Wsynth
-Wtemplate-debugging -Wtraditional -Wtrigraphs
-Wundef -Wuninitialized -Wunused -Wwrite-strings
```

## Debugging options

See “Options for debugging” on page 103

```
-a -ax -dletters -fpretend-float -fprofile-arcs
-ftest-coverage -g -glevel -gcoff -gdwarf -gdwarf-1
-gdwarf+1 -gdwarf-2 -ggdb -gstabs -gstabs+
-gxcoff -gxcoff+ -p -pg -print-file-name=library
-print-libgcc-file-name -print-prog-name=program
-print-search-dirs -save-temps
```

## Optimization options

See “Options that control optimization” on page 111.

```
-fbranch-probabilities -fcaller-saves
-fcse-follow-jumps -fcse-skip-blocks
-fdelayed-branch -fexpensive-optimizations
-ffast-math -ffloat-store -fforce-addr -fforce-mem
-ffunction-sections -finline-functions
-fkeep-inline-functions -fno-default-inline
-fno-defer-pop -fno-function-cse -fno-inline
-fno-peephole -fomit-frame-pointer -fregmove
-frerun-cse-after-loop -fschedule-insns
-fschedule-insns2 -fshorten-lifetimes
-fstrength-reduce -fthread-jumps -funroll-all-loops
-funroll-loops -O -O0 -O1 -O2 -O3 Os
```

## Preprocessor options

See “Options controlling the preprocessor” on page 117.

```
-Aquestion(answer) -C -dD -dM -dN -Dmacro[=defn]
```

```
-E -H -idirafter dir -include file -imacros file
-iprefix file -iwithprefix dir -iwithprefixbefore dir
-isystem dir -M -MD -MM -MMD -MG -nostdinc -P
-trigraphs -undef -Umacro -Wp, option
```

## Assembler option

See “Passing options to the assembler” on page 121.

```
-Wa, option
```

## Linker options

See “Options for linking” on page 123.

```
object-file-name -llibrary -nostartfiles
-nodfaultlibs -nostdlib -s -static -shared
-symbolic -Wl, option -Xlinker option -u symbol
```

## Directory options

See “Options for directory search” on page 127.

```
-Bprefix -Idir -I- -Ldir -specs=file
```

## Target options

See “Specifying target machine and compiler version” on page 129.

```
-bmachine -Vversion
```

## Machine dependent options

See “Hardware models and configurations” on page 131.

### AMD29K options

See “AMD29K options” on page 133.

```
-m29000 -m29050 -mbw -mnbw -mdw -mndw -mlarge -mnormal
-msmall -mkernel-registers -mno-reuse-arg-regs
-mno-stack-check -mno-storem-bug -mreuse-arg-regs
-msoft-float -mstack-check -mstorem-bug -muser-registers
```

### ARC options

See “ARC options” on page 135.

```
-EB -EL -mmangle-cpu-mcpu=cpu -mtext=text section
-mdata=data section -mrodata=readonly data section
```

## ARM options

See “ARM options” on page 136.

```
-mapcs-frame -mno-apcs-frame -mapcs-26 -mapcs-32
-mapcs-stack-check -mno-apcs-stack-check -mapcs-float
-mno-apcs-float -mapcs-reentrant -mno-apcs-reentrant
-msched-prolog -mno-sched-prolog
-mlittle-endian -mbig-endian -mwords-little-endian
-mshort-load-bytes -mno-short-load-bytes -mshort-load-words
-mno-short-load -msoft-float -mhard-float -mfpe
-mthumb-interwork -mno-thumb-interwork -mcpu= -march= -mfpe=
-mstructure-size-boundary= -mbsd -mxopen -mno-symrename
-mnop-fun-dllimport -mno-nop-fun-dllimport
```

## Clipper options

See “Clipper options” on page 140.

```
-mc300 -mc400
```

## Convex options

See “Convex options” on page 141.

```
-mc1 -mc2 -mc32 -mc34 -mc38 -margcount -mnoargcount
-mlong32 -mlong64 -mvolatile-cache -mvolatile-nocache
```

## D10V options

See “D10V options” on page 143.

```
-mint16 -mint32 -maddac3 -mno-addac3 -mdouble32 -mdouble64
-maccum -mno-accum -msim -masm-optimize -mno-asm-optimize
-msmall-insns -mno-small-insns -mbranch-cost=n -mcond-exec=n
-mno-cond-move -mcond-move
```

## DEC Alpha options

See “DEC Alpha options” on page 144.

```
-mfp-regs -mno-fp-regs -mno-soft-float -msoft-float
-malpha-as -mgas -mieee -mieee-with-inexact
-mieee-conformant -mfp-trap-mode -mfp-rounding-mode
-mtrap-precision -mbuild-constants -mcpu=cpu type
-mbwx -mno-bwx -mcix -mno-cix -mmax -mno-max
-mmmemory-latency=time
```

## Hitachi H8/300 options

See “Hitachi H8/300 options” on page 147.

```
-mrelax -mh -ms -mint32 -malign-300
```



## Hitachi SH options

See “Hitachi H8/300 options” on page 147.

`-ml -m2 -m3 -m3e -mb -ml -mrelax`

## HPPA options

See “HPPA options” on page 149.

`-mbig-switch -mdisable-fpregs -mdisable-indexing  
-mfast-indirect-calls -mgas -mjump-in-delay  
-mlong-load-store -mno-big-switch -mno-disable-fpregs  
-mno-disable-indexing -mfast-indirect-calls -mno-gas  
-mno-jump-in-delay -mno-long-load-store  
-mno-portable-runtime -mno-soft-float -mno-space  
-mno-space-regs -msoft-float -mpa-risc-1-0 -mpa-risc-1-1  
-mportable-runtime -mschedule=list -mspace -mspace-regs`

## IBM RS/6000 and PowerPC options

See “IBM RS/6000 and PowerPC options” on page 151.

`-mcpu=cpu type -mtune=cpu type -mpower -mno-power  
-mpower2 -mno-power2 -mpowerpc -mno-powerpc  
-mpowerpc-gpopt -mno-powerpc-gpopt -mpowerpc-gfxopt  
-mno-powerpc-gfxopt -mnew-mnemonics  
-mno-new-mnemonics -mfull-toc -mminimal-toc  
-mno-fop-in-toc -mno-sum-in-toc -mxl-call  
-mno-xl-call -mthreads -mpe -msoft-float -mhard-float  
-mmultiple -mno-multiple -mstring -mno-string  
-mupdate -mno-update -mfused-madd -mno-fused-madd  
-mbit-align -mno-bit-align -mstrict-align  
-mno-strict-align -mrelocatable -mno-relocatable  
-mrelocatable-lib -mno-relocatable-lib  
-mtoc -mno-toc -mtraceback -mno-traceback  
-mlittle -mlittle-endian -mbig -mbig-endian  
-mcall-aix -mcall-sysv -mprototype -mno-prototype  
-msim -mmvme -memb -mads -myellowknife  
-msdata=opt -G num`

## IBM RT options

See “IBM RT options” on page 159.

`-mcall-lib-mul -mfp-arg-in-fpregs -mfp-arg-in-gregs  
-mfull-fp-blocks -mhc-struct-return -min-line-mul  
-mminimum-fp-blocks -mnohc-struct-return`

## **i386 options**

See “Intel x86 options” on page 160.

```
-mcpu=cpu type -march=cpu type
-m486 -m386 -mieee-fp -mno-fancy-math-387
-mno-fp-ret-in-387 -msoft-float -msvr3-shlib
-mno-wide-multiply -mrtd -malign-double
-mreg-alloc=list -mregparm=num -malign-jumps=num
-malign-loops=num -malign-functions=num
```

## **i960 options**

See “Intel 960 options” on page 163.

```
-mcpu type -masm-compact -mclean-linkage
-mcode-align -mcomplex-addr -mleaf-procedures
-mic-compact -mic2.0-compact -mic3.0-compact
-mintel-asm -mno-clean-linkage -mno-code-align
-mno-complex-addr -mno-leaf-procedures
-mno-old-align -mno-strict-align -mno-tail-call
-mnumerics -mold-align -msoft-float -mstrict-align
-mtail-call
```

## **M32R/D options**

See “M32R/D options” on page 165.

```
-mcode-model=model type -msdata=sdata type -G num
```

## **MIPS options**

See “MIPS options” on page 166.

```
-mabicalls -mcpu=cpu type -membedded-data
-membedded-pic -mfp32 -mfp64 -mgas -mgp32 -mgp64
-mgpopt -mhalf-pic -mhard-float -mint64 -mips1
-mips2 -mips3 -mlong64 -mlong-calls -mmemcpy
-mmips-as -mmips-tfile -mno-abicalls
-mno-embedded-data -mno-embedded-pic -mno-gpopt
-mno-long-calls -mno-memcpy -mno-mips-tfile
-mno-rnames -mno-stats -mrnames -msoft-float -m4650
-msingle-float -mmad -mstats -EL -EB -G num -nocpp
```

## **MN10300 options**

See “MN10300 options” on page 170.

```
-mmult-bug -mno-mult-bug
```

## Motorola 68K options

See “Motorola 68K options” on page 171.

```
-m68000 -m68020 -m68020-40 -m68020-60 -m68030
-m68040 -m68060 -m68881 -mbitfield -mc68000
-mc68020 -mfpa -mnobitfield -mrtd -mshort -msoft-float
-malign-int
```

## Motorola 88K options

See “Motorola 88K options” on page 174.

```
-m88000 -m88100 -m88110 -mbig-pic -mcheck-zero-division
-mhandle-large-shift -midentify-revision
-mno-check-zero-division -mno-ocs-debug-info
-mno-ocs-frame-position -mno-optimize-arg-area
-mno-serialize-volatile -mno-underscores
-mocs-debug-info -mocs-frame-position
-moptimize-arg-area -mserialize-volatile
-mshort-data-num -msvr3 -msvr4 -mtrap-large-shift
-muse-div-instruction -mversion-03.00
-mwarn-passed-structs
```

## SPARC options

See “SPARC options” on page 178.

```
-mcpu=cpu type -mtune=cpu type -mapp-regs
-mbroken-saverestore -mcypress -mepilogue -mflat -mfpu
-mfullany -mhard-float -mhard-quad-float -mimpure-text
-mint32 -mint64 -mliveg0 -mlong32 -mlong64
-mmedlow -mmedany -mno-app-regs -mno-epilogue -mno-flat
-mno-fpu -mno-impure-text -mno-stack-bias
-mno-unaligned-doubles -msoft-float -msoft-quad-float
-msparclite -mstack-bias -msupersparc
-munaligned-doubles -mv8
```

## System V options

See “System V options” on page 182.

```
-Qy -Qn -YP,paths -Ym,dir
```

## Thumb options

See “Thumb options” on page 183.

```
-mtpcs-frame -mno-tpcs-frame -mtpcs-leaf-frame
-mno-tpcs-leaf-frame -mlittle-endian -mbig-endian
-mthumb-interwork -mno-thumb-interwork
-mstructure-size-boundary= -mnop-fun-dllimport
-mno-nop-fun-dllimport -mcallee-super-interworking
-mno-callee-super-interworking -mcaller-super-interworking
-mno-caller-super-interworking
```

## **Vax options**

See “Vax options” on page 185.

`-mg -mgnu -munix`

## **Code generation options**

See “Options for code generation conventions” on page 187.

`-fcall-saved-reg -fcall-used-reg -fexceptions -ffixed-reg`  
`-finhibit-size-directive -fcheck-memory-usage`  
`-fprefix-function-name -fno-common -fno-ident`  
`-fno-gnu-linker -fpcc-struct-return -fpic -fPIC`  
`-freg-struct-return -fshared-data -fshort-enums`  
`-fshort-double -fvolatile -fvolatile-global`  
`-funaligned-pointers -funaligned-struct-hack`  
`-fverbose-asm -fpack-struct -fstack-check +e0 +e1`  
`-fargument-alias -fargument-noalias -fargument-noalias-global`

# 4

## Options controlling the kind of output

---

The following documentation discusses types of output and source files. Compilation can involve up to four of the following stages.

- *preprocessing*
- *compiling*
- *assembling*
- *linking*

The stages are always in that order. The first three stages apply to an individual source file. The compilation process ends with a last stage by producing an object file; linking combines all the object files (those newly compiled, and those specified as input) into an executable file. For any given input file, the *filename* suffix determines what kind of compilation is done, as the following descriptions help to clarify.

*file.c*

C source code which must be preprocessed.

*file.i*

C source code which should not be preprocessed.

*file.ii*

C++ source code which should not be preprocessed.

*file.m*

Objective-C source code.

**NOTE:** You must link with the `libobjc.a` library to make an Objective-C

program work.

*file.h*

C header file (not to be compiled or linked).

*file.cc*

*file.cxx*

*file.cpp*

*file.C*

C++ source code which must be preprocessed. Note that in ‘.cxx’, the last two letters must both be literally ‘x.’ Likewise, ‘.C’ refers to a literal capital C.

*file.s*

Assembler code.

*file.S*

Assembler code which must be preprocessed.

*other*

An object file to be fed straight into linking. Any filename without a *recognized* suffix is treated this way.

You can specify the input language explicitly with the ‘-x’ option:

-x *language*

Specify explicitly the *language* for the following input files (rather than letting the compiler choose a default based on the filename suffix). This option applies to all following input files until the next ‘-x’ option.

Possible values for *language* are the following.

|           |                    |                |
|-----------|--------------------|----------------|
| c         | objective-c        | c++            |
| c-header  | cpp-output         | c++-cpp-output |
| assembler | assembler-with-cpp |                |

-x none

Turn off any specification of a language, so that subsequent files are handled according to their filename suffixes (as they are if ‘-x’ has not been used at all).

If you only want some of the stages of compilation, you can use ‘-x’ (or filename suffixes) to tell `gcc` where to start, and one of the options ‘-c’, ‘-S’, or ‘-E’ to say where `gcc` is to stop.

**NOTE:** Some combinations (for example, -x cpp-output -E) instruct `gcc` to do nothing.

-c

Compile or assemble the source files, but do not link. The linking stage simply is not done. The ultimate output is in the form of an object file for each source file. By default, the object filename for a source file is made by replacing the suffix ‘.c’, ‘.i’, ‘.s’, etc., with ‘.o’. Unrecognized input files, not requiring compilation or assembly, are ignored.

- S  
Stop after the stage of compilation proper; do not assemble. The output is in the form of an assembler code file for each non-assembler input file specified. By default, the assembler filename for a source file is made by replacing the suffix `‘.c’`, `‘.i’`, etc., with `‘.s’`.  
Input files that don’t require compilation are ignored.
- E  
Stop after the preprocessing stage; do not run the compiler proper. The output is in the form of preprocessed source code, which is sent to the standard output. Input files which don’t require preprocessing are ignored.
- o *file*  
Place output in file *file*. This applies regardless to whatever sort of output is being produced, whether it be an executable file, an object file, an assembler file or preprocessed C code. Since only one output file can be specified, it does not make sense to use `‘-o’` when compiling more than one input file, unless you are producing an executable file as output.  
If `‘-o’` is not specified, the default is to put an executable file in `‘a.out’`, the object file for `‘source.suffix’` in `‘source.o’`, its assembler file in `‘source.s’`, and all preprocessed C source on standard output.
- v  
Print (on standard error output) the commands executed to run the stages of compilation. Also print the version number of the compiler driver program and of the preprocessor and the compiler proper.
- pipe  
Use pipes rather than temporary files for communication between the various stages of compilation. This fails to work on some systems where the assembler is unable to read from a pipe; but the GNU assembler has no trouble.





## Options controlling C dialect

---

The following options control the dialect of C (or languages derived from C, such as C++ and Objective C) that the compiler accepts:

`-ansi`

Support all ANSI standard C programs.

This turns off certain features of GNU C that are incompatible with ANSI C, such as the `asm`, `inline` and `typeof` keywords, and predefined macros such as `unix` and `vax` that identify the type of system you are using. It also enables the undesirable and rarely used ANSI trigraph feature, disallows the dollar sign symbol (\$) as part of identifiers, and disables recognition of C++ style comments with double forward-slash (//).

The alternate keywords, `__asm__`, `__extension__`, `__inline__` and `__typeof__`, continue to work despite `-ansi`. You would not want to use them in an ANSI C program, of course, but it is useful to put them in header files that might be included in compilations done with `-ansi`. Alternate predefined macros, such as `__unix__` and `__vax__`, are also available, with or without `-ansi`.

The `-ansi` option does not cause non-ANSI programs to be rejected gratuitously. For that, `-pedantic` is required in addition to `-ansi`. See “Options to request or suppress warnings” on page 93.

The macro, `__STRICT_ANSI__`, is predefined when the `-ansi` option is used. Some header files may notice this macro and refrain from declaring certain functions or

defining certain macros that the ANSI standard doesn't call for; this is to avoid interfering with any programs that might use these names for other things.

The functions, `alloca`, `abort`, `exit`, and `_exit`, are not built-in functions when `-ansi` is used.

### `-fno-asm`

Do not recognize `asm`, `inline` or `typeof` as a keyword, so that code can use these words as identifiers. You can use, instead, the keywords, `__asm__`, `__inline__` and `__typeof__`. `-ansi` implies `-fno-asm`.

In C++, this switch only affects the `typeof` keyword, since `asm` and `inline` are standard keywords.

You may want to use the `-fno-gnu-keywords` flag instead, as it also disables the other, C++-specific, extension keywords such as `headof`.

### `-fno-builtin`

Don't recognize built-in functions that do not begin with two leading underscores. Currently, the functions affected include `abort`, `abs`, `alloca`, `cos`, `exit`, `fabs`, `ffs`, `labs`, `memcmp`, `memcpy`, `sin`, `sqrt`, `strcmp`, `strcpy`, and `strlen`. GCC normally generates special code to handle certain built-in functions more efficiently; for instance, calls to `alloca` may become single instructions that adjust the stack directly, and calls to `memcpy` may become inline copy loops. The resulting code is often both smaller and faster, but since the function calls no longer appear as such, you cannot set a breakpoint on those calls, nor can you change the behavior of the functions by linking with a different library.

### `-fhosted`

Assert that compilation takes place in a hosted environment. This implies `'-fbuiltin'`. A hosted environment is one in which the entire standard library is available, and in which `main` has a return type of `int`. Examples are nearly everything except a kernel.

This is equivalent to `'-fno-freestanding'`.

### `-ffreestanding`

Assert that compilation takes place in a freestanding environment. This implies `'-fno-builtin'`. A freestanding environment is one in which the standard library may not exist, and program startup may not necessarily be at `main`. The most obvious example is an OS kernel. This is equivalent to `'-fno-hosted'`.

The `-ansi` option prevents `alloca` and `ffs` from being built-in functions, since these functions do not have an ANSI standard meaning.

### `-trigraphs`

Support ANSI C trigraphs. You don't want to know about this brain-damage. The `-ansi` option implies `-trigraphs`.

`-traditional`

Attempt to support some aspects of traditional C compilers. Specifically:

- All `extern` declarations take effect globally even if they are written inside of a function definition. This includes implicit declarations of functions.
- The newer keywords, `typeof`, `inline`, `signed`, `const` and `volatile`, are not recognized. (You can still use the alternative keywords, such as `__typeof__`, `__inline__`, and so on.)
- Comparisons between pointers and integers are always allowed.
- Integer types, `unsigned short` and `unsigned char`, promote to `unsigned int`.
- Out-of-range floating point literals are not an error.
- Certain constructs which ANSI regards as a single invalid preprocessing number, such as `0xe-0xd`, are treated as expressions instead.
- String constants are not necessarily constant; they are stored in writable space, and identical looking constants are allocated separately. (This is the same as the effect of `-fwritable-strings`.)
- All automatic variables not declared `register` are preserved by `longjmp`. Ordinarily, GNU C follows ANSI C: automatic variables not declared `volatile` may be clobbered.
- The character escape sequences, `\x` and `\a`, evaluate as the literal characters 'x' and 'a' respectively. Without `-traditional`, `\x` is a prefix for the hexadecimal representation of a character, and `\a` produces a bell.
- In C++ programs, assignment to `this` is permitted with `-traditional`. (The option, `-fthis-is-variable`, also has this effect.)

You may wish to use `'-fno-builtin'` as well as `'-traditional'` if your program uses names that are normally GNU C builtin functions for other purposes of its own.

You cannot use `'-traditional'` if you include any header files that rely on ANSI C features. Some vendors are starting to ship systems with ANSI C header files and you cannot use `-traditional` on such systems to compile files that include any system headers.

The `'-traditional'` option also enables the `'-traditional-cpp'` option which is described in the following discussion.

`-traditional-cpp`

Attempts to support some aspects of traditional C preprocessors. Specifically:

- Comments convert to nothing at all rather than to a space. This allows traditional token concatenation.

- In preprocessing directive, the pound/number symbol (#) must appear as the first character of a line.
- Macro arguments are recognized within string constants in a macro definition (and their values are stringified, though without additional quote marks, when they appear in such a context). The preprocessor always considers a string constant to end at a newline.
- The predefined macro, `__STDC__`, is not defined when you use `-traditional`, but `__GNUC__` is (since the GNU extensions which `__GNUC__` indicates are not affected by `-traditional`). If you need to write header files that work differently depending on whether `-traditional` is in use, by testing both of these predefined macros you can distinguish four situations: GNU C, traditional GNU C, other ANSI C compilers, and other old C compilers. The predefined macro, `__STDC_VERSION__`, is also not defined when you use `-traditional`. For more discussion of these and other predefined macros, see the “Standard predefined macros” on page 353 in *The C Preprocessor*.
- The preprocessor considers a string constant to end at a newline (unless the newline is escaped with a backslash) . Without `-traditional`, string constants can contain the newline backslash character.

`-fcond-mismatch`

Allow conditional expressions with mismatched types in the second and third arguments. The value of such an expression is void.

`-funsigned-char`

Let the type, `char`, be unsigned, like `unsigned char`.

Each kind of machine has a default for what `char` should be. It is either like `unsigned char`, by default, or like `signed char`, by default.

Ideally, a portable program should always use `signed char` or `unsigned char` when it depends on the *signedness* of an object. But many programs have been written to use plain `char` and expect it to be signed, or expect it to be unsigned, depending on the machines for which they were written. This option, and its inverse, let you make such a program work with the opposite default.

The type, `char`, is always a distinct type from each of `signed char` or `unsigned char`, even though its behavior is always just like one of those two dialect options.

`-fsigned-char`

Let the type, `char`, be signed, like `signed char`.

**NOTE:** This is equivalent to `-fno-unsigned-char`, which is the negative form of `-funsigned-char`. Likewise, the option, `-fno-signed-char`, is equivalent to `-funsigned-char`.

You may wish to use ‘-fno-builtin’ as well as ‘-traditional’ if your program uses names that are normally GNU C builtin functions for other purposes of its own.

You cannot use ‘-traditional’ if you include any header files that rely on ANSI C features. Some vendors are starting to ship systems with ANSI C header files and you cannot use ‘-traditional’ on such systems to compile files that include any system headers.

-fsigned-bitfields  
-funsigned-bitfields  
-fno-signed-bitfields  
-fno-unsigned-bitfields

These options control whether a bitfield is signed or unsigned, when the declaration does not use either `signed` or `unsigned`. By default, such a bitfield is signed, because this is consistent: the basic integer types such as `int` are signed types.

However, when ‘-traditional’ is used, bitfields are all, no matter what, unsigned.

-fwritable-strings

Store string constants in the writable data segment and don’t uniquize them. This is for compatibility with old programs which assume they can write into string constants. The ‘-traditional’ option also has this effect. Writing into string constants is a very bad idea; *constants* should be constant.

-fallow-single-precision

Do not promote single precision math operations to double precision, even when compiling with ‘-traditional’.

Traditional K&R C promotes all floating point operations to double precision, regardless of the sizes of the operands. On the architecture for which you are compiling, single precision may be faster than double precision. If you must use ‘-traditional’, but want to use single precision operations when the operands are single precision, use this option. This option has no effect when compiling with ANSI or GNU C conventions (the default).



# 6

## Options controlling C++ dialect

---

The following documentation describes the command-line options that are only meaningful for C++ programs; you can also use most of the GNU compiler options regardless of what language your program uses. For instance, you might compile a file, `firstClass.C`, like the following example.

```
g++ -g -felide-constructors -O -c firstClass.C
```

In the example, only `-felide-constructors` is an option meant only for C++ programs; you can use the other options with any language supported by GCC.

See also “Compiling C++ programs” on page 91.

The following discussion lists options that are only for compiling C++ programs.

`-fno-access-control`

Turn off all access checking. Mainly useful for working around bugs in the access control code.

`-fall-virtual`

Treat all possible member functions as virtual, implicitly. All member functions (except for constructor functions and `new` or `delete` member operators) are treated as virtual functions of the class where they appear.

This does not mean that all calls to these member functions will be made through the internal table of virtual functions. In some circumstances, the compiler can determine that a call to a given virtual function can be made directly; in these cases the calls are direct in any case.

### `-fcheck-new`

Default. Check the return value of the global operator `new` for 0 and do not initialize variables if call to `new` fails.

### `-fconserve-space`

Put uninitialized or runtime-initialized global variables into the common segment, as C does. This saves space in the executable at the cost of not diagnosing duplicate definitions. If you compile with this flag and your program mysteriously crashes after `main()` has completed, you may have an object that is being destroyed twice because two definitions were merged.

### `-fdollars-in-identifiers`

Accept the dollar sign (\$) in identifiers. You can also explicitly prohibit use of '\$' with the `-fno-dollars-in-identifiers` option. GNU C++ allows '\$' by default on some target systems but not others. Traditional C allowed the character '\$', to form part of identifiers. However, ANSI C and C++ forbids '\$' in identifiers.

### `-fembedded-cxx`

In compliance with the Embedded C++ specification, makes the use of templates, exception handling, multiple inheritance, or RTTI illegal. This makes the use of these keywords result in warnings by default: `template`, `typename`, `catch`, `throw`, `try`, `using`, `namespace`, `dynamic_cast`, `static_cast`, `reinterpret_cast`, `const_cast`, and `typeid`.

Add the flag, `-pedantic-errors`, to make the warnings for these things be given as errors.

### `-fenum-int-equiv`

Anachronistically permit implicit conversion of `int` to enumeration types. Current C++ allows conversion of `enum` to `int`, but not the reverse.

### `-fexternal-templates`

Cause template instantiations to obey `#pragma interface` and `implementation`; template instances are emitted or not according to the location of the template definition. See “Where’s the template?” on page 278 for more explanation of templates. This option is deprecated.

### `-falt-external-templates`

Similar to `-fexternal-templates`, but template instances are emitted or not according to the place where they are first instantiated. See “Where’s the template?” on page 278 for more explanation on templates. This option is deprecated.

### `-ffor-scope`

### `-fno-for-scope`

If `-ffor-scope` is specified, the scope of variables declared in a *for-init-statement* is limited to the ‘`for`’ loop itself, as specified by the draft C++ standard. If

`-fno-for-scope` is specified, the scope of variables declared in a



*for-init-statement* extends to the end of the enclosing scope, as was the case in old versions of `gcc`, and other (traditional) implementations of C++.

The default if neither flag is given is to follow the standard, but to allow and give a warning for old-style code that would otherwise be invalid, or have different behavior.

#### `-fno-gnu-keywords`

Do not recognize `classof`, `headof`, `signature`, `sigof` or `typeof` as a keyword, so that code can use these words as identifiers. Instead, use the keywords, `__classof__`, `__headof__`, `__signature__`, `__sigof__`, and `__typeof__`. `-ansi` implies `-fno-gnu-keywords`.

#### `-fguiding-decls`

Treat a function declaration with the same type as a potential function template instantiation as though it declares that instantiation, not a normal function. If a definition is given for the function later in the translation unit (or another translation unit if the target supports weak symbols), that definition will be used; otherwise the template will be instantiated.

This behavior reflects the C++ language prior to September 1996, when guiding declarations were removed.

This option implies `-fname-mangling-version-0`, and will not work with other name mangling versions.

#### `-fno-implicit-templates`

Never emit code for templates which are instantiated implicitly (i.e., by use); only emit code for explicit instantiations. See “Where’s the template?” on page 278 for more explanation on templates.

#### `-fhandle-signatures`

Recognize the `signature` and `sigof` keywords for specifying abstract types. The default (`-fno-handle-signatures`) is not to recognize them. See “Type abstraction using signatures” on page 282.

#### `-fhuge-objects`

Support virtual function calls for objects that exceed the size representable by a `short int`. Users should not use this flag by default; if you need to use it, the compiler will tell you so. If you compile any of your code with this flag, you must compile all of your code with this flag (including `libg++`, if you use it). This flag is not useful when compiling with `-fvtable-thunks`.

#### `-fno-implement-inlines`

To save space, do not emit out-of-line copies of inline functions controlled by `#pragma implementation`. This will cause linker errors if these functions are not inlined everywhere they are called.

`-fmemoize-lookups`  
`-fsave-memoized`

`-fsave-memoized` uses heuristics to compile faster. These heuristics are not enabled by default, since they are only effective for certain input files. Other input files compile more slowly.

The first time the compiler must build a call to a member function (or reference to a data member), it must (1) determine whether the class implements member functions of that name; (2) resolve which member function to call (which involves figuring out what sorts of type conversions need to be made); and (3) check the visibility of the member function to the caller. All of this adds up to slower compilation. Normally, the second time a call is made to that member function (or reference to that data member), it must go through the same lengthy process again. This means that code like the following makes six passes through all three steps.

```
cout << "This " << p << " has " << n << " legs.\n";
```

By using a software cache, a *hit* significantly reduces this cost. Unfortunately, using the cache introduces another layer of mechanisms which must be implemented, and so incurs its own overhead.

`-fmemoize-lookups` enables the software cache.

Because access privileges (visibility) to members and member functions may differ from one function context to the next, G++ may need to flush the cache. With the `-fmemoize-lookups` flag, the cache is flushed after every function that is compiled. The `-fsave-memoized` flag enables the same software cache, but when the compiler determines that the context of the last function compiled would yield the same access privileges of the next function to compile, it preserves the cache. This is most helpful when defining many member functions for the same class: with the exception of member functions which are friends of other classes, each member function has exactly the same access privileges as every other, and the cache need not be flushed.

The code that implements these flags has rotted; you should probably avoid using them.

`-fstrict-prototype`

Within an ‘extern “C”’ linkage specification, treat a function declaration with no arguments, such as `int foo();`, as declaring the function to take no arguments. Normally, such a declaration means that the function `foo` can take any combination of arguments, as in C. `-pedantic` implies `-fstrict-prototype` unless overridden with `-fno-strict-prototype`.

This flag no longer affects declarations with C++ linkage.

**-fname-mangling-version-*n***

Control the way in which names are mangled. Version 0 is compatible with versions of G++ before 2.8. Version 1 is the default. Version 1 will allow correct mangling of function templates. For example, version 0 mangling does not mangle `foo<int, double>` and `foo<int, char>` given the following declaration.

```
template <class T, class U> void foo(T t);
```

**-fno-nonnull-objects**

Don't assume that a reference is initialized to refer to a valid object. Although the current C++ Working Paper prohibits null references, some old code may rely on them; you can use `-fno-nonnull-objects` to turn on checking.

At the moment, the compiler only does this checking for conversions to virtual base classes.

**-foperator-names**

Recognize the operator name keywords `and`, `bitand`, `bitor`, `compl`, `not`, `or` and `xor` as synonyms for the symbols they refer to. `-ansi` implies `-foperator-names`.

**-frepo**

Enable automatic template instantiation. This option also implies `-fno-implicit-templates`. See “Where's the template?” on page 278 for more explanation of templates.

**-fsquangle****-fno-squangle**

`-fsquangle` will enable a compressed form of name mangling for identifiers. In particular, it helps to shorten very long names by recognizing types and class names which occur more than once, replacing them with special short ID codes. `-fsquangle` also requires any C++ libraries being used to be compiled with this option as well. The compiler has this functionality disabled (the equivalent of using `-fno-squangle`) by default.

**-fthis-is-variable**

Permit assignment to `'this'`. The incorporation of user-defined free store management into C++ has made assignment to `'this'` an anachronism. Therefore, by default it is invalid to assign to `this` within a class member function; that is, GNU C++ treats `'this'` in a member function of class `'x'` as a non-lvalue of type `'x*'`.

However, for backwards compatibility, you can make it valid with

`-fthis-is-variable`.

**-fvtable-thunks**

Use `thunks` to implement the virtual function dispatch table (`vtable`). The traditional (`cfront`-style) approach to implementing `vtables` was to store a pointer to the function and two offsets for adjusting the `'this'` pointer at the call site. Newer implementations store a single pointer to a `'thunk'` function which does any necessary adjustment and then calls the target function.

This option also enables a heuristic for controlling emission of `vtables`; if a class has any non-inline virtual functions, the `vtable` will be emitted in the translation unit containing the first one of those.

**-ftemplate-depth-n**

Set the maximum instantiation depth for template classes to `n`. A limit on the template instantiation depth is needed to detect endless recursions during template class instantiation. ANSI/ISO C++ conforming programs must not rely on a maximum depth greater than 17.

**-nostdinc**

Do not search for header files in the standard directories specific to C++, but do still search the other standard directories. (This option is used when building `libg++`.)

**-traditional**

For C++ programs (in addition to the effects that apply to both C and C++), this has the same effect as `-fthis-is-variable`. See “Options controlling C dialect” on page 79.

In addition, the following options for optimization, warning, and code generation have meanings only for C++ programs:

**-fno-default-inline**

Do not assume inline for functions defined inside a class scope. See “Options that control optimization” on page 111.

**-Wold-style cast****-Woverloaded-virtual****-Wtemplate-debugging**

Warnings that apply only to C++ programs. See “Options to request or suppress warnings” on page 93.

**-Weffc++**

Warn about violation of some style rules from “*Effective C++*” by Scott Myers.

**+en**

Control how virtual function definitions are used, in a fashion compatible with `'cfront 1.x'`. For more explanation of compatibility conventions for this option, see the description for `'+e'` in “Options for code generation conventions” on page 187.

## Compiling C++ programs

C++ source files conventionally use one of the suffixes `.c`, `.cc`, `.cpp`, or `.cxx`; preprocessed C++ files use the suffix `.ii`. GCC recognizes files with these names and compiles them as C++ programs even if you call the compiler the same way as for compiling C programs (usually with the name `gcc`).

However, C++ programs often require class libraries as well as a compiler that understands the C++ language—and under some circumstances, you might want to compile programs from standard input, or otherwise without a suffix that flags them as C++ programs. `g++` is a program that calls GCC with the default language set to C++, and automatically specifies linking against the GNU class library `libg++`<sup>†</sup>. On many systems, the script `g++` is also installed with the name `c++`.

When you compile C++ programs, you may specify many of the same command-line options that you use for compiling programs in any language; or command-line options meaningful for C and related languages; or options that are meaningful only for C++ programs.

See “Options controlling C++ dialect” on page 85 for explanations of options for languages related to C.

---

<sup>†</sup> Prior to release 2 of the compiler, there was a separate G++ compiler. That version was based on GCC, but not integrated with it. Versions of G++ with a 1.xx version number—for example, G++ version 1.37 or 1.42—are much less reliable than the versions integrated with GCC 2. Moreover, combining G++ 1.xx with a version GCC 2 will simply not work.



# Options to request or suppress warnings

---

Warnings are diagnostic messages reporting constructions which are not inherently erroneous; they may warn of risky constructions or constructions, actually, in error.

You can request many specific warnings with options beginning with `-w`; for instance, use `-Wimplicit` to request warnings on implicit declarations. Each specific warning option also has a negative form beginning `'-Wno-'` to turn off warnings; for instance, `-Wno-implicit`. The following documentation lists only one of the two forms, whichever is not the default.

The following options control the amount and kinds of warnings produced by GNU CC.

`-fsyntax-only`

Check the code for syntax errors, but don't do anything beyond that.

`-pedantic`

Issue all the warnings demanded by strict ANSI standard C; reject all programs that use forbidden extensions.

Valid ANSI standard C programs should compile properly with or without this option (though a rare few will require `-ansi`). However, without this option, certain GNU extensions and traditional C features are supported as well. With this option, they are rejected.

`-pedantic` does not cause warning messages for use of the alternate keywords whose names begin and end with `'__'`. Pedantic warnings are also disabled in the expression that follows `__extension`. However, only system header files should

use these escape routes; application programs should avoid them. See “Alternate keywords” on page 266.

This option is not intended to be *useful*; it exists only to satisfy pedants who would otherwise claim that GNU CC fails to support the ANSI standard.

Some users try to use `-pedantic` to check programs for strict ANSI C conformance. They soon find that it does not do quite what they want: it finds some non-ANSI practices, but not all—only those for which ANSI C *requires* a diagnostic.

A feature to report any failure to conform to ANSI C might be useful in some instances, but would require considerable additional work and would be quite different from `-pedantic`. We recommend, rather, that users take advantage of the extensions of GNU C and disregard the limitations of compilers. Aside from certain supercomputers and obsolete small machines, there is less and less reason ever to use any other C compiler other than for bootstrapping GNU CC.

`-pedantic-errors`

Like `-pedantic`, except that errors are produced rather than warnings.

`-w`

Inhibit all warning messages.

`-Wno-import`

Inhibit warning messages about the use of `#import`.

`-Wchar-subscripts`

Warn if an array subscript has type `char`. This is a common cause of error, as programmers often forget that this type is signed on some machines.

`-Wcomment`

Warn whenever a comment-start sequence `‘/*’` appears in a `‘/*’` comment, or whenever a Backslash-Newline appears in a `‘//’` comment.

`-Wformat`

Check calls to `printf` and `scanf`, etc., to make sure that the arguments supplied have types appropriate to the format string specified.

`-Wimplicit-int`

Warn when a declaration does not specify a type.

`-Wimplicit-function-declarations`

Warn whenever a function is used before being declared.

`-Wimplicit`

Warn whenever a function or parameter is implicitly declared, or when a type implicitly defaults to `int`.



**-Wmain**

Warn if the type of `'main'` is suspicious. `'main'` should be a function with external linkage, returning `int`, taking either zero arguments, two, or three arguments of appropriate types.

**-Wparentheses**

Warn if parentheses are omitted in certain contexts, for instance, when there is an assignment in a context where a *truth value* is expected, or when operators are *nested* whose precedence may be confusing.

Also warn about constructions where there may be confusion to which `if` statement an `else` branch belongs. The following is an example of such a case.

```
{
 if (a)
 if (b)
 foo ();
 else
 bar ();
}
```

In C, every `else` branch belongs to the innermost possible `if` statement, which in this example is `if (b)`. This is often not what the programmer expected, as illustrated in the previous example of indentation that the programmer chose.

When there is the potential for this confusion, GNU C will issue a warning when this flag is specified.

To eliminate the warning, add explicit braces around the innermost `if` statement so there is no way the `else` could belong to the enclosing `if`.

The resulting code would look like the following declaration.

```
{
 if (a)
 {
 if (b)
 foo ();
 else
 bar ();
 }
}
```

**-Wreturn-type**

Warn whenever a function is defined with a return-type that defaults to `int`. Also warn about any `return` statement with no return-value in a function whose return-type is not `void`.

### -Wswitch

Warn whenever a `switch` statement has an index of enumerational type and lacks a `case` for one or more of the named codes of that enumeration. (The presence of a default label prevents this warning.)

`case` labels outside the enumeration range also provoke warnings when this option is used.

### -Wtrigraphs

Warn if any trigraphs are encountered (assuming they are enabled).

### -Wunused

Warn whenever a variable is unused aside from its declaration, whenever a function is declared static but never defined, whenever a label is declared but not used, and whenever a statement computes a result that is explicitly not used.

To suppress this warning for an expression, simply cast it to `void`. For unused variables and parameters, use the `unused` attribute (for explanation of specifications for this attribute, see “Specifying attributes of variables” on page 235).

### -Wuninitialized

An automatic variable is used without first being initialized.

These warnings are possible only in optimizing compilation, because they require data flow information that is computed only when optimizing. If you don’t specify ‘-O’, you simply won’t get these warnings.

These warnings occur only for variables that are candidates for register allocation.

Therefore, they do not occur for a variable that is declared `volatile`, or whose address is taken, or whose size is other than 1, 2, 4 or 8 bytes. Also, they do not occur for structures, unions or arrays, even when they are in registers.

**NOTE:** There may be no warning about a variable that is used only to compute a value that itself is never used.

Such computations may be deleted by data flow analysis before the warnings display.

These warnings are made optional because GNU CC is not smart enough to see all the reasons why the code might be correct despite appearing to have an error. The following is one example of how this occurs.

```
{
 int x;
 switch (y)
 {
 case 1: x = 1;
 break;
 case 2: x = 4;
```

```

 break;
 case 3: x = 5;
 }
 foo (x);
}

```

If the value of ‘y’ is always 1, 2 or 3, then ‘x’ is always initialized, but GNU CC doesn’t make this determination. The following is another common case.

```

{
 int save_y; if (change_y) save_y = y, y = new_y;
 ...
 if (change_y) y = save_y;
}

```

This has no bug because `save_y` is used only if it is set. Some spurious warnings can be avoided if you declare all the functions you use that never return as `noreturn`. See “Declaring attributes of functions” on page 226.

#### -Wreorder

For C++ only. Warn when the order of member initializers given in the code does not match the order in which they must be executed. For instance, in the following example, the compiler will warn that the member initializers for `i` and `j` will be rearranged to match the declaration order of the members.

```

struct A {
 int i;
 int j;
 A(): j (0), i (1) { }
};

```

#### -Wtemplate-debugging

When using templates in a C++ program, warn if debugging is not yet fully available (C++ only).

#### -Wunknown-pragmas

Warn when a `#pragma` directive is encountered which is not understood by GCC. If this command line option is used, warnings will even be issued for unknown pragmas in system header files. This is not the case if the warnings were only enabled by the `-Wall` command line option.

#### -Wall

All of the previous `-w` options combined. These are all the options which pertain to usage that we recommend avoiding and that we believe is easy to avoid, even in conjunction with macros.

The remaining ‘-w...’ options are not implied by `-Wall` because they warn about constructions that we consider reasonable to use, on occasion, in clean programs. Some of them warn about constructions that users generally do not consider questionable, but which occasionally you might wish to check for; others warn about constructions that are necessary or hard to avoid in some cases, and there is no simple

way to modify the code to suppress the warning.

-W

Print extra warning messages for the following events.

- A nonvolatile automatic variable might be changed by a call to `longjmp`. These warnings as well are possible only in optimizing compilation. The compiler sees only the calls to `setjmp`. It cannot know where `longjmp` will be called; in fact, a signal handler could call it at any point in the code. As a result, you may get a warning even when there is in fact no problem because `longjmp` cannot in fact be called at the place which would cause a problem.
- A function can return either with or without a value. (Falling off the end of the function body is considered returning without a value.) For instance, the following example shows how a function would evoke such a warning.

```
foo (a)
{
 if (a > 0)
 return a;
}
```

- An expression-statement or the left-hand side of a comma expression contains no side effects. To suppress the warning, cast the unused expression to `void`. For example, an expression such as `x[i, j]` will cause a warning, but one such as `x[(void)i, j]` will not.
- An unsigned value is compared against zero with '`<`' or '`<=`'.
- A comparison like `x<=y<=z` appears; this is equivalent to `(x<=y ? 1 : 0) <= z`, which is a different interpretation from that of ordinary mathematical notation.
- Storage-class specifiers like `static` are not the first things in a declaration. According to the C Standard, this usage is obsolescent.
- If `-Wall` or `-Wunused` is also specified, warn about unused arguments.
- An aggregate has a partly bracketed initializer. For example, the following code would evoke such a warning, because braces are missing around the initializer for `x.h`.

```
struct s { int f, g; };
struct t { struct s h; int i; };
struct t x = { 1, 2, 3 };
```

-Wtraditional

Warn about certain constructs that behave differently in traditional and ANSI C.

- Macro arguments occurring within string constants in the macro body. These would substitute the argument in traditional C, but are part of the constant in ANSI C.

- A function declared external in one block and then used after the end of the block.
- A `switch` statement that has an operand of type `long`.

**-Wundef**

Warn if an undefined identifier is evaluated in an `#if` directive.

**-Wshadow**

Warn whenever a local variable shadows another local variable.

**-Wid-clash-len**

Warn whenever two distinct identifiers match in the first *len* characters. This may help you prepare a program that will compile with certain obsolete, brain-damaged compilers.

**-Wlarger-than-len**

Warn whenever an object of larger than *len* bytes is defined.

**-Wpointer-arith**

Warn about anything that depends on the “*size of*” a function type or of `void`. GNU C assigns these types a size of 1, for convenience in calculations with `void*` pointers and pointers to functions.

**-Wbad-function-cast**

Warn whenever a function call is cast to a non-matching type. For example, warn if `int malloc()` is cast to `anything*`.

**-Wcast-qual**

Warn whenever a pointer is cast so as to remove a type qualifier from the target type. For example, warn if a `const char*` is cast to an ordinary `char*`.

**-Wcast-align**

Warn whenever a pointer is cast such that the required alignment of the target is increased. For example, warn if a `char*` is cast to an `int*` on machines where integers can only be accessed at two- or four-byte boundaries.

**-Wwrite-strings**

Give string constants the type `const char[length]` so that copying the address of one into a non-`const char*` pointer will get a warning.

These warnings will help you find code, at compile time, that can try to write into a string constant, but only if you have been very careful about using `const` in declarations and prototypes.

Otherwise, it will just be a nuisance; this is why `-Wall` doesn't request these warnings.

**-Wconversion**

Warn if a prototype causes a type conversion that is different from what would happen to the same argument in the absence of a prototype.

This includes conversions of fixed point to floating and vice versa, and

conversions changing the width or signedness of a fixed point argument except when the same as the default promotion.

Also, warn if a negative integer constant expression is implicitly converted to an unsigned type. For example, warn about the assignment `x = -1` if `x` is unsigned. But do not warn about explicit casts like `(unsigned) -1`.

**-Wsign-compare**

Warn when a comparison between signed and unsigned values could produce an incorrect result when the signed value is converted to unsigned. This warning is also enabled by `-w`; to get the other warnings of `-w` without this warning, use `-Wno-sign-compare`.

**-Waggregate-return**

Warn if any functions that return structures or unions are defined or called. (In languages where you can return an array, this also elicits a warning.)

**-Wstrict-prototypes**

Warn if a function is declared or defined without specifying the argument types. (An old-style function definition is permitted without a warning if preceded by a declaration which specifies the argument types.)

**-Wmissing-prototypes**

Warn if a global function is defined without a previous prototype declaration. This warning is issued even if the definition itself provides a prototype. The aim is to detect global functions that fail to be declared in header files.

**-Wmissing-declarations**

Warn if a global function is defined without a previous declaration. Do so even if the definition itself provides a prototype. Use this option to detect global functions that are not declared in header files.

**-Wredundant-decls**

Warn if anything is declared more than once in the same scope, even in cases where multiple declaration is valid and changes nothing.

**-Wnested-externs**

Warn if an extern declaration is encountered within a function.

**-Winline**

Warn if a function can not be inlined, and either it was declared as inline, or else the `-finline-functions` option was given.

**-Wold-style-cast**

Warn if an old-style (C-style) cast is used within a program.

**-Woverloaded-virtual**

For C++ only, warns when a derived class function declaration may be an error in defining a virtual function.

In a derived class, the definitions of virtual functions must match the type

signature of a virtual function declared in the base class.

With this option, the compiler warns when you define a function with the same name as a virtual function, but with a type signature that does not match any declarations from the base class.

`-Wsynth`

For C++ only, warns when G++'s synthesis behavior does not match that of `cfront`. For instance, see the following declaration:

```
struct A {
 operator int ();
 A& operator = (int);
};
main ()
{
 A a,b;
 a = b;
}
```

In this previous example, G++ will synthesize a default:

```
A& operator = (const A&);,
```

`cfront` will use the user-defined expression, `operator=`.

`-Werror`

Make all warnings into errors.





# Options for debugging

---

The following documentation discusses some special options for debugging either your program or the GNU compiler, GCC.

`-g`

Produce debugging information in the operating system's native format (stabs, COFF, XCOFF, or DWARF). GDB can work with this debugging information.

On most systems that use *stabs* format, '`-g`' enables use of extra debugging information that only GDB can use; this extra information makes debugging work better in GDB but will probably make other debuggers crash or refuse to read the program.

If you want to control for certain whether to generate the extra information, use `-gstabs+`, `-gstabs`, `-gxcoff+`, `-gxcoff`, `-gdwarf+1`, or `-gdwarf-1` (see the explanations for each option in the following discussions).

Unlike most other C compilers, GCC allows you to use '`-g`' with '`-o`'. The shortcuts taken by optimized code may occasionally produce surprising results: some variables you declared may not exist at all; flow of control may briefly move where you did not expect it; some statements may not be executed because they compute constant results or their values were already at hand; some statements may execute in different places because they were moved out of loops.

Nevertheless it proves possible to debug optimized output. This makes it reasonable to use the optimizer for programs that might have bugs.

The following options are useful when GCC is generated with the capability for more than one debugging format.

`-ggdb`

Produce debugging information in the native format (if that is supported), including GDB extensions if at all possible.

`-gstabs`

Produce debugging information in `stabs` format (if that is supported), without GDB extensions. This is the format used by DBX on most BSD systems. On MIPS, Alpha and System V Release 4 systems this option produces `stabs` debugging output which is not understood by DBX or SDB. On System V Release 4 systems this option requires the GNU assembler.

`-gstabs+`

Produce debugging information in `stabs` format (if that is supported), using GNU extensions understood only by the GNU debugger (GDB). The use of these extensions is likely to make other debuggers crash or refuse to read the program.

`-gcoff`

Produce debugging information in COFF format (if that is supported). This is the format used by SDB on most System V systems prior to System V Release 4.

`-gxcoff`

Produce debugging information in XCOFF format (if that is supported). This is the format used by the DBX debugger on IBM RS/6000 systems.

`-gxcoff+`

Produce debugging information in XCOFF format (if that is supported), using GNU extensions understood only by the GNU debugger (GDB). The use of these extensions is likely to make other debuggers crash or refuse to read the program, and may cause assemblers other than the GNU assembler (GAS) to fail with an error.

`-gdwarf`

Produce debugging information in DWARF version 1 format (if that is supported). This is the format used by SDB on most System V Release 4 systems.

`-gdwarf+`

Produce debugging information in DWARF version 1 format (if that is supported), using GNU extensions understood only by the GNU debugger (GDB). The use of these extensions is likely to make other debuggers crash or refuse to read the program.

`-gdwarf-2`

Produce debugging information in DWARF version 2 format (if that is supported). This is the format used by DBX on IRIX 6.

```
-glevel
-ggdblevel
-gstabslevel
-gcofflevel
-gxcofflevel
-gdwarflevel
-gdwarf-2level
```

Request debugging information and also use *level* to specify how much information. The default level is 2.

Level 1 produces minimal information, enough for making backtraces in parts of the program that you don't plan to debug. This includes descriptions of functions and external variables, but no information about local variables and no line numbers. Level 3 includes extra information, such as all the macro definitions present in the program. Some debuggers support macro expansion when you use the '-g3' option.

```
-p
```

Generate extra code to write profile information suitable for the analysis program *prof*. You must use this option when compiling the source files you want data about, and you must also use it when linking.

```
-pg
```

Generate extra code to write profile information suitable for the analysis program *gprof*. You must use this option when compiling the source files you want data about, and you must also use it when linking.

```
-a
```

Generate extra code to write profile information for basic blocks, which will record the number of times each basic block is executed, the basic block start address, and the function name containing the basic block. If '-g' is used, the line number and filename of the start of the basic block will also be recorded. If not overridden by the machine description, the default action is to append to the text file 'bb.out'.

This data could be analyzed by a program like *tcov*. However, the format of the data is not what *tcov* expects. Eventually GNU *gprof* should be extended to process this data.

```
-Q
```

Makes the compiler print out each function name as it is compiled, and print some statistics about each pass when it finishes.

```
-ax
```

Generate extra code to profile basic blocks. Your executable will produce output that is a superset of that produced when '-a' is used. Additional output is the source and target address of the basic blocks where a jump takes place, the

number of times a jump is executed, and (optionally) the complete sequence of basic blocks being executed. The output is appended to file 'bb.out'.

You can examine different profiling aspects without recompilation. Your executable will read a list of function names from file 'bb.in'. Profiling starts when a function on the list is entered and stops when that invocation is exited. To exclude a function from profiling, use a hyphen (-), to prefix its name.

If a function name is not unique, specify its location.

For example, use the following example's declaration (where *functionname* designates the unique function to specify for the location).

```
/path/file-name.d: functionname
```

Your executable will write the available paths and filenames in file 'bb.out'.

The following function names have a special meaning.

- `__bb_jumps__`  
Write source, target and frequency of jumps to file 'bb.out'.
- `__bb_hidecall__`  
Exclude function calls from frequency count.
- `__bb_showret__`  
Include function returns in frequency count.
- `__bb_trace__`  
Write the sequence of basic blocks to the 'bbtrace.gz' file. The file will be compressed using the program `gzip`, which must exist in your `PATH`. On systems *without* the `popen` function, the file will be named 'bbtrace' and will not be compressed.

**WARNING!** *Profiling for even a few seconds on these systems will produce a very large file.*

**NOTE:** `__bb_hidecall__` and `__bb_showret__` will not affect the sequence written to 'bbtrace.gz'.

What follows is a short example using different profiling parameters in file 'bb.in'. Assume function `foo` consists of basic blocks 1 and 2 and is called twice from block 3 of function `main`. After the calls, block 3 transfers control to block 4 of `main`. With `__bb_trace__` and `main` contained in file 'bb.in', the following sequence of blocks is written to file 'bbtrace.gz': 0 3 1 2 1 2 4. The return from block 2 to block 3 is not shown, because the return is to a point inside the block and not to the top. The block address 0 always indicates that control is transferred to the trace from somewhere outside the observed functions. With `-foo` added to 'bb.in', the blocks of function `foo` are removed from the trace, so only 0 3 4 remains. With `__bb_jumps__` and `main` contained in file 'bb.in', jump frequencies will be written to file 'b.out'. The frequencies are obtained by

constructing a trace of blocks and incrementing a counter for every neighboring pair of blocks in the trace. The trace 0 3 1 2 1 2 4 displays the following frequencies.

```
Jump from block 0x0 to block 0x3 executed 1 time(s)
Jump from block 0x3 to block 0x1 executed 1 time(s)
Jump from block 0x1 to block 0x2 executed 2 time(s)
Jump from block 0x2 to block 0x1 executed 1 time(s)
Jump from block 0x2 to block 0x4 executed 1 time(s)
```

With `__bb_hidecall__`, due to call instructions, control transfer is removed from the trace; that is, the trace is cut into three parts: 0 3 4, 0 1 2, and 0 1 2. With `__bb_showret__`, control transfer is added to the trace.

The trace becomes: 0 3 1 2 3 1 2 3 4.

**NOTE:** The previous trace is not the same as the sequence written to ‘`bbtrace.gz`’. It is solely used for counting jump frequencies.

#### `-fprofile-arcs`

Instrument *arcs* during compilation. For each function of your program, GNU CC creates a program flow graph, then finds a spanning tree for the graph. Only arcs that are not on the spanning tree have to be instrumented: the compiler adds code to count the number of times that these arcs are executed. When an arc is the only exit or only entrance to a block, the instrumentation code can be added to the block; otherwise, a new basic block must be created to hold the instrumentation code.

Since not every arc in the program must be instrumented, programs compiled with this option run faster than programs compiled with ‘`-a`’, which adds instrumentation code to every basic block in the program. The tradeoff: since `gcov` does not have execution counts for all branches, it must start with the execution counts for the instrumented branches, and then iterate over the program flow graph until the entire graph has been solved. Hence, `gcov` runs a little more slowly than a program which uses information from ‘`-a`’.

`-fprofile-arcs` also makes it possible to estimate branch probabilities, and to calculate basic block execution counts. In general, basic block execution counts do not give enough information to estimate all branch probabilities. When the compiled program exits, it saves the arc execution counts to a file called *sourcename.da*. Use the compiler option, `-fbranch-probabilities`, when recompiling, to optimize using estimated branch probabilities (see “Options that control optimization” on page 111).

#### `-ftest-coverage`

Create data files for the `gcov` code-coverage utility (see “`gcov`: a test coverage

program” on page 285). The data filenames begin with the name of your source file:

- *sourcename.bb*  
A mapping from basic blocks to line numbers, which `gcov` uses to associate basic block execution counts with line numbers.
- *sourcename.bbg*  
A list of all arcs in the program flow graph. This allows `gcov` to reconstruct the program flow graph, so that it can compute all basic block and arc execution counts from the information in the *sourcename.da* file (this last file is the output from `-fprofile-arcs`).

`-dletters`

This option says to make debugging dumps during compilation at times specified by *letters*. This is used for debugging the compiler. The filenames for most of the dumps are made by appending a word to the source filename (e.g., *file.c.rtl* or *file.c.jump*). What follows are the possible letters for use in *letters*, and their meanings for use in a filename, *file*, with their appropriate extensions.

- a  
Produce all the dumps previously listed.
- A  
Annotate the assembler output with miscellaneous debugging information.
- b  
Dump after computing branch probabilities, to *file.bp*.
- c  
Dump after instruction combination, to *file.combine*.
- d  
Dump after delayed branch scheduling, to *file.dbr*.
- D  
Dump after purging `ADDRESSOF`, to *file.addressof*.
- f  
Dump after flow analysis, to *file.flow*.
- g  
Dump after global register allocation, to *file.greg*.
- G  
Dump after GCSE (Global Common Sub-expression Elimination), to *file.gcse*.
- j  
Dump after first jump optimization, to *file.jump*.

- J Dump after last jump optimization, to *file.jump2*.
  - k Dump after conversion from registers to stack, to *file.stack*.
  - l Dump after local register allocation, to *file.lreg*.
  - L Dump after loop optimization, to *file.loop*.
  - m Print statistics on memory usage, at the end of the run, to standard error.
  - M Dump after performing the machine dependent reorganization pass, to *file.mach*.
  - N Dump after the register move pass, to *file.regmove*.
  - P Annotate the assembler output with a comment indicating which pattern and alternative was used.
  - r Dump after RTL generation, to *file.rtl*.
  - R Dump after the second instruction scheduling pass, to *file.sched2*.
  - s Dump after CSE (including the jump optimization that sometimes follows CSE, the Common Sub-expression Elimination), to *file.cse*.
  - S Dump after the first instruction scheduling pass, to *file.sched*.
  - t Dump after the second CSE pass (including the jump optimization that sometimes follows CSE), to *file.cse2*.
  - x Just generate RTL for a function instead of compiling it. Usually used with 'r'.
  - Y Dump debugging information during parsing, to standard error.
- fpretend-float**  
 When running a cross-compiler, pretend that the target machine uses the same floating point format as the host machine. This causes incorrect output of the actual floating constants, but the actual instruction sequence will probably be the same as GCC would make when running on the target machine.

### `-save-temps`

Store the usual *temporary* intermediate files permanently; place them in the current directory and name them based on the source file. Thus, compiling `'foo.c'` with `-c-save-temps` would produce files `'foo.i'` and `'foo.s'`, as well as `'foo.o'`.

### `-print-file-name=library`

Print the full absolute name of the library file, *library*, that would be used when linking—and don't do anything else. With this option, GCC does not compile or link anything; it just prints the filename.

### `-print-prog-name=program`

Like `-print-file-name`, but searches for a program such as `'cpp'`.

### `-print-libgcc-file-name`

Same as `-print-file-name=libgcc.a`.

This is useful when you use `-nostdlib` or `-nodefaultlibs` but you do want to link with `'libgcc.a'`. You can use the following example's command, where `files` signifies the files that need to link.

```
gcc -nostdlib files
...
gcc -print-libgcc-file-name
```

### `-print-search-dirs`

Print the name of the configured installation directory and a list of program and library directories `gcc` will search—and don't do anything else.

This is useful when `gcc` prints the following error message:

**installation problem, cannot exec cpp: No such file or directory**

To resolve this you either need to put `'cpp'` and the other compiler components where GCC expects to find them, or set the environment variable, `GCC_EXEC_PREFIX`, to the directory where you installed them. Don't forget the trailing slash; see "Environment variables affecting GCC" on page 195.



# Options that control optimization

---

The following options control various sorts of optimizations:

`-O`

`-O1`

Optimize. Optimizing compilation takes somewhat more time, and a lot more memory for a large function.

Without ‘`-O`’, the compiler’s goal is to reduce the cost of compilation and to make debugging produce the expected results.

Statements are independent: if you stop the program with a breakpoint between statements, you can then assign a new value to any variable or change the program counter to any other statement in the function and get exactly the results you would expect from the source code.

Without ‘`-O`’, the compiler only allocates variables declared `register` in registers. The resulting compiled code is a little worse than produced by PCC without ‘`-O`’.

With ‘`-O`’, the compiler tries to reduce code size and execution time.

When you specify ‘`-O`’, the compiler turns on `-fthread-jumps` and `-fdefer-pop` on all machines.

The compiler turns on `-fdelayed-branch` on machines that have delay slots, and `-fomit-frame-pointer` on machines that can support debugging even without a frame pointer. On some machines the compiler also turns on other flags.

`-O2`

Optimize even more. GNU CC performs nearly all supported optimizations that do not involve a space-speed tradeoff. The compiler does not perform loop unrolling or function inlining when you specify `'-O2'`. As compared to `'-O'`, this option increases both compilation time and the performance of the generated code.

`'-O2'` turns on all optional optimizations except for loop unrolling function inlining, life shortening, and static variable optimizations. It also turns on frame pointer elimination on machines where doing so does not interfere with debugging.

`-O3`

Optimize yet more. `'-O3'` turns on all optimizations specified by `'-O2'` and also turns on the option, `inline-functions`.

`-O0`

Do not optimize. If you use multiple `'-O'` options, with or without level numbers, the last such option is the one that is effective.

`-Os`

Optimize for size. `-Os` enables all `-O2` optimizations that do not typically increase code size. It also performs further optimizations designed to reduce code size.

If you use multiple `-O` options, with or without level numbers, the last such option is the one that is effective.

Options of the form, `-fflag`, specify machine-independent flags. Most flags have both positive and negative forms; the negative form of `-ffoo` would be `-fno-foo`. In the following options, only one of the forms is listed—the one which is not the default.

You can figure out the other form by either removing `'no-'` or adding it.

`-ffloat-store`

Do not store floating point variables in registers, and inhibit other options that might change whether a floating point value is taken from a register or memory.

This option prevents undesirable excess precision on machines such as the 68000 where the floating registers (of the 68881) keep more precision than a `double` is supposed to have. For most programs, the excess precision does only good, but a few programs rely on the precise definition of IEEE floating point. Use

`-ffloat-store` for such programs.

`-fno-default-inline`

Do not make member functions inline by default merely because they are defined inside the class scope (C++ only). Otherwise, when you specify `'-O'`, member functions defined inside class scope are compiled inline by default; i.e., you don't need to add `inline` in front of the member function name.

**-fno-defer-pop**

Always pop the arguments to each function call as soon as that function returns. For machines which must pop arguments after a function call, the compiler normally lets arguments accumulate on the stack for several function calls and pops them all at once.

**-fforce-mem**

Force memory operands to be copied into registers before doing arithmetic on them. This produces better code by making all memory references potential common subexpressions. When they are not common subexpressions, instruction combination should eliminate the separate register-load. The ‘-O2’ option turns on this option.

**-fforce-addr**

Force memory address constants to be copied into registers before doing arithmetic on them.

This may produce better code just as `-fforce-mem` may.

**-fomit-frame-pointer**

Don’t keep the frame pointer in a register for functions that don’t need one. This avoids the instructions to save, set up and restore frame pointers; it also makes an extra register available in many functions.

**WARNING:** It also makes debugging impossible on some machines.

On some machines, such as the VAX, this flag has no effect because the standard calling sequence automatically handles the frame pointer and nothing is saved by pretending it doesn’t exist. The machine-description macro, `FRAME_POINTER_REQUIRED`, controls whether a target machine supports this flag. See “Constraints for particular machines” on page 253 to determine *register usage* with your target machine.

**-fno-inline**

Don’t pay attention to the `inline` keyword. Normally this option is used to keep the compiler from expanding any functions inline.

**NOTE:** If you are not optimizing, no functions can be expanded inline.

**-finline-functions**

Integrate all simple functions into their callers. The compiler heuristically decides which functions are simple enough to be worth integrating in this way.

If all calls to a given function are integrated, and the function is declared `static`, then the function is normally not output as assembler code in its own right.

**-fkeep-inline-functions**

Even if all calls to a given function are integrated, and the function is declared `static`, output a separate run-time callable version of the function. This switch

does not affect `extern inline` functions.

### `-fkeep-static-consts`

Emit variables declared `static const` when optimization isn't turned on, even if the variables weren't referenced. This option is enabled by default.

`-fno-keep-static-consts` will force the compiler to check if the variable was referenced, regardless of whether or not optimization is turned on.

### `-fno-function-cse`

Do not put function addresses in registers; make each instruction that calls a constant function contain the function's address explicitly.

The `fno-function-cse` option results in less efficient code, but some strange hacks that alter the assembler output may be confused by the optimizations performed when this option is not used.

### `-ffast-math`

This option allows GCC to violate some ANSI or IEEE rules and/or specifications in the interest of optimizing code for speed. For example, it allows the compiler to assume arguments to the `sqrt` function are non-negative numbers and that no floating-point values are NaNs.

This option should never be turned on by any `'-O'` option since it can result in incorrect output for programs that depend on an exact implementation of IEEE or ANSI rules/specifications for math functions.

The following options control specific optimizations. The `'-O2'` option turns on all of these optimizations except `-funroll-loops` and `-funroll-all-loops`. On most machines, the `'-O'` option turns on the `-fthread-jumps` and `-fdelayed-branch` options, but specific machines may handle it differently. Use the following flags in the rare cases when you want to *fine-tune* optimizations.

### `-fstrength-reduce`

Perform the optimizations of loop strength reduction and elimination of iteration variables.

### `-fthread-jumps`

Perform optimizations where we check to see if a jump branches to a location where another comparison subsumed by the first is found. If so, the first branch is redirected to either the destination of the second branch or a point immediately following it, depending on whether the condition is known to be true or false.

### `-fcse-follow-jumps`

In common subexpression elimination, scan through jump instructions when the target of the jump is not reached by any other path. For example, when CSE encounters an `if` statement with an `else` clause, CSE will follow the jump when the condition tested is false.

### `-fcse-skip-blocks`

This is similar to `'-fcse-follow-jumps'`, but causes CSE to follow jumps which

conditionally skip over blocks. When CSE encounters a simple `if` statement with no `else` clause, `-fcse-skip-blocks` causes CSE to follow the jump around the body of the `if`.

`-frerun-cse-after-loop`

Re-run common sub-expression elimination after loop optimizations has been performed.

`-frerun-loop-opt`

Run the loop optimizer twice.

`-fgcse`

Perform a global common sub-expression elimination pass. This pass also performs global constant and copy propagation.

`-fexpensive-optimizations`

Perform a number of minor optimizations that are relatively expensive.

`-fdelayed-branch`

If supported for the target machine, attempt to reorder instructions to exploit instruction slots available after delayed branch instructions.

`-fschedule-insns`

If supported for the target machine, attempt to reorder instructions to eliminate execution stalls due to required data being unavailable. This helps machines that have slow floating point or memory load instructions by allowing other instructions to be issued until the result of the load or floating point instruction is required.

`-fschedule-insns2`

Similar to `-fschedule-insns`, but requests an additional pass of instruction scheduling after register allocation has been done. This is especially useful on machines with a relatively small number of registers and where memory load instructions take more than one cycle.

`-ffunction-sections`

Place each function into its own section in the output file if the target supports arbitrary sections. The function's name determines the section's name in the output file.

Use this option on systems where the linker can perform optimizations to improve locality of reference in the instruction space. HPPA processors running HP-UX and SPARC processors running Solaris 2 have linkers with such optimizations. Other systems using the ELF object format as well as AIX may have these optimizations in the future.

Only use this option when there are significant benefits from doing so. When you specify this option, the assembler and linker will create larger object and executable files and will also be slower. You will not be able to use `gprof` on all

systems if you specify this option and you may have problems with debugging if you specify both this option and ‘-g’.

**-fcaller-saves**

Enable values to be allocated in registers that will be clobbered by function calls, by emitting extra instructions to save and restore the registers around such calls. Such allocation is done only when it seems to result in better code than would otherwise be produced. This option is enabled by default on certain machines, usually those which have no call-preserved registers to use instead.

**-funroll-loops**

Perform the optimization of loop unrolling. This is only done for loops whose number of iterations can be determined at compile time or run time.

-funroll-loop implies both -fstrength-reduce and

-frerun-cse-after-loop.

**-funroll-all-loops**

Perform the optimization of loop unrolling. This is done for all loops and usually makes programs run more slowly. -funroll-all-loops implies

-fstrength-reduce as well as -frerun-cse-after-loop.

**-freduce-all-givs**

Forces all general-induction variables in loops to be strength-reduced.

**NOTE:** When compiling programs written in Fortran, -fmove-all-moveables and -freduce-all-givs are enabled by default when you use the optimizer.

These options may generate better or worse code; results are highly dependent on the structure of loops within the source code. These two options are intended to be removed someday, once they have helped determine the efficacy of various approaches to improving loop optimizations.

**IMPORTANT:** Please let us know how use of these options affects the performance of your production code. We’re very interested in code that runs slower when these options are enabled (email: [egcs@cygnus.com](mailto:egcs@cygnus.com) and [fortran@gnu.org](mailto:fortran@gnu.org)).

**-fno-peephole**

Disable any machine-specific peephole optimizations.

**-fbranch-probabilities**

After running a program compiled with -fprofile-arcs (see “Options for debugging” on page 103), you can compile it a second time using

-fbranch-probabilities, to improve optimizations based on guessing the path a branch might take.

**-fregmove**

Some machines only support 2 operands per instruction. On such machines, GNU CC might have to do extra copies. The ‘-fregmove’ option overrides the default for the machine to do the copy before register allocation.

# Options controlling the preprocessor

---

The following options control the C preprocessor, run on each C source file before actual compilation.

If you use the `-E` option, nothing is done except preprocessing. Some of these options make sense only together with `-E` because they cause the preprocessor output to be unsuitable for actual compilation.

`-include file`

Process *file* as input before processing the regular input file. In effect, the contents of *file* are compiled first. Any `-D` and `-U` options on the command line are always processed before `-include file`, regardless of the order in which they are written. All the `-include` and `-imacros` options are processed in the order in which they are written.

`-imacros file`

Process *file* as input, discarding the resulting output, before processing the regular input file. Because the output generated from *file* is discarded, the only effect of `-imacros file` is to make the macros defined in *file* available for use in the main input.

Any `-D` and `-U` options on the command line are always processed before the `-imacros file`, regardless of the order in which they are written. All the `-include` and `-imacros` options are processed in the order in which they are written.

- `-idirafter dir`

Add the directory *dir* to the second include path. The directories on the second include path are searched when a header file is not found in any of the directories in the main include path (the one that `-I` adds to).
- `-iprefix prefix`

Specify *prefix* as the prefix for subsequent `-iwithprefix` options.
- `-iwithprefix dir`

Add a directory to the second include path. The directory's name is made by concatenating *prefix* and *dir*, where *prefix* was specified previously with `-iprefix`. If you have not specified a prefix yet, the directory containing the installed passes of the compiler is used as the default.
- `-iwithprefixbefore dir`

Add a directory to the main include path. The directory's name is made by concatenating *prefix* and *dir*, as in the case of `-iwithprefix`.
- `-isystem dir`

Add a directory to the beginning of the second include path, marking it as a system directory, so that it gets the same special treatment as is applied to the standard system directories.
- `-nostdinc`

Do not search the standard system directories for header files. Only the directories you have specified with `-I` options (and the current directory, if appropriate) are searched. See “Options for directory search” on page 127 for information on `-I`.  
By using both `-nostdinc` and `-I-`, you can limit the include-file search path to only those directories you specify explicitly.
- `-undef`

Do not predefine any nonstandard macros, including architecture flags.
- `-E`

Run only the C preprocessor. Preprocess all the C source files specified and output the results to standard output or to the specified output file.
- `-C`

Tell the preprocessor not to discard comments. Used with the `-E` option.
- `-P`

Tell the preprocessor not to generate `#line` directives. Used with the `-E` option.
- `-M`

Tell the preprocessor to output a rule suitable for `make` describing the dependencies of each object file. For each source file, the preprocessor outputs one `make`-rule whose target is the object filename for that source file and whose dependencies are all the `#include` header files it uses.  
This rule may be a single line or may be continued with ``'-newline` if it is long.



The list of rules is printed on standard output instead of the preprocessed C program. `-M` implies `-E`.

Another way to specify output of a `make` rule is by setting the environment variable, `DEPENDENCIES_OUTPUT` (see “Environment variables affecting GCC” on page 195).

`-MM`

Like `-M` but the output mentions only the user header files included with `'include <file>'`.

System header files included with `'#include <file>'` are omitted.

`-MD`

Like `-M` but the dependency information is written to a file made by replacing `.c` with `.d` at the end of the input filenames. This is in addition to compiling the file as specified. `'-MD'` does not inhibit ordinary compilation the way `'-M'` does.

In Mach, you can use the utility `'md'` to merge multiple dependency files into a single dependency file suitable for using with the `make` command.

`-MMD`

Like `'-MD'` except output mentions only user header files, not system header files.

`-MG`

Treat missing header files as generated files and assume they live in the same directory as the source file. If you specify `'-MG'`, you must also specify either `'-M'` or `'-MM'`. `'-MG'` is not supported with `'-MD'` or `'-MMD'`.

`-H`

Print the name of each header file used, in addition to other normal activities.

`-Aquestion(answer)`

Assert the answer *answer* for *question*, in case it is tested with a preprocessing conditional such as `#if #question(answer)`. `'-A-'` disables the standard assertions that normally describe the target machine.

`-Dmacro`

Define macro *macro* with the string `'1'` as its definition.

`-Dmacro=defn`

Define macro *macro* as *defn*. All instances of `'-D'` on the command line are processed before any `'-U'` options.

`-Umacro`

Undefine macro *macro*. The `'-U'` options are evaluated after all `'-D'` options, but before any `-include` and `-imacros` options.

`-dM`

Tell the preprocessor to output only a list of the macro definitions that are in effect at the end of preprocessing. Used with the `'-E'` option.

`-dD`

Tell the preprocessing to pass all macro definitions into the output, in their proper sequence in the rest of the output.

`-dN`

Like ‘`-dD`’ except that the macro arguments and contents are omitted. Only `#define name` is included in the output.

`-trigraphs`

Support ANSI C trigraphs. The `-ansi` option also has this effect.

`-Wp, option`

Pass *option* as an option to the preprocessor. If *option* contains commas, it is split into multiple options at the commas.

# Passing options to the assembler

---

You can pass options to the assembler.

The following option is the only one regularly in use by the assembler.

`-Wa, option`

Pass *option* as an option to the assembler. If *option* contains commas, it is split into multiple options at the commas.



## Options for linking

---

The following options come into play when the compiler links object files into an executable output file. They are meaningless if the compiler is not doing a link step.

*object-file-name*

A filename that does not end in a special recognized suffix is considered to name an object file or library. (Object files are distinguished from libraries by the linker according to the file contents.) If linking is done, these object files are used as input to the linker.

-C  
-S  
-E

If any of these options is used, then the linker is not run, and object filenames should not be used as arguments. See “Options controlling the kind of output” on page 75.

*-llibrary*

Search the library named *library* when linking.

It makes a difference where in the command you write this option; the linker searches processes libraries and object files in the order they are specified. Thus, ‘foo.o -lz bar.o’ searches library ‘z’ after file ‘foo.o’ but before ‘bar.o’. If ‘bar.o’ refers to functions in ‘z’, those functions may not be loaded.

The linker searches a standard list of directories for the library, which is actually a file named `liblibrary.a`. The linker then uses this file as if it had been specified

precisely by name.

The directories searched include several standard system directories plus any that you specify with ‘-L’.

Normally the files found this way are library files—archive files whose members are object files. The linker handles an archive file by scanning through it for members which define symbols that have so far been referenced but not defined. But if the file that is found is an ordinary object file, it is linked in the usual fashion.

The only difference between using an ‘-l’ option and specifying a filename is that ‘-l’ surrounds library with ‘lib’ and ‘.a’ and searches several directories.

**-lobjc**

You need this special case of the ‘-l’ option in order to link an Objective C program.

**-nostartfiles**

Do not use the standard system startup files when linking. The standard system libraries are used normally, unless **-nostdlib** or **-nodefaultlibs** is used.

**-nodefaultlibs**

Do not use the standard system libraries when linking. Only the libraries you specify will be passed to the linker. The standard startup files are used normally, unless ‘-nostartfiles’ is used.

**-nostdlib**

Do not use the standard system startup files or libraries when linking. No startup files and only the libraries you specify will be passed to the linker.

One of the standard libraries bypassed by ‘-nostdlib’ and ‘-nodefaultlibs’ is ‘libgcc.a’, a library of internal subroutines that GNU CC uses to overcome shortcomings of particular machines, or special needs for some languages. (See “Building and installing a cross-compiler” on page 48 and “libgcc.a and cross-compilers” on page 50 for more discussion of libgcc.a.)

In most cases, you need libgcc.a even when you want to avoid other standard libraries. In other words, when you specify **-nostdlib** or ‘-nodefaultlibs’ you should usually specify ‘-lgcc’ as well. This ensures that you have no unresolved references to internal GNU CC library subroutines. (For example, `__main`, used to ensure C++ constructors will be called; see “collect2 and cross-compiling” on page 54.)

**-s**

Remove all symbol table and relocation information from the executable.

**-static**

On systems that support dynamic linking, this prevents linking with the shared libraries. On other systems, this option has no effect.

- shared  
Produce a shared object which can then be linked with other objects to form an executable. Not all systems support this option. You must also specify `-fpic` or `-fPIC` on some systems when you specify this option.
- symbolic  
Bind references to global symbols when building a shared object. Warn about any unresolved references (unless overridden by the `-Xlinker -z -Xlinker defs` link editor option ). Only a few systems support this option.
- Xlinker *option*  
Pass *option* as an option to the linker. You can use this to supply system-specific linker options which GNU CC does not know how to recognize. If you want to pass an option that takes an argument, you must use `-Xlinker` twice, once for the option and once for the argument. For example, to pass `-assert definitions`, you must write `-Xlinker -assert -Xlinker definitions`. It does not work to write `-Xlinker "-assert definitions"`, because this passes the entire string as a single argument, which is not what the linker expects.
- Wl, *option*  
Pass *option* as an option to the linker. If *option* contains commas, it is split into multiple options at the commas.
- u *symbol*  
Pretend the symbol, *symbol*, is undefined, to force linking of library modules to define it. You can use `-u` multiple times with different symbols to force loading of additional library modules.





## Options for directory search

---

The following options specify directories to search for header files, for libraries and for parts of the compiler.

`-I $\textit{dir}$`

Add the directory, *dir*, to the head of the list of directories to be searched for header files. This can be used to override a system header file, substituting your own version, since these directories are searched before the system header file directories. If you use more than one ‘-I’ option, the directories are scanned in left-to-right order; the standard system directories come after.

`-I-`

Any directories you specify with `-I` options before the ‘-I-’ option are searched only for the case of `#include "file"`; they are not searched for `#include <file>`. If additional directories are specified with ‘-I’ options after the ‘-I-’ options, these directories are searched for all `#include` directives. (Ordinarily all ‘-I’ directories are used this way.) In addition, the ‘-I-’ option inhibits the use of the current directory (where the current input file came from) as the first search directory for `#include "file"`. There is no way to override this ‘-I-’ effect. With ‘-I’, you can specify searching the directory which was current when the compiler was invoked. That is not exactly the same as what the preprocessor does by default, but it is often satisfactory.

‘-I-’ does not inhibit the use of the standard system directories for header files. Thus, ‘-I-’ and ‘-nostdinc’ are independent.

### `-Ldir`

Add directory, *dir*, to the list of directories to be searched for ‘-l’.

### `-Bprefix`

This option specifies where to find the executables, libraries, include files, and data files of the compiler itself.

The compiler driver program runs one or more of the subprograms `cpp`, `ccl`, `as` and `ld`. It tries *prefix* as a prefix for each program it tries to run, both with and without ‘*machine/version/*’ (see “Specifying target machine and compiler version” on page 129).

For each subprogram to be run, the compiler driver first tries, if any, the ‘-B’ prefix. If that name is not found, or if ‘-B’ was not specified, the driver tries two standard prefixes, which are `/usr/lib/gcc/` and `/usr/local/lib/gcc-lib/`. If neither of those results in a filename that is found, the unmodified program name is searched for using the directories specified in your `PATH` environment variable. ‘-B’ prefixes that effectively specify directory names also apply to libraries in the linker, because the compiler translates these options into ‘-L’ options for the linker. They also apply to includes files in the preprocessor, because the compiler translates these options into `-isystem` options for the pre-processor. In this case, the compiler appends `include` to the prefix.

The run-time support file, `libgcc.a`, can also be searched for by using the ‘-B’ prefix, if needed. If it is not found there, the two standard prefixes discussed in the previous discussion are tried, and that is all. The file is left out of the link if it is not found by those means.

Another way to specify a prefix much like the ‘-B’ prefix, is to use the environment variable, `GCC_EXEC_PREFIX`. See “The offset-info option” on page 193.

### `-specs=file`

Process *file* after the compiler reads in the standard ‘*specs*’ file, in order to override the defaults that the ‘`gcc`’ driver program uses when determining what switches to pass to ‘`ccl`’, ‘`cclplus`’, ‘`as`’, ‘`ld`’, etc. More than one `-specs=file` can be specified on the command line, and they are processed in order, from left to right.

# Specifying target machine and compiler version

---

By default, GCC compiles code for the same type of machine that you are using. However, it can also be installed as a cross-compiler, to compile for some other type of machine. In fact, several different configurations of GCC, for different target machines, can be installed side by side. Then you specify which one to use with the `-b` option.

In addition, older and newer versions of GCC can be installed side by side. One of them (probably the newest) will be the default, but you may sometimes wish to use another version, using the following arguments.

`-b machine`

The argument, *machine*, specifies the target machine for compilation. This is useful when you have installed GCC as a cross-compiler.

The value to use for machine is the same as was specified as the machine type when configuring GCC as a cross-compiler. For example, if a cross-compiler was configured with '`configure i386v`', meaning to compile for an 80386 running System V, then you would specify '`-b i386v`' to run that cross compiler.

When you do not specify '`-b`', it normally means to compile for the same type of machine that you are using.

`-V version`

The argument, *version*, specifies which version of GCC to run. This is useful when multiple versions are installed. For example, version might be `2.0`, meaning

to run GCC version 2.0. The default, *version*, when you do not specify `-v`, is the last version of GCC that you installed.

The `-b` and `-v` options actually work by controlling part of the filename used for the executable files and libraries used for compilation. A given version of GCC, for a given target machine, is normally kept in a standard directory, similarly named as in the following example, where *machine* and *version* signify the machine and its compliant version that you need to specify.

```
/usr/local/lib/gcc-lib/machine/version.
```

Thus, sites can customize the effect of `-b` or `-v` either by changing the names of these directories or adding alternate names (or symbolic links). If in a directory, `/usr/local/lib/gcc-lib/`, where the `'80386'` file is a link to the `'i386v'` file, then `'-b 80386'` becomes an alias for `'-b i386v'`.

In one respect, the `-b` or `-v` do not completely change to a different compiler: the top-level driver for the compiler that you originally invoked continues to run and invoke the other executables (preprocessor, compiler per se, assembler and linker) which do the real work. However, since no real work is done in the driver program, it usually does not matter that the driver program in use is not the one for the specified target and version.

The only way that the driver program depends on the target machine is in the parsing and handling of special machine-specific options. However, this is controlled by a file which is found, along with the other executables, in the directory for the specified version and target machine. As a result, a single installed driver program adapts to any specified target machine and compiler version.

The driver program executable does control one significant thing, however: the default version and target machine. Therefore, you can install different instances of the driver program, compiled for different targets or versions, under different names.

For example, if the driver for version 2.0 is installed as `other-gcc` and that for version 2.1 is installed as `gcc`, then the command `gcc` will use version 2.1 by default, while `ogcc` will use 2.0 by default. However, you can choose either version with either command with the `-v` option.

# Hardware models and configurations

---

The following documentation discusses the options for the GNU compiler, GCC, using the following hardware processors and their configurations. Some of these configurations may not have support. Contact Cygnus if you need further support.

- “AMD29K options” on page 133
- “ARC options” on page 135
- “ARM options” on page 136
- “Clipper options” on page 140
- “Convex options” on page 141
- “D10V options” on page 143
- “DEC Alpha options” on page 144
- “Hitachi H8/300 options” on page 147
- “Hitachi SH options” on page 148
- “HPPA options” on page 149
- “IBM RS/6000 and PowerPC options” on page 151
- “IBM RT options” on page 159
- “Intel x86 options” on page 160
- “Intel 960 options” on page 163
- “M32R/D options” on page 165
- “MIPS options” on page 166

- “MN10300 options” on page 170
- “Motorola 68K options” on page 171
- “Motorola 88K options” on page 174
- “SPARC options” on page 178
- “System V options” on page 182
- “Thumb options” on page 183
- “Vax options” on page 185

In “Specifying target machine and compiler version” on page 129, we discuss the standard option, `-b`, which chooses among different installed compilers for completely different target machines, such as Vax or Motorola 68K or Hitachi SH series.

In addition, each of these target machine types can have its own special options, starting with `-m`, to choose among various hardware *models*, or configurations—for example, Motorola’s 68010 or 68020 processors, floating coprocessor or none. A single installed version of the compiler can compile for any model or configuration, according to the options specified.

Some configurations of the compiler also support additional special options, usually for compatibility with other compilers on the same platform.

# AMD29K options

The following ‘-m’ options are defined for the AMD Am29000 processor’s implementations.

`-mdw`

Generate code that assumes the `DW` bit is set, i.e., that byte and halfword operations are directly supported by the hardware. This is the default.

`-mndw`

Generate code that assumes the `DW` bit is not set.

`-mbw`

Generate code that assumes the system supports byte and halfword write operations. This is the default.

`-mnbw`

Generate code that assumes the systems does not support byte and halfword write operations. `-mnbw` implies `-mndw`.

`-msmall`

Use a small memory model that assumes that all function addresses are either within a single 256 KB segment or at an absolute address of less than 256k. This allows the `call` instruction to be used instead of a `const`, `consth`, `calli` sequence.

`-mnormal`

Use the normal memory model: Generate `call` instructions only when calling functions in the same file and `calli` instructions otherwise. This works if each file occupies less than 256 KB but allows the entire executable to be larger than 256 KB. This is the default.

`-mlarge`

Always use `calli` instructions. Specify this option if you expect a single file to compile into more than 256 KB of code.

`-m29050`

Generate code for the Am29050.

`-m29000`

Generate code for the Am29000. This is the default.

`-mkernel-registers`

Generate references to registers `gr64-gr95` instead of to `gr96-gr127` registers. This option can be used when compiling kernel code that wants a set of global registers disjoint from that used by user-mode code.

**NOTE:** When this option is used, register names in ‘-f’ flags must use the normal, user-mode, names.

`-muser-registers`

Use the normal set of global registers, `gr96-gr127`. This is the default.

`-mstack-check`

`-mno-stack-check`

Insert (or do not insert) a call to `__msp_check` after each stack adjustment. This is often used for kernel code.

`-mstorem-bug`

`-mno-storem-bug`

`-mstorem-bug` handles 29k processors which cannot handle the separation of a `mtsrin` instruction and a `storem` instruction (most 29000 chips to date, but not the 29050).

`-mno-reuse-arg-regs`

`-mreuse-arg-regs`

`-mno-reuse-arg-regs` tells the compiler to only use incoming argument registers for copying out arguments. This helps detect calling a function with fewer arguments than those with which it was declared.

`-mno-impure-text`

`-mimpure-text`

`-mimpure-text`, used in addition to `-shared`, tells the compiler to not pass `-assert pure-text` to the linker when linking a shared object.

`-msoft-float`

Generate output containing library calls for floating point.

**WARNING:** The requisite libraries are not part of GCC. Normally the facilities of the machine's usual C compiler are used, but this can't be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation.



# ARC options

The following options are for the ARC processor's implementations.

`-EL`

Compile code for little endian mode. This is the default.

`-EB`

Compile code for big endian mode.

`-mmangle-cpu`

Prepend the name of the `cpu` to all public symbol names. In multiple-processor systems, there are many ARC variants with different instruction and register set characteristics. This flag prevents code compiled for one CPU to be linked with code compiled for another CPU. No facility exists for handling variants that are “almost identical.” This is an all or nothing option.

`-mcpu=cpu`

Compile code for ARC variant, `cpu`. Which variants are supported depend on the configuration. All variants support the default, `-mcpu=base`.

`-mtext=text section`

`-mdata=data section`

`-mrodata=readonly data section`

Put functions, data, and read-only data in `text section`, `data section`, and `readonly data section`, respectively, by default. This can be overridden with the `section` attribute. See “Specifying attributes of variables” on page 235.

## ARM options

The following ‘-m’ options are defined for Advanced RISC Machines (ARM) architectures.

`-mapcs-frame`

Generate a stack frame that is compliant with the ARM Procedure Call Standard for all functions, even if this is not strictly necessary for correct execution of the code.

`-mapcs`

This is a synonym for ‘-mapcs-frame’.

`-mapcs-26`

Generate code for a processor running with a 26-bit program counter, and conforming to the function calling standards for the APCS 26-bit option. This option replaces the ‘-m2’ and ‘-m3’ options of previous releases of the compiler.

`-mapcs-32`

Generate code for a processor running with a 32-bit program counter, and conforming to the function calling standards for the APCS 32-bit option. This option replaces the ‘-m6’ option of previous releases of the compiler.

`-mapcs-stack-check`

Generate code to check the amount of stack space available upon entry to every function (that actually uses some stack space). If there is insufficient space available then either the function ‘\_\_rt\_stkovf\_split\_small’ or ‘\_\_rt\_stkovf\_split\_big’ will be called, depending upon the amount of stack space required. The run time system is required to provide these functions. The default is ‘-mno-apcs-stack-check’ since this produces smaller code.

`-mapcs-float`

Pass floating point arguments using the float point registers. This is one of the variants of the APCS. This option is recommended if the target hardware has a floating point unit or if a lot of floating point arithmetic is going to be performed by the code. The default is ‘-mno-apcs-float’, since integer only code is slightly increased in size if ‘-mapcs-float’ is used.

`-mapcs-reentrant`

Generate reentrant, position independent code. This is the equivalent to specifying the ‘-fpic’ option. The default is ‘-mno-apcs-reentrant’.

`-mthumb-interwork`

Generate code supporting calls between the ARM and THUMB instruction sets. Without this option the two instruction sets cannot be reliably used inside one program. The default is ‘-mno-thumb-interwork’, since slightly larger code is generated when ‘-mthumb-interwork’ is specified.

`-mno-sched-prolog`

Prevent the reordering of instructions in the function prolog, or the merging of those instruction with the instructions in the function's body. This means that all functions will start with a recognisable set of instructions (or in fact one of a choice from a small set of different function prologues), and this information can be used to locate the start of functions inside an executable piece of code.

The default is `'-msched-prolog'`.

`-mhard-float`

Generate output containing floating point instructions. This is the default.

`-msoft-float`

Generate output containing library calls for floating point.

**WARNING!** The requisite libraries are not available for all ARM targets. Normally the facilities of the machine's usual C compiler are used, but this cannot be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation.

`-msoft-float` changes the calling convention in the output file; therefore, it is only useful if you compile all of a program with this option. In particular, you need to compile `libgcc.a`, the library that comes with GCC, with `-msoft-float` in order for this to work.

`-mlittle-endian`

Generate code for a processor running in little-endian mode. This is the default for all standard configurations.

`-mbig-endian`

Generate code for a processor running in big-endian mode; the default is to compile code for a little-endian processor.

`-mwords-little-endian`

This option only applies when generating code for big-endian processors. Generate code for a little-endian word order but a big-endian byte order. That is, a byte order of the form 32107654.

**NOTE:** This option should only be used if you require compatibility with code for big-endian ARM processors generated by versions of the compiler prior to 2.8.

`-mshort-load-bytes`

Do not try to load half-words (e.g., `short`) by loading a word from an unaligned address. For some targets the MMU is configured to trap unaligned loads; use this option to generate code that is safe in these environments.

`-mno-short-load-bytes`

Use unaligned word loads to load half-words (e.g., `short`'s). This option produces more efficient code, but the MMU is sometimes configured to trap these instructions.

`-mshort-load-words`

This is a synonym for the `'-mno-short-load-bytes'`.

`-mno-short-load-words`

This is a synonym for the `'-mshort-load-bytes'`.

`-mbsd`

This option only applies to RISC iX. Emulate the native BSD-mode compiler.

This is the default if `'-ansi'` is not specified.

`-mxopen`

This option only applies to RISC iX. Emulate the native X/Open-mode compiler.

`-mbsd`

This option only applies to RISC iX. Emulate the native BSD-mode compiler.

This is the default if `-ansi` is not specified.

`-mxopen`

This option only applies to RISC iX. Emulate the native X/Open-mode compiler.

`-mno-symrename`

This option only applies to RISC iX. Do not run the assembler post-processor, `symrename`, after code has been assembled. Normally it is necessary to modify some of the standard symbols in preparation for linking with the RISC iX C library; this option suppresses this pass. The post-processor is never run when the compiler is built for cross-compilation.

`-mcpu=<name>`

This specifies the name (*name*) of the target ARM processor. GCC uses this name to determine what kind of instructions it can use when generating assembly code.

Permissible names are: `arm2`, `arm250`, `arm3`, `arm6`, `arm60`, `arm600`, `arm610`, `arm620`, `arm7`, `arm7m`, `arm7d`, `arm7dm`, `arm7di`, `arm7dmi`, `arm70`, `arm700`, `arm700i`, `arm710`, `arm710c`, `arm7100`, `arm7500`, `arm7500fe`, `arm7tdmi`, `arm8`, `strongarm`, `strongarm110`.

`-march=<name>`

This specifies the name (*name*) of the target ARM architecture. GCC uses this name to determine what kind of instructions it can use when generating assembly code.

This option can be used in conjunction with or instead of the `'-mcpu='` option.

Permissible names are: `armv2`, `armv2a`, `armv3`, `armv3m`, `armv4`, `armv4t`.

`-mfpe=<number>`

This specifies the version (*number*) of the floating point emulation available on the target.

Permissible values are 2 and 3.

`-mstructure-size-boundary=<n>`

The size of all structures and unions will be rounded up to a multiple of the number of bits (*n*) set by this option. Permissible values are 8 and 32. The default value varies for different toolchains. For the COFF targeted toolchain the default value is 8. Specifying the larger number can produce faster, more efficient code, but can also increase the size of the program. The two values are potentially incompatible. Code compiled with one value cannot necessarily expect to work with code or libraries compiled with the other value, if they exchange information using structures or unions. Programmers are encouraged to use the 32 value as future versions of the toolchain may default to this value.

`-mnop-fun-dllimport`

Disable the support for the *dllimport* attribute.

## Clipper options

The following ‘-m’ options are defined for the Clipper processor’s implementations.

-mc300

Produce code for a C300 Clipper processor. This is the default.

-mc400

Produce code for a C400 Clipper processor; i.e., use floating point registers, `fr8` through `fr15`.

# Convex options

The following ‘-m’ options are defined for the Convex processor’s implementations.

-mc1

Generate output for C1. The code will run on any Convex machine. The preprocessor symbol, `__convex_c1__`, is defined.

-mc2

Generate output for C2. Uses instructions not available on C1. Scheduling and other optimizations are chosen for max performance on C2. The preprocessor symbol, `__convex_c2__`, is defined.

-mc32

Generate output for C32xx. Uses instructions not available on C1. Scheduling and other optimizations are chosen for max performance on C32. The preprocessor symbol, `__convex_c32__`, is defined.

-mc34

Generate output for C34xx. Uses instructions not available on C1. Scheduling and other optimizations are chosen for max performance on C34. The preprocessor symbol, `__convex_c34__`, is defined.

-mc38

Generate output for C38xx. Uses instructions not available on C1. Scheduling and other optimizations are chosen for max performance on C38. The preprocessor symbol, `__convex_c38__`, is defined.

-margcount

Generate code which puts an argument count in the word preceding each argument list. This is compatible with regular CC, and a few programs may need the argument count word. GDB and other source-level debuggers do not need it; this info is in the symbol table.

-mnoargcount

Omit the argument count word. This is the default.

-mvolatile-cache

Allow volatile references to be cached. This is the default.

-mvolatile-nocache

Volatile references bypass the data cache, going all the way to memory. This is only needed for multi-processor code that does not use standard synchronization instructions. Making non-volatile references to volatile locations will not necessarily work.

-mlong32

Type `long` is 32 bits, the same as type `int`. This is the default.

`-mlong64`

Type `long` is 64 bits, the same as type `long long`. This option is useless, because no library support exists for it.



# D10V options

These ‘-m’ options are defined for the D10V processor’s implementations.

`-mint32`

`-mint16`

Make `int` data 32 (or 16) bits by default. The default is ‘-mint16’.

`-mdouble64`

`-mdouble32`

Make `double` data 64 (or 32) bits by default. The default is ‘-mdouble32’.

`-maddac3`

`-mno-addac3`

Enable (disable) the use of `addac3` and `subac3` instructions. The ‘-maddac3’ instruction also enables the ‘-maccum’ instruction.

`-maccum`

`-mno-accum`

Enable (disable) the use of the 32-bit accumulators in compiler generated code.

`-mno-asm-optimize`

`-masm-optimize`

Disable (enable) passing ‘-o’ to the assembler when optimizing. The assembler uses the ‘-o’ option to parallelize adjacent short instructions automatically where possible.

`-mno-small-insns`

`-msmall-insns`

Disable (enable) converting some long instructions into two short instructions, which can eliminate some nops and enable more code to be conditionally executed.

`-mno-cond-move`

`-mcond-move`

Disable (or enable) conditional move instructions, eliminating short branches.

`-mbranch-cost=n`

Increase the internal costs of branches to *n*. Higher costs means that the compiler will issue more instructions to avoid doing a branch. The default is 1.

`-mcond-exec=n`

Specify the maximum number of conditionally executed instructions that replace a branch. The default is 4.

## DEC Alpha options

The following ‘-m’ options are defined for the DEC Alpha family of processor’s implementations.

`-mno-soft-float`  
`-msoft-float`

Use (do not use) the hardware floating-point instructions for floating-point operations. When `-msoft-float` is specified, functions in `libgcc1.c` will be used to perform floating-point operations. Unless they are replaced by routines that emulate the floating-point operations, or compiled in such a way as to call such emulation routines, these routines will issue floating-point operations. If you are compiling for an Alpha without floating-point operations, you must ensure that the library is built so as not to call them.

**NOTE:** Alpha implementations without floating-point operations are required to have floating-point registers.

`-mfp-reg`  
`-mno-fp-regs`

Generate code that uses (does not use) the floating-point register set.

`-mno-fp-regs` implies `-msoft-float`. If the floating-point register set is not used, floating point operands are passed in integer registers as if they were integers and floating-point results are passed in `$0` instead of `$f0`.

This is a non-standard calling sequence, so any function with a floating-point argument or return value called by code compiled with `-mno-fp-regs` must also be compiled with that option. A typical use of this option is building a kernel that does not use, and hence need not save and restore, any floating-point registers.

`-mieee`

The Alpha architecture implements floating-point hardware optimized for maximum performance. It is mostly compliant with the IEEE floating point standard. However, for full compliance, software assistance is required. This option generates code fully IEEE compliant code *except* that the `inexact` flag is not maintained (compare following description for `-mieee-with-inexact`). If this option is turned on, the CPP macro, `_IEEE_FP`, is defined during compilation.

The option is a shorthand for ‘`-D_IEEE_FP -D_IEEE_FP_INEXACT`’ plus ‘`-mieee-conformant`’, and ‘`-mfp-trap-mode=sui`’, and ‘`-mtrap-precision=i`’.

The resulting code is less efficient but is able to correctly support denormalized numbers and exceptional IEEE values such as not-a-number and plus/minus infinity.

Other Alpha compilers call this option `-ieee_with_no_inexact`.

**-mieee-with-inexact**

This is like ‘-mieee’ except the generated code also maintains the IEEE *inexact* flag. Turning on this option causes the generated code to implement fully-compliant IEEE math. The option is a shorthand; it signifies the declaration of ‘-D\_IEEE\_FP -D\_IEEE\_FP\_INEXACT’ plus ‘-mieee-conformant’, and ‘-mfp-trap-mode=sui’, and ‘-mtrap-precision=i’. On some Alpha implementations the resulting code may execute significantly slower than the code generated by default. Since there is very little code that depends on the *inexact* flag, you should normally not specify this option. Other Alpha compilers call this option ‘-ieee\_with\_inexact’.

**-mfp-trap-mode=trap mode**

This option controls what floating-point related traps are enabled. Other Alpha compilers call this option ‘-fptm’ *trap mode*, where *trap mode* is the definition of what can be set by one of the following four values.

- **n**  
This is the default (normal) setting. The only traps that are enabled are the ones that cannot be disabled in software (such as division by zero trap).
- **u**  
In addition to the traps enabled by **n**, underflow traps are enabled as well.
- **su**  
Like **u**, but the instructions are marked to be safe for software completion (see Alpha architecture manuals for details).
- **sui**  
Like **su**, but *inexact* traps are enabled as well.

**-mfp-rounding-mode=rounding mode**

Selects the IEEE rounding mode. Other Alpha compilers call this option ‘-fprm’ *rounding mode*. The rounding mode can be one of the following four values.

- **n**  
Normal IEEE rounding mode. Floating point numbers are rounded towards the nearest machine number or towards the even machine number in case of a tie.
- **m**  
Round towards minus infinity.
- **c**  
Chopped rounding mode. Floating point numbers are rounded towards zero.
- **d**  
Dynamic rounding mode. A field in the floating point control register (*fpcr*; see Alpha architecture reference manuals for details) controls the rounding

mode in effect. The C library initializes this register for rounding towards plus infinity. Thus, unless your program modifies the `fpcr`, 'd' corresponds to round towards plus infinity.

`-mtrap-precision=trap precision`

In the Alpha architecture, floating point traps are imprecise. This means without software assistance it is impossible to recover from a floating trap and program execution normally needs to be terminated. GCC can generate code that can assist operating system trap handlers in determining the exact location that caused a floating point trap. Depending on the requirements of an application, different levels of precision can be selected, such as with the following options.

- `p`  
Program precision. This option is the default and means a trap handler can only identify which program caused a floating point exception.
- `f`  
Function precision. The trap handler can determine the function that caused a floating point exception.
- `i`  
Instruction precision. The trap handler can determine the exact instruction that caused a floating point exception.

Other Alpha compilers provide the equivalent options, '`-scope_safe`' and '`-resumption_safe`'.

`-mieee-conformant`

This option marks the generated code as IEEE conformant. You must not use this option unless you also specify '`-mtrap-precision=i`' and either '`-mfp-trap-mode=su`' or '`-mfp-trap-mode=sui`'. Its only effect is to emit the '`.eflag 48`' line in the function prologue of the generated assembly file. Under DEC UNIX, this has the effect that IEEE-conformant math library routines will be linked in.

`-mbuild-constants`

Normally GCC examines a 32- or 64-bit integer constant to see if it can construct it from smaller constants in two or three instructions. If it cannot, it will output the constant as a literal and generate code to load it from the data segment at runtime. Use this option to require GCC to construct all integer constants using code, even if it takes more instructions (the maximum is six).

You would typically use this option to build a shared library dynamic loader. Itself a shared library, it must relocate itself in memory before it can find the variables and constants in its own data segment.

# Hitachi H8/300 options

The following `-m` options are defined for the H8/300 family of processor's implementations.

`-mrelax`

Shorten some address references at link time, when possible; this option uses the linker option, `-relax`. See “ld and the H8/300 processors” on page 246 in *Using LD of GNUPro Utilities* for a fuller description.

`-ms`

Generate code for the H8/S.

`-mh`

Generate code for the H8/300H.

`-mint32`

Make `int` data 32 bits by default.

`-malign-300`

On the H8/300H, use the same alignment rules as for the H8/300. The default for the H8/300H is to align longs and floats on 4 byte boundaries. `-malign-300` causes them to be aligned on 2 byte boundaries. This option has no effect on the H8/300.

## Hitachi SH options

The following `-m` options are defined for the Hitachi SH family processor's implementations.

`-m1`

Generate code for the SH1.

`-m2`

Generate code for the SH2.

`-m3`

Generate code for the SH3.

`-m3e`

Generate code for the SH3e.

`-mb`

Compile code for the processor in big endian mode.

`-ml`

Compile code for the processor in little endian mode.

`-mrelax`

Shorten some addresses at link time, when possible; uses the linker option, `-relax`.

# HPPA options

The following ‘-m’ options are defined for the HPPA family of computers.

`-mpa-risc-1-0`

Generate code for a PA 1.0 processor.

`-mpa-risc-1-1`

Generate code for a PA 1.1 processor.

`-mbig-switch`

Generate code suitable for big switch tables. Use this option only if the assembler/linker complain about out of range branches within a switch table.

`-mjump-in-delay`

Fill delay slots of function calls with unconditional jump instructions by modifying the return pointer for the function call to be the target of the conditional jump.

`-mdisable-fpregs`

Prevent floating point registers from being used in any manner. This is necessary for compiling kernels which perform lazy context switching of floating point registers. If you use this option and attempt to perform floating point operations, the compiler will abort.

`-mdisable-indexing`

Prevent the compiler from using indexing address modes. This avoids some rather obscure problems when compiling MIG generated code under MACH.

`-mno-space-regs`

Generate code that assumes the target has no space registers.

This allows GCC to generate faster indirect calls and use unscaled index address modes. Such code is suitable for level 0 PA systems and kernels.

`-mfast-indirect-calls`

Generate code that assumes calls never cross space boundaries. This allows GCC to emit code which performs faster indirect calls.

This option will not work in the presence of shared libraries or nested functions.

`-mspace`

Optimize for space rather than execution time. Currently this only enables out of line function prologues and epilogues. This option is incompatible with PIC code generation and profiling.

`-mlong-load-store`

Generate 3-instruction load and store sequences as some-times required by the HP/UX 10 linker. This is equivalent to the ‘+k’ option to the HP compilers.

`-mportable-runtime`

Use the portable calling conventions proposed by HP for ELF systems.

`-mgas`

Enable the use of assembler directives only GAS understands.

`-mschedule=cpu type`

Schedule code according to the constraints for the machine type (signified by the *cpu type*). The choices for *cpu type* are 700 for 7n0 machines, 7100 for 7n5 machines, and 7100 for 7n2 machines. 7100 is the default for *cpu type*.

**NOTE:** The 7100LC scheduling information is incomplete and using 7100LC often leads to bad schedules. For now it's probably best to use 7100 instead of 7100LC for the 7n2 machines.

`-mlinker-opt`

Enable the optimization pass in the HP/UX linker.

**NOTE:** This makes symbolic debugging impossible. It also triggers a bug in the HP/UX 8 and HP/UX 9 linkers in which they give bogus error messages when linking some programs.

`-msoft-float`

Generate output containing library calls for floating point.

**WARNING:** The requisite libraries are not available for all HPPA targets. Normally the facilities of the machine's usual C compiler are used, but this cannot be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation. The embedded target `hppa1.1-*-pro` does provide software floating point support.

`-msoft-float` changes the calling convention in the output file; therefore, it is only useful if you compile *all* of a program with this option. In particular, you need to compile `libgcc.a`, the library that comes with GCC, with `-msoft-float` in order for this to work.



## IBM RS/6000 and PowerPC options

These ‘-m’ options are defined for the IBM RS/6000 and PowerPC processor’s implementations.

```
-mpower
-mno-power
-mpower2
-mno-power2
-mpowerpc
-mno-powerpc
-mpowerpc-gpopt
-mno-powerpc-gpopt
-mpowerpc-gfxopt
-mno-powerpc-gfxopt
```

GCC supports two related instruction set architectures for the RS/6000 and PowerPC. The POWER instruction set are those instructions supported by the *rios* chip set used in the original RS/6000 systems and the PowerPC instruction set is the architecture of the Motorola MPC5<sub>xx</sub>, MPC6<sub>xx</sub>, MCP8<sub>xx</sub> and the IBM 4<sub>xx</sub> microprocessors. The PowerPC architecture defines 64-bit instructions, but they are not supported by any current processors.

Neither architecture is a subset of the other. However there is a large common subset of instructions supported by both. An MQ register is included in processors supporting the POWER architecture.

You use these options to specify which instructions are available on the processor you are using. The default value of these options is determined when configuring GCC. Specifying the ‘-mcpu=*cpu\_type*’ overrides the specification of these options.

We recommend you use the ‘-mcpu=*cpu\_type*’ option rather than any of these options.

The ‘-mpower’ option allows GCC to generate instructions that are found only in the POWER architecture and to use the MQ register. Specifying ‘-mpower2’ implies ‘-power’ and also allows GCC to generate instructions that are present in the POWER2 architecture but not the original POWER architecture.

The ‘-mpowerpc’ option allows GCC to generate instructions that are found only in the 32-bit subset of the PowerPC architecture. Specifying ‘-mpowerpc-gpopt’ implies ‘-mpowerpc’ and also allows GCC to use the optional PowerPC architecture instructions in the General Purpose group, including floating-point square root. Specifying ‘-mpowerpc-gfxopt’ implies ‘-mpowerpc’ and also allows GCC to use the optional PowerPC architecture instructions in the Graphics group, including floating-point select. If you specify both ‘-mno-power’ and

'-mno-powerpc' options, GCC will use only the instructions in the common subset of both architectures plus some special AIX common-mode calls, and will not use the MQ register. Specifying both '-mpower' and '-mpowerpc' permits GCC to use any instruction from either architecture and to allow use of the MQ register; specify this for the Motorola MPC601.

-mnew-mnemonics  
-mold-mnemonics

Select which mnemonics to use in the generated assembler code.

'-mnew-mnemonics' requests output that uses the assembler mnemonics defined for the PowerPC architecture, while '-mold-mnemonics' requests the assembler mnemonics defined for the POWER architecture. Instructions defined in only one architecture have only one mnemonic; GCC uses that mnemonic irrespective of which of these options is specified.

PowerPC assemblers support both the old and new mnemonics, as will later POWER assemblers. Current POWER assemblers only support the old mnemonics. Specify '-mnew-mnemonics' if you have an assembler that supports them, otherwise specify '-mold-mnemonics'.

The default value of these options depends on how GCC was configured. Specifying '-mcpu=cpu\_type' sometimes overrides the value of these options. Unless you are building a cross-compiler, you should normally not specify either '-mnew-mnemonics' or '-mold-mnemonics', but should instead accept the default.

-mcpu=cpu\_type

Set architecture type, register usage, choice of mnemonics, and instruction scheduling parameters for machine type *cpu\_type*. Supported values for *cpu\_type* are 'rs6000', 'rios1', 'rios2', 'rsc', '601', '602', '603', '603e', '604', '604e', '620', 'power', 'power2', 'powerpc', '403', '505', '801', '821', '823', '860' and 'common'.

The '-mcpu=power', '-mcpu=power2', and '-mcpu=powerpc' specify generic POWER, POWER2 and pure PowerPC (i.e., not MPC601) architecture machine types, with an appropriate, generic processor model assumed for scheduling purposes.

Specifying '-mcpu=rios1', '-mcpu=rios2', '-mcpu=rsc', '-mcpu=power', or '-mcpu=power2' enables the '-mpower' option and disables the '-mpowerpc' option; '-mcpu=601' enables both the '-mpower' and '-mpowerpc' options; '-mcpu=602', '-mcpu=603', '-mcpu=603e', '-mcpu=604', '-mcpu=620'; '-mcpu=403', '-mcpu=505', '-mcpu=821', '-mcpu=860' and '-mcpu=powerpc' enable the '-mpowerpc' option and disable the '-mpower' option; '-mcpu=common' disables both the '-mpower' and '-mpowerpc' options.

AIX versions 4 or greater selects '-mcpu=common' by default, so that code will

operate on all members of the RS/6000 and PowerPC families. In that case, GCC will use only the instructions in the common subset of both architectures plus some special AIX common-mode calls, and will not use the MQ register. GCC assumes a generic processor model for scheduling purposes.

Specifying `'-mcpu=rios1'`, `'-mcpu=rios2'`, `'-mcpu=rsc'`, `'-mcpu=power'`, or `'-mcpu=power2'` also disables the `'new-mnemonics'` option.

Specifying `'-mcpu=601'`, `'-mcpu=602'`, `'-mcpu=603'`, `'-mcpu=603e'`, `'-mcpu=604'`, `'-mcpu=620'`, `'-mcpu=403'`, or `'-mcpu=powerpc'` also enables the `'new-mnemonics'` option.

Specifying `'-mcpu=403'`, `'-mcpu=821'`, or `'-mcpu=860'` also enables the `'msoft-float'` option.

`-mtune=cpu_type`

Set the instruction scheduling parameters for machine type, *cpu\_type*, but do not set the architecture type, register usage, choice of mnemonics like `'-mcpu=cpu_type'` would. The same values for *cpu\_type* are used for `'-mtune=cpu_type'` as for `'-mcpu=cpu_type'`. The `'-mtune=cpu_type'` option overrides the `'-mcpu=cpu_type'` option in terms of instruction scheduling parameters.

`-mfull-toc`

`-mno-fp-in-toc`

`-mno-sum-in-toc`

`-mminimal-toc`

Modify generation of the TOC (Table Of Contents), which is created for every executable file. The `'-mfull-toc'` option is selected by default. In that case, GCC will allocate at least one TOC entry for each unique non-automatic variable reference in your program. GCC will also place floating-point constants in the TOC. However, only 16,384 entries are available in the TOC.

If you receive a linker error message that saying you have overflowed the available TOC space, you can reduce the amount of TOC space used with the `'-mno-fp-in-toc'` and `'-mno-sum-in-toc'` options.

`'-mno-fp-in-toc'` prevents GCC from putting floating-point constants in the TOC and `'-mno-sum-in-toc'` forces GCC to generate code to calculate the sum of an address and a constant at run-time instead of putting that sum into the TOC. You may specify one or both of these options. Each causes GCC to produce very slightly slower and larger code at the expense of conserving TOC space.

If you still run out of space in the TOC even when you specify both of these options, specify `'-mminimal-toc'` instead. This option causes GCC to make only one TOC entry for every file. When you specify this option, GCC will produce code that is slower and larger but which uses extremely little TOC space. You may wish to use this option only on files that contain less frequently executed

code.

`-mxl-call`

`-mno-xl-call`

On AIX, pass floating-point arguments to prototyped functions beyond the register save area (RSA) on the stack in addition to floating point register arguments. The AIX calling convention was extended but not initially documented to handle an obscure K&R C case of calling a function that takes the address of its arguments with fewer arguments than declared. AIX XL compilers assume that floating point arguments which do not fit in the RSA are on the stack when they compile a subroutine without optimization. Because always storing floating-point arguments on the stack is inefficient and rarely needed, this option is not enabled by default and only is necessary when calling subroutines compiled by AIX XL compilers without optimization.

`-mthreads`

Support *AIX Threads*. Link an application written to use *pthread*s with special libraries and startup code to enable the application to run.

`-mpe`

Support *IBM RS/6000 SP Parallel Environment* (PE). Link an application written to use message passing with special startup code to enable the application to run. The system must have PE installed in the standard location (`/usr/lpp/ppe.poe/`), or the `'specs'` file must be overridden with the `'-specs='` option to specify the appropriate directory location.

The Parallel Environment does not support threads, so the `'-mpe'` option and the `'-mthreads'` option are incompatible.

`-msoft-float`

`-mhard-float`

Generate code that does not use (uses) the floating-point register set. Software floating point emulation is provided if you use the `'-msoft-float'` option, and pass the option to GCC when linking.

`-mmultiple`

`-mno-multiple`

Generate code that uses (does not use) the load multiple word instructions and the store multiple word instructions. These instructions are generated by default on POWER systems, and not generated on PowerPC systems. Do not use `'-mmultiple'` on little endian PowerPC systems, since those instructions do not work when the processor is in little endian mode.

`-mstring`

`-mno-string`

Generate code that uses (does not use) the load string instructions and the store string word instructions to save multiple registers and do small block moves.

These instructions are generated by default on POWER systems, and not generated on PowerPC systems.

**WARNING:** Do not use `-mstring` on little endian PowerPC systems, since those instructions do not work when the processor is in little endian mode.

`-mupdate`  
`-mno-update`

Generate code that uses (or does not use) the load or store instructions that update the base register to the address of the calculated memory location. These instructions are generated by default.

If you use `'-mno-update'`, there is a small window between the time that the stack pointer is updated and the address of the previous frame is stored, which means code that walks the stack frame across interrupts or signals may get corrupted data.

`-mfused-madd`  
`-mno-fused-madd`

Generate code that uses (does not use) the floating point multiply and accumulate instructions. These instructions are generated by default if hardware floating is used.

`-mno-bit-align`  
`-mbit-align`

On System V.4 and embedded PowerPC systems do not (do) force structures and unions that contain bit fields to be aligned to the base type of the bit field. For example, by default a structure containing nothing but 8 unsigned bitfields of length 1 would be aligned to a 4 byte boundary and have a size of 4 bytes. By using `-mno-bit-align`, the structure would be aligned to a 1 byte boundary and be one byte in size.

`-mno-strict-align`  
`-mstrict-align`

On System V.4 and embedded PowerPC systems do not (do) assume that unaligned memory references will be handled by the system.

`-mrelocatable`  
`-mno-relocatable`

On embedded PowerPC systems generate code that allows (does not allow) the program to be relocated to a different address at runtime. If you use `-mrelocatable` on any module, all objects linked together must be compiled with `-mrelocatable` or `-mrelocatable-lib`.

`-mrelocatable-lib`  
`-mno-relocatable-lib`

On embedded PowerPC systems generate code that allows (does not allow) the program to be relocated to a different address at runtime. Modules compiled with `'-mrelocatable-lib'` can be linked with either modules compiled without

`'-mrelocatable'` and `'-mrelocatable-lib'` or with modules compiled with the `'-mrelocatable'` options.

`-mno-toc`  
`-mtoc`

On System V.4 and embedded PowerPC systems do not (do) assume that register 2 contains a pointer to a global area pointing to the addresses used in the program.

`-mno-traceback`  
`-mtraceback`

On embedded PowerPC systems do not (do) generate a trace-back tag before the start of the function. This tag can be used by the debugger to identify where the start of a function is.

`-mlittle`  
`-mlittle-endian`

On System V.4 and embedded PowerPC systems compile code for the processor in little endian mode. The `'-mlittle-endian'` option is the same as `'-mlittle'`.

`-mbig`  
`-mbig-endian`

On System V.4 and embedded PowerPC systems compile code for the processor in big endian mode. The `'-mbig-endian'` option is the same as `'-mbig'`.

`-mcall-sysv`

On System V.4 and embedded PowerPC systems compile code using calling conventions that adheres to the March 1995 draft of the System V ABI, PowerPC processor supplement. This is the default unless you configured GCC using `'powerpc-*-eabiaix'`.

`-mcall-sysv-eabi`

Specify both `'-mcall-sysv'` and `'-meabi'` options.

`-mcall-sysv-noeabi`

Specify both `'-mcall-sysv'` and `'-mnoeabi'` options.

`-mcall-aix`

On System V.4 and embedded PowerPC systems compile code using calling conventions that are similar to those used on AIX. This is the default if you configured GCC using `'powerpc-*-eabiaix'`.

`-mcall-solaris`

On System V.4 and embedded PowerPC systems, compile code for the Solaris operating system.

`-mcall-linux`

On System V.4 and embedded PowerPC systems, compile code for the Linux operating system.

`-mprototype`  
`-mno-prototype`

On System V.4 and embedded PowerPC systems assume that all calls to variable argument functions are properly prototyped. Otherwise, the compiler must insert

an instruction before every non prototyped call to set or clear bit 6 of the condition code register (*CR*) to indicate whether floating point values were passed in the floating point registers in case the function takes a variable arguments.

With ‘`-mprototype`’, only calls to prototyped variable argument functions will set or clear the bit.

#### `-msim`

On embedded PowerPC systems, assume that the startup module is called `sim-crt0.o` and the standard C libraries are `libsim.a` and `libc.a`. This is default for ‘`powerpc-*-eabisim`’ configurations.

#### `-mmvme`

On embedded PowerPC systems, assume that the startup module is called `mvme-crt0.o` and the standard C libraries are ‘`libmvme.a`’ and ‘`libc.a`’.

#### `-memb`

On embedded PowerPC systems, set the `PPC_EMB` bit in the ELF flags header to indicate that `eabi` extended relocations are used.

#### `-mads`

On embedded PowerPC systems, assume that the startup module is called ‘`crt0.o`’ and the standard C libraries are ‘`libads.a`’ and ‘`libc.a`’.

#### `-myellowknife`

On embedded PowerPC systems, assume that the startup module is called ‘`crt0.o`’ and ‘`libyk.a`’ and ‘`libc.a`’ are the standard C libraries.

#### `-meabi`

#### `-mno-eabi`

On System V.4 and embedded PowerPC systems do (do not) adhere to the Embedded Applications Binary Interface (EABI) which is a set of modifications to the System V.4 specifications. Selecting `-meabi` means that the stack is aligned to an 8 byte boundary, a function, `__eabi`, is called to from `main` to set up the EABI environment, and the ‘`-msdata`’ option can use both `r2` and `r13` to point to two separate small data areas.

Selecting `-mno-eabi` means that the stack is aligned to a 16 byte boundary, do not call an initialization function from `main`, and the ‘`-msdata`’ option will only use `r13` to point to a single small data area. The ‘`-meabi`’ option is on by default if you configured GCC using one of the ‘`powerpc*-*-eabi*`’ options.

`-msdata=eabi`

On System V.4 and embedded PowerPC systems, put small initialized const global and static data in the `‘.sdata2’` section, which is pointed to by register `r2`. Put small initialized non-const global and static data in the `‘.sdata’` section, which is pointed to by register `r13`. Put small uninitialized global and static data in the `‘.sbss’` section, which is adjacent to the `‘.sdata’` section. The `‘-msdata=eabi’` option is incompatible with the `‘-mrelocatable’` option. The `‘-msdata=eabi’` option also sets the `‘-memb’` option.

`-msdata=sysv`

On System V.4 and embedded PowerPC systems, put small global and static data in the `‘.sdata’` section, which is pointed to by register `r13`. Put small uninitialized global and static data in the `‘.sbss’` section, which is adjacent to the `‘.sdata’` section. The `‘-msdata=sysv’` option is incompatible with the `‘-mrelocatable’` option.

`-msdata=default`

`-msdata`

On System V.4 and embedded PowerPC systems, if `‘-meabi’` is used, compile code the same as `‘-msdata=eabi’`, otherwise compile code the same as `‘-msdata=sysv’`.

`-msdata-data`

On System V.4 and embedded PowerPC systems, put small global and static data in the `‘.sdata’` section. Put small uninitialized global and static data in the `‘.sbss’` section. Do not use register `r13` to address small data however.

This is the default behavior unless other `‘-msdata’` options are used.

`-msdata=none`

`-mno-sdata`

On embedded PowerPC systems, put all initialized global and static data in the `‘.data’` section, and all uninitialized data in the `‘.bss’` section.

`-G num`

On embedded PowerPC systems, put global and static items less than or equal to *num* bytes into the small data or bss sections instead of the normal data or bss section. By default, *num* is 8. The `‘-G num’` switch is also passed to the linker. All modules should be compiled with the same `‘-G num’` value.

`-mregnames`

`-mno-regnames`

On System V.4 and embedded PowerPC systems, do (do not) emit register names in the assembly language output using symbolic forms.



# IBM RT options

The following ‘-m’ options are defined for the IBM RT PC architectures.

`-min-line-mul`

Use an in-line code sequence for integer multiplies. This is the default.

`-mcall-lib-mul`

Call `lmul$$` for integer multiplies.

`-mfull-fp-blocks`

Generate full-size floating point data blocks, including the minimum amount of scratch space recommended by IBM. This is the default.

`-mminimum-fp-blocks`

Do not include extra scratch space in floating point data blocks. This results in smaller code, but slower execution, since scratch space must be allocated dynamically.

`-mfp-arg-in-fpregs`

Use a calling sequence incompatible with the IBM calling convention in which floating point arguments are passed in floating point registers.

**NOTE:** ‘`varargs.h`’ and ‘`stdargs.h`’ will not work with floating point operands if this option is specified.

`-mfp-arg-in-gregs`

Use the normal calling convention for floating point arguments. This is the default.

`-mhc-struct-return`

Return structures of more than one word in memory, rather than in a register. This provides compatibility with the MetaWare HighC (`hc`) compiler. Use the option ‘`fpcc-struct-return`’ for compatibility with the Portable C Compiler (`pcc`).

`-mnohc-struct-return`

Return some structures of more than one word in registers, when convenient. This is the default. For compatibility with the IBM-supplied compilers, use the option ‘`fpcc-struct-return`’ or the option ‘`-mhc-struct-return`’.

## Intel x86 options

The following ‘-m’ options are defined for the i<sub>x</sub>86 family of computers.

`-m486`

`-m386`

Control whether or not code is optimized for a 486 instead of an 386. Code generated for a 486 will run on a 386 and vice versa.

`-mieee-fp`

`-mno-ieee-fp`

Control whether or not the compiler uses IEEE floating point comparisons. These handle correctly the case where the result of a comparison is unordered.

`-msoft-float`

Generate output containing library calls for floating point.

**WARNING:** The requisite libraries are not part of GCC. Normally the facilities of the machine’s usual C compiler are used, but this can’t be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation. On machines where a function returns floating point results in the 80387 register stack, some floating point opcodes may be emitted even if `-msoft-float` is used.

`-mno-fp-ret-in-387`

Do not use the FPU registers for return values of functions.

The usual calling convention has functions return values of types `float` and `double` in an FPU register, even if there is no FPU. The idea is that the operating system should emulate an FPU.

The option `-mno-fp-ret-in-387` causes such values to be returned in ordinary CPU registers instead.

`-mno-fancy-math-387`

Some 387 emulators do not support the `sin`, `cos` and `sqrt` instructions for the 387. Specify this option to avoid generating those instructions. This option is the default on FreeBSD. As of revision 2.6.1, these instructions are not generated unless you also use the `-ffast-math` switch.

`-malign-double`

`-mno-align-double`

Control whether GCC aligns `double`, `long double`, and `long long` variables on a two word boundary or a one word boundary. Aligning double variables on a two word boundary will produce code that runs somewhat faster on a Pentium at the expense of more memory.

**WARNING:** If you use the `-malign-double` switch, structures containing the above types will be aligned differently than the published application binary interface

specifications for the 386.

`-msvr3-shlib`  
`-mno-svr3-shlib`

Control whether GCC places uninitialized locals into `bss` or `data`. ‘`-msvr3-shlib`’ places these locals into `bss`. These options are meaningful only on System V Release 3.

`-mno-wide-multiply`  
`-mwide-multiply`

Control whether GCC uses the `mul` and `imul` that produce 64 bit results in `eax:edx` from 32 bit operands to do `long long` multiplies and 32-bit division by constants.

`-mrtcd`

Use a different function-calling convention, in which functions that take a fixed number of arguments return with the `ret num` instruction, which pops their arguments while returning. This saves one instruction in the caller since there is no need to pop the arguments there.

You can specify that an individual function is called with this calling sequence with the function attribute `stdcall`. You can also override the option, `-mrtcd`, by using the function attribute, `cdecl`. See “Declaring attributes of functions” on page 226.

**WARNING:** This calling convention is incompatible with the one normally used on UNIX, so you cannot use it if you need to call libraries compiled with the UNIX compiler.

You must provide function prototypes for all functions that take variable numbers of arguments (including `printf`); otherwise incorrect code will be generated for calls to those functions.

In addition, seriously incorrect code will result if you call a function with too many arguments. Normally, extra arguments are harmlessly ignored.

`-mreg-alloc=regs`

Control the default allocation order of integer registers. The string `regs` is a series of letters specifying a register. The supported letters are: `a` allocate `EAX`; `b` allocate `EBX`; `c` allocate `ECX`; `d` allocate `EDX`; `s` allocate `ESI`; `D` allocate `EDI`; `B` allocate `EBP`.

`-mregparm=num`

Control how many registers are used to pass integer arguments. By default, no registers are used to pass arguments, and at most 3 registers can be used. You can control this behavior for a specific function by using the function attribute `regparm`. See “Declaring attributes of functions” on page 226.

**WARNING:** If you use this switch, and `num` is nonzero, then you must build all modules

with the same value, including any libraries. This includes the system libraries and startup modules.

`-malign-loops=num`

Align loops to a 2 raised to a *num* byte boundary. If `-malign-loops` is not specified, the default is 2.

`-malign-jumps=num`

Align instructions that are only jumped to a 2 raised to a *num* byte boundary. If `-malign-jumps` is not specified, the default is 2 if optimizing for a 386, and 4 if optimizing for a 486.

`-malign-functions=num`

Align the start of functions to a 2 raised to *num* byte boundary.

If `-malign-jumps` is not specified, the default is 2 if optimizing for a 386, and 4 if optimizing for a 486.

# Intel 960 options

The following ‘-m’ options are defined for the Intel 960 implementations.

`-mcpu type`

Assume the defaults for the machine type *cpu type* for some of the other options, including instruction scheduling, floating-point support, and addressing modes.

The choices for *cpu type* are *ka*, *kb*, *mc*, *ca*, *cf*, *sa*, and *sb*. The default is *kb*.

`-mnumerics`

`-msoft-float`

The `-mnumerics` option indicates that the processor does support floating-point instructions. The `-msoft-float` option indicates that floating-point support should not be assumed.

`-mleaf-procedures`

`-mno-leaf-procedures`

Do (or do not) attempt to alter leaf procedures to be callable with the `bal` instruction as well as `call`. This will result in more efficient code for explicit calls when the `bal` instruction can be substituted by the assembler or linker, but less efficient code in other cases, such as calls via function pointers, or using a linker that doesn’t support this optimization.

`-mtail-call`

`-mno-tail-call`

Do (or do not) make additional attempts (beyond those of the machine-independent portions of the compiler) to optimize tail-recursive calls into branches. You may not want to do this because the detection of cases where this is not valid is not totally complete. The default is `-mno-tail-call`.

`-mcomplex-addr`

`-mno-complex-addr`

Assume (or do not assume) that the use of a complex addressing mode is a win on this implementation of the i960. Complex addressing modes may not be worthwhile on the K-series, but they definitely are on the C-series. The default is currently `-mcomplex-addr` for all processors except the CB and CC.

`-mcode-align`

`-mno-code-align`

Align code to 8-byte boundaries for faster fetching (or don’t bother). Currently turned on by default for C-series implementations only.

`-mic-compat`

`-mic2.0-compat`

`-mic3.0-compat`

Enable compatibility with iC960 v2.0 or v3.0.

`-masm-compat`

`-mintel-asm`

Enable compatibility with the iC960 assembler.

`-mstrict-align`

`-mno-strict-align`

Do not permit (do permit) unaligned accesses.

`-mold-align`

Enable structure-alignment compatibility with Intel's GCC release version 1.3 (based on GCC 1.37). This option implies `-mstrict-align`.

# M32R/D options

The following ‘-m’ options are defined for Mitsubishi M32R/D architectures.

`-mcode-model=small`

Assume all objects live in the lower 16MB of memory (so that their addresses can be loaded with the `ld24` instruction), and assume all subroutines are reachable with the `bl` instruction. This is the default. The addressability of a particular object can be set with the `model` attribute.

`-mcode-model=medium`

Assume objects may be anywhere in the 32 bit address space (the compiler will generate `seth/add3` instructions to load their addresses), and assume all subroutines are reachable with the `bl` instruction.

`-mcode-model=large`

Assume objects may be anywhere in the 32 bit address space (the compiler will generate `seth/add3` instructions to load their addresses), and assume subroutines may not be reachable with the `bl` instruction (the compiler will generate the much slower `seth/add3/jl` instruction sequence).

`-msdata=none`

Disable use of the small data area. Variables will be put into one of ‘.data’, ‘bss’, or ‘.rodata’ (unless the section attribute has been specified). This is the default. The small data area consists of sections ‘.sdata’ and ‘.sbss’. Objects may be explicitly put in the small data area with the `section` attribute using one of these sections.

`-msdata=sdata`

Put small global and static data in the small data area, but do not generate special code to reference them.

`-msdata=use`

Put small global and static data in the small data area, and generate special instructions to reference them.

`-G num`

Put global and static objects less than or equal to *num* bytes into the small data or bss sections instead of the normal data or bss sections. The default value of *num* is 8. The ‘-msdata’ option must be set to one of ‘sdata’ or ‘use’ for this option to have any effect. All modules should be compiled with the same ‘-G *num*’ value. Compiling with different values of *num* may or may not work; if it doesn’t, the linker will give an error message: `incorrect code will not be generated`.

## MIPS options

The following ‘-m’ options are defined for the MIPS family of computers.

`-mcpu=cpu type`

Assume the defaults for the machine type *cpu type* when scheduling instructions. The choices for *cpu type* are ‘r2000’, ‘r3000’, ‘r4000’, ‘r4400’, ‘r4600’, and ‘r6000’. While picking a specific *cpu type* will schedule things appropriately for that particular chip, the compiler will not generate any code that does not meet level 1 of the MIPS ISA (instruction set architecture) without the ‘-mips2’ or ‘-mips3’ switches being used.

`-mips1`

Issue instructions from level 1 of the MIPS ISA. This is the default. ‘r3000’ is the default *cpu type* at this ISA level.

`-mips2`

Issue instructions from level 2 of the MIPS ISA (branch likely, square root instructions). ‘r6000’ is the default CPU type at this ISA level.

`-mips3`

Issue instructions from level 3 of the MIPS ISA (64 bit instructions). r4000 is the default *cpu type* at this ISA level. This option does not change the sizes of any of the C data types.

`-mfp32`

Assume that 32 32-bit floating point registers are available. This is the default.

`-mfp64`

Assume that 32 64-bit floating point registers are available. This is the default when the `-mips3` option is used.

`-mgp32`

Assume that 32 32-bit general purpose registers are available. This is the default.

`-mgp64`

Assume that 32 64-bit general purpose registers are available. This is the default when the `-mips3` option is used.

`-mint64`

Types `long`, `int`, and `pointer` are 64 bits. This works only if `-mips3` is also specified.

`-mlong64`

Types `long` and `pointer` are 64 bits, and type `int` is 32 bits. This works only if `-mips3` is also specified.



**-mmips-as**

Generate code for the MIPS assembler, and invoke `mips-tfile` to add normal debug information. This is the default for all platforms except for the OSF/1 reference platform, using the OSF/rose object format. If either of the `-gstabs` or `-gstabs+` switches are used, the `mips-tfile` program will encapsulate the `stabs` within MIPS ECOFF.

**-mgas**

Generate code for the GNU assembler. This is the default on the OSF/1 reference platform, using the OSF/rose object format.

**-msplit-addresses****-mno-split-addresses**

Generate code to load the high and low parts of address constants separately. This allows `gcc` to optimize away redundant loads of the high order bits of addresses. This optimization requires GNU `as` and GNU `ld`. This optimization is enabled by default for some embedded targets where GNU `as` and GNU `ld` are standard.

**-mrnames****-mno-rnames**

The `-mrnames` switch says to output code using the MIPS software names for the registers, instead of the hardware names (i.e., `a0` instead of `$4`). The only known assembler that supports this option is the Algorithmics assembler.

**-mgpopt****-mno-gpopt**

The `-mgpopt` switch says to write all of the data declarations before the instructions in the text section, this allows the MIPS assembler to generate one word memory references instead of using two words for short global or static data items. This is on by default if optimization is selected.

**-mstats****-mno-stats**

For each non-inline function processed, the `-mstats` switch causes the compiler to emit one line to the standard error file to print statistics about the program (number of registers saved, stack size, etc.).

**-mmemcpy****-mno-memcpy**

The `-mmemcpy` switch makes all block moves call the appropriate string function (`memcpy` or `bcopy`) instead of possibly generating inline code.

**-mmips-tfile****-mno-mips-tfile**

The `-mno-mips-tfile` switch causes the compiler not post-process the object file with the `mips-tfile` program, after the MIPS assembler has generated it to add debug support. If `mips-tfile` is not run, then no local variables will be available to the debugger. In addition, `stage2` and `stage3` objects will have the temporary filenames passed to the assembler embedded in the object file, which means the objects will not compare the same. The `-mno-mips-tfile` switch should only be

used when there are bugs in the `mips-tfile` program that prevents compilation.

`-msoft-float`

Generate output containing library calls for floating point.

**WARNING:** The requisite libraries are not part of GCC. Normally the facilities of the machine's usual C compiler are used, but this can't be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation.

`-mhard-float`

Generate output containing floating point instructions. This is the default if you use the unmodified sources.

`-mabicalls`

`-mno-abicalls`

Emit (or do not emit) the pseudo operations `.abicalls`, `.cpload`, and `.cprestore` that some System V.4 ports use for position independent code.

`-mlong-calls`

`-mno-long-calls`

Do all calls with the `JALR` instruction, which requires loading up a function's address into a register before the call. You need to use this switch, if you call outside of the current 512 megabyte segment to functions that are not through pointers.

`-mhalf-pic`

`-mno-half-pic`

Put pointers to extern references into the data section and load them up, rather than put the references in the text section.

`-membedded-pic`

`-mno-embedded-pic`

Generate PIC code suitable for some embedded systems. All calls are made using PC relative address, and all data is addressed using the register, `$gp`. This requires GNU `as` and GNU `ld` which do most of the work.

`-membedded-data`

`-mno-embedded-data`

Allocate variables to the read-only data section first if possible, then next in the small data section if possible, otherwise in data. This gives slightly slower code than the default, reducing the amount of RAM required when executing, and thus may be preferable for some embedded systems.

`-msingle-float`

`-mdouble-float`

The `-msingle-float` switch tells `gcc` to assume that the floating point coprocessor only supports single precision operations, as on the `r4650` chip. The `-mdouble-float` switch permits `gcc` to use double precision operations. This is the default.

- mmad
- mno-mad
  - Permit use of the `mad`, `madu` and `mul` instructions, as on the `r4650` chip.
- m4650
  - Turns on `-msingle-float`, `-mmad`, and, at least for now, `-mcpu=r4650`.
- EL
  - Compile code for the processor in little endian mode. The requisite libraries are assumed to exist.
- EB
  - Compile code for the processor in big endian mode. The requisite libraries are assumed to exist.
- G *num*
  - Put global and static items less than or equal to *num* bytes into the sections, `small data` or `bss`, instead of the normal `data` or `bss` section. This allows the assembler to emit one word memory reference instructions based on the global pointer (`gp` or `$28`), instead of the normal two words used. By default, *num* is 8. The `'-G num'` switch is also passed to the assembler and linker. All modules should be compiled with the same `'-G num'` value.
- nocpp
  - Tell the MIPS assembler to not run its preprocessor over user assembler files (with a suffix, `'.s'`) when assembling them.

## MN10300 options

These ‘-m’ options are defined for Matsushita MN10300 architectures.

`-mmult-bug`

Generate code to avoid bugs in the multiply instructions for the MN10300 processors. This is the default.

`-mno-mult-bug`

Do not generate code to avoid bugs in the multiply instructions for the MN10300 processors.

# Motorola 68K options

The following explains the ‘-m’ options defined for the Motorola 68000 series. The default values for these options depends on which style of 68000 was selected when the compiler was configured; the defaults for the most common choices are also given with the following options.

`-m68000`  
`-mc68000`

Generate output for a 68000. This is the default when the compiler is configured for 68000-based systems.

`-m68020`  
`-mc68020`

Generate output for a 68020. This is the default when the compiler is configured for 68020-based systems.

`-m68881`

Generate output containing 68881 instructions for floating point. This is the default for most 68020 systems unless ‘-nfp’ was specified when the compiler was configured.

`-m68030`

Generate output for a 68030. This is the default when the compiler is configured for 68030-based systems.

`-m68040`

Generate output for a 68040. This is the default when the compiler is configured for 68040-based systems.

This option inhibits the use of 68881/68882 instructions that have to be emulated by software on the 68040. If your 68040 does not have code to emulate those instructions, use ‘-m68040’.

`-m68060`

Generate output for a 68060. This is the default when the compiler is configured for 68060-based systems. This option inhibits the use of 68020 and 68881/68882 instructions that have to be emulated by software on the 68060. If your 68060 does not have code to emulate those instructions, use ‘-m68060’.

`-m68020-40`

Generate output for a 68040, without using any of the new instructions. This results in code which can run relatively efficiently on either a 68020/68881 or a 68030 or a 68040. The generated code does use the 68881 instructions that are emulated on the 68040.

`-m68020-60`

Generate output for a 68060, without using any of the new instructions. This results in code which can run relatively efficiently on either a 68020/68881 or a 68030 or a 68040. The generated code does use the 68881 instructions that are emulated on the 68060.

`-mfpa`

Generate output containing SunFPA instructions for floating point.

`-msoft-float`

Generate output containing library calls for floating point.

**WARNING:** The requisite libraries are not available for all `m68k` targets. Normally the facilities of the machine's usual C compiler are used, but this can't be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation. `m68k-*-aout` and `m68k-*-coff` do provide software floating point support.

`-mshort`

Consider type `int` to be 16 bits wide, like `short int`.

`-mnobitfield`

Do not use the bit-field instructions. The `-m68000` option implies `-mnobitfield`.

`-mbitfield`

Do use the bit-field instructions. The `-m68020` option implies `-mbitfield`. This is the default if you use a configuration designed for a 68020.

`-mrtd`

Use a different function-calling convention, in which functions that take a fixed number of arguments return with the `rtd` instruction, which pops their arguments while returning. This saves one instruction in the caller as there is no need to pop the arguments there.

The calling convention using the `-mrtd` option is incompatible with the one normally used on UNIX, so you cannot use it if you need to call libraries compiled with the UNIX compiler.

Also, you must provide function prototypes for all functions that take variable numbers of arguments (including `printf`); otherwise incorrect code will be generated for calls to those functions.

In addition, seriously incorrect code will result if you call a function with too many arguments. (Normally, extra arguments are harmlessly ignored.)

The `rtd` instruction is supported by the 68010, 68020, 68030, 68040, and 68060 processors but not by the 68000 or the 5200 processors.

---

`-malign-int`  
`-mno-align-int`

Control whether GCC aligns `int`, `long`, `long long`, `float`, `double`, and `long double` variables on a 32-bit boundary (`'-malign-int'`) or a 16-bit boundary (`'-mno-align-int'`). Aligning variables on 32-bit boundaries produces code that runs somewhat faster on processors with 32-bit busses at the expense of more memory.

**WARNING:** Using the `'-malign-int'` switch, GCC will align structures having the above types differently than most published ABI specifications for the `m68k`.

## Motorola 88K options

The following ‘-m’ options are defined for Motorola 88k architectures.

`-m88000`

Generate code that works well on both the m88100 and the m88110.

`-m88100`

Generate code that works best for the m88100, but that also runs on the m88110.

`-m88110`

Generate code that works best for the m88110, and may not run on the m88100.

`-mbig-pic`

Obsolete option to be removed from the next revision. Use `-fPIC`.

`-midentify-revision`

Include an `ident` directive in the assembler output recording the source filename, compiler name and version, timestamp, and compilation flags.

`-mno-underscores`

In assembler output, emit symbol names without adding an underscore character at the beginning of each name. The default is to use an underscore as prefix on each name.

`-mocs-debug-info`

`-mno-ocs-debug-info`

Include (or omit) additional debugging information (about registers used in each stack frame) as specified in the 88open Object Compatibility Standard, “OCS.”

This extra information allows debugging of code that has had the frame pointer eliminated. The default for DG/UX, SVr4, and Delta 88 SVr3.2 is to include this information; other 88k configurations omit this information by default.

`-mocs-frame-position`

When emitting COFF debugging information for automatic variables and parameters stored on the stack, use the offset from the canonical frame address which is the stack pointer (register 31) on entry to the function. The DG/UX, SVr4, Delta88 SVr3.2, and BCS configurations use `-mocs-frame-position`; other 88k configurations have the default, `-mno-ocs-frame-position`.

`-mno-ocs-frame-position`

When emitting COFF debugging information for automatic variables and parameters stored on the stack, use the offset from the frame pointer register (register 30). When this option is in effect, the frame pointer is not eliminated when debugging information is selected by the `-g` switch.

`-moptimize-arg-area`

`-mno-optimize-arg-area`

Control how function arguments are stored in stack frames.

`-moptimize-arg-area` saves space by optimizing them, but this conflicts with the



88open specifications. The opposite alternative, `-mno-optimize-arg-area`, agrees with 88open standards. By default GCC does not optimize the argument area.

`-mshort-data-num`

Generate smaller data references by making them relative to `r0`, which allows loading a value using a single instruction (rather than the usual two). You control which data references are affected by specifying *num* with this option. For example, if you specify `-mshort-data-512`, then the data references affected are those involving displacements of less than 512 bytes. `-mshort-data-num` is not effective for *num* greater than 64k.

`-mserialize-volatile`

`-mno-serialize-volatile`

Do, or don't, generate code to guarantee sequential consistency of volatile memory references. By default, consistency is guaranteed.

The order of memory references made by the MC88110 processor does not always match the order of the instructions requesting those references. In particular, a load instruction may execute before a preceding store instruction. Such reordering violates sequential consistency of volatile memory references, when there are multiple processors. When consistency must be guaranteed, GNU C generates special instructions, as needed, to force execution in the proper order.

The MC88100 processor does not reorder memory references and so always provides sequential consistency. However, by default, GNU C generates the special instructions to guarantee consistency even when you use `-m88100`, so that the code may be run on an MC88110 processor. If you intend to run your code only on the MC88100 processor, you may use `-mno-serialize-volatile`.

The extra code generated to guarantee consistency may affect the performance of your application. If you know that you can safely forgo this guarantee, you may use `-mno-serialize-volatile`.

`-msvr4`

`-msvr3`

Turn on (`-msvr4`) or off (`-msvr3`) compiler extensions related to System V release 4 (SVr4).

This controls the following implementations.

- Which variant of the assembler syntax to emit.
- `-msvr4` makes the C preprocessor recognize `'#pragma weak'` that is used on System V release 4.
- `-msvr4` makes GCC issue additional declaration directives used in SVr4.

`-msvr4` is the default for the `m88k-motorola-sysv4` and `m88k-dg-dgux` M88K configurations. `-msvr3` is the default for all other M88K configurations.

`-mversion-03.00`

This option is obsolete, and is ignored.

`-mno-check-zero-division`

`-mcheck-zero-division`

Do, or don't, generate code to guarantee that integer division by zero will be detected. By default, detection is guaranteed.

Some models of the MC88100 processor fail to trap upon integer division by zero under certain conditions. By default, when compiling code that might be run on such a processor, GNU C generates code that explicitly checks for zero-valued divisors and traps with exception number 503 when one is detected. Use of `-mno-check-zero-division` suppresses such checking for code generated to run on an MC88100 processor.

GNU C assumes that the MC88110 processor correctly detects all instances of integer division by zero. When `-m88110` is specified, both `-mcheck-zero-division` and `-mno-check-zero-division` are ignored, and no explicit checks for zero-valued divisors are generated.

`-muse-div-instruction`

Use the `div` instruction for signed integer division on the MC88100 processor. By default, the `div` instruction is not used.

On the MC88100 processor the signed integer division instruction `div` traps to the operating system on a negative operand. The operating system transparently completes the operation, but at a large cost in execution time. By default, when compiling code that might be run on an MC88100 processor, GNU C emulates signed integer division using the unsigned integer division instruction `divu`, thereby avoiding the large penalty of a trap to the operating system. Such emulation has its own, smaller, execution cost in both time and space. To the extent that your code's important signed integer division operations are performed on two nonnegative operands, it may be desirable to use the `div` instruction directly.

On the MC88110 processor the `div` instruction (also known as the `divs` instruction) processes negative operands without trapping to the operating system. When `-m88110` is specified, `-muse-div-instruction` is ignored, and the `div` instruction is used for signed integer division.

**NOTE:** The result of dividing `INT_MIN` by `-1` is undefined. In particular, the behavior of such a division with and without `-muse-div-instruction` may differ.

`-mtrap-large-shift`

`-mhandle-large-shift`

Include code to detect bit-shifts of more than 31 bits; respectively, trap such shifts or emit code to handle them properly. By default GCC makes no special provision for large bit shifts.

---

`-mwarn-passed-structs`

Warn when a function passes a struct as an argument or result. Structure-passing conventions have changed during the evolution of the C language, and are often the source of portability problems. By default, GCC issues no such warning.

## SPARC options

The following ‘-m’ switches are supported on the SPARC architectures.

`-mno-app-regs`  
`-mapp-regs`

Specify `-mapp-regs` to generate output using the global registers 2 through 4, which the SPARC SVR4 ABI reserves for applications. These options are the default.

To be fully SVR4 ABI compliant at the cost of some performance loss, specify `-mno-app-regs`. You should compile libraries and system software with `-mno-app-regs`.

`-mfp`  
`-mhard-float`

Generate output containing floating point instructions. These options are the default.

`-mno-fpu`  
`-msoft-float`

Generate output containing library calls for floating point.

**WARNING:** The requisite libraries are not available for all SPARC targets. Normally the facilities of the machine’s usual C compiler are used, but this cannot be done directly in cross-compilation. You must make your own arrangements to provide suitable library functions for cross-compilation. The embedded targets, `sparc-*-aout` and `sparclite-*-*`, do provide software floating point support.

`-msoft-float` changes the calling convention in the output file; therefore, it is only useful if you compile *all* of a program with this option. In particular, you need to compile `libgcc.a`, the library that comes with GCC, with `-msoft-float` in order for this to work.

`-mhard-quad-float`

Generate output containing quad-word (`long double`) floating point instructions.

`-msoft-quad-float`

Generate output containing library calls for quad-word (`long double`) floating point instructions. The functions called are those specified in the SPARC ABI. This is the default.

As of this writing, there are no SPARC implementations that have hardware support for the quad-word floating point instructions. They all invoke a trap handler for one of these instructions, and then the trap handler emulates the effect of the instruction. Because of the trap handler overhead, this is much slower than calling the ABI library routines. Thus the `-msoft-quad-float` option is the default.

`-mno-epilogue`  
`-mepilogue`

With `-mepilogue` (the default), the compiler always emits code for function exit at the end of each function. Any function exit in the middle of the function (such as a return statement in C) will generate a jump to the exit code at the end of the function. With `-mno-epilogue`, the compiler tries to emit exit code inline at every function exit.

`-mno-flat`  
`-mflat`

With `-mflat`, the compiler does not generate save/restore instructions and will use a *flat* or single register window calling convention. This model uses `%i7` as the frame pointer and is compatible with the normal register window model. Code from either may be intermixed. The local registers and the input registers (0-5) are still treated as “*call saved*” registers and will be saved on the stack as necessary.

With `-mno-flat` (the default), the compiler emits save/restore instructions (except for leaf functions) and is the normal mode of operation.

`-mno-unaligned-doubles`  
`-munaligned-doubles`

Assume that doubles have 8 byte alignment. `-mno-unaligned-doubles` is the default.

With `-munaligned-doubles`, GCC assumes that doubles have 8 byte alignment only if they are contained in another type, or if they have an absolute address.

Otherwise, `-munaligned-doubles` assumes they have 4 byte alignment.

Specifying `-munaligned-doubles` avoids some rare compatibility problems with code generated by other compilers. It is not the default because a performance loss results, especially for floating point code.

`-mv8`  
`-msparclite`

These two options select variations on the SPARC architecture.

By default (unless specifically configured for the Fujitsu SPARClite), GCC generates code for the v7 variant of the SPARC architecture.

`-mv8` will give you SPARC v8 code. The only difference from v7 code is that the compiler emits the integer multiply and integer divide instructions which exist in SPARC v8 but not in SPARC v7.

`-msparclite` will give you SPARClite code. This adds the `integer multiply`, `integer divide step` and `scan (ffs)` instructions which exist in SPARClite but not in SPARC v7.

These options are deprecated and will be deleted in GCC 2.9. They have been replaced with `-mcpu=xxx`.

`-mcypress`  
`-msupersparc`

These two options select the processor for which the code is optimized.

With `-mcypress` (the default), the compiler optimizes code for the Cypress CY7C602 chip, as used in the SparcStation/ SparcServer 3xx series. This is also appropriate for the older SparcStation 1, 2, IPX, and so forth.

With `-msupersparc` the compiler optimizes code for the SuperSparc CPU, as used in the SparcStation 10, 1000 and 2000 series. This flag also enables use of the full SPARC v8 instruction set. These options are deprecated and will be deleted in GCC 2.9. They have been replaced with `-mcpu=xxx`.

`-mcpu=cpu_type`

Set architecture type and instruction scheduling parameters for machine type *cpu\_type*. Supported values for *cpu\_type* are `common`, `cypress`, `v8`, `supersparc`, `sparclite`, `f930`, `f934`, `sparclet`, `90c701`, `v8plus`, `v9`, and `ultrasparc`. Specifying ‘v9’ is only supported on true 64 bit targets.

`-mtune=cpu_type`

Set the instruction scheduling parameters for machine type *cpu\_type*, but do not set the architecture type like `-mcpu=cpu_type` would. The same values for `-mcpu=cpu_type` are used for `-tune=cpu_type`.

The following ‘-m’ switches are supported in addition to the previous switches on the SPARClet processors.

`-mlittle-endian`

Generate code for a processor running in little-endian mode.

`-mlive-g0`

Treat register `%g0` as a normal register. GCC will continue to clobber it as necessary but will not assume it always reads as 0.

`-mbroken-saverestore`

Generate code that does not use non-trivial forms of the `save` and `restore` instructions. Early versions of the SPARClet processor do not correctly handle `save` and `restore` instructions used with arguments. They correctly handle them used without arguments. A `save` instruction used without arguments increments the current window pointer but does not allocate a new stack frame. It is assumed that the window overflow trap handler will properly handle this case as will interrupt handlers.

The following `-m` switches are supported in addition to the previous switches on SPARC V9 processors in 64 bit environments.

`-mlittle-endian`

Generate code for a processor running in little-endian mode.

`-mmedlow`

Generate code for the Medium/Low code model: assume a 32 bit address space. Programs are statically linked, PIC is not supported. Pointers are still 64 bits. It is very likely that a future version of GCC will rename this option.

`-mmedany`

Generate code for the Medium/Anywhere code model: assume a 32 bit text and a 32 bit data segment, both starting anywhere (determined at link time). Programs are statically linked, PIC is not supported. Pointers are still 64 bits.

It is very likely that a future version of GCC will rename this option.

`-mfullany`

Generate code for the Full/Anywhere code model: assume a full 64 bit address space. PIC is not supported.

It is very likely that a future version of GCC will rename this option.

`-mint64`

Types `long` and `int` are 64 bits.

`-mlong32`

Types `long` and `int` are 32 bits.

`-mlong64`

`-mint32`

Type `long` is 64 bits, and type `int` is 32 bits.

`-mstack-bias`

`-mno-stack-bias`

With `-mstack-bias`, assume that the stack pointer, and frame pointer if present, are offset by `-2047` which must be added back when making stack frame references. Otherwise, assume no such offset is present.

## System V options

The following additional options are available on System V Release 4 architectures for compatibility with other compilers on those systems.

- G  
Create a shared object. It is recommended that ‘-symbolic’ or ‘-shared’ be used instead.
- Qy  
Identify the versions of each tool used by the compiler, in a ‘.ident’ assembler directive in the output.
- Qn  
Refrain from adding .ident directives to the output file (this is the default).
- YP, *dirs*  
Search the directories, *dirs*, and no others, for libraries specified with ‘-l’.
- Ym, *dir*  
Look in the directory, *dir*, to find the M4 preprocessor. The assembler uses this option.



# Thumb options

The following ‘-m’ options are available on Thumb architectures.

`-mthumb-interwork`

Generate code which supports calling between the Thumb and ARM instruction sets. Without this option the two instruction sets cannot be reliably used inside one program. The default is ‘`-mno-thumb-interwork`’, since slightly smaller code is generated with this option.

`-mtpcs-frame`

Generate a stack frame that is compliant with the Thumb Procedure Call Standard for all non-leaf functions. (A leaf function is one that does not call any other functions). The default is ‘`-mno-apcs-frame`’.

`-mtpcs-leaf-frame`

Generate a stack frame that is compliant with the Thumb Procedure Call Standard for all leaf functions. (A leaf function is one that does not call any other functions). The default is ‘`-mno-apcs-leaf-frame`’.

`-mlittle-endian`

Generate code for a processor running in little-endian mode. This is the default for all standard configurations.

`-mbig-endian`

Generate code for a processor running in big-endian mode.

`-mstructure-size-boundary=<n>`

The size of all structures and unions will be rounded up to a multiple of the number of bits (*n*) set by this option. Permissible values are 8 and 32. The default value varies for different toolchains. For the COFF targeted toolchain, the default value is 8. Specifying the larger number can produce faster, more efficient code, but can also increase the size of the program. The two values are potentially incompatible. Code compiled with one value cannot necessarily expect to work with code or libraries compiled with the other value, if they exchange information using structures or unions. Programmers are encouraged to use the 32 value as future versions of the toolchain may default to this value.

`-mno-fun-dllimport`

Disable the support for the *dllimport* attribute.

`-mcallee-super-interworking`

Gives all externally visible functions in the file being compiled an ARM instruction set header which switches to Thumb mode before executing the rest of the function. This allows these functions to be called from non-interworking code.

## Thumb options

---

### `-mcaller-super-interworking`

Allows calls via function pointers (including virtual functions) to execute correctly regardless of whether the target code has been compiled for interworking or not. There is a small overhead in the cost of executing a function pointer if this option is enabled.

---

# Vax options

The following ‘-m’ options are defined for the Vax compiler.

-munix

Do not output certain jump instructions (aobleg, and so on) that the UNIX assembler for the Vax cannot handle across long ranges.

-mgnu

Do output those jump instructions, on the assumption that you will assemble with the GNU assembler.

-mg

Output code for g-format floating point numbers instead of d-format.



# Options for code generation conventions

---

These machine-independent options control the interface conventions used in code generation. Most of them have both positive and negative forms; the negative form of `-ffoo` would be `-fno-foo`. In the following options, only one of the forms is listed—the one which is not the default. You can figure out the other form by either removing ‘no-’ or adding it.

`-fpcc-struct-return`

Return “*short*” struct and union values in memory like longer ones, rather than in registers. This convention is less efficient, but it has the advantage of allowing intercallability between GCC-compiled files and files compiled with other compilers.

The precise convention for returning structures in memory depends on the target configuration macros.

Short structures and unions are those whose size and alignment match that of some integer type.

`-freg-struct-return`

Use the convention that struct and union values are returned in registers when possible. This is more efficient for small structures than ‘`-fpcc-struct-return`’.

If you specify neither ‘`-fpcc-struct-return`’ nor its contrary ‘`-freg-struct-return`’, GCC defaults to whichever convention is standard for the target. If there is no standard convention, GCC defaults to ‘`-fpcc-struct-return`’, except on targets where GCC is the principal compiler.

In those cases, we can choose the standard, and we chose the more efficient register return alternative.

`-fshort-enums`

Allocate to an `enum` type only as many bytes as it needs for the declared range of possible values. Specifically, the `enum` type will be equivalent to the smallest integer type which has enough room.

`-fshort-double`

Use the same size for `double` as for `float`.

`-fshared-data`

Requests that the data and non-`const` variables of this compilation be shared data rather than private data. The distinction makes sense only on certain operating systems, where shared data is shared between processes running the same program, while private data exists in one copy per process.

`-fno-common`

Allocate even uninitialized global variables in the `bss` section of the object file, rather than generating them as common blocks. This has the effect that if the same variable is declared (without `extern`) in two different compilations, you will get an error when you link them. The only reason this might be useful is if you wish to verify that the program will work on other systems which always work this way.

`-fno-ident`

Ignore the `#ident` directive.

`-fno-gnu-linker`

Do not output global initializations (such as C++ constructors and destructors) in the form used by the GNU linker (on systems where the GNU linker is the standard method of handling them). Use this option when you want to use a non-GNU linker, which also requires using the `collect2` program to make sure the system linker includes constructors and destructors. (`collect2` is included in the GCC distribution.) For systems which must use `collect2`, the compiler driver GCC is configured to do this automatically.

`-finhibit-size-directive`

Don't output a `.size` assembler directive, or anything else that would cause trouble if the function is split in the middle, and the two halves are placed at locations far apart in memory. This option is used when compiling `'crtstuff.c'`; you should not need to use it for anything else.

`-fverbose-asm`

Put extra commentary information in the generated assembly code to make it more readable. This option is generally only of use to those who actually need to read the generated assembly code (perhaps while debugging the compiler itself).

'`-fverbose-asm`' is the default. '`-fno-verbose-asm`' causes the extra information to be omitted and is useful when comparing two assembler files.

**-fvolatile**

Consider all memory references through pointers to be volatile.

**-fvolatile-global**

Consider all memory references to extern and global data items to be volatile.

**-fpic**

Generate position-independent code (PIC) suitable for use in a shared library, if supported for the target machine. Such code accesses all constant addresses through a global offset table (GOT). The dynamic loader resolves the GOT entries when the program starts (the dynamic loader is not part of GCC; it is part of the operating system). If the GOT size for the linked executable exceeds a machine-specific maximum size, you get an error message from the linker indicating that ‘-fpic’ does not work; in that case, recompile with ‘-fPIC’ instead. (These maximums are 16K on the m88k, 8K on the SPARC, and 32K on the m68k and RS/6000. The 386 has no such limit.)

Position-independent code requires special support, and therefore works only on certain machines. For the 386, GCC supports PIC for System V but not for the Sun 386i. Code generated for the IBM RS/6000 is always position-independent.

**-fPIC**

If supported for the target machine, emit position-independent code, suitable for dynamic linking and avoiding any limit on the size of the global offset table. This option makes a difference on the m68k, m88k, and the SPARC.

Position-independent code requires special support, and therefore works only on certain machines.

**-ffixed-reg**

Treat the register named *reg* as a fixed register; generated code should never refer to it (except perhaps as a stack pointer, frame pointer or in some other fixed role).

*reg* must be the name of a register. The register names accepted are machine-specific and are defined in the REGISTER\_NAMES macro in the machine description macro file.

This flag does not have a negative form, because it specifies a three-way choice.

**-fcall-used-reg**

Treat the register named *reg* as an allocatable register that is clobbered by function calls. It may be allocated for temporaries or variables that do not live across a call. Functions compiled this way will not save and restore the register *reg*.

Use of this flag for a register that has a fixed pervasive role in the machine’s execution model, such as the stack pointer or frame pointer, will produce disastrous results. This flag does not have a negative form, because it specifies a three-way choice.

### `-fcall-saved-reg`

Treat the register named *reg* as an allocatable register saved by functions. It may be allocated even for temporaries or variables that live across a call. Functions compiled this way will save and restore the register *reg* if they use it. Use of this flag for a register that has a fixed pervasive role in the machine's execution model, such as the stack pointer or frame pointer, will produce disastrous results.

A different sort of disaster will result from the use of this flag for a register in which function values may be returned. This flag does not have a negative form, because it specifies a three-way choice.

### `-fexceptions`

Enable exception handling, and generate extra code needed to propagate exceptions. If you do not specify this option, GCC enables it by default for languages like C++ that normally require exception handling, and disabled for languages like C that do not normally require it. However, when compiling C code that needs to interoperate properly with exception handlers written in C++, you may need to enable this option. You may also wish to disable this option if you are compiling older C++ programs that don't use exception handling.

### `-fexceptions`

Enable exception handling. For some targets, this implies generation of frame unwind information for all functions, which can produce significant data size overhead, though it does not affect execution.

This option is on by default for languages that support exception handling (such as C++), and off for those that don't (such as C).

### `-fpack-struct`

Pack all structure members together without holes. Usually you would not want to use this option, since it makes the code suboptimal, and the offsets of structure members won't agree with system libraries.

### `-fcheck-memory-usage`

Generate extra code to check each memory access. GCC will generate code that is suitable for a detector of bad memory accesses such as 'Checker'. If you specify this option, you can not use the `asm` or `__asm__` keywords.

You must also specify this option when you compile functions you call that have side effects. If you do not, you may get erroneous messages from the detector. Normally, you should compile all your code with this option.

If you use functions from a library that have side-effects (such as `read`), you may not be able to recompile the library and specify this option. In that case, you can enable the '`-fprefix-function-name`' option, which requests GCC to encapsulate your code and make other functions look as if they were compiled with '`-fcheck-memory-usage`'. This is done by calling "stubs" which are



provided by the detector. If you cannot find or build stubs for every function you call, you may have to specify `-fcheck-memory-usage` with `-fprefix-function-name`.

#### `-fprefix-function-name`

Request GCC to add a prefix to the symbols generated for function names. GCC adds a prefix to the names of functions defined as well as functions called. Code compiled with this option and code compiled without the option can't be linked together, unless stubs are used.

Consider compiling the following code with `-fprefix-function-name`.

```
extern void bar (int);
void
foo (int a)
{
 return prefix_bar (a + 5);
}
```

GCC will then compile the code written like the following example.

```
extern void prefix_bar (int);
void
prefix_foo (int a)
{
 return prefix_bar (a + 5);
}
```

#### `-fstack-check`

Generate code to verify that you do not go beyond the boundary of the stack. You should specify this flag if you are running in an environment with multiple threads, but only rarely need to specify it in a single-threaded environment since stack overflow is automatically detected on nearly all systems if there is only one stack.

#### `+e0`

#### `+e1`

C++ only. Control whether virtual function definitions in classes are used to generate code, or only to define interfaces for their callers.

These options are provided for compatibility with `cfront 1.x` usage; the recommended alternative G++ usage is in flux. See “Declarations and definitions in one header” on page 276.

With `+e0`, virtual function definitions in classes are declared `extern`; the declaration is used only as an interface specification, not to generate code for the virtual functions (in this compilation).

With `+e1`, G++ actually generates the code implementing virtual functions defined

in the code, and makes them publicly visible.

`-funaligned-pointers`

Assume that all pointers contain unaligned addresses. On machines where unaligned memory accesses trap, this will result in much larger and slower code for all pointer dereferences, but the code will work even if addresses are unaligned.

`-funaligned-struct-hack`

Always access structure fields using loads and stores of the declared size. This option is useful for code that dereferences pointers to unaligned structures, but only accesses fields that are themselves aligned. Without this option, GCC may try to use a memory access larger than the field. This might give an unaligned access fault on some hardware. This option makes some invalid code work at the expense of disabling some optimizations. It is strongly recommended that this option not be used.

## The offset-info option

---

The `-offset-info output-file` option simplifies access to C struct's from assembler. For each member of each structure the compiler will output a `.equ` directive to associate with a symbol with the member's offset in bytes into the structure. The symbol itself is the concatenation of the structure's tag name and the member's name, separated by an underscore.

This option will output to the specified *output-file* an assembler directive, `.equ`, for each member of each structure found in each manipulation. The `.equ` directives for the structures in the header file can be obtained by using the following input: `gcc -fsyntax-only -offset-info m.s -x c m.h`.

`m.h` is the header containing the structures, and `m.s` holds the directives.

The following is a short example of output produced by `-offset-info`.  
input file (for example `m.h`):

```
struct W {
 double d;
 int I;
};

struct X {
 int a;
 int b;

 struct Y {
```

## The `offset-info` option

---

```
int a;
int b;
};

struct Y y;
struct Y yy[10];
struct Y* p;
}
```

output file (for example `m.s`):

```
.equ W_d,0
.equ W_i,8
.equ Y_a,0
.equ Y_b,4
.equ X_a,0
.equ X_b,4
.equ X_y,8
.equ X_yy,16
.equ X_p,96
```

`-offset-info` has the following caveats.

- No directives are output for bit-field members.
- No directives are output for members whose offsets (as measured in bits) is greater than the word size of the host.
- No directives are output for members whose offsets are not constants. This can happen only in structures that use some GCC specific extensions allowing for variable sized members.

# Environment variables affecting GCC

---

The following documentation describes several special *environment variables* that affect how GCC operates. They work by specifying directories or prefixes to use when searching for various kinds of files.

**NOTE:** You can also specify places to search using options such as ‘-B’, ‘-I’ and ‘-L’ (see “Options for directory search” on page 127). These take precedence over places specified using environment variables, which in turn take precedence over those specified by the GCC configuration.

## `TMPDIR`

If `TMPDIR` is set, it specifies the directory to use for temporary files. GCC uses temporary files to hold the output of one stage of compilation which is to be used as input to the next stage: for example, the output of the preprocessor, which is the input to the compiler proper.

## `GCC_EXEC_PREFIX`

If `GCC_EXEC_PREFIX` is set, it specifies a prefix to use in the names of the subprograms executed by the compiler. No slash is added when this prefix is combined with the name of a subprogram, but you can specify a prefix that ends with a slash if you wish.

If GCC cannot find the subprogram using the specified prefix, it tries looking in the usual places for the subprogram. The default value of `GCC_EXEC_PREFIX` is `prefix/lib/gcc-lib/` where `prefix` is the value of `prefix` when you ran the

`configure` script. Other prefixes specified with ‘-B’ take precedence over this prefix.

This prefix is also used for finding files such as `crt0.o` that are used for linking. In addition, the prefix is used in an unusual way in finding the directories to search for header files.

For each of the standard directories whose name normally begins with `/usr/local/lib/gcc-lib` (more precisely, with the value of `GCC_INCLUDE_DIR`), GCC tries replacing that beginning with the specified prefix to produce an alternate directory name. So, with `-Bfoo/`, GCC will search `foo/bar` where it would normally search `/usr/local/lib/bar`. These alternate directories are searched first; the standard directories come next.

### COMPILER\_PATH

The value of `COMPILER_PATH` is a colon-separated list of directories, much like `PATH`. GCC tries the directories thus specified when searching for subprograms, if it can’t find the subprograms using `GCC_EXEC_PREFIX`.

### LIBRARY\_PATH

The value of `LIBRARY_PATH` is a colon-separated list of directories, much like `PATH`.

When configured as a native compiler, GCC tries the directories thus specified when searching for special linker files, if it can’t find them using `GCC_EXEC_PREFIX`.

Linking using GCC also uses these directories when searching for ordinary libraries for the `-l` option (but directories specified with ‘-L’ come first).

### C\_INCLUDE\_PATH

### CPLUS\_INCLUDE\_PATH

### OBJC\_INCLUDE\_PATH

These environment variables pertain to particular languages. Each variable’s value is a colon-separated list of directories, much like `PATH`. When GCC searches for header files, it tries the directories listed in the variable for the language you are using, after the directories specified with ‘-I’ but before the standard header file directories.

### DEPENDENCIES\_OUTPUT

If this variable is set, its value specifies how to output dependencies for Make based on the header files processed by the compiler. This output looks much like the output from the ‘-M’ option (see “Options controlling the preprocessor” on page 117), but it goes to a separate file, and is in addition to the usual results of compilation. The value of `DEPENDENCIES_OUTPUT` can be just a filename, in which case the Make rules are written to that file, guessing the target name from the source filename. Or the value can have the form, *file target*, in which case the rules are written to file, *file*, using *target* as the target name.

## Running the protoize program

---

The `protoize` program is an optional part of GCC. You can use it to add *prototypes* to a program, thus converting the program to ANSI C in one respect. The companion program, `unprotoize`, does the reverse: it removes argument types from any prototypes that are found.

When you run these programs, you must specify a set of source files as command line arguments. The conversion programs start out by compiling these files to see what functions they define. The information gathered about a file, `foo`, is saved in a file named `foo.X`.

After scanning comes actual conversion. The specified files are all eligible to be converted; any files they include (whether sources or just headers) are eligible as well.

But not all the eligible files are converted. By default, `protoize` and `unprotoize` convert only source and header files in the current directory.

You can specify additional directories whose files should be converted with the `-d` *directory* option. You can also specify particular files to exclude with the `-x file` option. A file is converted if it is eligible, its directory name matches one of the specified directory names, and its name within the directory has not been excluded.

Basic conversion with `protoize` consists of rewriting most function definitions and function declarations to specify the types of the arguments. The only ones not rewritten are those for `varargs` functions.

`protoize` optionally inserts prototype declarations at the beginning of the source file, to make them available for any calls that precede the function's definition. Or it can

insert prototype declarations with block scope in the blocks where undeclared functions are called.

Basic conversion with `unprotoize` consists of rewriting most function declarations to remove any argument types, and rewriting function definitions to the old-style pre-ANSI form.

Both conversion programs print a warning for any function declaration or definition they can't convert. You can suppress these warnings with `-q`.

The output from `protoize` or `unprotoize` replaces the original source file. The original file is renamed to a name ending with `.save`. If the `.save` file already exists, then the source file is simply discarded.

`protoize` and `unprotoize` both depend on GCC itself to scan the program and collect information about the functions it uses. So neither of these programs will work until GCC is installed.

You can use the following options with `protoize` and `unprotoize`. Each option works with both programs unless otherwise stated.

`-B directory`

Look for the file, `SYSCALLS.c.x`, in the specified directory, *directory*, instead of the usual directory (normally `/usr/local/lib`). This file contains prototype information about standard system functions. This option applies only to `protoize`.

`-c compilation-options`

Use *compilation-options* as the options when running GCC to produce the `.x` files. The special option `-aux-info` is always passed in addition, to tell GCC to write a `.x` file.

**WARNING:** The compilation options must be given as a single argument to `protoize` or `unprotoize`. If you want to specify several GCC options, *you must quote the entire set of compilation options to make them a single word in the shell*. You *cannot* use certain GCC arguments because they produce the *wrong kind of output*. These include `-g`, `-O`, `-c`, `-S`, and `-o`. If you include these in the *'compilation-options'*, they are ignored.

`-C`

Rename files to end in `.c` instead of `.c'`. This is convenient if you are converting a C program to C++. This option applies only to `protoize`.

`-g`

Add explicit global declarations. This means inserting explicit declarations at the beginning of each source file for each function that is called in the file and was not declared. These declarations precede the first function definition that contains a call to an undeclared function. This option applies only to `protoize`.



- `-i string`  
Indent old-style parameter declarations with the string *string*. This option applies only to `protoize`.  
`unprotoize` converts prototyped function definitions to old-style function definitions, where the arguments are declared between the argument list and the initial brace, '{'. By default, `unprotoize` uses five spaces as the indentation. If you want to indent with just one space instead, use '`-i " "`'.
- `-k`  
Keep the `.x` files. Normally, they are deleted after conversion is finished.
- `-l`  
Add explicit local declarations. `protoize` with '`-l`' inserts a prototype declaration for each function in each block which calls the function without any declaration. This option applies only to `protoize`.
- `-n`  
Make no real changes. This mode just prints information about the conversions that would have been done without '`-n`'.
- `-N`  
Make no `.save` files. The original files are simply deleted. Use this option with caution.
- `-p program`  
Use the program, *program*, as the compiler. Normally, GCC, is used.
- `-q`  
Work quietly. Most warnings are suppressed.
- `-v`  
Print the version number, just like '`-v`' for GCC.

If you need special compiler options to compile one of your program's source files, then you should generate that file's `.x` file specially, by running GCC on that source file with the `-aux-info` option and the appropriate options. Then run `protoize` on the entire set of files. `protoize` will use the existing `.x` file because it is newer than the source file. Use the following example for `protoize`.

```
gcc -Dfoo=bar file1.c -aux-info
protoize *.c
```

You need to include the special files along with the rest in the `protoize` command, even though their `.x` files already exist, because otherwise they won't get converted.

See "protoize and unprotoize warnings" on page 313 for more information on how to use `protoize` successfully.



# Extensions to the C language family

---

GCC provides several language features not found in ANSI standard C. (The option, `-pedantic`, directs GCC to print a warning message if any of these features is used.) To test for the availability of these features in conditional compilation, check for a predefined macro, `__GNUC__`, which is always defined under GCC.

These extensions are available in C and Objective C. Most of them are also available in C++. See “Extensions to the C++ language” on page 271 for extensions that apply *only* to C++.

See the following documentation for discussion on the subject of extensions to C.

- “Statements and declarations in expressions” on page 203
- “Locally declared labels” on page 204
- “Labels as values” on page 205
- “Nested functions” on page 206
- “Constructing function calls” on page 208
- “Naming an expression’s type” on page 209
- “Referring to a type with the `typeof` keyword” on page 210
- “Generalized lvalues” on page 211
- “Conditionals with omitted operands” on page 213
- “Double-word integers” on page 214
- “Complex numbers” on page 215

- “Arrays of length zero” on page 216
- “Arrays of variable length” on page 217
- “Macros with variable numbers of arguments” on page 219
- “Non-lvalue arrays may have subscripts” on page 220
- “Arithmetic on void- and function-pointers” on page 221
- “Non-constant initializers” on page 222
- “Constructor expressions” on page 223
- “Labeled elements in initializers” on page 224
- “Declaring attributes of functions” on page 226
- “Prototypes and old-style function definitions” on page 232
- “Compiling functions for interrupt calls” on page 233
- “Inquiring on alignment of types or variables” on page 234
- “Specifying attributes of variables” on page 235
- “Specifying attributes of types” on page 239
- “An inline function is as fast as a macro” on page 243
- “Assembler instructions with C expression operands” on page 245
- “Constraints for asm operands” on page 249
- “Controlling names used in assembler code” on page 262
- “Variables in specified registers” on page 263
- “Alternate keywords” on page 266
- “Incomplete enum types” on page 267
- “Function names as strings” on page 268

# Statements and declarations in expressions

A compound statement enclosed in parentheses may appear as an expression in GCC. This allows you to use loops, switches, and local variables within an expression.

Recall that a compound statement is a sequence of statements surrounded by braces; in the following construct, parentheses go around the braces.

```
({ int y = foo (); int z;
 if (y > 0) z = y;
 else z = - y;
 z; })
```

The previous construct is a valid (though slightly more complex than necessary) expression for the absolute value of `foo()`.

The last thing in the compound statement should be an expression followed by a semicolon; the value of this subexpression serves as the value of the entire construct. (If you use some other kind of statement last within the braces, the construct has type `void`, and thus effectively no value.)

This feature is especially useful in making macro definitions “*safe*” (so that they evaluate each operand exactly once). For example, the “*maximum*” function is commonly defined as a macro in standard C as follows.

```
#define max(a,b) ((a) > (b) ? (a) : (b))
```

But this definition computes either ‘*a*’ or ‘*b*’ twice, with bad results if the operand has side effects. In GCC, if you know the type of the operands (in the following example, `int`), you can define the macro safely as follows.

```
#define maxint(a,b) ({int _a = (a), _b = (b); _a > _b ? _a : _b; })
```

Embedded statements are not allowed in constant expressions, such as the value of an enumeration constant, the width of a bit field, or the initial value of a static variable.

If you don’t know the type of the operand, you can still do this, but you must use `typeof` (see “Referring to a type with the `typeof` keyword” on page 210) or type naming (see “Naming an expression’s type” on page 209).

## Locally declared labels

Each statement expression is a scope in which *local labels* can be declared. A local label is simply an identifier.

You can jump to it with an ordinary `goto` statement, but only from within the statement expression it belongs to; a local label declaration looks like

`__label__ label` or `__label__ label1, label2, ...`

Local label declarations must come at the beginning of the statement expression, right after the parenthesis and brace, ‘( {’, before any ordinary declarations.

The label declaration defines the label *name*, but does not define the label itself. You must do this in the usual way, with `label:`, within the statements of the statement expression.

The local label feature is useful because statement expressions are often used in macros. If the macro contains nested loops, a `goto` can be useful for breaking out of them. However, an ordinary label whose scope is the whole function cannot be used: if the macro can be expanded several times in one function, the label will be multiply defined in that function. A local label avoids this problem. Use the following example, for instance.

```
#define SEARCH(array, target) \
({ \
 __label__ found; \
 typeof (target) _SEARCH_target = (target); \
 typeof (*(array)) *_SEARCH_array = (array); \
 int i, j; \ int value; \
 for (i = 0; i < max; i++) \
 for (j = 0; j < max; j++) \
 if (_SEARCH_array[i][j] == _SEARCH_target) \
 { value = i; goto found; } \
 value = -1; \
found: \
 value; \
})
```

# Labels as values

You can get the address of a label defined in the current function (or a containing function) with the unary operator, `&&`.

The value has type `void *`. This value is a constant and can be used wherever a constant of that type is valid. Use the following example, for instance.

```
void *ptr;
...
ptr = &&foo;
```

To use these values, you need to be able to jump to one. This is done with the computed `goto` statement<sup>†</sup>, `goto * exp;` as in `goto *ptr;`.

Any expression of type `void *` is allowed.

One way of using these constants is in initializing a static array that will serve as a jump table:

```
static void *array[] = { &&foo, &&bar, &&hack };
```

Then you can select a label with indexing, as in the following example.

```
goto *array[i];.
```

**NOTE:** This does not check whether the subscript is in bounds—array indexing in C never does that.

Such an array of label values serves a purpose much like that of the `switch` statement. The `switch` statement is cleaner, so use that rather than an array unless the problem does not fit a `switch` statement very well.

Another use of label values is in an interpreter for threaded code. The labels within the interpreter function can be stored in the threaded code for super-fast dispatching.

You can use this mechanism to jump to code in a different function. If you do that, totally unpredictable things will happen. The best way to avoid this is to store the label address only in automatic variables and never pass it as an argument.

<sup>†</sup> The analogous feature in FORTRAN is called an assigned `goto`, but that name seems inappropriate in C, where one can do more than simply store label addresses in label variables.

## Nested functions

A *nested function* is a function defined inside another function. (Nested functions are not supported for GNU C++.) The nested function's name is local to the block where it is defined. For example, in the following, we define a nested function named `square`, and call it twice.

```
foo (double a, double b)
{
 double square (double z) { return z * z; }

 return square (a) + square (b);
}
```

The nested function can access all the variables of the containing function that are visible at the point of its definition. This is called *lexical scoping*. For example, in the following, we show a nested function which uses an inherited variable named `offset`.

```
bar (int *array, int offset, int size)
{
 int access (int *array, int index)
 { return array[index + offset]; }
 int i;
 ...
 for (i = 0; i < size; i++)
 ... access (array, i) ...
}
```

Nested function definitions are permitted within functions in the places where variable definitions are allowed; that is, in any block, before the first statement in the block.

It is possible to call the nested function from outside the scope of its name by storing its address or passing the address to another function, as in the following declaration.

```
hack (int *array, int size)
{
 void store (int index, int value)
 { array[index] = value; }

 intermediate (store, size);
}
```

Within this block, the function `intermediate` receives the address of `store` as an argument. If `intermediate` calls `store`, the arguments given to `store` are used to store into `array`. But this technique works only so long as the containing function (`hack`, in the previous example) does not exit.

If you try to call the nested function through its address after the containing function has exited, all hell will break loose. If you try to call it after a containing scope level has exited, and if it refers to some of the variables that are no longer in scope, you may be lucky, but it's not wise to take the risk. If, however, the nested function does not



refer to anything that has gone out of scope, you should be safe.

GCC implements taking the address of a nested function using a technique called *trampolines*. A paper describing trampolines is available from ‘maya.idiapi.ch’, under ‘pub/tmb’, in ‘usenix88-lexic.ps.Z’.

A nested function can jump to a label inherited from a containing function, provided the label was explicitly declared in the containing function (see “Locally declared labels” on page 204). Such a jump returns instantly to the containing function, exiting the nested function which did the `goto` and any intermediate functions as well. The following is an example.

```
bar (int *array, int offset, int size)
{
 __label__ failure;
 int access (int *array, int index)
 {
 if (index > size)
 goto failure;
 return array[index + offset];
 }
 int i;
 ...
 for (i = 0; i < size; i++)
 ...access (array, i) ...
 ...
 return 0;

 /* Control comes here from access if it detects an error. */
failure:
 return -1;
}
```

A nested function always has internal linkage. Declaring one with `extern` is erroneous. If you need to declare the nested function before its definition, use `auto` (which is otherwise meaningless for function declarations).

```
bar (int *array, int offset, int size)
{
 __label__ failure;
 auto int access (int *, int);
 ...
 int access (int *array, int index)
 {
 if (index > size)
 goto failure;
 return array[index + offset];
 }
 ...
}
```

## Constructing function calls

Using the built-in functions described in the following, you can record the arguments a function received, and call another function with the same arguments, without knowing the number or types of the arguments.

You can also record the return value of that function call, and later return that value, without knowing what data type the function tried to return (as long as your caller expects that data type).

`__builtin_apply_args ()`

This built-in function returns a pointer of type `void *` to data describing how to perform a call with the same arguments as were passed to the current function.

The function saves the `arg` pointer register, structure value address, and all registers that might be used to pass arguments to a function into a block of memory allocated on the stack. Then it returns the address of that block.

`__builtin_apply (function, arguments, size)`

This built-in function invokes `function` (type `void (*)()`) with a copy of the parameters that `arguments` (type `void *`) and `size` (type `int`) describe.

The value of `arguments` should be the value returned by `__builtin_apply_args`. The argument `size` specifies the size of the stack argument data, in bytes.

This function returns a pointer of type `void *` to data describing how to return whatever value was returned by `function`. The data is saved in a block of memory allocated on the stack.

It is not always simple to compute the proper value for `size`. The value is used by `__builtin_apply` to compute the amount of data that should be pushed on the stack and copied from the incoming argument area.

`__builtin_return (result)`

This built-in function returns the value described by `result` from the containing function. You should specify, for `result`, a value returned by `__builtin_apply`.

## Naming an expression's type

You can give a name to the type of an expression using a `typedef` declaration with an initializer. To define *name* as a type name for the type of *exp*, use the following example's guide as input.

```
typedef name = exp;
```

This is useful in conjunction with the statements-within-expressions feature.

The following shows how the two together can be used to define a safe “*maximum*” macro that operates on any arithmetic type.

```
#define max(a,b) \
 ({typedef _ta = (a), _tb = (b); _ta _a = (a); _tb _b = (b); \
 _a > _b ? _a : _b; })
```

The reason for using names that start with underscores for the local variables is to avoid conflicts with variable names that occur within the expressions that are substituted for ‘a’ and ‘b’. Eventually we hope to design a new form of declaration syntax that allows you to declare variables whose scopes start only after their initializers; this will be a more reliable way to prevent such conflicts.

## Referring to a type with the `typeof` keyword

Another way to refer to the type of an expression is with `typeof`. The syntax of using of this keyword looks like `sizeof`, but the construct acts semantically like a type name defined with `typedef`.

There are two ways of writing the argument to `typeof`: with an expression, or with a type. The following is an example with an expression.

```
typeof (x[0](1))
```

This input assumes that `'x'` is an array of functions; the type described is that of the values of the functions. The following is an example with a type name as the argument.

```
typeof (int *)
```

In the following, the type described is that of pointers to `int`.

If you are writing a header file that must work when included in ANSI C programs, write `__typeof__` instead of `typeof`. See “Alternate keywords” on page 266. A `typeof`-construct can be used anywhere a `typedef` name could be used. For example, you can use it in a declaration, in a cast, or inside of `sizeof` or `typeof`.

```
typeof (*x) y;
```

- The following declares `'y'` as an array of such values.

```
typeof (*x) y[4];
```

- The following declares `'y'` as an array of pointers to characters.

```
typeof (typeof (char *)[4]) y;
```

It is equivalent to the traditional C declaration, `char *y[4];`.

To see the meaning of the declaration using `typeof`, and why it might be a useful way to write, let's rewrite it with the following macros.

```
#define pointer(T) typeof(T *)
#define array(T, N) typeof(T [N])
```

Now the declaration can be rewritten the following way.

```
array (pointer (char), 4) y;
```

Thus, `array (pointer (char), 4)` is the type of arrays of 4 pointers to `char`.

# Generalized lvalues

Compound expressions, conditional expressions and casts are allowed as *lvalues* provided their operands are lvalues. This means that you can take their addresses or store values into them.

Standard C++ allows compound expressions and conditional expressions as lvalues, and permits casts to reference type, so use of this extension is deprecated for C++ code.

For example, a compound expression can be assigned, provided the last expression in the sequence is an lvalue.

The following two expressions are equivalent.

```
(a, b) += 5
a, (b += 5)
```

Similarly, the address of the compound expression can be taken. The following two expressions are equivalent.

```
&(a, b)
a, &b
```

A conditional expression is a valid lvalue if its type is not void and the true and false branches are both valid lvalue. For example, the following two expressions are equivalent.

```
(a ? b : c) = 5
(a ? b = 5 : (c = 5))
```

A cast is a valid lvalue if its operand is an lvalue. A simple assignment whose left-hand side is a cast works by converting the right-hand side first to the specified type, then to the type of the inner left-hand side expression. After this is stored, the value is converted back to the specified type to become the value of the assignment. Thus, if 'a' has type `char *`, the following two expressions are equivalent.

```
(int)a = 5
(int)(a = (char *) (int)5)
```

An assignment-with-arithmetic operation such as `+=` applied to a cast performs the arithmetic using the type resulting from the cast, and then continues as in the previous case. Therefore, the following two expressions are equivalent.

```
(int)a += 5
(int)(a = (char *) (int) ((int)a + 5))
```

You cannot take the address of an lvalue cast, because the use of its address would not work out coherently. Suppose that `&(int)f` were permitted, where 'f' has type `float`. Then the following statement would try to store an integer bit-pattern where a floating point number belongs.

```
*&(int)f = 1;
```

This is quite different from what `(int)f = 1` would do—that would convert 1 to floating point and store it. Rather than cause this inconsistency, we think it is better to prohibit use of `&` on a cast. If you really do want an `int *` pointer with the address of `f`, you can simply write `(int *)&f`.

## Conditionals with omitted operands

The middle operand in a conditional expression may be omitted. Then if the first operand is nonzero, its value is the value of the conditional expression. Therefore, the following expression has the value of 'x' if that is nonzero; otherwise, the value of 'y'.

`x ? :y`

This example is perfectly equivalent to the following.

`x ? x : y`

In this simple case, the ability to omit the middle operand is not especially useful. When it becomes useful is when the first operand does, or may (if it is a macro argument), contain a side effect. Then repeating the operand in the middle would perform the side effect twice. Omitting the middle operand uses the value already computed without the undesirable effects of recomputing it.

## Double-word integers

GCC supports data types for integers that are twice as long as `int`. Simply write `long long int` for a signed integer, or `unsigned long long int` for an unsigned integer. To make an integer constant of type, `long long int`, add the suffix `LL` to the integer. To make an integer constant of type, `unsigned long long int`, add the suffix `ULL` to the integer.

You can use these types in arithmetic like any other integer types. Addition, subtraction, and bitwise Boolean operations on these types are open-coded on all types of machines. Multiplication is open-coded if the machine supports fullword-to-doubleword (a widening multiply instruction). Division and shifts are open-coded only on machines that provide special support.

The operations that are not open-coded use special library routines that come with GCC.

There may be pitfalls when you use `long long` types for function arguments, unless you declare function prototypes. If a function expects type `int` for its argument, and you pass a value of type `long long int`, confusion will result because the caller and the subroutine will disagree about the number of bytes for the argument. Likewise, if the function expects `long long int` and you pass `int`. The best way to avoid such problems is to use prototypes.



# Complex numbers

GCC supports complex data types. You can declare both complex integer types and complex floating types, using the keyword, `__complex__`. For example, `'__complex__double x;'` declares 'x' as a variable whose real part and imaginary part are both of type `double`.

`'__complex__short int y;'` declares 'y' to have real and imaginary parts of type `short int`; this is not likely to be useful, but it shows that the set of complex types is complete.

To write a constant with a complex data type, use the suffix 'i' or 'j' (either one; they are equivalent). For example, `'2.5fi'` has type, `__complex__float`, and `'3i'` has type, `__complex__int`. Such a constant always has a pure imaginary value, but you can form any complex value you like by adding one to a real constant.

To extract the real part of a complex-valued expression `exp`, write `__real__exp`. Likewise, use `__imag__` to extract the imaginary part.

The tilde operator, (`~`), performs complex conjugation when used on a value with a complex type.

GCC can allocate complex automatic variables in a noncontiguous fashion; it's even possible for the real part to be in a register while the imaginary part is on the stack (or vice-versa). None of the supported debugging info formats has a way to represent noncontiguous allocation like this, so GCC describes a noncontiguous complex variable as if it were two separate variables of noncomplex type. If the variable's actual name is `foo`, the two fictitious variables are named `foo$real` and `foo$imag`. You can examine and set these two fictitious variables with your debugger.

A future version of GDB will know how to recognize such pairs and treat them as a single variable with a complex type.

## Arrays of length zero

Zero-length arrays are allowed in GCC. They are very useful as the last element of a structure which is really a header for a variable-length object.

```
struct line {
 int length;
 char contents[0];
};

{
 struct line *thisline = (struct line *)
 malloc (sizeof (struct line) + this_length);
 thisline->length = this_length;
}
```

In standard C, you would have to give `contents` a length of 1, which means either you waste space or complicate the argument to `malloc`.

# Arrays of variable length

Variable-length automatic arrays are allowed in GCC. These arrays are declared like any other automatic arrays, but with a length that is not a constant expression. The storage is allocated at the point of declaration and deallocated when the brace-level is exited. Use the following for example.

```
FILE *
concat_fopen (char *s1, char *s2, char *mode)
{
 char str[strlen (s1) + strlen (s2) + 1];
 strcpy (str, s1);
 strcat (str, s2);
 return fopen (str, mode);
}
```

Jumping or breaking out of the scope of the array name deallocates the storage. Jumping into the scope is not allowed; you get an error message for it. You can use the function `alloca` to get an effect much like variable-length arrays. The function `alloca` is available in many other C implementations (but not in all).

On the other hand, variable-length arrays are more elegant.

There are other differences between these two methods. Space allocated with `alloca` exists until the containing *function* returns.

The space for a variable-length array is deallocated as soon as the array name's scope ends. (If you use both variable-length arrays and `alloca` in the same function, deallocation of a variable-length array will also deallocate anything more recently allocated with `alloca`.)

You can also use variable-length arrays as arguments to functions, as in the following example.

```
struct entry
tester (int len, char data[len][len])
{
 ...
}
```

The length of an array is computed once when the storage is allocated and is remembered for the scope of the array in case you access it with `sizeof`. If you want to pass the array first and the length afterward, you can use a forward declaration in the parameter list—another GNU extension.

```
struct entry
tester (int len; char data[len][len], int len)
{
 ...
}
```

The `'int len'` before the semicolon is a *parameter forward declaration*, and it serves

the purpose of making the name `len` known when the declaration of `data` is parsed. You can write any number of such parameter forward declarations in the parameter list. They can be separated by commas or semicolons, but the last one must end with a semicolon, which is followed by the “*real*” parameter declarations. Each forward declaration must match a “*real*” declaration in parameter name and data type.

# Macros with variable numbers of arguments

In GCC, a macro can accept a variable number of arguments, much as a function can. The syntax for defining the macro looks much like that used for a function. The following is an example.

```
#define eprintf(format, args...) \
 fprintf (stderr, format , ## args)
```

`args` is a rest argument: it takes in zero or more arguments, as many as the call contains. All of them plus the commas between them form the value of `args`, which is substituted into the macro body where `args` is used. Thus, we have the following expansion.

```
eprintf ("%s:%d: ", input_file_name, line_number)
?
fprintf (stderr, "%s:%d: " , input_file_name, line_number)
```

**NOTE:** The comma after the string constant comes from the definition of `eprintf`, whereas the last comma comes from the value of `args`.

The reason for using ‘##’ is to handle the case when `args` matches no arguments at all. In this case, `args` has an empty value. In this case, the second comma in the definition becomes an embarrassment: if it got through to the expansion of the macro, we would get something like the following.

```
fprintf (stderr, "success!\n" ,)
```

The previous example shows invalid C syntax. ‘##’ gets rid of the comma, so, instead, we get the following.

```
fprintf (stderr, "success!\n")
```

This is a special feature of the GCC preprocessor: ‘##’ before a rest argument that is empty discards the preceding sequence of non-whitespace characters from the macro definition. (If another macro argument precedes, none of it is discarded.) It might be better to discard the last preprocessor token instead of the last preceding sequence of non-whitespace characters; in fact, we may someday change this feature to do so. We advise you to write the macro definition so that the preceding sequence of non-whitespace characters is just a single token, so that the meaning will not change if we change the definition of this feature.

## Non-lvalue arrays may have subscripts

Subscripting is allowed on arrays that are not lvalues, even though the unary ‘&’ operator is not. For example, the following declaration is valid in GCC though not valid in other C dialects.

```
struct foo {int a[4];};

struct foo f();

bar (int index)
{
 return f().a[index];
}
```

## Arithmetic on `void`- and function-pointers

In GCC, addition and subtraction operations are supported on pointers to `void` and on pointers to functions. This is done by treating the size of a `void` or of a function as 1. A consequence of this is that `sizeof` is also allowed on `void` and on function types, and returns 1. The option, `-Wpointer-arith`, requests a warning if these extensions are used.

## Non-constant initializers

As in standard C++, the elements of an aggregate initializer for an automatic variable are not required to be constant expressions in GCC. The following is an example of an initializer with run-time varying elements:

```
foo (float f, float g)
{
 float beat_freqs[2] = { f-g, f+g };
 ...
}
```



# Constructor expressions

GCC supports constructor expressions. A constructor looks like a cast containing an initializer. Its value is an object of the type specified in the cast, containing the elements specified in the initializer.

Usually, the specified type is a structure. Assume that `struct foo` and `structure` are declared as shown in the following declaration.

```
struct foo {int a; char b[2];} structure;
```

The following is an example of constructing a `struct foo` with a constructor.

```
structure = ((struct foo) {x + y, 'a', 0});
```

The previous declaration is equivalent to writing the following input.

```
{
 struct foo temp = {x + y, 'a', 0};
 structure = temp;
}
```

You can also construct an array. If all the elements of the constructor are (made up of) simple constant expressions, suitable for use in initializers, then the constructor is an lvalue and can be coerced to a pointer to its first element, as shown in the following.

```
char **foo = (char *[]) { "x", "y", "z" };
```

Array constructors whose elements are not simple constants are not very useful, because the constructor is not an lvalue. There are only two valid ways to use it: to subscript it, or initialize an array variable with it. The former is probably slower than a `switch` statement, while the latter does the same thing an ordinary C initializer would do. The following is an example of subscripting an array constructor.

```
output = ((int[]) { 2, x, 28 }) [input];
```

Constructor expressions for scalar types and union types are also allowed, but then the constructor expression is equivalent to a cast.

## Labeled elements in initializers

Standard C requires the elements of an initializer to appear in a fixed order, the same as the order of the elements in the array or structure being initialized. In GCC you can give the elements in any order, specifying the array indices or structure field names they apply to. This extension is not implemented in GCC++. To specify an array index, write `'[index]'` or `'[index]='` before the element value. Use the following example.

```
int a[6] = { [4] 29, [2] = 15 };
```

The previous specification is equivalent to the following.

```
int a[6] = { 0, 0, 15, 0, 29, 0 };
```

The index values must be constant expressions, even if the array being initialized is automatic. To initialize a range of elements to the same value, write `'[ first...last]=value'`. Use the following example's input.

```
int widths[] = { [0 ... 9] = 1, [10 ... 99] = 2, [100] = 3 };
```

**NOTE:** The length of the array is the highest value specified + 1.

In a structure initializer, specify the name of a field to initialize with `'fieldname:'` before the element value.

For example, given `struct point { int x, y; };`, the initialization, `struct point p = { y: yvalue, x: xvalue };` is equivalent to the following input:

```
point p = { xvalue, yvalue struct };
```

Another syntax which has the same meaning is `'.fieldname='` as in the following input:

```
struct point p = { .y = yvalue, .x = xvalue };
```

You can also use an element label (with either the colon syntax or the period-equal syntax) when initializing a union, to specify which element of the union should be used. For example, the following will convert 4 to a double to store it in the union using the second line's element.

```
union foo { int i; double d; };
```

```
union foo f = { d: 4 };
```

By contrast, casting 4 to type `union foo` would store it into the union as the integer 'i', since it is an integer. (See "Cast to a union type" on page 225.)

You can combine this technique of naming elements with ordinary C initialization of successive elements. Each initializer element that does not have a label applies to the next consecutive element of the array or structure. For example, `int a[6] = { [1] = v1, v2, [4] = v4 };` is equivalent to the following input.

```
int a[6] = { 0, v1, v2, 0, v4, 0 };
```

Labeling the elements of an array initializer is especially useful when the indices are characters or belong to an `enum` type, as in the following example.

```
int whitespace[256] = { [' '] = 1, ['\t'] = 1, ['\h'] = 1, \
 ['\f'] = 1, ['\n'] = 1, ['\r'] = 1 };
```

## Case ranges

You can specify a range of consecutive values in a single case label, like `case low ... high:`. This has the same effect as the proper number of individual `case` labels, one for each integer value from *low* to *high*, inclusive. This feature is especially useful for ranges of ASCII character codes, as in `case 'A' ... 'Z':`.

**WARNING:** Write spaces around the `...`, for otherwise it may be parsed wrong when you use it with integer values. For example, use `case 1 ... 5:` rather than `case 1...5:`.

## Cast to a union type

A cast to union type is similar to other casts, except that the type specified is a union type. You can specify the type either with `union tag` or with a typedef name. A cast to union is actually a constructor though, not a cast, and hence does not yield an lvalue like normal casts. (See “Constructor expressions” on page 223.)

The types that may be cast to the union type are those of the members of the union. Thus, given the following union and variables:

```
union foo { int i; double d; };
int x;
double y;
```

both `'x'` and `'y'` can be cast to type `union foo`.

Using the cast as the right-hand side of an assignment to a variable of union type is equivalent to storing in a member of the union, like the following.

```
union foo u;
...
u = (union foo) x ≡ u.i = x
u = (union foo) y ≡ u.d = y
```

You can also use the union cast as a function argument:

```
void hack (union foo);
...
hack ((union foo) x);
```

## Declaring attributes of functions

In GCC, you declare certain things about functions called in your program which help the compiler optimize function calls and check your code more carefully.

The keyword, `__attribute__`, allows you to specify special attributes when making a declaration. This keyword is followed by an attribute specification inside double parentheses. Eight attributes, `noreturn`, `const`, `format`, `section`, `constructor`, `destructor`, `unused` and `weak` are currently defined for functions. Other attributes, including `section` are supported for variables declarations (see “Specifying attributes of variables” on page 235) and for types (see “Specifying attributes of types” on page 239).

You may also specify attributes with ‘`_`’ preceding and following each keyword. This allows you to use them in header files without being concerned about a possible macro of the same name. For example, you may use `__noreturn__` instead of `noreturn`.

`noreturn`

A few standard library functions, such as `abort` and `exit`, cannot return. GCC knows this automatically. Some programs define their own functions that never return. You can declare them `noreturn` to tell the compiler this fact. For example,

```
void fatal () __attribute__ ((noreturn));

void
fatal (...)
{
 ... /* Print error message.*/ ...
 exit (1);
}
```

The `noreturn` keyword tells the compiler to assume that `fatal` cannot return. It can then optimize without regard to what would happen if `fatal` ever did return. This makes slightly better code. More importantly, it helps avoid spurious warnings of uninitialized variables.

Do not assume that registers saved by the calling function are restored before calling the `noreturn` function. It does not make sense for a `noreturn` function to have a return type other than `void`. The attribute `noreturn` is not implemented in GCC versions earlier than 2.5.

An alternative way to declare that a function does not return, which works in the current version and in some older versions, is as follows:

```
typedef void voidfn ();
volatile voidfn fatal;
```

`const`

Many functions do not examine any values except their arguments, and have no effects except the return value. Such a function can be subject to common subexpression elimination and loop optimization just as an arithmetic operator would be. These functions should be declared with the attribute, `const`. For example, the following says that the hypothetical function, `square`, is safe to call fewer times than the program says.

```
int square (int) __attribute__ ((const));
```

The attribute `const` is not implemented in GCC versions earlier than 2.5. An alternative way to declare that a function has no side effects, which works in the current version and in some older versions, is as follows:

```
typedef int intfn ();
```

```
extern const intfn square;
```

This approach does not work in GCC++ from 2.6.0 on, since the language specifies that `const` must be attached to the return value.

**NOTE:** A function that has pointer arguments and examines the data pointed to must not be declared `const`. Likewise, a function that calls a non-`const` function usually must not be `const`. It does not make sense for a `const` function to return `void`.

`format (archetype, string-index, first-to-check)`

The `format` attribute specifies that a function takes `printf` or `scanf` style arguments which should be type-checked against a format string.

For example, the following declaration causes the compiler to check the arguments in calls to `my_printf` for consistency with the `printf` style format string argument `my_format`.

```
extern int
my_printf (void *my_object, const char *my_format, ...)
__attribute__ ((format (printf, 2, 3)));
```

The parameter `archetype` determines how the format string is interpreted, and should be either `printf` or `scanf`.

The parameter `string-index` specifies which argument is the format string argument (starting from 1), while `first-to-check` is the number of the first argument to check against the format string. For functions where the arguments are not available to be checked (such as `vprintf`), specify the third parameter as zero. In this case the compiler only checks the format string for consistency.

In the previous example, the format string (`my_format`) is the second argument of the function, `my_print`, and the arguments to check start with the third argument, so the correct parameters for the `format` attribute are 2 and 3.

The `format` attribute allows you to identify your own functions which take format strings as arguments, so that GCC can check the calls to these functions for errors. The compiler always checks formats for the ANSI library functions, `printf`, `fprintf`, `sprintf`, `scanf`, `fscanf`, `sscanf`, `vprintf`, `vfprintf` and `vsprintf` whenever such warnings are re-requested (using `-Wformat`), so there is no need to modify the header file, `stdio.h`.

`format_arg (string-index)`

The `format_arg` attribute specifies that a function takes `printf` or `scanf` style arguments, modifies it (for example, to translate it into another language), and passes it to a `printf` or `scanf` style function. For example, the following declaration causes the compiler to check the arguments in calls to `my_dgettext`, whose result is passed to a `printf` or `scanf` type function for consistency with the `printf` style format string argument, `my_format`.

```
extern char *
my_dgettext (char *my_domain, const char *my_format)
 __attribute__((format_arg (2)));
```

The parameter, `string-index`, specifies which argument is the format string argument (starting from 1).

The `format-arg` attribute allows you to identify your own functions which modify format strings, so that GCC can check the calls to `printf` and `scanf` function whose operands are a call to one of your own function. The compiler always treats `gettext`, `dgettext`, and `dcgettext` in this manner.

`section ("section-name")`

Normally, the compiler places the code it generates in the `text` section. Sometimes, however, you need additional sections, or you need certain particular functions to appear in special sections. The `section` attribute specifies that a function lives in a particular section. For example, the following declaration puts the function, `foobar`, in the `bar` section.

```
extern void foobar (void) __attribute__((section ("bar")));
```

Some file formats do not support arbitrary sections so the `section` attribute is not available on all platforms. If you need to map the entire contents of a module to a particular section, consider using the facilities of the linker instead.

`constructor`  
`destructor`

The `constructor` attribute causes the function to be called automatically before execution enters `main()`. Similarly, the `destructor` attribute causes the function to be called automatically after `main()` has completed or `exit()` has been called. Functions with these attributes are useful for initializing data that will be used implicitly during the execution of the program. These attributes are not currently implemented for Objective C.

**unused**

This attribute, attached to a function, means that the function is meant to be possibly unused. GCC will not produce a warning for this function. GCC++ does not currently support this attribute as definitions without parameters are valid in C++.

**weak**

The `weak` attribute causes the declaration to be emitted as a weak symbol rather than a global. This is primarily useful in defining library functions which can be overridden in user code, though it can also be used with non-function declarations. Weak symbols are supported for ELF targets, and also for a.out targets when using the GNU assembler and linker.

**alias ("target")**

The `alias` attribute causes the declaration to be emitted as an alias for another symbol, which must be specified. For instance, the following declares 'f' to be a weak alias for '\_\_f'. In C++, the mangled name for the target must be used.

```
void __f () { /* do something */; }
void f () __attribute__((weak, alias("__f")));
```

Not all target machines support this attribute.

**regparm (number)**

On the Intel 386, the `regparm` attribute causes the compiler to pass up to *number* integer arguments in registers *EAX*, *EDX*, and *ECX* instead of on the stack. Functions that take a variable number of arguments will continue to be passed all of their arguments on the stack.

**stdcall**

On the Intel 386, the `stdcall` attribute causes the compiler to assume that the called function will pop off the stack space used to pass arguments, unless it takes a variable number of arguments. The PowerPC compiler for Windows NT currently ignores the `stdcall` attribute.

**cdecl**

On the Intel 386, the `cdecl` attribute causes the compiler to assume that the called function will pop off the stack space used to pass arguments. This is useful to override the effects of the switch, '-mrtld'. The PowerPC compiler for Windows NT currently ignores the `cdecl` attribute.

**longcall**

On the RS/6000 and PowerPC, the `longcall` attribute causes the compiler to always call the function via a pointer, so that functions which reside further than 64 megabytes (67,108,864 bytes) from the current location can be called.

### `dllimport`

On the PowerPC running Windows NT, the `dllimport` attribute causes the compiler to call the function via a global pointer to the function pointer that is set up by the Windows NT dll library. The pointer name is formed by combining `__imp_` and the function name.

### `dllexport`

On the PowerPC running Windows NT, the `dllexport` attribute causes the compiler to provide a global pointer to the function pointer, so that it can be called with the `dllimport` attribute. The pointer name is formed by combining `__imp_` and the function name.

### `exception (except-func[, except-arg])`

On the PowerPC running Windows NT, the `exception` attribute causes the compiler to modify the structured exception table entry it emits for the declared function. The string or identifier, `except-func`, is placed in the third entry of the structured exception table. It represents a function which is called by the exception handling mechanism if an exception occurs. If it was specified, the string or identifier, `except-arg`, is placed in the fourth entry of the structured exception table.

### `function_vector`

Use this option on the H8/300 and H8/300H to indicate that the specified function should be called through the function vector. Calling a function through the function vector will reduce code size, however; the function vector has a limited size (maximum 128 entries on the H8/300 and 64 entries on the H8/300H) and shares space with the interrupt vector.

### `interrupt_handler`

Use this option on the H8/300 and H8/300H to indicate that the specified function is an interrupt handler. The compiler will generate function entry and exit sequences suitable for use in an interrupt handler when this attribute is present.

### `eightbit_data`

Use this option on the H8/300 and H8/300H to indicate that the specified variable should be placed into the eight bit data section. The compiler will generate more efficient code for certain operations on data in the eight bit data area. Note the eight bit data area is limited to 256 bytes of data.

### `tiny_data`

Use this option on the H8/300H to indicate that the specified variable should be placed into the tiny data section. The compiler will generate more efficient code for loads and stores on data in the tiny data section. Note the tiny data area is limited to slightly under 32kbytes of data.



**interrupt**

Use this option on the M32R/D to indicate that the specified function is an interrupt handler. The compiler will generate a function entry and exit sequences suitable for use in an interrupt handler when this attribute is present.

**model** (*model-name*)

Use this attribute on the M32R/D to set the addressability of an object, and the code generated for a function. The identifier, *model-name*, is one of `small`, `medium`, or `large`, representing each of the code models.

Small model objects live in the lower 16MB of memory (so that their addresses can be loaded with the `ld24` instruction), and are callable with the `b1` instruction.

Medium model objects may live anywhere in the 32 bit address space (the compiler will generate `seth/add3` instructions to load their addresses), and are callable with the `b1` instruction.

You can specify multiple attributes in a declaration by separating them by commas within the double parentheses or by immediately following an attribute declaration with another attribute declaration.

Some people object to the `__attribute__` feature, suggesting that ANSI C's `#pragma` should be used instead.

There are two reasons for not doing this:

- It is impossible to generate `#pragma` commands from a macro.
- There is no telling what the same `#pragma` might mean in another compiler.

These two reasons apply to almost any application that might be proposed for `#pragma`. It is basically a mistake to use `#pragma` for anything.

# Prototypes and old-style function definitions

GCC extends ANSI C to allow a function prototype to override a later old-style non-prototype definition. Consider the following example.

```

/* Use prototypes unless the compiler is old-fashioned. */
#if __STDC__
#define P(x) x
#else
#define P(x) ()
#endif

/* Prototype function declaration. */
int isroot P((uid_t));

/* Old-style function definition. */
int isroot (x) /* ??? lossage here ??? */
 uid_t x;
{
 return x == 0;
}

```

Suppose the type `uid_t` happens to be `short`. ANSI C does not allow this example, because subword arguments in old-style non-prototype definitions are promoted. Therefore in this example the function definition's argument is really an `int`, which does not match the prototype argument type of `short`.

This restriction of ANSI C makes it hard to write code that is portable to traditional C compilers, because the programmer does not know whether the `uid_t` type is `short`, `int`, or `long`.

Therefore, in cases like these GCC allows a prototype to override a later old-style definition. More precisely, in GCC, a function prototype argument type overrides the argument type specified by a later old-style definition if the former type is the same as the latter type before promotion. Thus, in GCC, the previous example is equivalent to the following declaration.

```

int isroot (uid_t);
int
isroot (uid_t x)
{
 return x == 0;
}

```

GNU C++ does not support old-style function definitions, so the previous example's extension is irrelevant.

# Compiling functions for interrupt calls

When compiling code for certain platforms (currently the Hitachi H8/300 and the Tandem ST-2000), you can instruct {No value for "GCC"} that certain functions are meant to be called from hardware interrupts.

To mark a function as callable from interrupt, include the line `#pragma interrupt` somewhere before the beginning of the function's definition. (For maximum readability, you might place it immediately before the definition of the appropriate function.) `#pragma interrupt` will affect only the next function defined; if you want to define more than one function with this property, include `#pragma interrupt` before each of them.

When you define a function with `#pragma interrupt`, {No value for "GCC"} alters its usual calling convention, to provide the right environment when the function is called from an interrupt. *Such functions cannot be called in the usual way from your program.*

You must use other facilities to actually associate these functions with particular interrupts; {No value for "GCC"} can only compile them in the appropriate way.

## C++ style comments

In GCC, you may use C++ style comments, which start with `'//'` and continue until the end of the line. Many other C implementations allow such comments, and they are likely to be in a future C standard. However, C++ style comments are not recognized if you specify `-ansi` or `-traditional`, since they are incompatible with traditional constructs like `dividend/*comment*/divisor`.

## Dollar signs in identifier names

In GCC, you may use dollar signs in identifier names. This is because many traditional C implementations allow such identifiers. However, dollar signs are not supported on a few target machines, typically because the target assembler does not allow them.

## The character ESC in constants

You can use the sequence, `\e`, in a string or character constant to stand for the ASCII character, *ESC*.

## Inquiring on alignment of types or variables

The keyword `__alignof__` allows you to inquire about how an object is aligned, or the minimum alignment usually required by a type. Its syntax is just like `sizeof`.

For example, if the target machine requires a double value to be aligned on an 8-byte boundary, then `__alignof__ (double)` is 8. This is true on many RISC machines.

On more traditional machine designs, `__alignof__ (double)` is 4 or even 2.

Some machines never actually require alignment; they allow reference to any data type even at an odd addresses. For these machines, `__alignof__` reports the *recommended* alignment of a type.

When the operand of `__alignof__` is an lvalue rather than a type, the value is the largest alignment that the lvalue is known to have. It may have this alignment as a result of its data type, or because it is part of a structure and inherits alignment from that structure. For example, after this declaration:

```
struct foo { int x; char y; } foo1;
```

The value of `__alignof__ (foo1.y)` is probably 2 or 4, the same as `__alignof__ (int)`, even though the data type of `foo1.y` does not itself demand any alignment. A related feature which lets you specify the alignment of an object is `__attribute__((aligned ( alignment)))`; see “Specifying attributes of variables” on page 235.

# Specifying attributes of variables

The keyword, `__attribute__`, allows you to specify special attributes of variables or structure fields. This keyword is followed by an attribute specification inside double parentheses. Eight attributes are currently defined for variables: `aligned`, `mode`, `nocommon`, `packed`, `section`, `transparent_union`, `unused`, and `weak`. Other attributes are available for functions (see “Declaring attributes of functions” on page 226) and for types (see “Specifying attributes of types” on page 239). You may also specify attributes with ‘`_`’ preceding and following each keyword. This allows you to use them in header files without being concerned about a possible macro of the same name. For example, you may use `__aligned__` instead of `aligned`.

`aligned (alignment)`

This attribute specifies a minimum alignment for the variable or structure field, measured in bytes. For example, the following declaration causes the compiler to allocate the global variable `x` on a 16-byte boundary.

```
int x __attribute__ ((aligned (16))) = 0;
```

On a 68040, this could be used in conjunction with an `asm` expression to access the `move16` instruction which requires 16-byte aligned operands. You can also specify the alignment of structure fields.

For example, to create a double-word aligned `int` pair, you could write:

```
struct foo { int x[2] __attribute__ ((aligned (8))); };
```

This is an alternative to creating a union with a `double` member that forces the union to be double-word aligned. It is not possible to specify the alignment of functions; the alignment of functions is determined by the machine’s requirements and cannot be changed.

You cannot specify alignment for a `typedef` name because such a name is just an alias, not a distinct type.

As in the preceding examples, you can explicitly specify the alignment (in bytes) that you wish the compiler to use for a given variable or structure field.

Alternatively, you can leave out the alignment factor and just ask the compiler to align a variable or field to the maximum useful alignment for the target machine you are compiling for. For example, you could write:

```
short array[3] __attribute__ ((aligned));
```

Whenever you leave out the alignment factor in an `aligned` attribute specification, the compiler automatically sets the alignment for the declared variable or field to the largest alignment which is ever used for any data type on the target machine you are compiling for. Doing this can often make copy operations more efficient, because the compiler can use whatever instructions

copy the biggest chunks of memory when performing copies to or from the variables or fields that you have aligned this way.

The `aligned` attribute can only increase the alignment; but you can decrease it by specifying `packed` as well. (See attribute specifications for `packed`.)

**NOTE:** The effectiveness of `aligned` attributes may be limited by inherent limitations in your linker. On many systems, the linker is only able to arrange for variables to be aligned up to a certain maximum alignment. (For some linkers, the maximum supported alignment may be very small.)

If your linker is only able to align variables up to a maximum of 8 byte alignment, then specifying `aligned(16)` in an `__attribute__` will still only provide you with 8 byte alignment.

See *Using LD* in **GNUPro Utilities** for further information.

`mode (mode)`

This attribute specifies the data type for the declaration— whichever type corresponds to the mode `mode`. This in effect lets you request an integer or floating point type according to its width.

You may also specify a mode of `byte` or `__byte__` to indicate the mode corresponding to a one-byte integer, `word` or `__word__` for the mode of a one-word integer, and `pointer` or `__pointer__` for the mode used to represent pointers.

`nocommon`

This attribute specifies requests GCC not to place a variable “*common*” but instead to allocate space for it directly. If you specify the ‘`-fno-common`’ flag, GCC will do this for all variables.

Specifying the ‘`nocommon`’ attribute for a variable provides an initialization of zeros. A variable may only be initialized in one source file.

`packed`

The `packed` attribute specifies that a variable or structure field should have the smallest possible alignment—one byte for a variable, and one bit for a field, unless you specify a larger value with the `aligned` attribute. The following example is a structure in which the field ‘`x`’ is packed, so that it immediately follows ‘`a`’:

```
struct foo
{
 char a;
 int x[2] __attribute__((packed));
};
```

```
section ("section-name")
```

Normally, the compiler places the objects it generates in sections like `data` and `bss`. Sometimes, however, you need additional sections, or you need certain particular variables to appear in special sections, for example to map to special hardware.

The `section` attribute specifies that a variable (or function) lives in a particular section. For example, the following small program uses several specific `section` names.

```
struct duart a __attribute__((section("DUART_A"))) = { 0 };
struct duart b __attribute__((section("DUART_B"))) = { 0 };
char stack[10000] __attribute__((section("STACK"))) = {0 };
int init_data_copy __attribute__((section("INITDATACOPY"))) = 0;
main()
{
 /* Initialize stack pointer */
 init_sp (stack + sizeof (stack));

 /* Initialize initialized data */
 memcpy (&init_data_copy, &data, &edata - &data);

 /* Turn on the serial ports */
 init_duart (&a); init_duart (&b);
}
```

Use the `section` attribute with an *initialized* definition of a *global* variable, as shown in the previous example. GCC issues a warning and otherwise ignores the `section` attribute in uninitialized variable declarations.

You may only use the `section` attribute with a fully initialized global definition because of the way linkers work. The linker requires each object be defined once, with the exception that uninitialized variables tentatively go in the `common` (or `bss`) section and can be multiply-defined. You can force a variable to be initialized with the `-fno-common` flag or the `nocommon` attribute. Some file formats do not support arbitrary sections so the `section` attribute is not available on all platforms. If you need to map the entire contents of a module to a particular section, consider using the facilities of the linker instead.

```
transparent_union
```

This attribute, attached to a function parameter which is a union, means that the corresponding argument may have the type of any union member, but the argument is passed as if its type were that of the first union member. For more details, see “Specifying attributes of types” on page 239. You can also use this attribute on a `typedef` for a union data type; then it applies to all function parameters with that type.

### `unused`

This attribute, attached to a variable, means that the variable is meant to be possibly unused. GCC will not produce a warning for this variable.

### `weak`

See the descriptions for the `weak` attribute with “Declaring attributes of functions” on page 226.

### `model ( model-name )`

Use this attribute on the M32R/D to set the addressability of an object. The identifier *model-name* is one of `small`, `medium`, or `large`, representing each of the code models.

Small model objects live in the lower 16MB of memory (so that their addresses can be loaded with the `ld24` instruction).

Medium and large model objects may live anywhere in the 32 bit address space (the compiler will generate `seth/add3` instructions to load their addresses).

To specify multiple attributes, separate them by commas within double parentheses; for example, `__attribute__((aligned (16), packed))`.



## Specifying attributes of types

The keyword `__attribute__` allows you to specify special attributes of `struct` and `union` types when you define such types. This keyword is followed by an attribute specification inside double parentheses. Three attributes are currently defined for types: `aligned`, `packed`, and `transparent_union`. Other attributes are defined for functions (see “Declaring attributes of functions” on page 226) and for variables (see “Specifying attributes of variables” on page 235).

You may also specify any one of these attributes with ‘`__`’ preceding and following its keyword. This allows you to use these attributes in header files without being concerned about a possible macro of the same name. For example, you may use `__aligned__` instead of `aligned`.

You may specify the `aligned` and `transparent_union` attributes either in a `typedef` declaration or just past the closing curly brace of a complete `enum`, `struct` or `union` type definition and the `packed` attribute only past the closing brace of a definition.

`aligned (alignment)`

This attribute specifies a minimum alignment (in bytes) for variables of the specified type. For example, the following declarations force the compiler to insure (as fast as it can) that each variable whose type is `struct S` or `more_aligned_int` will be allocated and aligned *at least* on an 8-byte boundary.

```
struct S { short f[3]; } __attribute__((aligned (8)));
typedef int more_aligned_int __attribute__((aligned (8)));
```

On a Sparc, having all variables of type `struct S` aligned to 8-byte boundaries allows the compiler to use the `ldd` and `std` (doubleword load and store) instructions when copying one variable of type `struct S` to another, thus improving run-time efficiency.

**NOTE:** The alignment of any given `struct` or `union` type is required by the ANSI C standard to be at least a perfect multiple of the lowest common multiple of the alignments of all of the members of the `struct` or `union` in question. This means that you can effectively adjust the alignment of a `struct` or `union` type by attaching an `aligned` attribute to any one of the members of such a type, but the notation illustrated in the last example is a more obvious, intuitive, and readable way to request the compiler to adjust the alignment of an entire `struct` or `union` type.

As in the preceding example, you can explicitly specify the alignment (in bytes) that you wish the compiler to use for a given `struct` or `union` type. Alternatively, you can leave out the alignment factor and just ask the compiler to align a type to the maximum alignment for the target machine for which you’re compiling:

```
struct S { short f[3]; } __attribute__((aligned));
```

Whenever you leave out the alignment factor in an `aligned` attribute specification, the compiler automatically sets the alignment for the type to the largest alignment which is ever used for any data type on the target machine you are compiling for. Doing this can often make copy operations more efficient, because the compiler can use whatever instructions copy the biggest chunks of memory when performing copies to or from the variables which have types that you have aligned this way. In the example above, if the size of each `short` is 2 bytes, then the size of the entire `struct S` type is 6 bytes. The smallest power of two which is greater than or equal to that is 8, so the compiler sets the alignment for the entire `struct S` type to 8 bytes.

**NOTE:** Although you can ask the compiler to select a time-efficient alignment for a given type and then declare only individual stand-alone objects of that type, the compiler's ability to select a time-efficient alignment is primarily useful only when you plan to create arrays of variables having the relevant (efficiently aligned) type. If you declare or use arrays of variables of an efficiently-aligned type, then it is likely that your program will also be doing pointer arithmetic (or subscripting, which amounts to the same thing) on pointers to the relevant type, and the code that the compiler generates for these pointer arithmetic operations will often be more efficient for efficiently-aligned types than for other types.

The `aligned` attribute can only increase the alignment; but you can decrease it by specifying `packed` as well. See the following discussion for `packed`.

**NOTE:** The effectiveness of `aligned` attributes may be limited by inherent limitations in your linker. On many systems, the linker is only able to arrange for variables to be aligned up to a certain maximum alignment. (For some linkers, the maximum supported alignment may be very very small.) If your linker is only able to align variables up to a maximum of 8 byte alignment, then specifying `aligned(16)` in an `__attribute__` will still only provide you with 8 byte alignment. See *Using LD* in **GNUPro Utilities** for more information on alignment.

`packed`

This attribute, attached to an `enum`, `struct`, or `union` type definition, specified that the minimum required memory be used to represent the type. Specifying this attribute for `struct` and `union` types is equivalent to specifying the `packed` attribute on each of the structure or union members.

Specifying the `'-fshort-enums'` flag on the line is equivalent to specifying the `packed` attribute on all `enum` definitions.

You may only specify this attribute after a closing curly brace on an `enum` definition, not in a `typedef` declaration, unless that declaration also contains the definition of the `enum`.

#### `transparent_union`

This attribute, attached to a `union` type definition, indicates that any function parameter having that `union` type causes calls to that function to be treated in a special way.

First, the argument corresponding to a `transparent union` type can be of any type in the union; no cast is required. Also, if the union contains a pointer type, the corresponding argument can be a null pointer constant or a void pointer expression; and if the union contains a void pointer type, the corresponding argument can be any pointer expression. If the union member type is a pointer, qualifiers like `const` on the referenced type must be respected, just as with normal pointer conversions.

Second, the argument is passed to the function using the calling conventions of first member of the transparent union, not the calling conventions of the union itself. All members of the union must have the same machine representation; this is necessary for this argument passing to work properly.

Transparent unions are designed for library functions that have multiple interfaces for compatibility reasons. For example, suppose the `wait` function must accept either a value of type `int *` to comply with Posix, or a value of type `'union wait *'` to comply with the 4.1 BSD interface. If the `wait` function's parameter were `'void *'`, `wait` would accept both kinds of arguments, but it would also accept any other pointer type and this would make argument type checking less useful. Instead, `<sys/wait.h>` might define the interface as follows.

```
typedef union
{
 int *__ip;
 union wait *__up;
} wait_status_pointer_t __attribute__((transparent_union));

pid_t wait (wait_status_pointer_t);
```

This interface allows either `'int *'` or `'union wait *'` arguments to be passed, using the `'int *'` calling convention. The program can call `wait` with arguments of either of the following types.

```
int w1 () { int w; return wait (&w); }
int w2 () { union wait w; return wait (&w); }
```

With this interface, the `wait` implementation might look like the following example's declaration.

```
pid_t wait (wait_status_pointer_t p)
{
 return waitpid (-1, p.__ip, 0);
}
```

To specify multiple attributes, separate them by commas within the double parentheses, as in the following example.

```
__attribute__ ((aligned (16), packed))
```

# An inline function is as fast as a macro

By declaring a function `inline`, you can direct GCC to integrate that function's code into the code for its callers. This makes execution faster by eliminating the function-call overhead; in addition, if any of the actual argument values are constant, their known values may permit simplifications at compile time so that not all of the inline function's code needs to be included.

The effect on code size is less predictable; object code may be larger or smaller with function inlining, depending on the particular case. Inlining of functions is an optimization and it really “works” only in optimizing compilation. If you don't use ‘-o’, no function is really inline.

To declare a function inline, use the `inline` keyword in its declaration, like the following example shows.

```
inline int
inc (int *a)
{
 (*a)++;
}
```

(If you are writing a header file to be included in ANSI C programs, write `__inline__` instead of `inline`. See “Alternate keywords” on page 266.)

You can also make all “*simple enough*” functions inline with the option ‘-finline-functions’. Certain usage in a function definition can make it unsuitable for inline substitution.

**NOTE:** In C and Objective C, unlike C++, the `inline` keyword does not affect the linkage of the function.

GCC automatically inlines member functions defined within the class body of C++ programs even if they are not explicitly declared `inline`. (You can override this with ‘-fno-default-inline’; see “Options controlling C++ dialect” on page 85.)

When a function is both `inline` and `static`, if all calls to the function are integrated into the caller, and the function's address is never used, then the function's own assembler code is never referenced. In this case, GCC does not actually output assembler code for the function, unless you specify the option, ‘-fkeep-inline-functions’. Some calls cannot be integrated for various reasons (in particular, calls that precede the function's definition cannot be integrated, and neither can recursive calls within the definition). If there is a nonintegrated call, then the function is compiled to assembler code as usual. The function must also be compiled as usual if

the program refers to its address, because that can't be inlined.

When an `inline` function is not static, then the compiler must assume that there may be calls from other source files; since a global symbol can be defined only once in any program, the function must not be defined in the other source files, so the calls therein cannot be integrated. Therefore, a non-static inline function is always compiled on its own in the usual fashion.

If you specify both `inline` and `extern` in the function definition, then the definition is used only for inlining. In no case is the function compiled on its own, not even if you refer to its address explicitly. Such an address becomes an external reference, as if you had only declared the function, and had not defined it.

This combination of `inline` and `extern` has almost the effect of a macro. The way to use it is to put a function definition in a header file with these keywords, and put another copy of the definition (lacking `inline` and `extern`) in a library file. The definition in the header file will cause most calls to the function to be inlined. If any uses of the function remain, they will refer to the single copy in the library.

GCC does not `inline` any functions when not optimizing. It is not clear whether it is better to `inline` or not, in this case, but we found that a correct implementation when not optimizing was difficult. So we did the easy thing, and turned it off.

# Assembler instructions with C expression operands

In an assembler instruction, using `asm`, you can now specify the operands of the instruction using C expressions. This means no more guessing which registers or memory locations will contain the data you want to use.

You must specify an assembler instruction template much like what appears in a machine description, plus an operand constraint string for each operand.

For example, here is how to use the 68881's `fsinx` instruction:

```
asm ("fsinx %1,%0" : "=f" (result) : "f" (angle));
```

`angle` is the C expression for the input operand while `result` is that of the output operand. Each has `'f'` as its operand constraint, saying that a floating point register is required. The `'='` in `'=f'` indicates that the operand is an output; all output operands' constraints must use `'='`. The constraints use the same language used in the machine description (see “Constraints for `asm` operands” on page 249).

Each operand is described by an operand-constraint string followed by the C expression in parentheses. A colon separates the assembler template from the first output operand, and another separates the last output operand from the first input, if any. Commas separate output operands and separate inputs. The total number of operands is limited to ten or to the maximum number of operands in any instruction pattern in the machine description, whichever is greater.

If there are no output operands, and there are input operands, then there must be two consecutive colons surrounding the place where the output operands would go.

Output operand expressions must be lvalues; the compiler can check this. The input operands need not be lvalues. The compiler cannot check whether the operands have data types that are reasonable for the instruction being executed. It does not parse the assembler instruction template and does not know what it means, or whether it is valid assembler input. The extended `asm` feature is most often used for machine instructions that the compiler itself does not know exist. If the output expression cannot be directly addressed (for example, it is a bit field), your constraint must allow a register. In that case, GCC will use the register as the output of the `asm`, and then store that register into the output.

The output operands must be write-only; GCC will assume that the values in these operands before the instruction are dead and need not be generated. Extended `asm` supports input-output or read-write operands.

Use the constraint character, `'+'`, to indicate such an operand, and list it with the output operands.

When the constraints for the read-write operand (or an operand in which only some of the bits are to be changed) allows a register, you may, as an alternative logically split its function into two separate operands, one input operand and one write-only output operand. The connection between them is expressed by constraints which say they need to be in the same location when the instruction executes. You can use the same C expression for both operands, or different expressions. For example, in the following declaration, we write the (fictitious) ‘combine’ instruction with `bar` as its read-only source operand and `foo` as its read-write destination.

```
asm ("combine %2,%0" : "=r" (foo) : "0" (foo), "g" (bar));
```

The constraint ‘0’ for operand 1 says that it must occupy the same location as operand ‘0’. A digit in constraint is allowed only in an input operand, and it must refer to an output operand. Only a digit in the constraint can guarantee that one operand will be in the same place as another. The mere fact that `foo` is the value of both operands is not enough to guarantee that they will be in the same place in the generated assembler code. The following declaration would not work.

```
asm ("combine %2,%0" : "=r" (foo) : "r" (foo), "g" (bar));
```

Various optimizations or reloading could cause operands 0 and 1 to be in different registers; GCC knows no reason not to do so. For example, the compiler might find a copy of the value of `foo` in one register and use it for operand 1, but generate the output operand 0 in a different register (copying it afterward to `foo`’s own address). Of course, since the register for operand 1 is not even mentioned in the assembler code, the result will not work, but GCC can’t tell that. Some instructions clobber specific hard registers. To describe this, write a third colon after the input operands, followed by the names of the clobbered hard registers (given as strings). The following is a realistic example for the Vax.

```
asm volatile ("movc3 %0,%1,%2"
: /* no outputs */
: "g" (from), "g" (to), "g" (count)
: "r0", "r1", "r2", "r3", "r4", "r5");
```

If you refer to a particular hardware register from the assembler code, then you will probably have to list the register after the third colon to tell the compiler that the register’s value is modified.

In many assemblers, the register names begin with ‘%’; to produce one ‘%’ in the assembler code, you must write ‘%%’ in the input. If your assembler instruction can alter the condition code register, add ‘cc’ to the list of clobbered registers.

GCC on some machines represents the condition codes as a specific hardware register; ‘cc’ serves to name this register. On other machines, the condition code is handled differently, and specifying ‘cc’ has no effect. But it is valid no matter what the machine. If your assembler instruction modifies memory in an unpredictable fashion, add `memory` to the list of clobbered registers. This will cause GCC to not keep memory



values cached in registers across the assembler instruction. You can put multiple assembler instructions together in a single `asm` template, separated either with newlines (written as: `'\n'`) or with semi-colons if the assembler allows such semicolons. The GNU assembler allows semicolons and all UNIX assemblers seem to do so. The input operands are guaranteed not to use any of the clobbered registers, and neither will the output operands' addresses, so you can read and write the clobbered registers as many times as you like. The following is an example of multiple instructions in a template, assuming that the subroutine, `_foo`, accepts arguments in registers 9 and 10.

```
asm ("movl %0,r9;movl %1,r10;call _foo"
 : /* no outputs */
 : "g" (from), "g" (to)
 : "r9", "r10");
```

Unless an output operand has the `'&'` constraint modifier, GCC may allocate it in the same register as an unrelated input operand, on the assumption that the inputs are consumed before the outputs are produced. This assumption may be false if the assembler code actually consists of more than one instruction. In such a case, use `'&'` for each output operand that may not overlap an input. See “Constraint modifier characters” on page 252.

If you want to test the condition code produced by an assembler instruction, you must include a branch and a label in the `asm` construct, as follows.

This assumes your assembler supports local labels, as the GNU assembler and most UNIX assemblers do.

```
asm ("clr %0;frob %1;beq 0f;mov #1,%0;0:"
 : "g" (result)
 : "g" (input));
```

Speaking of labels, jumps from one `asm` to another are not supported. The compiler's optimizers do not know about these jumps, and therefore they cannot take account of them when deciding how to optimize.

Usually the most convenient way to use these `asm` instructions is to encapsulate them in macros that look like functions, as in the following example.

```
#define sin(x) ({ double __value, __arg = (x); \
 asm ("fsinx %1,%0": "=f" (__value): "f" (__arg)); \
 __value; })
```

In the previous example, the variable, `__arg`, is used to make sure that the instruction operates on a proper `double` value, and to accept only those arguments `'x'` which can convert automatically into a `double`.

Another way to make sure the instruction operates on the correct data type is to use a cast in the `asm`. This is different from using a variable `__arg` in that it converts more different types. For example, if the desired type were `int`, casting the argument to `int`

would accept a pointer with no complaint, while assigning the argument to an `int` variable named `__arg` would warn about using a pointer unless the caller explicitly casts it.

If an `asm` has output operands, GCC assumes for optimization purposes that the instruction has no side effects except to change the output operands. This does not mean that instructions with a side effect cannot be used, but you must be careful, because the compiler may eliminate them if the output operands aren't used, or move them out of loops, or replace two with one if they constitute a common subexpression. Also, if your instruction does have a side effect on a variable that otherwise appears not to change, the old value of the variable may be reused later if it happens to be found in a register.

You can prevent an `asm` instruction from being deleted, moved significantly, or combined, by writing the keyword `volatile` after the `asm`.

For example:

```
#define set_priority(x) \
 asm volatile ("set_priority %0": /* no outputs */ : "g" (x))
```

An instruction without output operands will not be deleted or moved significantly, regardless, unless it is unreachable.

**NOTE:** Even a `volatile asm` instruction can be moved in ways that appear insignificant to the compiler, such as across jump instructions. You can't expect a sequence of `volatile asm` instructions to remain perfectly consecutive. If you want consecutive output, use a single `asm`.

It is a natural idea to look for a way to give access to the condition code left by the assembler instruction. However, when we attempted to implement this, we found no way to make it work reliably. The problem is that output operands might need reloading, which would result in additional following “*store*” instructions. On most machines, these instructions would alter the condition code before there was time to test it. This problem doesn't arise for ordinary “*test*” and “*compare*” instructions because they don't have any output operands.

If you are writing a header file that should be included in ANSI C programs, write `__asm__` instead of `asm`. See “Alternate keywords” on page 266.

# Constraints for `asm` operands

The following details discuss what constraint letters you can use with `asm` operands. Constraints can say whether an operand may be in a register, and which kinds of register; whether the operand can be a memory reference, and which kinds of address; whether the operand may be an immediate constant, and which possible values it may have. Constraints can also require two operands to match.

## Simple constraints

The simplest kind of constraint is a string full of letters, each of which describes one kind of operand that is permitted. Here are the letters that are allowed:

`m`  
A memory operand is allowed, with any kind of address that the machine supports in general.

`o`  
A memory operand is allowed, but only if the address is *offsettable*. This means that adding a small integer (actually, the width in bytes of the operand, as determined by its machine mode) may be added to the address and the result is also a valid memory address.

For example, an address which is constant is offsettable; so is an address that is the sum of a register and a constant (as long as a slightly larger constant is also within the range of address-offsets supported by the machine); but an autoincrement or autodecrement address is not offsettable. More complicated indirect/indexed addresses may or may not be offsettable depending on the other addressing modes that the machine supports.

**NOTE:** In an output operand which can be matched by another operand, the constraint letter ‘`o`’ is valid only when accompanied by both ‘`<`’ (if the target machine has predecrement addressing) and ‘`>`’ (if the target machine has preincrement addressing).

`v`  
A memory operand that is not offsettable. In other words, anything that would fit the ‘`m`’ constraint but not the ‘`o`’ constraint.

`<`  
A memory operand with autodecrement addressing (either predecrement or postdecrement) is allowed.

`>`  
A memory operand with autoincrement addressing (either preincrement or postincrement) is allowed.

- `r`  
A register operand is allowed provided that it is in a general register.
- `d, a, f, ...`  
Other letters can be defined in machine-dependent fashion to stand for particular classes of registers. ‘`d`’, ‘`a`’ and ‘`f`’ are defined on the 68000/68020 to stand for data, address and floating point registers.
- `i`  
An immediate integer operand (one with constant value) is allowed. This includes symbolic constants whose values will be known only at assembly time.
- `n`  
An immediate integer operand with a known numeric value is allowed. Many systems cannot support assembly-time constants for operands less than a word wide. Constraints for these operands should use ‘`n`’ rather than ‘`i`’.
- `I, J, K, ... P`  
Other letters in the range ‘`I`’ through ‘`P`’ may be defined in a machine-dependent fashion to permit immediate integer operands with explicit integer values in specified ranges. For example, on the 68000, ‘`I`’ is defined to stand for the range of values 1 to 8. This is the range permitted as a shift count in the shift instructions.
- `E`  
An immediate floating operand (expression code `const_ double`) is allowed, but only if the target floating point format is the same as that of the host machine (on which the compiler is running).
- `F`  
An immediate floating operand (expression code `const_ double`) is allowed.
- `G, H`  
‘`G`’ and ‘`H`’ are defined in a machine-dependent fashion to permit immediate floating operands in particular ranges of values.
- `s`  
An immediate integer operand whose value is not an explicit integer is allowed. This might appear strange; if an `insn` allows a constant operand with a value not known at compile time, it certainly must allow any known value. So why use ‘`s`’ instead of ‘`i`’? Sometimes it allows better code to be generated. For example, on the 68000 in a fullword instruction it is possible to use an immediate operand; but if the immediate value is between -128 and 127, better code results from loading the value into a register and using the register. This is because the load into the register can be done with a `moveq` instruction. We arrange for this to happen by defining the letter ‘`K`’ to mean “any integer outside the range -128 to 127”, and then specifying ‘`Ks`’ in the operand constraints.

- g** Any register, memory or immediate integer operand is allowed, except for registers that are not general registers.
- x** Any operand whatsoever is allowed.
- 0, 1, 2, ... 9** An operand that matches the specified operand number is allowed. If a digit is used together with letters within the same alternative, the digit should come last. This is called a matching constraint and what it really means is that the assembler has only a single operand that fills two roles which `asm` distinguishes. For example, an `add` instruction uses two input operands and an output operand, but on most CISC machines an `add` instruction really has only two operands, one of them an input-output operand:
- ```
addl #35,r12
```
- Matching constraints are used in these circumstances. More precisely, the two operands that match must include one input-only operand and one output-only operand. Moreover, the digit must be a smaller number than the number of the operand that uses it in the constraint.
- p** An operand that is a valid memory address is allowed. This is for “load address” and “push address” instructions. ‘p’ in the constraint must be accompanied by `address_` operand as the predicate in the `match_operand`. This predicate interprets the mode specified in the `match_operand` as the mode of the memory reference for which the address would be valid.
- Q, R, S, ... U** Letters in the range ‘Q’ through ‘U’ may be defined in a machine-dependent fashion to stand for arbitrary operand types.

Multiple alternative constraints

Sometimes a single instruction has multiple alternative sets of possible operands. For example, on the 68000, a logical-or instruction can combine register or an immediate value into memory, or it can combine any kind of operand into a register; but it cannot combine one memory location into another.

These constraints are represented as multiple alternatives. An alternative can be described by a series of letters for each operand. The overall constraint for an operand is made from the letters for this operand from the first alternative, a comma, the letters for this operand from the second alternative, a comma, and so on until the last alternative.

If all the operands fit any one alternative, the instruction is valid. Otherwise, for each

alternative, the compiler counts how many instructions must be added to copy the operands so that that alternative applies. The alternative requiring the least copying is chosen. If two alternatives need the same amount of copying, the one that comes first is chosen. These choices can be altered with the ‘?’ and ‘!’ characters:

?

Disparage slightly the alternative that the ‘?’ appears in, as a choice when no alternative applies exactly. The compiler regards this alternative as one unit more costly for each ‘?’ that appears in it.

!

Disparage severely the alternative that the ‘!’ appears in. This alternative can still be used if it fits without reloading, but if reloading is needed, some other alternative will be used.

Constraint modifier characters

The following are constraint modifier characters.

=

Means that this operand is write-only for this instruction: the previous value is discarded and replaced by output data.

+

Means that this operand is both read and written by the instruction.

When the compiler fixes up the operands to satisfy the constraints, it needs to know which operands are inputs to the instruction and which are outputs from it. ‘=’ identifies an output; ‘+’ identifies an operand that is both input and output; all other operands are assumed to be input only.

&

Means (in a particular alternative) that this operand is an *earlyclobber* operand which is modified before the instruction is finished using the input operands. Therefore, this operand may not lie in a register that is used as an input operand or as part of any memory address.

‘&’ applies only to the alternative in which it is written. In constraints with multiple alternatives, sometimes one alternative requires ‘&’ while others do not. See, for example, the ‘`movdf`’ insn of the 68000.

‘&’ does not obviate the need to write ‘=’.

%

Declares the instruction to be commutative for this operand and the following operand. This means that the compiler may interchange the two operands if that is the cheapest way to make all operands fit the constraints.

`#`
Says that all following characters, up to the next comma, are to be ignored as a constraint. They are significant only for choosing register preferences.

Constraints for particular machines

Whenever possible, you should use the general-purpose constraint letters in `asm` arguments, since they will convey meaning more readily to people reading your code. Failing that, use the constraint letters that usually have very similar meanings across architectures. The most commonly used constraints are ‘`m`’ and ‘`r`’ (for memory and general-purpose registers respectively; see “Simple constraints” on page 249), and ‘`I`’, usually the letter indicating the most common immediate-constant format. For each machine architecture, the `config/ machine.h` file defines additional constraints. These constraints are used by the compiler itself for instruction generation, as well as for `asm` statements; therefore, some of the constraints are not particularly interesting for `asm`.

The constraints are defined through the following macros.

`REG_CLASS_FROM_LETTER`

Register class constraints (usually lower case).

`CONST_OK_FOR_LETTER_P`

Immediate constant constraints, for non-floating point constants of word size or smaller precision (usually upper case).

`CONST_DOUBLE_OK_FOR_LETTER_P`

Immediate constant constraints, for all floating point constants and for constants of greater than word size precision (usually upper case).

`EXTRA_CONSTRAINT`

Special cases of registers or memory. This macro is not required, and is only defined for some machines.

Inspecting these macro definitions in the compiler source for your machine is the best way to be certain you have the right constraints. However, the following is a summary of the machine-dependent constraints available on some particular machines.

ARM family—`arm.h` constraints

The following is a summary of the machine-dependent constraints available on ARM machines.

`f`
Floating-point register

`F`
One of the floating-point constants 0.0, 0.5, 1.0, 2.0, 3.0, 4.0, 5.0 or 10.0

G	Floating-point constant that would satisfy the constraint 'F' if it were negated
I	Integer that is valid as an immediate operand in a data processing instruction. That is, an integer in the range 0 to 255 rotated by a multiple of 2
J	Integer in the range -4095 to 4095
K	Integer that satisfies constraint 'I' when inverted (ones complement)
L	Integer that satisfies constraint 'I' when negated (twos complement)
M	Integer in the range 0 to 32
Q	A memory reference where the exact address is in a single register ('m' is preferable for <code>asm</code> statements)
R	An item in the constant pool
S	A symbol in the text segment of the current file

AMD 29000 family—`a29k.h` constraints

The following is a summary of the machine-dependent constraints available on AMD 29K machines.

l	Local register 0
b	Byte Pointer ('BP') register
q	'Q' register
h	Special purpose register
A	First accumulator register
a	Other accumulator register
f	Floating point register

I	Constant greater than 0, less than 0x100
J	Constant greater than 0, less than 0x10000
K	Constant whose high 24 bits are on (1)
L	16 bit constant whose high 8 bits are on (1)
M	32 bit constant whose high 16 bits are on (1)
N	32 bit negative constant that fits in 8 bits
O	The constant 0x80000000 or, on the 29050, any 32 bit constant whose low 16 bits are 0.
P	16 bit negative constant that fits in 8 bits
G, H	A floating point constant (in <code>asm</code> statements, use the machine independent ‘E’ or ‘F’ instead)

IBM RS6000—`rs6000.h` constraints

The following is a summary of the machine-dependent constraints available on IBM RS6000 machines.

b	Address base register
f	Floating point register
h	‘MQ’, ‘CTR’, or ‘LINK’ register
q	‘MQ’ register
c	‘CTR’ register
l	‘LINK’ register
x	‘CR’ register (condition register) number 0

<code>Y</code>	'CR' register (condition register)
<code>I</code>	Signed 16 bit constant
<code>J</code>	Constant whose low 16 bits are 0
<code>K</code>	Constant whose high 16 bits are 0
<code>L</code>	Constant suitable as a mask operand
<code>M</code>	Constant larger than 31
<code>N</code>	Exact power of 2
<code>O</code>	Zero
<code>P</code>	Constant whose negation is a signed 16 bit constant
<code>G</code>	Floating point constant that can be loaded into a register with one instruction per word
<code>Q</code>	Memory operand that is an offset from a register ('m' is preferable for <code>asm</code> statements)
<code>R</code>	AIX TOC entry
<code>S</code>	Windows NT SYMBOL REF
<code>T</code>	Windows NT LABEL REF
<code>U</code>	System V Release 4 small data area reference

Intel 386—`i386.h` constraints

The following is a summary of the machine-dependent constraints available on Intel 386 machines.

<code>r</code>	'a', 'b', 'c', or 'd' register
----------------	--------------------------------

A	'a', or 'd' register (for 64-bit ints)
f	Floating point register
t	First (top of stack) floating point register
u	Second floating point register
a	'a' register
b	'b' register
c	'c' register
d	'd' register
D	'di' register
S	'si' register
I	Constant in range 0 to 31 (for 32 bit shifts)
J	Constant in range 0 to 63 (for 64 bit shifts)
K	'0xff'
L	'0xffff'
M	0, 1, 2, or 3 (shifts for <code>leaq</code> instruction)
N	Constant in range 0 to 255 (for <code>out</code> instruction)
G	Standard 80387 floating point constant

Intel 960—`i960.h` constraints

The following is a summary of the machine-dependent constraints available on Intel 960 machines.

<code>f</code>	Floating point register (<code>fp0</code> to <code>fp3</code>)
<code>l</code>	Local register (<code>r0</code> to <code>r15</code>)
<code>b</code>	Global register (<code>g0</code> to <code>g15</code>)
<code>d</code>	Any local or global register
<code>I</code>	Integers from 0 to 31
<code>J</code>	0
<code>K</code>	Integers from -31 to 0
<code>G</code>	Floating point 0
<code>H</code>	Floating point 1

MIPS—`mips.h` constraints

The following is a summary of the machine-dependent constraints available on MIPS machines.

<code>d</code>	General-purpose integer register
<code>f</code>	Floating-point register (if available)
<code>h</code>	'Hi' register
<code>l</code>	'Lo' register
<code>x</code>	'Hi' or 'Lo' register
<code>Y</code>	General-purpose integer register
<code>z</code>	Floating-point status register
<code>I</code>	Signed 16 bit constant (for arithmetic instructions)

J	Zero
K	Zero-extended 16-bit constant (for logic instructions)
L	Constant with low 16 bits zero (can be loaded with <code>lui</code>)
M	32 bit constant which requires two instructions to load (a constant which is not 'I', 'K', or 'L')
N	Negative 16 bit constant
O	Exact power of two
P	Positive 16 bit constant
G	Floating point zero
Q	Memory reference that can be loaded with more than one instruction ('m' is preferable for <code>asm</code> statements)
R	Memory reference that can be loaded with one instruction ('m' is preferable for <code>asm</code> statements)
S	Memory reference in external OSF/rose PIC for-mat ('m' is preferable for <code>asm</code> statements)

Motorola 680x0—`m68k.h` constraints

The following is a summary of the machine-dependent constraints available on Motorola 68K machines.

a	Address register
d	Data register
f	68881 floating-point register, if available
x	Sun FPA (floating-point) register, if available

<code>Y</code>	First 16 Sun FPA registers, if available
<code>I</code>	Integer in the range 1 to 8
<code>J</code>	16 bit signed number
<code>K</code>	Signed number whose magnitude is greater than 0x80
<code>L</code>	Integer in the range -8 to -1
<code>M</code>	Signed number whose magnitude is greater than 0x100.
<code>G</code>	Floating point constant that is not a 68881 constant
<code>H</code>	Floating point constant that can be used by Sun FPA

SPARC—`sparc.h` constraints

The following is a summary of the machine-dependent constraints available on SPARC machines.

<code>f</code>	Floating-point register
<code>e</code>	Floating point register that can hold 64 or 128 bit values.
<code>I</code>	Signed 13 bit constant
<code>J</code>	Zero
<code>K</code>	32 bit constant with the low 12 bits clear (a constant that can be loaded with the <code>sethi</code> instruction)
<code>G</code>	Floating-point zero
<code>H</code>	Signed 13 bit constant, sign-extended to 32 or 64 bits
<code>Q</code>	Memory reference that can be loaded with one instruction (<code>'m'</code> is more appropriate for <code>asm</code> statements)

- S Constant, or memory address
- T Memory address aligned to an 8-byte boundary
- U Even register

Controlling names used in assembler code

You can specify the name to be used in the assembler code for a C function or variable by writing the `asm` (or `__asm__`) keyword after the declarator as follows:

```
int foo asm ("myfoo") = 2;
```

This specifies that the name to be used for the variable `foo` in the assembler code should be `myfoo` rather than the usual `_foo`.

On systems where an underscore is normally prepended to the name of a C function or variable, this feature allows you to define names for the linker that do not start with an underscore.

You cannot use `asm` in this way in a function definition; but you can get the same effect by writing a declaration for the function before its definition and putting `asm` there, like this:

```
extern func () asm ("FUNC");
```

```
func (x, y)
int x, y;
...
```

It is up to you to make sure that the assembler names you choose do not conflict with any other assembler symbols. Also, you must not use a register name; that would produce completely invalid assembler code. GCC does not as yet have the ability to store static variables in registers.

Variables in specified registers

GCC allows you to put a few global variables into specified hardware registers. You can also specify the register in which an ordinary register variable should be allocated.

- Global register variables reserve registers throughout the program. This may be useful in programs such as programming language interpreters which have a couple of global variables that are accessed very often.
- Local register variables in specific registers do not reserve the registers. The compiler's data flow analysis is capable of determining where the specified registers contain live values, and where they are available for other uses.

These local variables are sometimes convenient for use with the extended `asm` feature (see “Assembler instructions with C expression operands” on page 245), if you want to write one output of the assembler instruction directly into a particular register. (This will work provided the register you specify fits the constraints specified for that operand in the `asm`.)

Defining global register variables

You can define a global register variable in GCC, using the following example's input.

```
register int *foo asm ("a5");
```

In the previous example, `a5` is the name of the register which should be used. Choose a register which is normally saved and restored by function calls on your machine, so that library routines will not clobber it.

Naturally the register name is CPU-dependent, so you would need to conditionalize your program according to CPU type. The register `a5` would be a good choice on a 68000 for a variable of pointer type. On machines with register windows, be sure to choose a *global* register that is not affected magically by the function call mechanism.

In addition, operating systems on one type of CPU may differ in how they name the registers; then you would need additional conditionals. For example, some 68000 operating systems call the register, `%a5`.

Eventually there may be a way of asking the compiler to choose a register automatically, but first we need to figure out how it should choose and how to enable you to guide the choice. No solution is evident.

Defining a global register variable in a certain register reserves that register entirely for this use, at least within the current compilation. The register will not be allocated for any other purpose in the functions in the current compilation. The register will not be saved and restored by these functions. Stores into this register are never deleted even if they would appear to be dead, but references may be deleted or moved or

simplified. It is not safe to access the global register variables from signal handlers, or from more than one thread of control, because the system library routines may temporarily use the register for other things (unless you recompile them specially for the task at hand).

It is not safe for one function that uses a global register variable to call another such function, `foo`, by way of a third function, `lose`, that was compiled without knowledge of this variable (i.e., in a different source file in which the variable wasn't declared). This is because `lose` might save the register and put some other value there. For example, you can't expect a global register variable to be available in the comparison-function that you pass to `qsort`, since `qsort` might have put something else in that register. (If you are prepared to recompile `qsort` with the same global register variable, you can solve this problem.)

If you want to recompile `qsort` or other source files which do not actually use your global register variable, so that they will not use that register for any other purpose, then it suffices to specify the compiler option, `'-ffixed-reg.'` You need not actually add a global register declaration to their source code.

A function which can alter the value of a global register variable cannot safely be called from a function compiled without this variable, because it could clobber the value the caller expects to find there on return. Therefore, the function which is the entry point into the part of the program that uses the global register variable must explicitly save and restore the value which belongs to its caller.

On most machines, `longjmp` will restore to each global register variable the value it had at the time of the `setjmp`. On some machines, however, `longjmp` will not change the value of global register variables.

To be portable, the function that called `setjmp` should make other arrangements to save the values of the global register variables, and to restore them in a `longjmp`. This way, the same thing will happen regardless of what `longjmp` does. All global register variable declarations must precede all function definitions. If such a declaration could appear after function definitions, the declaration would be too late to prevent the register from being used for other purposes in the preceding functions.

Global register variables may not have initial values, because an executable file has no means to supply initial contents for a register. On the Sparc, there are reports that `g3 ... g7` are suitable registers, but certain library functions, such as `getwd`, as well as the subroutines for division and remainder, modify `g3` and `g4`. `g1` and `g2` are local temporaries. On the 68000, `a2 ... a5` should be suitable, as should `d2 ... d7`. Of course, it will not do to use more than a few of those.

Specifying registers for local variables

You can define a local register variable with a specified register like the following.

```
register int *foo asm ("a5");
```

Here `a5` is the name of the register which should be used. This is the same syntax used for defining global register variables, but for a local variable it would appear within a function.

Naturally the register name is CPU-dependent, but this is not a problem, since specific registers are most often useful with explicit assembler instructions (see “Assembler instructions with C expression operands” on page 245). Both of these things generally require that you conditionalize your program according to CPU type.

In addition, operating systems on one type of CPU may differ in how they name the registers; then you would need additional conditionals. For example, some 68000 operating systems call this register, `%a5`.

Eventually there may be a way of asking the compiler to choose a register automatically, but first we need to figure out how it should choose and how to enable you to guide the choice. No solution is evident.

Defining such a register variable does not reserve the register; it remains available for other uses in places where flow control determines the variable’s value is not live. However, these registers are made unavailable for use in the reload pass. I would not be surprised if excessive use of this feature leaves the compiler too few available registers to compile certain functions.

Alternate keywords

The option, `-traditional`, disables certain keywords; `-ansi` disables certain others. This causes trouble when you want to use GCC extensions, or ANSI C features, in a general-purpose header file that should be usable by all programs, including ANSI C programs and traditional ones.

The keywords, `asm`, `typeof` and `inline`, cannot be used since they won't work in a program compiled with `'-ansi'`, while the keywords, `const`, `volatile`, `signed`, `typeof` and `inline`, won't work in a program compiled with `'-traditional'`.

The way to solve these problems is to put `'__'` at the beginning and end of each problematical keyword. For example, use `__asm__` instead of `asm`, `__const__` instead of `const`, and `__inline__` instead of `inline`.

Other C compilers won't accept these alternative keywords; if you want to compile with another compiler, you can define the alternate keywords as macros to replace them with the customary keywords.

It looks like the following declaration..

```
#ifndef __GNUC__
#define __asm__ asm
#endif
```

`-pedantic` causes warnings for many GCC extensions. You can prevent such warnings within one expression by writing `__extension__` before the expression. `__extension__` has no effect aside from this problem.

Incomplete `enum` types

You can define an `enum` tag without specifying its possible values. This results in an incomplete type, much like what you get if you write `struct foo` without describing the elements. A later declaration which does specify the possible values completes the type.

You can't allocate variables or storage using the type while it is incomplete. However, you can work with pointers to that type.

This extension may not be very useful, but it makes the handling of `enum` more consistent with the way `struct` and `union` are handled.

This extension is not supported by GNU C++.

Function names as strings

GCC predefines two string variables to be the name of the current function. The variable, `__FUNCTION__`, is the name of the function as it appears in the source. The variable, `__PRETTY_FUNCTION__`, is the name of the function pretty printed in a language specific fashion. These names are always the same in a C function; in a C++ function, they may be different, like the following program.

```
extern "C" {
extern int printf (char *, ...);
}
class a {
public:
    sub (int i)
    {
        printf ("__FUNCTION__ = %s\n", __FUNCTION__);
        printf ("__PRETTY_FUNCTION__ = %s\n", __PRETTY_FUNCTION__);
    }
};
int
main (void)
{
    a ax;
    ax.sub (0);
    return 0;
}
```

The program, then, gives the following output.

```
__FUNCTION__ = sub
__PRETTY_FUNCTION__ = int a::sub (int)
```

These names are not macros: they are predefined string variables. For example, `#ifdef __FUNCTION__` does not have any special meaning inside a function, since the preprocessor does not do anything special with the identifier, `__FUNCTION__`.

Getting the return or frame address of a function

The following calls may be used to get information about the callers of a function.

`__builtin_return_address (level)`

This function returns the return address of the current function, or of one of its callers. The *level* argument is number of frames to scan up the call stack. A value of 0 yields the return address of the current function, a value of 1 yields the return address of the caller of the current function, and so forth. The *level* argument must be a constant integer. On some machines it may be impossible to determine the return address of any function other than the current one; in such cases, or when the top of the stack has been reached, this function will return 0. This function should only be used with a non-zero argument for debugging purposes.

```
__builtin_frame_address ( level)
```

This function is similar to `__builtin_return_address`, but it returns the address of the function frame rather than the return address of the function. Calling `__builtin_frame_address` with a value of 0 yields the frame address of the current function, a value of 1 yields the frame address of the caller of the current function, and so forth. The frame is the area on the stack which holds local variables and saved registers. The frame address is normally the address of the first word pushed on to the stack by the function. However, the exact definition depends upon the processor and the calling convention. If the processor has a dedicated frame pointer register, and the function has a frame, then `__builtin_frame_address` will return the value of the frame pointer register. The caveats that apply to `__builtin_return_address` apply to this function as well.

Extensions to the C++ language

The GNU compiler provides extensions to the C++ language (and you can also use most of the C language extensions in your C++ programs). If you want to write code that checks whether these features are available, you can test for the GNU compiler the same way as for C programs: check for a predefined macro, `__GNUC__`. You can also use `__GNUG__` to test specifically for the GNU C++ compiler tools, G++ (see “Standard predefined macros” on page 353 in *The C Preprocessor* in **GNUPro Compiler Tools**).

- “Named return values in C++” on page 272
- “Minimum and maximum operators in C++” on page 274
- “The `goto` and destructors in GNU C++” on page 275
- “Declarations and definitions in one header” on page 276
- “Where’s the template?” on page 278
- “Type abstraction using signatures” on page 282

Named return values in C++

GNU C++ extends the function-definition syntax, which allows you to specify a name for the result of a function outside the body of the definition, in C++ programs

```
type
functionname (args) return resultname;
{
    ...
    body
    ...
}
```

You can use this feature to avoid an extra constructor call when a function result has a class type. For example, consider a function `m`, declared as

`'X v =m ();'`, whose result is of class `x`:

```
X
m ()
{
    X b;
    b.a = 23;
    return b;
}
```

Although `'m'` appears to have no arguments, in fact it has one implicit argument: the address of the return value. At invocation, the address of enough space to hold `'v'` is sent in as the implicit argument. Then `'b'` is constructed and its `'a'` field is set to the value 23. Finally, a copy constructor (a constructor of the form `'x(x&)'`) is applied to `'b'`, with the (implicit) return value location as the target, so that `'v'` is now bound to the return value.

But this is wasteful. The local `'b'` is declared just to hold something that will be copied right out. While a compiler that combined an *elision* algorithm with interprocedural data flow analysis could conceivably eliminate all of this, it is much more practical to allow you to assist the compiler in generating efficient code by manipulating the return value explicitly, thus avoiding the local variable and copy constructor altogether.

Using the extended GNU C++ function-definition syntax, you can avoid the temporary allocation and copying by naming `r` as your return value at the outset, and assigning to its `'a'` field directly the following declaration.

```
X
m () return r;
{
    r.a = 23;
}
```

The declaration of `'r'` is a standard, proper declaration, whose effects are executed *before* any of the body of `'m'`.

Functions of this type impose no additional restrictions; in particular, you can execute return statements, or return implicitly by reaching the end of the function body (*falling off the edge*). Cases like the following declaration (or even the statement, `X m ()` `return r (23) ; { }`) are unambiguous, since the return value 'r' has been initialized in either case.

```
X
m () return r (23);
{
    return;
}
```

The following code may be hard to read, but also works predictably.

```
X
m () return r;
{
    X b;
    return b;
}
```

The return value slot denoted by 'r' is initialized at the outset, but the statement 'return b;' overrides this value. The compiler deals with this by destroying 'r' (calling the destructor if there is one, or doing nothing if there is not), and then reinitializing 'r' with 'b'.

This extension is provided primarily to help people who use overloaded operators, where there is a great need to control not just the arguments, but the return values of functions. For classes where the copy constructor incurs a heavy performance penalty (especially in the common case where there is a quick default constructor), this is a major savings. The disadvantage of this extension is that you do not control when the default constructor for the return value is called: it is always called at the beginning.

Minimum and maximum operators in C++

It is very convenient to have operators which return the *minimum* or the *maximum* of two arguments. For instance, in GNU C++ (but not in GNU C), operations perform the following returns.

$a < ? b$ is the *minimum*, returning the smaller of the numeric values a and b ;

$a > ? b$ is the *maximum*, returning the larger of the numeric values a and b .

These operations are not primitive in ordinary C++, since you can use a macro to return the minimum of two things in C++, as in the following example.

```
#define MIN(X,Y) ((X) < (Y) ? (X) : (Y))
```

You might then use `'int min = MIN (i, j);'` to set `'min'` to the minimum value of variables `'i'` and `'j'`. However, side effects in `'x'` or `'y'` may cause unintended behavior. For example, `MIN (i++, j++)` will fail, incrementing the smaller counter twice. A GNU C extension allows you to write safe macros that avoid this kind of problem (see “Naming an expression’s type” on page 209). However, writing `MIN` and `MAX` as macros also forces you to use function-call notation for a fundamental arithmetic operation.

Using GNU C++ extensions, you can write `'int min = i < ? j;'` instead.

Since `< ?` and `> ?` are built into the compiler, they properly handle expressions with side-effects; `'int min = i++ < ? j++;'` works correctly.

The `goto` and destructors in GNU C++

In C++ programs, you can safely use the `goto` statement. When you use it to exit a block which contains aggregates requiring destructors, the destructors will run before the `goto` transfers control. (In ANSI C++, `goto` is restricted to targets within the current function.) The compiler still forbids using `goto` to *enter* a scope that requires constructors.

Declarations and definitions in one header

C++ object definitions can be quite complex. In principle, your source code will need two kinds of things for each object that you use across more than one source file.

First, you need an *interface* specification, describing its structure with type declarations and function prototypes. Second, you need the *implementation* itself. It can be tedious to maintain a separate interface description in a header file, in parallel to the actual implementation. It is also dangerous, since separate interface and implementation definitions may not remain parallel. With GNU C++, you can use a single header file for both purposes.

Warning: The mechanism to specify this is in transition. For the nonce, you must use one of two `#pragma` commands; in a future release of GNU C++, an alternative mechanism will make these `#pragma` commands unnecessary.

The header file contains the full definitions, but is marked with `'#pragma interface'` in the source code. This allows the compiler to use the header file only as an interface specification when ordinary source files incorporate it with `#include`. In the single source file where the full implementation belongs, you can use either a naming convention or

`'#pragma implementation'` to indicate this alternate use of the header file.

```
#pragma interface
#pragma interface "subdir/objects.h"
```

Use this directive in *header files* that define object classes, to save space in most of the object files that use those classes. Normally, local copies of certain information (backup copies of inline member functions, debugging information, and the internal tables that implement virtual functions) must be kept in each object file that includes class definitions. You can use this pragma to avoid such duplication.

When a header file containing `'#pragma interface'` is included in a compilation, this auxiliary information will not be generated (unless the main input source file itself uses `'#pragma implementation'`). Instead, the object files will contain references to be resolved at link time.

The second form of this directive is useful for the case where you have multiple headers with the same name in different directories. If you use this form, you must specify the same string to `'#pragma implementation'`.

```
#pragma implementation
#pragma implementation "objects.h"
```

Use this pragma in a *main input file*, when you want full output from included header files to be generated (and made globally visible). The included header file, in turn, should use `'#pragma interface'`. Backup copies of inline member

functions, debugging information, and the internal tables used to implement virtual functions are all generated in implementation files.

If you use `#pragma implementation` with no argument, it applies to an include file with the same *basename*[†] as your source file. For example, in `'allclass.cc'`, `'#pragma implementation'`, by itself, is equivalent to `'#pragma implementation "allclass.h"'`.

In versions of GNU C++ prior to 2.6.0, `'allclass.h'` was treated as an implementation file whenever you would include it from `'allclass.cc'` even if you never specified `'#pragma implementation'`. This was deemed to be more trouble than it was worth, however, and disabled. If you use an explicit `'#pragma implementation'`, it must appear in your source file before you include the affected header files.

Use the string argument if you want a single implementation file to include code from multiple header files. (You must also use `#include` to include the header file; `'#pragma implementation'` only specifies how to use the file—it doesn't actually include it.) There is no way to split up the contents of a single header file into multiple implementation files.

`#pragma implementation` and `'#pragma interface'` also have an effect on function inlining.

If you define a class in a header file marked with `'#pragma interface'`, the effect on a function defined in that class is similar to an explicit `extern` declaration—the compiler emits no code at all to define an independent version of the function. Its definition is used only for inlining with its callers.

Conversely, when you include the same header file in a main source file that declares it as `'#pragma implementation'`, the compiler emits code for the function itself; this defines a version of the function that can be found via pointers (or by callers compiled without inlining). If all calls to the function can be inlined, you can avoid emitting the function by compiling with `'-fno-implement-inlines'`. If any calls were not inlined, you will get linker errors.

[†] A file's *basename* was the name stripped of all leading path information and of trailing suffixes (such as: `.h` or `.C` or `.cc`).

Where's the template?

C++ templates are the first language feature to require more intelligence from the environment than one usually finds on a UNIX system. Somehow the compiler and linker have to make sure that each template instance occurs exactly once in the executable if it is needed, and not at all otherwise.

There are two basic approaches to this problem, which we will refer to as the *Borland model* and the *Cfront model*.

Borland model

Borland C++ solved the template instantiation problem by adding the code equivalent of common blocks to their linker; the compiler emits template instances in each translation unit that uses them, and the linker collapses them together at run time. The advantage of this model is that the linker only has to consider the object files themselves; there is no external complexity to worry about. This disadvantage is that compilation time is increased because the template code is being compiled repeatedly. Code written for this model tends to include definitions of all member templates in the header file, since they must be seen to be instantiated.

Cfront model

The AT&T C++ translator, “Cfront”, solved the template instantiation problem by creating the notion of a template repository, an automatically maintained place where template instances are stored. A more modern version of the repository works as follows.

As individual object files are built, the compiler places any template definitions and instantiations encountered in this repository. At link time, the link wrapper adds in the objects in the repository and compiles any needed instances that were not previously emitted. The advantages of this model are more optimal compilation speed and the ability to use the system linker; to implement the Borland model, a compiler vendor also needs to replace the linker. The disadvantages are vastly increased complexity, and thus potential for error; for some code, this can be just as transparent, but in practice it can be very difficult to build multiple programs in one directory and one program in multiple directories using Cfront. Code written for this model tends to separate definitions of non-inline member templates into a separate file, which should be separately compiled.

When used with GNU 1d version 2.8 or later on an ELF system such as Linux/GNU or Solaris 2, or on Microsoft Windows, g++ supports the Borland model.

A future version of `g++` will support a hybrid model whereby the compiler will emit any instantiations for which the template definition is included in the compile, and store template definitions and instantiation context information into the object file for the rest. The link wrapper will extract that information as necessary and invoke the compiler to produce the remaining instantiations. The linker will then combine duplicate instantiations.

In the mean time, you have the following options for dealing with template instantiations.

- Compile your template-using code with `-frepo`. The compiler will generate files with the extension `.rpo` listing all of the template instantiations used in the corresponding object files which could be instantiated there; the link wrapper, `collect2`, will then update the `.rpo` files to tell the compiler where to place those instantiations and rebuild any affected object files. The link-time overhead is negligible after the first pass, as the compiler will continue to place the instantiations in the same files.

This is your best option for application code written for the Borland model, as it will just work. Code written for the Cfront model will need to be modified so that the template definitions are available at one or more points of instantiation; usually this is as simple as adding `#include <tmethods.cc>` to the end of each template header.

For library code, if you want the library to provide all of the template instantiations it needs, just try to link all of its object files together; the link will fail, but cause the instantiations to be generated as a side effect. Be warned, however, that this may cause conflicts if multiple libraries try to provide the same instantiations. For greater control, use explicit instantiation as described in the next option.

- Compile your code with `-fno-implicit-templates` to disable the implicit generation of template instances, and explicitly instantiate all the ones you use. This approach requires more knowledge of exactly which instances you need than do the others, but it's less mysterious and allows greater control. You can scatter the explicit instantiations throughout your program, perhaps putting them in the translation units where the instances are used or the translation units that define the templates themselves; you can put all of the explicit instantiations you need into one big file; or you can create small files for each of the instances you need, like the following examples define, and create a template instantiation library from those files.

```
#include "Foo.h"
#include "Foo.cc"

template class Foo<int>;
template ostream& operator <<
    (ostream&, const Foo<int>&);
```

If you are using Cfront-model code, you can probably get away with not using `'-fno-implicit-templates'` when compiling files that don't `'#include'` the member template definitions.

If you use one big file to do the instantiations, you may want to compile it without `'-fno-implicit-templates'` so you get all of the instances required by your explicit instantiations (but not by any other files) without having to specify them as well.

g++ has extended the template instantiation syntax outlined in the Working Paper to allow forward declaration of explicit instantiations, explicit instantiation of members of template classes and instantiation of the compiler support data for a template class (i.e., the vtable) without instantiating any of its members as the following example shows.

```
extern template int max (int, int);
template void Foo<int>::f ();
inline template class Foo<int>;
```

- Do nothing. Pretend G++ does implement automatic instantiation management. Code written for the Borland model will work fine, but each translation unit will contain instances of each of the templates it uses. In a large program, this can lead to an unacceptable amount of code duplication.
- Add `'#pragma interface'` to all files containing template definitions. For each of these files, add `'#pragma implementation "filename"'` to the top of some `'.c'` file which `'#include's'` it. Then compile everything with `-fexternal-templates`. The templates will then only be expanded in the translation unit which implements them (i.e., the translation unit has a `#pragma implementation` line for the file where they live); all other files will use external references. If you're lucky, everything should work properly. If you get undefined symbol errors, you need to make sure that each template instance which is used in the program is used in the file which implements that template. If you don't have any use for a particular instance in that file, you can just instantiate it explicitly, using the syntax from the latest C++ working paper:

```
template class A<int>;
template ostream& operator << (ostream&, const A<int>&);
```

This strategy will work with code written for either model. If you are using code written for the Cfront model, the file containing a class template and the file containing its member templates should be implemented in the same translation unit.

A slight variation on this approach is to use the flag `-falt-external-templates` instead; this flag causes template instances to be emitted in the translation unit that implements the header where they are first instantiated, rather than the one which implements the file where the templates are defined. This header must be the same in all translation units, or things are likely to break.

See “Declarations and definitions in one header” on page 276 for more discussion of these pragmas.

Type abstraction using signatures

In GNU C++, you can use the keyword `signature` to define a completely abstract class interface as a datatype. You can connect this abstraction with actual classes using signature pointers. If you want to use signatures, run the GNU compiler with the `-fhandle-signatures` command-line option. (With this option, the compiler reserves a second keyword, `sigof`, as well, for a future extension.)

Roughly, signatures are type abstractions or interfaces of classes. Some other languages have similar facilities. C++ signatures are related to ML's signatures, Haskell's type classes, definition modules in Modula-2, interface modules in Modula-3, abstract types in Emerald, type modules in Trellis/Owl, categories in Scratchpad II, and types in POOL-I. For a more detailed discussion of signatures, see *Signatures: A Language Extension for Improving Type Abstraction and Subtype Polymorphism in C++* by Gerald Baumgartner and Vincent F. Russo (Tech report CSD-TR-95-051, Dept. of Computer Sciences, Purdue University, August 1995, a slightly improved version appeared in *Software—Practice & Experience*, **25**(8), pp. 863–889, August 1995). You can get the tech report by anonymous FTP from `ftp.cs.purdue.edu` in `'pub/gb/Signature-design.ps.gz'`.

Syntactically, a signature declaration is a collection of member function declarations and nested type declarations. For example, the following signature declaration defines a new abstract type 's' with member functions, 'int foo()' and 'int bar(int)'.

```
signature S
{
    int foo ();
    int bar (int);
};
```

Since signature types do not include implementation definitions, you cannot write an instance of a signature directly. Instead, you can define a pointer to any class that contains the required interfaces as a *signature pointer*. Such a class *implements* the signature type.

To use a class as an implementation of `s`, you must ensure that the class has public member functions 'int foo()' and 'int bar(int)'. The class can have other member functions as well, public or not; as long as it offers what's declared in the signature, it is suitable as an implementation of that signature type.

For example, suppose that 'c' is a class that meets the requirements of signature 's' (c *conforms to s*). Then the following statement defines a signature pointer 'p' and initializes it to point to an object of type 'c'.

```
C obj;
S * p = &obj;
```

The member function call, `int i = p->foo ();`, executes `obj.foo ()`.

Abstract virtual classes provide somewhat similar facilities in standard C++. There are two main advantages to using signatures instead:

- Subtyping becomes independent from inheritance.
A class or signature type ‘*T*’ is a subtype of a signature type ‘*S*’ independent of any inheritance hierarchy as long as all the member functions declared in ‘*S*’ are also found in ‘*T*’. So you can define a subtype hierarchy that is completely independent from any inheritance (implementation) hierarchy, instead of being forced to use types that mirror the class inheritance hierarchy.
- Signatures allow you to work with existing class hierarchies as implementations of a signature type. If those class hierarchies are only available in compiled form, you’re out of luck with abstract virtual classes, since an abstract virtual class cannot be retrofitted on top of existing class hierarchies.

So you would be required to write interface classes as subtypes of the abstract virtual class.

There is one more detail about signatures. A signature declaration can contain member function *definitions* as well as member function declarations. A signature member function with a full definition is called a *default implementation*; classes need not contain that particular interface in order to conform.

For example, a class ‘*C*’ can conform to the following signature.

```
signature T
{
    int f (int);
    int f0 () { return f (0); };
};
```

This happens whether *C* implements the member function, ‘*int f0()*’, or not. If you define *C* : : *f0*, that definition takes precedence; otherwise, the default implementation, *S* : : *f0*, applies.

gcov: a test coverage program

`gcov` is a tool you can use, together with GCC, to test code coverage in your programs. `gcov` is free software, but for the moment it is only available from Cygnus (pending discussions with the FSF about how they think Cygnus should *really* write it). The following documentation describes version 1.5 of `gcov`.

- “Introduction to gcov test coverage” on page 286
- “Invoking the gcov program” on page 287
- “Using gcov with GCC optimization” on page 290
- “Brief description of gcov data files” on page 291

Jim Wilson wrote `gcov`, and the original form of this documentation. Pat McGregor edited the documentation.

Introduction to `gcov` test coverage

`gcov` is a test coverage program. Use it in concert with GCC to analyze your programs to help create more efficient, faster running code. You can use `gcov` as a profiling tool, to help discover where your optimization efforts will best affect your code. You can also use `gcov` in concert with the other profiling tool, `gprof`, to assess which parts of your code use the greatest amount of computing time.

Profiling tools help you analyze your code's performance. Using a profiler, such as `gcov` or `gprof`, you can find out some basic performance statistics, such as:

- how often each line of code executes
- what lines of code are actually executed
- how much computing time each section of code uses

Once you know these things about how your code works when compiled, you can look at each module to see which modules should be optimized. `gcov` helps you determine where to work on optimization.

Software developers also use coverage testing in concert with testsuites, to make sure software is actually good enough for a release. Testsuites can verify that a program works as expected; a coverage program tests to see how much of the program is exercised by the testsuite. Developers can then determine what kinds of test cases need to be added to the testsuites to create both better testing and a better final product.

You should compile your code without optimization if you plan to use `gcov`, because the optimization, by combining some lines of code into one function, may not give you as much information as you need to look for *hot spots*, where the code is using a great deal of computer time. Likewise, because `gcov` accumulates statistics by line (at the lowest resolution), it works best with a programming style that places only one statement on each line.

If you use complicated macros that expand to loops or to other control structures, the statistics are less helpful—they only report on the line where the macro call appears. If your complex macros behave like functions, you can replace them with inline functions to solve this problem.

`gcov` creates a logfile called '`sourcefile.gcov`' which indicates how many times each line of a source file, '`sourcefile.c`', has executed. You can use these logfiles in conjunction with `gprof` to aid in fine-tuning the performance of your programs. `gprof` gives timing information you can use along with the information you get from `gcov`.

`gcov` works only on code compiled with GCC; it is not compatible with any other profiling or test coverage mechanism.

Invoking the `gcov` program

The following declaration is an example of invoking the `gcov` program

```
gcov [-b] [-v] [-n] [-l] [-f] [-o directory] sourcefile
```

- b Write branch frequencies to the output file. Write branch summary info to standard output. This option allows you to see how often each branch was taken.
- v Display the `gcov` version number (on the standard error stream).
- n Do not create the `gcov` output file.
- l Create long filenames for included source files. For example, if the header file, 'x.h', contains code, and was included in the file, 'a.c', then running `gcov` on the file, 'a.c', will produce an output file 'a.c.x.h.gcov' instead of 'x.h.gcov'. This can be useful if 'x.h' is included in multiple source files.
- f Output summaries for each function in addition to the file level summary.
- o The directory where the object files live. `gcov` will search for '.bb', '.bbg', and '.da' files in this directory.

To use `gcov`, first compile your program with two special GCC options:

`-fprofile-arcs` and `-ftest-coverage`. This tells the compiler to generate additional information needed by `gcov` (basically a flow graph of the program) and also includes additional code in the object files for generating the extra profiling information needed by `gcov`. These additional files are placed in the directory where the source code is located.

Running `gcov` will cause profile output to be generated. For each source file compiled with `-fprofile-arcs`, an accompanying file, `.da`, will be placed in the source directory.

Running `gcov` with your program's source file names as arguments will now produce a listing of the code along with frequency of execution for each line. For example, if your program is called 'tmp.c', you'll see the following when you use the `gcov` facility.

```
% gcc -fprofile-arcs -ftest-coverage tmp.c
% a.out
% gcov tmp.c
87.50% of 8 source lines executed in file tmp.c
Creating tmp.c.gcov.
```

The file, 'tmp.c.gcov', contains output from `gcov`.

The following is a sample of `gcov` output.

```

main()
{
1      int i, total;

1      total = 0;
11     for (i = 0; i < 10; i++)
10     total += i;
1      if (total != 45)
#####    printf ("Failure\n");
        else
1         printf ("Success\n");
1      }

```

When you use the ‘b’ option, your output looks like the following statement.

```

% gov -b tmp.c
% a.out
% gcov tmp.c
87.50% of 8 source lines executed in file tmp.c
80.00% of 5 branches executed in file tmp.c
80.00% of 5 branches taken at least once in file tmp.c
50.00% of 2 calls executed in file tmp.c
Creating tmp.c.gcov.

```

The following is an example of a resulting ‘`tmp.c.gcov`’ file.

```

main()
{
1      int i, total;

1      total = 0;

11     for (i = 0; i < 10; i++)
branch 0 taken = 91%
branch 1 taken = 100%
branch 2 taken = 100%
10     total += i;

1      if (total != 45)
branch 0 taken = 100%
#####    printf ("Failure\n");
call 0 never executed
branch 1 never executed
        else
1         printf ("Success\n");

call 0 returns = 100%
1      }

```

For each basic block, a line is printed after the last line of the basic block describing the branch or call that ends the basic block. There can be multiple branches and calls listed for a single source line if there are multiple basic blocks that end on that line. In this case, the branches and calls are each given a number. There is no simple way to map these branches and calls back to source constructs. In general, though, the lowest

numbered branch or call will correspond to the leftmost construct on the source line.

For a branch, if it was executed at least once, then a percentage indicating the number of times the branch was taken divided by the number of times the branch was executed will be printed. Otherwise, the message “**never executed**” is printed.

For a call, if it was executed at least once, then a percentage indicating the number of times the call returned divided by the number of times the call was executed will be printed.

This will usually be 100%, but may be less for functions, `exit` or `longjmp`, and thus may not return every time they are called.

The execution counts are cumulative. If the example program were executed again without removing the `.da` file, the count for the number of times each line in the source was executed would be added to the results of the previous run(s). This is potentially useful in several ways. For example, it could be used to accumulate data over a number of program runs as part of a test verification suite, or to provide more accurate long-term information over a large number of program runs.

The data in the `.da` files is saved immediately before the program exits. For each source file compiled with `-fprofile-arcs`, the profiling code first attempts to read in an existing `.da` file; if the file doesn't match the executable (differing number of basic block counts) it will ignore the contents of the file. It then adds in the new execution counts and finally writes the data to the file.

Using `gcov` with GCC optimization

If you plan to use `gcov` to help optimize your code, you must first compile your program with two options: `-fprofile-arcs` and `-ftest-coverage`. Aside from that, you can use any other GCC options; but if you want to prove that every single line in your program was executed, you should not compile with optimization at the same time. On some machines the optimizer can eliminate some simple code lines by combining them with other lines. For example, code like the following can be compiled into one instruction on some machines.

```
if (a != b)
    c =1;
else
    c =0;
```

In this case, there is no way for `gcov` to calculate separate execution counts for each line because there isn't separate code for each line. Hence, the `gcov` output looks like the following declaration if you compiled the program with optimization.

```
100 if (a != b)
100 c = 1;
100 else
100 c = 0;
```

The output shows that this block of code, combined by optimization, executed 100 times. In one sense this result is correct, because there was only one instruction representing all four of these lines. However, the output does not indicate how many times the result was 0 and how many times the result was 1.

Brief description of gcov data files

gcov uses three files for doing profiling. The names of these files are derived from the original *source* file by substituting the file suffix with either *.bb*, *.bbg*, or *.da*. All of these files are placed in the same directory as the source file, and contain data stored in a platform-independent method.

The *.bb* and *.bbg* files are generated when the source file is compiled with the GCC ‘-ftest-coverage’ option. The *.bb* file contains a list of source files (including headers), functions within those files, and line numbers corresponding to each basic block in the source file.

The *.bb* file format consists of several lists of 4-byte integers which correspond to the line numbers of each basic block in the file. Each list is terminated by a line number of 0. A line number of -1 is used to designate that the source file name (padded to a 4-byte boundary and followed by another -1) follows. In addition, a line number of -2 is used to designate that the name of a function follows (also padded to a 4-byte boundary and followed by a -2).

The *.bbg* file is used to reconstruct the program flow graph for the source file. It contains a list of the program flow arcs (possible branches taken from one basic block to another) for each function which, in combination with the *.bb* file, enables gcov to reconstruct the program flow.

In the *.bbg* file, the format is the following declaration.

```

number of basic blocks for function #0 (4-byte number)
total number of arcs for function #0 (4-byte number)
count of arcs in basic block #0 (4-byte number)
destination basic block of arc #0 (4-byte number)
flag bits (4-byte number)
destination basic block of arc #1 (4-byte number)
flag bits (4-byte number)
...
destination basic block of arc #N (4-byte number)
flag bits (4-byte number)
count of arcs in basic block #1 (4-byte number)
destination basic block of arc #0 (4-byte number)
flag bits (4-byte number)
...
```

A -1 (stored as a 4-byte number) is used to separate each function’s list of basic blocks, and to verify that the file has been read correctly.

The *.da* file is generated when a program containing object files built with the GCC ‘-fprofile-arcs’ option is executed. A separate *.da* file is created for each source file compiled with this option, and the name of the *.da* file is stored as an absolute pathname in the resulting object file. This path name is derived from the source file name by substituting a *.da* suffix.

The format of the *.da* file is fairly simple. The first 8-byte number is the number of

counts in the file, followed by the counts (stored as 8-byte numbers). Each count corresponds to the number of times each arc in the program is executed. The counts are cumulative; each time the program is executed, it attempts to combine the existing `.da` files with the new counts for this invocation of the program. It ignores the contents of any `.da` files whose number of arcs doesn't correspond to the current program, and merely overwrites them instead.

All three of these files use the functions in `gcov-io.h` to store integers; the functions in this header provide a machine-independent mechanism for storing and retrieving data from a stream.

Known causes of trouble with GCC

The following documentation describes known problems that affect users of GCC. Most of these are not GCC bugs per se—if they were, we would fix them. But the result for a user may be like the result of a bug.

- “Actual bugs we haven’t fixed yet” on page 294
- “Installation problems” on page 294
- “Cross-compiler problems” on page 299
- “Interoperation” on page 300
- “Problems compiling certain programs” on page 305
- “Incompatibilities of GCC” on page 306
- “Fixed header files” on page 309
- “Standard libraries” on page 310
- “Disappointments and misunderstandings” on page 310
- “Common misunderstandings with GNU C++” on page 312
- “protoize and unprotoize warnings” on page 313
- “Certain changes we don’t want to make” on page 315
- “Warning messages and error messages” on page 318

Some of these problems were due to bugs in other software, some are missing features too problematic to add, and some are due to conflicts of opinion.

Actual bugs we haven't fixed yet

The following documentation describes known problems that affect users of GCC. Most of these are not GCC bugs per se—if they were, we would fix them. But the result for a user may be like the result of a bug.

Some of these problems are due to bugs in other software, some are missing features that are too much work to add, and some are places where people's opinions differ as to what is best.

- The `fixincludes` script interacts badly with automounters; if the directory of system header files is automounted, it tends to be unmounted while `fixincludes` is running. This would seem to be a bug in the automounter. We don't know any good way to work around it.
- The `fixproto` script will sometimes add prototypes for the `sigsetjmp` and `siglongjmp` functions that reference the `jmp_buf` type before that type is defined. To work around this, edit the offending file and place the `typedef` in front of the prototypes.
- There are several obscure cases of misusing `struct`, `union`, and `enum` tags that are not detected as errors by the compiler.
- When `'-pedantic-errors'` is specified, GNU C will incorrectly give an error message when a function name is specified in an expression involving the comma operator.
- Loop unrolling doesn't work properly for certain C++ programs. This is a bug in the C++ front end. It sometimes emits incorrect debug information, and the loop unrolling code is unable to recover from this error.

```
#!/bin/ksh
```

Installation problems

The following documentation describes problems (and some apparent problems that don't really mean anything is wrong) showing up during installation of GCC.

- On certain systems, defining certain environment variables such as `CC` can interfere with the functioning of `make`.
- If you encounter seemingly strange errors when trying to build the compiler in a directory other than the source directory, it could be because you have previously configured the compiler in the source directory. Make sure you have done all the necessary preparations. See “Compilation in a separate directory” on page 47.
- If you build GCC on a BSD system using a directory stored in a System V file system, problems may occur in running `fixincludes` if the System V file system doesn't support symbolic links. These problems result in a failure to fix the declaration of `size_t` in `'sys/types.h'`. If you find that `size_t` is a signed type and that type mismatches occur, this could be the cause. The solution is to use a different directory for building GCC.

- In previous versions of GCC, the `gcc` driver program looked for `as` and `ld` in various places; for example, in files beginning with `‘/usr/local/lib/gcc-’`. GCC version 2 looks for them in the following path, signifying machine for the *machine* and *version* that you need to specify.

```
/usr/local/lib/gcc-lib/target/version
```

To use a version of the assembler or linker that is not the default, for example, `gas` or `ld`, you must put them in that path (or make links to them from that path).

- Some commands executed when making the compiler may fail (return a non-zero status) and be ignored by `make`. These failures, which are often due to files that were not found, are expected, and can safely be ignored.
- It is normal to have warnings in compiling certain files about unreachable code and about enumeration type clashes. The names of these files begin with `‘insn-’`. Also, `‘real.c’` may get some warnings that you can ignore.
- Sometimes `make` recompiles parts of the compiler when installing the compiler. In one case, this was traced down to a bug in `make`. Either ignore the problem or switch to GNU Make.
- If you have installed a program known as `purify`, you may find that it causes errors while linking `enquire`, which is part of building GCC. The fix is to get rid of the file `real-ld` which `purify` installs—so that GCC won’t try to use it.
- On Linux SLS 1.01, there is a problem with `‘libc.a’`; it does not contain the `obstack` functions. However, GCC assumes that the `obstack` functions are in `‘libc.a’` when it is the GNU C library. To work around this problem, change the `__GNU_LIBRARY__` conditional (around line 31) to `‘#if 1’`.
- On some 386 systems, building the compiler never finishes because `enquire` hangs due to a hardware problem in the motherboard—it reports floating point exceptions to the kernel incorrectly. You can install GCC except for `‘float.h’` by patching out the command to run `enquire`. You may also be able to fix the problem for real by getting a replacement motherboard. This problem was observed in Revision E of the Micronics motherboard, and is fixed in Revision F. It has also been observed in the MYLEX MXA-33 motherboard. If you encounter this problem, you may also want to consider removing the FPU from the socket during the compilation. Alternatively, if you are running SCO UNIX, you can reboot and force the FPU to be ignored. To do this, type `‘hd(40)unix auto ignorefpu’`.
- On some 386 systems, GCC crashes trying to compile the `‘enquire.c’` class file. This happens on machines that don’t have a 387 FPU chip. On 386 machines, the system kernel is supposed to emulate the 387 when you don’t have one. The crash is due to a bug in the emulator.

One of these systems is the UNIX from Interactive Systems: 386/ix. On this system, an alternate emulator is provided, and it does work. To use it, execute the following command as super-user and then reboot the system. (The default

emulator file remains present under the name, `emulator.dflt.`)

```
ln /etc/emulator.rell /etc/emulator
```

Try using `/etc/emulator.att`, if you have such a problem on the SCO system.

Another system which has this problem is Esix. We don't know whether it has an alternate emulator that works.

On NetBSD 0.8, a similar problem manifests itself as the following error messages.

```
enquire.c: In function 'fprop':
enquire.c:2328: floating overflow
```

- On SCO systems, when compiling GCC with the system's compiler, do not use `-o`. Some versions of the system's compiler miscompile GCC with `-o`.
- Sometimes on a Sun 4 you may observe a crash in the program `genflags` or `genoutput` while building GCC. This is said to be due to a bug in `'sh'`. You can probably get around it by running `genflags` or `genoutput` manually and then retrying the `make`.
- On Solaris 2, executables of GCC version 2.0.2 are commonly available, but they have a bug that shows up when compiling current versions of GCC: undefined symbol errors occur during assembly if you use the `-g` option. The solution is to compile the current version of GCC without the `-g` option. That makes a working compiler which you can use to recompile with the `-g` option.
- Solaris 2 comes with a number of optional OS packages. Some of these packages are needed to use GCC fully. If you did not install all optional packages when installing Solaris, you will need to verify that the packages that GCC needs are installed.

To check whether an optional package is installed, use the `pkginfo` command. To add an optional package, use the `pkgadd` command. For further details, see the Solaris documentation. For Solaris 2.0 and 2.1, GCC needs six packages: `'SUNWarc'`, `'SUNWbtool'`, `'SUNWesu'`, `'SUNWhea'`, `'SUNWlibm'`, and `'SUNWtoo'`. For Solaris 2.2, GCC needs an additional seventh package: `'SUNWsprot'`.

- On Solaris 2, trying to use the linker and other tools in `/usr/ucb` to install GCC has been observed to cause trouble. For example, the linker may hang indefinitely. The fix is to remove `/usr/ucb` from your `PATH`.
- If you use the 1.31 version of the MIPS assembler (such as was shipped with Ultrix 3.1), you need to use the `-fno-delayed-branch` switch when optimizing floating point code. Otherwise, the assembler will complain when the GCC compiler fills a branch delay slot with a floating point instruction, such as `'add.d'`.
- If, on a MIPS system, you get an error message saying `"does not have gp sections for all it's [sic] sections [sic]"`, don't worry about it. This happens whenever you use GAS with the MIPS linker, but there is not really anything wrong, and it is okay to use the output file. You can stop such warnings by

installing the GNU linker. It would be nice to extend GAS to produce the `gp` tables, but they are optional, and there should not be a warning about their absence.

- In Ultrix 4.0 on the MIPS machine, `'stdio.h'` does not work with GCC at all unless it has been fixed with `fixincludes`. This causes problems in building GCC. Once GCC is installed, the problems go away. To work around this problem, when making the stage 1 compiler, specify the following option to make.

```
GCC_FOR_TARGET="./xgcc -B./ -I./include"
```

When making stage 2 and stage 3, specify the following option.

```
CFLAGS="-g -I./include"
```

- Users have reported some problems with version 2.0 of the MIPS compiler tools that were shipped with Ultrix 4.1. Version 2.10 which came with Ultrix 4.2 seems to work fine.

Users have also reported some problems with version 2.20 of the MIPS compiler tools that were shipped with RISC/OS 4.x. The earlier version 2.11 seems to work fine.

- Some versions of the MIPS linker will issue an assertion failure when linking code that uses `alloca` against shared libraries on RISC-OS 5.0, and DEC's OSF/1 systems.

This is a bug in the linker that is supposed to be fixed in future revisions. To protect against this, GCC passes `'-non_shared'` to the linker unless you pass an explicit `'-shared'` or `'-call_shared'` switch.

- On System V release 3, you may get the following error message while linking:

```
ld fatal: failed to write symbol name something
in strings table for file whatever
```

This probably indicates that the disk is full or your `ULIMIT` won't allow the file to be as large as it needs to be. This problem can also result because the kernel parameter `MAXUMEM` is too small. If so, regenerate the kernel and make the value much larger. The default value is reported to be 1024; a value of 32768 is said to work. Smaller values may also work.

- On System V, if you get an error like this,

```
/usr/local/lib/bison.simple: In function 'yyparse':
/usr/local/lib/bison.simple:625: virtual memory exhausted
```

This indicates a problem with disk space, `ULIMIT`, or `MAXUMEM`.

- Current GCC versions probably do not work on version 2 of the NeXT operating system.
- On NeXTStep 3.0, the Objective C compiler does not work, due, apparently, to a kernel bug that it happens to trigger. This problem does not happen on 3.1.

- On the Tower models 4n0 and 6n0, by default a process is not allowed to have more than one megabyte of memory. GCC cannot compile itself (or many other programs) with '-o' in that much memory.

To solve this problem, reconfigure the kernel adding the following line to the configuration file:

```
MAXMEM = 4096
```

- On HP 9000 series 300 or 400 running HP/UX release 8.0, there is a bug in the assembler that must be fixed before GCC can be built. This bug manifests during the first stage of compilation, while building 'libgcc2.a':

```
_floatdisf
ccl: warning: '-g' option not supported on this of GCC
ccl: warning: '-gl' option not supported on this version of GCC
./xgcc: Internal compiler error: program as got fatal signal 11
```

archive/cph/hpux-8.0-assembler, a patched version of the assembler, is available by anonymous ftp from altdorf.ai.mit.edu. If you have HP software support, the patch can also be obtained directly from HP, as described in the following note:

```
This is the patched assembler, to patch SR#1653-010439,
where the assembler aborts on floating point constants.
```

```
The bug is not really in the assembler, but in the shared
library version of the function "cvtnum(3c)". The bug
on "cvtnum(3c)" is SR#4701-078451. Anyway, the attached
assembler uses the archive library version of "cvtnum(3c)" and
thus does not exhibit the bug.
```

This patch is also known as PHCO 4484.

- On HP-UX version 8.05, but not on 8.07 or more recent versions, the fixproto shell script triggers a bug in the system shell. If you encounter this problem, upgrade your operating system or use BASH (the GNU shell) to run fixproto.
- Some versions of the Pyramid C compiler are reported to be unable to compile GCC. You must use an older version of GCC for bootstrapping. One indication of this problem is if you get a crash when GCC compiles the function, muldi3, in file, 'libgcc2.c'.

You may be able to succeed by getting GCC version 1, installing it, and using it to compile GCC version 2. The bug in the Pyramid C compiler does not seem to affect GCC version 1.

- On the Intel Paragon (an i860 machine), if you are using operating system version 1.0, you will get warnings or errors about redefinition of va_arg when you build GCC. If this happens, then you need to link most programs with the library, 'iclib.a'. You must also modify 'stdio.h' as follows.

```
#if defined(__i860__) && !defined(_VA_LIST)
#include <va_list.h>
```

Between the previous lines, insert the line, `#if __PGC__`. Then, after the following lines, insert the line, `#endif /* __PGC__ */`.

```
extern int vprintf(const char *, va_list );
extern int vsprintf(char *, const char *, va_list );
#endif
```

These problems don't exist in operating system version 1.1.

- On the Altos 3068, programs compiled with GCC won't work unless you fix a kernel bug. This happens using system versions V.2.2 1.0gT1 and V.2.2 1.0e and perhaps later versions as well. See the file, 'README.ALTOS'.
- You will get several sorts of compilation and linking errors on the we32k if you don't follow the special instructions. See "Configurations supported by GCC" on page 33.
- A bug in the HP/UX 8.05 (and earlier) shell will cause the `fixproto` program to report an error of the following form.

```
./fixproto: sh internal 1K buffer overflow
```

To fix this, change the first line of the `fixproto` script to look like the following declaration.

```
#!/bin/ksh
```

Cross-compiler problems

You may run into problems with cross compilation on certain machines, for many reasons.

- Cross compilation can run into trouble for certain machines because some assemblers on some target machines require floating point numbers to be written as `integer` constants in certain contexts.

The compiler writes these integer constants by examining the floating point value as an integer and printing that integer, because this is simple to write and independent of the details of the floating point representation. But this does not work if the compiler is running on a different machine with an incompatible floating point format, or even a different byte-ordering. In addition, correct constant folding of floating point values requires representing them in the target machine's format. (The C standard does not quite require this, but in practice it is the only way to win.) It is now possible to overcome these problems by defining macros such as `REAL_VALUE_TYPE`. But doing so is a substantial amount of work for each target machine. See "libgcc.a and cross-compilers" on page 50 and "Actually building the cross-compiler" on page 53 for more on macros.

- At present, the program, `mips-tfile`, which adds debug support to object files on MIPS systems, does not work in a cross compile environment.

Interoperation

The following documentation discusses various difficulties in using GNU C or GNU C++ together with other compilers or with the assemblers, linkers, libraries and debuggers on certain systems.

- Objective C does not work on the RS/6000.
- GNU C++ does not do name mangling in the same way as other C++ compilers. This means that object files compiled with one compiler cannot be used with another compiler.

This effect is intentional, to protect you from more subtle problems. Compilers differ as to many internal details of C++ implementation, including: how class instances are laid out, how multiple inheritance is implemented, and how virtual function calls are handled.

If the name encoding were made the same, your programs would link against libraries provided from other compilers—but the programs would then crash when run. Incompatible libraries are then detected at link time, rather than at run time.

- Older GDB versions sometimes fail to read the output of GCC version 2. If you have trouble, get GDB version 4.4 or later.
- DBX rejects some files produced by GCC, though it accepts similar constructs in output from PCC. Until someone can supply a coherent description of what is valid DBX input and what is not, there is nothing we can do about these problems. You are on your own.
- The GNU assembler (GAS) does not support PIC. To generate PIC code, you must use some other assembler, such as `/bin/as`.
- On some BSD systems, including some versions of Ultrix, use of profiling causes static variable destructors (currently used only in C++) not to be run.
- Use of `-I/usr/include` may cause trouble.

Many systems come with header files that won't work with GCC unless corrected by `fixincludes`. The corrected header files go in a new directory; GCC searches this directory before `/usr/include`. If you use `-I/usr/include`, this tells GCC to search `/usr/include` earlier on, before the corrected headers. The result is that you get the uncorrected header files.

Instead, you should use the following options when compiling C programs (where *target* and *version* signify what you must specify for the particular machine and its version).

```
-I/usr/local/lib/gcc-lib/target/version/include  
-I/usr/include
```

For C++ programs, GCC also uses a special directory that defines C++ interfaces to standard C subroutines. This directory is meant to be searched *before* other

standard include directories, so that it takes precedence. If you are compiling C++ programs and specifying include directories explicitly, use this option first, then the two previous options.

```
-I/usr/local/lib/g++-include
```

- On some SGI systems, when you use `'-lg1_s'` as an option, it gets translated magically to `'-lg1_s -lx11_s -lc_s'`. Naturally, this does not happen when you use GCC. You must specify all three options explicitly.
- On a SPARC, GCC aligns all values of type `double` on an 8-byte boundary, and it expects every `double` to be so aligned. The Sun compiler usually gives `double` values 8-byte alignment, with one exception: function arguments of type `double` may not be aligned.

As a result, if a function compiled with Sun CC takes the address of an argument of type `double` and passes this pointer of type `double*` to a function compiled with GCC, de-referencing the pointer may cause a fatal signal.

One way to solve this problem is to compile your entire program with GCC. Another solution is to modify the function that is compiled with Sun CC to copy the argument into a local variable; local variables are always properly aligned.

A third solution is to modify the function that uses the pointer to de-reference it using the function, `access_double`, instead of directly with `'*'` as in the following declaration.

```
inline double
access_double (double *unaligned_ptr)
{
    union d2i { double d; int i[2]; };

    union d2i *p = (union d2i *) unaligned_ptr;
    union d2i u;

    u.i[0] = p->i[0];
    u.i[1] = p->i[1];

    return u.d;
}
```

Storing into the pointer can be done likewise with the same union.

- On Solaris, the `malloc` function in the `'libmalloc.a'` library may allocate memory that is only 4 byte aligned. Since GCC on the SPARC assumes that doubles are 8 byte aligned, this may result in a fatal signal if doubles are stored in memory allocated by the `'libmalloc.a'` library.

The solution is not to use the `'libmalloc.a'` library. Instead, use `malloc` and related functions from `'libc.a'`; they do not have this problem.

- Sun forgot to include a static version of `'libdl.a'` with some versions of SunOS (mainly 4.1). This results in undefined symbols when linking static binaries (that is, if you use `'-static'`). If you see undefined symbols `_dlclose`, `_dlsym` or `_dlopen` when linking, compile and link against the file, `'mit/util/misc/dlsym.c'`, from the MIT version of X windows.
- The 128-bit long double format that the Sparc port supports currently works by using the architecturally defined quad-word floating point instructions. Since there is no hardware that supports these instructions they must be emulated by the operating system.

Long doubles do not work in Sun OS versions 4.0.3 and earlier, because the kernel emulator uses an obsolete and incompatible format. Long doubles do not work in Sun OS version 4.1.1 due to a problem in a Sun library. Long doubles do work on Sun OS versions 4.1.2 and higher, but GCC does not enable them by default. Long doubles appear to work in Sun OS 5.x (Solaris 2.x).

- On HP-UX version 9.01 on the HPPA, the HP compiler, `cc`, does not compile GCC correctly. We do not yet know why. However, GCC compiled on earlier HP-UX versions works properly on HP-UX 9.01 and can compile itself properly on 9.01.
- On the HPPA machine, ADB sometimes fails to work on functions compiled with GCC. Specifically, it fails to work on functions that use `alloca` or variable-size arrays. This is because GCC doesn't generate HP-UX unwind descriptors for such functions. It may even be impossible to generate them.
- Debugging `'-g'` is not supported on the HPPA machine, unless you use the preliminary GNU tools (see "Installing GCC" on page 23 for descriptions of the `--with-gnu-as` and `--with-gnu-ld`).
- Taking the address of a label may generate errors from the HP-UX PA assembler. GAS for the PA does not have this problem.
- Using floating point parameters for indirect calls to static functions will not work when using the HP assembler. There simply is no way for GCC to specify what registers hold arguments for static functions when using the HP assembler. GAS for the PA does not have this problem.
- In extremely rare cases involving some very large functions you may receive errors from the HP linker complaining about an out of bounds unconditional branch offset. This used to occur more often in previous versions of GCC, but is now exceptionally rare. If you should run into it, you can work around by making your function smaller.
- GCC compiled code sometimes emits warnings from the HP/UX assembler of the following form; such warnings can be safely ignored.

(warning) Use of GR3 when frame >= 8192 may cause conflict.

- The current version of the assembler (`/bin/as`) for the RS/6000 has certain problems that prevent the `-g` option in GCC from working. Note that `Makefile.in` uses `-g` by default when compiling `libgcc2.c`.

IBM has produced a fixed version of the assembler. The upgraded assembler unfortunately was not included in any of the AIX 3.2 update PTF releases (3.2.2, 3.2.3, or 3.2.3e). Users of AIX 3.1 should request PTF U403044 from IBM and users of AIX 3.2 should request PTF U416277. See the file `README.RS6000` for more details on these updates.

Test for the presence of a fixed assembler using the following command.

```
as -u < /dev/null
```

If the command exits normally, the assembler fix already is installed. If the assembler complains that `-u` is an unknown flag, you need to order the fix.

- On the IBM RS/6000, compiling code of the following form will cause the linker to report an undefined symbol, `foo`.

```
extern int foo;

... foo ...

static int foo;
```

Although this behavior differs from most other systems, it is not a bug because redefining an `extern` variable as `static` is undefined in ANSI C.

- AIX on the RS/6000 provides support (NLS) for environments outside of the United States. Compilers and assemblers use NLS to support locale-specific representations of various objects including floating-point numbers (`'.'` vs `','` for separating decimal fractions). There have been problems reported where the library linked with GCC does not produce the same floating-point formats that the assembler accepts. If you have this problem, set the `LANG` environment variable to `'C'` or `'En US'`.
- Even if you specify `'-fdollars-in-identifiers'`, you cannot successfully use `'$'` in identifiers on the RS/6000 due to a restriction in the IBM assembler. GAS supports these identifiers.
- On the RS/6000, XLC version 1.3.0.0 will miscompile `'jump.c'`. XLC version 1.3.0.1 or later fixes this problem. You can obtain XLC-1.3.0.2 by requesting PTF 421749 from IBM.
- There is an assembler bug in versions of DG/UX prior to 5.4.2.01 that occurs when the `fldcr` instruction is used. GCC uses `fldcr` on the 88100 to serialize volatile memory references. Use the option, `-mno-serialize-volatile`, if your version of the assembler has this bug.
- On VMS, GAS versions 1.38.1 and earlier may cause spurious warning messages from the linker. These warning messages complain of mismatched `psect` attributes. You can ignore them. See “Installing GCC on VMS” on page 55.

- On NewsOS version 3, if you include both of the files, `stddef.h`, and `sys/types.h`, you get an error because there are two typedefs of `size_t`. You should change '`sys/types.h`' by adding the following lines around the definition of `size_t`.

```
#ifndef _SIZE_T
#define _SIZE_T
    /* actual typedef here */
#endif
```

- On the Alliant, the system's own convention for returning structures and unions is unusual, and is not compatible with GCC no matter what options are used.
- On the IBM RT PC, the MetaWare HighC compiler (`hc`) uses a different convention for structure and union returning.

Use the option, '`-mhc-struct-return`', to tell GCC to use a convention compatible with it.

- On Ultrix, the Fortran compiler expects registers 2 through 5 to be saved by function calls. However, the C compiler uses conventions compatible with BSD UNIX: registers 2 through 5 may be clobbered by function calls. GCC uses the same convention as the Ultrix C compiler. Use these options to produce code compatible with the Fortran compiler.

```
-fcall-saved-r2      -fcall-saved-r3
-fcall-saved-r4      -fcall-saved-r5
```

- On the WE32k, you may find that programs compiled with GCC do not work with the standard shared C library. You may need to link with the ordinary C compiler. If you do so, you must specify the following options:

```
-L/usr/local/lib/gcc-lib/we32k-att-sysv/2.7.1 -lgcc -lc_s
```

The first specifies where to find the library, `libgcc.a`, specified with the '`-lgcc`' option. GCC does linking by invoking `ld`, just as `cc` does, and there is no reason why it should matter which compilation program invokes `ld`. If someone tracks this problem down, it can probably be fixed easily.

- On the Alpha, you may get assembler errors about invalid syntax as a result of floating point constants. This is due to a bug in the C library functions, `ecvt`, `fcvt` and `gcvt`. Given valid floating point numbers, they sometimes print '`NaN`'.
- On Irix 4.0.5F (and perhaps in some other versions), an assembler bug sometimes reorders instructions incorrectly when optimization is turned on. If you think this may be happening to you, try using the GNU assembler; GAS version 2.1 supports ECOFF on Irix.

Or use the '`-noasmopt`' option when you compile GCC with itself, and then again when you compile your program. (This is a temporary kludge to turn off assembler optimization on Irix.) If this proves to be what you need, edit the assembler spec in the file, `specs` so that it unconditionally passes '`-O0`' to the assembler, and never passes '`-O2`' or '`-O3`'.

Problems compiling certain programs

Certain programs have problems compiling.

- Parse errors may occur compiling X11 on a Decstation running Ultrix 4.2 because of problems in DEC's versions of the X11 header files 'x11/xlib.h' and 'x11/xutil.h'. People recommend adding '-I/usr/include/mit' to use the MIT versions of the header files, using the '-traditional' switch to turn off ANSI C, or fixing the header files by adding the following:

```
#ifdef __STDC__
#define NeedFunctionPrototypes 0
#endif
```

- If you have trouble compiling Perl on a SunOS 4 system, it may be because Perl specifies '-I/usr/ucbinclude'. This accesses the unfixed header files. Perl specifies the following options most of which are unnecessary with GCC 2.4.5 and newer versions.

```
-traditional -Dvolatile=__volatile__
-I/usr/include/sun -I/usr/ucbinclude
-fpcc-struct-return
```

You can make a properly working Perl by setting `ccflags` to '-fwritable-strings' (implied by the option, '-traditional', in the original options) and `cppflags` to empty in 'config.sh', then using the following declaration.

```
./doSH; make depend; make
```

- On various 386 UNIX systems derived from System V, including SCO, ISC, and ESIX, you may get error messages about running out of virtual memory while compiling certain programs.

You can prevent this problem by linking GCC with the GNU `malloc` (which thus replaces the `malloc` that comes with the system). GNU `malloc` is available as a separate package, and also in the file, 'src/gmalloc.c' in the GNU Emacs 19 distribution.

If you have installed GNU `malloc` as a separate library package, use the following option when you relink GCC.

```
MALLOC=/usr/local/lib/libgmalloc.a
```

Alternatively, if you have compiled 'gmalloc.c' from Emacs 19, copy the object file to 'gmalloc.o' and use the following option when you relink GCC.

```
MALLOC=gmalloc.o
```

Incompatibilities of GCC

There are several noteworthy incompatibilities between GNU C and most existing (non-ANSI) versions of C. The `'-traditional'` option eliminates many of these incompatibilities, *but not all*, by telling GNU C to behave like the other C compilers.

- GCC normally makes string constants read-only. If several identical-looking string constants are used, GCC stores only one copy of the string.

One consequence is that you cannot call `mktemp` with a string constant argument. The function, `mktemp`, always alters the string its argument points to. Another consequence is that `sscanf` does not work on some systems when passed a string constant as its format control string or input. This is because `'sscanf'` incorrectly tries to write into the string constant. Likewise `fscanf` and `scanf`.

The best solution to these problems is to change the program to use `char`-array variables with initialization strings for these purposes instead of string constants. But if this is not possible, you can use the `'-fwritable-strings'` flag, which directs GCC to handle string constants the same way most C compilers do. `'-traditional'` also has this effect, among others.

- `-2147483648` is positive.

This is because `2147483648` cannot fit in the type, `int`; so (following the ANSI C rules), its data type is `unsigned long int`. Negating this value yields `2147483648` again. GCC does not substitute macro arguments when they appear within string constants. For example, the following macro in GCC will produce output `"a"` regardless of what the argument `'a'` is.

```
#define foo(a) "a"
```

The `'-traditional'` option directs GCC to handle such cases (among others) in the old-fashioned (non-ANSI) fashion.

- When you use `setjmp` and `longjmp`, the only automatic variables guaranteed to remain valid are those declared `volatile`. This is a consequence of automatic register allocation. Consider the following function statement. In the statement, `'a'` may or may not be restored to its first value when the `longjmp` occurs.

```
jmp_buf j;

foo ()
{
    int a, b;

    a = fun1 ();

    if (setjmp (j))
        return a;
```

```

    a = fun2 ();
    /* longjmp (j) may occur in fun3. */
    return a + fun3 ();
}

```

If ‘a’ is allocated in a register, then its first value is restored; otherwise, it keeps the last value stored in it.

If you use the ‘-w’ option with the ‘-o’ option, you will get a warning when GCC thinks such a problem might be possible.

The ‘-traditional’ option directs GNU C to put variables in the stack by default, rather than in registers, in functions that call `setjmp`. This results in the behavior found in traditional C compilers.

- Programs that use preprocessing directives in the middle of macro arguments do not work with GCC. For example, a program like the following will not work:

```

foobar (
#define luser
        hack)

```

ANSI C does not permit such a construct. It would make sense to support it when ‘-traditional’ is used, but it is too much work to implement.

- Declarations of external variables and functions within a block apply only to the block containing the declaration. In other words, they have the same scope as any other declaration in the same place.

In some other C compilers, an `extern` declaration affects all the rest of the file even if it happens within a block. The ‘-traditional’ option directs GNU C to treat all `extern` declarations as global, like traditional compilers.

- In traditional C, you can combine `long`, etc., with a typedef name, as shown by the following declaration.

```

typedef int foo;
typedef long foo bar;

```

In ANSI C, this is not allowed: `long` and other type modifiers require an explicit `int`. Because this criterion is expressed by Bison grammar rules rather than C code, the ‘-traditional’ flag cannot alter it.

- PCC allows typedef names to be used as function parameters. The difficulty described previously applies here too.
- PCC allows whitespace in the middle of compound assignment operators such as ‘+=’. GCC, following the ANSI standard, does not allow this. The difficulty described previously applies here too.

- GCC complains about unterminated character constants inside of preprocessing conditionals that fail. Some programs have English comments enclosed in conditionals that are guaranteed to fail; if these comments contain apostrophes, GCC will probably report an error. For example, the following such code would produce an error.

```
#if 0
You can't expect this to work.
#endif
```

The best solution to such a problem is to put the text into an actual C comment delimited by `/*...*/`. However, `-traditional` suppresses these error messages.

- Many user programs contain the declaration `'long time()'`. In the past, the system header files on many systems did not actually declare `time`, so it did not matter what type your program declared it to return. But in systems with ANSI C headers, `time` is declared to return `time_t`, and if that is not the same as `long`, then `'long time()'` is erroneous. The solution is to change your program to use `time_t` as the return type of `time`.
- When compiling functions that return `float`, PCC converts it to a `double`. GCC actually returns a `float`. If you are concerned with PCC compatibility, you should declare your functions to return `double`; you might as well say what you mean.
- When compiling functions that return structures or unions, GCC output code normally uses a method different from that used on most versions of UNIX. As a result, code compiled with GCC cannot call a structure-returning function compiled with PCC, and vice versa.

The method used by GCC is as follows: a structure or union which is 1, 2, 4 or 8 bytes long is returned like a scalar. A structure or union with any other size is stored into an address supplied by the caller (usually in a special, fixed register, but on some machines it is passed on the stack). The machine-description macros, `STRUCT_VALUE` and `STRUCT_INCOMING_VALUE`, tell GCC where to pass this address.

By contrast, PCC on most target machines returns structures and unions of any size by copying the data into an area of static storage, and then returning the address of that storage as if it were a pointer value. The caller must copy the data from that memory area to the place where the value is wanted. GCC does not use this method because it is slower and non re-entrant.

On some newer machines, PCC uses a reentrant convention for all structure and union returning. GCC on most of these machines uses a compatible convention when returning structures and unions in memory, but still returns small structures and unions in registers.

You can tell GCC to use a compatible convention for all structure and union

returning with the option `'-fpcc-struct-return'`.

- GNU C complains about program fragments such as `0x74ae-0x4000` which appear to be two hexadecimal constants separated by the minus operator. Actually, this string is a single preprocessing token. Each such token must correspond to one token in C. Since this does not, GNU C prints an error message. Although it may appear obvious that what is meant, is an operator and two values, the ANSI C standard specifically requires that this be treated as erroneous. A preprocessing token is a preprocessing number if it begins with a digit and is followed by letters, underscores, digits, periods and `'e+'`, `'e-'`, `'E+'`, or `'E-'` character sequences. To make the previous program fragment valid, place whitespace in front of the minus sign. This whitespace will end the preprocessing number.

Fixed header files

GCC needs to install corrected versions of some system header files. This is because most target systems have some header files that won't work with GCC unless they are changed. Some have bugs, some are incompatible with ANSI C, and some depend on special features of other compilers.

Installing GCC automatically creates and installs the fixed header files, by running a program called `fixincludes` (or for certain targets an alternative such as `'fixinc.svr4'`). Normally, you don't need to pay attention to this. But there are cases where it doesn't do the right thing automatically.

- If you update the system's header files, such as by installing a new system version, the fixed header files of GCC are not automatically updated. The easiest way to update them is to reinstall GCC. (If you want to be clever, look in the makefile and you can find a shortcut.)
- On some systems, in particular SunOS 4, header file directories contain machine-specific symbolic links in certain places. This makes it possible to share most of the header files among hosts running the same version of SunOS 4 on different machine models.

The programs that fix the header files do not understand this special way of using symbolic links; therefore, the directory of fixed header files is good only for the machine model used to build it.

In SunOS 4, only programs that look inside the kernel will notice the difference between machine models. Therefore, for most purposes, you need not be concerned about this.

It is possible to make separate sets of fixed header files for the different machine models, and arrange a structure of symbolic links so as to use the proper set, but you'll have to do this by hand.

- On Lynxos, GCC by default does not fix the header files. This is because bugs in the shell cause the `fixincludes` script to fail.

This means you will encounter problems due to bugs in the system header files. It may be no comfort that they aren't GCC's fault, but it does mean that there's nothing for us to do about them.

Standard libraries

GCC by itself attempts to be what the ISO/ANSI C standard calls a *conforming freestanding implementation*. This means all ANSI C language features are available, as well as the contents of `'float.h'`, `'limits.h'`, `'stdarg.h'`, and `'stddef.h'`. The rest of the C library is supplied by the vendor of the operating system. If that C library doesn't conform to the C standards, then your programs might get warnings (especially when using `'-Wall'`) that you don't expect.

For example, the `sprintf` function on SunOS 4.1.3 returns `char *` while the C standard says that `sprintf` returns an `int`. The `fixincludes` program could make the prototype for this function match the Standard, but that would be wrong, since the function will still return `char *`.

If you need a Standard compliant library, then you need to find one, as GCC does not provide one. The GNU C library (called `glibc`) has been ported to a number of operating systems, and provides ANSI/ISO, POSIX, BSD and SystemV compatibility. You could also ask your operating system vendor if newer libraries are available.

Disappointments and misunderstandings

These problems are regrettable, but we don't know any practical way around them.

- Certain local variables aren't recognized by debuggers when you compile with optimization.

This occurs because sometimes GCC optimizes the variable out of existence. There is no way to tell the debugger how to compute the value such a variable "would have had" and it is not clear that would be desirable anyway. So GCC simply does not mention the eliminated variable when it writes debugging information.

You have to expect a certain amount of disagreement between the executable and your source code, when you use optimization.

- Users often think it is a bug when GCC reports an error for code like the following example.

```
int foo (struct mumble *);

struct mumble { ... };

int foo (struct mumble *x)
{ ... }
```


This code really is erroneous, because the scope of `struct mumble` in the prototype is limited to the argument list containing it. (It does not refer to the `struct mumble` defined with file scope in the following descriptions—they are two unrelated types with similar names in different scopes.)

In the definition of `foo`, the file-scope type is used because that is available to be inherited. Thus, the definition and the prototype do not match, and you get an error.

This behavior may seem silly, but it's what the ANSI standard specifies. It is easy enough for you to make your code work by moving the definition of `struct mumble` above the prototype. It's not worth being incompatible with ANSI C just to avoid an error like the previous example.

- Accesses to bitfields even in volatile objects works by accessing larger objects, such as a byte or a word. You cannot rely on what size of object is accessed in order to read or write the bitfield; it may even vary for a given bitfield according to the precise usage.

If you care about controlling the amount of memory that is accessed, use volatile but do not use bitfields.

- GCC comes with shell scripts to fix certain known problems in system header files. They install corrected copies of various header files in a special directory where only GCC will normally look for them. The scripts adapt to various systems by searching all the system header files for the problem cases that we know about.

If new system header files are installed, nothing automatically arranges to update the corrected header files. You will have to reinstall GCC to fix the new header files. More specifically, go to the build directory and delete the files `'stmp-fixinc'` and `'stmp-headers'`, and the subdirectory `include`; then do `'make install'` again.

- On 68000 systems and x86 systems, you can get paradoxical results if you test the precise values of floating point numbers. For example, you can find that a floating point value, which is not a NaN, is not equal to itself. This results from the fact that the floating point registers hold a few more bits of precision than fit in a `double` in memory. Compiled code moves values between memory and floating point registers at its convenience, and moving them into memory truncates them. You can partially avoid this problem by using the `'-ffloat-store'` option (see “Options that control optimization” on page 111).
- On the MIPS, variable argument functions using `'varargs.h'` cannot have a floating point value for the first argument. The reason for this is that in the absence of a prototype in scope, if the first argument is a floating point, it is passed in a floating point register, rather than an integer register.

If the code is rewritten to use the ANSI standard `'stdarg.h'` method of variable

arguments, and the prototype is in scope at the time of the call, everything will work fine.

- On the H8/300 and H8/300H, variable argument functions must be implemented using the ANSI standard ‘`stdarg.h`’ method of variable arguments. Furthermore, calls to functions using ‘`stdarg.h`’ variable arguments must have a prototype for the called function in scope at the time of the call.

Common misunderstandings with GNU C++

C++ is a complex language and an evolving one, and its standard definition (the ANSI C++ draft standard) is also evolving. As a result, your C++ compiler may occasionally surprise you, even when its behavior is correct. The following documentation discusses some areas that frequently give rise to questions of this nature.

Declare and define static members

When a class has static data members, it is not enough to declare the static member; you must also define it. Use the following example, for instance, for such declarations.

```
class Foo
{
: : :
void method();
static int bar;
};
```

This declaration only establishes that the class `Foo` has an `int` named `Foo::bar`, and a member function named `Foo::method`. But you still need to define both `method` and `bar` elsewhere. According to the draft ANSI standard, you must supply an initializer in one (and only one) source file, such as the following example shows.

```
int Foo::bar = 0;
```

Other C++ compilers may not correctly implement the standard behavior. As a result, when you switch to `g++` from one of these compilers, you may discover that a program that appeared to work correctly in fact does not conform to the standard: `g++` reports as undefined symbols any static data members that lack definitions.

Temporaries may vanish before you expect

It is dangerous to use pointers or references to portions of a temporary object.

The compiler may very well delete the object before you expect it to, leaving a pointer to garbage. The most common place where this problem crops up is in classes like the `libg++` `string` class, that define a conversion function to `type`, `char *`, or to `type`, `const char *`.

However, any class that returns a pointer to some internal structure is potentially subject to this problem. For instance, a program may use a function, `strfunc`, that returns `String` objects, and another function, `charfunc`, that operates on pointers to

char, as in the following example's declaration.

```
String strfunc ();
void charfunc (const char *);
```

In such a situation, it may seem natural to write `'charfunc (strfunc());'` based on the knowledge that class `String` has an explicit conversion to `char` pointers. However, what really happens is akin to `'charfunc (strfunc().convert());'` where the conversion method uses a function to do the same data conversion normally performed by a cast. Since the last use of the temporary `String` object is the call to the conversion function, the compiler may delete that object before actually calling `charfunc`. The compiler has no way of knowing that deleting the `String` object will invalidate the pointer. The pointer then points to garbage, so that by the time `charfunc` is called, it gets an invalid argument. Code like this may run successfully under some other compilers, especially those that delete temporaries relatively late. However, the GNU C++ behavior is also standard-conforming, so if your program depends on late destruction of temporaries it is not portable.

If you think this is surprising, you should be aware that the ANSI C++ committee continues to debate the lifetime-of-temporaries problem.

For now, at least, for instance, the following declaration defines the safe way to write such code by giving the temporary a name, forcing it to remain until the end of the scope of the name.

```
String& tmp = strfunc ();
charfunc (tmp);
```

protoize and unprotoize warnings

The conversion programs, `protoize` and `unprotoize`, can sometimes change a source file in a way that won't work unless you rearrange it.

- `protoize` can insert references to a type name or type tag before the definition, or in a file where they are not defined.

If this happens, compiler error messages should show you where the new references are, so fixing the file by hand is straightforward.

- There are some C constructs which `protoize` cannot figure out. For example, it can't determine argument types for declaring a pointer-to- function variable; this you must do by hand. `protoize` inserts a comment containing '???' each time it finds such a variable; so you can find all such variables by searching for this string. ANSI C does not require declaring the argument types of pointer-to-function types.
- Using `unprotoize` can easily introduce bugs. If the program relied on prototypes to bring about conversion of arguments, these conversions will not take place in the program without prototypes. One case in which you can be sure `unprotoize` is safe is when you are removing prototypes that were made with `protoize`; if the

program worked before without any prototypes, it will work again without them. You can find all the places where this problem might occur by compiling the program with the `-Wconversion` option. It prints a warning whenever an argument is converted.

- Both conversion programs can be confused if there are macro calls in and around the text to be converted. In other words, the standard syntax for a declaration or definition must not result from expanding a macro. This problem is inherent in the design of C and cannot be fixed. If only a few functions have confusing macro calls, you can easily convert them manually.
- `protoize` cannot get the argument types for a function whose definition was not actually compiled due to preprocessing conditionals. When this happens, `protoize` changes nothing in regard to such a function. `protoize` tries to detect such instances and warn about them.

You can generally work around this problem by using `protoize` step by step, each time specifying a different set of `-D` options for compilation, until all of the functions have been converted. There is no automatic way to verify that you have got them all, however.

- Confusion may result if there is an occasion to convert a function declaration or definition in a region of source code where there is more than one formal parameter list present. Thus, attempts to convert code containing multiple (conditionally compiled) versions of a single function header (in the same vicinity) may not produce the desired (or expected) results.

If you plan on converting source files which contain such code, it is recommended that you first make sure that each conditionally compiled region of source code which contains an alternative function header also contains at least one additional follower token (past the final right parenthesis of the function header). This should circumvent the problem.

- `unprotoize` can become confused when trying to convert a function definition or declaration which contains a declaration for a pointer-to-function formal argument which has the same name as the function being defined or declared. We recommend you avoid such choices of formal parameter names.
- You might also want to correct some of the indentation by hand and break long lines. (The conversion programs don't write lines longer than eighty characters in any case.)

Certain changes we don't want to make

The following documentation lists changes that people frequently request, but which we do not make because we think GCC is better without them.

- Checking the number and type of arguments to a function which has an old-fashioned definition and no prototype.
Such a feature would work only occasionally—only for calls that appear in the same file as the called function, following the definition. The only way to check all calls reliably is to add a prototype for the function. But adding a prototype eliminates the motivation for this feature. So the feature is not worthwhile.
- Warning about using an expression whose type is signed as a shift count.
Shift count operands are probably signed more often than unsigned. Warning about this would cause far more annoyance than good.
- Warning about assigning a signed value to an unsigned variable.
Such assignments must be very common; warning about them would cause more annoyance than good.
- Warning about unreachable code.
It's very common to have unreachable code in machine-generated programs. For example, this happens normally in some files of GNU C itself.
- Warning when a non-void function value is ignored.
Coming from a Lisp background, the idea seems silly that there is something dangerous about discarding a value. There are functions that return values which some callers may find useful; it makes no sense to clutter the program with a cast to `void` whenever the value isn't useful.
- Assuming (for optimization) that the address of an external symbol is never zero.
This assumption is false on certain systems when `#pragma weak` is used.
- Making `'-fshort-enums'` the default.
This would cause storage layout to be incompatible with most other C compilers. And it doesn't seem very important, given that you can get the same result in other ways. The case where it matters most is when the enumeration-valued object is inside a structure, and in that case you can specify a field width explicitly.
- Making bitfields unsigned by default on particular machines where “the ABI standard” says to do so.
The ANSI C standard leaves it up to the implementation whether a bitfield declared plain `int` is signed or not. This in effect creates two alternative dialects of C.
The GNU C compiler supports both dialects; you can specify the signed dialect with `'-fsigned-bitfields'` and the unsigned dialect with

'`funsigned-bitfields`'. However, this leaves open the question of which dialect to use by default.

Currently, the preferred dialect makes plain bitfields signed, because this is simplest. Since `int` is the same as `signed int` in every other context, it is cleanest for them to be the same in bitfields as well.

Some computer manufacturers have published Application Binary Interface standards which specify that plain bitfields should be unsigned. It is a mistake, however, to say anything about this issue in an ABI. This is because the handling of plain bitfields distinguishes two dialects of C. Both dialects are meaningful on every type of machine. Whether a particular object file was compiled using signed bitfields or unsigned is of no concern to other object files, even if they access the same bitfields in the same data structures.

A given program is written in one or the other of these two dialects. The program stands a chance to work on most any machine if it is compiled with the proper dialect. It is unlikely to work at all if compiled with the wrong dialect.

Many users appreciate the GNU C compiler because it provides an environment that is uniform across machines. These users would be inconvenienced if the compiler treated plain bitfields differently on certain machines.

Occasionally users write programs intended only for a particular machine type. On these occasions, the users would benefit if the GNU C compiler were to support by default the same dialect as the other compilers on that machine. But such applications are rare. And users writing a program to run on more than one type of machine cannot possibly benefit from this kind of compatibility.

This is why GCC does and will treat plain bitfields in the same fashion on all types of machines (by default). There are some arguments for making bitfields unsigned by default on all machines. If, for example, this becomes a universal de facto standard, it would make sense for GCC to go along with it. This is a future consideration.

(Of course, users strongly concerned about portability should indicate explicitly in each bitfield whether it is signed or not. In this way, they write programs which have the same meaning in both C dialects.)

- Undefined `__STDC__` when '`-ansi`' is not used.

Currently, GCC defines `__STDC__` as long as you don't use '`-traditional`'. This provides good results in practice.

Programmers normally use conditionals on `__STDC__` to ask whether it is safe to use certain features of ANSI C, such as function prototypes or ANSI token concatenation. Since plain GCC supports all the features of ANSI C, the correct answer to these questions is "yes." Some users try to use `__STDC__` to check for

the availability of certain library facilities. This is actually incorrect usage in an ANSI C program, because the ANSI C standard says that a conforming freestanding implementation should define `__STDC__` even though it does not have the library facilities. `'gcc -ansi -pedantic'` is a conforming freestanding implementation, and it is therefore required to define `__STDC__`, even though it does not come with an ANSI C library.

Sometimes people say that defining `__STDC__` in a compiler that does not completely conform to the ANSI C standard somehow violates the standard. This is illogical. The standard is a standard for compilers that claim to support ANSI C, such as `'gcc -ansi'`—not for other compilers such as plain `gcc`. Whatever the ANSI C standard says is relevant to the design of plain `'gcc'` without `'-ansi'` only for pragmatic reasons, not as a requirement.

- Undefined `__STDC__` in C++: programs written to compile with C++-to-C translators get the value of `__STDC__` that goes with the C compiler that is subsequently used. These programs must test `__STDC__` to determine what kind of C preprocessor that compiler uses: whether they should concatenate tokens in the ANSI C fashion or in the traditional fashion. These programs work properly with GNU C++ if `__STDC__` is defined. They would not work otherwise.

In addition, many header files are written to provide prototypes in ANSI C but not in traditional C. Many of these header files can work without change in C++ provided `__STDC__` is defined. If `__STDC__` is not defined, they will all fail, and will all need to be changed to test explicitly for C++ as well.

- Deleting empty loops: GCC does not delete empty loops because the most likely reason you would put one in a program is to have a delay. Deleting them will not make real programs run any faster, so it would be pointless.

It would be different if optimization of a non-empty loop could produce an empty one. But this generally can't happen.

- Making side effects happen in the same order as in some other compiler. It is never safe to depend on the order of evaluation of side effects. For example, a function call like this may very well behave differently from one compiler to another:

```
void func (int, int);

int i = 2;
func (i++, i++);
```

There is no guarantee (in either the C or the C++ standard language definitions) that the increments will be evaluated in any particular order. Either increment might happen first. `func` might get the arguments `'2, 3'`, or it might get `'3, 2'`, or even `'2, 2'`.

- Not allowing structures with volatile fields in registers. Strictly speaking, there is no prohibition in the ANSI C standard against allowing structures with volatile

fields in registers, but it does not seem to make any sense and is probably not what you wanted to do. So the compiler will give an error message in this case.

Warning messages and error messages

The GNU compiler can produce two kinds of diagnostics: errors and warnings. Each kind has a different purpose:

- Errors report problems that make it impossible to compile your program. GCC reports errors with the source filename and line number where the problem is apparent.
- Warnings report other unusual conditions in your code that may indicate a problem, although compilation can (and does) proceed. Warning messages also report the source filename and line number, but include the text, `warning:`, to distinguish them from error messages.

Warnings may indicate danger points where you should check to make sure that your program really does what you intend; or the use of obsolete features; or the use of nonstandard features of GNU C or C++. Many warnings are issued only if you ask for them, with one of the `-w` options (for instance, `-Wall` requests a variety of useful warnings). GCC always tries to compile your program if possible; it never gratuitously rejects a program whose meaning is clear merely because (for instance) it fails to conform to a standard. In some cases, however, the C and C++ standards specify that certain extensions are forbidden, and a diagnostic must be issued by a conforming compiler. The `-pedantic` option tells GCC to issue warnings in such cases; `-pedantic-errors` says to make them errors instead. This does not mean that all non-ANSI constructs get warnings or errors. See “Options to request or suppress warnings” on page 93 for more detail on these and related command-line options.

Reporting bugs

Your reporting of problems or *bugs* plays an essential role in making GCC reliable. When you encounter a problem, the first thing to do is to see if it is already known. See “Known causes of trouble with GCC” on page 293. If it isn’t known, report the problem.

Reporting a bug may have a solution, or someone may have to fix it. If it does not have a solution, look in the service directory; see “How to get help with GCC” on page 329. In any case, the principal function of a bug report is to help the entire community by making the next version of GCC work better. Bug reports are your contribution to the maintenance of GCC.

Since the maintainers are very overloaded, we cannot respond to every bug report. However, if the bug has not been fixed, we are likely to send you a patch and ask you to tell us whether it works.

In order for a bug report to serve its purpose, you must include the information that makes for fixing the bug.

Have you found a bug?

If you are not sure whether you have found a bug, the following guidelines will help:

- If the compiler gets a fatal signal, for any input whatever, that is a compiler bug. Reliable compilers never crash.
- If the compiler produces invalid assembly code, for any input whatever (except an `asm` statement), that is a compiler bug, unless the compiler reports errors (not just warnings) which would ordinarily prevent the assembler from being run.
- If the compiler produces valid assembly code that does not correctly execute the input source code, that is a compiler bug.

However, you must double-check to make sure, because you may have run into an incompatibility between GCC and traditional C (see “Incompatibilities of GCC” on page 306). These incompatibilities might be considered bugs, but they are inescapable consequences of valuable features. Or you may have a program whose behavior is undefined, which happened by chance to give the desired results with another C or C++ compiler.

For example, in many nonoptimizing compilers, you can write `‘x;’` at the end of a function instead of `‘return x;’`, with the same results. But the value of the function is undefined if return is omitted; it is not a bug when GCC produces different results.

Problems often result from expressions with two increment operators, as in `‘f(*p++, *p++)’`. Your previous compiler might have interpreted that expression the way you intended; GCC might interpret it another way. Neither compiler is wrong. The bug is in your code.

After you have localized the error to a single source line, it should be easy to check for these things. If your program is correct and well defined, you have found a compiler bug.

- If the compiler produces an error message for valid input, that is a compiler bug.
- If the compiler does not produce an error message for invalid input, that is a compiler bug. However, you should note that your idea of “invalid input” might be my idea of “an extension” or “support for traditional practice.”
- If you are an experienced user of C or C++ compilers, your suggestions for improvement of GCC or G++ are welcome in any case.

Where to report bugs

Send bug reports for GCC to:

`bug-gcc@prep.ai.mit.edu`

Send bug reports for G++ to:

`bug-g++@prep.ai.mit.edu`

If your bug involves the C++ class library `libg++`, send mail to:

`bug-lib-g++@prep.ai.mit.edu`

If you're not sure, you can send the bug report to both lists.

Do not send bug reports to 'help-gcc@prep.ai.mit.edu' or to the newsgroup, 'gnu.gcc.help'. Most users of GCC do not want to receive bug reports. Those that do, have asked to be on 'bug-gcc' and/or 'bug-g++'.

The mailing lists, 'bug-gcc' and 'bug-g++', both have newsgroups which serve as repeaters: 'gnu.gcc.bug' and 'gnu.g++.bug'. Each mailing list and its newsgroup carry exactly the same messages.

Often people think of posting bug reports to the newsgroup instead of mailing them. This appears to work, but it has one problem which can be crucial: *a newsgroup posting does not contain a mail path back to the sender*. Thus, if maintainers need more information, they may be unable to reach you. For this reason, you should always send bug reports by mail to the proper mailing list.

As a last resort, send bug reports on paper to:

GNU Compiler Bugs
Free Software Foundation
59 Temple Place Suite 330
Boston, MA 02111-1307
USA

How to report bugs

The fundamental principle of reporting bugs usefully is this: *report all the facts*. If you are not sure whether to state a fact or leave it out, state it!

Often people omit facts because they think they know what causes the problem and they conclude that some details don't matter. Thus, you might assume that the name of the variable you use in an example does not matter. Well, probably it doesn't, but one cannot be sure. Perhaps the bug is a stray memory reference which happens to fetch from the location where that name is stored in memory; perhaps, if the name were different, the contents of that location would fool the compiler into doing the right thing despite the bug. Play it safe and give a specific, complete example. That is the easiest thing for you to do, and the most helpful.

Keep in mind that the purpose of a bug report is to enable someone to fix the bug if it is not known. It isn't very important what happens if the bug is already known. Therefore, always write your bug reports on the assumption that the bug is not known.

Sometimes people give a few sketchy facts and ask, "Does this ring a bell?" This cannot help us fix a bug, so it is basically useless. We respond by asking for enough details to enable us to investigate. You might as well expedite matters by sending them to begin with.

Try to make your bug report self-contained. If we have to ask you for more information, it is best if you include all the previous information in your response, as well as the information that was missing.

Please report each bug in a separate message. This makes it easier for us to track which bugs have been fixed and to forward your bugs reports to the appropriate maintainer.

Do not compress and encode any part of your bug report using programs such as 'uuencode'. If you do so it will slow down the processing of your bug. If you must submit multiple large files, use 'shar', which allows us to read your message without having to run any decompression programs.

To enable someone to investigate the bug, you should include all the following things.

- ***The version of GCC***

You can get this by running the program with the '-v' option.

Without this, we won't know whether there is any point in looking for the bug in the current version of GCC.

- ***A complete input file that will reproduce the bug***

If the bug is in the C preprocessor, send a source file and any header files that it requires. If the bug is in the compiler proper (cc1), run your source file through the C preprocessor by doing 'gcc -E sourcefile > outfile', then include the contents of outfile in the bug report. (When you do this, use the same '-I', '-D' or '-U' options that you used in actual compilation.)

A single statement is not enough of an example. In order to compile it, it must be embedded in a complete file of compiler input; and the bug might depend on the details of how this is done.

Without a real example one can compile, all anyone can do about your bug report is wish you luck. It would be futile to try to guess how to provoke the bug. For example, bugs in register allocation and reloading frequently depend on every little detail of the function they happen in.

Even if the input file that fails comes from a GNU program, you should still send the complete test case. Don't ask the *GCC* maintainers to do the extra work of obtaining the program in question—they are all overworked as it is. Also, the problem may depend on what is in the header files on your system; it is unreliable for the *GCC* maintainers to try the problem with the header files available to them. By sending CPP output, you can eliminate this source of uncertainty and save us a certain percentage of wild goose chases.

- ***The command arguments you gave GCC or G++ to compile that example and observe the bug***

For example, did you use '-o'? To guarantee you won't omit something important, list all the options.

If we were to try to guess the arguments, we would probably guess wrong and then we would not encounter the bug.

- ***The type of machine you are using, and the operating system name and version number***

- ***The operands you gave to the configure command when you installed the compiler***

- ***A complete list of any modifications you have made to the compiler source***

There is no promise to investigate the bug unless it happens in an unmodified compiler. But if you've made modifications and don't tell us, then you are sending us on a wild goose chase.

Be precise about these changes. A description in English is not enough—send a context `diff` for them.

Adding files of your own (such as a machine description for a machine we don't support) is a modification of the compiler source.

- ***Details of any other deviations from the standard procedure for installing GCC***

- ***A description of what behavior you observe that you believe is incorrect***

For example, "The compiler gets a fatal signal" or "The assembler instruction at line 208 in the output is incorrect" are messages of a fatal signal.

Of course, if the bug is of the compiler getting a fatal signal, then one can't miss it. But if the bug is incorrect output, the maintainer might not notice unless it is

glaringly wrong. None of us has time to study all the assembler code from a 50-line C program just on the chance that one instruction might be wrong. *We need you to do this part!*

Even if the problem you experience is a fatal signal, you should still say so explicitly. Suppose something strange is going on, such as, your copy of the compiler is out of synch, or you have encountered a bug in the C library on your system. (This has happened!) Your copy might crash and the copy here would not. If you *said* to expect a crash, then when the compiler here fails to crash, we would know that the bug was not happening. If you don't say to expect a crash, then we would not know whether the bug was happening. We would not be able to draw any conclusion from our observations.

If the problem is a diagnostic when compiling GCC with some other compiler, say whether it is a warning or an error.

Often the observed symptom is incorrect output when your program is run. Sad to say, this is not enough information unless the program is short and simple. None of us has time to study a large program to figure out how it would work if compiled correctly, much less which line of it was compiled wrong. So you will have to do that. Tell us which source line it is, and what incorrect result happens when that line is executed. A person who understands the program can find this as easily as finding a bug in the program itself.

- ***If you send examples of assembler code output from GCC or G++, please use ‘-g’ when you make them***

The debugging information includes source line numbers which are essential for correlating the output with the input.

- ***If you wish to mention something in the GCC source, refer to it by context, not by line number***

The line numbers in the development sources don't match those in your sources. Your line numbers would convey no useful information to the maintainers.

- ***Additional information from a debugger***

This might enable someone to find a problem on a machine which he does not have available. However, you need to think when you collect this information if you want it to have any chance of being useful.

For example, many people send just a backtrace, but that is never useful by itself. A simple backtrace with arguments conveys little about GCC because the compiler is largely data-driven; the same functions are called over and over for different RTL insns, doing different things depending on the details of the insn.

Most of the arguments listed in the backtrace are useless because they are pointers to RTL list structure. The numeric values of the pointers, which the debugger prints in the backtrace, have no significance whatever; all that matters is the contents of the objects they point to (and most of the contents are other such

pointers).

In addition, most compiler passes consist of one or more loops that scan the RTL insn sequence. The most vital piece of information about such a loop—which insn it has reached—is usually in a local variable, not in an argument.

What you need to provide in addition to a backtrace are the values of the local variables for several stack frames up. When a local variable or an argument is an RTX, first print its value and then use the GDB command `pr` to print the RTL expression that it points to. (If GDB doesn't run on your machine, use your debugger to call the function `debug_rtx` with the RTX as an argument.) In general, whenever a variable is a pointer, its value is no use without the data it points to.

What follows are some things that are not necessary.

- ***A description of the envelope of the bug***

Often people who encounter a bug spend a lot of time investigating which changes to the input file will make the bug go away and which changes will not affect it. This is often time consuming and not very useful, because the way we will find the bug is by running a single example under the debugger with breakpoints, not by pure deduction from a series of examples. You might as well save your time for something else.

Of course, if you can find a simpler example to report instead of the original one, that is a convenience.

Errors in the output will be easier to spot, running under the debugger will take less time, etc.

Most GCC bugs involve just one function, so the most straightforward way to simplify an example is to delete all the function definitions except the one where the bug occurs. Those earlier in the file may be replaced by external declarations if the crucial function depends on them. (Exception: inline functions may affect compilation of functions defined later in the file.)

However, simplification is not vital; if you don't want to do this, report the bug anyway and send the entire test case you used.

- ***Conditionals***

In particular, some people insert conditionals (`#ifdef BUG`) around a statement which, if removed, makes the bug not happen. These are just clutter; we won't pay any attention to them anyway. Besides, you should send us cpp output, and that can't have conditionals.

- ***A patch for the bug***

A patch for the bug is useful if it is a good one. But don't omit the necessary information, such as the test case, on the assumption that a patch is all we need.

We might see problems with your patch and decide to fix the problem another way, or we might not understand it at all.

Sometimes with a program as complicated as GCC it is very hard to construct an example that will make the program follow a certain path through the code.

If you don't send the example, we won't be able to construct one, so we won't be able to verify that the bug is fixed.

And if we can't understand what bug you are trying to fix, or why your patch should be an improvement, we won't install it. A test case will help us to understand.

See "Sending patches for GCC" on page 327 for guidelines on how to make it easy for us to understand and install your patches.

- ***A guess about what the bug is or what it depends on***

Such guesses are usually wrong. Even I can't guess right about such things without first using the debugger to find the facts.

- ***A core dump file***

We have no way of examining a core dump for your type of machine unless we have an identical system—and if we do have one, we should be able to reproduce the crash ourselves.

Sending patches for GCC

If you would like to write bug fixes or improvements for the GCC compiler, that is very helpful.

Send suggested fixes to the bug report mailing list:

`bug-gcc@prep.ai.mit.edu`

Please follow these guidelines so we can study your patches efficiently. If you don't follow these guidelines, your information might still be useful, but using it will take extra work. Maintaining GCC is a lot of work in the best of circumstances, and we can't keep up unless you do your best to help.

- Send an explanation with your changes of what problem they fix or what improvement they bring about. For a bug fix, just include a copy of the bug report, and explain why the change fixes the bug.

(Referring to a bug report is not as good as including it, because then we will have to look it up, and we have probably already deleted it if we've already fixed the bug.)

- Always include a proper bug report for the problem you think you have fixed. We need to convince ourselves that the change is right before installing it. Even if it is right, we might have trouble judging it if we don't have a way to reproduce the problem.
- Include all the comments that are appropriate to help people reading the source in the future understand why this change was needed.
- Don't mix together changes made for different reasons. Send them *individually*.

If you make two changes for separate reasons, then we might not want to install them both. We might want to install just one. If you send them all jumbled together in a single set of diffs, we have to do extra work to disentangle them—to figure out which parts of the change serve which purpose. If we don't have time for this, we might have to ignore your changes entirely.

If you send each change as soon as you have written it, with its own explanation, then the two changes never get tangled up, and we can consider each one properly without any extra work to disentangle them. Ideally, each change you send should be impossible to subdivide into parts that we might want to consider separately, because each of its parts gets its motivation from the other parts.

- Send each change as soon as that change is finished. Sometimes people think they are helping us by accumulating many changes to send them all together. As explained above, this is absolutely the worst thing you could do.

Since you should send each change separately, you might as well send it right away. That gives us the option of installing it immediately if it is important.

- Use `'diff -c'` to make your diffs. Diffs without context are hard for us to install reliably. More than that, they make it hard for us to study the diffs to decide whether we want to install them. Unidiff format is better than diffs without context, but not as easy to read as `'-c'` format.

If you have GNU `diff`, use `'diff -cp'`, in order to show the name of the function in which each change occurs.

- Write the change log entries for your changes. We get lots of changes, and we don't have time to do all the change log writing ourselves.

Read the `'ChangeLog'` file to see what sorts of information to put in, and to learn the style that we use. The purpose of the change log is to show people where to find what was changed. So you need to be specific about what functions you changed; in large functions, it's often helpful to indicate where within the function the change was.

On the other hand, once you have shown people where to find the change, you need not explain its purpose. Thus, if you add a new function, all you need to say about it is that it is new. If you feel that the purpose needs explaining, it probably does—but the explanation will be much more useful if you put it in comments in the code.

If you would like your name to appear in the header line for who made the change, send us the header line.

- When you write the fix, keep in mind that we can't install a change that would break other systems.

People often suggest fixing a problem by changing machine-independent files such as `'toplev.c'` to do something special that a particular system needs. Sometimes it is totally obvious that such changes would break GCC for almost all users. We can't possibly make a change like that. At best it might tell us how to write another patch that would solve the problem acceptably.

Sometimes people send fixes that *might* be an improvement in general—but it is hard to be sure of this. It's hard to install such changes because we have to study them very carefully. Of course, a good explanation of the reasoning by which you concluded the change was correct can help convince us.

The safest changes are changes to the configuration files for a particular machine. These are safe because they can't create new bugs on other machines.

Please help us keep up with the workload by designing the patch in a form that is good to install.

How to get help with GCC

If you still need help installing, using or changing GCC, there are several ways to find it, if they aren't in this documentation.

- Send a message to a suitable network mailing list. First try:

`bug-gcc@prep.ai.mit.edu`

If that brings no response, try:

`help-gcc@prep.ai.mit.edu.`

- Look in the service directory for someone who might help you for a fee. The service directory is found in the file named 'SERVICE' in the GCC distribution.
- Send a message to `bugs@cygnus.com`.

The C Preprocessor

Copyright © 1988-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the sections entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Published by: **Free Software Foundation**
59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

1

Overview of the C preprocessor

The C preprocessor is a *macro processor* that is used automatically by the C compiler to transform your program before actual compilation. It is called a macro processor because it allows you to define macros, which are abbreviations for longer constructs.

The following documentation discusses the GNU C preprocessor, the C-compatible compiler preprocessor.

- “What the C preprocessor provides” on page 334
- “Transformations made globally” on page 335
- “Preprocessing directives” on page 337
- “Header files” on page 339
- “Macros” on page 347
- “Conditionals” on page 371
- “Combining source files” on page 381
- “Other preprocessing directives” on page 383
- “C preprocessor output” on page 385e
- “Invoking the C preprocessor” on page 387

What the C preprocessor provides

The C preprocessor provides the following four separate facilities that you can use as you see fit.

- ***Inclusion of header files***

These are files of declarations that can be substituted into your program.

- ***Macro expansion***

You can define *macros*, which are abbreviations for arbitrary fragments of C code, and then the C preprocessor will replace the macros with their definitions throughout the program.

- ***Conditional compilation***

Using special preprocessing directives, you can include or exclude parts of the program according to various conditions.

- ***Line control***

If you use a program to combine or rearrange source files into an intermediate file, which is then compiled, you can use line control to inform the compiler of each source line's origin.

C preprocessors vary in some details. The GNU C preprocessor provides a superset of the features of ANSI Standard C.

ANSI Standard C requires the rejection of many harmless constructs commonly used by today's C programs. Such incompatibility would be inconvenient for users, so the GNU C preprocessor is configured to accept these constructs by default. To get ANSI Standard C, you must use the options, `-trigraphs`, `-undef` and `-pedantic`, but in practice the consequences of having strict ANSI Standard C make it undesirable to follow this practice. For more details, see "Invoking the C preprocessor" on page 387.

2

Transformations made globally

Most C preprocessor features are inactive unless you give specific directives to request their use. (Preprocessing directives are lines starting with #; see “Preprocessing directives” on page 337). But there are three transformations that the preprocessor always makes on all the input it receives, even in the absence of directives.

- All C comments are replaced with single spaces.
- Backslash-Newline sequences are deleted, no matter where. This feature allows you to break long lines for cosmetic purposes without changing their meaning.
- Predefined macro names are replaced with their expansions (see “Predefined macros” on page 353).

The first two transformations are done *before* nearly all other parsing and before preprocessing directives are recognized. Thus, for example, you can split a line cosmetically with Backslash-Newline anywhere (except when trigraphs are in use; see the following example and its description).

```
/*
 * / # /*
 */ defi\
ne FO\
O 10\
20
```

This input has the equivalent of `#define FOO 1020`. You can split an escape

sequence with Backslash-Newline.

For example, you can split `"foo\bar"` between the `\` and the `b` to get the following sequence.

```
"foo\  
bar"
```

This behavior is unclear: in all other contexts, a Backslash can be inserted in a string constant as an ordinary character by writing a double Backslash, and this creates an exception. But the ANSI C standard requires it. (Strict ANSI C does not allow Newlines in string constants, so they do not consider this a problem.)

There are a few exceptions to all three transformations.

- C comments and predefined macro names are not recognized inside a `#include` directive in which the file name is delimited with `<` and `>`.
- C comments and predefined macro names are never recognized within a character or string constant. (Strictly speaking, this is the rule, not an exception, but it is worth noting here anyway.)
- Backslash-Newline may not safely be used within an ANSI trigraph. Trigraphs are converted before Backslash-Newline is deleted. If you write what looks like a trigraph with a Backslash-Newline inside, the Backslash-Newline is deleted as usual, but it is then too late to recognize the trigraph.

This exception is relevant only if you use the `-trigraphs` option to enable trigraph processing. See “Invoking the C preprocessor” on page 387.

3

Preprocessing directives

Most preprocessor features are active only if you use preprocessing directives to request their use. Preprocessing directives are lines in your program that start with `#`. The `#` is followed by an *identifier*, which is the directive name. For example, `#define` is the directive that defines a macro. Whitespace is also allowed before and after the `#`. The set of valid directive names is fixed. Programs cannot define new preprocessing directives. Some directive names require arguments; these make up the rest of the directive line and must be separated from the directive name by whitespace. For example, `#define` must be followed by a macro name and the intended expansion of the macro. See “Simple macros” on page 348.

A preprocessing directive cannot be more than one line in normal circumstances. It may be split cosmetically with Backslash-Newline, but that has no effect on its meaning. Comments containing Newlines can also divide the directive into multiple lines, but the comments are changed to Spaces before the directive is interpreted. The only way a significant Newline can occur in a preprocessing directive is within a string constant or character constant.

NOTE: C compilers that are applied to the output from the preprocessor do not accept string or character constants containing Newlines.

The `#` and the directive name cannot come from a macro expansion; if `foo` is defined as a macro expanding to `define`, `#foo` does not become a valid preprocessing directive.

4

Header files

A header file is a file containing C declarations and macro definitions (see “Macros” on page 347) to be shared between several source files. You request the use of a header file in your program with the C preprocessing directive `#include`.

The following documentation describes more about header files.

- “Uses of header files” on page 340
- “The `#include` directive” on page 341
- “How `#include` works” on page 343
- “Once-only include files” on page 344
- “Inheritance and header files” on page 345

Uses of header files

Header files serve two kinds of purposes.

- System header files declare the interfaces to parts of the operating system. You include them in your program to supply the definitions and declarations you need to invoke system calls and libraries.
- Your own header files contain declarations for interfaces between the source files of your program. Each time you have a group of related declarations and macro definitions all or most of which are needed in several different source files, it is a good idea to create a header file for them.

Including a header file produces the same results in C compilation as copying the header file into each source file that needs it. But such copying would be time-consuming and error-prone. With a header file, the related declarations appear in only one place. If they need to be changed, they can be changed in one place, and programs that include the header file will automatically use the new version when next recompiled. The header file eliminates the labor of finding and changing all the copies as well as the risk that a failure to find one copy will result in inconsistencies within a program.

The usual convention is to give header files names that end with `.h`. Avoid unusual characters in header file names, as they reduce portability.

The #include directive

Both user and system header files are included using the preprocessing directive `#include`. It has three variants:

- `#include <file>`
This variant is used for system header files. It searches for a file named *file* in a list of directories specified by you, then in a standard list of system directories. You specify directories to search for header files with the command option `-I` (see “Invoking the C preprocessor” on page 387). The option `-nostdinc` inhibits searching the standard system directories; in this case only the directories you specify are searched.

The parsing of this form of `#include` is slightly special because comments are not recognized within the `<...>`.

Thus, in `#include <x/*y>`, the `/*` does not start a comment and the directive specifies inclusion of a system header file named `x/*y`. Of course, a header file with such a name is unlikely to exist on UNIX, where shell wildcard features would make it hard to manipulate.

The argument *file* may not contain a `>` character. It may, however, contain a `<` character.

- `#include "file"`
This variant is used for header files of your own program. It searches for a file named *file*, first in the current directory, then in the same directories used for system header files. The current directory is the directory of the current input file. It is tried first because it is presumed to be the location of the files that the current input file refers to. (If the `-I-` option is used, the special treatment of the current directory is inhibited.)

The argument, *file*, may not contain `"` characters. If backslashes occur within *file*, they are considered ordinary text characters, not escape characters. None of the character escape sequences appropriate to string constants in C are processed. Thus, `#include "x\n\\y"` specifies a filename containing three backslashes. It is not clear why this behavior is ever useful, although the ANSI standard specifies it.

- `#include anything else`
This variant is called a computed `#include`. Any `#include` directive whose argument does not fit the above two forms is a computed include. The text *anything else*, is checked for macro calls, which are expanded (see “Macros” on page 347). When this is done, the result must fit one of the previous two variants—in particular, the expanded text must in the end be surrounded by either quotes or angle braces.

This feature allows you to define a macro which controls the file name to be used

at a later point in the program. One application of this is to allow a site-specific configuration file for your program to specify the names of the system `include` files to be used. This can help in porting the program to various operating systems in which the necessary system header files are found in different places.

How #include works

The `#include` directive works by directing the C preprocessor to scan the specified file as input before continuing with the rest of the current file. The output from the preprocessor contains the output already generated, followed by the output resulting from the included file, followed by the output that comes from the text after the `#include` directive. The example, given a header file `header.h`, follows.

```
char *test ();
```

Then, there is a main program called `program.c` that uses the header file, like the following.

```
int x;
#include "header.h"

main ()
{
    printf (test ());
}
```

The output generated by the C preprocessor for `program.c` as input would be as follows.

```
int x;
char *test ();
main ()
{
    printf (test ());
}
```

Included files are not limited to declarations and macro definitions; those are merely the typical uses. Any fragment of a C program can be included from another file. The include file could even contain the beginning of a statement that is concluded in the containing file, or the end of a statement that was started in the including file.

However, a comment or a string or character constant may not start in the included file and finish in the including file. An unterminated comment, string constant or character constant in an included file is considered to end (with an error message) at the end of the file.

It is possible for a header file to begin or end a syntactic unit such as a function definition, but that would be very confusing, so don't do it.

The line following the `#include` directive is always treated as a separate line by the C preprocessor even if the included file lacks a final newline.

Once-only include files

Very often, one header file includes another header file. It can easily result that a certain header file is included more than once in a file. This may lead to errors, if the header file defines structure types or typedefs, and is certainly wasteful. Therefore, we often try to prevent multiple inclusion of a header file. The standard way to do this is to enclose the entire real contents of the file in a conditional, like the following example's input demonstrates.

```
#ifndef FILE_FOO_SEEN
#define FILE_FOO_SEEN

    the entire file

#endif /* FILE_FOO_SEEN */
```

The macro, `FILE_FOO_SEEN`, indicates that the file has been included *once already*. In a user header file, the macro name should not begin with an underscore, `'_'`. In a system header file, this name should begin with `'__'` (two underscores) to avoid conflicts with user programs. In any kind of header file, the macro name should contain the name of the file and some additional text, to avoid conflicts with other header files.

The GNU C preprocessor is programmed to notice when a header file uses this particular construct and handle it efficiently. If a header file is contained entirely in a `#ifndef` conditional, then it records that fact. If a subsequent `#include` specifies the same file, and the macro in the `#ifndef` is already defined, then the file is entirely skipped, without even reading it.

There is also an explicit directive to tell the preprocessor that it need not include a file more than once. This is called `#pragma once`, and was used in addition to the `#ifndef` conditional around the contents of the header file. `#pragma once` is now obsolete and should not be used at all.

In the Objective C language, there is a variant of `#include` called `#import` which includes a file, but does so at most once. If you use `#import` instead of `#include`, then you don't need the conditionals inside the header file to prevent multiple execution of the contents.

`#import` is obsolete because it is not a well designed feature. Using `#ifndef` meets the requirement in a more straightforward manner.

Inheritance and header files

Inheritance is what happens when one object or file derives some of its contents by virtual copying from another object or file. In the case of C header files, inheritance means that one header file includes another header file and then replaces or adds something.

If the inheriting header file and the base header file have different names, then inheritance is straightforward: simply write `#include "base"` in the inheriting file (where *base* stands for the base file in use).

Sometimes it is necessary to give the inheriting file the same name as the base file. This is less straightforward.

For example, suppose an application program uses the system header file `sys/signal.h`, but the version of `/usr/include/sys/signal.h` on a particular system doesn't do what the application program expects. It might be convenient to define a *local* version, perhaps under the name `/usr/local/include/sys/signal.h`, to override or add to the one supplied by the system. Use the option, `-I.`, for compilation, and writing a file `sys/signal.h` that does what the application program expects. But making this file include the standard `sys/signal.h` is not so easy—writing `#include <sys/signal.h>` in that file doesn't work, because it includes your own version of the file, not the standard system version.

Used in that file itself, this leads to an infinite recursion and a fatal error in compilation.

`#include </usr/include/sys/signal.h>` would find the proper file, but that is not clean, since it makes an assumption about where the system header file is found. This is bad for maintenance, since it means that any change in where the system's header files are kept requires a change somewhere else.

The clean way to solve this problem is to use `#include_next`, which means, "Include the *next* file with this name." This directive works like `#include` except in searching for the specified file: it starts searching the list of header file directories after the directory in which the current file was found.

Suppose you specify `-I /usr/local/include`, and the list of directories to search also includes `/usr/include`; and suppose that both directories contain a file named `sys/signal.h`. Ordinary `#include <sys/signal.h>` finds the file under `/usr/local/include`. If that file contains `#include_next <sys/signal.h>`, it starts searching after that directory, and finds the file in `/usr/include`.

5

Macros

A macro is a sort of abbreviation which you can define once and then use later. There are many complicated features associated with macros in the C preprocessor.

The following documentation describes more about macros.

- “Simple macros” on page 348
- “Macros with arguments” on page 350
- “Predefined macros” on page 353
- “Stringification” on page 358
- “Concatenation” on page 360
- “Undefining macros” on page 362
- “Redefining macros” on page 363
- “Pitfalls and subtleties of macros” on page 364

Simple macros

A *simple macro* is a kind of abbreviation. It is a name which stands for a fragment of code. Some people refer to these as *manifest constants*. Before you can use a macro, you must define it explicitly with the `#define` directive. `#define` is followed by the name of the macro and then the code it should be an abbreviation for. For example, use the following statement.

```
#define BUFFER_SIZE 1020
```

This input defines a macro named `BUFFER_SIZE` as an abbreviation for the text `1020`. If somewhere after this `#define` directive there comes a C statement of the form of the following example, then the C preprocessor will recognize and expand the macro `BUFFER_SIZE`.

```
foo = (char *) xmalloc (BUFFER_SIZE);
```

This then gives the following result.

```
foo = (char *) xmalloc (1020);
```

The use of all uppercase for macro names is a standard convention so that it is possible to tell at a glance which names are macros.

Normally, a macro definition must be a single line, like all C preprocessing directives. (You can split a long macro definition cosmetically with Backslash-Newline.) There is one exception: *newlines* can be included in the macro definition if within a string or character constant. This is because it is not possible for a macro definition to contain an unbalanced quote character; the definition automatically extends to include the matching quote character that ends the string or character constant. Comments within a macro definition may contain newlines, which make no difference since the comments are entirely replaced with spaces regardless of their contents.

Aside from the previous explanation, there is no restriction on what can go in a macro body. Parentheses need not balance. The body need not resemble valid C code. (But if it does not, you may get error messages from the C compiler when you use the macro.)

The C preprocessor scans your program sequentially, so macro definitions take effect at the place you write them. First, input the following.

```
foo = X;
#define X 4
bar = X;
```

This produces the following output.

```
foo = X;

bar = 4;
```

After the preprocessor expands a macro name, the macro's definition body is appended to the front of the remaining input, and the check for macro calls continues. Therefore, the macro body can contain calls to other macros. Use the following, for example, as input.

```
#define BUFSIZE 1020
#define TABLESIZE BUFSIZE
```

Afterwards, name `TABLESIZE` when used in the program would go through two stages of expansion, resulting ultimately in `1020`.

This is not at all the same as defining `TABLESIZE` to be `1020`. The `#define` for `TABLESIZE` uses exactly the body you specify—in this case, `BUFSIZE`—and does not check to see whether it too is the name of a macro. It's only when you use `TABLESIZE` that the result of its expansion is checked for more macro names. See .

Macros with arguments

A simple macro always stands for exactly the same text, each time it is used. Macros can be more flexible when they accept *arguments*. Arguments are fragments of code that you supply each time the macro is used. These fragments are included in the expansion of the macro according to the directions in the macro definition. A macro that accepts arguments is called a function-like macro because the syntax for using it looks like a function call.

To define a macro that uses arguments, you write a `#define` directive with a list of *argument names* in parentheses after the name of the macro. The argument names may be any valid C identifiers, separated by commas and optionally whitespace. The open parenthesis must follow the macro name immediately, with no space in between. For example, here is a macro that computes the minimum of two numeric values, as it is defined in many C programs:

```
#define min(X, Y) ((X) < (Y) ? (X) : (Y))
```

This is not the best way to define a *minimum* macro in GNU C; see “Duplication of side effects” on page 366 for more information. To use a macro that expects arguments, you write the name of the macro followed by a list of *actual arguments* in parentheses, separated by commas. The number of actual arguments you give must match the number of arguments the macro expects.

Examples of using the macro `min` include `min (1, 2)` and `min (x + 28, *p)`.

The expansion text of the macro depends on the arguments you use. Each of the argument names of the macro is replaced, throughout the macro definition, with the corresponding actual argument. Using the same macro, `min` (as defined in the previous input example), `min (1, 2)` expands into the following output.

```
((1) < (2) ? (1) : (2))
```

1 has been substituted for `x` and 2 for `y`. Likewise, `min (x + 28, *p)` expands into the following output.

```
((x + 28) < (*p) ? (x + 28) : (*p))
```

Parentheses in the actual arguments must balance; a comma within parentheses does not end an argument. However, there is no requirement for brackets or braces to balance, and they do not prevent a comma from separating arguments. Use the following input as an example.

```
macro (array[x = y,x +1])
```

This passes two arguments to macro: `array[x = y` and `x + 1]`. If you want to supply `array[x = y, x + 1]` as an argument, you must write it as equivalent C code with `array[(x = y, x + 1)]`.

After the actual arguments are substituted into the macro body, the entire result is appended to the front of the remaining input, and the check for macro calls continues. Therefore, the actual arguments can contain calls to other macros, either with or without arguments, or even to the same macro. The macro body can also contain calls to other macros.

For example, `min (min (a, b), c)` expands into the following output.

```
((((a) < (b) ? (a) : (b))) < (c)
 ? (((a) < (b) ? (a) : (b)))
 : (c))
```

If a macro `foo` takes one argument, and you want to supply an empty argument, you must write at least some whitespace between the parentheses, like this: `foo ()`. Just `foo ()` is providing no arguments, which is an error if `foo` expects an argument.

But `foo0 ()` is the correct way to call a macro defined to take zero arguments, like the following example for input.

```
#define foo0() : : :
```

If you use the macro name followed by something other than an open-parenthesis (after ignoring any spaces, tabs and comments that follow), it is not a call to the macro, and the preprocessor does not change what you have written. Therefore, it is possible for the same name to be a variable or function in your program as well as a macro, and you can choose in each instance whether to refer to the macro (if an actual argument list follows) or the variable or function (if an argument list does not follow).

Such dual use of one name could be confusing and should be avoided except when the two meanings are effectively synonymous: that is, when the name is both a macro and a function and the two have similar effects. You can think of the name simply as a function; use of the name for purposes other than calling it (such as, to take the address) will refer to the function, while calls will expand the macro and generate better but equivalent code. For example, you can use a function named `min` in the same source file that defines the macro. If you write `&min` with no argument list, you refer to the function. If you write `min (x, bb)`, with an argument list, the macro is expanded. If you write `(min) (a, bb)`, where the name `min` is not followed by an open-parenthesis, the macro is not expanded, so you wind up with a call to the function `min`.

You may not define the same name as both a simple macro and a macro with arguments.

In the definition of a macro with arguments, the list of argument names must follow the macro name immediately with no space in between your input. If there is a space after the macro name, the macro is defined as taking no arguments, and all the rest of the line is taken to be the *expansion*. The reason for this is that it is often useful to define a macro that takes no arguments and whose definition begins with an identifier in parentheses. This rule about spaces would use input like either one of the two following examples.

```
#define FOO(x) - 1 / (x)
```

The previous example defines `FOO` to take an argument and expand into minus the reciprocal of that argument.

The following example defines `BAR` to take no argument and always expand into `(x) - 1 / (x)`.

```
#define BAR (x) - 1 / (x)
```

NOTE: The uses of a macro with arguments can have spaces before the left parenthesis; it's the *definition* where it matters whether there is a space.

Predefined macros

Several simple macros are predefined. You can use them without giving definitions for them.

They fall into two classes: standard macros and system-specific macros.

Standard predefined macros

The standard predefined macros are available with the same meanings regardless of the machine or operating system on which you are using GNU C. Their names all start and end with double underscores. Those preceding `__GNUC__` in this table are standardized by ANSI C; the rest are GNU C extensions.

`__FILE__`

This macro expands to the name of the current input file, in the form of a C string constant. The precise name returned is the one that was specified in `#include` or as the input file name argument.

`__LINE__`

This macro expands to the current input line number, in the form of a decimal integer constant. While we call it a predefined macro, it's a pretty strange macro, since its "definition" changes with each new line of source code.

This and `__FILE__` are useful in generating an error message to report an inconsistency detected by the program; the message can state the source line at which the inconsistency was detected. Use the following output as an example.

```
fprintf (stderr, "Internal error: "
          "negative string length "
          "%d at %s, line %d.",
          length, __FILE__, __LINE__);
```

A `#include` directive changes the expansions of `__FILE__` and `__LINE__` to correspond to the included file. At the end of that file, when processing resumes on the input file that contained the `#include` directive, the expansions of `__FILE__` and `__LINE__` revert to the values they had before the `#include` (but `__LINE__` is then incremented by one as processing moves to the line after the `#include`).

The expansions of both `__FILE__` and `__LINE__` are altered if a `#line` directive is used. See "Combining source files" on page 381.

`__DATE__`

This macro expands to a string constant that describes the date on which the preprocessor is being run. The string constant contains eleven characters and looks like "Jan 29 1987" or "Apr 1 1905".

`__TIME__`

This macro expands to a string constant that describes the time at which the preprocessor is being run. The string constant contains eight characters and looks like "23:59:01".

`__STDC__`

This macro expands to the constant 1, to signify that this is ANSI Standard C. (Whether that is actually true depends on what C compiler will operate on the output from the preprocessor.)

`__STDC_VERSION__`

This macro expands to the C Standard's version number, a long integer constant of the form `yyyymmL` where `yyyy` and `mm` are the year and month of the Standard version. This signifies which version of the C Standard the preprocessor conforms to. Like `__STDC__`, whether this version number is accurate for the entire implementation depends on what C compiler will operate on the output from the preprocessor.

`__GNUC__`

This macro is defined if and only if this is GNU C. This macro is defined only when the entire GNU C compiler is in use; if you invoke the preprocessor directly, `__GNUC__` is undefined. The value identifies the major version number of GNU CC (1 for GNU CC version 1, which is now obsolete, and 2 for version 2).

`__GNUC_MINOR__`

The macro contains the minor version number of the compiler. This can be used to work around differences between different releases of the compiler (for example, if GCC 2.6.3 supports a feature, you can test for `__GNUC__ > 2 || (__GNUC__ == 2 && __GNUC_MINOR__ >= 6)`). The last number, 3, denotes the bugfix level of the compiler; no macro contains this value.

`__GNUG__`

The GNU C compiler defines this when the compilation language is C++; use `__GNUG__` to distinguish between GNU C and GNU C++.

`__cplusplus`

The draft ANSI standard for C++ *used to* require predefining this variable. Though it is no longer required, GNU C++ continues to define it, as do other popular C++ compilers. You can use `__cplusplus` to test whether a header is compiled by a C compiler or a C++ compiler.

`__STRICT_ANSI__`

This macro is defined if and only if the `-ansi` switch was specified when GNU C was invoked. Its definition is the null string. This macro exists primarily to direct certain GNU header files not to define certain traditional UNIX constructs which are incompatible with ANSI C.

__BASE_FILE__

This macro expands to the name of the main input file, in the form of a C string constant. This is the source file that was specified as an argument when the C compiler was invoked.

__INCLUDE_LEVEL__

This macro expands to a decimal integer constant that represents the depth of nesting in `include` files. The value of this macro increments on every `#include` directive and decremented at every end of file. For input files specified by command line arguments, the nesting level is zero.

__VERSION__

This macro expands to a string which describes the version number of GNU C. The string is normally a sequence of decimal numbers separated by periods, such as `"2.6.0"`. The only reasonable use of this macro is to incorporate it into a string constant.

__OPTIMIZE__

This macro is defined in optimizing compilations. It causes certain GNU header files to define alternative macro definitions for some system library functions. It is unwise to refer to or test the definition of this macro unless you make very sure that programs will execute with the same effect regardless.

__CHAR_UNSIGNED__

This macro is defined if and only if the data type `char` is unsigned on the target machine. It exists to cause the standard header file `limit.h` to work correctly. It is bad practice to refer to this macro yourself; instead, refer to the standard macros defined in `limit.h`. The preprocessor uses this macro to determine whether or not to sign-extend large character constants written in octal; see “The `#if` directive” on page 373.

__REGISTER_PREFIX__

This macro expands to a string describing the prefix applied to cpu registers in assembler code. It can be used to write assembler code that is usable in multiple environments. For example, in the `m68k-aout` environment it expands to the doublequote string (`" "`). In the `m68k-coff` environment, it expands to the string, `"%"`.

__USER_LABEL_PREFIX__

This macro expands to a string describing the prefix applied to user generated labels in assembler code. It can be used to write assembler code that is usable in multiple environments. For example, in the `m68k-aout` environment it expands to the string `"_"`, but in the `m68k-coff` environment, it expands to the string `"` `"`.

Non-standard predefined macros

The C preprocessor normally has several predefined macros that vary between machines because their purpose is to indicate what type of system and machine is in use.

This documentation, being for all systems and machines, cannot tell you exactly what their names are. Instead, we offer a list of some typical ones.

You can use `cpp -dM` to see the values of predefined macros; for more information, see “Invoking the C preprocessor” on page 387. Some nonstandard predefined macros describe the operating system in use, with more or less specificity, as in the following two examples.

`unix`

`unix` is normally predefined on all UNIX systems.

`BSD`

`BSD` is predefined on recent versions of Berkeley UNIX (perhaps only in version 4.3).

Other nonstandard predefined macros describe the kind of CPU, with more or less specificity, as in the following six examples.

`vax`

`vax` is predefined on Vax computers.

`mc68000`

`mc68000` is predefined on most computers whose CPU is a Motorola 68000, 68010 or 68020.

`m68k`

`m68k` is also predefined on most computers whose CPU is a 68000, 68010 or 68020; however, some makers use `mc68000` and some use `m68k` as the names for the macros. Some predefine both names. What happens in GNU C depends on the system you are using.

`M68020`

`M68020` has been observed to be predefined on some systems that use 68020 CPUs—in addition to `mc68000` and `m68k`, which are less specific.

`_AM29K`

`_AM29000`

Both `_AM29K` and `_AM29000` are predefined for the AMD 29000 CPU family.

`ns32000`

`ns32000` is predefined on computers which use the National Semiconductor 32000 series CPU.

Yet other nonstandard predefined macros describe the manufacturer of the system, as in the following three examples.

`sun`

`sun` is predefined on all models of Sun computers.

`pyr`

`pyr` is predefined on all models of Pyramid computers.

`sequent`

`sequent` is predefined on all models of Sequent computers.

These predefined symbols are not only nonstandard, they are contrary to the ANSI standard because their names do not start with underscores. Therefore, the option `-ansi` inhibits the definition of these symbols.

This tends to make `-ansi` useless, since many programs depend on the customary nonstandard predefined symbols. Even system header files check them and will generate incorrect declarations if they do not find the names that are expected. You might think that the header files supplied for the Uglix computer would not need to test what machine they are running on, because they can simply assume it is the Uglix; but often they do, and they do so using the customary names. As a result, very few C programs will compile with `-ansi`. We intend to avoid such problems on the GNU system.

What, then, should you do in an ANSI C program to test the type of machine it will run on? GNU C offers a parallel series of symbols for this purpose, whose names are made from the customary ones by adding `'__'` at the beginning and end. Thus, the symbol `__vax__` would be available on a Vax, and so on. The set of nonstandard predefined names in the GNU C preprocessor is controlled (when `cpp` is itself compiled) by the macro `CPP_PREDEFINES`, which should be a string containing `-D` options, separated by spaces. For example, on the Sun 3, we use the following definition.

```
#define CPP_PREDEFINES "-Dmc68000 -Dsun -Dunix -Dm68k"
```

This macro is usually specified in `tm.h`.

Stringification

Stringification means turning a code fragment into a string constant whose contents are the text for the code fragment. For example, stringifying `foo (z)` results in `"foo (z)"`.

In the C preprocessor, stringification is an option available when macro arguments are substituted into the macro definition. In the body of the definition, when an argument name appears, the `#` character before the name specifies stringification of the corresponding actual argument when it is substituted at that point in the definition. The same argument may be substituted in other places in the definition without stringification if the argument name appears in those places with no `#`.

What follows is an example of a macro definition that uses stringification.

```
#define WARN_IF(EXP) \
do { if (EXP) \
    fprintf (stderr, "Warning: " #EXP "\n"); } \
while (0)
```

The actual argument for `EXP` is substituted once as given, into the `if` statement, and once as stringified, into the argument to `fprintf`. The `do` and `while (0)` are a work-around to make it possible to write `WARN_IF(arg);`, which the resemblance of `WARN_IF` to a function would make C programmers want to do; see “Swallowing the semicolon” on page 365.

The stringification feature is limited to transforming one macro argument into one string constant; there is no way to combine the argument with other text and then stringify it all together. The previous example shows how an equivalent result can be obtained in ANSI Standard C using the feature that adjacent string constants are concatenated as one string constant. The preprocessor stringifies the actual value of `EXP` into a separate string constant, resulting in text like the following output.

```
do {if (x==0)\
    fprintf (stderr, "Warning: " "x == 0" "\n"); } \
while (0)
```

The C compiler then sees three consecutive string constants and concatenates them into one, producing, effectively, the following output.

```
do {if (x==0)\
    fprintf (stderr, "Warning: x == 0\n"); } \
while (0)
```

Stringification in C involves more than putting doublequote characters around the fragment; it is necessary to put backslashes in front of all doublequote characters, and all backslashes in string and character constants, in order to get a valid C string constant with the proper contents. Thus, stringifying `p="foo\n";` results in `"p=\"foo\\n\";"`. However, backslashes that are not inside of string or character constants are not duplicated: `\n` by itself stringifies to `"\n"`.

Whitespace (including comments) in the text being stringified is handled according to precise rules. All leading and trailing whitespace is ignored. Any sequence of whitespace in the middle of the text is converted to a single space in the stringified result.

Concatenation

In the context of macro expansion, *concatenation* refers to joining two lexical units or two strings into one longer one. Specifically, an actual argument to the macro can be concatenated with another actual argument or with fixed text to produce a longer name. The longer name might be the name of a function, variable or type, or a C keyword; it might even be the name of another macro, in which case it will expand.

When you define a macro, you request concatenation with the special operator `##` in the macro body. When the macro is called, after actual arguments are substituted, all `##` operators are deleted, and any whitespace next to them (including whitespace that was part of an actual argument). The result is to concatenate the syntactic tokens on either side of the `##`. Consider a C program interpreting named commands. There likely needs to be a table of commands, or an array of structures, declared as follows.

```
struct command
{
    char *name;
    void (*function) ();
};

struct command commands[] =
{
    { "quit", quit_command},
    { "help", help_command},
    ...
};
```

It would be cleaner not to have to give each command name twice, once in the string constant and once in the function name. A macro which takes the name of a command as an argument can make this unnecessary. The string constant can be created with stringification, and the function name by concatenating the argument with `_command`. What follows is an example of how it is done.

```
#define COMMAND(NAME)    { #NAME, NAME ## _command }

struct command commands[] =
{
    COMMAND (quit),
    COMMAND (help),
    ...
};
```

The usual case of concatenation is concatenating two names (or a name and a number) into a longer name. But this isn't the only valid case. It is also possible to concatenate two numbers (or a number and a name, such as `1.5` and `e3`) into a number. Also, multi-character operators such as `+=` can be formed by concatenation. In some cases it is even possible to piece together a string constant. However, two pieces of text that

don't together form a valid lexical unit cannot be concatenated. For example, concatenation with `x` on one side and `+` on the other is not meaningful because those two characters can't fit together in any lexical unit of C. The ANSI standard says that such attempts at concatenation are undefined, but in the GNU C preprocessor it is well defined: it puts `x` and `+` side by side with no particular special results.

Keep in mind that the C preprocessor converts comments to whitespace before macros are even considered. Therefore, you cannot create a comment by concatenating `/` and `*`: the `/*` sequence that starts a comment is not a lexical unit, but rather the beginning of a “long” space character. Also, you can freely use comments next to `##` in a macro definition, or in actual arguments that will concatenate, because the comments will be convert to spaces at first sight, and concatenation will later discard the spaces.

Undefining macros

To *undefine* a macro means to cancel its definition. This is done with the `#undef` directive. `#undef` is followed by the macro name to be undefined.

Like definition, *undefinition* occurs at a specific point in the source file, and it applies starting from that point. The name ceases to be a macro name, and from that point on it is treated by the preprocessor as if it had never been a macro name. For clarification, use the following example.

```
#define FOO 4 x = FOO; #undef FOO x = FOO;
```

This input expands into the following output.

```
x =4;  
x = FOO;
```

In the previous example, `FOO` had better be a variable or function as well as (temporarily) a macro, in order for the result of the expansion to be valid C code. The same form of `#undef` directive will cancel definitions with arguments or definitions that don't expect arguments. The `#undef` directive has no effect when used on a name not currently defined as a macro.

Redefining macros

Redefining a macro means defining (with `#define`) a name that is already defined as a macro. A redefinition is trivial if the new definition is transparently identical to the old one. You probably wouldn't deliberately write a trivial redefinition, but they can happen automatically when a header file is included more than once (see "Header files" on page 339), so they are accepted silently and without effect.

Nontrivial redefinition is considered likely to be an error, so it provokes a warning message from the preprocessor. However, sometimes it is useful to change the definition of a macro in mid-compilation. You can inhibit the warning by undefining the macro with `#undef` before the second definition. In order for a redefinition to be trivial, the new definition must exactly match the one already in effect, with two possible exceptions, as in the following.

- Whitespace may be added or deleted at the beginning or the end.
- Whitespace may be changed in the middle (but not inside strings). However, it may not be eliminated entirely, and it may not be added where there was no whitespace at all.

Recall that a comment counts as whitespace.

Pitfalls and subtleties of macros

In this section we describe some special rules that apply to macros and macro expansion, and point out certain cases in which the rules have counterintuitive consequences that you must watch out for.

Improperly nested constructs

Recall that when a macro is called with arguments, the arguments are substituted into the macro body and the result is checked, together with the rest of the input file, for more macro calls. It is possible to piece together a macro call coming partially from the macro body and partially from the actual arguments. Use the following input as an example.

```
#define double(x) (2*(x))
#define call_with_1(x) x(1)
```

This input would expand `call_with_1 (double)` into `(2*(1))`.

Macro definitions do not have to have balanced parentheses. By writing an unbalanced open parenthesis in a macro body, it is possible to create a macro call that begins inside the macro body but ends outside of it. Use the following input as an example.

```
#define strange(file) fprintf (file, "%s %d",
. . .
strange(stderr) p, 35)
```

The previous bizarre example expands to the following output.

```
fprintf (stderr, "%s %d", p, 35)
```

Unintended grouping of arithmetic

You may have noticed that in most of the macro definition examples, each occurrence of a macro argument name had parentheses around it. In addition, another pair of parentheses usually surround the entire macro definition. The following is why it is best to write macros that way.

Suppose you define a macro as follows.

```
#define ceil_div(x, y) (x + y - 1) / y
```

This produces macro output whose purpose is to divide, rounding up. (One use for this operation is to compute how many `int` objects are needed to hold a certain number of `char` objects.) Then suppose it is used as follows.

```
a = ceil_div (b & c, sizeof (int));
```

This expands into output like the following.

```
a = (b & c + sizeof (int) - 1) / sizeof (int);
```

This output does not do what is intended. The operator-precedence rules of C make it equivalent to an operation like the following.

```
a = (b & (c + sizeof (int) - 1)) / sizeof (int);
```

But what we want is the following result.

```
a = ((b & c) + sizeof (int) - 1) / sizeof (int);
```

This would mean defining the macro as the following.

```
#define ceil_div(x, y) ((x) + (y) - 1) / (y)
```

This provides the desired result. However, unintended grouping can result in another way. Consider `sizeof ceil_div(1, 2)`. That has the appearance of a C expression that would compute the size of the type of `ceil_div (1, 2)`, but in fact it means something very different. Use the following output as an example of how it expands.

```
sizeof ((1) + (2) - 1) / (2)
```

This would take the size of an integer and divide it by two. The precedence rules have put the division outside the `sizeof` when it was intended to be inside.

Parentheses around the entire macro definition can prevent such problems. What follows, then, is the recommended way to define `ceil_div`.

```
#define ceil_div(x, y) (((x) + (y) - 1) / (y))
```

Swallowing the semicolon

Often it is desirable to define a macro that expands into a compound statement. Consider, for example, the following macro, that advances a pointer (the argument `p` says where to find it) across whitespace characters:

```
#define SKIP_SPACES (p, limit) \
{ register char *lim = (limit); \
  while (p != lim) { \
    if (*p++ != ' ') { \
      p--; break; } } }
```

Note: In the previous example, Backslash-Newline is used to split the macro definition, which must be a single line, so that it resembles the way such C code would be laid out if not part of a macro definition.

A call to this macro might be `SKIP_SPACES (p, lim)`. Strictly speaking, the call expands to a compound statement, which is a complete statement with no need for a semicolon to end it. But it looks like a function call. So it minimizes confusion if to use it like a function call, writing a semicolon afterward, as in the following example.

```
SKIP_SPACES (p, lim);
```

But this can cause trouble before `else` statements, because the semicolon is actually a null statement. Suppose you have the following input.

```
if (*p != 0)
  SKIP_SPACES (p, lim);
else. . .
```

The presence of two statements—the compound statement and a null statement—in between the `if` condition and the `else` makes invalid C code. The definition of the macro `SKIP_SPACES` can be altered to solve this problem, using a `do...while` statement. Use the following input as an example.

```
#define SKIP_SPACES (p, limit) \
do { register char *lim = (limit); \
    while (p != lim) { \
        if (*p++ != ' ') { \
            p--; break; } } \
    while (0)
```

Now `SKIP_SPACES (p, lim);` expands into one output statement as the following example shows.

```
do { : : :
} while (0);
```

Duplication of side effects

Many C programs define a macro `min` (standing for *minimum*) like the following.

```
#define min(X, Y) ((X) < (Y) ? (X) : (Y))
```

When you use this macro with an argument containing a *side effect* like the statement, `next = min (x + y, foo (z));`, it expands into the following output.

```
next = ((x + y) < (foo (z)) ? (x + y) : (foo (z)));
```

`x + y` has been substituted for `x` and `foo (z)` for `y`. The function `foo` is used only once in the statement as it appears in the program, but the expression `foo (z)` has been substituted twice into the macro expansion. As a result, `foo` might be called two times when the statement is executed. If it has side effects or if it takes a long time to compute, the results might not be what you intended. We say that `min` is an unsafe macro. The best solution to this problem is to define `min` in a way that computes the value of `foo (z)` only once. The C language offers no standard way to do this, but it can be done with GNU C extensions as the following example shows.

```
#define min(X, Y) \
({ typeof (X) __x = (X), __y = (Y); (__x < __y) ? __x : __y;
})
```

If you do not wish to use GNU C extensions, the only solution is to be careful when using the macro, `min`. For instance, you can calculate the value of `foo (z)`, save it in a variable, and use that variable in `min`, as in the following example.

```
#define min(X, Y) ((X) < (Y) ? (X) : (Y))
. . .
{ int tem = foo (z); next = min (x + y, tem); }
```

This operation assumes that `foo` returns type `int`.

Self-referential macros

A self-referential macro is one whose name appears in its definition. A special feature of ANSI Standard C is that the self-reference is not considered a macro call. It is passed into the preprocessor output unchanged.

Let's consider an example.

```
#define foo (4 + foo)
```

`foo`, then, is also a variable in your program.

Following the ordinary rules, each reference to `foo` will expand into `(4+foo)`; then this will be re-scanned and will expand into `(4+(4+foo))`; and so on until it causes a fatal error (memory full) in the preprocessor.

However, the special rule about self-reference cuts this process short after one step at `(4 + foo)`. So, this macro definition has the possibly useful effect of causing the program to add 4 to the value of `foo` wherever `foo` is referred to. In most cases, it is a bad idea to take advantage of this feature. A person reading the program who sees that `foo` is a variable will not expect that it is a macro as well. The reader will come across the identifier `foo` in the program and think its value should be that of the variable `foo`, whereas in fact the value is four greater. The special rule for self-reference applies also to *indirect* self-reference. This is the case where a macro `x` expands to use a macro `y`, and the expansion of `y` refers to the macro `x`. The resulting reference to `x` comes indirectly from the expansion of `x`, so it is a self-reference and is not further expanded. Suppose you used the following input.

```
#define x (4 + y)
#define y (2 * x)
```

`x`, then, would expand into `(4+(2*x))`.

But suppose `y` is used elsewhere, not from the definition of `x`. Then the use of `x` in the expansion of `y` is not a self-reference because `x` is not *in progress*. So it does expand. However, the expansion of `x` contains a reference to `y`, and that is an indirect self-reference now because `y` is in progress. The result is that `y` expands to `(2*(4+y))`. It is not clear that this behavior would ever be useful, but it is specified by the ANSI C standard, so you may need to understand it.

Separate expansion of macro arguments

We have explained that the expansion of a macro, including the substituted actual arguments, is scanned over again for macro calls to be expanded.

What really happens is more subtle: first each actual argument text is scanned separately for macro calls. Then the results of this are substituted into the macro body to produce the macro expansion, and the macro expansion is scanned again for macros to expand. The result is that the actual arguments are scanned twice to expand macro calls in them. Most of the time, this has no effect. If the actual argument contained any macro calls, they are expanded during the first scan. The result therefore contains no

macro calls, so the second scan does not change it.

If the actual argument were substituted as given, with no pre-scan, the single remaining scan would find the same macro calls and produce the same results. You might expect the double scan to change the results when a self-referential macro is used in an actual argument of another macro (see “Self-referential macros” on page 367). The self-referential macro would be expanded once in the first scan, and a second time in the second scan. But this is not what happens. The self-references that do not expand in the first scan are marked so that they will not expand in the second scan either.

The pre-scan is not done when an argument is stringified or concatenated. Use the following input as an example.

```
#define str(s) #s
#define foo 4
str (foo)
```

This, then, expands to `"foo"`. Once more, prescan has been prevented from having any noticeable effect. More precisely, stringification and concatenation use the argument as written, in un-prescanned form. The same actual argument would be used in pre-scanned form if it is substituted elsewhere without stringification or concatenation.

```
#define str(s) #s lose(s)
#define foo 4
str (foo)
```

This, then, expands to `"foo" lose(4)`.

You might now ask, “Why mention the pre-scan, if it makes no difference? And why not skip it and make the preprocessor faster?” The answer is that the pre-scan does make a difference in three special cases:

- Nested calls to a macro.
- Macros that call other macros that stringify or concatenate.
- Macros whose expansions contain unshielded commas.

We say that nested calls to a macro occur when a macro’s actual argument contains a call to that very macro.

For example, if `f` is a macro expecting one argument, `f (f (1))` is a nested pair of calls to `f`. The desired expansion is made by expanding `f (1)` and substituting that into the definition of `f`. The pre-scan causes the expected result to happen.

Without the prescan, `f (1)` itself would be substituted as an actual argument, and the inner use of `f` would appear during the main scan as an indirect self-reference and would not be expanded.

Here, the pre-scan cancels an undesirable side effect (in the medical, not computational, sense of the term) of the special rule for self-referential macros. But pre-scan causes trouble in certain other cases of nested macro calls, as in the following

example.

```
#define foo a,b
#define bar(x) lose(x)
#define lose(x) (1 + (x))

bar(foo)
```

We would like `bar(foo)` to turn into `(1 + (foo))`, which would then turn into `(1 + (a,b))`. But instead, `bar(foo)` expands into `lose(a,b)`, and you get an error because `lose` requires a single argument. In this case, the problem is easily solved by the same parentheses that ought to be used to prevent mis-nesting of arithmetic operations.

```
#define foo (a,b)
#define bar(x) lose((x))
```

The problem is more serious when the operands of the macro are not expressions; for example, when they are statements. Then parentheses are unacceptable because they would make for invalid C code.

```
#define foo { int a, b; . . .
}
```

In GNU C you can shield the commas using the `({ . . . })` construct which turns a compound statement into an expression like the following.

```
#define foo ({ int a, b; . . .
})
```

Or you can rewrite the macro definition to avoid such commas, using the following input.

```
#define foo { int a; int b; . . .
}
```

There is also one case where pre-scan is useful. It is possible to use pre-scan to expand an argument and then stringify it—if you use two levels of macros. Let's add a new macro, `xstr`, to the previous definition.

```
#define xstr(s) str(s)
#define str(s) #s
#define foo 4
xstr (foo)
```

This expands into `"4"`, not `"foo"`. The reason for the difference is that the argument of `xstr` is expanded at pre-scan (because `xstr` does not specify stringification or concatenation of the argument). The result of pre-scan then forms the actual argument for `str`. `str` uses its argument without pre-scan because it performs stringification; but it cannot prevent or undo the pre-scanning already done by `xstr`.

Cascaded use of macros

A cascade of macros is when one macro's body contains a reference to another macro. This is very common practice, as in the following example.

```
#define BUFSIZE 1020
#define TABLESIZE BUFSIZE
```

This is not at all the same as defining `TABLESIZE` to be 1020. The `#define` for `TABLESIZE` uses exactly the body you specify—in this case, `BUFSIZE`—and does not check to see whether it too is the name of a macro.

It's only when you use `TABLESIZE` that the result of its expansion is checked for more macro names. This makes a difference if you change the definition of `BUFSIZE` at some point in the source file. `TABLESIZE`, defined as in the following example, will always expand using the definition of `BUFSIZE` that is currently in effect.

```
#define BUFSIZE 1020
#define TABLESIZE BUFSIZE
#undef BUFSIZE #define BUFSIZE 37
```

Now `TABLESIZE` expands (in two stages) to 37. (The `#undef` is to prevent any warning about the nontrivial redefinition of `BUFSIZE`.)

Newlines in macro arguments

Traditional macro processing carries forward all newlines in macro arguments into the expansion of the macro. This means that, if some of the arguments are substituted more than once, or not at all, or out of order, newlines can be duplicated, lost, or moved around within the expansion. If the expansion consists of multiple statements, then the effect is to distort the line numbers of some of these statements. The result can be incorrect line numbers, in error messages or displayed in a debugger.

The GNU C preprocessor operating in ANSI C mode adjusts appropriately for multiple use of an argument—the first use expands all the newlines, and subsequent uses of the same argument produce no new-lines. But even in this mode, it can produce incorrect line numbering if arguments are used out of order, or not used at all. What follows is an example illustrating this problem.

```
#define ignore_second_arg(a,b,c) a; c

ignore_second_arg (foo (),
                  ignored (),
                  syntax error);
```

The syntax error triggered by the tokens “`syntax error`” results in an error message citing *the line containing* “`ignored ()`,” even though the statement of “`syntax error` ;” is *the line containing the error*.

6

Conditionals

In a macro processor, a *conditional* is a directive that allows a part of the program to be ignored during compilation, on some conditions. In the C preprocessor, a conditional can test either an arithmetic expression or whether a name is defined as a macro. A conditional in the C preprocessor resembles in some ways an `if` statement in C, but it is important to understand the difference between them.

The condition in an `if` statement is tested during the execution of your program. Its purpose is to allow your program to behave differently from run to run, depending on the data it is operating on. The condition in a preprocessing conditional directive is tested when your program is compiled. Its purpose is to allow different code to be included in the program depending on the situation at the time of compilation.

See the following documentation for more details.

- “Why conditionals are used” on page 372
- “The `#if` directive” on page 373
- “The `#else` directive” on page 374
- “The `#elif` directive” on page 374
- “Keeping deleted code for future reference” on page 375
- “Conditionals and macros” on page 376
- “Assertions” on page 378
- “The `#error` and `#warning` directives” on page 380

Why conditionals are used

Generally there are three kinds of reason to use a conditional.

- A program may need to use different code depending on the machine or operating system it is to run on. In some cases the code for one operating system may be erroneous on another operating system; for example, it might refer to library routines that do not exist on the other system. When this happens, it is not enough to avoid executing the invalid code: merely having it in the program makes it impossible to link the program and run it. With a preprocessing conditional, the offending code can be effectively excised from the program when it is not valid.
- You may want to be able to compile the same source file into two different programs. Sometimes the difference between the programs is that one makes frequent time-consuming consistency checks on its intermediate data, or prints the values of those data for debugging, while the other does not.
- A conditional whose condition is always false is a good way to exclude code from the program but keep it as a sort of comment for future reference.

Most simple programs that are intended to run on only one machine will not need to use preprocessing conditionals.

Syntax of conditionals

A conditional in the C preprocessor begins with one of three conditional directives: `#if`, `#ifdef` or `#ifndef`. See “Conditionals and macros” on page 376 for information on `#ifdef` and `#ifndef`; see the explanations for `#if` with “The `#if` directive” in the following discussion.

The `#if` directive

The `#if` directive in its simplest form consists of the following statement.

```
#if expression
    controlled text
#endif /* expression */
```

The comment following `#endif` is not required, but it is a good practice because it helps people match the `#endif` to the corresponding `#if`. Such comments should always be used, except in short conditionals that are not nested. In fact, you can put anything at all after the `#endif` and it will be ignored by the GNU C preprocessor, but only comments are acceptable in ANSI Standard C. *expression* is a C expression of integer type, subject to stringent restrictions. It may contain:

- *Integer constants*, which are all regarded as `long` or unsigned `long`.
- *Character constants*, which are interpreted according to the character set and conventions of the machine and operating system on which the preprocessor is running. The GNU C preprocessor uses the C data type, `char`, for these character constants; therefore, whether some character codes are negative is determined by the C compiler used to compile the preprocessor. If it treats `char` as signed, then character codes large enough to set the sign bit will be considered negative; otherwise, no character code is considered negative.
- *Arithmetic operators* for addition, subtraction, multiplication, division, bitwise operations, shifts, comparisons, and logical operations (`&&` and `||`).
- *Identifiers that are not macros*, which are all treated as zero(!).
- *Macro calls*; all macro calls in the expression are expanded before actual computation of the expression’s value begins.

NOTE: `sizeof` operators and enum-type values are not allowed. enum-type values, like all other identifiers that are not taken as macro calls and expanded, are treated as zero.

The *controlled text* inside of a conditional can include preprocessing directives. Then the directives inside the conditional are obeyed only if that branch of the conditional succeeds. The text can also contain other conditional groups. However, the `#if` and `#endif` directives must balance.

The #else directive

The `#else` directive can be added to a conditional to provide alternative text to be used if the condition is false. The following code example is what it resembles.

```
#if expression
text-if-true
#else /* Not expression */
text-if-false
#endif /* Not expression */
```

If *expression* is nonzero, and thus the *text-if-true* is active, then `#else` acts like a failing conditional and the *text-if-false* is ignored. Contrarily, if the `#if` conditional fails, the *text-if-false* is considered included.

The #elif directive

One common case of nested conditionals is used to check for more than two possible alternatives. For example, you might have the following statement.

```
#if X == 1
. . .
#else /* X != 1 */ #if X == 2
. . .
#else /* X != 2 */
. . .
#endif /* X != 2 */
#endif /* X != 1 */
```

Another conditional directive, `#elif`, allows this to be abbreviated as in the following example.

```
#if X == 1
. . .
#elif X == 2
. . .
#else /* X != 2 and X != 1 */
. . .
#endif /* X != 2 and X != 1 */
```

`#elif` stands for *else if*. Like `#else`, it goes in the middle of “`#if`” and “`#endif`” pairs, subdividing the pair; `#elif` does not require a matching `#endif`. Like `#if`, the `#elif` directive includes an expression to be tested.

The text following the `#elif` is processed only if the original `#if` condition failed and the `#elif` condition succeeds. More than one `#elif` can go in the same “`#if`”-“`#endif`” group. Then the text after each `#elif` is processed only if the `#elif` condition succeeds after the original `#if` and any previous `#elif` directives within it have failed. `#else` is equivalent to `#elif 1`, and `#else` is allowed after any number of `#elif` directives, but `#elif` may not follow `#else`.

Keeping deleted code for future reference

If you replace or delete a part of the program but want to keep the old code around as a comment for future reference, the easy way to do this is to put `#if 0` before it and `#endif` after it. This is better than using comment delimiters `/*` and `*/` since those won't work if the code already contains comments (C comments do not nest).

This works even if the code being turned off contains conditionals, but they must be entire conditionals (balanced `#if` and `#endif`).

Conversely, do not use `#if 0` for comments which are not C code. Use the comment delimiters `/*` and `*/` instead. The interior of `#if 0` must consist of complete tokens; in particular, single quote characters must balance. But comments often contain unbalanced single quote characters (known in English as apostrophes). These confuse `#if 0`. They do not confuse `/*`.

Conditionals and macros

Conditionals are useful in connection with macros or assertions, because those are the only ways that an expression's value can vary from one compilation to another. A `#if` directive whose expression uses no macros or assertions is equivalent to `#if 1` or `#if 0`; you might as well determine which one, by computing the value of the expression yourself, and then simplify the program. For example, what follows is a conditional statement that tests the expression `BUFSIZE==1020`, where `BUFSIZE` must be a macro.

```
#if BUFSIZE == 1020
    printf ("Large buffers!\n");
#endif /* BUFSIZE is large */
```

NOTE: Programmers often wish they could test the size of a variable or data type in `#if`, but this does not work. The preprocessor does not understand `sizeof`, or `typedef` names, or even the type keywords such as `int`.)

The special operator `defined` is used in `#if` expressions to test whether a certain name is defined as a macro. Either `defined name` or `defined(name)` is an expression whose value is 1 if `name` is defined as macro at the current point in the program, and 0 otherwise. For the `defined` operator it makes no difference what the definition of the macro is; all that matters is whether there is a definition, as in the following example.

```
#if defined (vax) || defined (ns16000)
```

This statement would succeed if either of the names `vax` and `ns16000` is defined as a macro. You can test the same condition using assertions (see “Assertions” on page 378), like the following example shows.

```
#if #cpu (vax) || #cpu (ns16000)
```

If a macro is defined and later undefined with `#undef`, subsequent use of the `defined` operator returns 0, because the name is no longer defined. If the macro is defined again with another `#define`, `defined` will recommence returning 1.

Conditionals that test whether just one name is defined are very common, so there are two special short conditional directives for this case.

- `#ifdef name` is equivalent to `#if defined (name)`.
- `#ifndef name` is equivalent to `#if ! defined (name)`.

Macro definitions can vary between compilations for several reasons.

- Some macros are predefined on each kind of machine. For example, on a Vax, the name `vax` is a predefined macro. On other machines, it would not be defined.
- Many more macros are defined by system header files. Different systems and machines define different macros, or give them different values. It is useful to test these macros with conditionals to avoid using a system feature on a machine where it is not implemented.

- Macros are a common way of allowing users to customize a program for different machines or applications. For example, the macro `BUFSIZE` might be defined in a configuration file for your program that is included as a header file in each source file. You would use `BUFSIZE` in a preprocessing conditional in order to generate different code depending on the chosen configuration.
- Macros can be defined or undefined with `-D` and `-U` command options when you compile the program. You can arrange to compile the same source file into two different programs by choosing a macro name to specify which program you want, writing conditionals to test whether or how this macro is defined, and then controlling the state of the macro with compiler command options. See “Invoking the C preprocessor” on page 387.

Assertions

Assertions are a more systematic alternative to macros in writing conditionals to test what sort of computer or system the compiled program will run on. Assertions are usually predefined, but you can define them with preprocessing directives or command-line options. The macros traditionally used to describe the type of target are not classified in any way according to which question they answer; they may indicate a hardware architecture, a particular hardware model, an operating system, a particular version of an operating system, or specific configuration options. These are jumbled together in a single namespace. In contrast, each assertion consists of a named question and an answer. The question is usually called the predicate. An assertion looks like the following statement.

```
#predicate (answer)
```

You must use a properly formed identifier for *predicate*. The value of *answer* can be any sequence of words; all characters are significant except for leading and trailing whitespace, and differences in internal whitespace sequences are ignored. Thus, `x +y` is different from `x+y` but equivalent to `x +y.` `)` is not allowed in an answer. What follows is a conditional to test whether the *answer* is asserted for *predicate*.

```
#if #predicate (answer)
```

There may be more than one answer asserted for a given predicate. If you omit the answer, you can test whether *any* answer is asserted for *predicate*:

```
#if #predicate
```

Most of the time, the assertions you test will be predefined assertions. GNU C provides three predefined predicates: `system`, `cpu`, and `machine`. `system` is for assertions about the type of software, `cpu` describes the type of computer architecture, and `machine` gives more information about the computer. For example, on a GNU system, the following assertions would be true.

```
#system (gnu)
#system (mach)
#system (mach 3)
#system (mach 3. subversion)
#system (hurd)
#system (hurd version)
```

Perhaps there are others. The alternatives with more or less version information let you ask more or less detailed questions about the type of system software. On a UNIX system, you would find `#system (unix)` and perhaps one of: `#system (aix)`, `#system (bsd)`, `#system (hpux)`, `#system (lynx)`, `#system (mach)`, `#system (posix)`, `#system (svr3)`, `#system (svr4)`, or `#system (xpg4)` with possible version numbers following.

Other values for `system` are `#system (mvs)` and `#system (vms)`.

NOTE: Many UNIX C compilers provide only one answer for the `system` assertion:

`#system (unix)`, if they support assertions at all. This is less than useful.

An assertion with a multi-word answer is completely different from several assertions with individual single-word answers. For example, the presence of `system (mach 3.0)` does not mean that `system (3.0)` is true. It also does not directly imply `system (mach)`, but in GNU C, that last will normally be asserted as well. The current list of possible assertion values for `cpu` is: `#cpu (a29k)`, `#cpu (alpha)`, `#cpu (arm)`, `#cpu (clipper)`, `#cpu (convex)`, `#cpu (elxsi)`, `#cpu (tron)`, `#cpu (h8300)`, `#cpu (i370)`, `#cpu (i386)`, `#cpu (i860)`, `#cpu (i960)`, `#cpu (m68k)`, `#cpu (m88k)`, `#cpu (mips)`, `#cpu (ns32k)`, `#cpu (hppa)`, `#cpu (pyr)`, `#cpu (ibm032)`, `#cpu (rs6000)`, `#cpu (sh)`, `#cpu (sparc)`, `#cpu (spur)`, `#cpu (tahoe)`, `#cpu (vax)`, `#cpu (we32000)`.

You can create assertions within a C program using `#assert`, with the following input.

```
#assert predicate (answer)
```

NOTE: `#` does not appear before *predicate*.

Each time you do this, you assert a new true answer for *predicate*. Asserting one answer does not invalidate previously asserted answers; they all remain true. The only way to remove an assertion is with `#unassert`. `#unassert` has the same syntax as `#assert`. You can also remove all assertions about *predicate* using the following example's statement.

```
#unassert predicate
```

You can also add or cancel assertions using command options when you run `gcc` or `cpp`. See “Invoking the C preprocessor” on page 387.

The `#error` and `#warning` directives

The `#error` directive causes the preprocessor to report a fatal error. The rest of the line that follows `#error` is used as the error message.

You would use `#error` inside of a conditional that detects a combination of parameters which you know the program does not properly support. For example, if you know that the program will not run properly on a Vax, you might use the following input.

```
#ifndef __vax__
#error Won't work on Vaxen. See comments at get_last_object.
#endif
```

See “Non-standard predefined macros” on page 356 for a description of why this works.

If you have several configuration parameters that must be set up by the installation in a consistent way, you can use conditionals to detect an inconsistency and report it with `#error`. For clarification, see the following example.

```
#if HASH_TABLE_SIZE % 2 == 0 || HASH_TABLE_SIZE % 3 == 0 \
|| HASH_TABLE_SIZE % 5 ==
0
#error HASH_TABLE_SIZE should not be divisible by a small
prime
#endif
```

The directive, `#warning`, is like the directive, `#error`, but causes the preprocessor to issue a warning and continue preprocessing. The rest of the line that follows `#warning` is used as the warning message.

You might use `#warning` in obsolete header files, with a message directing the user to the header file which should instead be used.

7

Combining source files

One of the jobs of the C preprocessor is to inform the C compiler of where each line of C code has as its origin, which source file and which line number.

C code can come from multiple source files if you use `#include`; both `#include` and the use of conditionals and macros can cause the line number of a line in the preprocessor output to be different from the line's number in the original source file. You will appreciate the value of making both the C compiler (in error messages) and symbolic debuggers such as GDB use the line numbers in your source file.

The C preprocessor builds on this feature by offering a directive by which you can control the feature explicitly. This is useful when a file for input to the C preprocessor is the output from another program such as the `bison` parser generator, which operates on another file that is the true source file. Parts of the output from `bison` are generated from scratch, other parts come from a standard parser file. The rest are copied nearly verbatim from the source file, but their line numbers in the `bison` output are not the same as their original line numbers. Naturally you would like compiler error messages and symbolic debuggers to know the original source file and line number of each line in the `bison` input.

`bison` arranges this by writing `#line` directives into the output file. `#line` is a directive that specifies the original line number and source file name for subsequent input in the current preprocessor input file. `#line` has three variants:

`#line linenum`

linenum is a decimal integer constant, for specifying that the line number of the

following line of input, in its original source file, was that same designation, *linenum*.

`#line linenum filename`

Here *linenum* is a decimal integer constant and *filename* is a string constant.

This specifies that the following line of input came originally from source file *filename* and its line number there was *linenum*. Keep in mind that *filename* is not just a file *name*.

`#line anything else`

anything else is checked for macro calls, which are expanded. The result should be a decimal integer constant followed optionally by a string constant.

`#line` directives alter the results of the `__FILE__` and `__LINE__` predefined macros from that point on. See “Standard predefined macros” on page 353.

The output of the preprocessor (which is the input for the rest of the compiler) contains directives that look much like `#line` directives. They start with just `#` instead of `#line`, but this is followed by a line number and file name as in `#line`. See “C preprocessor output” on page 385.

8

Other preprocessing directives

The following documentation describes three additional preprocessing directives.

- The *null directive* consists of a `#` followed by a Newline, with only whitespace (including comments) in between. A null directive is understood as a preprocessing directive but has no effect on the preprocessor output. The primary significance of the existence of the null directive is that an input line consisting of just a `#` will produce no output, rather than a line of output containing just a `#`. Supposedly some old C programs contain such lines.
- The ANSI standard specifies that the `#pragma` directive has an arbitrary, implementation-defined effect. In the GNU C preprocessor, `#pragma` directives are not used, except for `#pragma once` (see “Once-only include files” on page 344). However, they are left in the preprocessor output, so they are available to the compilation pass.
- The `#ident` directive is supported for compatibility with certain other systems. It is followed by a line of text. On some systems, the text is copied into a special place in the object file; on most systems, the text is ignored and this directive has no effect. Typically, `#ident` is only used in header files supplied with those systems where it is meaningful.

9

C preprocessor output

The output from the C preprocessor looks much like the input, except that all preprocessing directive lines have been replaced with blank lines and all comments with spaces. Whitespace within a line is not altered; however, a space is inserted after the expansions of most macro calls. Source file name and line number information is conveyed by lines of the following form.

```
# linenum filename flags
```

Spaces are inserted as needed into the middle of the input (but never within a string or character constant). Such a line means that the line following it originated in file, *filename*, at line, *linenum*. After the file name comes zero or more flags, which are 1, 2, 3, or 4. If there are multiple flags, spaces separate them. What the flags mean:

- 1 indicates the start of a new file.
- 2 indicates returning to a file (after having included another file).
- 3 indicates that the text following comes from a system header file, so certain warnings should be suppressed.
- 4 indicates that the its subsequent text should be treated as C.

10

Invoking the C preprocessor

The following documentation discusses the commands as options accepted by the C preprocessor.

Most often when using the C preprocessor, you will not have to invoke it explicitly: the C compiler will do so automatically. However, the preprocessor is sometimes useful on its own.

The C preprocessor expects two file names as arguments, referred to in this documentation as *infile* and *outfile*. The preprocessor reads *infile* together with any other files that it specifies with `#include`. All the output generated by the combined input files is written in a file that the documentation refers to as *outfile*.

Either *infile* or *outfile* may use a preceding hyphen or dash, which *infile* means to read from standard input and, as *outfile*, means to write to standard output.

Also, if *outfile* or both file names are omitted, the standard output and standard input are used for the omitted file names.

What follows is a list of command options accepted by the C preprocessor. These options can also be given when compiling a C program; they are passed along automatically to the preprocessor when it is invoked by the compiler.

-P

Inhibit generation of `#`-lines with line-number information in the output from the preprocessor (see “C preprocessor output” on page 385). This might be useful when running the preprocessor on something that is not C code and is sent to a program that might be confused with the `#`-lines.

`-C`

Do not discard comments: pass them through to the output file. Comments appearing in arguments of a macro call will be copied to the output before the expansion of the macro call.

`-traditional`

Try to imitate the behavior of old-fashioned C, as opposed to ANSI C.

- Traditional macro expansion pays no attention to single-quote or doublequote characters; macro argument symbols are replaced by the argument values even when they appear within apparent string or character constants.
- Traditionally, it is permissible for a macro expansion to end in the middle of a string or character constant. The constant continues into the text surrounding the macro call.
- However, traditionally the end of the line terminates a string or character constant, with no error.
- In traditional C, a comment is equivalent to no text at all. (In ANSI C, a comment counts as whitespace.)
- Traditional C does not have the concept of a “preprocessing number”. It considers `1.0e+4` to be three tokens: `1.0e`, `+`, and `4`.
- A macro is not suppressed within its own definition, in traditional C. Thus, any macro that is used recursively inevitably causes an error.
- The `#` character has no special meaning within a macro definition in traditional C.
- In traditional C, the text at the end of a macro expansion can run together with the text after the macro call, to produce a single token. (This is impossible in ANSI C.)
- Traditionally, `\` inside a macro argument suppresses the syntactic significance of the following character.

`-trigraphs`

Process ANSI standard trigraph sequences. These are three-character sequences, all starting with `??`, that are defined by ANSI C to stand for single characters. For example, `??/` stands for `\`, so `??/n` is a character constant for a new-line. Strictly speaking, the GNU C preprocessor does not support all programs in ANSI C Standard unless `-trigraphs` is used, but if you ever notice the difference it will be with relief. You don't want to know any more about trigraphs.

`-pedantic`

Issue warnings required by the ANSI C standard in certain cases such as when text other than a comment follows `#else` or `#endif`.

- pedantic-errors
Like -pedantic, except that errors are produced rather than warnings.
- Wtrigraphs
Warn if any trigraphs are encountered (assuming they are enabled).
- Wcomment
Warn whenever a comment-start sequence `/*` appears in a comment.
- Wall
Requests both -Wtrigraphs and -Wcomment (but not -Wtraditional).
- Wtraditional
Warn about certain constructs that behave differently in traditional and ANSI C.
- I *directory*
Add the directory *directory* to the head of the list of directories to be searched for header files (see “The #include directive” on page 341). This can be used to over-ride a system header file, substituting your own version, since these directories are searched before the system header file directories. If you use more than one -I option, the directories are scanned in left-to-right order; the standard system directories come after.
- I-
Any directories specified with -I options before the -I- option are searched only for the case of `#include "file"`; they are not searched for `#include <file>`. If additional directories are specified with -I options after the -I-, these directories are searched for all `#include` directives. In addition, the -I- option inhibits the use of the current directory as the first search directory for `#include "file"`. Therefore, the current directory is searched only if -I. is requested explicitly with it. Specifying both -I- and -I. allows you to control precisely which directories are searched before and which after the current one is searched.
- nostdinc
Do not search the standard system directories for header files. Only the directories you have specified with -I options (and the current directory, if appropriate) are searched.
- nostdinc++
Do not search for header files in the C++-specific standard directories, but do still search the other standard directories. (This option is used when building libg++.)
- D *name*
Predefine *name* as a macro, with definition 1.
- D *name=definition*
Predefine *name* as a macro, with definition *definition*. There are no restrictions on the contents of *definition*, but if you are invoking the preprocessor from a

shell or shell-like program you may need to use the shell's quoting syntax to protect characters such as spaces that have a meaning in the shell syntax. If you use more than one `-D` for the same *name*, the rightmost definition takes effect.

`-U name`

Do not predefine *name*. If both `-U` and `-D` are specified for one name, the `-U` beats the `-D` and the name is not predefined.

`-undef`

Do not predefine any nonstandard macros.

`-A predicate(answer)`

Make an assertion with the predicate *predicate* and answer *answer*. See “Assertions” on page 378. You can use `-A-` to disable all predefined assertions; it also undefines all predefined macros that identify the type of target system.

`-dM`

Instead of outputting the result of preprocessing, output a list of `#define` directives for all the macros defined during the execution of the preprocessor, including predefined macros. This gives you a way of finding out what is predefined in your version of the preprocessor; assuming you have no file `foo.h`, the following command will show the values of any predefined macros.

```
touch foo.h; cpp -dM foo.h
```

`-dD`

Like `-dM` except in two respects: it does *not* include the pre-defined macros, and it outputs *both* the `#define` directives and the result of preprocessing. Both kinds of output go to the standard output file.

`-M [-MG]`

Instead of outputting the result of preprocessing, output a rule suitable for `make` describing the dependencies of the main source file. The preprocessor outputs one `make` rule containing the object file name for that source file, a colon, and the names of all the included files. If there are many included files then the rule is split into several lines using `\-newline`. `-MG` says to treat missing header files as generated files and assume they live in the same directory as the source file. It must be specified in addition to `-M`. This feature is used in automatic updating of makefiles.

`-MM [-MG]`

Like `-M`, mentions only the files included with `#include "file"`. System header files included with `#include <file>` are omitted.

`-MD file`

Like `-M` but the dependency information is written to *file*. This is in addition to compiling the file as specified—`-MD` does not inhibit ordinary compilation the way `-M` does. When invoking `gcc`, do not specify the *file* argument. `gcc` will

create file names made by replacing “.c” with “.d” at the end of the input file names.

In Mach, you can use the utility `md` to merge multiple dependency files into a single dependency file suitable for using with the `make` command.

`-MMD file`

Like `-MD` except mention only user header files, not system header files.

`-H`

Print the name of each header file used, in addition to other normal activities.

`-imacros file`

Process *file* as input, discarding the resulting output, before processing the regular input file. Because the output generated from *file* is discarded, the only effect of `-imacros file` is to make the macros defined in *file* available for use in the main input.

`-include file`

Process *file* as input, and include all the resulting output, before processing the regular input file.

`-idirafter dir`

Add the directory *dir* to the second include path. The directories on the second include path are searched when a header file is not found in any of the directories in the main include path (the one that `-I` adds to).

`-iprefix prefix`

Specify *prefix* as a prefix for using subsequent `-iwithprefix` options.

`-iwithprefix dir`

Add a directory to the second include path. The directory's name is made by concatenating *prefix* and *dir*, where *prefix* was previously specified with `-iprefix`.

`-isystem dir`

Add a directory to the beginning of the second include path, marking it as a system directory, so that it gets the same special treatment as is applied to the standard system directories.

`-lang-c`

`-lang-c89`

`-lang-c++`

`-lang-objc`

`-lang-objc++`

Specify the source language. `-lang-c` is the default; it allows recognition of C++ comments (comments that begin with `//` and end at end of line), since this is a common feature and it will most likely be in the next C standard. `-lang-c89` disables recognition of C++ comments. `-lang-c++` handles C++ comment

syntax and includes extra default include directories for C++. `-lang-objc` enables the Objective C `#import` directive. `-lang-objc++` enables both C++ and Objective C extensions. These options are generated by the compiler driver `gcc`, but not passed from the `gcc` command line unless you use the driver's `-Wp` option.

`-lint`

Look for commands to the program checker `lint` embedded in comments, and emit them preceded by `#pragma lint`.

For example, `/* NOTREACHED */` becomes `#pragma lint NOTREACHED`.

This option is available only when you call `cpp` directly; `gcc` will not pass it from its command line.

`-$`

Forbid the use of `$` in identifiers. This is required for ANSI conformance. `gcc` automatically supplies this option to the preprocessor when you specify `-ansi`, but `gcc` doesn't recognize the `-$` option itself—to use it without the other effects of `-ansi`, you must call the preprocessor directly.

Index

Symbols

- 341, 345
- # 82, 379
- # character
 - in macros 388
- ##, concatenating 360
- ##, in a macro definition, in arguments 361
- ##, in arguments 219
- #, directives in programs 337
- #assert 379
- #define 337, 349, 390
- #define directive 348
- #elif directive 374
- #else 388
- #else directive 374
- #endif 373, 374, 388
- #error directive 380
 - inside of conditionals 380
- #foo 337
- #ident directive 383
- #if 375
- #if directive 373
- #if, #ifdef or #ifndef 373
- #ifndef 344
- #import 344
- #import directive 392
- #include 339, 341, 345, 389, 389
 - variants 341
- #include directive 343
- #include, use of conditionals 381
- #include_next 345
- #line 381
- #line directive 381
- #pragma 231
- #pragma directive 383
- #pragma implementation 276
- #pragma implementation filename 280
- #pragma interface 276, 280
- #pragma once 383
- #system 378
- #system (mvs) 378
- #system (vms) 378

#unassert 379
#undef directive 362
-\$ 392
\$, dollar signs in identifier names 233
\$, in identifiers 392
&&, unary operators 205
(), as expressions in C 203
{ 204
(quotes), in #include statements 341
+=, in a cast 211
, in asm template 247
, in constants 233
.bb 108
.bbg 108
.save files 199
.verstamp 35
/* and */ 375
/* within a comment 389
// comments 391
//, to constrain comments 233
/usr/local/lib 27
?/?/, in trigraph sequences 388
_, underscores in macro names 344
__aligned__ 239
__alignof__ 234
__arg 247
__attribute__ 226
__attribute__((aligned (alignment))) 234
__BASE_FILE__ 355
__bb_hidecall__ 106
__bb_jumps__ 106
__bb_showret__ 106
__bb_trace__ 106
__builtin_apply () 208
__builtin_apply_args () 208
__builtin_return () 208
__byte__ 236
__CHAR_UNSIGNED__ 355
__cplusplus 354
__DATE__ 353

__extension__ 266
__FILE__ 353
__GNUC__ 201, 354
__GNUC_MINOR__ 354
__GNUG__ 354
__imp__ 230
__INCLUDE_LEVEL__ 355
__LINE__ 353
__OPTIMIZE__ 355
__REGISTER_PREFIX__ 355
__STDC__ 316, 354
__STDC_VERSION__ 354
__STRICT_ANSI__ 354
__TIME__ 354
__typeof__ 210
__USER_LABEL_PREFIX__ 355
__VERSION__ 355
__word__ 236
_command, used with naming 360
_exit 80
_IEEE_FP 144
{ }, as expressions in C 203
~, with complex type 215

Numerics

0 in constraint 246
386 25
char *y 210

A

abort 80, 226
abs 80
abstract datatypes 282
ADB 302
addition operations 221
addressable unit 34
aggregate initializer 222
AIX 44
AIX Threads 154
aligned 235, 239

- alignment
 - maximum 240
- alignment of types, of variables 234
- allclass.cc 277
- allclass.h 277
- alloca 55, 80, 217, 297, 302
- alloca, for Unos 41
- Alpha 26
- alternative keywords, compiling 266
- Altos 299
- AMD29K options
 - m29000 133
 - m29050 133
 - mbw 133
 - mdw 133
 - mimpure-text 134
 - mkernel-registers 133
 - mlarge 133
 - mnbw 133
 - mndw 133
 - mno-impure-text 134
 - mno-reuse-arg-regs 134
 - mnormal 133
 - mno-stack-check 134
 - mno-storem-bug 134
 - mreuse-arg-regs 134
 - msmall 133
 - msoft-float 134
 - mstack-check 134
 - mstorem-bug 134
 - muser-registers 134
- AMD29K options list 69
- angle braces 341
- ansi 316, 357
- ANSI C Standard
 - string constants 358
 - using -trigraphs 388
- ANSI C standard
 - comments = whitespace 388
 - exceptions with using backslash 341
 - naming definitions 357
 - requirements 336
- ANSI conformance 392
- ANSI or GNU C conventions 83
- ANSI Standard C 334
- ANSI Standard C requirements 334
- ANSI support 79
- ansi with keywords 266
- answer 378
- AOS 44
- APCS 136
- ar 49
- archive file 124
- arcs during compilation 107
- arcs, gcov 108
- args 219
- argument types 313
- arguments
 - with braces 350
 - with brackets 350
 - with commas 350
- arithmetic operators 373
- ARM options
 - mapcs-26 136
 - mapcs-32 136
 - mapcs-frame 136
 - mbig-endian 137
 - mbsd 138
 - mhard-float 137
 - mlittle-endian 137
 - mno-short-load-bytes 138
 - mno-symrename 138
 - mshort-load-bytes 137
 - msoft-float 137
 - mwords-little-endian 137
 - mxopen 138
- ARM options list 70
- array 206, 360
- array constructors 223
- array index 224
- array indexing in C 205
- arrays

- non-lvalue, subscripting 220
- as 49, 295
- as C++ constructors and destructors 188
- asm 266
- assembler
 - Wa, 121
- assembler code
 - labels 355
- assembler code, controlling naming process 262
- assertion
 - named question, predicate 378
- assertions
 - cpu 378
 - machine 378
 - system 378
 - testing 378
- assertions, predefined 378
- assessment of computing time 286
- AT&T C++ translator, Cfront 278
- attribute 231
 - alias 229
 - aligned () 235
 - cdecl 229
 - dllexport 230
 - dllimport 230
 - longcall 229
 - mode 236
 - nocommon 236
 - packed 236
 - regparm 229
 - section 237
 - stdcall 229
 - transparent_union 237
 - unused 238
 - weak 238
- attribute of types 239
- attributes 226
 - const 226
 - constructor 226
 - destructor 226
 - format 226

- noreturn 226
- section 226
 - specifying with underscore 226
- unused 226
- weak 226
- attributes of variables 235
- auto 207
- automatic arrays 217
- automatic variables 81, 306
- automounters 294
- AXP systems (DEC Alpha) 35

B

- b 132
- back-ends 32
- backslashes 358
- backslashes in string and character constants 358
- backslash-newline 335
- backtrace 324
- basename 277
- basic block 288
- basic block execution counts 108
- basic block execution, gcov 108
- Berkeley 44
- Binutils 49
- Bison parser generator 28
- bison parser generator 381
- bit field 203
- bitfields, unsigned 315
- blocks to linkers 278
- Borland model 278
- braces 241, 341
- braces as expressions 203
- BSD
 - configuring 33
- BSD system 294
- BSD systems 300
- BSD-style stabs 26
- BUFSIZE 370
- bug guidelines 320

- bug reports 319
- bugs
 - reporting, by class 321
- build machine 24
- building a cross-compiler 53
- building in separate directories 47
- built-in functions 208
- byte 236
- byte writes (29K) 133
- C**
- C code 381
- C dialect options
 - ansi 79
 - fallow-single-precision 83
 - fcond-mismatch 82
 - fno-asm 80
 - fno-builtin 80
 - fno-signed-bitfields 83
 - fno-unsigned-bitfields 83
 - fsigned-bitfields 83
 - fsigned-char 82
 - funsigned-bitfields 83
 - funsigned-char 82
 - fwritable-strings 83
 - pedantic 79
 - traditional 81
 - trigraphs 80
- C expression
 - arithmetic operators 373
 - character constants 373
 - identifiers 373
 - integer constants 373
- C language extensions 271
- C preprocessor
 - #include 339
 - changes to input
 - deleted backslash-newline sequences 335
 - exceptions 336
 - replaced predefined macro names 335
 - spaces 335
 - conditional compilation 334
 - CPP_PREDEFINES 357
 - definition 333
 - directives to activate 335
 - header files 339
 - header files, including 334
 - infile, outfile arguments 387
 - invoking 387
 - line control 334
 - macro 333
 - macro definitions 348
 - macro expansion 334
 - necessity of -trigraphs 388
 - output 385
 - parsing 335
 - recognizing directives 335
 - source file requirements 340
 - source files, combining 381
 - stringification 358
 - variants of #include 341
- C thread implementation 32
- C++ 22, 354
 - source file suffixes 91
 - undefining __STDC__ 317
- C++ comments 391
- C++ dialect options
 - +en 90
 - fall-virtual 85
 - falt-external-templates 86
 - fcheck-new 86
 - fconserve-space 86
 - fdollars-in-identifiers 86
 - fenum-int-equiv 86
 - fexternal-templates 86
 - ffor-scope 86
 - fguiding-decls 87
 - fhandle-signatures 87
 - fhuge-objects 87
 - fmemoize-lookups 88
 - fno-access-control 85
 - fno-for-scope 86

- fno-gnu-keywords 87
- fno-implement-inlines 87
- fno-implicit-templates 87
- fno-nonnull-objects 89
- foperator-names 89
- frepo 89
- fsave-memoized 88
- fstrict-prototype 88
- fthis-is-variable 89
- fvtable-thunks 90
- nostdinc 90
- traditional 90
- Wold-style-cast 90
- Woverloaded-virtual 90
- Wtemplate-debugging 90
- C++ run-time library 31
- canonical configuration name 24
- canonical name, defined 24
- case ranges 225
- cast 223
- cast to void 315
- casts 211
- cdecl 161
- ceil_div 365
- cexp.c 28
- cexp.y 28
- cfront 101
- Cfront model 278
- changes, unnecessary 315
- character constants 373
- charfunc 312
- class function declaration 100
- classes
 - with static members 312
- Clipper options 140
- Clipper options list 70
- clobbering 81
- code
 - hot spots 286
- code generation conventions 187
- code generation options
 - +e0 191
 - +e1 191
 - fcall-saved 190
 - fcall-used 189
 - ffixed 189
 - finhibit-size-directive 188
 - fno-common 188
 - fno-gnu-linker 188
 - fno-ident 188
 - fpack-struct 190
 - fpcc-struct-return 187
 - fPIC 189
 - fpic 189
 - freg-struct-return 187
 - fshared-data 188
 - fshort-double 188
 - fshort-enums 188
 - funaligned-pointers 192
 - funaligned-struct-hack 192
 - fverbose-asm 188
 - fvolatile 189
 - fvolatile-global 189
- code size reduction 111
- COFF 103
- collect2 188
- collect2, for start time 54
- colon after the input operands 246
- commas 231
- commas in arguments 350
- commas in statements, in arguments 219
- comment delimiters 375
- comments 233, 363
- comments, embedded 392
- Common Sub-expression Elimination 109
- common sub-expression elimination 115
- compare instructions 248
- compilation errors 29
- compilation options 198
- compiler
 - passes, statistics on 105
- compiler driver 128

- compiler versions, specifying 129
 - compiler's search directories 54
 - compiling
 - conditionally 334
 - complex automatic variables 215
 - complex conjugation 215
 - complex data types 215
 - compound expressions 211
 - compound statement 203
 - concatenation 360
 - concatenation of names, of numbers 360
 - conditional compilation 201, 334
 - conditional expressions 211
 - conditional expressions with omitted operands 213
 - conditionals 265, 344, 373
 - config.h 27
 - configuration 195
 - configuration example 24
 - configuration file names 34
 - configuration name 24
 - configure
 - error messages 33
 - configuring
 - custom naming 33
 - conflicts with names in assembler code 262
 - conforming freestanding implementation 310
 - consecutive string constants 358
 - const 266
 - CONST_DOUBLE_OK_FOR_LETTER_P 253
 - CONST_OK_FOR_LETTER_P 253
 - constant addresses 189
 - constant expressions 224
 - constants 83
 - constraint modifier characters 252
 - constraints for asm operands 249
 - constructor call 272
 - constructor expressions 223
 - constructors 54
 - containing function 217
 - containing functions 206
 - contributors to GNU CC 11
 - conversion programs 313
 - Convex options
 - margcount 141
 - mc1 141
 - mc2 141
 - mc32 141
 - mc34 141
 - mc38 141
 - mlong32 141
 - mlong64 142
 - mnoargcount 141
 - mvolatile-cache 141
 - mvolatile-nocache 141
 - Convex options list 70
 - cos 80
 - cp 27
 - c-parse.c 28
 - c-parse.c, Bison versions 28
 - c-parse.y 28
 - CPP_PREDEFINES 357
 - CPU types, supported 33
 - cross compilation problems 299
 - cross compiling 129
 - CROSS_INCLUDE_DIR 53
 - cross-assembler 49
 - cross-linker 49
 - CSE 109, 114, 115
 - C-series 960 163
 - customary abbreviations for configuring 33
- D**
- D 377
 - D options 357
 - D10V options 70
 - data flow analysis, interprocedural 272
 - DBX 300
 - debugging 26, 103, 215, 320, 372
 - debugging options
 - a 105

- ax 105
- dletters 108
- fpretend-float 109
- fprofile-arcs 107
- ftest-coverage 107
- g 103
- gcoff 104
- gcofflevel 105
- gdwarf-1 104
- gdwarf-1+ 104
- gdwarf-2 104
- ggdb 104
- ggdblevel 105
- glevel 105
- gstabs 104
- gstabs+ 104
- gstabslevel 105
- gxcoff 104
- gxcoff+ 104
- gxcofflevel 105
- p 105
- pg 105
- print-file-name=library 110
- print-libgcc-file-name 110
- print-prog-name=program 110
- print-search-dirs 110
- Q 105
- save-temps 110
- DEC Alpha 35
 - floating point operations 144
- mbuild-constants 146
- mfp-reg 144
- mfp-rounding-mode=rounding mode 145
- mfp-trap-mode=trap mode 145
- mieee 144
- mieee-conformant 146
- mieee-with-inexact 145
- mno-fp-regs 144
- mtrap-precision=trap precision 146
- DEC Alpha options 144
 - mno-soft-float 144
 - m-soft-float 144
- DEC Alpha options list 70
- DEC Unix 35
- declarations as header files 334
- definition 362
- Delta 88
 - default debugging 174
- dependencies, make 118
- derived class 100
- destructor, calling 273
- destructors 275
- DG/UX
 - default debugging 174
- differentiating 351
- diffs 328
- dir 391
- directive names, defined 337
- directives, miscellaneous 383
- directories, searching 127
- directories, specifying 127
- directory options
 - Bprefix 128
 - I- 127
 - Idir 127
 - Ldir 128
- do...while statement 366
- double precision 83
- double scan 368
- doublequote characters 358
- double-word integers 214
- DW bit (29K) 133
- DWARF 103
- DWARF format 104
- E**
 - EABI, PowerPC 157
 - ECOFF symbol table 26
 - elision algorithm 272
 - else statements 365

Emacs 44
Embedded Applications Binary Interface 157
Embedded C++ specification 86
Embedded PowerPC
 default 45
embedded statements, unallowed in constants 203
Emerald type modules 282
Encore ns32000 44
enquire 53, 295
enum 241, 267
enumeration 295
enumeration constant 203
environment variables 294
 C_INCLUDE_PATH 196
 COMPILER_PATH 196
 CPLUS_INCLUDE_PATH 196
 DEPENDENCIES_OUTPUT 196
 GCC_EXEC_PREFIX 195
 LIBRARY_PATH 196
 OBJC_INCLUDE_PATH 196
 TMPDIR 195
environment variables, defined 195
errors 318
ESC, in constants 233
exception handling 86
execution time, reducing 111
exit 80, 226
expanding a macro 314
expression space, constraints 29
expression, as C expression 373
expressions
 naming a type 209
extensions 271
extensions to GNU CC language 201
extern 81, 207
extern declaration 307
external variables 307
EXTRA_CONSTRAINT 253

F

f68881 55
fabs 80
Fairchild F9450 C compiler 34
-falt-external-templates 280
-fdollars-in-identifiers 303
-fembedded-cxx 86
ffs 80
-fhandle-signatures 282
filenames 123
fixincludes 294, 300
fixproto 294
flags
 template instantiation 280
float 308
float.h 53
FLOAT_OPTION 55
floating point 112
floating point operations 83
floating point value 299
-fno-delayed-branch 296
-fpcc-struct-return 309
-fprofile-arcs 287
FPU 160
frame pointers 113
FRAME_POINTER_REQUIRED 113
-fshort-enums 315
-ftest-coverage 287
function 217
 declarations 282
 expanded macros 204
 inline 243
 inlining 277
 jumping to a different function's code 205
function calls 264
function inlining 112
function names, as strings 268
function pointer 230
function prototype 232
function value, non-void 315

function-definition extensions 272

functions

 compiling for interrupt calls 233

 constructing calls 208

 declaring attributes 226

 with noreturn 226

 within a block 307

G

g++ 91

G++, definition 22

GAS 25, 60, 300

gcc 295

 invoking with file 390

GCC_INCLUDE_DIR 52

gcov 107, 108, 285

gcov, arcs 108

gcov, basic block execution 108

gcov, invoking 286

 GCSE 108

GDB 300

 source files, debugging 381

genflags 296

genoutput 296

glibc 310

Global Common Sub-expression Elimination 108

global common sub-expression elimination 115

global declarations 307

global declarations, explicit 198

global initializations 188

global offset table 189

global pointer 230

global symbols 125

global variables 263

globalref, globaldef, globalvalue 60

GNU

 extensions 93

GNU C

 complex data types 215

 errors, warnings 318

 extensions 266

 function-definition syntax 272

 global variables 263

 incompatibilities 306

 library 295

 operators 274

 with other compilers 300

GNU C builtin functions 83

GNU C++

 defining with `__STDC__` 317

 errors, warnings 318

 extensions 271

 function-definition syntax 272

 goto statements 275

 headers 276

 operators 274

 templates 278

 unsupported extensions 267

 unsupported nested functions 206

 with other compilers 300

GNU CC

 BSD systems 294

 bugs 293, 294, 325

 building 297

 building & installing a cross-compiler 48

 errors, warnings 318

 help services 329

 installation problems 294

 machine-description macros 308

 malloc 305

 MIPS 168

 on VMS 59

 optimization with gcov 290

 standard directory 27

 standard libraries 310

 static variables in registers 262

 string constants 306

 strings 268

 supported configurations 33

 system types 33

 testing with gcov 285

GNU CC command options 65
 GNU diff 328
 GNU linker 25, 188
 GNU Make 295
 GOT 189
 goto 204, 275
 goto * exp 205
 goto statement 205
 GPLUS_INCLUDE_DIR 52
 gprof 105
 grouping 365
 gzip 106

H

h8/300 options 135, 147
 h8/300options list 70
 hard registers 246
 hardware interrupts 233
 hardware, configuring 132
 Haskell type classes 282
 header 216
 header file 339
 header file, defining a class 277
 header files 27, 52, 127, 277, 334

- base header files 345
- GCC_INCLUDE_DIR 52
- GPLUS_INCLUDE_DIR 52
- inheritance 345
- LOCAL_INCLUDE_DIR 52
- missing 390
- with #ident 383

 header files, with redefinition 363
 header.h 343
 help 329
 heuristic 113
 host machine 24
 hot spots, code 286
 House Subcommittee on Intellectual Property 20
 HPPA 25, 302
 HPPA options

-mdisable-fpregs 149
 -mdisable-indexing 149
 -mfast-indirect-calls 149
 -mgas 150
 -mjump-in-delay 149
 -mlinker-opt 150
 -mlong-load-store 149
 -mno-space-regs 149
 -mpa-risc-1-0 149
 -mpa-risc-1-1 149
 -mportable-runtime 149
 -mschedule=cpu type 150
 -msoft-float 150
 -mspace 149

HPPA options list 71

HPUX 37, 302

HPUX 10 149

HPUX linker 150

I

-I- 389
 goto *array 205
 -I directory 389
 -I. 345
 i386 options

- m386 160
- m486 160
- malign-double 160
- malign-functions=num 162
- malign-jumps=num 162
- malign-loops=num 162
- mieee-fp 160
- mno-align-double 160
- mno-fancy-math-387 160
- mno-fp-ret-in-387 160
- mno-ieee-fp 160
- mno-svr3-shlib 161
- mno-wide-multiply 161
- mreg-alloc=regs 161
- mregparm=num 161

- mrtd 161
- msoft-float 160
- msvr3-shlib 161
- mwide-multiply 161
- i386 options list 72
- i960
 - K-series, C-series 163
- IBM RS/6000 and PowerPC 151
- IBM RS/6000 options
 - G 158
 - mads 157
 - mbig 156
 - mbig-endian 156
 - mbit-align 155
 - mcall-aix 156
 - mcall-linux 156
 - mcall-solaris 156
 - mcall-sysv 156
 - mcall-sysv-eabi 156
 - mcall-sysv-noeabi 156
 - meabi 157
 - memb 157
 - mfull-toc 153
 - mhard-float 154
 - mlittle 156
 - mlittle-endian 156
 - mminimal-toc 153
 - mmultiple 154
 - mmvme 157
 - mnew-mnemonics 152
 - mno-bit-align 155
 - mno-eabi 157
 - mno-multiple 154
 - mno-power 151
 - mno-power2 151
 - mno-powerpc 151
 - mno-powerpc-gfxopt 151
 - mno-powerpc-gpopt 151
 - mno-prototype 156
 - mno-regnames 158
 - mno-relocatable 155
 - mno-relocatable-lib 155
 - mno-sdata 158
 - mno-strict-align 155
 - mno-string 154
 - mno-sum-in-toc 153
 - mno-toc 156
 - mno-traceback 156
 - mno-xl-call 154
 - mold-mnemonics 152
 - mpe 154
 - mpower 151
 - mpower2 151
 - mpowerpc 151
 - mpowerpc-gfxopt 151
 - mpowerpc-gpopt 151
 - mprototype 156
 - mregnames 158
 - mrelocatable 155
 - mrelocatable-lib 155
 - msdata 158
 - msdata=default 158
 - msdata=eabi 158
 - msdata=none 158
 - msdata=sysv 158
 - msdata-data 158
 - msim 157
 - msoft-float 154
 - mstrict-align 155
 - mstring 154
 - mthreads 154
 - mtoc 156
 - mtraceback 156
 - mtune= 153
 - mxl-call 154
 - myellowknife 157
- IBM RS/6000 SP Parallel Environment 154
- IBM RT options
 - mcall-lib-mul 159
 - mfp-arg-in-fpregs 159
 - mfp-arg-in-gregs 159
 - mfull-fp-blocks 159

- mhc-struct-return 159
- min-line-mul 159
- mminimum-fp-blocks 159
- mnohc-struct-return 159
- IBM RT PC
 - supported operating systems 44
- IBM XLC compiler 44
- identifier 337
- identifier names 233
- identifiers 373
- idirafter dir 391
- IEEE compliant code 144
- IEEE floating point standard 144
- IEEE rounding mode 145
- immediate integer operand 250
- implementation of object definition in headers 276
- index values 224
- infinite recursion 345
- inheritance 345
- inheritance in header files 345
- inherited variables 206
- initialization functions 54
- initialized global definition 237
- initializer 209
- inline 243, 266
- inline functions 100
- input operands 245
- inside a macro 388
- insn-emit.c 29
- install
 - binaries for Sun 55
 - compiler driver 31
 - libg++ distribution 31
- Installation 24
- installing 27
 - naming scheme 31
- installing GNU CC
 - Microsoft compilers 29
 - Ultrix compilers 29
- instantiation of templates 278
- instructions
 - compare 248
 - store 248
 - test 248
- int 214, 232, 306, 366, 376
- integer constant 214
- integer constants 299, 373
- Intel 960 options
 - m 163
 - masm-compat 163
 - mcode-align 163
 - mcomplex-addr 163
 - mic2.0-compat 163
 - mic3.0-compat 163
 - mic-compat 163
 - mintel-asm 163
 - mleaf-procedures 163
 - mno-code-align 163
 - mno-complex-addr 163
 - mno-leaf-procedures 163
 - mno-strict-align 164
 - mno-tail-call 163
 - mnumerics 163
 - mold-align 164
 - msoft-float 163
 - mstrict-align 164
 - mtail-call 163
- Intel 960options list 72
- intellectual property 20
- interface specification 276
- intermediate 206
- internal linkage 207
- interpreter function 205
- invalid assembler code 262
- invalid C code 366
- iprefix prefix 391
- IRIX thread support 32
- ISC 25
- ISO/ANSI C standard 310
- isystem dir 391
- iwithprefix dir 391

J

JALR instruction
 for MIPS 168
jmp_buf 294

K

K&R C 83
keyword
 with parenthesis(es) 226
K-series 960 163

L

-l 196
label
 with colon or period= syntax 224
labeled elements of an initializer 224
labels
 addresses stored in automatic variables 205
labels, as values, as constants 205
labels, with indexing 205
labs 80
-lang-c 391
 default source language 391
-lang-c++ 391
-lang-c89 391
-lang-objc 391
-lang-objc++ 391
LANGUAGES=c 28
ld 49, 54, 295
leaf procedures 163
League for Programming Freedom 19
lexical scoping 206
lexical units, valid usage 361
-lgcc 124
libg++ 31
libgcc.a 50
libgcc1.a 50
libgcc1.c 50
libgcc1-test 53
libraries
 for cross-compiling 49
libraries, linking 123, 124
library 123
limit.h 355
line control 334
linker 123
linker errors 277
linker options
 -c 123
 -E 123
 -llibrary 123
 -lobjc 124
 -nodfaultlibs 124
 -nostartfiles 124
 -nostdlib 124
 object-file-name 123
 -S 123
 -s 124
 -shared 125
 -static 124
 -symbolic 125
 -u 125
 -Wl 125
 -Xlinker 125
linking libraries 123
linking object files 123
-lint 392
Linux SLS 1.01 295
Linux-based GNU systems 32
LL, adding to an integer 214
local labels as identifiers 204
local register variables 263, 265
local variables 209
local variables, when compiling 310
local version of applications 345
LOCAL_INCLUDE_DIR 53
long 307, 373
long long int 214
longjmp 81, 98, 306
loop optimizer 115
loop unrolling 112

loops 286
 deleting 317
 empty 317
 lvalue 234
 lvalues 211, 220, 245
 unallowed unary operators (&) 220

M

M32R/D options 72
 M680x0 options
 -m68000 171
 -m68020 171
 -m68020-40 171
 -m68030 171
 -m68040 171
 -m68060 171
 -m68881 171
 -mbitfield 172
 -mc68000 171
 -mc68020 171
 -mfpa 172
 -mnobitfield 172
 -mrtd 172
 -mshort 172
 -msoft-float 172
 M680x0 options list 73
 M88k options
 -m88000 174
 -m88100 174
 -m88110 174
 -mbig-pic 174
 -mcheck-zero-division 176
 -mhandle-large-shift 176
 -midentify-revision 174
 -mno-check-zero-division 176
 -mno-ocs-debug-info 174
 -mno-ocs-frame-position 174
 -mno-optimize-arg-area 174
 -mno-serialize-volatile 175
 -mno-underscores 174

 -mocs-debug-info 174
 -mocs-frame-position 174
 -moptimize-arg-area 174
 -mserialize-volatile 175
 -mshort-data 175
 -msvr 175
 -msvr3 175
 -mtrap-large-shift 176
 -muse-div-instruction 176
 -mversion 176
 -mwarn-passed-structs 177
 M88K options list 73
 MACH 44
 Mach 391
 MACH thread support 32
 machine dependent options 129
 machine names 34
 macro
 arguments, expecting 350
 concatenation 360
 definitions 352
 double scan 368
 expansion text 350
 min 366
 names, functions 351
 stringification 358
 using parentheses 352
 using spaces 352
 macro calls 373
 macro definitions, safe, maximum 203
 macro expansion 334
 macro processor 333
 macro, substituted arguments 367
 macros 347
 AM29000K series 356
 appending 351
 arguments 350
 assertions 376
 calls, substituted 367
 combining source files 381
 compound statements 365

- conditionals 376
- conditionals as directives 371
- controlling file names 341
- customizing a program 377
- defining 333
- errors 370, 388
- expansion 367
- for Motorola 356
- for Vax 356
- function-like 350
- input
 - string, character constants 385
- long definitions 348
- M68000 series 356
- m68k 356
- names 348
- nested calls 368
- nested constructs 364
- newlines 370
- ns 32000 series 356
- parentheses 364
- pitfalls, subtleties 364
- predefined 353, 390
- predefined, non-standard 356
 - Pyramid 356
 - Sequent 356
 - Sun 356
- pre-scan 368
- quote characters 388
- recursive 388
- redefining 363
- self-referential 367
- simple 348
- source files 390
- standard 353
- string or character constants 388
- stringified, concatenated 368
- system-specific 353
- undefined 362
- unintended grouping 365
- using assertions 378
 - using parentheses 348
 - variables, functions 362
 - with # 388
 - with side effect 366
 - with text after 388
- macros with variable numbers of arguments 219
- macros, predefined, on Vax 376
- main 54
- make 118
 - environmental variables 294
- make install 31
- make LANGUAGES=c 28
- Makefile 27
- makefile variable 32
- malloc 37, 305
- matching quote characters 348
- maximum arguments in C operators 274
- mcpu 152
- mcpu= 179
- member functions 282
- memcmp 80
- memcpy 80
- memory model (29K) 133
- memory operand 249
- Metaware 44
- MetaWare HighC (hc) compiler 159
- min, unsafe macro 366
- minimum arguments in C operators 274
- MIPS 26, 297
 - requisite libraries for GNU CC 168
- mips
 - cross-compiling 48
 - in BSD mode 42
- mips C compiler
 - switch statements 42
- MIPS options
 - EB 169
 - EL 169
 - G num 169
 - m4650 169
 - mabiccalls 168

- mcpu= 166
 - mdouble-float 168
 - membedded-data 168
 - membedded-pic 168
 - mfp32 166
 - mfp64 166
 - mgas 167
 - mgrp32 166
 - mgrp64 166
 - mgpopt 167
 - mhalf-pic 168
 - mhard-float 168
 - mint64 166
 - mips1 166
 - mips2 166
 - mips3 166
 - mlong64 166
 - mlong-calls 168
 - mmad 169
 - mmemcpy 167
 - mmips-as 167
 - mmips-tfile 167
 - mno-abicalls 168
 - mno-embedded-data 168
 - mno-embedded-pic 168
 - mno-gpopt 167
 - mno-half-pic 168
 - mno-long-calls 168
 - mno-mad 169
 - mno-memcpy 167
 - mno-mips-tfile 167
 - mno-rnames 167
 - mno-stats 167
 - mrnames 167
 - msingle-float 168
 - msoft-float 168
 - mstats 167
 - nocpp 169
 - MIPS options list 72
 - mips-mips-riscosrev 43
 - mips-mips-riscosrevbsd 43
 - mips-mips-riscosrevsysv 43
 - mips-mips-riscosrevsysv4 43
 - mips-sgi-irix5.* 25
 - miscellaneous preprocessing directives 383
 - mktemp 306
 - ML signatures 282
 - mno-serialize-volatile 303
 - mode 236
 - Modula-2, Modula-3 282
 - modules, optimizing 286
 - Motorola 356
 - multiple alignments 239
 - multiple alternative operands 251
 - multiple attributes 238
 - multiple attributes, specifying 231
 - multiple basic blocks, branches 288
 - multiple inheritance 86
- ## N
- name mangling 63
 - National Semiconductor ns32000 44
 - nested conditionals 374
 - nested function 206
 - nested loops 204
 - nested type declarations 282
 - NeXT operating system 297
 - NEXTSTEP 32
 - nfp 25
 - nfp 26
 - nocommon 236
 - non-ANSI constructs, warnings and errors 318
 - non-constant initializers 222
 - noncontiguous complex variable 215
 - non-GNU linker 188
 - non-void function value 315
 - nostdinc 341
 - ns16000 376
 - null directive 383
 - null statements 365

O

- O 111, 296
- OBJC_THREAD_FILE 32
- Object Compatibility Standard 174
- object definitions 276
- object files, linking 123
- Objective C 392
- Objective C extensions 392
- Objective C thread implementation 32
- objects
 - temporary 312
- offset 206
- old-style (C-style) cast 100
- old-style function definitions 199
 - unsupported 232
- old-style non-prototype definition 232
- omitted operands, in conditionals 213
- once-only include files 344
- operands 213
- operands, types of 203
- operators, minimum, maximum 274
- optimization 286
- optimization options
 - fbranch-probabilities 116
 - fcaller-saves 116
 - fcse-follow-jumps 114
 - fcse-skip-blocks 114
 - fdelayed-branch 115
 - fexpensive-optimizations 115
 - ffast-math 114
 - ffloat-store 112
 - fforce-addr 113
 - fforce-mem 113
 - ffunction-sections 115
 - fgcse 115
 - finline-functions 113
 - fkeep-inline-functions 113
 - fkeep-static-consts 114
 - fno-default-inline 112
 - fno-defer-pop 113
 - fno-function-cse 114
 - fno-inline 113
 - fno-peephole 116
 - fomit-frame-pointer 113
 - freduce-all-givs 116
 - frerun-cse-after-loop 115
 - frerun-loop-opt 115
 - fschedule-insns 115
 - fschedule-insns2 115
 - fstrength-reduce 114
 - fthread-jumps 114
 - funroll-all-loops 116
 - funroll-loops 116
- O 111
- O0 112
- O1 111
- O2 112
- O3 112
- optimizations 114
- option
 - build 24
 - GNU assembler 25
- options
 - compiling 198
 - gcc 198
 - gcoff 26
 - gstabs+ 26
 - passing to assembler 121
- options list
 - code generation 74
- OS/2 thread support 32
- OSF/1 167
- OSF/1 thread support 32
- output file option 77
- output operand expressions 245
- output operands, write-only 245

P

- PA systems and kernels 149
- packed 236

-
- parameter forward declaration 217
 - parentheses 231, 348, 364
 - parentheses as expressions 203
 - parenthesis 364
 - parser files 381
 - parsing 335
 - patches for bugs 325
 - Patents, Trademarks and Copyrights 20
 - PCC 307
 - PCThreads 32
 - pedantic 201, 266, 318, 334, 388
 - pedantic-errors 86, 294, 389
 - Perl 305
 - PIC 189
 - PIC code 168
 - PIC support 189
 - pointer 236
 - pointer-to-function types 313
 - Portable C Compiler 159
 - porting 342
 - position-independent code 189
 - POSIX thread support 32
 - POWER 151
 - POWER instruction set 151
 - PowerPC 151
 - configuration 46
 - running System Version4 45
 - running Windows NT 46
 - PowerPC instruction set 151
 - predecrement addressing 249
 - predefined macros 201
 - prefix 391
 - preincrement addressing 249
 - preprocessing 307
 - preprocessing conditionals 372
 - preprocessing directives 373
 - preprocessing number 388
 - preprocessor
 - ## in rest arguments 219
 - directive name 337
 - preprocessor options 117
 - A 119
 - C 118
 - D 119
 - dD 120
 - dM 119
 - dN 120
 - E 118
 - H 119
 - idirafter 118
 - imacros 117
 - include 117
 - iprefix 118
 - isystem 118
 - iwithprefix 118
 - iwithprefixbefore 118
 - M 118
 - MD 119
 - MG 119
 - MM 119
 - MMD 119
 - nostdinc 118
 - P 118
 - trigraphs 120
 - Umacro 119
 - undef 118
 - Wp 120
 - pre-scan 368
 - nested calls 368
 - self-referent macros 368
 - unshielded commas 368
 - private data 188
 - processor selection (29K) 133
 - prof 105
 - profiling 106, 300
 - profiling tools 286
 - program checker 392
 - program.c 343
 - protoize 197, 313
 - global declarations 198
 - prototypes 197

PSIM simulator 46

pthread 154

purify 295

Pyramid 357

Q

qsort 264

quotes 341

quotes, balancing 348

R

r 272

r4650 169

ranlib 49

read-write operand 246

real.c 295

REAL_LD_FILE_NAME 54

real-ld 54

redefining macros (with #define) 363

redefinition 363

REG_CLASS_FROM_LETTER 253

register 246

register variables 265

REGISTER_NAMES 189

reporting a bug 319

rest arguments 219

RISC-OS

 reconfiguring 43

RS/6000 300

RS/6000, PowerPC options list 71

RT options list 71

rtd 172

RTTI 86

run-time varying elements 222

S

saveset

 for VMS 55

scalar types 223

scope

 unallowed jumping 217

scope, of variables 86

Scratchpad II 282

SDB 104

search path 127

second include path 118

self-reference, indirect 367

self-referents 367

semicolon, following an expression 203

semicolon, swallowing 365

semicolons in GNU assembler 247

Senate Subcommittee on Patents, Trademarks and
 Copyrights 20

Sequent 357

setjmp 306

sh 296

SH options 71, 148

shared data 188

shift count operands 315

side effect 366

signal.h 345

signature 282

signatures 282

signed 266

signed integer 214

Silicon Graphics

 compiling GNU CC on IRIX 43

simple constraints 249

sin 80

sizeof 210, 221, 234, 376

SKIP_SPACES 365

smallest addressable unit 34

Solaris

 installing GNU CC 55

Solaris thread support 32

Sony

 compiling 44

source

 line control 334

source file 362

SPARC options

- mapp-regs 178
- mcpu=cpu_type 180
- mcypress 180
- mepilogue 179
- mflat 179
- mfpv 178
- mfullany 181
- mhard-float 178
- mhard-quad-float 178
- mint32 181
- mint64 181
- mlong32 181
- mlong64 181
- mmedany 181
- mmedlow 181
- mno-app-regs 178
- mno-epilogue 179
- mno-flat 179
- mno-fpv 178
- mno-stack-bias 181
- mno-unaligned-doubles 179
- msoft-float 178
- msoft-quad-float 178
- msparclite 179
- mstack-bias 181
- msupersparc 180
- mtune=cpu_type 180
- munaligned-doubles 179
- mv8 179
- SPARC options list 73
- specify -ansi 392
- specifying targets 129
- sqrt 80
- stabs 103
- stabs debugging output 104
- standard C extensions 201
- start files 50
- STARTFILE_SPEC 50
- static data members, declaring, defining 312
- static variable 203
- static variables in registers 262
- stdcall 161
- stdio.h 297
- store 206
- store instructions 248
- storem bug (29) 134
- strcmp 80
- strcpy 80
- strfunc 312
- string constants 336, 358, 382
- string constants end at newline 82
- string or character constants 385
- string variables 268
- stringification 358
- strlen 80
- struct mumble 311
- STRUCT_VALUE 308
- structure and union returning 308
- structure initializer 224
- structure types, defining 344
- structures, arraying 360
- submodel options 132
- subscript 205
- subtraction operations 221
- Sun 357
- SVr4
 - default debugging 174
- switch statements 205
- symbolic links
 - unsupported 27
- symbols, assembler 262
- sys/signal.h 345
- System V 297, 305
- System V options 182
 - G 182
 - Qn 182
 - Qy 182
 - Ym 182
 - YP 182
- System V options list 73

T

- Table Of Contents, executable files 153
- tail-recursive calls 163
- target machine 24
- target machine, specifying 129
- target options
 - b 129
 - v 129
- tconfig.h 27
- tcov 105
- Tektronix
 - install 42
- templates 86, 278
- test coverage 286
- test instructions 248
- testsuites 286
- thread support 32
- threaded code 205
- Thumb options 183
- tilde operator, (~) 215
- time 308
- tm.h 27, 357
- tmp.c.gcov 287
- tokens, syntactic 360
- TOOL_INCLUDE_DIR 53
- tools
 - for cross-compiling 49
- trace 106
- traditional 306
- traditional with keywords 266
- trampolines 207
- transparent_union 237, 239
- Trellis/Owl type modules 282
- trigraphs 334, 388
- trigraphs 335
 - enabling 336
- type attributes 239
- typedef 210, 241, 294
- typedef name = 209
- typedefs 344
- typeof 203, 210, 266

types

- argument 313
- pointer-to-function 313

U

- U 377
- U and -D 390
- uid_t 232
- ULL, adding to an integer 214
- unaligned addresses 192
- unary operator, && 205
- undef 334
- undefined macros 362
- undefinition 362
- underscore, for C language standard 262
- underscores
 - using 209
- Unidiff format 328
- uninitialized global variables 188
- union 225
- union returning 308
- union type 224, 225
- union types 223
- unprotoize 313
- unresolved references
 - nodefaultlibs 124
 - nostdlib 124
- unsigned integer 214
- unsigned long 373
- unsigned long long int 214
- unused 238
- uppercase usage in macro names 348
- using '-g' when reporting bugs 324
- UTEK ns32000 44

V

- v with GAS 25
- variable attributes 235
- variable names 209
- variable-length arrays 217

variable-length automatic arrays 217
 variables
 noncomplex 215
 noncontiguous complex 215
 variables in macros 367
 variables, scope 86
 Vax 357
 compiling 46
 vax 356, 376
 Vax computers 356
 VAX options
 -mg 185
 -mgnu 185
 -munix 185
 VAX options list 74
 VAX-C 59, 60
 VAXCRTL library 58
 version of GNU CC 129
 virtual classes 283
 virtual function calls 300
 virtual function definitions in classes 191
 virtual functions 100
 VMS 59
 installing, with CLD file 55
 void and function pointers 221
 volatile 266, 306

W

-W with -O 307
 -Wall 389
 WARN_IF 358
 warning options
 -fsyntax-only 93
 -pedantic 93
 -pedantic-errors 94
 -W 98
 -w 94
 -Waggregate-return 100
 -Wall 97
 -Wbad-function-cast 99

-Wcast-align 99
 -Wcast-qual 99
 -Wchar-subscripts 94
 -Wcomment 94
 -Wconversion 99
 -Weffc++ 90
 -Werror 101
 -Wformat 94
 -Wid-clash-len 99
 -Wimplicit 94
 -Winline 100
 -Wlarger-than-len 99
 -Wmain 95
 -Wmissing-declarations 100
 -Wmissing-prototypes 100
 -Wnested-externs 100
 -Wno-import 94
 -Wold-style-cast 100
 -Woverloaded-virtual 100
 -Wparentheses 95
 -Wpointer-arith 99
 -Wredundant-decls 100
 -Wreorder 97
 -Wreturn-type 95
 -Wshadow 99
 -Wstrict-prototypes 100
 -Wswitch 96
 -Wsynth 101
 -Wtemplate-debugging 97
 -Wtraditional 98
 -Wtrigraphs 96
 -Wuninitialized 96
 -Wunused 96
 -Wwrite-strings 99
 warnings 318, 388
 warnings into errors 101
 -Wcomment 389
 -Wconversion 314
 we32k-*-* 47
 weak 238

Index

whitespace 337, 359, 365
whitespace, exceptions 363
whitespace, in input 385
-Wimplicit-function-declarations 94
-Wimplicit-int 94
Win32 API thread support 32
--with-gnu-as 25
--with-gnu-ld 25
--with-stabs 25
word 236
-Wpointer-arith 221
-Wsign-compare 100

-Wtraditional 389
-Wtrigraphs 389
-Wunknown-pragmas 97

X

X11 header files 305
XCOFF 103
XCOFF format 104

Z

zero-length arrays 216