

GNUPro Utilities

98r2

Using as
Using ld
Using binutils
Using make
Using diff & patch
Using info

CYGNUS

Copyright © 1988-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the documentation entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Free Software Foundation

59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

This documentation has been prepared by Cygnus Technical Publications; contact the Cygnus Technical Publications staff: doc@cygnus.com.

Part #: 300-400-1010045-98r2

GNUPro warranty

The GNUPro Toolkit is free software, covered by the GNU General Public License, and you are welcome to change it and/or distribute copies of it under certain conditions. This version of GNUPro Toolkit is supported for customers of Cygnus.

For non-customers, GNUPro Toolkit software has NO WARRANTY.

Because this software is licensed free of charge, there are no warranties for it, to the extent permitted by applicable law. Except when otherwise stated in writing, the copyright holders and/or other parties provide the software “as is” without warranty of any kind, either expressed or implied, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose. The entire risk as to the quality and performance of the software is with you. Should the software prove defective, you assume the cost of all necessary servicing, repair or correction.

In no event, unless required by applicable law or agreed to in writing, will any copyright holder, or any other party who may modify and/or redistribute the program as permitted above, be liable to you for damages, including any general, special, incidental or consequential damages arising out of the use or inability to use the program (including but not limited to loss of data or data being rendered inaccurate or losses sustained by you or third parties or a failure of the program to operate with any other programs), even if such holder or other party has been advised of the possibility of such damages.

How to contact Cygnus

Use the following means to contact Cygnus.

Cygnus Headquarters

1325 Chesapeake Terrace
Sunnyvale, CA 94089 USA
Telephone (toll free): +1 800 CYGNUS-1
Telephone (main line): +1 408 542 9600
Telephone (hotline): +1 408 542 9601
FAX: +1-408 542 9699
(Faxes are answered 8 a.m.–5 p.m., Monday through Friday.)
email: info@cygnus.com
Website: www.cygnus.com.

Cygnus United Kingdom

36 Cambridge Place
Cambridge CB2 1NS
United Kingdom
Telephone: +44 1223 728728
FAX: +44 1223 728728
email: info@cygnus.co.uk/

Cygnus Japan

Nihon Cygnus Solutions
Madre Matsuda Building
4-13 Kioi-cho Chiyoda-ku
Tokyo 102-0094
Telephone: +81 3 3234 3896
FAX: +81 3 3239 3300
email: info@cygnus.co.jp
Website: <http://www.cygnus.co.jp/>

Use the hotline (+1 408 542 9601) to get help, although the most reliable and most expedient means to resolve problems with GNUPro Toolkit is by using the Cygnus Web Support site:

<http://support.cygnus.com>

Contents

GNUPro warranty	iii
How to contact Cygnus	iv

Using **as**

<i>Overview of as, the GNU assembler</i>	5
Overview of special features for as	6
<i>Invoking as, the GNU assembler</i>	7
Summary of invocation of as options	8
Object file formats	9
Command line	10
Input files	11
<i>Filenames and line-numbers</i>	11
Output (object) file	12
Error and warning messages	13
<i>Command-line options</i>	15
Enable listings: -a[cdhlns] options	17
-D option	18
Work faster: -f option	19
.include search path: -I path option	20

Difference tables: -k option	21
Include local labels: -L option	22
Assemble in MRI compatibility mode: -m option	23
Dependency tracking: --md option	26
Name the object file: -o option	27
Join data and text sections: -R option	28
Display assembly statistics: --statistics option	29
Announce version: -v option	30
Suppress warnings: -w option	31
Generate object file in spite of errors: -z option	32
Syntax	33
Preprocessing	34
<i>Whitespace</i>	34
Comments	35
Symbols	36
Statements	37
Constants	38
<i>Character constants</i>	38
<i>Strings</i>	38
<i>Characters</i>	39
<i>Number constants</i>	39
<i>Bignums</i>	40
<i>Flonums</i>	40
Sections and relocation	43
ld sections	46
as internal sections	47
Sub-sections	48
bss section	50
Symbols for the GNU assembler	51
Labels	52
Giving symbols other values	53
Symbol names	54
<i>Local symbol names</i>	54
The special dot symbol	56
Symbol attributes	57
<i>Value</i>	57
<i>Type</i>	57
<i>Symbol attributes for a.out file format</i>	57
<i>Symbol attributes for COFF format</i>	58

<i>Symbol attributes for SOM format</i>	58
Expressions	59
Empty expressions	60
Integer expressions.....	61
Arguments	62
Operators	63
<i>Prefix operators</i>	63
<i>Infix operators</i>	63
Assembler directives	65
<i>.abort directive</i>	67
<i>.ABORT directive</i>	68
<i>.align abs-expr, .abs-expr directive</i>	68
<i>.app-file string directive</i>	68
<i>.ascii "string"... directive</i>	68
<i>.asciz "string"... directive</i>	69
<i>.balign[wl] abs-expr, abs-expr directive</i>	69
<i>.byte expressions directive</i>	69
<i>.comm symbol, length directive</i>	69
<i>.data subsection directive</i>	70
<i>.def name directive</i>	70
<i>.desc symbol, abs-expression directive</i>	70
<i>.dim directive</i>	70
<i>.double flonums directive</i>	70
<i>.eject directive</i>	70
<i>.else directive</i>	71
<i>.endif directive</i>	71
<i>.endif directive</i>	71
<i>.equ symbol, expression directive</i>	71
<i>.equiv symbol, expression directive</i>	71
<i>.err directive</i>	71
<i>.extern directive</i>	72
<i>.file string directive</i>	72
<i>.fill repeat, size, value directive</i>	72
<i>.float flonums directive</i>	72
<i>.global symbol, .globl symbol directive</i>	73
<i>.hword expressions directive</i>	73
<i>.ident directive</i>	73
<i>.if absolute expression directive</i>	73
<i>.include "file" directive</i>	74
<i>.int expressions directive</i>	74
<i>.irp symbol, values... directive</i>	74
<i>.irpc symbol, values... directive</i>	74
<i>.lcomm symbol, length directive</i>	75
<i>.lflags directive</i>	75
<i>.line line-number directive</i>	75
<i>.list directive</i>	76

<code>.ln</code> <i>line-number directive</i>	76
<code>.long</code> <i>expressions directive</i>	76
<code>.macro</code> <i>directive</i>	76
<code>.mri</code> <i>val directive</i>	77
<code>.nolist</code> <i>directive</i>	77
<code>.octa</code> <i>bignums directive</i>	77
<code>.org</code> <i>new-lc, fill directive</i>	78
<code>.p2align[wl]</code> <i>abs-expr, abs-expr directive</i>	78
<code>.psize</code> <i>lines, columns directive</i>	79
<code>.quad</code> <i>bignums directive</i>	79
<code>.rept</code> <i>count directive</i>	79
<code>.sbt1</code> <i>"subheading" directive</i>	79
<code>.scl</code> <i>class directive</i>	80
<code>.section</code> <i>name directive</i>	80
<code>.set</code> <i>symbol, expression directive</i>	81
<code>.short</code> <i>expressions directive</i>	82
<code>.single</code> <i>flonums directive</i>	82
<code>.size</code> <i>directive</i>	82
<code>.sleb128</code> <i>expressions directive</i>	82
<code>.skip</code> <i>size, fill directive</i>	82
<code>.space</code> <i>size, fill directive</i>	82
<code>.stabd</code> , <code>.stabs</code> , and <code>.stabs</code> <i>directives</i>	83
<code>.string</code> <i>"str" directive</i>	84
<code>.symver</code> <i>directive</i>	84
<code>.tag</code> <i>structname directive</i>	85
<code>.text</code> <i>subsection directive</i>	85
<code>.title</code> <i>"heading" directive</i>	85
<code>.type</code> <i>int directive</i>	85
<code>.val</code> <i>addr directive</i>	85
<code>.uleb128</code> <i>expressions directive</i>	85
<code>.word</code> <i>expressions directive</i>	86
Machine dependent features for as	89
AMD 29K dependent features	91
AMD 29K options.....	92
AMD 29K syntax	93
AMD 29K macros	93
AMD 29K special characters.....	93
AMD 29K register names	93
AMD 29K floating point	94
AMD 29K machine directives	95
AMD 29K opcodes.....	95
ARC dependent features	97
ARC command-line options.....	98
ARC floating point	98
ARC machine directives	99

ARM dependent features	101
ARM options	102
ARM syntax	103
ARM special characters.....	103
ARM floating point.....	103
ARM machine directives.....	104
ARM opcodes	104
AT&T/Intel dependent features	105
AT&T syntax versus Intel syntax	106
AT&T opcode naming.....	106
80386 register naming.....	107
80386 opcode prefixes	108
AT&T memory references.....	109
Handling of jump instructions for 80386 machines.....	109
AT&T floating point.....	110
Writing 16-bit Code for 80386 machines	111
Notes on AT& T and the 80386 instructions.....	111
D10V dependent features	113
D10V options	114
D10V syntax.....	115
D10V size modifiers.....	115
D10V sub-instructions	115
D10V special characters.....	116
D10V register names	117
D10V addressing modes	118
@WORD modifier for D10V.....	118
D10V floating point	118
D10V opcodes.....	119
H8/300 dependent features.....	121
H8/300 options	122
H8/300 syntax	123
Special characters.....	123
H8/300 register nnames.....	123
H8/300 addressing modes.....	123
H8/300 floating point.....	124
H8/300 machine directives.....	125
H8/300 opcodes.....	126
H8/500 dependent features.....	127
H8/500 options	128
H8/500 syntax	129
H8/500 special characters.....	129
H8/500 register names.....	129

<i>H8/500 addressing modes</i>	129
<i>H8/500 floating point</i>	130
H8/500 machine directives.....	131
<i>H8/500 opcodes</i>	131
<i>HPPA dependent features</i>	133
HPPA options.....	134
HPPA syntax	135
<i>HPPA floating point</i>	135
<i>HPPA assembler directives</i>	135
<i>HPPA opcodes</i>	138
<i>SH dependent features</i>	139
Hitachi SH options	140
Hitachi SH syntax	141
<i>Hitachi SH special characters</i>	141
<i>Hitachi SH register names</i>	141
<i>Hitachi SH addressing modes</i>	141
<i>Hitachi SH floating point</i>	142
Hitachi SH machine directives.....	143
<i>Hitachi SH opcodes</i>	143
<i>Intel 960 dependent features</i>	145
i960 command-line options.....	146
<i>i960 floating point</i>	147
i960 machine directives.....	148
<i>i960 opcodes</i>	148
The <code>callj</code> instruction.....	149
i960 Compare-and-Branch.....	150
<i>M68K dependent features</i>	151
M680x0 options	152
M680x0 syntax.....	154
Standard Motorola syntax differences	155
<i>M680x0 floating point</i>	156
M680x0 machine directives	157
<i>M680x0 opcodes</i>	157
<i>M680x0 branch improvement</i>	157
<i>M680x0 special characters</i>	159
<i>MIPS dependent features</i>	161
Assembler options for MIPS.....	162
<i>MIPS ECOFF object code</i>	163
Directives for debugging information for MIPS.....	164
Directives to override the ISA level for MIPS.....	165
Directives for extending MIPS 16 bit instructions.....	166

Directive to mark data as an instruction for MIPS.....	167
Directives to save and restore options for MIPS.....	168
SPARC dependent features	169
SPARC options	170
SPARC floating point.....	170
SPARC machine directives	171
Vax dependent features	173
Vax command-Line options.....	174
Vax floating point.....	175
Vax machine directives	176
Vax opcodes	176
Vax branch improvement	176
Vax operands	178
Unsupported Vax assembler properties.....	178
Acknowledgments for the GNU assembler	181

Using ld

Overview of ld, the GNU linker	187
Invocation of ld, the GNU linker	189
ld command line options	190
ld environment variables	202
Linker scripts	203
Basic linker script concepts.....	204
Linker script format.....	205
Simple linker script example.....	206
Simple linker script commands	207
Setting the entry point.....	207
Commands dealing with files.....	207
Commands dealing with object file formats	208
Other linker script commands.....	209
Assigning values to symbols	210
Simple assignments.....	210
PROVIDE command.....	211
SECTIONS command	212
Output section description.....	212
Output section name	213
Output section address.....	213
Input section description.....	214
Input section basics.....	214
Input section wildcard patterns	215

<i>Input section for common symbols</i>	216
<i>Input section example</i>	216
<i>Output section data</i>	217
<i>Output section keywords</i>	218
<i>Output section discarding</i>	219
<i>Output section attributes</i>	219
<i>Output section type</i>	220
<i>Output section LMA</i>	220
<i>Output section region</i>	221
<i>Output section phdr</i>	221
<i>Output section fill</i>	221
<i>Overlay description</i>	222
MEMORY command	224
PHDRS command	226
VERSION command	229
Expressions in linker scripts.....	232
<i>Constants</i>	232
<i>Symbol names</i>	232
<i>The location counter</i>	232
<i>Operators</i>	233
<i>Evaluation</i>	234
<i>The section of an expression</i>	234
<i>Builtin functions</i>	235
1d machine dependent features	239
1d and the H8/300 processors	240
1d and the Intel 960 processors	241
BFD libraries	243
How it works: an outline of BFD	245
Information loss	246
The BFD canonical object-file format	247
MRI compatible script files for the GNU linker	249

Using binutils

Overview of binutils, the GNU binary utilities	255
Controlling ar on the command line	257
Controlling ar with a script.....	260
Selecting the target system	293
Target selection	294
objdump target.....	294
objcopy and strip input target.....	294
objcopy and strip Output target.....	295
nm , size , and strings target	295

Linker input target 295
Linker output target 295
 Architecture selection..... 296
 objdump architecture 296
 objcopy, nm, size, strings architecture..... 296
 Linker input architecture 296
 Linker output architecture 296
 Linker emulation selection..... 297

Using make

Overview of make, a program for recompiling..... 301
Introduction to makefiles 303
 What a rule looks like 305
 A simple makefile 306
 How **make** processes a makefile 308
 Variables make makefiles simpler 309
 Letting **make** deduce the commands..... 310
 Another style of makefile..... 311
 Rules for cleaning the directory 312
Writing makefiles..... 313
 What makefiles contain..... 314
 What name to give your makefile 315
 Including other makefiles..... 316
 The **MAKEFILES** variable 318
 How makefiles are remade..... 319
 Overriding part of another makefile..... 321
Writing rules 323
 Rule syntax..... 325
 Using wildcard characters in file names 326
 Wildcard examples..... 327
 Pitfalls of using wildcards 328
 The **wildcard** function 329
 Searching directories for dependencies..... 330
 VPATH: *search path for all dependencies* 330
 The **vpath** directive..... 330
 Writing shell commands with directory search 331
 Directory search and implicit rules..... 332
 Directory search for link libraries..... 332

Phony targets	334
Rules without commands or dependencies	336
Empty target files to record events.....	337
Special built-in target names	338
Multiple targets in a rule	340
Multiple rules for one target.....	341
Static pattern rules.....	342
<i>Syntax of static pattern rules</i>	342
<i>Static pattern rules versus implicit rules</i>	344
Double-colon rules	345
Generating dependencies automatically.....	346
Writing the commands in rules	349
Command echoing	351
Command execution.....	352
Parallel execution	353
Errors in commands	355
Interrupting or killing the make utility.....	357
Recursive use of the make utility	358
<i>How the MAKE variable works</i>	358
<i>Communicating variables to a sub-make utility</i>	359
<i>Communicating options to a sub-make utility</i>	361
<i>The '--print-directory' option</i>	362
Defining canned command sequences	364
Using empty commands	366
How to use variables	367
Basics of variable references.....	369
The two flavors of variables.....	370
Advanced features for reference to variables.....	373
<i>Substitution references</i>	373
<i>Computed variable names</i>	373
How variables get their values	377
Setting variables	378
Appending more text to variables	379
The override directive.....	381
Defining variables verbatim.....	382
Variables from the environment.....	383
Conditional parts of makefiles	385
Example of a conditional	386

Syntax of conditionals.....	388
Conditionals that test flags	391
Functions for transforming text	393
Function call syntax	394
Functions for string substitution and analysis	395
Functions for file names.....	398
The <code>foreach</code> function	401
The <code>origin</code> function	403
The <code>shell</code> function	405
How to run the <code>make</code> utility.....	407
Arguments to specify the makefile	408
Arguments to specify the goals	409
Instead of executing the commands	411
Avoiding recompilation of some files.....	413
Overriding variables.....	414
Testing the compilation of a program	415
Summary of <code>make</code> options.....	417
Implicit rules.....	423
Using implicit rules	425
Catalogue of implicit rules	427
Variables used by implicit rules	432
Chains of implicit rules	435
Defining and redefining pattern rules	436
<i>Fundamentals of pattern rules.....</i>	<i>436</i>
<i>Pattern rule examples</i>	<i>437</i>
<i>Automatic variables</i>	<i>438</i>
<i>How patterns match.....</i>	<i>441</i>
<i>Match-anything pattern rules</i>	<i>441</i>
<i>Canceling implicit rules.....</i>	<i>443</i>
Defining last-resort default rules.....	444
Old-fashioned suffix rules.....	445
Implicit rule search algorithm	447
Using <code>make</code> to update archive files	449
Archive members as targets	450
Implicit rule for archive member targets.....	451
<i>Updating archive symbol directories.....</i>	<i>451</i>
Dangers when using archives.....	453
Suffix rules for archive files.....	454

<i>Summary of the features for the GNU make utility</i>	455
<i>GNU make's incompatibilities and missing features</i>	459
Problems and bugs with make tools.....	461
<i>Makefile conventions</i>	463
General conventions for makefiles.....	464
Utilities in makefiles	465
Standard targets for users	466
Variables for specifying commands.....	470
Variables for installation directories	471
<i>GNU make quick reference</i>	477
<i>Complex makefile example</i>	483

Using diff & patch

<i>Overview of diff & patch, the compare & merge tools</i>	491
<i>What comparison means</i>	493
Hunks	495
Suppressing differences in blank and tab spacing.....	496
Suppressing differences in blank lines.....	497
Suppressing case differences.....	498
Suppressing lines matching a regular expression.....	499
Summarizing which files differ.....	500
Binary files and forcing text comparisons	501
<i>diff output formats</i>	503
Two sample input files	504
Showing differences without context.....	505
<i>Detailed description of normal format</i>	505
<i>An example of normal format</i>	506
Showing differences in their context.....	507
<i>Context format</i>	507
<i>Unified format</i>	509
<i>Showing which sections differences are in</i>	510
<i>Showing alternate file names</i>	512
Showing differences side by side.....	513
Controlling side by side format.....	514
<i>An example of side by side format</i>	514
<i>Making edit scripts</i>	514
<i>ed scripts</i>	515
<i>Forward ed scripts</i>	516

<i>RCS scripts</i>	516
Merging files with if-then-else.....	518
<i>Line group formats</i>	518
<i>Line formats</i>	521
<i>Detailed description of if-then-else format</i>	523
<i>An example of if-then-else format</i>	523
Comparing directories	525
Making <code>diff</code> output prettier	527
Preserving tabstop alignment.....	528
Paginating <code>diff</code> output.....	529
<code>diff</code> performance tradeoffs	531
Comparing three files	533
A third sample input file.....	534
Detailed description of <code>diff3</code> normal format.....	535
<code>diff3</code> hunks.....	536
An example of <code>diff3</code> normal format.....	537
Merging from a common ancestor	539
Selecting which changes to incorporate.....	541
Marking conflicts.....	542
Generating the merged output directly.....	544
How <code>diff3</code> merges incomplete lines.....	545
Saving the changed file.....	546
<code>sdiff</code> interactive merging	547
Specifying <code>diff</code> options to the <code>sdiff</code> utility.....	548
Merge commands.....	549
Merging with the <code>patch</code> utility	551
Selecting the <code>patch</code> input format.....	553
Applying imperfect patches.....	554
<i>Applying patches with changed white space</i>	554
<i>Applying reversed patches</i>	554
Helping <code>patch</code> find inexact matches.....	555
Removing empty files.....	557
Multiple patches in a file.....	558
Messages and questions from the <code>patch</code> utility.....	559
Tips for making distributions with patches	561
Invoking the <code>cmp</code> utility	563
<code>cmp</code> options.....	564

<i>Invoking the diff utility</i>	565
<i>diff</i> options	566
<i>Invoking the diff3 utility</i>	573
<i>diff3</i> options	574
<i>Invoking the patch utility</i>	577
Applying patches in other directories	579
Backup file names	580
Naming reject files	582
<i>patch</i> options	583
<i>Invoking the sdiff utility</i>	587
<i>sdiff</i> options	588
<i>Incomplete lines</i>	591
<i>Future projects for diff and patch utilities</i>	593
Suggested projects for improving GNU <i>diff</i> and <i>patch</i> utilities	594
<i>Handling changes to the directory structure</i>	594
<i>Files that are neither directories nor regular files</i>	594
<i>File names that contain unusual characters</i>	594
<i>Arbitrary limits</i>	595
<i>Handling files that do not fit in memory</i>	595
<i>Ignoring certain changes</i>	595
<i>Reporting bugs</i>	595

Using info

<i>Overview of info, the GNU online documentation</i>	599
Using the <i>info</i> program	600
<i>Reading info files</i>	601
GNU <i>info</i> command line options.....	602
Moving the cursor	604
Moving text within a window	606
Selecting a new node.....	607
Searching an <i>info</i> file.....	609
Selecting cross references	610
<i>Parts of a cross reference</i>	610
<i>Selecting xrefs</i>	610
Manipulating multiple windows	612
<i>The mode line</i>	612
<i>Window commands</i>	612
<i>The Echo Area</i>	613
Printing out nodes	616

Miscellaneous commands	617
Manipulating variables.....	619
<i>Making info files from Texinfo files</i>	623
Controlling paragraph formats	624
<code>makeinfo</code> command line options	625
What makes a valid <code>info</code> file?	627
Defaulting the <code>Prev</code> , <code>Next</code> , and <code>Up</code> pointers	628
Index	631

Overview of GNUPro Utilities

The following list details the contents of the *GNUPro Utilities* documentation.

- “Using as”
(see “Overview of `as`, the GNU assembler” on page 5)
- “Using ld”
(see “Overview of `ld`, the GNU linker” on page 187)
- “Using binutils”
(see “Overview of `binutils`, the GNU binary utilities” on page 255)
- “Using make”
(see “Overview of `make`, a program for recompiling” on page 301)
- “Using diff & patch”
(see “Overview of `diff` & `patch`, the compare & merge tools” on page 491)
- “Using info”
(see “Overview of `info`, the GNU online documentation” on page 599)

Using as

Copyright © 1991-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the documentation entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Free Software Foundation

59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

This documentation has been prepared by Cygnus Technical Publications; contact the Cygnus Technical Publications staff: doc@cygnus.com.

Overview of `as`, the GNU assembler

The following documentation serves as a user guide to the GNU assembler, `as`. It describes what you need to know to use GNU `as`, covering the syntax expected in source files, including symbols, constants, expressions, and the general directives. This documentation is not intended as an introduction to programming in assembly language—let alone programming in general.

- “Invoking `as`, the GNU assembler” on page 7
- “Command-line options” on page 15
- “Syntax” on page 33
- “Sections and relocation” on page 43
- “Symbols for the GNU assembler” on page 51
- “Expressions” on page 59
- “Assembler directives” on page 65
- “Machine dependent features” on page 89
- “Acknowledgments for the GNU assembler” on page 181

See also “Overview of special features for `as`” on page 6 for where to find documentation concerning machine-dependent assembler issues.

Overview of special features for `as`

The following documentation points to some of the special machine-dependent features of the assembler.

This is not an attempt to introduce each machine's architecture, and neither is it a description of the instruction set, standard mnemonics, registers or addressing modes that are standard to the specific architecture. Consult the manufacturer's machine architecture documentation for such information.

- “AMD 29K dependent features” on page 91
- “ARC dependent features” on page 97
- “ARM dependent features” on page 101
- “AT&T/Intel dependent features” on page 105
- “D10V dependent features” on page 113
- “H8/300 dependent features” on page 121
- “H8/500 dependent features” on page 127
- “HPPA dependent features” on page 133
- “SH dependent features” on page 139
- “Intel 960 dependent features” on page 145
- “M68K dependent features” on page 151
- “MIPS dependent features” on page 161
- “SPARC dependent features” on page 169
- “Vax dependent features” on page 173

Invoking `as`, the GNU assembler

The GNU Assembler, `as`, is really a family of assemblers. If you use (or have used) the GNU assembler on one architecture, you should find a fairly similar environment when you use it on another architecture. Each version has much in common with the others, including object file formats, most assembler directives (often called *pseudo-ops*) and assembler syntax.

See the following documentation for more discussion of the GNU assembler.

- “Summary of invocation of `as` options” on page 8
- “Object file formats” on page 9
- “Command line” on page 10
- “Input files” on page 11
- “Output (object) file” on page 12
- “Error and warning messages” on page 13

`as` is primarily intended to assemble the output of the GNU C compiler, `gcc`, for use by the GNU linker, `ld`. Nevertheless, we’ve tried to make `as` assemble correctly everything that other assemblers for the same machine would assemble.

Any exceptions are documented explicitly (for specific processors and their families, see “Machine dependent features” on page 89). This doesn’t mean `as` always uses the same syntax as another assembler for the same architecture; for example, we know of several incompatible versions of Motorola 680x0 assembly language syntax.

Summary of invocation of `as` options

The following example summarizes how to invoke `as`. For further details, see “Command-line options” on page 15.

```
as [ -a[cdhlms][=file] ] [ -D ] [ --defsym sym=value ]
    [ -f ] [ --gstabs ] [ --help ] [ -I dir ] [ -J ] [ -K ] [ -L ]
    [ -o objfile ] [ -R ] [ --statistics ] [ -v ] [ -version ]
    [ --version ] [ -W ] [ -w ] [ -x ] [ -Z ]
    [ -mbig-endian | -mlittle-endian ]
    [ -m[arm]1 | -m[arm]2 | -m[arm]250 | -m[arm]3 | -m[arm]6
      | -m[arm]7[t][d]m[I] ]
    [ -m[arm]v2 | -m[arm]v2a | -m[arm]v3 | -m[arm]v3m | -m[arm]v4
      | -m[arm]v4t ]

    [ -mthumb | -mall ]
    [ -mfpa10 | -mfpa11 | -mfpe-old \ -mno-fpu
    [ -EB | -EL ]
    [ -mapcs-32 | -mapcs-26 ]
    [ -O ]
    [ -Av6 | -Av7 | -Av8 | -Asparclite | -Av9 | -Av9a ]
    [ -xarch=v8plus | -xarch=v8plusa ] [ -bump ]
    [ -ACA | -ACA_A | -ACB | -ACC | -AKA | -AKB | -AKC | -AMC ]
    [ -b ] [ -no-relax ]
    [ -l ] [ -m68000 | -m68010 | -m68020 | ... ]
    [ -nocpp ] [ -EL ] [ -EB ] [ -G num ] [ -mcpu=CPU ]
    [ -mips1 ] [ -mips2 ] [ -mips3 ] [ -m4650 ] [ -no-m4650 ]
    [ --trap ] [ --break ]
    [ --emulation= name ]
    [ -- | files... ]
```

Unlike older assemblers, `as` is designed to assemble a source program in one pass of the source file. This has a subtle impact on the `.org` directive (see “`.org new-lc`, fill directive” on page 78).

Object file formats

The GNU assembler can be configured to produce several alternative object file formats. For the most part, this does not affect how you write assembly language programs; but directives for debugging symbols are typically different in different file formats. See “Symbol attributes” on page 57.

Command line

After the program name `as`, the command line may contain options and file names. Options may appear in any order, and may be before, after, or between file names. The order of file names is significant.

`--` (two hyphens), by themselves, name the standard input file explicitly, as one of the files for `as` to assemble. Except for `--`, any command line argument that begins with a hyphen (`-`) is an option. Each option changes the behavior of `as`. No option changes the way another option works. An option is a `-` followed by one or more letters; the *case* of the letter is important. All options are optional.

Some options expect exactly one file name to follow them. The file name may either immediately follow the option's letter (compatible with older assemblers) or it may be the next command argument (GNU standard). The following two command lines are equivalent.

```
as -o my-object-file.o mumble.s
as -omy-object-file.o mumble.s
```

Input files

We use the phrase *source program*, abbreviated *source*, to describe the program input to one run of `as`. The program may be in one or more files; how the source is partitioned into files doesn't change the meaning of the source.

The source program is a concatenation of the text in all the files, in the order specified. Each time you run `as`, it assembles exactly one source program. The source program is made up of one or more files. (The standard input is also a file.)

You give `as` a command line that has zero or more input file names. The input files are read (from left file name to right). A command line argument (in any position) that has no special meaning is taken to be an input file name.

If you give `as` no file names it attempts to read one input file from the `as` standard input, which is normally your terminal. You may have to type Ctrl-D to tell `as` there is no more program to assemble.

Use `--` if you need to explicitly name the standard input file in your command line.

If the source is empty, `as` produces a small, empty object file.

Filename and line-numbers

There are two ways of locating a line in the input file (or files) and either may be used in reporting error messages. One way refers to a line number in a physical file; the other refers to a line number in a "logical" file. See "Error and warning messages" on page 13.

Physical files are those files named in the command line given to `as`.

Logical files are simply names declared explicitly by assembler directives; they bear no relation to physical files. Logical file names help error messages reflect the original source file, when `as` source is itself synthesized from other files. See ".app-file string directive" on page 68.

Output (object) file

Every time you run `as`, it produces an output file, which is your assembly language program translated into numbers. This file is the object file. Its default name is `a.out`, or `b.out`, when `as` is configured for the Intel 80960. You can give it another name by using the `-o` option. Conventionally, object file names end with `.o`. The default name is used for historical reasons: older assemblers were capable of assembling self-contained programs directly into a runnable program. (For some formats, this isn't currently possible, but it can be done for the `a.out` format.)

The object file is meant for input to the linker `ld`. It contains assembled program code, information to help `ld` integrate the assembled program into a runnable file, and (optionally) symbolic information for the debugger.

Error and warning messages

`as` may write warnings and error messages to the standard error file (usually your terminal). This should not happen when a compiler runs `as` automatically. Warnings report an assumption made so that `as` could keep assembling a flawed program; errors report a grave problem that stops the assembly. Warning messages have the following format (where *NNN* is a line number).

```
file_name:NNN:Warning Message Text
```

If a logical file name has been given (see “.app-file string directive” on page 68) it is used for the filename, otherwise the name of the current input file is used. If a logical line number was given (see “.line line-number directive” on page 75) then it is used to calculate the number printed, otherwise the actual line in the current source file is printed. The message text is intended to be self explanatory (in the Unix tradition).

Error messages have the following format.

```
file_name:NNN:FATAL>Error Message Text
```

The file name and line number are derived as for warning messages. The actual message text may be rather less explanatory because many of them aren’t supposed to happen.

Command-line options

The following documentation describes command-line options available in all versions of the GNU assembler; see “Summary of invocation of as options” on page 8 for a complete invocation declaration example.

- “Enable listings: -a[cdhlms] options” on page 17
- “-D option” on page 18
- “Work faster: -f option” on page 19
- “.include search path: -I path option” on page 20
- “Difference tables: -K option” on page 21
- “Include local labels: -L option” on page 22
- “Assemble in MRI compatibility mode: -M option” on page 23
- “Dependency tracking: --MD option” on page 26
- “Name the object file: -o option” on page 27
- “Join data and text sections: -R option” on page 28
- “Display assembly statistics: --statistics option” on page 29
- “Announce version: -v option” on page 30
- “Suppress warnings: -W option” on page 31
- “Generate object file in spite of errors: -Z option” on page 32

See also “Machine dependent features” on page 89 in order to find documentation

specific to each architecture's particular options.

If you are invoking `as` using the GNU C compiler (version 2), you can use the `-Wa` option to pass arguments through to the assembler. The assembler arguments must be separated from each other (and the `-Wa`) by commas. Use the following for example.

```
gcc -c -g -O -Wa,-alh,-L file.c
```

This input emits a listing to standard output with high-level and assembly source. Usually you do not need to use this `-Wa` mechanism, since many compiler command-line options are automatically passed to the assembler by the compiler. (You can call the GNU compiler driver with the `-v` option to see precisely what options it passes to each compilation pass, including the assembler.)

Enable listings: `-a[cdhlns]` options

These options enable listing output from the assembler. By itself, `-a` requests high-level, assembly, and symbols listing. You can use other letters to select specific options for the list: `-ah` requests a high-level language listing, `-al` requests an output-program assembly listing, and `-as` requests a symbol table listing. High-level listings require that a compiler debugging option like `-g` be used, and that also assembly listings (`-al`) be requested. Use the `-ad` option to omit debugging directives from the listing.

Use the `-ac` option to omit false conditionals from a listing. Any lines which are not assembled because of a false `.if` (or `.ifdef`, or any other conditional), or a true `.if` followed by a `.else`, will be omitted from the listing.

Use the `-ad` option to omit debugging directives from the listing.

Once you have specified one of these options, you can further control listing output and its appearance using the directives `.list`, `.nolist`, `.psize`, `.eject`, `.title`, and `.sbttl`. The `-an` option turns off all forms processing. If you do not request listing output with one of the `-a` options, the listing-control directives have no effect. The letters after `-a` may be combined into one option, such as `-aln`.

-D option

This option has no effect whatsoever, but it is accepted to make it more likely that scripts written for other assemblers also work with `as`.

Work faster: `-f` option

`-f` should only be used when assembling programs written by a (trusted) compiler. `-f` stops the assembler from doing whitespace and comment preprocessing on the input file(s) before assembling them. See “Preprocessing” on page 34.

WARNING: If you use `-f` when the files actually need to be pre-processed (if they contain comments, for example), as does not work correctly.

`.include search path: -I path option`

.include search path: -I *path* option

Use this option to add a *path* to the list of directories as searches for files specified in `.include` directives (see “`.include` “file” directive” on page 74). You may use `-I` as many times as necessary to include a variety of paths. The current working directory is always searched first; after that, as searches any `-I` directories in the same order as they were specified (left to right) on the command line.

Difference tables: -K option

as sometimes alters the code emitted for directives of the form `.word sym1-sym2`; see “`.word expressions directive`” on page 86. You can use the `-K` option if you want a warning issued when this is done.

Include local labels: `-L` option

Labels beginning with `L` (upper case only) are called *local labels*. See “Symbol names” on page 54. Normally, you do not see such labels when debugging, because they are intended for the use of programs (like compilers) that compose assembler programs, not for your notice. Normally both `as` and `ld` discard such labels, so you do not normally debug with them.

This option tells `as` to retain those `L . . .` symbols in the object file. Usually if you do this you also tell the linker, `ld`, to preserve symbols whose names begin with `L`. By default, a local label is any label beginning with `L`, but each target is allowed to redefine the local label prefix. On the HPPA, local labels begin with `L$; ;` for the ARM family.

Assemble in MRI compatibility mode: -M option

The `-M` or `--mri` option selects Microtec Research Inc.'s (MRI's) compatibility mode. This changes the syntax and pseudo-op handling of `as` to make it compatible with the ASM68K or the ASM960 (depending upon the configured target) assembler from MRI.

The exact nature of the MRI syntax will not be documented here; see the MRI manuals for more information.

The purpose of this option is to permit assembling existing MRI assembler code using `as`.

The MRI compatibility is not complete. Certain operations of the MRI assembler depend upon its object file format, and can not be supported using other object file formats.

Supporting these would require enhancing each object file format individually. These are the following.

- *global symbols in common section*
The m68k MRI assembler supports common sections which are merged by the linker. Other object file formats do not support this. `as` handles common sections by treating them as a single common symbol. It permits local symbols to be defined within a common section, but it can not support global symbols, since it has no way to describe them.
- *complex relocations*
The MRI assemblers support relocations against a negated section address, and relocations which combine the start addresses of two or more sections. These are not supported by other object file formats.
- *END pseudo-op specifying start address*
The MRI `END` pseudo-op permits the specification of a start address. This is not supported by other object file formats. The start address may instead be specified using the `-e` option to the linker, or in a linker script.
- *IDNT, .ident and NAME pseudo-ops*
The MRI `IDNT`, `.ident` and `NAME` pseudo-ops assign a module name to the output file. This is not supported by other object file formats.
- *ORG pseudo-op*
The m68k MRI `ORG` pseudo-op begins an absolute section at a given address. This differs from the usual `as .org` pseudo-op, which changes the location within the

current section. Absolute sections are not supported by other object file formats.

The address of a section may be assigned within a linker script.

There are some other features of the MRI assembler which are not supported by `as`, typically either because they are difficult or because they seem of little consequence. Some of these may be supported in future releases.

- *EBCDIC strings*
EBCDIC strings are not supported.
- *Packed binary coded decimal*
Packed binary coded decimal is not supported. This means that the `DC.P` and `DCB.P` pseudo-ops are not supported.
- *FEQU pseudo-op*
The m68k `FEQU` pseudo-op is not supported.
- *NOOBJ pseudo-op*
The m68k `NOOBJ` pseudo-op is not supported.
- *OPT branch control options*
The m68k `OPT` branch control options—`B`, `BRS`, `BRB`, `BRL`, and `BRW`—are ignored. `as` automatically relaxes all branches, whether forward or backward, to an appropriate size, so these options serve no purpose.
- *OPT list control options*
The following m68k `OPT` list control options are ignored: `C`, `CEX`, `CL`, `CRE`, `E`, `G`, `I`, `M`, `MEX`, `MC`, `MD`, `X`.
- *Other OPT options*
The following m68k `OPT` options are ignored: `NEST`, `O`, `OLD`, `OP`, `P`, `PCO`, `PCR`, `PCS`, `R`.
- *OPT D option is default*
The m68k `OPT D` option is the default, unlike the MRI assembler. `OPT NOD` may be used to turn it off.
- *XREF pseudo-op*
The m68k `XREF` pseudo-op is ignored.
- *.debug pseudo-op*
The i960 `.debug` pseudo-op is not supported.
- *.extended pseudo-op*
The i960 `.extended` pseudo-op is not supported.
- *.list pseudo-op*
The various options of the i960 `.list` pseudo-op are not supported.
- *.optimize pseudo-op*
The i960 `.optimize` pseudo-op is not supported.

- `.output` pseudo-op
The i960 `.output` pseudo-op is not supported.
- `.setreal` pseudo-op
The i960 `.setreal` pseudo-op is not supported.

Dependency tracking: --MD option

`as` can generate a dependency file for the file that it creates. This file consists of a single rule suitable for `make`, describing the dependencies for the main source file.

The rule is written to the file named in its argument.

This feature is used in the automatic updating of files.

Name the object file: `-o` option

There is always one object file output when you run `as`. By default it has the name `a.out` (or `b.out`, for Intel 960 targets only). You use this option (which takes exactly one filename) to give the object file a different name.

Whatever the object file is called, `as` overwrites any existing file of the same name.

Join data and text sections: `-R` option

`-R` tells `as` to write the object file as if all data-section data lives in the text section. This is only done at the very last moment: your binary data are the same, but data section parts are relocated differently. The data section part of your object file is zero bytes long because all its bytes are appended to the text section. (See “Sections and relocation” on page 43.) When you specify `-R`, it would be possible to generate shorter address displacements (because we do not have to cross between text and data section). We refrain from doing this simply for compatibility with older versions of `as`. In future, `-R` may work this way.

When `as` is configured for COFF output, this option is only useful if you use sections named `.text` and `.data`.

`-R` is not supported for any of the HPPA targets. Using `-R` generates a warning from `as`.

Display sssembly statistics: `--statistics` option

Use `--statistics` to display two statistics about the resources used by `as`: the maximum amount of space allocated during the assembly (in bytes), and the total execution time taken for the assembly (in CPU seconds).

Announce version: `-v` option

You can find out what version of `as` is running by including the `-v` option (which you can also spell as `-version`) on the command line.

Suppress warnings: -W option

`as` should never give a warning or error message when assembling compiler output. However, programs written by people often cause `as` to give a warning that a particular assumption was made. All such warnings are directed to the standard error file. If you use this option, `-w`, no warnings are issued. This option only affects the warning messages and it does not change any particular of how `as` assembles your file. Errors that stop the assembly are still reported.

Generate object file in spite of errors: -Z option

After an error message, `as` normally produces no output. If for some reason you are interested in object file output even after `as` gives an error message on your program, use the `-z` option. If there are any errors, `as` continues, and writes an object file after a final warning message of the following form.

```
n errors, m warnings, generating bad object file.
```

Syntax

The following documentation describes the machine-independent syntax allowed in a source file.

- “Preprocessing” on page 34
- “Comments” on page 35
- “Symbols” on page 36
- “Statements” on page 37
- “Constants” on page 38

`as` syntax is similar to what many other assemblers use; it is inspired by the BSD 4.2 assembler, except that `as` does not assemble Vax bit-fields.

Preprocessing

The `as` internal preprocessor has the following functionality.

- Adjusts and removes extra whitespace. It leaves one space or tab before the keywords on a line, and turns any other whitespace on the line into a single space.
- Removes all comments, replacing them with a single space, or an appropriate number of newlines.
- Converts character constants into the appropriate numeric values.

It does not do macro processing, include file handling, or anything else you may get from your C compiler's preprocessor. You can do include file processing with the `.include` directive (see “.include “file” directive” on page 74). You can use the GNU C compiler driver to get other `c++`-style preprocessing by giving the input file a `.s` suffix. See “Options Controlling the Kind of Output” in *Using GNU CC* in **GNUPro Compiler Tools**.

Excess whitespace, comments, and character constants cannot be used in the portions of the input text that are not preprocessed.

If the first line of an input file is `#NO_APP` or if you use the `-f` option, whitespace and comments are not removed from the input file. Within an input file, you can ask for whitespace and comment removal in specific portions of the by putting a line that says `#APP` before the text that may contain whitespace or comments, and putting a line that says `#NO_APP` after this text. This feature is mainly intend to support `asm` statements in compilers whose output is otherwise free of comments and whitespace.

Whitespace

Whitespace is one or more blanks or tabs, in any order. Whitespace is used to separate symbols, and to make programs neater for people to read. Unless within character constants (see “Character constants” on page 38), any whitespace means the same as exactly one space.

Comments

There are two ways of rendering comments to `as`. In both cases the comment is equivalent to one space.

Anything from `/*` through the next `*/` is a comment, as in the following statement.

```
/*
    The only way to include a newline ('\n') in a comment
    is to use this sort of comment.
*/
```

This means you may not nest the following types of comments.

```
/* This sort of comment does not nest. */
```

Anything from the *line comment* character to the next newline is considered a comment and is ignored.

The line comment character is `#` for the Vax; `#` for the i960; `!` for the SPARC family; `|` for the Motorola 68K family; `;` for the AMD 29K family; `;` for the H8/300 family; `!` for the H8/500 family; `;` for the HPPA family; `!` for the Hitachi SH family; `!` for the Z8000 family, and `;` for the ARC family; see “Machine dependent features” on page 89 to locate specific families.

On some machines there are two different line comment characters. One character only begins a comment if it is the first non-whitespace character on a line, while the other always begins a comment.

To be compatible with past assemblers, lines that begin with `#` have a special interpretation. Following the `#` should be an absolute expression (see “Expressions” on page 59): the logical line number of the *next* line. Then a string (see “Strings” on page 38) is allowed: if present, it is a new logical file name. The rest of the line, if any, should be whitespace.

If the first non-whitespace characters on the line are not numeric, the line is ignored. (Just like a comment.)

```
        # This is an ordinary comment.
# 42-6 "new_file_name"# New logical file name
        # This is logical line # 36.
```

This feature is deprecated, and may disappear from future versions of `as`.

Symbols

A *symbol* is one or more characters chosen from the set of all letters (both upper and lower case), digits and the three characters ‘_ . \$’. On most machines, you can also use \$ in symbol names; exceptions are noted for each architecture; see “Machine dependent features” on page 89 for specific processor families.

WARNING: No symbol may begin with a digit. Case is significant. There is no length limit: all characters are significant.

Symbols are delimited by characters not in that set, or by the beginning of a file (since the source program must end with a newline, the end of a file is not a possible symbol delimiter). See “Symbols for the GNU assembler” on page 51.

Statements

A statement ends at a newline character (`\n`) or line separator character. (The line separator is usually `(:)`, unless this conflicts with the comment character. Exceptions are noted for each architecture; see “Machine dependent features” on page 89.) The newline or separator character is considered part of its preceding statement.

NOTE: Newlines and separators within character constants are an exception: *they do not end statements.*

It is an error to end any statement with end-of-file: the last character of any input file should be a newline.

You may write a statement on more than one line if you put a backslash (`\`) immediately in front of any newlines within the statement. When `as` reads a backslashed newline both characters are ignored. You can even put backslashed newlines in the middle of symbol names without changing the meaning of your source program.

An empty statement is allowed, and may include whitespace. It is ignored.

A statement begins with zero or more labels, optionally followed by a key symbol which determines what kind of statement it is. The key symbol determines the syntax of the rest of the statement. If the symbol begins with a dot (`.`), then the statement is an assembler directive: typically valid for any computer. If the symbol begins with a letter the statement is an assembly language instruction: it assembles into a machine language *instruction*. Different versions of `as` for different computers recognize different instructions. In fact, the same symbol may represent a different instruction in a different computer’s assembly language.

A label is a symbol immediately followed by a colon (`:`). Whitespace before a label or after a colon is permitted, but you may not have whitespace between a label’s symbol and its colon. See “Labels” on page 52.

For HPPA targets, labels need not be immediately followed by a colon, but the definition of a label must begin in column zero, or the beginning of the line. This also implies that only one label may be defined on each line. Use the following statement as an example.

```
label:      .directivefollowed by something
another_label:      # This is an empty statement.
                    instruction    operand_1, operand_2, ...
```

Constants

A constant is a number, written so that its value is known by inspection, without knowing any context, like the following input example.

```
.byte    74, 0112, 092, 0x4A, 0X4a, 'J', '\J' # All the same value.
.ascii   "Ring the bell\7"                    # A string constant.
.octa     0x123456789abcdef0123456789ABCDEF0 # A bignum.
.float    0f-314159265358979323846264338327\
95028841971.693993751E-40                     # - pi, a flonum.
```

Character constants

There are two kinds of character constants. A character stands for one character in one byte and its value may be used in *numeric expressions*. *String constants* (properly called *string literals*) are potentially many bytes and their values may not be used in *arithmetic expressions*.

Strings

A string is written between double-quotes. It may contain double-quotes or null characters.

The way to get special characters into a string is to escape these characters: precede them with a backslash (\) character. For example \\ represents one backslash: the first \ is an escape which tells us to interpret the second character literally as a backslash (which prevents us from recognizing the second \ as an escape character).

The complete list of escapes follows.

```
\b
    Mnemonic for backspace; for ASCII this is octal code 010.
\f
    Mnemonic for FormFeed; for ASCII this is octal code 014.
\n
    Mnemonic for newline; for ASCII this is octal code 012.
\r
    Mnemonic for carriage-Return; for ASCII this is octal code 015.
\t
    Mnemonic for horizontal Tab; for ASCII this is octal code 011.
\ digit digit digit
    An octal character code. The numeric code is 3 octal digits. For compatibility with
    other Unix systems, 8 and 9 are accepted as digits: for example, \008 has the
    value 010, and \009 the value 011.
```

`\x hex-digit hex-digit`

A hex character code. The numeric code is 2 hexadecimal digits. Either upper or lower case `x` works.

`\\`

Represents one `\` character.

`\"`

Represents one `"` character. Needed in strings to represent this character, because an unescaped `"` would end the string.

`\ anything-else`

Any other character when escaped by `\` gives a warning, but assembles as if the `\` was not present. The idea is that if you used an escape sequence you clearly didn't want the literal interpretation of the following character. However `as` has no other interpretation, so `as` knows it is giving you the wrong code and warns you of the fact.

Which characters are escapable, and what those escapes represent, varies widely among assemblers. The current set is what we think the BSD 4.2 assembler recognizes, and is a subset of what most C compilers recognize. If you are in doubt, do not use an escape sequence.

Characters

A single character may be written as a single quote immediately followed by that character. The same escapes apply to characters as to strings. So if you want to write the character backslash, you must write `'\\` where the first `\` escapes the second `\`. As you can see, the quote is an 'acute' accent, not a 'grave' accent (```). A newline immediately following an acute accent is taken as a literal character and does not count as the end of a statement. The value of a character constant in a numeric expression is the machine's byte-wide code for that character. `as` assumes your character code is ASCII: `'A` means 65, `'B` means 66, and so on.

Number constants

`as` distinguishes three kinds of numbers according to how they are stored in the target machine. *Integers* are numbers that would fit into an int in the C language. *Bignums* are integers, but they are stored in more than 32 bits. *Flonums* are floating point numbers; see "Flonums" on page 40.

Integers

A binary integer is `0b` or `0B` followed by one or more of the binary digits, 01.

An octal integer is `0` followed by zero or more of the octal digits (0, 1, 2, 3, 4, 5, 6, 7).

0b invalid

0b0 valid

A decimal integer starts with a non-zero digit followed by zero or more digits (0, 1, 2, 3, 4, 5, 6, 7, 8, 9).

A hexadecimal integer is 0x, or 0X, followed by zero or more hexadecimal digits chosen from 0123456789abcdefABCDEF. Integers have the usual values. To denote a negative integer, use the prefix operator (-), discussed under expressions (see “Prefix operators” on page 63).

0x valid=0x0

0x1 valid

Bignums

A *bignum* has the same syntax and semantics as an integer except that the number (or its negative) takes more than 32 bits to represent in binary. The distinction is made because in some places integers are permitted while bignums are not.

Flonums

A *flonum* represents a floating point number. The translation is indirect: a decimal floating point number from the text is converted by `as` to a generic binary floating point number of more than sufficient precision. This generic floating point number is converted to a particular computer’s floating point format (or formats) by a portion of `as` specialized to that computer.

A flonum is written by writing (in order) the following:

- The digit 0. (0 is optional on the HPPA.)
- A letter, to tell `as` the rest of the number is a flonum. `e` is recommended. Case is not important.
 - On the H8/300, H8/500, Hitachi SH, and AMD 29K architectures, the letter must be one of the letters: `D`, `F`, `P`, `R`, `S`, or `X` (in uppercase or lower case).
 - On the ARC, the letter must be one of the letters: `D`, `F`, `R`, or `S` (in uppercase or lowercase).
 - On the Intel 960 architecture, the letter must be one of the letters: `D`, `F`, or `T` (in uppercase or lower case).
 - On the HPPA architecture, the letter must be `E` (uppercase only).
- An optional sign: either `+` or `-`.
- An optional integer part: zero or more decimal digits.

- An optional fractional part: (.) followed by zero or more decimal digits.
- An optional exponent, consisting of the following elements:
 - An E or e.
 - Optional sign: either + or -.
 - One or more decimal digits.

At least one of the integer part or the fractional part must be present. The floating point number has the usual base-10 value.

`as` does all processing using integers. Flonums are computed independently of any floating point hardware in the computer running `as`.

Sections and relocation

A *section* is a range of addresses, with no gaps; all data in those addresses is treated the same for some particular purpose (for example, there may be a *read-only* section).

The following documentation discusses sections and their relocation.

- “ld sections” on page 46
- “as internal sections” on page 47
- “Sub-sections” on page 48
- “bss section” on page 50

The linker, `ld`, reads many object files (partial programs) and combines their contents to form a runnable program. When `as` emits an object file, the partial program is assumed to start at address, 0. `ld` assigns the final addresses for the partial program, so that different partial programs do not overlap. This is actually an oversimplification of relocation, but it suffices to explain how `as` uses sections.

`ld` moves blocks of bytes of your program to their run-time addresses. These blocks slide to their run-time addresses as rigid units; their length does not change and neither does the order of bytes within them. Such a rigid unit is called a *section*. Assigning run-time addresses to sections is called *relocation*. It includes the task of adjusting mentions of object-file addresses so they refer to the proper run-time addresses. For the H8/300 and H8/500, and for the Hitachi SH, `as` pads sections if needed to ensure they end on a word (sixteen bit) boundary.

An object file written by `as` has at least three sections, any of which may be empty. These are named *text*, *data* and *bss* sections.

When it generates COFF output, `as` can also generate whatever other named sections you specify using the `.section` directive (see “`.section` name directive” on page 80). If you do not use any directives that place output in the `.text` or `.data` sections, these sections still exist, but are empty.

When `as` generates SOM or ELF output for the HPPA, `as` can also generate whatever other named sections you specify using the `.space` and `.subspace` directives.

See *HP9000 Series 800 Assembly Language Reference Manual* (HP 92432-90001) for details on the `.space` and `.subspace` assembler directives.

Additionally, `as` uses different names for the standard text, data, and bss sections when generating SOM output. Program text is placed into the `$CODE$` section, data into `$DATA$`, and BSS into `BSS`.

Within the object file, the text section starts at address 0, the data section follows, and the bss section follows the data section.

When generating either SOM or ELF output files on the HPPA, the text section starts at address 0, the data section at address `0x4000000`, and the bss section follows the data section.

To let `ld` know which data changes when the sections are relocated, and how to change that data, `as` also writes to the object file details of the relocation needed. To perform relocation, each time an address in the object file is mentioned, `ld` must know the following criteria.

- Where in the object file is the beginning of this reference to an address?
- How long (in bytes) is this reference?
- Which section does the address refer to? What is the numeric value of $(\text{address}) - (\text{start-address of section})$?
- Is the reference to an address “Program-Counter relative”?

In fact, every address `as` ever uses is expressed as the following.

$(\text{section}) + (\text{offset into section})$.

Further, most expressions `as` computes have this section-relative nature. (For some object formats, such as SOM for the HPPA, some expressions are symbol-relative instead.) In this documentation, we use the notation `{secname N}` to mean “offset *N* into a specified section, *secname*.”

Apart from text, data and bss sections you need to know about the *absolute* section. When `ld` mixes partial programs, addresses in the absolute section remain unchanged. For example, address `{absolute 0}` is “relocated” to run-time address 0 by `ld`. Although the linker never arranges two partial programs’ data sections with overlapping addresses after linking, *by definition* their absolute sections must overlap. Address `{absolute 239}` in one part of a program is always the same address when the program is running as address `{absolute 239}` in any other part of the program. The idea of sections is extended to the *undefined* section. Any address whose section is unknown at assembly time is by definition rendered `{undefined v}` – where *v* is filled in later.

Since numbers are always defined, the only way to generate an undefined address is to mention an undefined symbol. A reference to a named common block would be such a symbol: its value is unknown at assembly time so it has section *undefined*.

By analogy the word *section* is used to describe groups of sections in the linked program. `ld` puts all partial programs’ text sections in contiguous addresses in the linked program.

It is customary to refer to the text section of a program, meaning all the addresses of all partial programs’ *text sections*. Likewise for data and bss sections.

Some sections are manipulated by `ld`; others are invented for use by the assembler and have no meaning except during assembly.

ld sections

ld deals with just four kinds of sections, summarized by the following documentation.

named sections

text section

data section

These sections hold your program. as and ld treat them as separate but equal sections. Anything you can say of one section is true for another. When the program is running, however, it is customary for the text section to be unalterable. The text section is often shared among processes: it contains instructions, constants and the like. The data section of a running program is usually alterable; for example, C variables would be stored in the data section.

bss section

This section contains zeroed bytes when your program begins running. It is used to hold uninitialized variables or common storage. The length of each partial program's bss section is important, but because it starts out containing zeroed bytes there is no need to store explicit zero bytes in the object file. The bss section was invented to eliminate those explicit zeros from object files.

absolute section

Address 0 of this section is always *relocated* to runtime address 0. This is useful if you want to refer to an address that ld must not change when relocating. In this sense we speak of absolute addresses being *unrelocatable*: they do not change during relocation.

undefined section

This *section* is a catch-all for address references to objects not in the preceding sections.

An idealized example of three relocatable sections follows. The following example uses the traditional section names `.text` and `.data`. Memory addresses are on the horizontal axis.

Partial program #1

text	data	bss
ttttt	dddd	00

Partial program #2

text	data	bss
TTT	DDDD	000

Linked program

text	data	bss
TTT	ttttt	dddd
		DDDD
		00000 ...

Addresses: 0...

as internal sections

These sections are meant only for the internal use of `as`. They have no meaning at run-time. You do not really need to know about these sections for most purposes; but they can be mentioned in `as` warning messages, so it might be helpful to have an idea of their meanings to `as`. These sections are used to permit the value of every expression in your assembly language program to be a section-relative address.

ASSEMBLER-INTERNAL-LOGIC-ERROR!

An internal assembler logic error has been found. This means there is a bug in the assembler.

expr *section*

The assembler stores complex expression internally as combinations of symbols. When it needs to represent an expression as a symbol, it puts it in the `expr` section.

Sub-sections

Assembled bytes conventionally fall into two sections: text and data. You may have separate groups of data in named sections that you want to end up near to each other in the object file, even though they are not contiguous in the assembler source. `as` allows you to use *subsections* for this purpose. Within each section, there can be numbered subsections with values from 0 to 8192. Objects assembled into the same subsection go into the object file together with other objects in the same subsection. For example, a compiler might want to store constants in the text section, but might not want to have them interspersed with the program being assembled. In this case, the compiler could issue a `.text 0` before each section of code being output, and a `.text 1` before each group of constants being output.

Subsections are optional. If you do not use subsections, everything goes in subsection number zero.

Each subsection is zero-padded up to a multiple of four bytes. (Sub-sections may be padded a different amount on different flavors of `as`.)

Subsections appear in your object file in numeric order, lowest numbered to highest. (All this to be compatible with other people's assemblers.) The object file contains no representation of subsections; `ld` and other programs that manipulate object files see no trace of them. They just see all your text subsections as a text section, and all your data subsections as a data section.

To specify which subsection you want subsequent statements assembled into, use a numeric argument to specify it, in a `.text expression` or a `.data expression` statement. When generating COFF output, you can also use an extra subsection argument with arbitrary named sections: `.section name, expression. expression` should be an absolute expression. (See "Expressions" on page 59.) If you just say `.text`, then `.text 0` is assumed. Likewise, `.data` means `.data 0`. Assembly begins in `text 0`. For instance, use the following example.

```
.text    0      # The default subsection is text 0 anyway.
.ascii   "This lives in the first text subsection. *"
.text    1
.ascii   "But this lives in the second text subsection."
.data    0
.ascii   "This lives in the data section,"
.ascii   "in the first data subsection."
.text    0
.ascii   "This lives in the first text section,"
.ascii   "immediately following the asterisk (*)."
```

Each section has a *location counter* incremented by one for every byte assembled into that section. Because subsections are merely a convenience restricted to `as`, there is no concept of a subsection location counter. There is no way to directly manipulate a location counter—but the `.align` directive changes it, and any label definition captures its current value. The location counter of the section where statements are being assembled is said to be the *active* location counter.

bss section

The bss section is used for local common variable storage. You may allocate address space in the bss section, but you may not dictate data to load into it before your program executes. When your program starts running, all the contents of the bss section are zeroed bytes. Addresses in the bss section are allocated with special directives; you may not assemble anything directly into the bss section. Hence there are no bss subsections.

See “.comm symbol, length directive” on page 69.

Symbols for the GNU assembler

Symbols are a central concept: the programmer uses symbols to name things, the linker uses symbols to link, and the debugger uses symbols to debug. The following documentation discusses more about symbols and the assembler.

- “Labels” on page 52
- “Giving symbols other values” on page 53
- “Symbol names” on page 54
- “The special dot symbol” on page 56
- “Symbol attributes” on page 57

WARNING: `as` does not place symbols in the object file in the same order they were declared. This may break some debuggers.

Labels

A *label* is written as a symbol immediately followed by a colon (:). The symbol then represents the current value of the active location counter, and is, for example, a suitable instruction operand. You are warned if you use the same symbol to represent two different locations: the first definition overrides any other definitions.

On the HPPA, the usual form for a label need not be immediately followed by a colon, but instead must start in column zero. Only one label may be defined on a single line. To work around this, the HPPA version of `as` also provides a special directive, `.label`, for defining labels more flexibly.

Giving symbols other values

A symbol can be given an arbitrary value by writing a symbol, followed by an equal sign (=), followed by an expression (see “Expressions” on page 59). This is equivalent to using the `.set` directive. See “.set symbol, expression directive” on page 81.

Symbol names

Symbol names begin with a letter or with either a dot (`.`) or an underscore (`_`); on most machines, you can also use `$` in symbol names; exceptions are noted for each architecture in the documentation; see “Machine dependent features” on page 89 to locate a specific architecture. That character may be followed by any string of digits, letters, dollar signs (unless exceptions are noted for each architecture), and underscores. For the AMD 29K family, `?` is also allowed in the body of a symbol name, though not at its beginning. Case of letters is significant: `foo` is a different symbol name than `Foo`.

Each symbol has exactly one name. Each name in an assembly language program refers to exactly one symbol. You may use that symbol name any number of times in a program.

Local symbol names

Local symbols help compilers and programmers use names temporarily. There are ten local symbol names, which are re-used throughout the program. You may refer to them using the names `'0'` `'1'` . . . `'9'`. To define a local symbol, write a label of the form `'N:'` (where *N* represents any digit). To refer to the most recent previous definition of that symbol write `'Nb'`, using the same digit as when you defined the label. To refer to the next definition of a local label, write `'Nf'`—where *N* gives you a choice of 10 forward references. The `b` stands for *backwards* and the `f` stands for *forwards*.

Local symbols are not emitted by the current GNU C compiler.

There is no restriction on how you can use these labels, but remember that at any point in the assembly you can refer to at most 10 prior local labels and to at most 10 forward local labels.

Local symbol names are only a notation device. They are immediately transformed into more conventional symbol names before the assembler uses them. The symbol names stored in the symbol table, appearing in error messages and optionally emitted to the object file have the following parts.

`L`

All local labels begin with `L`. Normally both `as` and `ld` forget symbols that start with `L`. These labels are used for symbols you are never intended to see. If you use the `-L` option then `as` retains these symbols in the object file. If you also instruct `ld` to retain these symbols, you may use them in debugging.

digit

If the label is written `0:` then the digit is 0. If the label is written `1:` then the digit is 1. And so on up through `9:`.

`^A`

This unusual character is included so you do not accidentally invent a symbol of the same name. The character has ASCII value `\001`.

ordinal number

This is a serial number to keep the labels distinct. The first `0:` gets the number 1;

The 15th `0:` gets the number 15; etc.. Likewise for the other labels `1:` through `9:`.

For instance, the first `1:` is named `L1^A1`, the 44th `3:` is named `L3^A44`.

The special dot symbol

The special symbol “.” refers to the current address into which `as` is assembling. Thus, the expression, `melvin: .long`, defines `melvin` to contain its own address.

Assigning a value to “.” is treated the same as a `.org` directive. Thus, the expression of `‘.=.+4’` is the same as saying `‘.space 4’`.

Symbol attributes

Every symbol has, as well as its name, the attributes, *Value* and *Type*. Depending on output format, symbols can also have auxiliary attributes.

If you use a symbol without defining it, `as` assumes zero for all these attributes, and probably won't warn you. This makes the symbol an externally defined symbol, which is generally what you would want.

Value

The value of a symbol is (usually) 32 bits. For a symbol which labels a location in the text, data, bss or absolute sections the value is the number of addresses from the start of that section to the label. Naturally for text, data and bss sections the value of a symbol changes as `ld` changes section base addresses during linking. Absolute symbols' values do not change during linking: that is why they are called absolute.

The value of an undefined symbol is treated in a special way. If it is 0 then the symbol is not defined in this assembler source file, and `ld` tries to determine its value from other files linked into the same program. You make this kind of symbol simply by mentioning a symbol name without defining it. A non-zero value represents a `.comm` common declaration. The value is how much common storage to reserve, in bytes (addresses). The symbol refers to the first address of the allocated storage.

Type

The type attribute of a symbol contains relocation (section) information, any flag settings indicating that a symbol is external, and (optionally), other information for linkers and debuggers. The exact format depends on the object-code output format in use.

Symbol attributes for a.out file format

The following documentation discusses the symbol attributes for the `a.out` file format.

Descriptor

This is an arbitrary 16-bit value. You may establish a symbol's descriptor value by using a `.desc` statement (see “`.desc` symbol, abs-expression directive” on page 70). A descriptor value means nothing to `as`.

Other

This is an arbitrary 8-bit value. It means nothing to `as`.

Symbol attributes for COFF format

The COFF format supports a multitude of auxiliary symbol attributes; like the primary symbol attributes, they are set between `.def` and `.endef` directives.

Primary attributes

The symbol name is set with `.def`; the value and type, respectively, with `.val` and `.type`.

Auxiliary attributes

The `as` directives `.dim`, `.line`, `.scl`, `.size`, and `.tag` can generate auxiliary symbol table information for COFF.

Symbol attributes for SOM format

The SOM format for the HPPA supports a multitude of symbol attributes set with the `.EXPORT` and `.IMPORT` directives. The attributes are described in *HP9000 Series 800 Assembly Language Reference Manual* (HP 92432-90001) with the “IMPORT and EXPORT assembler directive” documentation.

Expressions

The following documentation describes the assembler syntax parts such as expressions, arguments, symbols and operators. An *expression* specifies an address or numeric value. Whitespace may precede and/or follow an expression. An *argument* is like an operand; however, with `as`, the argument has a different meaning. An *operator* is an arithmetic function. See the following documentation for more specific discussion.

- “Empty expressions” on page 60
- “Integer expressions” on page 61
- “Arguments” on page 62
- “Operators” on page 63

The result of an expression must be an absolute number, or else an off-set into a particular section. If an expression is not absolute, and there is not enough information when `as` sees the expression to know its section, a second pass over the source program might be necessary to interpret the expression—but the second pass is currently not implemented. `as` aborts with an error message in this situation.

Empty expressions

An empty expression has no value: it is just whitespace or null. Wherever an absolute expression is required, you may omit the expression, and `as` assumes a value of (absolute) 0. This is compatible with other assemblers.

Integer expressions

An *integer expression* is one or more *arguments* delimited by *operators*.

Arguments

Arguments are symbols, numbers or sub-expressions. In other contexts, *arguments* are sometimes called *arithmetic operands*. In this documentation, to avoid confusing them with the *instruction operands* of the machine language, we use the term *argument* to refer to parts of expressions only, reserving the word *operand* to refer only to machine instruction operands.

Symbols are evaluated to yield `{section NNN}` where *section* is one of text, data, bss, absolute, or undefined. *NNN* is a signed, 2's complement 32 bit integer. Numbers are usually integers.

A number can be a flonum or bignum. In this case, you are warned that only the low order 32 bits are used, and `as` pretends these 32 bits are an integer. You may write integer-manipulating instructions that act on exotic constants, compatible with other assemblers.

Sub-expressions are a left parenthesis '(' followed by an integer expression, followed by a right parenthesis ')'; or a prefix operator followed by an argument.

Operators

Operators are arithmetic functions, like + or %. Prefix operators are followed by an argument (see “Prefix operators”). Infix operators appear between their arguments (see “Infix operators”). Operators may be preceded and/or followed by whitespace.

Prefix operators

`as` has the following *prefix operators*. They each take one argument, which must be absolute.

- (*Negation*)
Two’s complement negation.
- ~ (*Complementation*)
Bitwise not.

Infix operators

Infix operators take two arguments, one on either side. Operators have precedence, but operations with equal precedence are performed left to right. Apart from + or -, both arguments must be absolute, and the result is absolute.

■ *Highest Precedence*

- * (*Multiplication*)
- / (*Division*)
Truncation is the same as the C operator ‘/’.
- % (*Remainder*)
- <
- << (*Shift Left*)
Same as the C operator ‘<<’.
- >
- >> (*Shift Right*)
Same as the C operator ‘>>’.

■ *Intermediate Precedence*

- | (*Bitwise Inclusive Or*)
- & (*Bitwise And*)
- ^ (*Bitwise Exclusive Or*)
- ! (*Bitwise Or Not*)

■ ***Lowest Precedence***

+ (*Addition*)

If either argument is absolute, the result has the section of the other argument.
You may not add together arguments from different sections.

- (*Subtraction*)

If the right argument is absolute, the result has the section of the left argument. If both arguments are in the same section, the result is absolute.
You may not subtract arguments from different sections.

In short, it's only meaningful to add or subtract the *offsets* in an address; you can only have a defined section in one of the two arguments.

Assembler directives

All assembler directives have names that begin with a period (.). The rest of the name is letters, usually in lower case.

The following documentation lists the directives that are available regardless of the target machine configuration for GAS, the GNU assembler. With each directive is the location of their descriptions. For additional directives pertinent to each architecture, see “Machine dependent features” on page 89.

- “.abort directive” on page 67
- “.ABORT directive” on page 68
- “.align abs-expr, .abs-expr directive” on page 68
- “.app-file string directive” on page 68
- “.ascii “string”... directive” on page 68
- “.asciz “string”... directive” on page 69
- “.balign[wl] abs-expr, abs-expr directive” on page 69
- “.byte expressions directive” on page 69
- “.comm symbol, length directive” on page 69
- “.data subsection directive” on page 70
- “.def name directive” on page 70
- “.desc symbol, abs-expression directive” on page 70

- “.dim directive” on page 70
- “.double flonums directive” on page 70
- “.eject directive” on page 70
- “.else directive” on page 71
- “.endif directive” on page 71
- “.endif directive” on page 71
- “.equ symbol, expression directive” on page 71
- “.equiv symbol, expression directive” on page 71
- “.err directive” on page 71
- “.extern directive” on page 72
- “.file string directive” on page 72
- “.fill repeat, size, value directive” on page 72
- “.float flonums directive” on page 72
- “.global symbol, .globl symbol directive” on page 73
- “.hword expressions directive” on page 73
- “.ident directive” on page 73
- “.if absolute expression directive” on page 73
- “.include “file” directive” on page 74
- “.int expressions directive” on page 74
- “.irp symbol, values... directive” on page 74
- “.irpc symbol, values... directive” on page 74
- “.lcomm symbol, length directive” on page 75
- “.lflags directive” on page 75
- “.line line-number directive” on page 75
- “.list directive” on page 76
- “.ln line-number directive” on page 76
- “.long expressions directive” on page 76
- “.macro directive” on page 76
- “.mri val directive” on page 77
- “.nolist directive” on page 77
- “.octa bignums directive” on page 77
- “.org new-lc, fill directive” on page 78
- “.p2align[w1] abs-expr, abs-expr directive” on page 78

- “.psize lines, columns directive” on page 79
- “.quad bignums directive” on page 79
- “.rept count directive” on page 79
- “.sbtll “subheading” directive” on page 79
- “.scl class directive” on page 80
- “.section name directive” on page 80
- “.set symbol, expression directive” on page 81
- “.short expressions directive” on page 82
- “.single flonums directive” on page 82
- “.size directive” on page 82
- “.sleb128 expressions directive” on page 82
- “.skip size, fill directive” on page 82
- “.space size, fill directive” on page 82
- “.stabd, .stabn, and .stabs directives” on page 83
- “.symver directive” on page 84
- “.tag structname directive” on page 85
- “.text subsection directive” on page 85
- “.title “heading” directive” on page 85
- “.type int directive” on page 85
- “.val addr directive” on page 85
- “.uleb128 expressions directive” on page 85
- “.word expressions directive” on page 86

One day the following directives won’t work. They are included for compatibility with older assemblers.

```
.abort
.app-file
.line
```

See the following documentation for specific information on the previously listed assembler directives.

.abort directive

This directive stops the assembly immediately. It is for compatibility with other assemblers. The original idea was that the assembly language source would be piped into the assembler. If the sender of the source quit, it could use this directive to tell as to quit also. One day `.abort` will not be supported.

.ABORT directive

When producing COFF output, GAS accepts this directive as a synonym for `.abort`.

When producing `b.out` output, GAS accepts this directive, but ignores it.

.align *abs-expr*, .*abs-expr* directive

Pad the location counter (in the current subsection) to a particular storage boundary. The first expression (which must be absolute) is the alignment required, as described in the following. The second expression (also absolute) gives the value to be stored in the padding bytes. It (and the comma) may be omitted. If it is omitted, the padding bytes are zero.

The way the required alignment is specified varies from system to system. For the a29k, hppa, m86k, m88k, w65, sparc, and Hitachi SH, and i386 using ELF format, the first expression is the alignment request in bytes. For example, `.align 8` advances the location counter until it is a multiple of 8. If the location counter is already a multiple of 8, no change is needed. For other systems, including the i386 using a.out format, it is the number of low-order zero bits the location counter must have after advancement. For example, `.align 3` advances the location counter until it is a multiple of 8. If the location counter is already a multiple of 8, no change is needed.

This inconsistency is due to the different behaviors of the various native assemblers for these systems which GAS must emulate. GAS also provides `.balign` and `.p2align` directives, which have a consistent behavior across all architectures (but are specific to GAS).

.app-file *string* directive

`.app-file` (which may also be spelled `‘.file’`) tells us that we are about to start a new logical file. *string* is the new file name. In general, the filename is recognized whether or not it is surrounded by quotes `“”`; but if you wish to specify an empty filename that is permitted, you must give the quotes—`“”`. This statement may go away in future: it is only recognized to be compatible with old GAS programs.

.ascii “*string*”... directive

`.ascii` expects zero or more string literals (see “Strings” on page 38) separated by commas. It assembles each string (with no automatic trailing zero byte) into consecutive addresses.

.asciz "string"... directive

`.asciz` is just like `.ascii`, but each string is followed by a zero byte. The `z` in `.asciz` stands for *zero*.

.balign[w1] *abs-expr*, *abs-expr* directive

Pad the location counter (in the current subsection) to a particular storage boundary. The first expression (which must be absolute) is the alignment request in bytes. For example `.balign 8` advances the location counter until it is a multiple of 8. If the location counter is already a multiple of 8, no change is needed.

The second expression (also absolute) gives the value to be stored in the padding bytes. It (and the comma) may be omitted. If it is omitted, the padding bytes are zero.

The `.balignw` and `.balignl` directives are variants of the `.balign` directive. The `.balignw` directive treats the fill pattern as a two byte word value. The `.balignl` directive treats the fill pattern as a four byte longword value. For example, `.balignw 4, 0x368d` will align to a multiple of 4. If it skips two bytes, they will be filled in with the value `0x368d` (the exact placement of the bytes depends upon the endianness of the processor). If it skips 1 or 3 bytes, the fill value is undefined.

.byte *expressions* directive

`.byte` expects zero or more expressions, separated by commas. Each expression is assembled into the next byte.

.comm *symbol*, *length* directive

`.comm` declares a common symbol named *symbol*. When linking, a common symbol in one object file may be merged with a defined or common symbol of the same name in another object file. If `ld` does not see a definition for the symbol—just one or more common symbols—it will allocate *length* bytes of uninitialized memory. *length* must be an absolute expression. If `ld` sees multiple common symbols with the same name, and they do not all have the same size, it will allocate space using the largest size.

When using ELF, the `.comm` directive takes an optional third argument. This is the desired alignment of the symbol, specified as a byte boundary (for example, an alignment of 16 means that the least significant 4 bits of the address should be 0). The alignment must be an absolute expression, and it must be a power of 2. If `ld` allocates uninitialized memory for the common symbol, it will use the alignment when placing the symbol. If no alignment is specified, `as` will set the alignment to the largest power of 2 less than or equal to the size of the symbol, up to a maximum of 16.

The syntax for `.comm` differs slightly on the HPPA. The syntax is `symbol.comm, length`; `symbol` is optional.

`.data subsection directive`

`.data` tells `as` to assemble the following statements onto the end of the data subsection numbered *subsection* (which is an absolute expression). If *subsection* is omitted, it defaults to zero.

`.def name directive`

Begin defining debugging information for a symbol *name*; the definition extends until the `.endef` directive is encountered.

This directive is only observed when `as` is configured for COFF format output; when producing `b.out`, `.def` is recognized, but ignored.

`.desc symbol, abs-expression directive`

This directive sets the descriptor of the symbol (see “Symbol attributes” on page 57) to the low 16 bits of an absolute expression.

The `.desc` directive is not available when `as` is configured for COFF output; it is only for `a.out` or `b.out` object format. For the sake of compatibility, `as` accepts it, but produces no output, when configured for COFF.

`.dim directive`

This directive is generated by compilers to include auxiliary debugging information in the symbol table. It is only permitted inside `.def/.endef` pairs. `.dim` is only meaningful when generating COFF format output; when `as` is generating `b.out`, it accepts this directive but ignores it.

`.double flonums directive`

`.double` expects zero or more flonums, separated by commas. It assembles floating point numbers. The exact kind of floating point numbers emitted depends on how `as` is configured. For each architecture’s documentation, see “Machine dependent features” on page 89.

`.eject directive`

Force a page break at this point, when generating assembly listings.

.else directive

`.else` is part of the `as` support for conditional assembly; see “.if absolute expression directive” on page 73. It marks the beginning of a section of code to be assembled if the condition for the preceding `.if` was false.

.endef directive

This directive flags the end of a symbol definition begun with `.def`.

`.endef` is only meaningful when generating COFF format output; if `as` is configured to generate `b.out`, it accepts this directive but ignores it.

.endif directive

`.endif` is part of the `as` support for conditional assembly; it marks the end of a block of code that is only assembled conditionally. See “.if absolute expression directive” on page 73.

.equ *symbol*, *expression* directive

This directive sets the value of *symbol* to *expression*. It is synonymous with `.set`; see “.set symbol, expression directive” on page 81. The syntax for `equ` on the HPPA is *symbol.equ expression*.

.equiv *symbol*, *expression* directive

The `.equiv` directive is like the `.equ` and the `.set` directives, except the assembler will signal an error if *symbol* is already defined.

Except for the contents of the error message, the following input is roughly equivalent.

```
.ifdef SYM
.err
.endif
.equ SYM,VAL
```

.err directive

If `as` assembles a `.err` directive, it will print an error message and, unless the `-z` option was used, it will not generate an object file. This can be used to signal error an unconditionally compiled code..

.extern directive

`.extern` is accepted in the source program—for compatibility with other assemblers—but it is ignored. `as` treats all undefined symbols as external.

.file *string* directive

`.file` (which may also be spelled `.app-file`) tells `as` that we are about to start a new logical file. *string* is the new file name. In general, the filename is recognized whether or not it is surrounded by quotes (""); but if you wish to specify an empty file name, you must give the quotes—"". This statement may go away in future; it is only recognized to be compatible with old `as` programs. In some configurations of `as`, `.file` has already been removed to avoid conflicts with other assemblers. For each architecture's documentation, see "Machine dependent features" on page 89.

.fill *repeat*, *size*, *value* directive

result, *size* and *value* are absolute expressions. This emits *repeat* copies of *size* bytes. *repeat* may be zero or more. *size* may be zero or more, but if it is more than 8, then it is deemed to have the value 8, compatible with other people's assemblers. The contents of each *repeat* bytes is taken from an 8-byte number. The highest order 4 bytes are zero. The lowest order 4 bytes are *value* rendered in the byte-order of an integer on the computer `as` is assembling for. Each *size* bytes in a repetition is taken from the lowest order *size* bytes of this number. Again, this bizarre behavior is compatible with other people's assemblers.

size and *value* are optional. If the second comma and *value* are absent, *value* is assumed zero. If the first comma and following tokens are absent, *size* is assumed to be 1.

.float *flonums* directive

`.float` assembles zero or more flonums, separated by commas. It has the same effect as `.single`. The exact kind of floating point numbers emitted depends on how `as` is configured. For each architecture's documentation, see "Machine dependent features" on page 89.

.global *symbol*, .globl *symbol* directive

`.global` makes the symbol visible to `ld`. If you define *symbol* in your partial program, its value is made available to other partial programs that are linked with it. Otherwise, *symbol* takes its attributes from a symbol of the same name from another file linked into the same program. Both spellings (`.globl` and `.global`) are accepted, for compatibility with other assemblers.

On the HPPA, `.global` is not always enough to make it accessible to other partial programs. You may need the HPPA-only `.EXPORT` directive as well. See “HPPA assembler directives” on page 135.

.hword *expressions* directive

`.hword` expects zero or more expressions, and emits a 16 bit number for each.

This directive is a synonym for `.short`; depending on the target architecture, it may also be a synonym for `.word`.

.ident directive

`.ident` is used by some assemblers to place tags in object files. `as` simply accepts the directive for source-file compatibility with such assemblers, but does not actually emit anything for it.

.if *absolute expression* directive

`.if` marks the beginning of a section of code which is only considered part of the source program being assembled if the argument (which must be an absolute expression) is non-zero. The end of the conditional section of code must be marked by `.endif` (see “`.endif` directive” on page 71); optionally, you may include code for the alternative condition, flagged by `.else` (see “`.else` directive” on page 71). The following variants of `.if` are also supported:

```
.ifdef symbol
```

Assembles the following section of code if the specified *symbol* has been defined.

```
.ifndef symbol
```

```
ifndef symbol
```

Assembles the following section of code if the specified *symbol* has not been defined. Both spelling variants are equivalent.

.include "file" directive

`.include` provides a way to include supporting files at specified points in your source program. The code from *file* is assembled as if it followed the point of the `.include`; when the end of the included file is reached, assembly of the original file continues. You can control the search paths used with the `-I` command-line option (see “`.include` search path: `-I path option`” on page 20). Quotation marks are required around *file*.

.int expressions directive

Expect zero or more *expressions*, of any section, separated by commas. For each expression, emit a number that, at run time, is the value of that expression. The byte order and bit size of the number depends on what kind of target the assembly uses.

.irp symbol, values... directive

Evaluate a sequence of statements assigning different values to *symbol*. The sequence of statements starts at the `.irp` directive, and is terminated by an `.endr` directive. For each *value*, *symbol* is set to *value*, and the sequence of statements is assembled.

If no *value* is listed, the sequence of statements is assembled once, with *symbol* set to the null string. To refer to *symbol* within the sequence of statements, use `\symbol`.

```
.irp      param,1,2,3
move     d\param,sp@-
.endr
```

For example, assembling the previous statement is equivalent to assembling the following.

```
move     d1,sp@-
move     d2,sp@-
move     d3,sp@-
```

.irpc symbol, values... directive

Evaluate a sequence of statements assigning different values to *symbol*. The sequence of statements starts at the `.irpc` directive, and is terminated by an `.endr` directive. For each character in *value*, *symbol* is set to the character, and the sequence of statements is assembled. If no *value* is listed, the sequence of statements is assembled once, with *symbol* set to the null string. To refer to *symbol* within the sequence of statements, use `\symbol`.

```
.irpc     param,123
move     d\param,sp@-
.endr
```

For example, assembling the previous statement is equivalent to assembling the following declaration.

```
move      d1, sp@-
move      d2, sp@-
move      d3, sp@-7.32
```

.lcomm *symbol*, *length* directive

.lcomm Reserve *length* (an absolute expression) bytes for a local common denoted by *symbol*. The section and value of *symbol* are those of the new local common. The addresses are allocated in the bss section, so that at run-time the bytes start off zeroed. *symbol* is not declared global (see “.global *symbol*, .globl *symbol* directive” on page 73), so it’s normally not visible to `ld`.

The syntax for **.lcomm** differs slightly on the HPPA. The syntax is *symbol.lcomm, length*; *symbol* is optional.

.lflags directive

`as` accepts this directive, for compatibility with other assemblers, but ignores it.

.line *line-number* directive

Change the logical line number. *line-number* must be an absolute expression. The next line has that logical line number. Therefore any other statements on the current line (after a statement separator character) are reported as on logical line number *line-number*-1. One day `as` will no longer support this directive: it is recognized only for compatibility with existing assembler programs.

WARNING: In the AMD29K configuration of `as`, this command is not available; use the synonym `.ln` in that context.

Even though this is a directive associated with the `a.out` or `b.out` object-code formats, `as` still recognizes it when producing COFF output, and treats **.line** as though it were the COFF `.ln` if it is found outside a `.def/.endef` pair. Inside a `.def`, **.line** is, instead, one of the directives used by compilers to generate auxiliary symbol information for debugging.

.list directive

Control (in conjunction with the `.nolist` directive) whether or not assembly listings are generated. These two directives maintain an internal counter (which is zero initially). `.list` increments the counter, and `.nolist` decrements it. Assembly listings are generated whenever the counter is greater than zero. By default, listings are disabled. When you enable them (with the `-a` command line option; see “Enable listings: `-a[cdhlms]` options” on page 17), the initial value of the listing counter is one.

.ln *line-number* directive

`.ln` is a synonym for `.line`.

.long *expressions* directive

`.long` is the same as `.int`, see “.int expressions directive” on page 74.

.macro directive

The commands `.macro` and `.endm` allow you to define macros that generate assembly output. For example, the following definition specifies a macro `sum` that puts a sequence of numbers into memory.

```
.macro      sum
.long      \from
.if        \to-\from
sum        "(\from+1)",\to
.endif     from=0,to=5
.endm
```

With that definition, `SUM 0,5` is equivalent to the following assembly input.

```
.long      0
.long      1
.long      2
.long      3
.long      4
.long      5
.macro macname
.macro macname macargs...
```

Begin the definition of a macro called *macname*. If your macro definition requires arguments, specify their names after the macro name, separated by commas or spaces. You can supply a default value for any macro argument by following the name with `=deflt`. For example, the following are all valid `.macro` statements:

```
.macro comm
```

Begin the definition of a macro called `comm`, which takes no arguments.


```
.macro plus1 p, p1
.macro plus1 p p1
```

Either statement begins the definition of a macro called `plus1`, which takes two arguments; within the macro definition, write `\p` or `\p1` to evaluate the arguments.

```
.macro reserve_str p1=0 p2
```

Begin the definition of a macro called `reserve_str`, with two arguments. The first argument has a default value, but not the second.

After the definition is complete, you can call the macro either as `reserve_str a, b` (with `\p1` evaluating to `a` and `\p2` evaluating to `b`), or as `reserve_str, b` (with `\p1` evaluating as the default, in this case `0`, and `\p2` evaluating to `b`).

When you call a macro, you can specify the argument values either by position, or by keyword.

For example, `sum 9,17` is equivalent to `sum to=17, from=9`.

```
.endm
```

Mark the end of a macro definition.

```
.exitm
```

Exit early from the current macro definition.

```
\@
```

`as` maintains a counter of how many macros it has executed in this pseudo-variable; you can copy that number to your output with `\@`, but *only within a macro definition*.

.mri val directive

If `val` is non-zero, this tells `as` to enter MRI mode. If `val` is zero, this tells `as` to exit MRI mode. This change affects code assembled until the next `.mri` directive, or until the end of the file. See “Assemble in MRI compatibility mode: `-M` option” on page 23.

.nolist directive

Control (in conjunction with the `.list` directive) whether or not assembly listings are generated. These two directives maintain an internal counter (which is zero initially). `.list` increments the counter, and `.nolist` decrements it. Assembly listings are generated whenever the counter is greater than zero.

.octa bignums directive

`.octa` This directive expects zero or more bignums, separated by commas. For each bignum, it emits a 16-byte integer.

The term *octa* comes from contexts in which a *word* is two bytes; hence *octa*-word for 16 bytes.

.org *new-lc*, *fill* directive

Advance the location counter of the current section to *new-lc*. *new-lc* is either an absolute expression or an expression with the same section as the current subsection. That is, you can't use `.org` to cross sections: if *new-lc* has the wrong section, the `.org` directive is ignored. To be compatible with former assemblers, if the section of *new-lc* is absolute, as issues a warning, then pretends the section of *new-lc* is the same as the current subsection.

`.org` may only increase the location counter, or leave it unchanged; you cannot use `.org` to move the location counter backwards.

Because `as` tries to assemble programs in one pass, *new-lc* may not be undefined. If you really detest this restriction we eagerly await a chance to share your improved assembler.

Beware that the origin is relative to the start of the section, not to the start of the subsection. This is compatible with other people's assemblers.

When the location counter (of the current subsection) is advanced, the intervening bytes are filled with *fill* which should be an absolute expression. If the comma and *fill* are omitted, *fill* defaults to zero.

.p2align[w1] *abs-expr*, *abs-expr* directive

`.p2align[w1]` Pad the location counter (in the current subsection) to a particular storage boundary.

The first expression (which must be absolute) is the number of low-order zero bits the location counter must have after advancement. For example `.p2align 3` advances the location counter until it is a multiple of 8. If the location counter is already a multiple of 8, no change is needed.

The second expression (also absolute) gives the value to be stored in the padding bytes. It (and the comma) may be omitted. If it is omitted, the padding bytes are zero.

The `.p2alignw` and `.p2alignl` directives are variants of the `.p2align` directive. The `.p2alignw` directive treats the fill pattern as a two byte word value. The `.p2alignl` directive treats the fill pattern as a four byte longword value. For example, `.p2alignw 2,0x368d` will align to a multiple of 4. If it skips two bytes, they will be filled in with the value 0x368d (the exact placement of the bytes depends upon the endianness of the processor). If it skips 1 or 3 bytes, the fill value is undefined.

.psize *lines, columns* **directive**

Use this directive to declare the number of lines—and, optionally, the number of columns—to use for each page, when generating listings.

If you do not use `.psize`, listings use a default line-count of 60. You may omit the comma and columns specification; the default width is 200 columns.

`as` generates formfeeds whenever the specified number of lines is exceeded (or whenever you explicitly request one, using `.eject`).

If you specify *lines* as 0, no formfeeds are generated save those explicitly specified with `.eject`.

.quad *bignums* **directive**

`.quad` expects zero or more bignums, separated by commas. For each bignum, it emits an 8-byte integer. If the bignum won't fit in 8 bytes, it prints a warning message; and just takes the lowest order 8 bytes of the bignum. The term *quad* comes from contexts in which a *word* is two bytes; hence *quad*-word for 8 bytes.

.rept *count* **directive**

Repeat the sequence of lines between the `.rept` directive and the next `.endr` directive *count* times (where *count* stands for the appropriate sequence).

```
.rept          3
.long         0
.endr
```

For example, assembling the previous statement is equivalent to assembling the following directive.

```
.long         0
.long         0
.long         0
```

.sbt1 "*subheading*" **directive**

Use *subheading* as the title (third line, immediately after the title line) when generating assembly listings.

This directive affects subsequent pages, as well as the current page if it appears within ten lines of the top of a page.

.scl *class* directive

Set the storage-class value for a symbol. This directive may only be used inside a `.def/.endef` pair. Storage class may flag whether a symbol is static or external, or it may record further symbolic debugging information.

The `.scl` directive is primarily associated with COFF output; when configured to generate `b.out` output format, it accepts this directive but ignores it.

.section *name* directive

Use the `.section` directive to assemble the following code into a section called *name*.

This directive is only supported for targets that actually support arbitrarily named sections; on `a.out` targets, for example, it is not accepted, even with a standard `a.out` section name.

For COFF targets, the `.section` directive is used in one of the following ways.

```
.section name[, "flags"]  
.section name[, subsegment]
```

If the optional argument is quoted, it is taken as flags to use for the section. Each flag is a single character. The following flags are recognized.

b	bss section (uninitialized data)
n	section is not loaded
w	writable section
d	data section
r	read-only section
x	executable section

If no flags are specified, the default flags depend upon the section name. If the section name is not recognized, the default will be for the section to be loaded and writable.

If the optional argument to the `.section` directive is not quoted, it is taken as a subsequent number (see “Sub-sections” on page 48).

For ELF targets, the `.section` directive is used in the following way.

```
.section name[, "flags"[, @type]]
```

The optional *flags* argument is a quoted string which may contain any combination of the following characters.

```
a
    section is allocatable
w
    section is writable
x
    section is executable
```

The optional *type* argument may contain one of the following constants.

```
@progbits
section contains data
@nobits
section does not contain data (that is, section only occupies space)
```

If no flags are specified, the default flags depend on the section name. If the section name is not recognized, the default will be for the section to have none of the previously described flags; it will not be allocated in memory, nor will it be writable or executable. The section will contain data.

For ELF targets, the assembler supports another type of `.section` directive for compatibility with the Solaris assembler, as with the following declaration.

```
.section "name"[, flags...]
```

Notice that the section name is quoted. There may be a sequence of the following comma-separated flags.

```
#alloc
    section is allocatable
#write
    section is writable
#excinstr
    section is executable
```

`.set` *symbol*, *expression* directive

Set the value of *symbol* to *expression*. This changes *symbol*'s value and type to conform to *expression*. If *symbol* was flagged as external, it remains flagged. (See “Symbol attributes” on page 57.) You may `.set` a symbol many times in the same assembly.

If you `.set` a global symbol, the value stored in the object file is the last value stored into it.

The syntax for `set` on the HPPA is `symbol.set expression`.

.short *expressions* directive

`.short` is normally the same as `.word`. See “`.word` expressions directive” on page 86.

In some configurations, however, `.short` and `.word` generate numbers of different lengths; for a specific architecture’s documentation, see “Machine dependent features” on page 89.

.single *flonums* directive

`.single` assembles zero or more flonums, separated by commas. It has the same effect as `.float`. The exact kind of floating point numbers emitted depends on how `as` is configured; for a specific architecture’s documentation, see “Machine dependent features” on page 89.

.size directive

This directive is generated by compilers to include auxiliary de-bugging information in the symbol table. It is only permitted inside `.def/.endef` pairs. `.size` is only meaningful when generating COFF format output; when `as` is generating `b.out`, it accepts this directive but ignores it.

.sleb128 *expressions* directive

`.sleb128` stands for “signed little endian base 128.” This is a compact, variable length representation of numbers used by the DWARF symbolic debugging format. See also “`.uleb128` expressions directive” on page 85.

.skip *size*, *fill* directive

This directive emits *size* bytes, each of value, *fill*. Both *size* and *fill* are absolute expressions. If the comma and *fill* are omitted, *fill* is assumed to be zero. This is the same as `.space`.

.space *size*, *fill* directive

This directive emits *size* bytes, each of value *fill*. Both *size* and *fill* are absolute expressions. If the comma and *fill* are omitted, *fill* is assumed to be zero.

WARNING: `.space` has a completely different meaning for HPPA targets; use `.block` as a substitute. See *HP9000 Series 800 Assembly Language Reference Manual* (HP 92432-90001) for the meaning of the `.space` directive. See “HPPA dependent features” on page 133.

On the AMD 29K, this directive is ignored; it is accepted for compatibility with other AMD 29K assemblers.

WARNING: In most versions of the GNU assembler, the `.space` directive has the effect of `.block`. For each architecture’s documentation, see “Machine dependent features” on page 89.

.stabd, .stabs, and .stabs directives

There are three directives that begin `.stab`. All emit symbols for use by symbolic debuggers (see “Symbols” on page 36). The symbols are not entered in the `as` hash table: they cannot be referenced elsewhere in the source file. Up to five fields are required, as the following descriptions clarify.

string

This is the symbol’s name. It may contain any character except `\000`, so is more general than ordinary symbol names. Some debuggers used to code arbitrarily complex structures into symbol names using this field.

type

An absolute expression. The symbol’s type is set to the low 8 bits of this expression. Any bit pattern is permitted, but `ld` and debuggers choke on silly bit patterns.

other

An absolute expression. The symbol’s “other” attribute is set to the low 8 bits of this expression.

desc

An absolute expression. The symbol’s descriptor is set to the low 16 bits of this expression.

value

An absolute expression which becomes the symbol’s value.

If a warning is detected while reading a `.stabd`, `.stabs`, or `.stabs` statement, the symbol has probably already been created; you get a half-formed symbol in your object file. This is compatible with earlier assemblers!

`.stabd type, other, desc`

The “name” of the symbol generated is not even an empty string. It is a null pointer, for compatibility. Older assemblers used a null pointer so they didn’t waste space in object files with empty strings.

The symbol’s value is set to the location counter, relocatability. When your program is linked, the value of this symbol is the address of the location counter when the `.stabd` was assembled.

```
.stabsn type, other , desc, value
```

The name of the symbol is set to the empty string "".

```
.stabs string , type, other , desc, value
```

All five fields are specified.

.string "*str*" directive

Copy the characters in *str* to the object file. You may specify more than one string to copy, separated by commas. Unless otherwise specified for a particular machine, the assembler marks the end of each string with a 0 byte. You can use any of the escape sequences described in “Strings” on page 38

.symver directive

Use the `.symver` directive to bind symbols to specific version nodes within a source file. This is only supported on ELF platforms, and is typically used when assembling files to be linked into a shared library. There are cases where it may make sense to use this in objects to be bound into an application itself so as to override a versioned symbol from a shared library.

For ELF targets, the `.symver` directive is used like the following declaration shows

```
.symver name, name2@nodename
```

In this case, the symbol, *name*, must exist and be defined within the file being assembled.

The `.versym` directive effectively creates a symbol alias with the name *name2@nodename*, and in fact the main reason that we just don't try and create a regular alias is that the `@` character isn't permitted in symbol names. The *name2* part of the name is the actual name of the symbol by which it will be externally referenced. The name, *name*, itself is merely a name of convenience that is used so that it is possible to have definitions for multiple versions of a function within a single source file, and so that the compiler can unambiguously know which version of a function is being mentioned. The *nodename* portion of the alias should be the name of a node specified in the version script supplied to the linker when building a shared library. If you are attempting to override a versioned symbol from a shared library, then *nodename* should correspond to the nodename of the symbol you are trying to override.it.

.tag *structname* directive

This directive is generated by compilers to include auxiliary debugging information in the symbol table. It is only permitted inside `.def/.endef` pairs. Tags are used to link structure definitions in the symbol table with instances of those structures. `.tag` is only used when generating COFF format output; when `as` is generating `b.out`, it accepts this directive but ignores it.

.text *subsection* directive

`.text` tells `as` to assemble the following statements onto the end of the text subsection numbered *subsection*, which is an absolute expression. If *subsection* is omitted, subsection number zero is used.

.title "*heading*" directive

Use *heading* as the title (second line, immediately after the source file name and pagenumber) when generating assembly listings. `.title` affects subsequent pages, as well as the current page if it appears within ten lines of the top of a page.

.type *int* directive

This directive, permitted only within `.def/.endef` pairs, records the integer *int* as the type attribute of a symbol table entry. `.type` is associated only with COFF format output; when `as` is configured for `b.out` output, it accepts this directive but ignores it.

.val *addr* directive

This directive, permitted only within `.def/.endef` pairs, records the address *addr* as the value attribute of a symbol table entry.

`.val` is used only for COFF output; when `as` is configured for `b.out`, it accepts this directive but ignores it.

.uleb128 *expressions* directive

`uleb128` stands for “unsigned little endian base 128.” This is a compact, variable length representation of numbers used by the DWARF symbolic debugging format. See “`.sleb128 expressions directive`” on page 82.

.word *expressions* **directive**

This directive expects zero or more *expressions*, of any section, separated by commas. The size of the number emitted, and its byte order, depend on what target computer for which the assembly builds.

WARNING: To support compilers on machines with a 32-bit address space, having less than 32-bit addressing, this requires special treatment. If the machine of interest to you does 32-bit addressing (or doesn't require it), ignore this issue. For documentation for specific processors, see "Machine dependent features" on page 89. In order to assemble compiler output into something that works, *as* occasionally does strange things to `.word` directives.

Directives of the form, `.word sym1-sym2`, are often emitted by compilers as part of jump tables. Therefore, when *as* assembles a directive of the form, `.word sym1-sym2`, and the difference between `sym1` and `sym2` does not fit in 16 bits, *as* creates a *secondary jump table*, immediately before the next label. This secondary jump table is preceded by a short-jump to the first byte after the secondary table. This short-jump prevents the flow of control from accidentally falling into the new table. Inside the table is a long-jump to `sym2`. The original `.word` contains `sym1` minus the address of the long-jump to `sym2`. If there were several occurrences of `.word sym1-sym2` before the secondary jump table, all of them are adjusted.

If there was a `.word sym3-sym4` that also did not fit in sixteen bits, a long-jump to `sym4` is included in the secondary jump table, and the `.word` directives are adjusted to contain `sym3` minus the address of the long-jump to `sym4`; and so on, for as many entries in the original jump table as necessary.

Machine dependent features

The machine instruction sets are (almost by definition) different on each machine where the GNU assembler, `as`, runs. Floating point representations vary as well, and `as` often supports a few additional directives or command-line options for compatibility with other assemblers on a particular platform. Finally, some versions of `as` support special pseudo-instructions for branch optimization.

There is discussion on most of these differences, though there aren't details on any machine's instruction set. For details on that subject, see the specific hardware manufacturer's documentation.

The following architectures are discussed.

- “AMD 29K dependent features” on page 91
- “ARC dependent features” on page 97
- “ARM dependent features” on page 101
- “AT&T/Intel dependent features” on page 105
- “D10V dependent features” on page 113
- “H8/300 dependent features” on page 121
- “H8/500 dependent features” on page 127
- “HPPA dependent features” on page 133
- “SH dependent features” on page 139

- “Intel 960 dependent features” on page 145
- “M68K dependent features” on page 151
- “MIPS dependent features” on page 161
- “SPARC dependent features” on page 169
- “Vax dependent features” on page 173

AMD 29K dependent features

The following documentation discusses the features pertinent to the AMD 29K machines regarding the GNU assembler.

- “AMD 29K options” on page 92
- “AMD 29K syntax” on page 93
- “AMD 29K macros” on page 93
- “AMD 29K special characters” on page 93
- “AMD 29K register names” on page 93
- “AMD 29K floating point” on page 94
- “AMD 29K machine directives” on page 95
- “AMD 29K opcodes” on page 95

AMD 29K options

as has no additional command-line options for the AMD 29K family.

AMD 29K syntax

The following documentation discusses the features pertinent to the AMD 29K regarding the syntax for the GNU assembler.

AMD 29K macros

The macro syntax used on the AMD 29K is like syntax in the AMD 29K Family Macro Assembler Specification. Normal `as` macros should still work.

AMD 29K special characters

`;` is the line comment character.

The question mark character, `?`, is permitted in identifiers (but *may not begin* an identifier).

AMD 29K register names

General-purpose registers are represented by predefined symbols of the form `GR nnn` (for global registers) or `LR nnn` (for local registers), where nnn represents a number between 0 and 127, written with no leading zeros. The leading letters may be in either upper or lower case; for example, `gr13` and `LR7` are both valid register names.

You may also refer to general-purpose registers by specifying the register number as the result of an expression (prefixed with `%%` to flag the expression as a register number), as in the following example.

```
%% expression
```

expression in the previous example's case must be an absolute expression evaluating to a number between 0 and 255. The range [0, 127] refers to global registers, and the range [128, 255] to local registers.

In addition, `as` understands the protected special-purpose register names for the AMD 29K family, as in the following.

<code>vab</code>	<code>chd</code>	<code>pc0</code>
<code>ops</code>	<code>chc</code>	<code>pc1</code>
<code>cps</code>	<code>rbp</code>	<code>pc2</code>
<code>cfg</code>	<code>tmc</code>	<code>mmu</code>
<code>cha</code>	<code>tmr</code>	<code>lru</code>

These unprotected special-purpose register names are also recognized:

<code>ipc</code>	<code>alu</code>	<code>fpe</code>
<code>ipa</code>	<code>bp</code>	<code>inte</code>
<code>ipb</code>	<code>fc</code>	<code>fps</code>
<code>q</code>	<code>cr</code>	<code>exop</code>

AMD 29K floating point

The AMD 29K family uses IEEE floating-point numbers.

AMD 29K machine directives

The following documentation discusses the features pertinent to the AMD 29K regarding the machine directives for the GNU assembler.

`.block size, fill`

This directive emits *size* bytes, each of value *fill*. Both *size* and *fill* are absolute expressions. If the comma and *fill* are omitted, *fill* is assumed to be zero. In other versions of the GNU assembler, this directive is called `.space`.

`.cputype`

This directive is ignored; it is accepted for compatibility with other AMD 29K assemblers.

`.file`

This directive is ignored; it is accepted for compatibility with other AMD 29K assemblers.

WARNING: In other versions of the GNU assembler, `.file` is used for the directive called `.app-file` in the AMD 29K support.

`.line`

This directive is ignored; it is accepted for compatibility with other AMD 29K assemblers.

`.sect`

This directive is ignored; it is accepted for compatibility with other AMD 29K assemblers.

`.use section name`

Establishes the section and subsection for the following code; section name may be one of `.text`, `.data`, `.data1`, or `.lit`. With one of the first three *section name* options, `.use` is equivalent to the machine directive, *section name*; the remaining case, `.use .lit`, is the same as `.data 200`.

AMD 29K opcodes

as implements all the standard AMD 29K opcodes. No additional pseudo-instructions are needed on this family.

For information on the 29K machine instruction set, see *AMD 29000 User's Manual*, published by Advanced Micro Devices, Inc.

ARC dependent features

The following documentation discusses the features pertinent to the ARC processor regarding the GNU assembler.

- “ARC command-line options” on page 98
- “ARC floating point” on page 98
- “ARC machine directives” on page 99

ARC command-line options

The ARC chip family includes several successive levels (or other variants) of chip, using the same core instruction set, but including a few additional instructions at each level.

By default, `as` assumes the core instruction set (ARC base). The `.cpu` pseudo-op is intended to be used to select the variant.

`-mbig-endian`

`-mlittle-endian`

Any ARC configuration of `as` can select big-endian or little-endian output at run time (unlike most other GNU development tools, which must be configured for one or the other). Use `'-mbig-endian'` to select big-endian output, and `'-mlittle-endian'` for little-endian.

ARC floating point

The ARC processor family currently does not have hardware floating point support. Software floating point support is provided by `gcc` and uses IEEE floating-point numbers.

ARC machine directives

The ARC version of `as` supports the following additional machine directive.

`.cpu`

This must be followed by the desired CPU. The ARC is intended to be customizable, `.cpu` is used to select the desired variant (although, currently, there are none).

ARM dependent features

The following documentation discusses the features pertinent to the ARM processor regarding the GNU assembler.

- “ARM options” on page 102
- “ARM syntax” on page 103
- “ARM special characters” on page 103
- “ARM floating point” on page 103
- “ARM machine directives” on page 104
- “ARM opcodes” on page 104

ARM options

`as` has no additional command-line options for the ARM processor family.

ARM syntax

as has no additional command-line options for the ARM processor family.

ARM special characters

‘;’ is the line comment character.

ARM floating point

The ARM family uses IEEE floating-point numbers.

ARM machine directives

The following assembler directives are for the ARM processor.

`.code [16| 32]`

This directive selects the instruction set being generated.

The value 16 selects Thumb, with the value 32 selecting ARM.

`.thumb`

This performs the same action as `.code 16`.

`.arm`

This performs the same action as `.code 32`.

ARM opcodes

as implements all the standard ARM opcodes.

For information on the ARM or Thumb instruction sets, see *ARM Software Development Toolkit Reference Manual*, Advanced RISC Machines, Ltd.

AT&T/Intel dependent features

The following documentation discusses the features pertinent to the AT&T/Intel processors regarding the GNU assembler.

- “AT&T syntax versus Intel syntax” on page 106
- “AT&T opcode naming” on page 106
- “80386 register naming” on page 107
- “80386 opcode prefixes” on page 108
- “AT&T memory references” on page 109
- “Handling of jump instructions for 80386 machines” on page 109
- “AT&T floating point” on page 110
- “Writing 16-bit Code for 80386 machines” on page 111
- “Notes on AT& T and the 80386 instructions” on page 111

NOTE: The 80386 has no machine dependent options.

AT&T syntax versus Intel syntax

In order to maintain compatibility with the output of GCC, `as` supports AT&T System V/386 assembler syntax. This is quite different from Intel syntax. We mention these differences because almost all 80386 documents used only Intel syntax. Notable differences between the two syntaxes are:

- AT&T immediate operands are preceded by `$`; Intel immediate operands are undelimited (Intel `push 4` is AT&T `pushl $4`). AT&T register operands are preceded by `%`; Intel register operands are undelimited. AT&T absolute (as opposed to PC relative) `jump/call` operands are prefixed by `*`; they are undelimited in Intel syntax.
- AT&T and Intel syntax use the opposite order for source and destination operands. Intel `add eax, 4` is AT&T `addl $4, %eax`. The *source, dest* convention is maintained for compatibility with previous Unix assemblers.
- In AT&T syntax the size of memory operands is determined from the last character of the opcode name. Opcode suffixes of `b`, `w`, and `l` specify byte (8-bit), word (16-bit), and long (32-bit) memory references. Intel syntax accomplishes this by prefixes memory operands (*not* the opcodes themselves) with `byte ptr`, `word ptr`, and `dword ptr`. Thus, Intel `mov al, byte ptr foo` is AT&T `movb foo, %al` in AT&T syntax.
- Immediate form long jumps and calls are `lcall/ljmp $ section, $offset` in AT&T syntax; the Intel syntax is `call/jmp far section: offset`. Also, the far return instruction is `lret $stack-adjust` in AT&T syntax; Intel syntax is `ret far stack-adjust`.
- The AT&T assembler does not provide support for multiple section programs. Unix style systems expect all programs to be single sections.

AT&T opcode naming

Opcode names are suffixed with one character modifiers that specify the size of operands.

The letters `b`, `w`, and `l` specify byte, word, and long operands. If no suffix is specified by an instruction and it contains no memory operands then `as` tries to fill in the missing suffix based on the destination register operand (the last one by convention).

Thus, `mov %ax, %bx` is equivalent to `movw %ax, %bx`; also, `mov $1, %bx` is equivalent to `movw $1, %bx`.

NOTE: This is incompatible with the AT&T Unix assembler which assumes that a missing opcode suffix implies long operand size. (This incompatibility does not affect compiler output since compilers

always explicitly specify the opcode suffix.)

Almost all opcodes have the same names in AT&T and Intel format. There are a few exceptions.

The sign extend and zero extend instructions need two sizes to specify them. They need a size to sign/zero extend from and a size to zero extend *to*. This is accomplished by using two opcode suffixes in AT&T syntax.

Base names for `sign extend` and `zero extend` are `movs...` and `movz...` in AT&T syntax (`movsx` and `movzx` in Intel syntax).

The opcode suffixes are tacked on to this base name, the *from* suffix before the *to* suffix. Thus, `movsbl %al, %edx` is AT&T syntax for “move sign extend *from* `%al` *to* `%edx`.” Possible suffixes, thus, are `b1` (from byte to long), `bw` (from byte to word), and `w1` (from word to long).

The following Intel-syntax conversion instructions are called `cbtw`, `cwtl`, `cwtd`, and `cltd` in AT&T naming. `as` accepts either naming for the following instructions.

- `cbw`
sign-extend byte in `%al` to word in `%ax`
- `cwde`
sign-extend word in `%ax` to long in `%eax`
- `cwd`
sign-extend word in `%ax` to long in `%dx:%ax`
- `cdq`
sign-extend dword in `%eax` to quad in `%edx:%eax`

Far call/jump instructions are `lcall` and `ljump` in AT&T syntax, but `call far` and `jump far` in Intel convention.

80386 register naming

Register operands are always prefixed with `%`. The 80386 registers consist of the following register operands.

- **8 32-bit registers:**
`%eax` (the accumulator), `%ebx`, `%ecx`, `%edx`, `%edi`, `%esi`, `%ebp` (the frame pointer), and `%esp` (the stack pointer).
- **8 16-bit low-ends:**
`%ax`, `%bx`, `%cx`, `%dx`, `%di`, `%si`, `%bp`, and `%sp`.
- **8 8-bit registers**
`%ah`, `%al`, `%bh`, `%bl`, `%ch`, `%cl`, `%dh`, and `%dl` (these are the high-bytes and low-bytes of `%ax`, `%bx`, `%cx`, and `%dx`)
- **6 section registers:**
`%cs` (*code section*), `%ds` (*data section*), `%ss` (*stack section*), `%es`, `%fs`, and `%gs`.

- *3 processor control registers:*
%cr0, %cr2, and %cr3.
- *6 debug registers:*
%db0, %db1, %db2, %db3, %db6, and %db7.
- *2 test registers:*
%tr6 and %tr7.
- *8 floating point register stack:*
%st, or, equivalently, %st(0), %st(1), %st(2), %st(3), %st(4), %st(5), %st(6), and %st(7).

80386 opcode prefixes

Opcode prefixes are used to modify the following opcode. They are used to repeat string instructions, to provide section overrides, to perform bus lock operations, and to give operand and address size (16-bit operands are specified in an instruction by prefixing what would normally be 32-bit operands with a *operand size* opcode prefix). Opcode prefixes are usually given as single-line instructions with no operands, and must directly precede the instruction they act upon. For example, the ‘scas’ (scan string) instruction is repeated with the following declaration.

```
repne  
scas
```

The following is a list of opcode prefixes:

- *Section override prefixes:*
cs, ds, ss, es, fs and gs. These are automatically added by using the *section: memory-operand* form for memory references.
- *Operand/Address size prefixes:*
data16 and addr16 change 32-bit operands/addresses into 16-bit operands/addresses.

NOTE: 16-bit addressing modes are not supported (i.e., 8086 and 80286 addressing modes).
- The *bus lock prefix*, lock, inhibits interrupts during execution of the instruction it precedes. (This is only valid with certain instructions; see a 80386 manual for details).
- The *wait for coprocessor prefix*, wait, waits for the coprocessor to complete the current instruction. This should never be needed for the 80386/80387 combination.
- The rep, repe, and repne prefixes are added to string instructions to make them repeat %ecx times.

AT&T memory references

An Intel syntax indirect memory reference of the following form translates into the AT&T syntax like the following declaration.

```
section:[base + index*scale + disp]
```

It then outputs like the following example.

```
section:disp(base, index, scale)
```

base and *index* are the optional 32-bit base and index registers, *disp* is the optional displacement, and *scale*, taking the values 1, 2, 4, and 8, multiplies *index* to calculate the address of the operand. If no *scale* is specified, *scale* is taken to be 1. *section* specifies the optional section register for the memory operand, and may override the default section register (see a 80386 manual for section register defaults). Section overrides in AT&T syntax *must* have be preceded by a %. If you specify a section override which coincides with the default section register, *as* does *not* output any section register override prefixes to assemble the given instruction. Thus, section overrides can be specified to emphasize which section register is used for a given memory operand.

Here are some examples of Intel and AT&T style memory references:

AT&T: `'-4(%ebp)'`, Intel: `'[ebp - 4]'`

base is `'%ebp'`; *disp* is `'-4'`. *section* is missing, and the default section is used (`'%ss'` for addressing with `'%ebp'` as the base register). *index*, *scale* are both missing.

AT&T: `'foo(,%eax,4)'`, Intel: `'[foo + eax*4]'`

index is `'%eax'` (scaled by a *scale*4); *disp* is `'foo'`. All other fields are missing. The section register here defaults to `'%ds'`.

AT&T: `'foo(,1)'`; Intel `'[foo]'`

This uses the value pointed to by `'foo'` as a memory operand. Note that *base* and *index* are both missing, but there is only *one* `','`. This is a syntactic exception.

AT&T: `'%gs:foo'`; Intel `'gs:foo'`

This selects the contents of the variable `'foo'` with section register *section* being `'%gs'`.

Absolute (as opposed to PC relative) call and jump operands must be prefixed with `*`. If no `*` is specified, *as* always chooses PC relative addressing for jump/call labels.

Any instruction that has a memory operand *must* specify its size (byte, word, or long) with an opcode suffix (`b`, `w`, or `l`, respectively).

Handling of jump instructions for 80386 machines

Jump instructions are always optimized to use the smallest possible displacements.

This is accomplished by using byte (8-bit) displacement jumps whenever the target is sufficiently close. If a byte displacement is insufficient a long (32-bit) displacement is used. We do not support word (16-bit) displacement jumps (i.e., prefixing the jump instruction with the `addr16` opcode prefix), since the 80386 insists upon masking `%eip` to 16 bits after the word displacement is added.

NOTE: ‘`jcxz`’, ‘`jecxz`’, ‘`loop`’, ‘`loopz`’, ‘`loope`’, ‘`loopnz`’ and ‘`loopne`’ instructions only come in byte displacements, so that if you use these instructions (GCC does not use them) you may get an error message (and incorrect code). The AT&T 80386 assembler tries to get around this problem by expanding `jcxz foo` to the following.

```
                jcxz cx_zero
                jmp  cx_nonzero
cx_zero:        jmp  foo
cx_nonzero:
```

AT&T floating point

All 80387 floating point types except packed BCD are supported. (BCD support may be added without much difficulty).

These data types are 16-, 32-, and 64- bit integers, and single (32-bit), double (64-bit), and extended (80-bit) precision floating point. Each supported type has an opcode suffix and a constructor associated with it. Opcode suffixes specify operand’s data types. Constructors build these data types into memory.

■ ***Floating point constructors*** are:

`.float` or `.single`, `.double`, and `.tfloat` for 32-, 64-, and 80-bit formats. These correspond to opcode suffixes, `s`, `l`, and `t`. `t` stands for *temporary real*, and that the 80387 only supports this format via the `fldt` (*load temporary real to stack top*) and `fstpt` (*store temporary real and pop stack*) instructions.

■ ***Integer constructors*** are:

`.word`, `.long` or `.int`, and `.quad` for the 16-, 32-, and 64-bit integer formats.

The corresponding *opcode suffixes* are `s` (single), `l` (long), and `q` (quad). As with the *temporary real* format, the 64-bit `q` format is only present in the `fildq` (*load quad integer to stack top*) and `fistpq` (*store quad integer and pop stack*) instructions.

Register to register operations do not require opcode suffixes, so that the statement, `fst %st, %st(1)`, is equivalent to `fstl %st, %st(1)`.

Since the 80387 automatically synchronizes with the 80386, `fwait` instructions are almost never needed (this is not the case for the 80286/80287 and 8086/8087 combinations). Therefore, `as` suppresses the `fwait` instruction whenever it is implicitly selected by one of the `fn` instructions. For example, `fsave` and `fnsave`

are treated identically. In general, all the `fn` the instructions are made equivalent to `f` the instructions. If `fwait` is desired, it must be explicitly coded.

Writing 16-bit Code for 80386 machines

While GAS normally writes only *pure* 32-bit i386 code, it has limited support for writing code to run in real mode or in 16-bit protected mode code segments. To do this, insert a `.code16` directive before the assembly language instructions to be run in 16-bit mode. You can switch GAS back to writing normal 32-bit code with the `.code32` directive.

GAS understands exactly the same assembly language syntax in 16-bit mode as in 32-bit mode. The function of any given instruction is exactly the same regardless of mode, as long as the resulting object code is executed in the mode for which GAS wrote it. So, for example, the `ret` mnemonic produces a 32-bit return instruction regardless of whether it is to be run in 16-bit or 32-bit mode. (If GAS is in 16-bit mode, it will add an operand size prefix to the instruction to force it to be a 32-bit return.)

This means, for one thing, that you can use GNU CC to write code to be run in real mode or 16-bit protected mode. Just insert the statement `asm(".code16");` at the beginning of your C source file, and while GNU CC will still be generating 32-bit code, GAS will automatically add all the necessary size prefixes to make that code run in 16-bit mode. Of course, since GNU CC only writes small-model code (it doesn't know how to attach segment selectors to pointers like native x86 compilers do), any 16-bit code you write with GNU CC will essentially be limited to a 64K address space. Also, there will be a code size and performance penalty due to all the extra address and operand size prefixes GAS has to add to the instructions.

NOTE: Placing GAS in 16-bit mode does not mean that the resulting code will necessarily run on a 16-bit pre-80386 processor. To write code that runs on such a processor, you would have to refrain from using *any* 32-bit constructs which require GAS to output address or operand size prefixes. At the moment this would be rather difficult, because GAS currently supports *only* 32-bit addressing modes: when writing 16-bit code, it *always* outputs address size prefixes for any instruction that uses a non-register addressing mode. So you can write code that runs on 16-bit processors, but only if that code never references memory.

Notes on AT&T and the 80386 instructions

There is some trickery concerning the `mul` and `imul` instructions that deserves mention. The 16-, 32-, and 64-bit expanding multiplies (base opcode `0xf6`; extension

4 for `mul` and 5 for `imul`) can be output only in the one operand form. Thus, `imul %ebx, %eax` does *not* select the expanding multiply; the expanding multiply would clobber the `%edx` register, and this would confuse gcc output. Use `imul %ebx` to get the 64-bit product in `%edx:%eax`.

We have added a two operand form of `imul` when the first operand is an immediate mode expression and the second operand is a register. This is just a shorthand, so that, multiplying `%eax` by 69, for example, can be done with `imul $69, %eax` rather than `imul $69, %eax, %eax`.

D10V dependent features

The following documentation discusses the features pertinent to the D10V processor regarding the GNU assembler.

- “D10V options” on page 114
- “D10V syntax” on page 115
- “D10V size modifiers” on page 115
- “D10V sub-instructions” on page 115
- “D10V special characters” on page 116
- “D10V register names” on page 117
- “D10V addressing modes” on page 118
- “@WORD modifier for D10V” on page 118
- “D10V floating point” on page 118
- “D10V opcodes” on page 119

D10V options

The Mitsubishi D10V version of `as` uses the following machine dependent option.

-O

The D10V can often execute two sub-instructions in parallel.

When this option is used, `as` will attempt to optimize its output by detecting when instructions can be executed in parallel.

D10V syntax

The D10V syntax is based on the syntax in the Mitsubishi document, *D10V Architecture: A VLIW Mediaprocessor for Multimedia Applications* (Mitsubishi Electric Corporation, version 1.00; October 24, 1996). The differences are detailed in the following discussions.

D10V size modifiers

The D10V version of `as` uses the instruction names in the Mitsubishi document, *D10V Architecture: A VLIW Mediaprocessor for Multimedia Applications* (Mitsubishi Electric Corporation, version 1.00; October 24, 1996).. However, the names in the documentation are sometimes ambiguous. There are instruction names that can assemble to a short or long form opcode. How does the assembler pick the correct form? `as` will always pick the smallest form if it can. When dealing with a symbol that is not defined yet when a line is being assembled, it will always use the long form. If you need to force the assembler to use either the short or long form of the instruction, you can append either `‘.s’` (short) or `‘.l’` (long) to it. For example, if you are writing an assembly program and you want to do a branch to a symbol that is defined later in your program, you can write `‘bra.s foo’`. `objdump` and `gdb` will always append `‘.s’` or `‘.l’` to instructions which have both short and long forms.

D10V sub-instructions

The D10V assembler takes as input a series of instructions, either one-per-line, or in the special two-per-line format described in the next section. Some of these instructions will be short-form or sub-instructions. These sub-instructions can be packed into a single instruction. The assembler will do this automatically. It will also detect when it should not pack instructions. For example, when a label is defined, the next instruction will never be packaged with the previous one. Whenever a branch and link instruction is called, it will not be packaged with the next instruction so the return address will be valid. Nops are automatically inserted when necessary.

If you do not want the assembler automatically making these decisions, you can control the packaging and execution type (parallel or sequential) with the special execution symbols described in the following documentation; see “D10V special characters” on page 116.

D10V special characters

‘;’ and ‘#’ are the line comment characters. Sub-instructions may be executed in order, in reverse-order, or in parallel. Instructions listed in the standard one-per-line format will be executed sequentially. To specify the executing order, use the following symbols.

->
Sequential with instruction on the left first.

<-
Sequential with instruction on the right first.

||
Parallel.

The D10V syntax allows either one instruction per line, one instruction per line with the execution symbol, or two instructions per line, as the following documentation discusses.

```
abs a1 -> abs r0
```

Execute these sequentially. The instruction on the right is in the right container and is executed second.

```
abs r0 <- abs a1
```

Execute these reverse-sequentially. The instruction on the right is in the right container, and is executed first.

```
ld2w r2,@r8+ || mac a0,r0,r7
```

Execute these in parallel.

```
ld2w r2,@r8+ ||  
mac a0,r0,r7
```

Two-line format. Execute these in parallel.

```
ld2w r2,@r8+  
mac a0,r0,r7
```

Two-line format. Execute these sequentially. Assembler will put them in the proper containers.

```
ld2w r2,@r8+ ->  
mac a0,r0,r7
```

Two-line format. Execute these sequentially. Same as previously described instructions except second instruction will always go into right container.

Since ‘\$’ has no special meaning, you can use it in symbol names.

D10V register names

You can use the predefined symbols ‘r0’ through ‘r15’ to refer to the D10V registers. You can also use ‘sp’ as an alias for ‘r15’. The accumulators are ‘a0’ and ‘a1’. There are special register-pair names that may optionally be used in opcodes that require even-numbered registers. Register names are not case sensitive.

r0-r1
r2-r3
r4-r5
r6-r7
r8-r9
r10-r11
r12-r13
r14-r15

The D10V also has predefined symbols for the following control registers and status bits.

psw	Processor Status Word
bpsw	Backup Processor Status Word
pc	Program Counter
bpc	Backup Program Counter
rpt_c	Repeat Count
rpt_s	Repeat Start address
rpt_e	Repeat End address
mod_s	Modulo Start address
od_e	Modulo End address
iba	Instruction Break Address
f0	Flag 0
f1	Flag 1
c	Carry flag

D10V addressing modes

as understands the following addressing modes for the D10V. *Rn* in the following documentation refers to any of the numbered registers, but not the control registers.

Rn

Register direct

@*Rn*

Register indirect

@*Rn*+

Register indirect with post-increment

@*Rn*-

Register indirect with post-decrement

@-SP

Register indirect with pre-decrement

@(*disp*, *Rn*)

Register indirect with displacement

addr

PC relative address (for branch or rep).

#*imm*

Immediate data (the '#' is optional and ignored)

@WORD modifier for D10V

Any symbol followed by **@word** will be replaced by the symbol's value shifted right by 2. This is used in situations such as loading a register with the address of a function (or any other code fragment). For example, if you want to load a register with the location of the function `main` then jump to that function, you could do it as follows.

```
ldi r2, main @WORD
jmp r2
```

D10V floating point

The D10V has no hardware floating point, but the `.float` and `.double` directives generates IEEE floating-point numbers for compatibility with other development tools.

D10V opcodes

For detailed information on the D10V machine instruction set, see the Mitsubishi document, *D10V Architecture: A VLIW Mediaprocessor for Multimedia Applications* (published by Mitsubishi Electric Corporation, version 1.00; October 24, 1996); `as` implements all the standard D10V opcodes found there. The only changes are those described in the section on size modifiers.

H8/300 dependent features

The following documentation discusses the features pertinent to the H8/300 processor regarding the GNU assembler.

- “H8/300 options” on page 122
- “H8/300 syntax” on page 123
- “H8/300 register names” on page 123
- “H8/300 addressing modes” on page 123
- “H8/300 floating point” on page 124
- “H8/300 machine directives” on page 125
- “H8/300 opcodes” on page 126

H8/300 options

`as` has no additional command-line options for the Hitachi H8/300 family.

H8/300 syntax

The following documentation discusses the features pertinent to the H8/300 regarding the syntax for the GNU assembler.

Special characters

`;` is the line comment character.

`$` can be used instead of a newline to separate statements. Therefore, *you may not use \$ in symbol names* on the H8/300.

H8/300 register names

You can use predefined symbols of the form `rnn` and `rnl` to refer to the H8/300 registers as sixteen 8-bit general-purpose registers. *n* is a digit from 0 to 7; for instance, both `r0h` and `r7l` are valid register names.

You can also use the eight predefined symbols `rn` to refer to the H8/300 registers as 16-bit registers (you must use this form for addressing).

On the H8/300H, you can also use the eight predefined symbols `ern` (`er0` . . . `er7`) to refer to the 32-bit general purpose registers.

The two control registers are called `pc` (program counter; a 16-bit register, except on the H8/300H where it is 24 bits) and `ccr` (condition code register; an 8-bit register). `r7` is used as the stack pointer, and can also be called `sp`.

H8/300 addressing modes

`as` understands the following addressing modes for the H8/300.

`rn`

Register direct

`@rn`

Register indirect

`@(d, rn)`

`@(d:16, rn)`

`@(d:24, rn)`

Register indirect: 16-bit or 24-bit displacement, *d*, from register, *n*. (24-bit displacements are only meaningful on the H8/300H.)

`@rn+`

Register indirect with post-increment.

`@-rn`

Register indirect with pre-decrement.

`@aa`

`@aa:8`

`@aa:16`

`@aa:24`

Absolute address `aa`. (The address size `:24` only makes sense on the H8/300H.)

`#xx`

`#xx:8`

`#xx:16`

`#xx:32`

Immediate data `xx`. You may specify the `:8`, `:16`, or `:32` for clarity, if you wish; but `as` neither requires this nor uses it—the data size required is taken from context.

`@@aa`

`@@aa:8`

Memory indirect. You may specify the `:8` for clarity, if you wish; but `as` neither requires this nor uses it.

H8/300 floating point

The H8/300 family has no hardware floating point, but the `.float` directive generates IEEE floating-point numbers for compatibility with other development tools.

H8/300 machine directives

`as` has only the following machine-dependent directive for the H8/300.

`.h8300h`

Recognize and emit additional instructions for the H8/300H variant, and also make `.int` emit 32-bit numbers rather than the usual (16-bit) for the H8/300 family.

On the H8/300 family (including the H8/300H) `.word` directives generate 16-bit numbers.

H8/300 opcodes

For detailed information on the H8/300 machine instruction set, see *H8/300 Series Programming Manual* (Hitachi ADE-602-025).

For information specific to the H8/300H, see *H8/300H Series Programming Manual* (Hitachi).

`as` implements all the standard H8/300 opcodes. No additional pseudo-instructions are needed on this family.

Four H8/300 instructions (`add`, `cmp`, `mov`, `sub`) are defined with variants using the suffixes `.b`, `.w`, and `.l` to specify the size of a memory operand.

`as` supports these suffixes, but does not require them; since one of the operands is always a register, `as` can deduce the correct size.

For instance, since `r0` refers to a 16-bit register, note the distinctions in the following examples.

```
mov    r0,@foo
```

```
mov    w r0,@foo
```

If you use the size suffixes, `as` issues a warning when the suffix and the register size do not match.

H8/500 dependent features

The following documentation discusses the features pertinent to the H8/500 processor regarding the GNU assembler.

- “H8/500 options” on page 128
- “H8/500 syntax” on page 129
- “H8/500 syntax” on page 129
- “H8/500 special characters” on page 129
- “H8/500 register names” on page 129
- “H8/500 addressing modes” on page 129
- “H8/500 floating point” on page 130
- “H8/500 machine directives” on page 131
- “H8/500 opcodes” on page 131

H8/500 options

`as` has no additional command-line options for the Hitachi H8/500 family.

H8/500 syntax

The following documentation discusses the features pertinent to the H8/500 regarding the syntax for the GNU assembler.

H8/500 special characters

`!` is the line comment character.

`;` can be used instead of a newline to separate statements.

Since `$` has no special meaning, you may use it in symbol names.

H8/500 register names

You can use the predefined symbols `r0`, `r1`, `r2`, `r3`, `r4`, `r5`, `r6`, and `r7` to refer to the H8/500 registers.

The H8/500 also has the following control registers.

<code>cp</code>	code pointer
<code>dp</code>	data pointer
<code>bp</code>	base pointer
<code>tp</code>	stack top pointer
<code>ep</code>	extra pointer
<code>sr</code>	status register
<code>ccr</code>	condition code register

All registers are 16 bits long. To represent 32 bit numbers, use two adjacent registers; for distant memory addresses, use one of the segment pointers (`cp` for the program counter; `dp` for `r0-r3`; `ep` for `r4` and `r5`; and `tp` for `r6` and `r7`).

H8/500 addressing modes

`as` understands the following addressing modes for the H8/500:

`Rn`

Register direct.

`@Rn`

Register indirect.

`@(d:8, Rn)`

Register indirect with 8 bit signed displacement.

`@(d:16, Rn)`

Register indirect with 16 bit signed displacement.

`@-Rn`
Register indirect with pre-decrement.

`@Rn+`
Register indirect with post-increment.

`@aa:8`
8 bit absolute address.

`@aa:16`
16 bit absolute address.

`#xx:8`
8 bit immediate.

`#xx:16`
16 bit immediate.

H8/500 floating point

The H8/500 family has no hardware floating point, but the `.float` directive generates IEEE floating-point numbers for compatibility with other development tools.

H8/500 machine directives

`as` has no machine-dependent directives for the H8/500. However, on this platform the `.int` and `.word` directives generate 16-bit numbers.

H8/500 opcodes

For detailed information on the H8/500 machine instruction set, see *H8/500 Series Programming Manual* (Hitachi M21T001).

`as` implements all the standard H8/500 opcodes. No additional pseudo-instructions are needed on this family.

HPPA dependent features

The following documentation discusses the features pertinent to the HPPA processor regarding the GNU assembler.

- “HPPA options” on page 134
- “HPPA syntax” on page 135
- “HPPA floating point” on page 135
- “HPPA assembler directives” on page 135
- “HPPA opcodes” on page 138

As a back end for GNU CC, `as` has been thoroughly tested and should work extremely well for the HPPA targets. We have tested it only minimally on hand-written assembly code and no one has tested it much on the assembly output from the HP compilers.

The format of the debugging sections has changed since the original `as` port (version 1.3x) was released; therefore, you must rebuild all HPPA objects and libraries with the new assembler so that you can debug the final executable.

The HPPA `as` port generates a small subset of the relocations available in the SOM and ELF object file formats. Additional relocation support will be added as it becomes necessary.

HPPA options

`as` has no machine-dependent command-line options for the HPPA.

HPPA syntax

The assembler syntax closely follows the HPPA instruction set reference manual; assembler directives and general syntax closely follow the HPPA assembly language reference manual, with a few noteworthy differences.

First, a colon may immediately follow a label definition. This is simply for compatibility with how most assembly language programmers write code.

Some obscure expression parsing problems may affect hand written code which uses the `sop` instructions, or code which makes significant use of the `!` line separator.

`as` is much less forgiving about missing arguments and other similar oversights than the HP assembler. `as` notifies you of missing arguments as syntax errors; this is regarded as a feature, not a bug.

Finally, `as` allows you to use an external symbol without explicitly importing the symbol.

WARNING: In the future this allowance will be an error for HPPA targets.

Special characters for HPPA targets include the following.

`;` is the line comment character.

`!` can be used instead of a newline to separate statements.

Since `$` has no special meaning, you may use it in symbol names.

HPPA floating point

The HPPA family uses IEEE floating-point numbers.

HPPA assembler directives

`as` for the HPPA supports many additional directives for compatibility with the native assembler. The following documentation only briefly describes them. For detailed information on HPPA-specific assembler directives, see *HP9000 Series 800 Assembly Language Reference Manual* (HP 92432-90001).

`as` does *not* support the following assembler directives described in the HP manual:

```
.endm      listoff      macro
.enter     liston
.leave     locct
```

Beyond those implemented for compatibility, `as` supports one additional assembler directive for the HPPA: `.param`. It conveys register argument locations for static functions. Its syntax closely follows the `.export` directive.

The following are the additional directives in `as` for the HPPA:

- `.block n`
- `.blockz n`
Reserve *n* bytes of storage, and initialize them to zero.
- `.call`
Mark the beginning of a procedure call. Only the special case with no arguments is allowed.
- `.callinfo [param=value, ...][flag, ...]`
Specify a number of parameters and flags that define the environment for a procedure. *param* may be any of *frame* (frame size), *entry_gr* (end of general register range), *entry_fr* (end of float register range), *entry_sr* (end of space register range). The values for *flag* are *calls* or *caller* (proc has subroutines), *no_calls* (proc does not call subroutines), *save_rp* (preserve return pointer), *save_sp* (proc preserves stack pointer), *no_unwind* (do not unwind this proc), *hpux_int* (proc is interrupt routine).
- `.code`
Assemble into the standard section called `$TEXT$`, subsection `$CODE$`.
- `.copyright "string"`
In the SOM object format, insert *string* into the object code, marked as a copyright string.
- `.enter`
Not yet supported; the assembler rejects programs containing this directive.
- `.entry`
Mark the beginning of a procedure.
- `.exit`
Mark the end of a procedure.
- `.export name[,typ][,param=r]`
Make a procedure *name* available to callers. *typ*, if present, must be one of absolute, code (ELF only, not SOM), data, entry, data, entry, millicode, plabel, *pri_prog*, or *sec_prog*.
param, if present, provides either relocation information for the procedure arguments and result, or a privilege level. *param* may be *argw n* (where *n* ranges from 0 to 3, and indicates one of four one-word arguments); *rtnval* (the procedure's result); or *priv_lev* (privilege level). For arguments or the result, *r* specifies how to relocate, and must be one of *no* (not relocatable), *gr* (argument is in general register), *fr* (in floating point register), or 'fu' (upper half of float register). For *priv_lev*, *r* is an integer.
- `.half n`
Define a two-byte integer constant *n*; synonym for the portable `as` directive, `.short`.

```
.import name[,typ]
    Converse of .export; make a procedure available to call. The arguments use the
    same conventions as the first two arguments for .export.
.label name
    Define name as a label for the current assembly location.
.leave
    Not yet supported; the assembler rejects programs containing this directive.
.origin lc
    Advance location counter to lc. Synonym for the {No value for ``as''}
    portable directive .org.
.param name[,typ][,param=r]
    Similar to .export, but used for static procedures.
.proc
    Use preceding the first statement of a procedure.
.procend
    Use following the last statement of a procedure.
label.reg expr
    Synonym for .equ; define label with the absolute expression expr as its value.
.space secname[,params]
    Switch to section secname, creating a new section by that name if necessary. You
    may only use params when creating a new section, not when switching to an
    existing one. secname may identify a section by number rather than by name. If
    specified, the list params declares attributes of the section, identified by
    keywords. The keywords recognized are spnum=exp (identify this section by the
    number exp, an absolute expression), sort=exp (order sections according to this
    sort key when linking; exp is an absolute ex-pression), unloadable (section
    contains no loadable data), notdefined (this section defined elsewhere), and
    private (data in this section not available to other programs).
.spnum secnam
    Allocate four bytes of storage, and initialize them with the section number of the
    section named secnam. (You can define the section number with the HPPA .space
    directive.)
.string "str"
    Copy the characters in the string str to the object file. See “Strings” on page 38
    for information on escape sequences you can use in as strings.

WARNING:    The HPPA version of .string differs from the usual as definition: it
                does not write a zero byte after copying str.

.stringz "str"
    Like .string, but appends a zero byte after copying str to object file.
.subspa name[,params]
```

`.nsubspa name[,params]`

Similar to `.space`, but selects a subsection name within the current section. You may only specify `params` when you create a subsection (in the first instance of `.subspa` for this `name`). If specified, the list `params` declares attributes of the subsection, identified by keywords. The keywords recognized are `quad=expr` (“quadrant” for this subsection), `align=expr` (alignment for beginning of this subsection; a power of two), `access=expr` (value for “access rights” field), `sort=expr` (sorting order for this subspace in link), `code_only` (subsection contains only code), `unloadable` (subsection cannot be loaded into memory), `common` (subsection is common block), `dup_comm` (initialized data may have duplicate names), or `zero` (subsection is all zeros, do not write in object file).

`.nsubspa` always creates a new subspace with the given name, even if one with the same name already exists.

`.version "str"`

Write `str` as version identifier in object code.

HPPA opcodes

For detailed information on the HPPA machine instruction set, see *PA-RISC Architecture and Instruction Set Reference Manual* (published by Hewlett-Packard [HP 09740- 90039]).

SH dependent features

The following documentation discusses the features pertinent to the Hitachi SH processor regarding the GNU assembler.

- “Hitachi SH options” on page 140
- “Hitachi SH syntax” on page 141
- “Hitachi SH register names” on page 141
- “Hitachi SH addressing modes” on page 141
- “Hitachi SH floating point” on page 142
- “Hitachi SH machine directives” on page 143
- “Hitachi SH opcodes” on page 143

Hitachi SH options

`as` has no additional command-line options for the Hitachi SH family.

Hitachi SH syntax

The following documentation discusses the features pertinent to the Hitachi SH regarding the syntax for the GNU assembler.

Hitachi SH special characters

`!` is the line comment character.

You can use `;` instead of a newline to separate statements.

Since `$` has no special meaning, you may use it in symbol names.

Hitachi SH register names

You can use the predefined symbols `'r0'`, `'r1'`, `'r2'`, `'r3'`, `'r4'`, `'r5'`, `'r6'`, `'r7'`, `'r8'`, `'r9'`, `'r10'`, `'r11'`, `'r12'`, `'r13'`, `'r14'`, and `'r15'` to refer to the SH registers.

The SH also has the following control registers.

```
pr
    procedure register (holds return address)
pc
    program counter
mach
macl
    high and low multiply accumulator registers
sr
    status register
gbr
global
    base register
vbr
    vector base register (for interrupt vectors)
```

Hitachi SH addressing modes

`as` understands the following addressing modes for the SH. `Rn` in the following refers to any of the numbered registers, but not the control registers.

```
Rn
    Register direct.
@Rn
    Register indirect.
@-Rn
    Register indirect with pre-decrement.
```

@Rn+

Register indirect with post-increment.

@(disp, Rn)

Register indirect with displacement.

@(R0, Rn)

Register indexed.

@(disp, GBR)

GBR offset.

@(R0, GBR)

GBR indexed.

addr

@(disp, PC)

PC relative address (for branch or for addressing memory). The `as` implementation allows you to use the simpler form *addr* anywhere a PC relative address is called for; the alternate form is supported for compatibility with other assemblers.

#imm

Immediate data.

Hitachi SH floating point

The SH family has no hardware floating point, but the `.float` directive generates IEEE floating-point numbers for compatibility with other development tools.

Hitachi SH machine directives

`as` has no machine-dependent directives for the SH.

Hitachi SH opcodes

For detailed information on the SH machine instruction set, see *SH-Microcomputer User's Manual* (published by Hitachi Micro Systems, Inc.).

`as` implements all the standard SH opcodes. No additional pseudo-instructions are needed on this family. Note, however, that because `as` supports a simpler form of PC-relative addressing, you may simply write the following (for example).

```
mov.l bar,r0
```

Other assemblers might require an explicit displacement to `bar` from the program counter:

```
mov.l @(disp, PC)
```


Intel 960 dependent features

The following documentation discusses the features pertinent to the Intel 960 processor regarding the GNU assembler.

- “i960 command-line options” on page 146
- “i960 floating point” on page 147
- “i960 machine directives” on page 148
- “i960 opcodes” on page 148
- “The callj instruction” on page 149.
- “i960 Compare-and-Branch” on page 150

i960 command-line options

`-ACA | -ACA_A | -ACB | -ACC | -AKA | -AKB | -AKC | -AMC`

Select the 80960 architecture. Instructions or features not supported by the selected architecture cause fatal errors.

`-ACA` is equivalent to `-ACA_A`; `-AKC` is equivalent to `-AMC`.

Synonyms are provided for compatibility with other tools.

If you do not specify any of these options, `as` generates code for any instruction or feature that is supported by *some* version of the i960 (even if this means mixing architectures!).

In principle, `as` attempts to deduce the minimal sufficient processor type if none is specified; depending on the object code format, the processor type may be recorded in the object file. If it is critical that the `as` output match a specific architecture, specify that architecture explicitly.

`-b`

Add code to collect information about conditional branches taken, for later optimization using branch prediction bits. (The conditional branch instructions have branch prediction bits in the CA, CB, and CC architectures.) If `BR` represents a conditional branch instruction, the following represents the code generated by the assembler when `-b` is specified.

```

                call      increment  routine
                word      0          # pre-counter
Label:  BR
        call  increment                routine
        .word      0          # post-counter
```

The counter following a branch records the number of times that branch was *not* taken; the difference between the two counters is the number of times the branch *was* taken.

A table of every such `Label` is also generated, so that the external postprocessor `gbr960` (supplied by Intel) can locate all the counters. This table is always labeled `__BRANCH_TABLE__`; this is a local symbol to permit collecting statistics for many separate object files. The table is word aligned, and begins with a two-word header. The first word, initialized to 0, is used in maintaining linked lists of branch tables. The second word is a count of the number of entries in the table, which follow immediately: each is a word, pointing to one of the labels illustrated in the previous example.

```

*NEXT      COUNT: N      *BRLAB 1      ...      *BRLAB N
          __BRANCH_TABLE__ layout
```

The first word of the header is used to locate multiple branch tables, since each object file may contain one. Normally the links are maintained with a call to an initialization routine, placed at the beginning of each function in the file. The GNU C compiler generates these calls automatically when you give it a `-b` option. For further details, see the documentation of `gbr960`.

`-no-relax`

Normally, Compare-and-Branch instructions with targets that require displacements greater than 13 bits (or that have external targets) are replaced with the corresponding compare (or `chkbit`) and branch instructions. You can use the `-no-relax` option to specify that `as` should generate errors instead, if the target displacement is larger than 13 bits.

This option does not affect the Compare-and-Jump instructions; the code emitted for them is *always* adjusted when necessary (depending on displacement size), regardless of whether you use `-no-relax`.

i960 floating point

`as` generates IEEE floating-point numbers for the directives `.float`, `.double`, `.extended`, and `.single`.

i960 machine directives

The following documentation discusses the features pertinent to the i960 regarding the machine directives for the GNU assembler.

`.bss symbol, length, align`

Reserve *length* bytes in the bss section for a local *symbol*, aligned to the power of two specified by *align*. *length* and *align* must be positive absolute expressions. This directive differs from `.lcomm` only in that it permits you to specify an alignment. See “.lcomm *symbol*, *length* directive” on page 75.

`.extended flonums`

`.extended` expects zero or more flonums, separated by commas; for each flonum, `.extended` emits an IEEE extended-format (80-bit) floating-point number.

`.leafproc call-lab, bal-lab`

You can use the `.leafproc` directive in conjunction with the optimized `callj` instruction to enable faster calls of leaf procedures. If a procedure is known to call no other procedures, you may define an entry point that skips procedure prolog code (and that does not depend on system-supplied saved context), and declare it as the *bal-lab* using `.leafproc`. If the procedure also has an entry point that goes through the normal prolog, you can specify that entry point as *call-lab*.

A `.leafproc` declaration is meant for use in conjunction with the optimized `call` instruction `callj`; the directive records the data needed later to choose between converting the `callj` into a `bal` or a `call`.

call-lab is optional; if only one argument is present, or if the two arguments are identical, the single argument is assumed to be the *bal* entry point.

`.sysproc name, index`

The `.sysproc` directive defines a name for a system procedure. After you define it using `.sysproc`, you can use *name* to refer to the system procedure identified by *index* when calling procedures with the optimized `call` instruction `callj`.

Both arguments are required; *index* must be between 0 and 31 (inclusive).

i960 opcodes

All Intel 960 machine instructions are supported; see “i960 command-line options” on page 146 for a discussion of selecting the instruction subset for a particular 960 architecture. Some opcodes are processed beyond simply emitting a single corresponding instruction: `callj`, and Compare-and-Branch or Compare-and- Jump instructions with target displacements larger than 13 bits.

The `callj` instruction

You can write `callj` to have the assembler or the linker determine the most appropriate form of subroutine call: `call`, `bal`, or `calls`. If the assembly source contains enough information—a `.leafproc` or `.sysproc` directive defining the operand—then `as` translates the `callj`; if not, it simply emits the `callj`, leaving it for the linker to resolve.

i960 Compare-and-Branch

The 960 architectures provide combined *Compare-and-Branch* instructions that permit you to store the branch target in the lower 13 bits of the instruction word itself. However, if you specify a branch target far enough away that its address won't fit in 13 bits, the assembler can either issue an error, or convert your Compare-and-Branch instruction into separate instructions to do the compare and the branch.

Whether `as` gives an error or expands the instruction depends on two choices you can make: whether you use the `-no-relax` option, and whether you use a “*Compare and Branch*” instruction or a “*Compare and Jump*” instruction. The “*Jump*” instructions are *always* expanded if necessary; the “*Branch*” instructions are expanded when necessary *unless* you specify `-no-relax` in which case `as` gives an error instead.

The following example shows the Compare-and-Branch instructions, their “Jump” variants, and the instruction pairs into which they may expand.

<i>Compare and</i>		
<i>Branch</i>	<i>Jump</i>	<i>Expanded to</i>
<code>bbc</code>		<code>chkbit; bno</code>
<code>bbs</code>		<code>chkbit; bo</code>
<code>cmpibe</code>	<code>cmpije</code>	<code>cmpi; be</code>
<code>cmpibg</code>	<code>cmpijg</code>	<code>cmpi; bg</code>
<code>cmpibge</code>	<code>cmpijge</code>	<code>cmpi; bge</code>
<code>cmpibl</code>	<code>cmpijl</code>	<code>cmpi; bl</code>
<code>cmpible</code>	<code>cmpijle</code>	<code>cmpi; ble</code>
<code>cmpibno</code>	<code>cmpijno</code>	<code>cmpi; bno</code>
<code>cmpibne</code>	<code>cmpijne</code>	<code>cmpi; bne</code>
<code>cmpibo</code>	<code>cmpijo</code>	<code>cmpi; bo</code>
<code>cmpobe</code>	<code>cmpoje</code>	<code>cmpo; be</code>
<code>cmpobg</code>	<code>cmpojg</code>	<code>cmpo; bg</code>
<code>cmpobge</code>	<code>cmpojge</code>	<code>cmpo; bge</code>
<code>cmpobl</code>	<code>cmpojl</code>	<code>cmpo; bl</code>
<code>cmpoble</code>	<code>cmpojle</code>	<code>cmpo; ble</code>
<code>cmpobne</code>	<code>cmpojne</code>	<code>cmpo; bne</code>

M68K dependent features

The following documentation discusses the features pertinent to the Motorola 68K series of processors regarding the GNU assembler.

- “M680x0 options” on page 152
- “M680x0 syntax” on page 154
- “Standard Motorola syntax differences” on page 155
- “M680x0 floating point” on page 156
- “M680x0 machine directives” on page 157
- “M680x0 opcodes” on page 157
- “M680x0 branch improvement” on page 157
- “Pseudo-ops for M68K branch instructions” on page 158
- “M680x0 special characters” on page 159

M680x0 options

The Motorola 680x0 version of `as` has a few machine dependent options.

You can use the `-l` option to shorten the size of references to undefined symbols. If you do not use the `-l` option, references to undefined symbols are wide enough for a full `long` (32 bits). (Since `as` cannot know where these symbols end up, `as` can only allocate space for the linker to fill in later. Since `as` does not know how far away these symbols are, it allocates as much space as it can.) If you use this option, the references are only one word wide (16 bits). This may be useful if you want the object file to be as small as possible, and you know that the relevant symbols are always less than 17 bits away.

For some configurations, especially those where the compiler normally does not prepend an underscore to the names of user variables, the assembler requires a `%` before any use of a register name. This is intended to let the assembler distinguish between C variables and functions named `a0` through `a7`, and so on. The `%` is always accepted, but is not required for certain configurations, notably `sun3`. The `--register-prefix-optional` option may be used to permit omitting the `%` even for configurations for which it is normally required. If this is done, it will generally be impossible to refer to C variables and functions with the same names as register names.

`as` can assemble code for several different members of the Motorola 680x0 family. The default depends upon how `as` was configured when it was built; normally, the default is to assemble code for the 68020 microprocessor. The following options may be used to change the default. These options control which instructions and addressing modes are permitted. The members of the 680x0 family are very similar. For detailed information about the differences, see the Motorola manuals.

`-m68000`

`-m68008`

`-m68302`

Assemble for the 68000. `-m68008` and `-m68302` are synonyms for `-m68000`, since the chips are the same from the point of view of the assembler.

`-m68010`

Assemble for the 68010.

`-m68020`

Assemble for the 68020. This is normally the default.

`-m68030`

Assemble for the 68030.

`-m68040`

Assemble for the 68040.

`-m68060`

Assemble for the 68060.

`-mcpu32`

`-m68331`

`-m68332`

`-m68333`

`-m68340`

`-m68360`

Assemble for the CPU32 family of chips.

`-m68881`

`-m68882`

Assemble 68881 floating point instructions. This is the default for the 68020, 68030, and the CPU32. The 68040 and 68060 always support floating point instructions.

`-mno-68881`

Do not assemble 68881 floating point instructions. This is the default for 68000 and the 68010. The 68040 and 68060 always support floating point instructions, even if this option is used.

`-m68851`

Assemble 68851 MMU instructions. This is the default for the 68020, 68030, and 68060. The 68040 accepts a somewhat different set of MMU instructions;

`-m68851` and `-m68040` should not be used together.

`-mno-68851`

Do not assemble 68851 MMU instructions. This is the default for the 68000, 68010, and the CPU32. The 68040 accepts a somewhat different set of MMU instructions.

This syntax for the Motorola 680x0 was developed at MIT.

The 680x0 version of `as` uses instructions names and syntax compatible with the Sun assembler. Intervening periods are ignored; for example, `movl` is equivalent to `mov.l`.

M680x0 syntax

The following documentation discusses the M680x0 syntax.

apc stands for any of the address registers (%a0 through %a7), the program counter (%pc), the zero-address relative to the program counter (%zpc), a suppressed address register (%za0 through %za7), or it may be omitted entirely. The use of *size* means one of w or l, and it may be omitted, along with the leading colon, unless a scale is also specified. The use of *scale* means one of 1, 2, 4, or 8, and it may always be omitted along with the leading colon.

The following addressing modes are understood.

Immediate

#*number*

Data Register

%d0 through %d7

Address Register

%a0 through %a7

%a7 is also known as %sp (i.e., the stack pointer).

%a6 is also known as %fp (i.e., the frame pointer).

Address Register Indirect

%a0@ through %a7@

Address Register Postincrement

%a0@+ through %a7@+

Address Register Predecrement

%a0@- through %a7@-

Indirect Plus Offset

apc@(*number*)

Index

apc@(*number*, *register*:*size*:*scale*)

The *number* may be omitted.

Postindex

apc@(*number*)@(*onumber*, *register*:*size*:*scale*)

The *onumber* or the *register*, but not both, may be omitted.

Preindex

apc@(*number*, *register*:*size*:*scale*)@(*onumber*)

The *number* may be omitted. Omitting the register produces the *Postindex* addressing mode.

Absolute

symbol, or *digits*, optionally followed by :b, :w, or :l.

Standard Motorola syntax differences

The standard Motorola syntax for this chip differs from the syntax already discussed (see “Syntax” on page 33). *as* can accept Motorola syntax for operands, even if MIT syntax is used for other operands in the same instruction.

The two kinds of syntax are fully compatible.

In the following, *apc* stands for any of the address registers (%a0 through %a7), the program counter (%pc), the zero-address relative to the program counter (%zpc), or a suppressed address register (%za0 through %za7). The use of *size* means one of *w* or *l*, and it may always be omitted along with the leading dot. The use of *scale* means one of 1, 2, 4, or 8, and it may always be omitted along with the leading asterisk. The following additional addressing modes are understood.

Address Register Indirect

%a0 through %a7

%a7 is also known as %sp (i.e., the stack pointer).

%a6 is also known as %fp (i.e., the frame pointer).

Address Register Postincrement

(%a0)+ through (%a7)+

Address Register Predecrement

-(%a0) through -(%a7)

Indirect Plus Offset

number(%a0) through *number*(%a7), or *number*(%pc).

The *number* may also appear within the parentheses, as in (*number*, %a0). When used with the *pc*, the *number* may be omitted (with an address register, omitting the *number* produces Address Register Indirect mode).

Index

number(*apc*, *register.size*scale*)

The *number* may be omitted, or it may appear within the parentheses. The *apc* may be omitted. The *register* and the *apc* may appear in either order. If both *apc* and *register* are address registers, and the *size* and *scale* are omitted, then the first register is taken as the base register, and the second as the index register.

Postindex

([*number*, *apc*], *register.size*scale*, *onumber*)

The *onumber*, or the *register*, or both, may be omitted. Either the *number* or the *apc* may be omitted, but not both.

Preindex

([*number*, *apc*, *register.size*scale*], *onumber*)

The *number*, or the *apc*, or the *register*, or any two of them, may be omitted. The *onumber* may be omitted. The *register* and the *apc* may appear in either order. If

both *apc* and *register* are address registers, and the *size* and *scale* are omitted, then the first register is taken as the base register, and the second as the index register.

M680x0 floating point

Packed decimal (P) format floating literals are not supported. Feel free to add the code!

The floating point formats generated by directives are the following.

`.float`

single precision floating point constants.

`.double`

double precision floating point constants.

`.extend`

`.ldouble`

extended precision (long double) floating point constants.

M680x0 machine directives

In order to be compatible with the Sun assembler, the 680x0 assembler understands the following directives.

- `.data1`
This directive is identical to a `.data 1` directive.
- `.data2`
This directive is identical to a `.data 2` directive.
- `.even`
This directive is a special case of the `.align` directive; it aligns the output to an even byte boundary.
- `.skip`
This directive is identical to a `.space` directive.

M680x0 opcodes

The following documentation discusses the features pertinent to the M680x0 regarding the opcodes for the GNU assembler.

M680x0 branch improvement

Certain pseudo opcodes are permitted for branch instructions. They expand to the shortest branch instruction that reach the target. Generally these mnemonics are made by substituting `j` for `b` at the start of a Motorola mnemonic. The following summarizes the pseudo-operations.

Table 1: Pseudo-ops for Motorola M68K

Pseudo-Op	Displacement				
	BYTE	WORD	LONG	LONG	non-PC relative
jbsr	bsrs	bsr	bsrl	jsr	jsr
jra	bras	bra	bral	jmp	jmp
* jXX	bXXs	bXX	bXXl	bNXs; jmp	bNXs; jmp
* dbXX	dbXX	dbXX	dbXX; bra; jmp		
* fjXX	fbXXw	fbXXw	fbXXl	fbNXw; jmp	

XX: condition
NX: negative of condition XX

Pseudo-ops (as shown with asterisks in Table 1: Pseudo-ops for Motorola M68K) have more specific discussion with the documentation, “Pseudo-ops for M68K branch instructions” on page 158.

Pseudo-ops for M68K branch instructions

The following documentation discusses the specific requirements for pseudo-ops for the M68K branch instructions.

jbsr
jra

These are the simplest jump pseudo-operations; they always map to one particular machine instruction, depending on the displacement to the branch target.

JXX

Here, `jxx` stands for an entire family of pseudo-operations, where `xx` is a conditional branch or condition-code test. The full list of pseudo-ops in this family is the following:

```
jhi    jls    jcc    jcs    jne    jeq    jvc
jvs    jpl    jmi    jge    jlt    jgt    jle
```

For the cases of non-PC relative displacements and long displacements on the 68000 or 68010, `as` issues a longer code fragment in terms of `nx`, the opposite condition to `xx`.

```
jXX foo
```

For the non-PC relative case in the previous example, `as` gives the following output.

```
bNXs oof
jmp foo
oof:
```

dbXX

The full family of pseudo-operations covered here is the following.

```
dbhi    dbls    dbcc    dbcs    dbne    dbeq    dbvc
dbvs    dbpl    dbmi    dbge    dblt    dbgt    dble
dbf     dbra    dbt
```

Other than for word and byte displacements, when the source reads `dbXX foo`, `as` emits the following output.

```
dbXX ool
bra oo2
ool:    jmp1 foo
oo2:
```

fjXX

This family includes the following.

```
fjne    fjeq    fjge    fjlt    fjgt    fjle    fjf
fjt     fjgl    fjgle    fjnge    fjngl    fjngle    fjngt
fjnle    fjnlt    fjoge    fjogl    fjogt    fjole    fjolt
fjor     fjseq    fjstf    fjstne    fjst    fjueq    fjuge
fjugt    fjule    fjult    fjun
```

For branch targets that are not PC relative, `as` emits the following output when it

```
encounters fjXX foo.  
        fbNX oof  
        jmp foo  
oof:
```

M680x0 special characters

The immediate character is # for Sun compatibility. The line-comment character is |. If a # appears at the beginning of a line, it is treated as a comment unless it looks like # *line file*, in which case it is treated normally.

MIPS dependent features

The following documentation discusses the features pertinent to the MIPS processors regarding the GNU assembler.

- “Assembler options for MIPS” on page 162
- “MIPS ECOFF object code” on page 163
- “Directives for debugging information for MIPS” on page 164
- “Directives to override the ISA level for MIPS” on page 165
- “Directives for extending MIPS 16 bit instructions” on page 166
- “Directive to mark data as an instruction for MIPS” on page 167
- “Directives to save and restore options for MIPS” on page 168

MIPS processors that have support are the MIPS R2000, R3000, R4000 and R6000. For information about the MIPS instruction set, see *MIPS RISC Architecture*, by Kane and Heindrich (Prentice-Hall). For an overview of MIPS assembly conventions, see “Appendix D: Assembly Language Programming” in the same work.

Assembler options for MIPS

The MIPS configurations of GNU `as` support the following special options:

`-G num`

This option sets the largest size of an object that can be referenced implicitly with the `gp` register. It is only accepted for targets that use ECOFF format. The default value is 8.

`-EB`

`-EL`

Any MIPS configuration of `as` can select big-endian or little-endian output at run time (unlike the other GNU development tools, which must be configured for one or the other). Use `-EB` to select big-endian output, and `-EL` for little-endian.

`-mips1`

`-mips2`

`-mips3`

Generate code for a particular MIPS Instruction Set Architecture level. `-mips1` corresponds to the R2000 and R3000 processors, `-mips2` to the R6000 processor, and `-mips3` to the R4000 processor. You can also switch instruction sets during the assembly; see “Directives to override the ISA level for MIPS” on page 165.

`-m4650`

`-no-m4650`

Generate code for the MIPS R4650 chip. This tells the assembler to accept the `mad` and `madu` instruction, and to not schedule `nop` instructions around accesses to the `HI` and `LO` registers. `-no-m4650` turns off this option.

`-m4010`

`-no-m4010`

Generate code for the LSI R4010 chip. This tells the assembler to accept the R4010 specific instructions (`addciu`, `ffc`, etc.), and to not schedule `nop` instructions around accesses to the `HI` and `LO` registers. `-no-m4010` turns off this option.

`-mcpu=CPU`

Generate code for a particular MIPS CPU. This has little effect on the assembler, but it is passed by `gcc`.

`-nocpp`

This option is ignored. It is accepted for command-line compatibility with other assemblers, which use it to turn off C style preprocessing. With GNU `as`, there is no need for `-nocpp`, because the GNU assembler itself never runs the C preprocessor.

`--trap`

`--no-break`

`as` automatically macro expands certain division and multiplication instructions to

check for overflow and division by zero. This option causes `as` to generate code to take a trap exception rather than a break exception when an error is detected. The trap instructions are only supported at Instruction Set Architecture level 2 and higher.

```
--break
--no-trap
```

Generate code to take a break exception rather than a trap exception when an error is detected. This is the default.

MIPS ECOFF object code

Assembling for a MIPS ECOFF target supports some additional sections besides the usual `.text`, `.data` and `.bss`. The additional sections are `.rdata`, used for read-only data, `.sdata`, used for small data, and `.sbss`, used for small common objects.

When assembling for ECOFF, the assembler uses the `$gp` (`$28`) register to form the address of a *small object*. Any object in the `.sdata` or `.sbss` sections is considered “small” in this sense. For external objects, or for objects in the `.bss` section, you can use the `gcc` ‘`-G`’ option to control the size of objects addressed via `$gp`; the default value is 8, meaning that a reference to any object eight bytes or smaller uses `$gp`.

Passing ‘`-G 0`’ to `as` prevents it from using the `$gp` register on the basis of object size (but the assembler uses `$gp` for objects in `.sdata` or `.sbss` in any case). The size of an object in the `.bss` section is set by the `.comm` or `.lcomm` directive that defines it. The size of an external object may be set with the `.extern` directive. For example, `.extern sym, 4` declares that the object at `sym` is 4 bytes in length, while leaving `sym` otherwise undefined.

Using small ECOFF objects requires linker support, and assumes that the `$gp` register is correctly initialized (normally done automatically by the startup code). MIPS ECOFF assembly code must not modify the `$gp` register.

Directives for debugging information for MIPS

MIPS ECOFF `as` supports several directives used for generating debugging information which are not support by traditional MIPS assemblers. These are `.def`, `.endef`, `.dim`, `.file`, `.scl`, `.size`, `.tag`, `.type`, `.val`, `.stabd`, `.stabn`, and `.stabs`. The debugging information generated by the three `.stab` directives can only be read by GDB, not by traditional MIPS debuggers (this enhancement is required to fully support C++ debugging). These directives are primarily used by compilers, not assembly language programmers!

Directives to override the ISA level for MIPS

GNU `as` supports an additional directive to change the MIPS Instruction Set Architecture level on the fly: `.set mipsn`. `n` should be a number from 0 to 3. A value from 1 to 3 makes the assembler accept instructions for the corresponding isa level, from that point on in the assembly. `.set mipsn` affects not only which instructions are permitted, but also how certain macros are expanded. `.set mips0` restores the ISA level to its original level: either the level you selected with command line options, or the default for your configuration. You can use this feature to permit specific R4000 instructions while assembling in 32-bit mode. Use this directive with care! Traditional MIPS assemblers do not support this directive.

Directives for extending MIPS 16 bit instructions

By default, MIPS 16 instructions are automatically extended to 32 bits when necessary. The directive `.set noautoextend` will turn this off. When `.set noautoextend` is in effect, any 32 bit instruction must be explicitly extended with the `.e` modifier (e.g., `li.e $4,1000`). The directive `.set autoextend` may be used to once again automatically extend instructions when necessary.

This directive is only meaningful when in MIPS 16 mode. Traditional MIPS assemblers do not support this directive.

Directive to mark data as an instruction for MIPS

The `.insn` directive tells `as` that the following data is actually instructions. This makes a difference in MIPS 16 mode: when loading the address of a label which precedes instructions, `as` automatically adds 1 to the value, so that jumping to the loaded address will do the right thing.

Directives to save and restore options for MIPS

The directives, `.set push` and `.set pop`, may be used to save and restore the current settings for all the options which are controlled by `.set`. The `.set push` directive saves the current settings on a stack. The `.set pop` directive pops the stack and restores the settings.

These directives can be useful inside an macro which must change an option such as the ISA level or instruction reordering but does not want to change the state of the code which invoked the macro.

Traditional MIPS assemblers do not support these directives.

SPARC dependent features

The following documentation discusses the features pertinent to the SPARCprocessors regarding the GNU assembler.

- “SPARC options” on page 170
- “SPARC floating point” on page 170
- “SPARC machine directives” on page 171

SPARC options

The SPARC chip family includes several successive levels (or other variants) of chip, using the same core instruction set, but including a few additional instructions at each level.

By default, `as` assumes the core instruction set (SPARC v6), but “bumps” the architecture level as needed: it switches to successively higher architectures as it encounters instructions that only exist in the higher levels.

If not configured for SPARC v9 (`sparc64-*-*`) `as` will not bump passed `sparclite` by default, an option must be passed to enable the v9 instructions.

`as` treats `sparclite` as being compatible with v8, unless an architecture is explicitly requested. SPARC v9 is always incompatible with `sparclite`.

`-Av6` | `-Av7` | `-Av8` | `-Asparclite` | `-Av9` | `-Av9a`

Use one of the `-A` options to select one of the SPARC architectures explicitly. If you select an architecture explicitly, `as` reports a fatal error if it encounters an instruction or feature requiring a higher level.

`-xarch=v8plus` | `-xarch=v8plusa`

For compatibility with the Solaris v9 assembler. These options are equivalent to `-Av9` and `-Av9a`, respectively.

`-bump`

Warn whenever it is necessary to switch to another level. If an architecture level is explicitly requested, `as` will not issue warnings until that level is reached, and will then bump the level as required (except between incompatible levels).

SPARC floating point

The SPARC uses IEEE floating-point numbers.

SPARC machine directives

The SPARC version of `as` supports the following additional machine directives:

`.align`

This must be followed by the desired alignment in bytes.

`.common`

This must be followed by a symbol name, a positive number, and `"bss"`. This behaves somewhat like `.comm`, but the syntax is different.

`.half`

This is functionally identical to `.short`.

`.proc`

This directive is ignored. Any text following it on the same line is also ignored.

`.reserve`

This must be followed by a symbol name, a positive number, and `"bss"`. This behaves somewhat like `.lcomm`, but the syntax is different.

`.seg`

This must be followed by `"text"`, `"data"`, or `"data1"`. It behaves like `.text`, `.data`, or `.data 1`.

`.skip`

This is functionally identical to the `.space` directive.

`.word`

On the SPARC, the `.word` directive produces 32 bit values, instead of the 16 bit values it produces on many other machines.

`.xword`

On the SPARC V9 processor, the `.xword` directive produces 64 bit values.

Vax dependent features

The following documentation discusses the features pertinent to the Vax processors regarding the GNU assembler.

- “Vax command-Line options” on page 174
- “Vax floating point” on page 175
- “Vax machine directives” on page 176
- “Vax opcodes” on page 176
- “Vax branch improvement” on page 176
- “Vax operands” on page 178
- “Unsupported Vax assembler properties” on page 178

Vax command-Line options

The Vax version of `as` accepts any of the following options, gives a warning message that the option was ignored and proceeds. These options are for compatibility with scripts designed for other people's assemblers.

`-D` (Debug)

`-S` (Symbol Table)

`-T` (Token Trace)

These are obsolete options used to debug old assemblers.

`-d` (Displacement size for JUMPs)

This option expects a number following the `-d`. Like options that expect filenames, the number may immediately follow the `-d` (old standard) or constitute the whole of the command line argument that follows `-d` (GNU standard).

`-V` (Virtualize Interpass Temporary File)

Some other assemblers use a temporary file. This option commanded them to keep the information in active memory rather than in a disk file. `as` always does this, so this option is redundant.

`-J` (JUMPify Longer Branches)

Many 32-bit computers permit a variety of branch instructions to do the same job. Some of these instructions are short (and fast) but have a limited range; others are long (and slow) but can branch anywhere in virtual memory. Often there are 3 flavors of branch: short, medium and long. Some other assemblers would emit short and medium branches, unless told by this option to emit short and long branches.

`-t` (Temporary File Directory)

Some other assemblers may use a temporary file, and this option takes a filename being the directory to site the temporary file. Since `as` does not use a temporary disk file, this option makes no difference. `-t` needs exactly one filename.

The Vax version of the assembler accepts two options when compiled for VMS. They are `-h`, and `++`. The `-h` option prevents `as` from modifying the symbol-table entries for symbols that contain lowercase characters (I think). The `++` option causes `as` to print warning messages if the `FILENAME` part of the object file, or any symbol name is larger than 31 characters.

The `++` option also inserts some code following the `_main` symbol so that the object file is compatible with Vax-11 "C".

Vax floating point

Conversion of flonums to floating point is correct, and compatible with previous assemblers. Rounding is towards zero if the remainder is exactly half the least significant bit. `D`, `F`, `G` and `H` floating point formats are understood. Immediate floating literals (e.g., `S`$6.9`) are rendered correctly. Again, rounding is towards zero in the boundary case. The `.float` directive produces `f` format numbers. The `.double` directive produces `d` format numbers.

Vax machine directives

The Vax version of the assembler supports four directives for generating Vax floating point constants. They are described in the following.

`.dfloat`

This expects zero or more flonums, separated by commas, and assembles Vax `d` format 64-bit floating point constants.

`.ffloat`

This expects zero or more flonums, separated by commas, and assembles Vax `f` format 32-bit floating point constants.

`.gfloat`

This expects zero or more flonums, separated by commas, and assembles Vax `g` format 64-bit floating point constants.

`.hfloat`

This expects zero or more flonums, separated by commas, and assembles Vax `h` format 128-bit floating point constants.

Vax opcodes

All DEC mnemonics are supported. Beware that `case...` instructions have exactly 3 operands. The dispatch table that follows the `case...` instruction should be made with `.word` statements. As far as we know, this is compatible with all Unix assemblers.

Vax branch improvement

Certain pseudo opcodes are permitted. They are for branch instructions. They expand to the shortest branch instruction that reaches the target. Generally these mnemonics are made by substituting `j` for `b` at the start of a DEC mnemonic.

This feature is included both for compatibility and to help compilers. If you do not need this feature, avoid these opcodes. What follows are the mnemonics, and the code into which they can expand.

`jbsb`

`Jsb` is already an instruction mnemonic, so we chose `jbsb`.

(byte displacement)

`bsbb...`

(word displacement)

`bsbw...`

(long displacement)

`jsb...`

```

jbr
  jr
  Unconditional branch.
  (byte displacement)
    brb...
  (word displacement)
    brw...
  (long displacement)
    jmp...
jCOND
  COND may be any one of the conditional branches neq, nequ, eql, eqlu, gtr, geq,
  lss, gtru, lequ, vc, vs, gequ, cc, lssu, cs. COND may also be one of the bit tests
  bs, bc, bss, bcs, bsc, bcc, bssi, bcci, lbs, lbc. NOTCOND is the opposite
  condition to COND.
  (byte displacement)
    bCOND...
  (word displacement)
    bNOTCOND foo ; brw...; foo:
  (long displacement)
    bNOTCOND foo ; jmp...; foo:
jacbX
  X may be one of b d f g h l w.
  (word displacement)
    OPCODE...
  (long displacement)
    OPCODE..., foo ;
    brb bar ; foo: jmp... ;
    bar:
jaobYYY
  YYY may be one of lss leq.
jsobZZZ
  ZZZ may be one of geq gtr.
  (byte displacement)
    OPCODE...
  (word displacement)
    OPCODE... , foo ;
    brb bar ;
    foo: brw destination ;
    bar:

```

```
(long displacement)
    OPCODE..., foo ;
    brb bar ;
    foo: jmp destination ;
    bar:
aobleq
aoblss
sobgeq
sobgtr
(byte displacement)
    OPCODE...
(word displacement)
    OPCODE... , foo ;
    brb bar ;
    foo: brw destination ;
    bar:
(long displacement)
    OPCODE..., foo ;
    brb bar ;
    foo: jmp destination ;
    bar:
```

Vax operands

The immediate character is \$ for Unix compatibility, not # as DEC writes it.

The indirect character is * for Unix compatibility, not @ as DEC writes it.

The displacement sizing character is ` (an accent grave) for Unix compatibility, not ^ (a circumflex) as DEC writes it. The letter preceding ` may have either case. g is not understood, but all other letters (b, i, l, s, w) are understood.

Register names understood are r0 r1 r2...r15 ap fp sp pc. Upper and lower case letters are equivalent. For instance

```
tstb *w`$4(r5)
```

Any expression is permitted in an operand. Operands are comma separated.

Unsupported Vax assembler properties

Vax bit fields can not be assembled with as. Someone can add the required code if they really need it.

Acknowledgments for the GNU assembler

The following individuals contributed to the development of the GNU assembler.

- Dean Elsner wrote the original GNU assembler for the VAX.
- Jay Fenlason maintained GAS for a while, adding support for GDB-specific debug information and the 68k series machines, most of the pre-processing pass, and extensive changes in `messages.c`, `input-file.c`, `write.c`.
- K. Richard Pixley maintained GAS for a while, adding various enhancements and many bug fixes, including merging support for several processors, breaking GAS up to handle multiple object file format back ends (including heavy rewrite, testing, an integration of the COFF and `b.out` back ends), adding configuration including heavy testing and verification of cross assemblers and file splits and renaming, converted GAS to strictly ANSI C including full prototypes, added support for m680[34]0 and `cpu32`, did considerable work on i960 including a COFF port (including considerable amounts of reverse engineering), a SPARC opcode file rewrite, DECstation, RS6000, and HP300/HPUX host ports, updated “`know`” assertions and made them work, much other reorganization, cleanup, and lint.
- Ken Raeburn wrote the high-level BFD interface code to replace most of the code in format-specific I/O modules.
- The original VMS support was contributed by David L. Kashtan. Eric Youngdale has done much work with it since.

- The Intel 80386 machine description was written by Eliot Dresselhaus.
- Minh Tran-Le at IntelliCorp contributed some AIX 386 support.
- The Motorola 88k machine description was contributed by Devon Bowen of Buffalo University and Torbjorn Granlund of the Swedish Institute of Computer Science.
- Keith Knowles at the Open Software Foundation wrote the original MIPS back end (`tc-mips.c`, `tc-mips.h`), and contributed Rose format support (which hasn't been merged in yet). Ralph Campbell worked with the MIPS code to support `a.out` format.
- Support for the Zilog Z8k and Hitachi H8/300 and H8/500 processors (`tc-z8k`, `tc-h8300`, `tc-h8500`), and IEEE 695 object file format (`obj-ieee`), was written by Steve Chamberlain of Cygnus. Steve also modified the COFF back end to use BFD for some low-level operations, for use with the H8/300 and AMD 29k targets.
- John Gilmore built the AMD 29000 support, added `.include` support, and simplified the configuration of which versions accept which directives. He updated the 68k machine description so that Motorola's opcodes always produced fixed-size instructions (e.g., `jsr`), while synthetic instructions remained shrinkable (`jsr`). John fixed many bugs, including true tested cross-compilation support, and one bug in relaxation that took a week and required the proverbial one-bit fix.
- Ian Lance Taylor of Cygnus merged the Motorola and MIT syntax for the 68k, completed support for some COFF targets (68k, i386 SVR3, and SCO Unix), added support for MIPS, ECOFF and ELF targets, and made a few other minor patches.
- Steve Chamberlain made `as` able to generate listings.
- Hewlett-Packard contributed support for the HP9000/300.
- Jeff Law wrote GAS and BFD support for the native HPPA object format (SOM) along with a fairly extensive HPPA testsuite (for both SOM and ELF object formats). This work was supported by both the Center for Software Science at the University of Utah and Cygnus.
- Support for ELF format files has been worked on by Mark Eichin of Cygnus (original, incomplete implementation for SPARC), Pete Hoogenboom and Jeff Law at the University of Utah (HPPA mainly), Michael Meissner of the Open Software Foundation (i386 mainly), and Ken Raeburn of Cygnus (SPARC, and some initial 64-bit support).
- Several engineers at Cygnus have also provided many small bug fixes and configuration enhancements. Many others have contributed large or small

bugfixes and enhancements. If you have contributed significant work and are not mentioned on this list, and want to be, let Cygnus know. Some of the history has been lost; Cygnus is not intentionally neglecting anyone's contributions.

Using Id

Copyright © 1991-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the documentation entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Free Software Foundation

59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

This documentation has been prepared by Cygnus Technical Publications; contact the Cygnus Technical Publications staff: doc@cygnus.com.

1

Overview of ld, the GNU linker

The GNU linker, `ld`, combines a number of object and archive files, relocates their data and ties up symbol references. Usually the last step in compiling a program is to run `ld`.

The following documentation discusses using the GNU linker.

- “Invocation of `ld`, the GNU linker” on page 189
- “Linker scripts” on page 203
- “`ld` machine dependent features” on page 239
- “BFD libraries” on page 243
- “MRI compatible script files for the GNU linker” on page 249

`ld` accepts Linker Command Language files written in a superset of AT&T’s *Link Editor Command Language* syntax, to provide explicit and total control over the linking process.

This version of `ld` uses the general purpose BFD libraries to operate on object files. This allows `ld` to read, combine, and write object files in many different formats—for example, COFF or `a.out`. Different formats may be linked together to produce any available kind of object file. See “BFD libraries” on page 243 for more information.

Aside from its flexibility, the GNU linker is more helpful than other linkers in providing diagnostic information. Many linkers abandon execution immediately upon encountering an error. Whenever possible, ld continues executing, allowing you to identify other errors (or, in some cases, to get an output file in spite of the error).

2

Invocation of ld, the GNU linker

The GNU linker ld is meant to cover a broad range of situations, and to be as compatible as possible with other linkers. As a result, you have many choices to control its behavior. The following documentation discusses using the GNU linker with such choices.

- “ld command line options” on page 190
- “ld environment variables” on page 202

See also “Linker scripts” on page 203.

ld command line options

The following example summarizes the options you can use on the ld command line.

```
ld [ -o output ] objfile : ::
  [ -Aarchitecture ] [ -b input-format ]
  [ -Bstatic ] [ -Bdynamic ] [ -Bsymbolic ]
  [ -c MRI-commandfile ] [ -d | -dc | -dp ]
  [ -defsym symbol=expression ]
  [ -dynamic-linker file ] [ -embedded-relocs ]
  [ -e entry ] [ -F ] [ -F format ]
  [ -format input-format ] [ -g ] [ -G size ]
  [ -help ] [ -i ] [ -l archive ] [ -Lsearchdir ]
  [ -M ] [ -Map mapfile ] [ -m emulation ]
  [ -N | -n ] [ -noinhibit-exec ] [ -no-keep-memory ]
  [ -oformat output-format ] [ -R filename ]
  [ -relax ] [ -retain-symbols-file filename ]
  [ -r | -Ur ] [ -rpath dir ] [ -rpath-link dir ]
  [ -S ] [ -s ] [ -soname name ] [ -shared ]
  [ -sort-common ] [ -stats ] [ -T commandfile ]
  [ -Ttext org ] [ -Tdata org ]
  [ -Tbss org ] [ -t ] [ -traditional-format ]
  [ -u symbol ] [ -V ] [ -v ] [ -verbose ] [ -version ]
  [ -warn-common ] [ -warn-constructors ] [ -warn-once ]
  [ -y symbol ] [ -X ] [ -x ]
  [ -( [ archives ] -) ]
  [ --start-group [ archives ] --end-group ]
  [ -split-by-reloc count ] [ -split-by-file ]
  [ --whole-archive ]
```

This plethora of command-line options may seem intimidating, but in actual practice few of them are used in any particular context. For instance, a frequent use of ld is to link standard Unix object files on a standard, supported Unix system. On such a system, to link a file, `hello.o`, use the following example's input.

```
ld -o output /lib/crt0.o hello.o -lc
```

This tells ld to produce a file called `output` as the result of linking the file `/lib/crt0.o` with `hello.o` and the library `libc.a`, which will come from the standard search directories. (See the discussion of search directories with the `-L` option.)

The command-line options to ld may be specified in any order, and may be repeated at will. Repeating most options with a different argument will either have no further effect, or override prior occurrences (those further to the left on the command line) of that specific option.

The exceptions—which may meaningfully be used more than once—are the

following: `-A`, `-b` (or its synonym, `-format`), `-defsym`, `-L`, `-l`, `-R`, `-u`, and `-` (or its synonym, `--start-group`).

The list of object files to be linked together, shown as *objfile...*, may follow, precede, or be mixed in with command-line options, except that an *objfile* argument may not be placed between an option and its argument.

Usually the linker is invoked with at least one object file, but you can specify other forms of binary input files using `‘-l’`, `‘-R’`, and the script command language. If no binary input files at all are specified, the linker does not produce any output, and issues the message ‘No input files’.

If the linker can not recognize the format of an object file, it will assume that it is a linker script. A script specified in this way augments the main linker script used for the link (either the default linker script or the one specified by using `‘-T’`). This feature permits the linker to link against a file which appears to be an object or an archive, but actually merely defines some symbol values, or uses `INPUT` or `GROUP` to load other objects.

For options whose names are a single letter, option arguments must either follow the option letter without intervening whitespace, or be given as separate arguments immediately following the option that requires them.

For options whose names are multiple letters, either one dash or two can precede the option name; for example, `--oformat` and `-oformat` are equivalent. Arguments to multiple-letter options must either be separated from the option name by an equals sign, or be given as separate arguments immediately following the option that requires them. For example, `--oformat srec` and `--oformat=srec` are equivalent. Unique abbreviations of the names of multiple-letter options are accepted.

`-Aarchitecture`

In the current release of `ld`, this option is useful only for the Intel 960 family of architectures. In that `ld` configuration, the *architecture* argument identifies the particular architecture in the 960 family, enabling some safeguards and modifying the archive-library search path. See “ld and the Intel 960 processors” on page 241 for details. Future releases of `ld` may support similar functionality for other architecture families.

`-b input-format`

`ld` may be configured to support more than one kind of object file. If your `ld` is configured this way, you can use the `‘-b’` option to specify the binary format for input object files that follow this option on the command line. Even when `ld` is configured to support alternative object formats, you don’t usually need to specify this, as `ld` should be configured to expect as a default input format the most usual format on each machine. *input-format* is a text string, the name of a particular format supported by the BFD libraries. (You can list the available binary formats

with `'objdump -i')` `'-format input-format'` has the same effect, as does the script command `TARGET`. See “BFD libraries” on page 243.

You may want to use this option if you are linking files with an unusual binary format. You can also use `'-b'` to switch formats explicitly (when linking object files of different formats), by including `'-b input-format'` before each group of object files in a particular format.

The default format is taken from the environment variable `GNUTARGET`. See “ld environment variables” on page 202. You can also define the input format from a script, using the command `TARGET`.

`-Bstatic`

Do not link against shared libraries. This is only meaningful on platforms for which shared libraries are supported.

`-Bdynamic`

Link against dynamic libraries. This is only meaningful on platforms for which shared libraries are supported. This option is normally the default on such platforms.

`-Bsymbolic`

When creating a shared library, bind references to global symbols to the definition within the shared library, if any. Normally, it is possible for a program linked against a shared library to override the definition within the shared library. This option is only meaningful on ELF platforms that support shared libraries.

`-c MRI-commandfile`

For compatibility with linkers produced by MRI, `ld` accepts script files written in an alternate, restricted command language; see “MRI compatible script files for the GNU linker” on page 249. Introduce MRI script files with the option, `'-c'`; use the `'-T'` option to run linker scripts written in the general-purpose `ld` scripting language. If `MRI-cmdfile` does not exist, `ld` looks for it in the directories specified by any `'-L'` options.

`-d`

`-dc`

`-dp`

These three options are equivalent; multiple forms are supported for compatibility with other linkers. They assign space to common symbols even if a relocatable output file is specified (with `'-r'`). The script command,

`FORCE_COMMON_ALLOCATION`, has the same effect.

`-defsym symbol=expression`

Create a global symbol in the output file, containing the absolute address given by *expression*. You may use this option as many times as necessary to define multiple symbols in the command line. A limited form of arithmetic is supported for the *expression* in this context: you may give a hexadecimal constant or the

name of an existing symbol, or use + and - to add or subtract hexadecimal constants or symbols. If you need more elaborate expressions, consider using the linker command language from a script (see “Simple assignments” on page 210).

NOTE: There should be no white space between *symbol*, the equals sign (=), and *expression*.

-dynamic-linker *file*

Set the name of the dynamic linker. This is only meaningful when generating dynamically linked ELF executables. The default dynamic linker is normally correct; don't use this unless you know what you are doing.

-embedded-relocs

This option is only meaningful when linking MIPS embedded PIC code, generated by the -membedded-pic option to the GNU compiler and assembler. It causes the linker to create a table which may be used at runtime to relocate any data which was statically initialized to pointer values. See the code in testsuite/ld-empic for details.

-e *entry*

Use *entry* as the explicit symbol for beginning execution of your program, rather than the default entry point. See “Setting the entry point” on page 207 for a discussion of defaults and other ways of specifying the entry point.

-F

-F *format*

Ignored. Some older linkers used this option throughout a compilation toolchain for specifying object-file format for both input and output object files.

The mechanisms ld uses for this purpose (the '-b' or '-format' options for input files, '-oformat' option or the TARGET command in linker scripts for output files, the GNUTARGET environment variable) are more flexible, but ld accepts the '-F' option for compatibility with scripts written to call the old linker.

-format *input-format*

Synonym for '-b *input-format*'.

-g

Ignored. Provided for compatibility with other tools.

-G *value*

-G *value*

Set the maximum size of objects to be optimized using the GP register to *size* under MIPS ECOFF. Ignored for other object file formats.

-help

Print a summary of the command-line options on the standard output and exit.

-I

Perform an incremental link (same as option '-r').

- `-lar`
Add archive file *archive* to the list of files to link. This option may be used any number of times. `ld` will search its path-list for occurrences of `libar.a` for every *archive* specified.
- `-Lsearchdir`
`-L searchdir`
Add path, *searchdir*, to the list of paths that `ld` will search for archive libraries and `ld` control scripts. You may use this option any number of times. The directories are searched in the order in which they are specified on the command line. Directories specified on the command line are searched before the default directories. All `-L` options apply to all `-l` options, regardless of the order in which the options appear. The default set of paths searched (without being specified with `-L`) depends on which emulation mode `ld` is using, and in some cases also on how it was configured. See “ld environment variables” on page 202. The paths can also be specified in a link script with the `SEARCH_DIR` command. Directories specified this way are searched at the point in which the linker script appears in the command line.
- `-M`
Print (to the standard output) a link map—diagnostic information about where symbols are mapped by `ld`, and information on global common storage allocation.
- `-Map mapfile`
Print to the file *mapfile* a link map—diagnostic information about where symbols are mapped by `ld`, and information on global common storage allocation.
- `-memulation`
`-m emulation`
Emulate the *emulation* linker. You can list the available emulations with the `--verbose` or `-V` options. The default depends on `ld`’s configuration.
- `-N`
Set the text and data sections to be readable and writable. Also, do not page-align the data segment. If the output format supports Unix style magic numbers, mark the output as `OMAGIC`.
- `-n`
Set the text segment to be read only, and mark the output as `NMAGIC` if possible.
- `-noinhibit-exec`
Retain the executable output file whenever it is still usable. Normally, the linker will not produce an output file if it encounters errors during the link process; it exits without writing an output file when it issues any error whatsoever.
- `-no-keep-memory`
`ld` normally optimizes for speed over memory usage by caching the symbol tables of input files in memory. This option tells `ld` to instead optimize for memory

usage, by rereading the symbol tables as necessary. This may be required if `ld` runs out of memory space while linking a large executable.

`-o output`

Use *output* as the name for the program produced by `ld`; if this option is not specified, the name `a.out` is used by default. The script command `OUTPUT` can also specify the output file name.

`-oformat output-format`

`ld` may be configured to support more than one kind of object file. If your `ld` is configured this way, you can use the `'-oformat'` option to specify the binary format for the output object file. Even when `ld` is configured to support alternative object formats, you don't usually need to specify this, as `ld` should be configured to produce as a default output format the most usual format on each machine.

output-format is a text string, the name of a particular format supported by the BFD libraries. (You can list the available binary formats with `'objdump -i'`.) The script command `OUTPUT_FORMAT` can also specify the output format, but this option overrides it. See “BFD libraries” on page 243.

`-R filename`

Read symbol names and their addresses from *filename*, but do not relocate it or include it in the output. This allows your output file to refer symbolically to absolute locations of memory defined in other programs.

`-relax`

An option with machine dependent effects. Currently this option is only supported on the H8/300 and the Intel 960. See “ld and the H8/300 processors” on page 240 and “ld and the Intel 960 processors” on page 241.

On some platforms, the `'-relax'` option performs global optimizations that become possible when the linker resolves addressing in the program, such as relaxing address modes and synthesizing new instructions in the output object file.

On platforms where this is not supported, `'-relax'` is accepted, but ignored.

`-retain-symbols-file filename`

Retain only the symbols listed in the file *filename*, discarding all others.

filename is simply a flat file, with one symbol name per line. This option is especially useful in environments (such as VxWorks) where a large global symbol table is accumulated gradually, to conserve runtime memory.

`'-retain-symbols-file'` does not discard undefined symbols, or symbols needed for relocations.

You may only specify `'-retain-symbols-file'` once in the command line. It overrides `'-s'` and `'-S'`.

`-rpath dir`

Add a directory to the runtime library search path. This is used when linking an

ELF executable with shared objects. All `-rpath` arguments are concatenated and passed to the runtime linker, which uses them to locate shared objects at runtime.

The `-rpath` option is also used when locating shared objects which are needed by shared objects explicitly included in the link; see the description of the `-rpath-link` option. If `-rpath` is not used when linking an ELF executable, the contents of the environment variable `LD_RUN_PATH` will be used if it is defined.

The `-rpath` option may also be used on SunOS. By default, on SunOS, the linker will form a runtime search path out of all the `-L` options it is given. If a `-rpath` option is used, the runtime search path will be formed exclusively using the `-rpath` options, ignoring the `-L` options. This can be useful when using `gcc`, which adds many `-L` options which may be on NFS mounted filesystems.

`-rpath-link DIR`

When using ELF or SunOS, one shared library may require another. This happens when an `ld -shared` link includes a shared library as one of the input files. When the linker encounters such a dependency when doing a non-shared, non-relocateable link, it will automatically try to locate the required shared library and include it in the link, if it is not included explicitly. In such a case, the `-rpath-link` option specifies the first set of directories to search. The `-rpath-link` option may specify a sequence of directory names either by specifying a list of names separated by colons, or by appearing multiple times. The linker uses the following search paths to locate required shared libraries.

- Any directories specified by `-rpath-link` options.
- Any directories specified by `-rpath` options. The difference between `-rpath` and `-rpath-link` is that directories specified by `-rpath` are included in the executable to use at runtime, whereas the `-rpath-link` is only effective at link time.
- On an ELF system, if the `-rpath` and `rpath-link` options were not used, search the contents of the environment variable, `LD_RUN_PATH`.
- On SunOS, if the `-rpath` option was not used, search any directories specified using `-L` options.
- For a native linker, the contents of the environment variable `LD_LIBRARY_PATH`.
- The default directories, normally `‘/lib’` and `‘/usr/lib’`.

If the required shared library is not found, the linker will issue a warning and continue with the link.

`-r`

Generate relocatable output—i.e., generate an output file that can in turn serve as input to `ld`. This is often called partial linking. As a side effect, in environments

that support standard Unix magic numbers, this option also sets the output file's magic number to `OMAGIC`. If this option is not specified, an absolute file is produced. When linking C++ programs, this option will not resolve references to constructors; to do that, use `-Ur`.

This option does the same thing as `-i`.

`-S`

Omit debugger symbol information (but not all symbols) from the output file.

`-s`

Omit all symbol information from the output file.

`-soname name`

When creating an ELF shared object, set the internal `DT SONAME` field to the specified name. When an executable is linked with a shared object which has a `DT SONAME` field, then when the executable is run the dynamic linker will attempt to load the shared object specified by the `DT SONAME` field rather than the using the file name given to the linker.

`-shared`

Create a shared library. This is currently only supported on ELF and SunOS platforms. On SunOS, the linker will automatically create a shared library if the `-e` option is not used and there are undefined symbols in the link.

`-sort-common`

Normally, when `ld` places the global common symbols in the appropriate output sections, it sorts them by size. First come all the one byte symbols, then all the two bytes, then all the four bytes, and then everything else. This is to prevent gaps between symbols due to alignment constraints. This option disables that sorting.

`-split-by-reloc count`

Tries to create extra sections in the output file so that no single output section in the file contains more than *count* relocations. This is useful when generating huge relocatable for downloading into certain real time kernels with the COFF object file format; since COFF cannot represent more than 65535 relocations in a single section. Note that this will fail to work with object file formats which do not support arbitrary sections. The linker will not split up individual input sections for redistribution, so if a single input section contains more than *count* relocations one output section will contain that many relocations.

`-split-by-file`

Similar to `-split-by-reloc` but creates a new output section for each input file.

`-stats`

Compute and display statistics about the operation of the linker, such as execution time and memory usage.

`-Tbss org`

`-Tdata org`

`-Ttext org`

Use *org* as the starting address for—respectively—the `bss`, `data`, or the `text` segment of the output file. *org* must be a single hexadecimal integer; for compatibility with other linkers, you may omit the leading ‘0x’ usually associated with hexadecimal values.

`-Tcommandfile`

`-T commandfile`

Read link commands from the file, *commandfile*. These commands replace ld’s default link script (rather than adding to it), so *commandfile* must specify everything necessary to describe the target format. If *commandfile* does not exist, ld looks for it in the directories specified by any preceding ‘-L’ options. Multiple ‘-T’ options accumulate.

`-t`

Print the names of the input files as ld processes them.

`-traditional-format`

For some targets, the output of ld is different in some ways from the output of some existing linker. This switch requests ld to use the traditional format instead. For example, on SunOS, ld combines duplicate entries in the symbol string table. This can reduce the size of an output file with full debugging information by over 30 percent. Unfortunately, the SunOS `dbx` program can not read the resulting program (`gdb` has no trouble). The ‘-traditional-format’ switch tells ld to not combine duplicate entries.

`-u symbol`

Force *symbol* to be entered in the output file as an undefined symbol. Doing this may, for example, trigger linking of additional modules from standard libraries. ‘-u’ may be repeated with different option arguments to enter additional undefined symbols.

`-Ur`

For anything other than C++ programs, this option is equivalent to ‘-r’: it generates relocatable output—i.e., an output file that can in turn serve as input to ld. When linking C++ programs, ‘-Ur’ does resolve references to constructors, unlike ‘-r’. It does not work to use ‘-Ur’ on files that were themselves linked with ‘-Ur’; once the constructor table has been built, it cannot be added to. Use ‘-Ur’ only for the last partial link, and ‘-r’ for the others.

`--verbose`

Display the version number for ld and list the linker emulations supported.
Display which input files can and cannot be opened.

-v

-V

Display the version number for ld. The -v option also lists the supported emulations.

-version

Display the version number for ld and exit.

-warn-common

Warn when a common symbol is combined with another common symbol or with a symbol definition. Unix linkers allow this somewhat sloppy practice, but linkers on some other operating systems do not. This option allows you to find potential problems from combining global symbols. Unfortunately, some C libraries use this practice, so you may get some warnings about symbols in the libraries as well as in your programs.

There are three kinds of global symbols, illustrated here by C examples:

■ `'int i = 1;'`

A definition, which goes in the initialized data section of the output file.

■ `'extern int i;'`

An undefined reference, which does not allocate space. There must be either a definition or a common symbol for the variable somewhere.

■ `'int i;'`

A common symbol. If there are only (one or more) common symbols for a variable, it goes in the uninitialized data area of the output file. The linker merges multiple common symbols for the same variable into a single symbol. If they are of different sizes, it picks the largest size. The linker turns a common symbol into a declaration, if there is a definition of the same variable.

The `'-warn-common'` option can produce the following five kinds of warnings. Each warning consists of a pair of lines: the first describes the symbol just encountered, and the second describes the previous symbol encountered with the same name. One or both of the two symbols will be a common symbol.

■ Turning a common symbol into a reference, because there is already a definition for the symbol.

```
file( section): warning: common of ' symbol'
                overridden by definition
file( section): warning: defined here
```

■ Turning a common symbol into a reference, because a later definition for the symbol is encountered. This is the same as the previous case, except that the symbols are encountered in a different order.

```
file( section): warning: definition of `symbol'
overriding common
file( section): warning: common is here
```

■ Merging a common symbol with a previous same-sized common symbol.

```
file( section): warning: multiple common
of `symbol'
file( section): warning: previous common is here
```

■ Merging a common symbol with a previous larger common symbol.

```
file( section): warning: common of `symbol'
overridden by larger common
file( section): warning: larger common is here
```

■ Merging a common symbol with a previous smaller common symbol. The following is the same as the previous case, except that the symbols are encountered in a different order.

```
file( section): warning: common of `symbol'
overriding smaller common
file( section): warning: smaller common is here
```

-warn-constructors

Warn if any global constructors are used. This is only useful for a few object file formats. For formats like COFF or ELF, the linker can not detect the use of global constructors.

-warn-once

Only warn once for each undefined symbol, rather than once per module which refers to it.

For each archive mentioned on the command line, include every object file in the archive in the link, rather than searching the archive for the required object files. This is normally used to turn an archive file into a shared library, forcing every object to be included in the resulting shared library.

-X

Delete all temporary local symbols. For most targets, this is all local symbols whose names begin with 'L'.

-x

Delete all local symbols.

-y symbol

Print the name of each linked file in which symbol appears. This option may be given any number of times. On many systems it is necessary to prepend an underscore. This option is useful when you have an undefined symbol in your link but don't know where the reference is coming from.

```
-(archives-)  
--start-group archives--end-group
```

The *archives* should be a list of archive files. They may be either explicit file names, or ‘-l’ options.

The specified archives are searched repeatedly until no new undefined references are created. Normally, an archive is searched only once in the order that it is specified on the command line. If a symbol in that archive is needed to resolve an undefined symbol referred to by an object in an archive that appears later on the command line, the linker would not be able to resolve that reference. By grouping the archives, they all be searched repeatedly until all possible references are resolved.

Using this option has a significant performance cost. Use it only when there are unavoidable circular references between two or more archives.

ld environment variables

You can change the behavior of `ld` with the environment variable, `GNUTARGET`.

`GNUTARGET` determines the input-file object format if you don't use `-b` (or its synonym, `'-format'`). Its value should be one of the BFD names for an input format (see “BFD libraries” on page 243). If there is no `GNUTARGET` in the environment, `ld` uses the natural format of the target. If `GNUTARGET` is set to `default`, then BFD attempts to discover the input format by examining binary input files; this method often succeeds, but there are potential ambiguities, since there is no method of ensuring that the magic number used to specify object-file formats is unique. However, the configuration procedure for BFD on each system places the conventional format for that system first in the search-list, so ambiguities are resolved in favor of convention.

Linker scripts

A *linker script* controls every link. Such a script is written in the linker command language. The main purpose of the linker script is to describe how the sections in the input files should be mapped into the output file, and to control the memory layout of the output file. However, when necessary, the linker script can also direct the linker to perform many other operations, using the linker commands.

The linker always uses a linker script. If you do not supply one yourself, the linker will use a default script that is compiled into the linker executable. You can use the `--verbose` command line option to display the default linker script. Certain command line options, such as `-r` or `-N`, will affect the default linker script. You may supply your own linker script by using the `-T` command line option. When you do this, your linker script will replace the default linker script.

You may also use linker scripts implicitly by naming them as input files to the linker, as though they were files to be linked. If the linker opens a file, which it can not recognize as an object file or as an archive file, it will try to read it as a linker script. If the file can not be parsed as a linker script, the linker will report an error. An implicit linker script will not replace the default linker script. Typically an implicit linker script would contain only the `INPUT`, `GROUP`, or `VERSION` commands.

Basic linker script concepts

We need to define some basic concepts and vocabulary in order to describe the linker script language.

The linker combines input files into a single output file. The output file and each input file are in a special data format known as an *object file format*. Each file is called an *object file*. The output file is often called an *executable*, but for our purposes we will also call it an object file. Each object file has, among other things, a list of *sections*. We sometimes refer to a section in an input file as an *input section*; similarly, a section in the output file is an *output section*.

Each section in an object file has a name and a size. Most sections also have an associated block of data, known as the *section contents*. A section may be marked as *loadable*, meaning that the contents should be loaded into memory when the output file is run. A section with no contents may be *allocatable*, which means that an area in memory should be set aside, but nothing in particular should be loaded there (in some cases this memory must be zeroed out). A section, which is neither loadable nor allocatable, typically contains some sort of debugging information.

Every loadable or allocatable output section has two addresses. The first is the *VMA*, or *virtual memory address*. This is the address the section will have when the output file is run. The second is the *LMA*, or *load memory address*. This is the address at which the section will be loaded. In most cases the two addresses will be the same. An example of when they might be different is when a data section is loaded into ROM, and then copied into RAM when the program starts up (this technique is often used to initialize global variables in a ROM based system). In this case the ROM address would be the LMA, and the RAM address would be the VMA.

You can see the sections in an object file by using the ‘`objdump`’ program with the ‘`-h`’ option.

Every object file also has a list of *symbols*, known as the *symbol table*. A symbol may be defined or undefined. Each symbol has a name, and each defined symbol has an address, among other information. If you compile a C or C++ program into an object file, you will get a defined symbol for every defined function and global or static variable. Every undefined function or global variable, which is referenced in the input file, will become an undefined symbol. You can see the symbols in an object file by using the ‘`nm`’ program, or by using the ‘`objdump`’ program with the ‘`-t`’ option.

Linker script format

Linker scripts are text files. You write a linker script as a series of commands. Each command is either a keyword, possibly followed by arguments or an assignment to a symbol. You may separate commands using semicolons. Whitespace is generally ignored.

Strings such as file or format names can normally be entered directly. If the file name contains a character such as a comma, which would otherwise serve to separate file names, you may put the file name in double quotes. There is no way to use a double quote character in a file name.

You may include comments in linker scripts just as in C, delimited by `/*` and `*/`. As in C, comments are syntactically equivalent to whitespace.

Simple linker script example

Many linker scripts are fairly simple. The simplest possible linker script has just one command: `SECTIONS`. You use the `SECTIONS` command to describe the memory layout of the output file. The `SECTIONS` command is a powerful command. Here we will describe a simple use of it. Let's assume your program consists only of code, initialized data, and uninitialized data. These will be in the `.text`, `.data`, and `.bss` sections, respectively. Let's assume further that these are the only sections, which appear in your input files.

For this example, let's say that the code should be loaded at address `0x10000`, and that the data should start at address `0x8000000`. The following linker script will do this function.

```
SECTIONS
{
    . = 0x10000;
    .text : { *(.text) }
    . = 0x8000000;
    .data : { *(.data) }
    .bss : { *(.bss) }
}
```

You write the `SECTIONS` command as the keyword `SECTIONS`, followed by a series of symbol assignments and output section descriptions enclosed in curly braces. The first line in the above example sets the special symbol `.`, which is the location counter. If you do not specify the address of an output section in some other way (other ways are described later), the address is set from the current value of the location counter. The location counter is then incremented by the size of the output section. The second line defines an output section, `.text`. The colon is required syntax, which may be ignored for now. Within the curly braces after the output section name, you list the names of the input sections, which should be placed into this output section. The `*` is a wildcard which matches any file name. The expression `*(.text)` means all `.text` input sections in all input files.

Since the location counter is `0x10000` when the output section `.text` is defined, the linker will set the address of the `.text` section in the output file to be `0x10000`. The remaining lines define the `.data` and `.bss` sections in the output file. The `.data` output section will be at address `0x8000000`. When the `.bss` output section is defined, the value of the location counter will be `0x8000000` plus the size of the `.data` output section. The effect is that the `.bss` output section will follow immediately after the `.data` output section in memory.

That's it! That's a simple and complete linker script.

Simple linker script commands

In the following documentation, the discussion describes the simple linker script commands. See also “ld command line options” on page 190 and “BFD libraries” on page 243.

Setting the entry point

The first instruction to execute in a program is called the *entry point*. You can use the ‘ENTRY’ linker script command to set the entry point. The argument is a symbol name:

```
ENTRY (symbol)
```

There are several ways to set the entry point. The linker will set the entry point by trying each of the following methods in order, and stopping when one of them succeeds:

- The ‘-e’ *entry* command-line option;
- The ‘ENTRY (*symbol*)’ command in a linker script;
- The value of the symbol, *start*, if defined;
- The address of the first byte of the ‘.text’ section, if present;
- The address ‘0’.

Commands dealing with files

Several linker script commands deal with files. See also “ld command line options” on page 190 and “BFD libraries” on page 243.

```
INCLUDE filename
```

Include the linker script *filename* at this point. The file will be searched for in the current directory, and in any directory specified with the ‘-L’ option. You can nest calls to ‘INCLUDE’ up to 10 levels deep.

```
INPUT (file, file, ...)
```

```
INPUT (file file ...)
```

The ‘INPUT’ command directs the linker to include the named files in the link, as though they were named on the command line.

For example, if you always want to include ‘*subr.o*’ any time you do a link, but you can’t be bothered to put it on every link command line, then you can put ‘INPUT (*subr.o*)’ in your linker script. In fact, if you like, you can list all of your input files in the linker script, and then invoke the linker with nothing but a ‘-T’ option. The linker will first try to open the file in the current directory. If it is not found, the linker will search through the archive library search path. See the description of ‘-L’. If you use ‘INPUT (*-lfile*)’, ld will transform the name to

`'libfile.a'`, as with the command line argument `'-l'`. When you use the `'INPUT'` command in an implicit linker script, the files will be included in the link at the point at which the linker script file is included. This can affect archive searching.

`GROUP(FILE, FILE, ...)`

`GROUP (file file ...)`

The `'GROUP'` command is like `'INPUT'`, except that the named files should all be archives, and they are searched repeatedly until no new undefined references are created.

`OUTPUT (filename)`

The `'OUTPUT'` command names the output file. Using `'OUTPUT(FILENAME)'` in the linker script is exactly like using `'-o filename'` on the command line. If both are used, the command line option takes precedence.

You can use the `'OUTPUT'` command to define a default name for the output file other than the usual default of `'a.out'`.

`SEARCH_DIR (path)`

The `'SEARCH_DIR'` command adds *path* to the list of paths where `'ld'` looks for archive libraries. Using `'SEARCH_DIR (path)'` is exactly like using `'-L path'` on the command line; see “ld command line options” on page 190. If both are used, then the linker will search both paths. Paths specified using the command line option are searched first.

`STARTUP (filename)`

The `'STARTUP'` command is just like the `'INPUT'` command, except that *filename* will become the first input file to be linked, as though it were specified first on the command line. This may be useful when using a system in which the entry point is always the start of the first file.

Commands dealing with object file formats

A couple of linker script commands deal with object file formats. See also “ld command line options” on page 190 and “BFD libraries” on page 243.

`OUTPUT_FORMAT (bfdname)`

`OUTPUT_FORMAT(default, big, little)`

The `'OUTPUT_FORMAT'` command names which BFD format to use for the output file. Using `'OUTPUT_FORMAT (bfdname)'` is exactly like using `'-oformat bfdname'` on the command line. If both are used, the command line option takes precedence.

You can use `'OUTPUT_FORMAT'` with three arguments to use different formats based on the `'-EB'` and `'-EL'` command line options. This permits the linker script to set the output format based on the desired endianness. If neither `'-EB'` nor `'-EL'` is used, then the output format will be the first argument, `'DEFAULT'`. If `'-EB'` is used, the output format will be the second argument, `'BIG'`. If `'-EL'` is used, the

output format will be the third argument, 'LITTLE'. For example, the default linker script for the MIPS ELF target uses the following command:

```
OUTPUT_FORMAT(elf32-bigmips, elf32-bigmips, elf32-littlemips)
```

This says that the default format for the output file is 'elf32-bigmips', but if the user uses the '-EL' command line option, the output file will be created in the 'elf32-littlemips' format.

TARGET(*bfdname*) The 'TARGET' command names which BFD format to use when reading input files. It affects subsequent 'INPUT' and 'GROUP' commands. This command is like using '-b *bfdname*' on the command line. If the 'TARGET' command is used but 'OUTPUT_FORMAT' is not, then the last 'TARGET' command is also used to set the format for the output file.

Other linker script commands

There are a few other linker scripts commands. See also "ld command line options" on page 190 and "BFD libraries" on page 243.

FORCE_COMMON_ALLOCATION

This command has the same effect as the '-d' command-line option: to make 'ld' assign space to common symbols even if a relocatable output file is specified ('-r').

NOCROSSREFS(*section section ...*)

This command may be used to tell 'ld' to issue an error about any references among certain output sections.

In certain types of programs, particularly on embedded systems when using overlays, when one section is loaded into memory, another section will not be. Any direct references between the two sections would be errors. For example, it would be an error if code in one section called a function defined in the other section. The 'NOCROSSREFS' command takes a list of output section names. If 'ld' detects any cross-references between the sections, it reports an error and returns a non-zero exit status. Remember that the 'NOCROSSREFS' command uses output section names, not input section names.

OUTPUT_ARCH(*bfdarch*)

Specify a particular output machine architecture, *bfdarch*. The argument is one of the names used by the BFD library. You can see the architecture of an object file by using the 'objdump' program with the '-f' option.

Assigning values to symbols

You may assign a value to a symbol in a linker script. This will define the symbol as a global symbol.

Simple assignments

You may assign to a symbol using any of the C assignment operators:

```
symbol = expression ;
symbol += expression ;
symbol -= expression ;
symbol *= expression ;
symbol /= expression ;
symbol <= expression ;
symbol >= expression ;
symbol &= expression ;
symbol |= expression ;
```

- The first case will define ‘*symbol*’ to the value of ‘*expression*’. In the other cases, ‘*symbol*’ must already be defined, and the value will be accordingly adjusted.
- The special symbol name ‘.’ indicates the location counter. You may only use this within a ‘SECTIONS’ command.
- The semicolon after ‘*expression*’ is required.
- See “Expressions in linker scripts” on page 232.
- You may write symbol assignments as commands in their own right, or as statements within a ‘SECTIONS’ command, or as part of an output section description in a ‘SECTIONS’ command.
- The section of the symbol will be set from the section of the expression; for more information, see “Expressions in linker scripts” on page 232.
- The following is an example showing the three different places that symbol assignments may be used:

```
floating_point = 0;
SECTIONS
{
    .text :
    {
        *(.text)
        _etext = .;
    }
    _bdata = (. + 3) & ~ 4;
    .data : { *(.data) }
}
```

In the previous example, the `'floating_point'` symbol will be defined as zero. The `'_etext'` symbol will be defined as the address following the last `'.text'` input section. The symbol `'_bdata'` will be defined as the address following the `'.text'` output section aligned upward to a 4 byte boundary.

PROVIDE command

In some cases, it is desirable for a linker script to define a symbol only if it is referenced and is not defined by any object included in the link. For example, traditional linkers defined the symbol `'etext'`. However, ANSI C requires that the user be able to use `'etext'` as a function name without encountering an error. The `'PROVIDE'` keyword may be used to define a symbol, such as `'etext'`, only if it is referenced but not defined. The syntax is `'PROVIDE(symbol = expression)'`.

Here is an example of using `'PROVIDE'` to define `'etext'`:

```
SECTIONS
{
    .text :
    {
        *(.text)
        _etext = .;
        PROVIDE(etext = .);
    }
}
```

In the previous example, if the program defines `'_etext'`, the linker will give a multiple definition error. If, on the other hand, the program defines `'etext'`, the linker will silently use the definition in the program. If the program references `'etext'` but does not define it, the linker will use the definition in the linker script.

SECTIONS command

The ‘SECTIONS’ command tells the linker how to map input sections into output sections, and how to place the output sections in memory. The format of the ‘SECTIONS’ command is:

```
SECTIONS
{
    sections-command
    sections-command
    ...
}
```

Each ‘*sections-command*’ may be one of the following:

- An ‘ENTRY’ command (see “Setting the entry point” on page 207.)
- A symbol assignment (see “Simple assignments” on page 210)
- An output section description
- An overlay description

The ‘ENTRY’ command and symbol assignments are permitted inside the ‘SECTIONS’ command for convenience in using the location counter in those commands. This can also make the linker script easier to understand because you can use those commands at meaningful points in the layout of the output file. See “Output section description” on page 212 and “Overlay description” on page 222.

If you do not use a ‘SECTIONS’ command in your linker script, the linker will place each input section into an identically named output section in the order that the sections are first encountered in the input files. If all input sections are present in the first file, for example, the order of sections in the output file will match the order in the first input file. The first section will be at address-zero.

Output section description

The full description of an output section looks like this:

```
SECTION [address] [(type)] : [AT(LMA)]
{
    output-sections-command
    output-sections-command
    ...
} [>region] [:phdr :phdr ...] [=fillexp]
```

Most output sections do not use most of the optional section attributes. The whitespace around ‘SECTION’ is required, so that the section name is unambiguous. The colon and the curly braces are also required. The line breaks and other white space are optional.

Each `'output-sections-command'` may be one of the following:

- A symbol assignment (see “Simple assignments” on page 210)
- An input section description (see “Input section description” on page 214)
- Data values to include directly (see “Output section data” on page 217)
- A special output section keyword (see “Output section keywords” on page 218)

Output section name

The name of the output section is `'section'`. `'section'` must meet the constraints of your output format. In formats which only support a limited number of sections, such as `'a.out'`, the name must be one of the names supported by the format (`'a.out'`, for example, allows only `'text'`, `'data'` or `'bss'`). If the output format supports any number of sections, but with numbers and not names (as is the case for Oasys), the name should be supplied as a *quoted numeric string*. A section name may consist of any sequence of characters, but a name, which contains any unusual characters such as commas, must be quoted. The output section name `'/DISCARD/'` is special. See “Output section discarding” on page 219.

Output section address

The `'address'` is an expression for the VMA (the virtual memory address) of the output section. If you do not provide `'address'`, the linker will set it based on `'REGION'` if present, or otherwise based on the current value of the location counter.

If you provide `'address'`, the address of the output section will be set to precisely that specification. If you provide neither `'address'` nor `'region'`, then the address of the output section will be set to the current value of the location counter aligned to the alignment requirements of the output section. The alignment requirement of the output section is the strictest alignment of any input section contained within the output section.

For example:

```
.text . : { *(.text) }
```

and

```
.text : { *(.text) }
```

are subtly different. The first will set the address of the `'text'` output section to the current value of the location counter. The second will set it to the current value of the location counter aligned to the strictest alignment of a `'text'` input section.

The `'address'` may be an arbitrary expression. See “Expressions in linker scripts” on page 232. For example, if you want to align the section on a 0x10 byte boundary, so that the lowest four bits of the section address are zero, you could do something like

the following declaration:

```
.text ALIGN(0x10) : { *(.text) }
```

This works because ‘ALIGN’ returns the current location counter aligned upward to the specified value.

Specifying ‘*address*’ for a section will change the value of the location counter.

Input section description

The most common output section command is an *input section description*. The input section description is the most basic linker script operation. You use output sections to tell the linker how to lay out your program in memory. You use input section descriptions to tell the linker how to map the input files into your memory layout.

Input section basics

An *input section description* consists of a file name optionally followed by a list of section names in parentheses. The file name and the section name may be wildcard patterns, which we describe; see “Input section wildcard patterns” on page 215. The most common input section description is to include all input sections with a particular name in the output section. For example, to include all input ‘.text’ sections, you would write:

```
*(.text)
```

Here the ‘*’ is a wildcard which matches *any* file name.

There are two ways to include more than one section:

```
*(.text .rdata)
*(.text) *(.rdata)
```

The difference between these is the order in which the ‘.text’ and ‘.rdata’ input sections will appear in the output section. In the first example, they will be intermingled. In the second example, all ‘.text’ input sections will appear first, followed by all ‘.rdata’ input sections.

You can specify a file name to include sections from a particular file. You would do this if one or more of your files contain special data that needs to be at a particular location in memory. For example:

```
data.o(.data)
```

If you use a file name without a list of sections, then all sections in the input file will be included in the output section. This is not commonly done, but it may be useful on occasion. For example:

```
data.o
```

When you use a file name, which does not contain any wild card characters, the linker will first see if you also specified the file name on the linker command line or in an

‘INPUT’ command. If you did not, the linker will attempt to open the file as an input file, as though it appeared on the command line. Note that this differs from an ‘INPUT’ command, because the linker will not search for the file in the archive search path.

Input section wildcard patterns

In an input section description, either the file name or the section name or both may be wildcard patterns. The file name of ‘*’ seen in many examples is a simple wildcard pattern for the file name. The wildcard patterns are like those used by the Unix shell.

‘*’

Matches any number of characters.

‘?’

Matches any single character.

‘[*chars*]’

Matches a single instance of any of the *chars*; the ‘-’ character may be used to specify a range of characters, as in ‘[a-z]’ to match any lower case letter.

‘\’

Quotes the following character.

When a file name is matched with a wildcard, the wildcard characters will not match a ‘/’ character (used to separate directory names on Unix). A pattern consisting of a single ‘*’ character is an exception; it will always match any file name, whether it contains a ‘/’ or not. In a section name, the wildcard characters will match a ‘/’ character.

File name wildcard patterns only match files which are explicitly specified on the command line or in an ‘INPUT’ command. The linker does not search directories to expand wildcards.

If a file name matches more than one wildcard pattern, or if a file name appears explicitly and is also matched by a wildcard pattern, the linker will use the first match in the linker script. For example, this sequence of input section descriptions is probably in error, because the ‘data.o’ rule will not be used:

```
.data : { *(.data) }
.data1 : { data.o(.data) }
```

If you ever get confused about where input sections are going, use the ‘-M’ linker option to generate a map file. The map file shows precisely, how input sections are mapped to output sections.

This example shows how wildcard patterns might be used to partition files. This linker script directs the linker to place all ‘.text’ sections in ‘.text’ and all ‘.bss’ sections in ‘.bss’. The linker will place the ‘.data’ section from all files beginning with an upper case character in ‘.DATA’; for all other files, the linker will place the

```
'data' section in '.data'.
SECTIONS {
    .text : { *(.text) }
    .DATA : { [A-Z]*(.data) }
    .data : { *(.data) }
    .bss : { *(.bss) }
}
```

Input section for common symbols

A special notation is needed for common symbols, because in many object-file formats common symbols do not have a particular input section. The linker treats common symbols as though they are in an input section named 'COMMON'.

You may use file names with the 'COMMON' section just as with any other input sections. You can use this to place common symbols from a particular input file in one section while common symbols from other input files are placed in another section.

In most cases, common symbols in input files will be placed in the '.bss' section in the output file. For example:

```
.bss { *(.bss) *(COMMON) }
```

Some object file formats have more than one type of common symbol. For example, the MIPS ELF object file format distinguishes standard common symbols and small common symbols. In this case, the linker will use a different special section name for other types of common symbols. In the case of MIPS ELF, the linker uses 'COMMON' for standard common symbols and '.scommon' for small common symbols. This permits you to map the different types of common symbols into memory at different locations.

You will sometimes see '[COMMON]' in old linker scripts. This notation is now considered obsolete. It is equivalent to '* (COMMON)'.

Input section example

The following example is a complete linker script. It tells the linker to read all of the sections from file 'all.o' and place them at the start of output section 'outputa', which starts at location '0x10000'. All of section '.input1' from file 'foo.o' follows immediately, in the same output section. All of section '.input2' from 'foo.o' goes into output section 'outputb', followed by section '.input1' from 'foo1.o'. All of the remaining '.input1' and '.input2' sections from any files are written to output section 'outputc'.

```
SECTIONS {
    outputa 0x10000 :
    {
```

```

    all.o
    foo.o (.input1)
}
outputb :
{
    foo.o (.input2)
    fool.o (.input1)
}
outputc :
{
    *(.input1)
    *(.input2)
}
}

```

Output section data

You can include explicit bytes of data in an output section by using ‘BYTE’, ‘SHORT’, ‘LONG’, ‘QUAD’, or ‘SQUAD’ as an output section command. Each keyword is followed by an expression in parentheses providing the value to store; see “Expressions in linker scripts” on page 232. The value of the expression is stored at the current value of the location counter.

The ‘BYTE’, ‘SHORT’, ‘LONG’, and ‘QUAD’ commands store one, two, four, and eight bytes (respectively). After storing the bytes, the location counter is incremented by the number of bytes stored. For example, this will store the byte 1 followed by the four byte value of the symbol ‘addr’:

```

    BYTE(1)
    LONG(addr)

```

When using a 64-bit host or target, ‘QUAD’ and ‘SQUAD’ are the same; they both store an 8-byte, or 64-bit, value. When both host and target are 32 bits, an expression is computed as 32 bits. In this case ‘QUAD’ stores a 32-bit value zero extended to 64 bits, and ‘SQUAD’ stores a 32-bit value sign extended to 64 bits.

If the object file format of the output file has an explicit endianness, which is the normal case, the value will be stored in that endianness. When the object file format does not have an explicit endianness, as is true of, for example, S-records, the value will be stored in the endianness of the first input object file.

You may use the ‘FILL’ command to set the fill pattern for the current section. It is followed by an expression in parentheses. Any otherwise unspecified regions of memory within the section (for example, gaps left due to the required alignment of input sections) are filled with the two least significant bytes of the expression, repeated as necessary. A ‘FILL’ statement covers memory locations after the point at which it occurs in the section definition; by including more than one ‘FILL’ statement,

you can have different fill patterns in different parts of an output section.

This example shows how to fill unspecified regions of memory with the value ‘0x9090’:

```
FILL(0x9090)
```

The ‘FILL’ command is similar to the ‘=*fillexp*’ output section attribute (see “Output section fill” on page 221); but it only affects the part of the section following the ‘FILL’ command, rather than the entire section. If both are used, the ‘FILL’ command takes precedence.

Output section keywords

There are a couple of keywords, which can appear as output section commands.

CREATE_OBJECT_SYMBOLS

The command tells the linker to create a symbol for each input file. The name of each symbol will be the name of the corresponding input file. The section of each symbol will be the output section in which the ‘CREATE_OBJECT_SYMBOLS’ command appears.

This is conventional for the ‘a.out’ object file format. It is not normally used for any other object file format.

CONSTRUCTORS

When linking, using the ‘a.out’ object file format, the linker uses an unusual set construct to support C++ global constructors and destructors. When linking object file formats, which do not support arbitrary sections, such as ‘ECOFF’ and ‘XCOFF’, the linker will automatically recognize C++ global constructors and destructors by name. For these object file formats, the ‘CONSTRUCTORS’ command tells the linker to place constructor information in the output section where the ‘CONSTRUCTORS’ command appears. The ‘CONSTRUCTORS’ command is ignored for other object file formats.

The symbol ‘__CTOR_LIST__’ marks the start of the global constructors, and the symbol ‘__DTOR_LIST’ marks the end. The first word in the list is the number of entries, followed by the address of each constructor or destructor, followed by a zero word. The compiler must arrange to actually run the code. For these object file formats GNU C++ normally calls constructors from a subroutine ‘__main’; a call to ‘__main’ is automatically inserted into the startup code for ‘main’. GNU C++ normally runs destructors either by using ‘atexit’, or directly from the function ‘exit’.

For object file formats such as ‘COFF’ or ‘ELF’, which support arbitrary section names, GNU C++ will normally arrange to put the addresses of global constructors and destructors into the ‘.ctors’ and ‘.dtors’ sections. Placing the following sequence into your linker script will build the sort of table that the GNU

C++ runtime code expects to see.

```
__CTOR_LIST__ = .;
LONG((__CTOR_END__ - __CTOR_LIST__) / 4 - 2)
*(.ctors)
LONG(0)
__CTOR_END__ = .;
__DTOR_LIST__ = .;
LONG((__DTOR_END__ - __DTOR_LIST__) / 4 - 2)
*(.dtors)
LONG(0)
__DTOR_END__ = .;
```

Normally the compiler and linker will handle these issues automatically, and you will not need to concern yourself with them. However, you may need to consider this occurrence, if you are using C++ and writing your own linker scripts.

Output section discarding

The linker will not create output section which do not have any contents. This is for convenience when referring to input sections that may or may not be present in any of the input files. For example:

```
.foo { *(.foo) }
```

Will only create a ‘.foo’ section in the output file if there is a ‘.foo’ section in at least one input file.

If you use anything other than an input section description as an output section command, such as a symbol assignment, then the output section will always be created, even if there are no matching input sections.

The special output section name ‘/DISCARD/’ may be used to discard input sections. Any input sections which are assigned to an output section named ‘/DISCARD/’ are not included in the output file.

Output section attributes

We showed above that the full description of an output section looked like this:

```
SECTION [address] [(type)] : [AT(LMA)]
{
    output-sections-command
    output-sections-command
    ...
} [>region] [:phdr :phdr ...] [=fillexp]
```

We’ve already described ‘section’, ‘address’, and ‘output-sections-command’. In the following discussion, we will describe the remaining section attributes.

Output section type

Each output section may have a type. The type is a keyword in parentheses. The following types are defined:

NOLOAD

The section should be marked as not loadable, so that it will not be loaded into memory when the program is run.

DSECT

COPY

INFO

OVERLAY

These type names are supported for backward compatibility, and are rarely used. They all have the same effect: the section should be marked as not allocatable, so that no memory is allocated for the section when the program is run.

The linker normally sets the attributes of an output section, based on the input sections, which map into it. You can override this by using the section type. For example, in the script sample below, the ‘ROM’ section is addressed at memory location ‘0’ and does not need to be loaded when the program is run. The contents of the ‘ROM’ section will appear in the linker output file as usual.

```
SECTIONS {
    ROM 0 (NOLOAD) : { ... }
    ...
}
```

Output section LMA

Every section has a virtual address (VMA) and a load address (LMA); see “Basic linker script concepts” on page 204. The address expression that, may appear in an output section description sets the VMA. The linker will normally set the LMA equal to the ‘VMA’. You can change that by using the ‘AT’ keyword. The expression, LMA, that follows the ‘AT’ keyword specifies the load address of the section. This feature is designed to make it easy to build a ROM image. For example, the following linker script creates three output sections: one called ‘.text’, which starts at ‘0x1000’, one called ‘.mdata’, which is loaded at the end of the ‘.text’ section even though its VMA is ‘0x2000’, and one called ‘.bss’ to hold uninitialized data at address ‘0x3000’. The symbol ‘_data’ is defined with the value ‘0x2000’, which shows that the location counter holds the VMA value, not the LMA value.

```
SECTIONS
{
    .text 0x1000 : { *(.text) _etext = . ; }
    .mdata 0x2000 :
        AT ( ADDR (.text) + SIZEOF (.text) )
        { _data = . ; *(.data); _edata = . ; }
```



```

        .bss 0x3000 :
        { _bstart = . ; *(.bss) *(COMMON) ; _bend = . ; }
    }

```

The run-time initialization code for use with a program generated with this linker script would include something like the following, to copy the initialized data from the ROM image to its runtime address. Notice how this code takes advantage of the symbols defined by the linker script.

```

extern char _etext, _data, _edata, _bstart, _bend;
char *src = &_etext;
char *dst = &_data;

/* ROM has data at end of text; copy it. */
while (dst < &_edata) {
    *dst++ = *src++;
}

/* Zero bss */
for (dst = &_bstart; dst < &_bend; dst++)
    *dst = 0;

```

Output section region

You can assign a section to a previously defined region of memory by using ‘>REGION’. Here is a simple example:

```

MEMORY { rom : ORIGIN = 0x1000, LENGTH = 0x1000 }
SECTIONS { ROM : { *(.text) } >rom }

```

Output section phdr

You can assign a section to a previously defined program segment by using ‘:phdr’. If a section is assigned to one or more segments, then all subsequent allocated sections will be assigned to those segments as well, unless they use an explicitly ‘:phdr’ modifier. To prevent a section from being assigned to a segment when it would normally default to one, use ‘:NONE’. See “PHDRS command” on page 226.

Here is a simple example:

```

PHDRS { text PT_LOAD ; }
SECTIONS { .text : { *(.text) } :text }

```

Output section fill

You can set the fill pattern for an entire section by using *=fillexp*. ‘*fillexp*’ is an expression; see “Expressions in linker scripts” on page 232. Any otherwise unspecified regions of memory within the output section (for example, gaps left due to the required alignment of input sections) will be filled with the two least significant

bytes of the value, repeated as necessary.

You can also change the fill value with a `'FILL'` command in the output section commands. See “Output section data” on page 217.

Here is a simple example:

```
SECTIONS { .text : { *(.text) } =0x9090 }
```

Overlay description

An overlay description provides an easy way to describe sections, which are to be loaded as part of a single memory image but are to be run at the same memory address. At run time, some sort of overlay manager will copy the overlaid sections in and out of the runtime memory address as required, perhaps by simply manipulating addressing bits. This approach can be useful, for example, when a certain region of memory is faster than another region of memory.

Overlays are described using the `'OVERLAY'` command. The `'OVERLAY'` command is used within a `'SECTIONS'` command, like an output section description. The full syntax of the `'OVERLAY'` command is as follows:

```
OVERLAY [start] : [NOCROSSREFS] [AT ( ldaddr )]
{
    secname1
    {
        output-section-command
output-section-command
        ...
    } [:PHDR...] [=FILL]
    secname2
    {
        output-section-command
output-section-command
        ...
    } [:phdr...] [=fill]
    ...
} [>region] [:phdr...] [=fill]
```

Everything is optional except `'OVERLAY'` (a keyword), and each section must have a name (`'secname1'` and `'secname2'` above). The section definitions within the `'OVERLAY'` construct are identical to those within the general `'SECTIONS'` construct, except that no addresses and no memory regions may be defined for sections within an `'OVERLAY'`. See “SECTIONS command” on page 212.

The sections are all defined with the same starting address. The load addresses of the sections are arranged, so that they are consecutive in memory, starting at the load address used for the `'OVERLAY'` as a whole (as with normal section definitions. The load address is optional, and defaults to the start address. The start address is also

optional, and defaults to the current value of the location counter).

If the 'NOCROSSREFS' keyword is used, and there any references among the sections, the linker will report an error. Since the sections all run at the same address, it normally does not make sense for one section to refer directly to another.

For each section within the 'OVERLAY', the linker automatically defines two symbols. The symbol '__load_start_secname' is defined as the starting load address of the section. The symbol '__load_stop_secname' is defined as the final load address of the section. Any characters within 'secname' that are not legal within C identifiers are removed. C (or assembler) code may use these symbols to move the overlaid sections around as necessary.

At the end of the overlay, the value of the location counter is set to the start address of the overlay plus the size of the largest section.

Here is an example. Remember that this would appear inside a 'SECTIONS' construct.

```
OVERLAY 0x1000 : AT (0x4000)
{
    .text0 { o1/*.o(.text) }
    .text1 { o2/*.o(.text) }
}
```

This will define both '.text0' and '.text1' to start at address '0x1000'. '.text0' will be loaded at address '0x4000', and '.text1' will be loaded immediately after '.text0'. The following symbols will be defined: '__load_start_text0', '__load_stop_text0', '__load_start_text1', '__load_stop_text1'.

C code to copy overlay '.text1' into the overlay area might look like the following.

```
extern char __load_start_text1, __load_stop_text1;
memcpy ((char *) 0x1000, &__load_start_text1,
        &__load_stop_text1 - &__load_start_text1);
```

Note that the 'OVERLAY' command is just syntactic sugar, since everything it does can be done using the more basic commands. The above example could have been written identically as follows.

```
.text0 0x1000 : AT (0x4000) { o1/*.o(.text) }
__load_start_text0 = LOADADDR (.text0);
__load_stop_text0 = LOADADDR (.text0) + SIZEOF (.text0);
.text1 0x1000 : AT (0x4000 + SIZEOF (.text0)) { o2/*.o(.text) }
__load_start_text1 = LOADADDR (.text1);
__load_stop_text1 = LOADADDR (.text1) + SIZEOF (.text1);
. = 0x1000 + MAX (SIZEOF (.text0), SIZEOF (.text1));
```

MEMORY command

The linker's default configuration permits allocation of all available memory. You can override this by using the 'MEMORY' command.

The 'MEMORY' command describes the location and size of blocks of memory in the target. You can use it to describe which memory regions may be used by the linker, and which memory regions it must avoid. You can then assign sections to particular memory regions. The linker will set section addresses based on the memory regions, and will warn about regions that become too full. The linker will not shuffle sections around to fit into the available regions.

A linker script may contain at most one use of the 'MEMORY' command. However, you can define as many blocks of memory within it as you wish. The syntax is:

```
MEMORY
{
    name [(attr)] : ORIGIN = origin, LENGTH = len
    ...
}
```

The '*name*' is a name used in the linker script to refer to the region. The region name has no meaning outside of the linker script. Region names are stored in a separate name space, and will not conflict with symbol names, file names, or section names. Each memory region must have a distinct name.

The '*attr*' string is an optional list of attributes that specify whether to use a particular memory region for an input section, which is not explicitly mapped in the linker script. If you do not specify an output section for some input section, the linker will create an output section with the same name as the input section. If you define region attributes, the linker will use them to select the memory region for the output section that it creates. See "SECTIONS command" on page 212.

The '*attr*' string must consist only of the following characters:

R	Read-only section
W	Read/write section
X	Executable section
A	Allocatable section
I	Initialized section

L

Same as 'I'

!

Invert the sense of any of the preceding attributes

If an unmapped section matches any of the listed attributes other than '!', it will be placed in the memory region. The '!' attribute reverses this test, so that an unmapped section will be placed in the memory region only if it does not match any of the listed attributes.

The 'ORIGIN' is an expression for the start address of the memory region. The expression must evaluate to a constant before memory allocation is performed, which means that you may not use any section relative symbols. The keyword 'ORIGIN' may be abbreviated to 'org' or 'o' (but not, for example, 'ORG').

The 'len' is an expression for the size in bytes of the memory region. As with the 'origin' expression, the expression must evaluate to a constant before memory allocation is performed. The keyword 'LENGTH' may be abbreviated to 'len' or 'l'.

In the following example, we specify that there are two memory regions available for allocation: one starting at '0' for 256 kilobytes, and the other starting at '0x40000000' for four megabytes. The linker will place into the 'rom' memory region every section, which is not explicitly mapped into a memory region, and is either read-only or executable. The linker will place other sections, which are not explicitly mapped into a memory region into the 'ram' memory region.

```
MEMORY
{
    rom (rx) : ORIGIN = 0, LENGTH = 256K
    ram (!rx) : org = 0x40000000, l = 4M
}
```

If you have defined a memory region named 'mem', you can direct the linker to place specific output sections into that memory region by using the '>region' output section attribute. If no address was specified for the output section, the linker will set the address to the next available address within the memory region. If the combined output sections directed to a memory region are too large for the region, the linker will issue an error message. See "Output section region" on page 221.

PHDRS command

The ELF object file format uses *program headers*, also known as *segments*. The program headers describe how the program should be loaded into memory. You can print them out by using the ‘objdump’ program with the ‘-p’ option.

When you run an ELF program on a native ELF system, the system loader reads the program headers in order to figure out how to load the program. This will only work if the program headers are set correctly. This documentation does not describe the details of how the system loader interprets program headers; for more information, see the ELF ABI.

The linker will create reasonable program headers by default. However, in some cases, you may need to specify the program headers more precisely. You may use the ‘PHDRS’ command for this purpose. When the linker sees the ‘PHDRS’ command in the linker script, it will not create any program headers other than the ones specified.

The linker only pays attention to the ‘PHDRS’ command when generating an ELF output file. In other cases, the linker will simply ignore ‘PHDRS’.

This is the syntax of the ‘PHDRS’ command. The words ‘PHDRS’, ‘FILEHDR’, ‘AT’, and ‘FLAGS’ are keywords.

```
PHDRS
{
    name type[ FILEHDR ] [ PHDRS ] [ AT ( address ) ]
        [ FLAGS ( flags ) ] ;
}
```

The ‘name’ is used only for reference in the ‘SECTIONS’ command of the linker script. It is not put into the output file. Program header names are stored in a separate name space, and will not conflict with symbol names, file names, or section names. Each program header must have a distinct name.

Certain program header types describe segments of memory, which the system loader will load from the file. In the linker script, you specify the contents of these segments by placing allocatable output sections in the segments. You use the ‘:phdr’ output section attribute to place a section in a particular segment. See “Output section phdr” on page 221.

It is normal to put certain sections in more than one segment. This merely implies that one segment of memory contains another. You may repeat ‘:phdr’, using it once for each segment which should contain the section.

If you place a section in one or more segments using ‘:phdr’, then the linker will place all subsequent allocatable sections which do not specify ‘:phdr’ in the same segments. This is for convenience, since generally a whole set of contiguous sections

will be placed in a single segment. To prevent a section from being assigned to a segment when it would normally default to one, use `:NONE`.

You may use the `FILEHDR` and `PHDRS` keywords appear after the program header type to further describe the contents of the segment. The `FILEHDR` keyword means that the segment should include the ELF file header. The `PHDRS` keyword means that the segment should include the ELF program headers themselves.

The `type` may be one of the following. The numbers indicate the value of the keyword.

`PT_NULL (0)`

Indicates an unused program header.

`PT_LOAD (1)`

Indicates that this program header describes a segment to be loaded from the file.

`PT_DYNAMIC (2)`

Indicates a segment where dynamic linking information can be found.

`PT_INTERP (3)`

Indicates a segment where the name of the program interpreter may be found.

`PT_NOTE (4)`

Indicates a segment holding note information.

`PT_SHLIB (5)`

A reserved program header type, defined but not specified by the ELF ABI.

`PT_PHDR (6)`

Indicates a segment where the program headers may be found.

expression

An expression giving the numeric type of the program header. This may be used for types not defined above.

You can specify that a segment should be loaded at a particular address in memory by using an `AT` expression. This is identical to the `AT` command used as an output section attribute. The `AT` command for a program header, overrides the output section attribute. See “Output section LMA” on page 220.

The linker will normally set the segment flags based on the sections, which comprise the segment. You may use the `FLAGS` keyword to explicitly specify the segment flags. The value of *flags* must be an integer. It is used to set the `p_flags` field of the program header.

Here is an example of `PHDRS`. This shows a typical set of program headers used on a native ELF system.

```
PHDRS
{
    headers PT_PHDR PHDRS ;
    interp PT_INTERP ;
    text PT_LOAD FILEHDR PHDRS ;
```

```
    data PT_LOAD ;
    dynamic PT_DYNAMIC ;
}

SECTIONS
{
    . = SIZEOF_HEADERS;
    .interp : { *(.interp) } :text :interp
    .text : { *(.text) } :text
    .rodata : { *(.rodata) } /* defaults to :text */
    ...
    . = . + 0x1000; /* move to a new page in memory */
    .data : { *(.data) } :data
    .dynamic : { *(.dynamic) } :data :dynamic
    ...
}
```


VERSION command

The linker supports symbol versions when using ELF. Symbol versions are only useful when using shared libraries. The dynamic linker can use symbol versions to select a specific version of a function when it runs a program that may have been linked against an earlier version of the shared library.

You can include a version script directly in the main linker script, or you can supply the version script as an implicit linker script. You can also use the ‘`--version-script`’ linker option.

The syntax of the ‘`VERSION`’ command is simply

```
VERSION { version-script-commands }
```

The format of the version script commands is identical to that used by Sun’s linker in Solaris 2.5. The version script defines a tree of version nodes. You specify the node names and interdependencies in the version script. You can specify which symbols are bound to which version nodes, and you can reduce a specified set of symbols to local scope so that they are not globally visible outside of the shared library.

The easiest way to demonstrate the version script language is with a few examples.

```
VERS_1.1 {
    global:
    foo1;
    local:
    old*;
    original*;
    new*;
};

VERS_1.2 {
    foo2;
} VERS_1.1;

VERS_2.0 {
    bar1; bar2;
} VERS_1.2;
```

This example version script defines three version nodes. The first version node defined is ‘`VERS_1.1`’; it has no other dependencies. The script binds the symbol ‘`foo1`’ to ‘`VERS_1.1`’. It reduces a number of symbols to local scope so that they are not visible outside of the shared library.

Next, the version script defines node ‘`VERS_1.2`’. This node depends upon ‘`VERS_1.1`’. The script binds the symbol ‘`foo2`’ to the version node ‘`VERS_1.2`’.

Finally, the version script defines node ‘`VERS_2.0`’. This node depends upon

`'VERS_1.2'`. The script binds the symbols `'bar1'` and `'bar2'` to the version node `'VERS_2.0'`.

When the linker finds a symbol defined in a library, which is not specifically bound to a version node, it will effectively bind it to an unspecified base version of the library. You can bind all otherwise unspecified symbols to a given version node by using `'global: *'` somewhere in the version script.

The names of the version nodes have no specific meaning other than what they might suggest to the person reading them. The `'2.0'` version could just as well have appeared in between `'1.1'` and `'1.2'`. However, this would be a confusing way to write a version script.

When you link an application against a shared library that has versioned symbols, the application itself knows which version of each symbol it requires, and it also knows which version nodes it needs from each shared library it is linked against. Thus at runtime, the dynamic loader can make a quick check to make sure that the libraries you have linked against do in fact supply all of the version nodes that the application will need to resolve all of the dynamic symbols. In this way it is possible for the dynamic linker to know with certainty that all external symbols that it needs will be resolvable without having to search for each symbol reference.

The symbol versioning is in effect a much more sophisticated way of doing minor version checking that SunOS does. The fundamental problem that is being addressed here is that typically references to external functions are bound on an as-needed basis, and are not all bound when the application starts up. If a shared library is out of date, a required interface may be missing; when the application tries to use that interface, it may suddenly and unexpectedly fail. With symbol versioning, the user will get a warning when they start their program if the libraries being used with the application are too old.

There are several GNU extensions to Sun's versioning approach. The first of these is the ability to bind a symbol to a version node in the source file where the symbol is defined instead of in the versioning script. This was done mainly to reduce the burden on the library maintainer. You can do this by putting something like this in the C source file:

```
__asm__( ".symver original_foo,foo@VERS_1.1" );
```

This renames the function `'original_foo'` to be an alias for `'foo'` bound to the version node `'VERS_1.1'`. The `'local:'` directive can be used to prevent the symbol `'original_foo'` from being exported.

The second GNU extension is to allow multiple versions of the same function to appear in a given, shared library. In this way you can make an incompatible change to an interface without increasing the major version number of the shared library, while

still allowing applications linked against the old interface to continue to function.

To do this, you must use multiple `‘.symver’` directives in the source file. Here is an example:

```
__asm__( ".symver original_foo,foo@" );  
__asm__( ".symver old_foo,foo@VERS_1.1" );  
__asm__( ".symver old_fool,foo@VERS_1.2" );  
__asm__( ".symver new_foo,foo@@VERS_2.0" );
```

In this example, `‘foo@’` represents the symbol `‘foo’` bound to the unspecified base version of the symbol. The source file that contains this example would define four C functions: `‘original_foo’`, `‘old_foo’`, `‘old_fool’`, and `‘new_foo’`.

When you have multiple definitions of a given symbol, there needs to be some way to specify a default version to which external references to this symbol will be bound. You can do this with the `‘foo@@VERS_2.0’` type of `‘.symver’` directive. You can only declare one version of a symbol as the default in this manner; otherwise you would effectively have multiple definitions of the same symbol.

If you wish to bind a reference to a specific version of the symbol within the shared library, you can use the aliases of convenience (i.e. `‘old_foo’`), or you can use the `‘.symver’` directive to specifically bind to an external version of the function in question.

Expressions in linker scripts

The syntax for expressions in the linker script language is identical to that of C expressions. All expressions are evaluated as integers. All expressions are evaluated in the same size, which is 32 bits if both the host and target are 32 bits, and is otherwise 64 bits. You can use and set symbol values in expressions. The linker defines several special purpose builtin functions for use in expressions.

Constants

All constants are integers. As in C, the linker considers an integer beginning with '0' to be octal, and an integer beginning with '0x' or '0X' to be hexadecimal.

The linker considers other integers to be decimal.

In addition, you can use the suffixes 'K' and 'M' to scale a constant by '1024' or '1024*1024' respectively. For example, the following all refer to the same quantity:

```
_fourk_1 = 4K;  
_fourk_2 = 4096;  
_fourk_3 = 0x1000;
```

Symbol names

Unless quoted, symbol names start with a letter, underscore, or period and may include letters, digits, underscores, periods, and hyphens. Unquoted symbol names must not conflict with any keywords. You can specify a symbol, which contains odd characters or has the same name as a keyword by surrounding the symbol name in double quotes:

```
"SECTION" = 9;  
"with a space" = "also with a space" + 10;
```

Since symbols can contain many non-alphabetic characters, it is safest to delimit symbols with spaces. For example, 'A-B' is one symbol, whereas 'A - B' is an expression involving subtraction.

The location counter

The special linker *dot* '.' variable always contains the current output location counter. Since the '.' always refers to a location in an output section, it may only appear in an expression within a 'SECTIONS' command. The '.' symbol may appear anywhere that an ordinary symbol is allowed in an expression.

Assigning a value to '.' will cause the location counter to be moved. This may be used to create holes in the output section. The location counter may never be moved backwards.

```

SECTIONS
{
    output :
    {
        file1(.text)
        . = . + 1000;
        file2(.text)
        . += 1000;
        file3(.text)
    } = 0x1234;
}

```

In the previous example, the `.text` section from `file1` is located at the beginning of the output section `output`. It is followed by a 1000 byte gap. Then the `.text` section from `file2` appears, also with a 1000 byte gap following before the `.text` section from `file3`. The notation `= 0x1234` specifies data to write in the gaps.

Operators

The linker recognizes the standard C set of arithmetic operators, with the standard bindings and precedence levels; see Table 2.

Table 2: Arithmetic operators with precedence levels and bindings associations

<i>Precedence</i>	<i>Association</i>	<i>Operators</i>	<i>Notes</i>
(highest)			
1	left	! - ~	†
2	left	* / %	
3	left	+ -	
4	left	>> <<	
5	left	== != > < <= >=	
6	left	&	
7	left		
8	left	&&	
9	left		‡
10	right	? :	
11	right	&= += -= *= /=	
(lowest)			

† Prefix operators

‡ See “Assigning values to symbols” on page 210.

Evaluation

The linker evaluates expressions lazily. It only computes the value of an expression when absolutely necessary.

The linker needs some information, such as the value of the start address of the first section, and the origins and lengths of memory regions, in order to do any linking at all. These values are computed as soon as possible when the linker reads in the linker script.

However, other values (such as symbol values) are not known or needed until after storage allocation. Such values are evaluated later, when other information (such as the sizes of output sections) is available for use in the symbol assignment expression.

The sizes of sections cannot be known until after allocation, so assignments dependent upon these are not performed until after allocation.

Some expressions, such as those depending upon the location counter ‘.’, must be evaluated during section allocation.

If the result of an expression is required, but the value is not available, then an error results. For example, a script, like the following, will cause the error message ‘**non constant expression for initial address**’:

```
SECTIONS
{
    .text 9+this_isnt_constant :
        { *(.text) }
}
```

The section of an expression

When the linker evaluates an expression, the result is either absolute or relative to some section. A relative expression is expressed as a fixed offset from the base of a section.

The position of the expression within the linker script determines whether it is absolute or relative. An expression, which appears within an output section definition, is relative to the base of the output section. An expression, which appears elsewhere, will be absolute.

A symbol set to a relative expression will be relocatable if you request relocatable output using the ‘-r’ option. That means that a further link operation may change the value of the symbol. The symbol’s section will be the section of the relative expression.

A symbol set to an absolute expression will retain the same value through any further link operation. The symbol will be absolute, and will not have any particular associated section.

You can use the builtin function ‘ABSOLUTE’ to force an expression to be absolute when it would otherwise be relative. For example, to create an absolute symbol set to the address of the end of the output section ‘.data’:

```
SECTIONS
{
    .data : { *(.data) _edata = ABSOLUTE(.); }
}
```

If ‘ABSOLUTE’ were not used, ‘_edata’ would be relative to the ‘.data’ section.

Builtin functions

The linker script language includes a number of builtin functions for use in linker script expressions.

ABSOLUTE(*exp*)

Return the absolute (non-relocatable, as opposed to non-negative) value of the ‘*exp*’ expression. Primarily useful to assign an absolute value to a symbol within a section definition, where symbol values are normally section relative. See “Expressions in linker scripts” on page 232.

ADDR(*section*)

Return the absolute address (the VMA) of the named *section*. Your script must previously have defined the location of that section. In the following example, ‘symbol_1’ and ‘symbol_2’ are assigned identical values:

```
SECTIONS { ...
    .output1 :
    {
        start_of_output_1 = ABSOLUTE(.);
        ...
    }
    .output :
    {
        symbol_1 = ADDR(.output1);
        symbol_2 = start_of_output_1;
    }
    ...
}
```

ALIGN(*exp*)

Return the location counter (‘.’) aligned to the next ‘*exp*’ boundary. ‘*exp*’ must be an expression whose value is a power of two. This is equivalent to:

$$(. + exp - 1) \& \sim(exp - 1)$$

‘ALIGN’ doesn’t change the value of the location counter, it just does arithmetic on it. Here is an example which aligns the output ‘.data’ section to the next ‘0x2000’ byte boundary after the preceding section and sets a variable within the section to the next ‘0x8000’ boundary after the input sections:

```
SECTIONS { ...  
  .data ALIGN(0x2000): {  
    *(.data)  
    variable = ALIGN(0x8000);  
  }  
  ...  
}
```

The first use of ‘ALIGN’ in this example specifies the location of a section because it is used as the optional ‘ADDRESS’ attribute of a section definition. The second use of ‘ALIGN’ is to define the value of a symbol. The builtin function ‘NEXT’ is closely related to ‘ALIGN’. See “Output section address” on page 213.

BLOCK(*exp*)

This is a synonym for ‘ALIGN’, for compatibility with older linker scripts. It is most often seen when setting the address of an output section.

DEFINED(*symbol*)

Return ‘1’ if ‘*symbol*’ is in the linker global symbol table and is defined, otherwise return ‘0’. You can use this function to provide default values for symbols. For example, the following script fragment shows how to set a global symbol ‘begin’ to the first location in the ‘.text’ section, but if a symbol called ‘begin’ already existed, its value is preserved:

```
SECTIONS{ ...  
  .text : {  
    begin = DEFINED(begin) ? begin : . ;  
    ...  
  }  
  ...  
}
```

LOADADDR(*section*)

Return the absolute ‘LMA’ of the named ‘SECTION’. This is normally the same as ‘ADDR’, but it may be different if the ‘AT’ attribute is used in the output section definition.

MAX(*exp1*, *exp2*)

Returns the maximum of ‘*exp1*’ and ‘*exp2*’.

MIN(*exp1*, *exp2*)

Returns the minimum of ‘*exp1*’ and ‘*exp2*’.

NEXT(*exp*)

Return the next unallocated address that is a multiple of ‘*exp*’. This function is closely related to ‘ALIGN(*exp*)’; unless you use the ‘MEMORY’ command to define discontinuous memory for the output file, the two functions are equivalent.

SIZEOF(*section*)

Return the size in bytes of the named ‘*section*’, if that section has been allocated. If the section has not been allocated when this is evaluated, the linker will report

an error. See “PHDRS command” on page 226. In the following example, ‘symbol_1’ and ‘symbol_2’ are assigned identical values:

```
SECTIONS{ ...  
  .output {  
    .start = . ;  
  ...
```


4

ld machine dependent features

The following documentation describes some machine independent features for the GNU linker for the H8/300 and the Intel 960 processors. See the following documentation for details.

- “ld and the H8/300 processors” on page 240
- “ld and the Intel 960 processors” on page 241

Machines with ld having no additional functionality have no documentation.

ld and the H8/300 processors

For the H8/300 processors, ld can perform these global optimizations when you specify the `-relax` command-line option.

relaxing address modes

ld finds all `jsr` and `jmp` instructions whose targets are within eight bits, and turns them into eight-bit program-counter relative `bsr` and `bra` instructions, respectively.

synthesizing instructions

ld finds all `mov.b` instructions which use the sixteen-bit absolute address form, but refer to the top page of memory, and changes them to use the eight-bit address form. (That is, the linker turns `'mov.b @ aa:16'` into `'mov.b @ aa:8'` whenever the address `aa` is in the top page of memory).

ld and the Intel 960 processors

You can use the `‘-Aarchitecture’` command line option to specify one of the two-letter names identifying members of the 960 processors; the option specifies the desired output target, and warns of any incompatible instructions in the input files. It also modifies the linker’s search strategy for archive libraries, to support the use of libraries specific to each particular architecture, by including in the search loop names suffixed with the string identifying the architecture.

For example, if your ld command line included `‘-ACA’` as well as `‘-ltry’`, the linker would look (in its built-in search paths, and in any paths you specify with `‘-L’`) for a library with the names

```
try
libtry.a
tryca
libtryca.a
```

The first two possibilities would be considered in any event; the last two are due to the use of `‘-ACA’`.

You can meaningfully use `‘-A’` more than once on a command line, since the 960 architecture family allows combination of target architectures; each use will add another pair of name variants to search for when `‘-l’` specifies a library.

ld supports the `‘-relax’` option for the i960 family. If you specify `‘-relax’`, ld finds all `balx` and `calx` instructions whose targets are within 24 bits, and turns them into 24-bit program-counter relative `bal` and `cal` instructions, respectively. ld also turns `cal` instructions into `bal` instructions when it determines that the target subroutine is a leaf routine (that is, the target subroutine does not itself call any subroutines).

BFD libraries

The linker accesses object and archive files using the BFD libraries (libraries whose name comes from binary file descriptors).

The following documentation discusses the BFD libraries and how to use them.

- “How it works: an outline of BFD” on page 245
- “Information loss” on page 246
- “The BFD canonical object-file format” on page 247

BFD libraries allow the linker to use the same routines to operate on object files whatever the object file format. A different object file format can be supported simply by creating a new BFD back end and adding it to the library. To conserve runtime memory, however, the linker and associated tools are usually configured to support only a subset of the object file formats available. To list all the formats available for your configuration, use `objdump -i` (see “`objdump`” in *The GNU Binary Utilities*).

As with most implementations, BFD is a compromise between several conflicting requirements. The major factor influencing BFD design was efficiency: any time used converting between formats is time which would not have been spent had BFD not been involved. This is partly offset by abstraction payback; since BFD simplifies applications and back ends, more time and care may be spent optimizing algorithms for a greater speed.

One minor artifact of the BFD solution which you should bear in mind is the potential

for information loss. There are two places where useful information can be lost using the BFD mechanism: during conversion and during output. See “Information loss” on page 246.

How it works: an outline of BFD

When an object file is opened, BFD subroutines automatically determine the format of the input object file. They then build a descriptor in memory with pointers to routines that will be used to access elements of the object file's data structures.

As different information from the object files is required, BFD reads from different sections of the file and processes them. For example, a very common operation for the linker is processing symbol tables. Each BFD back end provides a routine for converting between the object file's representation of symbols and an internal canonical format. When the linker asks for the symbol table of an object file, it calls through a memory pointer to the routine from the relevant BFD back end which reads and converts the table into a canonical form. The linker then operates upon the canonical form. When the link is finished and the linker writes the output file's symbol table, another BFD back end routine is called to take the newly created symbol table and convert it into the chosen output format.

Information loss

Information can be lost during output. The output formats supported by BFD do not provide identical facilities, and information which can be described in one form has nowhere to go in another format. One example of this is alignment information in `b.out`. There is nowhere in an `a.out` format file to store alignment information on the contained data, so when a file is linked from `b.out` and an `a.out` image is produced, alignment information will not propagate to the output file. (The linker will still use the alignment information internally, so the link is performed correctly).

Another example is COFF section names. COFF files may contain an unlimited number of sections, each one with a textual section name. If the target of the link is a format which does not have many sections (such as `a.out`) or has sections without names (such as the Oasys format), the link cannot be done simply. You can circumvent this problem by describing the desired input-to-output section mapping with the linker command language.

Information can be lost during canonicalization. The BFD internal canonical form of the external formats is not exhaustive; there are structures in input formats for which there is no direct representation internally. This means that the BFD back ends cannot maintain all possible data richness through the transformation between external to internal and back to external formats.

This limitation is only a problem when an application reads one format and writes another. Each BFD back end is responsible for maintaining as much data as possible, and the internal BFD canonical form has structures which are opaque to the BFD core, and exported only to the back ends. When a file is read in one format, the canonical form is generated for BFD and the application. At the same time, the back end saves away any information which may otherwise be lost. If the data is then written back in the same format, the back end routine will be able to use the canonical form provided by the BFD core as well as the information it prepared earlier. Since there is a great deal of commonality between back ends, there is no information lost when linking or copying big endian COFF to little endian COFF, or `a.out` to `b.out`. When a mixture of formats is linked, the information is only lost from the files whose format differs from the destination.

The BFD canonical object-file format

The greatest potential for loss of information occurs when there is the least overlap between the information provided by the source format, that stored by the canonical format, and that needed by the destination format. A brief description of the canonical form may help you understand which kinds of data you can count on preserving across conversions.

files

Information stored on a per-files basis includes target machine architecture, particular implementation format type, a demand pageable bit, and a write protected bit. Information like Unix magic numbers is not stored here—only the magic numbers' meaning, so a `ZMAGIC` file would have both the demand pageable bit and the write protected text bit set. The byte order of the target is stored on a per-file basis, so that big- and little-endian object files may be used with one another.

sections

Each section in the input file contains the name of the section, the section's original address in the object file, size and alignment information, various flags, and pointers into other BFD data structures.

symbols

Each symbol contains a pointer to the information for the object file which originally defined it, its name, its value, and various flag bits. When a BFD back end reads in a symbol table, it relocates all symbols to make them relative to the base of the section where they were defined.

Doing this ensures that each symbol points to its containing section. Each symbol also has a varying amount of hidden private data for the BFD back end. Since the symbol points to the original file, the private data format for that symbol is accessible. `ld` can operate on a collection of symbols of wildly different formats without problems.

Normal global and simple local symbols are maintained on output, so an output file (no matter its format) will retain symbols pointing to functions and to global, static, and common variables. Some symbol information is not worth retaining; in `a.out`, type information is stored in the symbol table as long symbol names.

This information would be useless to most COFF debuggers; the linker has command line switches to allow users to throw it away.

There is one word of type information within the symbol, so if the format supports symbol type information within symbols (for example, COFF, IEEE, Oasys) and the type is simple enough to fit within one word (nearly everything but

aggregates), the information will be preserved.

relocation level

Each canonical BFD relocation record contains a pointer to the symbol to relocate to, the offset of the data to relocate, the section the data is in, and a pointer to a relocation type descriptor. Relocation is performed by passing messages through the relocation type descriptor and the symbol pointer. Therefore, relocations can be performed on output data using a relocation method that is only available in one of the input formats. For instance, Oasys provides a byte relocation format. A relocation record requesting this relocation type would point indirectly to a routine to perform this, so the relocation may be performed on a byte being written to a 68k COFF file, even though 68k COFF has no such relocation type.

line numbers

Object formats can contain, for debugging purposes, some form of mapping between symbols, source line numbers, and addresses in the output file. These addresses have to be relocated along with the symbol information. Each symbol with an associated list of line number records points to the first record of the list. The head of a line number list consists of a pointer to the symbol, which allows finding out the address of the function whose line number is being described. The rest of the list is made up of pairs: offsets into the section and line numbers. Any format which can simply derive this information can pass it successfully between formats (COFF, IEEE and Oasys).

6

MRI compatible script files for the GNU linker

To aid users making the transition to GNU `ld` from the MRI linker, `ld` can use MRI compatible linker scripts as an alternative to the more general-purpose linker scripting language described in “`ld` command line options” on page 190. The following documentation describes the use of the scripts.

MRI compatible linker scripts have a much simpler command set than the scripting language otherwise used with `ld`. GNU `ld` supports the most commonly used MRI linker commands; these commands are described in the following.

In general, MRI scripts aren’t of much use with the `a.out` object file format, since it only has three sections and MRI scripts lack some features to make use of them.

You can specify a file containing an MRI-compatible script using the ‘`-c`’ command-line option.

Each command in an MRI-compatible script occupies its own line; each command line starts with the keyword that identifies the command (though blank lines are also allowed for punctuation). If a line of an MRI-compatible script begins with an unrecognized keyword, `ld` issues a warning message, but continues processing the script. Lines beginning with ‘`*`’ are comments. You can write these commands using all upper-case letters, or all lower case; for example, ‘`chip`’ is the same as ‘`CHIP`’. The following list shows only the upper-case form of each command.

ABSOLUTE *secname*

ABSOLUTE *secname, secname, ... secname*

Normally, `ld` includes in the output file all sections from all the input files. However, in an MRI-compatible script, you can use the **ABSOLUTE** command to restrict the sections that will be present in your output program. If the **ABSOLUTE** command is used at all in a script, then only the sections named explicitly in **ABSOLUTE** commands will appear in the linker output. You can still use other input sections (whatever you select on the command line, or using **LOAD**) to resolve addresses in the output file.

ALIAS *out-secname, in-secname*

Use this command to place the data from input section *in-secname* in a section called *out-secname* in the linker output file. *in-secname* may be an integer.

ALIGN *secname=expression*

Align the section called *secname* to *expression*. The *expression* should be a power of two.

BASE *expression*

Use the value of *expression* as the lowest address (other than absolute addresses) in the output file.

CHIP *expression*

CHIP *expression, expression*

This command does nothing; it is accepted only for compatibility.

END

This command does nothing whatever; it's only accepted for compatibility.

FORMAT *output-format*

Similar to the **OUTPUT_FORMAT** command in the more general linker language, but restricted to one of these output formats:

- S-records, if *output-format* is 'S'
- IEEE, if *output-format* is 'IEEE'
- COFF (the 'coff-m68k' variant in BFD), if *output-format* is 'COFF'

LIST *anything...*

Print (to the standard output file) a link map, as produced by the `ld` command-line option '-M'.

The keyword **LIST** may be followed by anything on the same line, with no change in its effect.

LOAD *filename*

LOAD *filename, filename, ... filename*

Include one or more object file *filename* in the link; this has the same effect as specifying *filename* directly on the `ld` command line.

`NAME output-name`

output-name is the name for the program produced by `ld`; the MRI-compatible command `NAME` is equivalent to the command-line option ‘`-o`’ or the general script language command `OUTPUT`.

`ORDER secname, secname, ... secname`

`ORDER secname secname secname`

Normally, `ld` orders the sections in its output file in the order in which they first appear in the input files. In an MRI-compatible script, you can override this ordering with the `ORDER` command. The sections you list with `ORDER` will appear first in your output file, in the order specified.

`PUBLIC name=expression`

`PUBLIC name, expression`

`PUBLIC name expression`

Supply a value (*expression*) for external symbol name used in the linker input files.

`SECT secname, expression`

`SECT secname=expression`

`SECT secname expression`

You can use any of these three forms of the `SECT` command to specify the start address (*expression*) for section *sec-name*. If you have more than one `SECT` statement for the same *sec-name*, only the *first* sets the start address.

Using binutils

Copyright © 1991-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the documentation entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Free Software Foundation

59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

This documentation has been prepared by Cygnus Technical Publications; contact the Cygnus Technical Publications staff: doc@cygnus.com.

1

Overview of `binutils`, the GNU binary utilities

The following documentation contains basic descriptions for the GNU binary utilities:

- `ar`
Creates, modifies, and extracts from archives; see “`ar` utility” on page 257
- `nm`
Lists symbols from object files; see “`nm` utility” on page 263
- `objcopy`
Copies and translates object files; see “`objcopy` utility” on page 267
- `objdump`
Displays information from object files; see “`objdump` utility” on page 272
- `ranlib`
Generates index to archive contents; see “`ranlib` utility” on page 277
- `size`
Lists file section sizes and total size; see “`size` utility” on page 278
- `strings`
Lists printable strings from files; see “`strings` utility” on page 280
- `strip`
Discards symbols; see “`strip` utility” on page 281
- `c++filt`
Demangles encoded C++ symbols; see “`c++filt` utility” on page 283

- `addr2line`
Converts addresses into file names and line numbers; see “`addr2line` utility” on page 285
- `nlmconv`
Converts object code into a Netware Loadable Module (NLM); see “`nlmconv` utility” on page 287
- `windres`
Manipulates Windows resources; see “`windres` utility” on page 289

ar utility

```
ar [-]p[mod [relpos]] archive [member ...]
ar -M [ <mri-script ]
```

The GNU `ar` program creates, modifies, and extracts from archives. An *archive* is a single file holding a collection of other files in a structure that makes it possible to retrieve the original individual files (called *members* of the archive).

The original files' contents, mode (permissions), timestamp, owner, and group are preserved in the archive, and can be restored on extraction.

GNU `ar` can maintain archives whose members have names of any length; however, depending on how `ar` is configured on your system, a limit on member-name length may be imposed for compatibility with archive formats maintained with other tools. If it exists, the limit is often 15 characters (typical of formats related to `a.out`) or 16 characters (typical of formats related to `coff`).

`ar` is considered a binary utility because archives of this sort are most often used as libraries holding commonly needed subroutines.

`ar` creates an index to the symbols defined in relocatable object modules in the archive when you specify the modifier 's'. Once created, this index is updated in the archive whenever `ar` makes a change to its contents (save for the 'q' update operation). An archive with such an index speeds up linking to the library, and allows routines in the library to call each other without regard to their placement in the archive.

You may use 'nm -s' or 'nm --print-armap' to list this index table. If an archive lacks the table, another form of `ar` called `ranlib` can be used to add only the table.

GNU `ar` is designed to be compatible with two different facilities. You can control its activity using command-line options like the different varieties of `ar` on Unix systems; or, if you specify the single command-line option '-M', you can control it with a script supplied via standard input, like the MRI *librarian* program.

Controlling ar on the command line

```
ar [-]p[mod [relpos]] archive [member ...]
```

When you use `ar` in the Unix style, `ar` insists on at least two arguments to execute: one keyletter specifying the *operation* (optionally accompanied by other keyletters specifying *modifiers*), and the archive name to act on.

Most operations can also accept further member arguments, specifying particular files to operate on. GNU `ar` allows you to mix the operation code `p` and modifier flags `mod` in any order, within the first command-line argument. If you wish, you may begin the first command-line argument with a dash.

The `p` keyletter specifies what operation to execute; it may be any of the following, but

you must specify only one of them:

d

Delete modules from the archive. Specify the names of modules to be deleted as *member* . . . ; the archive is untouched if you specify no files to delete.

If you specify the 'v' modifier, ar lists each module as it is deleted.

m

Move members in an archive.

The ordering of members in an archive can make a difference in how programs are linked using the library, if a symbol is defined in more than one member.

If no modifiers are used with m, any members you name in the *member* arguments are moved to the *end* of the archive; you can use the 'a', 'b', or 'i' modifiers to move them to a specified place instead.

p

Print the specified members of the archive, to the standard output file. If the 'v' modifier is specified, show the member name before copying its contents to standard output.

If you specify no *member* arguments, all the files in the archive are printed.

q

Quick append; add the files *member* . . . to the end of *archive*, without checking for replacement.

The modifiers 'a', 'b', and 'i' do *not* affect this operation; new members are always placed at the end of the archive. The modifier 'v' makes ar list each file as it is appended. Since the point of this operation is speed, the archive's symbol table index is not updated, even if it already existed; you can use 'ar s' or ranlib explicitly to update the symbol table index.

r

Replacement; inserts the files *member* . . . into *archive*. This operation differs from 'q' in that any previously existing members are deleted if their names match those being added.

If one of the files named in *member* . . . does not exist, ar displays an error message, and leaves undisturbed any existing members of the archive matching that name.

By default, new members are added at the end of the file; but you may use one of the modifiers 'a', 'b', or 'i' to request placement relative to some existing member.

The modifier 'v' used with this operation elicits a line of output for each file inserted, along with one of the letters 'a' or 'r' to indicate whether the file was appended (no old member deleted) or replaced.

- t Display a *table* listing the contents of *archive*, or those of the files listed in *member*... that are present in the archive. Normally only the member name is shown; if you also want to see the modes (permissions), timestamp, owner, group, and size, you can request that by also specifying the ‘v’ modifier.
- If you do not specify a *member*, all files in the archive are listed.
- If there is more than one file with the same name (for instance, ‘file’) in an archive (for instance, ‘b.a’), ‘ar t b.a file’ lists only the first instance; to see them all, you must ask for a complete listing—in the earlier example, ‘ar t b.a’.
- x Extract members (named *member*) from the archive. You can use the ‘v’ modifier with this operation, to request that ar list each name as it extracts it. If you do not specify a *member*, all files in the archive are extracted.
- A number of modifiers (*mod*) may immediately follow the *p* keyletter, to specify variations on an operation’s behavior:
- a Add new files *after* an existing member of the archive. If you use the modifier ‘a’, the name of an existing archive member must be present as the *relpos* argument, before the archive specification.
- b Add new files *before* an existing member of the archive. If you use the modifier ‘b’, the name of an existing archive member must be present as the *relpos* argument, before the archive specification. (same as ‘i’).
- c *Create* the archive. The specified *archive* is always created if it did not exist, when you request an update. But a warning is issued unless you specify in advance that you expect to create it, by using this modifier.
- f Truncate names in the archive. GNU ar will normally permit file names of any length. This will cause it to create archives which are not compatible with the native ar program on some systems. If this is a concern, the ‘f’ modifier may be used to truncate file names when putting them in the archive.
- i Insert new files *before* an existing member of the archive. If you use the modifier ‘i’, the name of an existing archive member must be present as the *relpos* argument, before the *archive* specification. (same as ‘b’).
- l This modifier is accepted but not used.
- o Preserve the *original* dates of members when extracting them. If you do not

specify this modifier, files extracted from the archive are stamped with the time of extraction.

s

Write an object-file index into the archive, or update an existing one, even if no other change is made to the archive. You may use this modifier flag either with any operation, or alone. Running ‘ar s’ on an archive is equivalent to running ‘ranlib’ on it.

u

Normally, ‘ar r’ ... inserts all files listed into the archive. If you would like to insert *only* those of the files you list that are newer than existing members of the same names, use this modifier. The ‘u’ modifier is allowed only for the operation ‘r’ (*replace*). In particular, the combination ‘qu’ is not allowed, since checking the timestamps would lose any speed advantage from the operation ‘q’.

v

This modifier requests the verbose version of an operation. Many operations display additional information, such as file-names processed, when the modifier ‘v’ is appended.

V

This modifier shows the version number of ar.

Controlling ar with a script

```
ar -M [ < script ]
```

If you use the single command-line option ‘-M’ with ar, you can control its operation with a rudimentary command language. This form of ar operates interactively if standard input is coming directly from a terminal. During interactive use, ar prompts for input (the prompt is ‘AR >’), and continues executing even after errors. If you redirect standard input to a script file, no prompts are issued, and ar abandons execution (with a nonzero exit code) on any error.

The ar command language is *not* designed to be equivalent to the command-line options; in fact, it provides somewhat less control over archives. The only purpose of the command language is to ease the transition to GNU ar for developers who already have scripts written for the MRI *librarian* program.

The syntax for the ar command language is straightforward:

- Commands are recognized in upper or lower case; for example, LIST is the same as list. In the following descriptions, commands are shown in upper case for clarity.
- A single command may appear on each line; it is the first word on the line.
- Empty lines are allowed, and have no effect.

- Comments are allowed; text after either of the characters ‘*’ or ‘;’ is ignored.
- Whenever you use a list of names as part of the argument to an `ar` command, you can separate the individual names with either commas or blanks. Commas are shown in the following explanations, for clarity.
- ‘+’ is used as a line continuation character; if ‘+’ appears at the end of a line, the text on the following line is considered part of the current command.

The following are the commands you can use in `ar` scripts, or when using `ar` interactively. Three of them have special significance:

`OPEN` or `CREATE` specify a *current* archive, which is a temporary file required for most of the other commands.

`SAVE` commits the changes so far specified by the script. Prior to `SAVE`, commands affect only the temporary copy of the current archive.

`ADDLIB archive`

`ADDLIB archive(module, module, ...module)`

Add all the contents of *archive* (or, if specified, each named *module* from *archive*) to the current archive.

Requires prior use of `OPEN` or `CREATE`.

`ADDMOD member, member, ...member`

Add each named *member* as a module in the current archive.

Requires prior use of `OPEN` or `CREATE`.

`CLEAR`

Discard the contents of the current archive, canceling the effect of any operations since the last `SAVE`. May be executed (with no effect) even if no current archive is specified.

`CREATE archive`

Creates an archive, and makes it the current archive (required for many other commands). The new archive is created with a temporary name; it is not actually saved as *archive* until you use `SAVE`. You can overwrite existing archives; similarly, the contents of any existing file named *archive* will not be destroyed until `SAVE`.

`DELETE module, module, ...module`

Delete each listed *module* from the current archive.

Equivalent to ‘`ar -d archive module ...module`’.

Requires prior use of `OPEN` or `CREATE`.

`DIRECTORY archive(module, ...module)`

`DIRECTORY archive(module, ...module) outputfile`

List each named *module* present in *archive*. The separate command `VERBOSE` specifies the form of the output: when verbose output is off, output is like that of

`'ar -t archive module ...'`. When verbose output is on, the listing is like `'ar -tv archive module ...'`.

Output normally goes to the standard output stream; however, if you specify *outputfile* as a final argument, `ar` directs the output to that file.

END

Exit from `ar`, with a 0 exit code to indicate successful completion. This command does not save the output file; if you have changed the current archive since the last `SAVE` command, those changes are lost.

EXTRACT *module, module, ...module*

Extract each named *module* from the current archive, writing them into the current directory as separate files.

Equivalent to `'ar -x archive module ...'`.

Requires prior use of `OPEN` or `CREATE`.

LIST

Display full contents of the current archive, in *verbose* style regardless of the state of `VERBOSE`. The effect is like `'ar tv archive'`. This single command is a GNU `ld` enhancement, rather than present for MRI compatibility.

Requires prior use of `OPEN` or `CREATE`.

OPEN *archive*

Opens an existing archive for use as the current archive (required for many other commands). Any changes as the result of subsequent commands will not actually affect *archive* until you next use `SAVE`.

REPLACE *module, module, ...module*

In the current archive, replace each existing *module* (named in the `REPLACE` arguments) from files in the current working directory. To execute this command without errors, both the file, and the module in the current archive, must exist.

Requires prior use of `OPEN` or `CREATE`.

VERBOSE

Toggle an internal flag governing the output from `DIRECTORY`. When the flag is on, `DIRECTORY` output matches output from `'ar -tv' ...`

SAVE

Commit your changes to the current archive, and actually save it as a file with the name specified in the last `CREATE` or `OPEN` command.

Requires prior use of `OPEN` or `CREATE`.

nm utility

```
nm [ -a | --debug-syms ] [ -g | --extern-only ]
    [ -B ] [ -C | --demangle ] [ -D | --dynamic ]
    [ -s | --print-armap ] [ -A | -o | --print-file-name ]
    [ -n | -v | --numeric-sort ] [ -p | --no-sort ]
    [ -r | --reverse-sort ] [ --size-sort ]
    [ -u | --undefined-only ]
    [-t radix | --radix= radix ]
    [ -P | --portability ] [ --target= bfdname ]
    [-f format | --format= format ]
    [ --defined-only ] [ -l | --line-numbers ]
    [ --no-demangle ] [ -V | --version ]
    [ --help ] [ objfile ...]
```

GNU `nm` lists the symbols from object files *objfile* ... If no object files are listed as arguments, `nm` assumes ‘a.out’.

For each symbol, `nm` shows:

- The symbol value, in the radix selected by the following options, or hexadecimal by default.
- The symbol type. At least the following types are used; others are, as well, depending on the object file format. If lowercase, the symbol is local; if uppercase, the symbol is global (external).

- A The symbol’s value is absolute, and will not be changed by further linking.
- B The symbol is in the uninitialized data section (known as BSS).
- C The symbol is common. Common symbols are uninitialized data. When linking, multiple common symbols may appear with the same name. If the symbol is defined anywhere, the common symbols are treated as undefined references. For more details on common symbols, see the discussion of `--warn-common` in “ld command line options” on page 190 in *Using LD*.
- D The symbol is in the initialized data section.
- G The symbol is in an initialized data section for small objects. Some object file formats permit more efficient access to small data objects, such as a global `int` variable as opposed to a large global array.
- I The symbol is an indirect reference to another symbol. This is a GNU extension to the a.out object file format which is rarely used.

- N
The symbol is a debugging symbol.
- R
The symbol is in a read only data section.
- S
The symbol is in an uninitialized data section for small objects.
- T
The symbol is in the text (code) section.
- U
The symbol is undefined.
- W
The symbol is weak. When a weak defined symbol is linked with a normal defined symbol, the normal defined symbol is used with no error. When a weak undefined symbol is linked and the symbol is not defined, the value of the weak symbol becomes zero with no error.
- The symbol is a `stabs` symbol in an `a.out` object file. In this case, the next values printed are the `stabs other` field, the `stabs desc` field, and the `stab type`. `stabs` symbols are used to hold debugging information; for more information, see “`.stabd`, `.stabn`, and `stabs` directives” on page 83 in *Using as*.
- ?
The symbol type is unknown, or object file format specific.

■ The symbol name.

The long and short forms of options, shown here as alternatives, are equivalent.

-A

-O

--print-file-name

Precede each symbol by the name of the input file (or archive element) in which it was found, rather than identifying the input file once only, before all of its symbols.

-a

--debug-syms

Display all symbols, even debugger-only symbols; normally these are not listed.

-B

The same as ‘--format=bsd’ (for compatibility with the MIPS `nm`).

-C

--demangle

Decode (*demangle*) low-level symbol names into user-level names. Besides removing any initial underscore prepended by the system, this makes C++

function names readable. See “c++filt utility” on page 283 for more information on demangling.

--no-demangle

Do not *demangle* low-level symbol names. This is the default.

-D

--dynamic

Display the dynamic symbols rather than the normal symbols. This is only meaningful for dynamic objects, such as certain types of shared libraries.

-f *format*

--format=*format*

Use the output format, *format*, which can be `bsd`, `sysv`, or `posix`. The default is `bsd`. Only the first character of *format* is significant; it can be either upper or lower case.

-g

--extern-only

Display only external symbols.

-l

--line-numbers

For each symbol, use debugging information to try to find a filename and line number. For a defined symbol, look for the line number of the address of the symbol. For an undefined symbol, look for the line number of a relocation entry that refers to the symbol. If line number information can be found, print it after the other symbol information.

-n

-v

--numeric-sort

Sort symbols numerically by their addresses, rather than alphabetically by their names.

-p

--no-sort

Do not bother to sort the symbols in any order; print them in the order encountered.

-P

--portability

Use the POSIX.2 standard output format instead of the de-fault format. Equivalent to ‘-f `posix`’.

-s

--print-armap

When listing symbols from archive members, include the index: a mapping (stored in the archive by `ar` or `ranlib`) of which modules contain definitions for which names.

-r

`--reverse-sort`

Reverse the order of the sort (whether numeric or alphabetic); let the last come first.

`--size-sort`

Sort symbols by `size`. `size` is computed as the difference between the value of the symbol and the value of the symbol with the next higher value. The size of the symbol is printed, rather than the value.

`-t radix`

`--radix=radix`

Use `radix` as the radix for printing the symbol values. It must be 'd' for decimal, 'o' for octal, or 'x' for hexadecimal.

`--target=bfdname`

Specify an object code format other than your system's default format. See "Target selection" on page 294, for more information.

`-u`

`--undefined-only`

Display only undefined symbols (those external to each object file).

`--defined-only`

Display only defined symbols for each object file.

`-V`

`--version`

Show the version number of `nm` and exit.

`--help`

Show a summary of the options to `nm` and exit.

objcopy utility

```
objcopy [ -F bfdname | --target=bfdname ]
[ -I bfdname | --input-target=bfdname ]
[ -O bfdname | --output-target=bfdname ]
[ -S | --strip-all ] [ -g | --strip-debug ]
[ -K symbolname | --keep-symbol=symbolname ]
[ -N symbolname | --strip-symbol=symbolname ]
[ -L symbolname | --localize-symbol=symbolname ]
[ -W symbolname | --weaken-symbol=symbolname ]
[ -x | --discard-all ] [ -X | --discard-locals ]
[ -b byte | --byte=byte ]
[ -i interleave | --interleave=interleave ]
[ -R sectionname | --remove-section=sectionname ]
[ -p | --preserve-dates ] [ --debugging ]
[ --gap-fill=val ] [ --pad-to=address ]
[ --set-start=val ] [ --adjust-start=incr ]
[ --adjust-vma=incr ]
[ --adjust-section-vma=section{=,+,-}val ]
[ --adjust-warnings ] [ --no-adjust-warnings ]
[ --set-section-flags=section=flags ]
[ --add-section=sectionname=filename ]
[ --change-leading-char ] [ --remove-leading-char ]
[ --weaken ]
[ -v | --verbose ] [ -V | --version ] [ --help ]
infile [outfile]
```

The GNU `objcopy` utility copies the contents of an object file to another. `objcopy` uses the GNU BFD Library to read and write the object files. It can write the destination object file in a format different from that of the source object file. The exact behavior of `objcopy` is controlled by command-line options.

`objcopy` creates temporary files to do its translations and deletes them afterward. `objcopy` uses BFD to do all its translation work; it has access to all the formats described in BFD and is able to recognize most formats without being told explicitly. See “BFD libraries” on page 243 and “The BFD canonical object-file format” on page 247 in *Using LD*.

`objcopy` can be used to generate S-records by using an output target of ‘srec’ (use ‘-O srec’).

`objcopy` can be used to generate a raw binary file by using an output target of ‘binary’ (meaning ‘-O binary’). When `objcopy` generates a raw binary file, it will essentially produce a memory dump of the contents of the input object file. All symbols and relocation information will be discarded. The memory dump will start at the virtual address of the lowest section copied into the output file.

When generating an S-record or a raw binary file, it may be helpful to use `-s` to

remove sections containing debugging information. In some cases ‘-R’ will be useful to remove sections containing information that is not needed by the binary file.

infile

outfile

The source and output files, respectively. If you do not specify *outfile*, *objcopy* creates a temporary file and destructively renames the result with the name of *infile*.

-I *bfdname*

--input-target=*bfdname*

Consider the source file’s object format to be *bfdname*, rather than attempting to deduce it. See “Target selection” on page 294 for more information.

-O *bfdname*

--output-target=*bfdname*

Write the output file using the object format, *bfdname*. See “Target selection” on page 294 for more information.

-F *bfdname*

--target=*bfdname*

Use *bfdname* as the object format for both the input and the output file; i.e., simply transfer data from source to destination with no translation. See “Target selection” on page 294 for more information.

-R *sectionname*

--remove-section=*sectionname*

Remove any section named *sectionname* from the output file. This option may be given more than once.

NOTE: Using this option inappropriately may make the output file unusable.

-S

--strip-all

Do not copy relocation and symbol information from the source file.

-g

--strip-debug

Do not copy debugging symbols from the source file.

--strip-unnneeded

Strip all symbols that are not needed for relocation processing.

-K *symbolname*

--keep-symbol=*symbolname*

Copy only symbol *symbolname* from the source file. This option may be given more than once.

```

-N symbolname
--strip-symbol=symbolname
    Do not copy symbol symbolname from the source file. This option may be given
    more than once, and may be combined with strip options other than -K.
-L symbolname
--localize-symbol=symbolname
    Make symbol symbolname local to the file, so that it is not visible externally. This
    option may be given more than once.
-W symbolname
--weaken-symbol=symbolname
    Make symbol symbolname weak. This option may be given more than once.
-x
--discard-all
    Do not copy non-global symbols from the source file.
-X
--discard-locals
    Do not copy compiler-generated local symbols. (These usually start with 'L' or
    '.')
-b byte
--byte=byte
    Keep only every byte of the input file (header data is not affected). byte can be
    in the range from 0 to interleave-1, where interleave is given by the '-i' or
    '--interleave' option, or the default of 4. This option is useful for creating files
    to program ROM. It is typically used with an srec output target.
-i interleave
--interleave=interleave
    Only copy one out of every interleave bytes. Select which byte to copy with the
    '-b' or '--byte' option. The default is 4. objcopy ignores this option if you do
    not specify either '-b' or '--byte'.
-p
--preserve-dates
    Set the access and modification dates of the output file to be the same as those of
    the input file.
--debugging
    Convert debugging information, if possible. This is not the default because only
    certain debugging formats are supported, and the conversion process can be time
    consuming.
--gap-fill val
    Fill gaps between sections with val. This is done by increasing the size of the
    section with the lower address, and filling in the extra space created with val.
--pad-to address
    Pad the output file up to the virtual address address. This is done by increasing

```

the size of the last section. The extra space is filled in with the value specified by `--gap-fill` (default zero).

`--set-start val`

Set the address of the new file to *val*. Not all object file formats support setting the start address.

`--adjust-start incr`

Adjust the start address by adding *incr*. Not all object file formats support setting the start address.

`--adjust-vma incr`

Adjust the address of all sections, as well as the start address, by adding *incr*. Some object file formats do not permit section addresses to be changed arbitrarily. Note that this does not relocate the sections; if the program expects sections to be loaded at a certain address, and this option is used to change the sections such that they are loaded at a different address, the program may fail.

`--adjust-section-vma section{=,+,-} val`

Set or adjust the address of the named *section*. If `'='` is used, the section address is set to *val*. Otherwise, *val* is added to or subtracted from the section address.

See the comments under `--adjust-vma`. If *section* does not exist in the input file, a warning will be issued, unless `--no-adjust-warnings` is used.

`--adjust-warnings`

If `--adjust-section-vma` is used, and the named section does not exist, issue a warning. This is the default.

`--no-adjust-warnings`

Do not issue a warning if `--adjust-section-vma` is used, even if the named section does not exist.

`--set-section-flags section=flags`

Set the flags for the named section. The *flags* argument is a comma separated string of flag names. The recognized names are `'alloc'`, `'load'`, `'readonly'`, `'code'`, `'data'`, and `'rom'`. Not all flags are meaningful for all object file formats.

`--add-section sectionname=filename`

Add a new section named *sectionname* while copying the file. The contents of the new section are taken from the file *filename*. The size of the section will be the size of the file. This option only works on file formats which can support sections with arbitrary names.

`--change-leading-char`

Some object file formats use special characters at the start of symbols. The most common such character is underscore, which compilers often add before every symbol. This option tells `objcopy` to change the leading character of every symbol when it converts between object file formats. If the object file formats use the same leading character, this option has no effect. Otherwise, it will add a character, or remove a character, or change a character, as appropriate.

`--remove-leading-char`
If the first character of a global symbol is a special symbol leading character used by the object file format, remove the character. The most common symbol leading character is underscore. This option will remove a leading underscore from all global symbols. This can be useful if you want to link together objects of different file formats with different conventions for symbol names.

`--weaken`
Change all global symbols in the file to be weak. This can be useful when building an object that will be linked against other objects using the `-R` option to the linker. This option is only effective when using an object file format that supports weak symbols.

`-V`
`--version`
Show the version number of `objcopy`.

`-v`
`--verbose`
Verbose output: list all object files modified. In the case of archives, `'objcopy -v'` lists all members of the archive.

`--help`
Show a summary of the options to `objcopy`.

objdump utility

```
objdump [ -a | --archive-headers ]
        [ -b bfdname | --target=bfdname ] [ --debugging ]
        [ -d | --disassemble ] [ -D | --disassemble-all ]
        [ -EB | -EL | --endian={big | little} ]
        [ -f | --file-headers ]
        [ -h | --section-headers | --headers ]
        [ -i | --info ]
        [ -j section | --section=section ]
        [ -l | --line-numbers ] [ -S | --source ]
        [ -m machine | --architecture=machine ]
        [ -r | --reloc ] [ -R | --dynamic-reloc ]
        [ -s | --full-contents ] [ --stabs ]
        [ -t | --syms ] [ -T | --dynamic-syms ]
        [ -x | --all-headers ]
        [ -w | --wide ] [ --start-address=address ]
        [ --stop-address=address ]
        [ --prefix-addresses ] [ --show-raw-insn ]
        [ --adjust-vma=offset ]
        [ --version ]
        [ --help ]
objfile...
```

`objdump` displays information about one or more object files. The options control what particular information to display. This information is mostly useful to programmers who are working on the compilation tools, as opposed to programmers who just want their program to compile and work.

objfile... are the object files to be examined. When you specify archives, `objdump` shows information on each of the member object files.

The long and short forms of options, shown here as alternatives, are equivalent. At least one option besides ‘-l’ must be given.

-a

--archive-header

If any of the *objfile* files are archives, display the archive header information (in a format similar to ‘ls -l’). Besides the information you could list with ‘ar tv’, ‘`objdump -a`’ shows the object file format of each archive member.

--adjust-vma=*offset*

When dumping information, first add *offset* to all the section addresses. This is useful if the section addresses do not correspond to the symbol table, which can happen when putting sections at particular addresses when using a format that can not represent section addresses, such as `a.out`.

`-b bfdname`

`--target=bfdname`

Specify that the object-code format for the object files is BFD-name. This option may not be necessary; `objdump` can automatically recognize many formats.

`objdump -b oasys -m vax -h fu.o`

The previous example displays summary information from the section headers (`-h`) of `fu.o`, which is explicitly identified (`-m`) as a Vax object file in the format produced by Oasys compilers. You can list the formats available with the `-i` option. See “Target selection” on page 294 for more information.

`-C`

`--demangle`

Decode (demangle) low-level symbol names into user-level names. Besides removing any initial underscore prepended by the system, this makes C++ function names readable. See “c++filt utility” on page 283 for more information on demangling.

`--debugging`

Display debugging information. This attempts to parse debugging information stored in the file and print it out using a C like syntax. Only certain types of debugging information have been implemented.

`-d`

`--disassemble`

Display the assembler mnemonics for the machine instructions from *objfile*. This option only disassembles those sections which are expected to contain instructions.

`-D`

`--disassemble-all`

Like `-d`, but disassembles the contents of all sections, not just those expected to contain instructions.

`--prefix-addresses`

When disassembling, print the complete address on each line. This is the older disassembly format.

`--disassemble-zeroes`

Normally the disassembly output will skip blocks of zeroes.

This option directs the disassembler to disassemble those blocks, just like any other data.

`-EB`

`-EL`

`--endian={big | little}`

Specify the endianness of the object files. This only affects disassembly. This can be useful when disassembling a file format that does not describe endianness information, such as S-records.

-f
--file-header
 Display summary information from the overall header of each of the *objfile* files.

-h
--section-header
--header
 Display summary information from the section headers of the object file.
 File segments may be relocated to nonstandard addresses, for example by using the '-Ttext', '-Tdata', or '-Tbss' options to ld. However, some object file formats, such as a.out, do not store the starting address of the file segments. In those situations, although ld relocates the sections correctly, using 'objdump -h' to list the file section headers cannot show the correct addresses. Instead, it shows the usual addresses, which are implicit for the target.

--help
 Print a summary of the options to objdump and exit.

-i
--info
 Display a list showing all architectures and object formats available for specification with '-b' or '-m'.

-j *name*
--section=*name*
 Display information only for section *name*.

-l
--line-numbers
 Label the display (using debugging information) with the filename and source line numbers corresponding to the object code shown. Only useful with '-d' or '-D'.

-m *machine*
--architecture=*machine*
 Specify that the object files, *objfile*, are for the intended architecture, *machine*. You can list available architectures using the '-i' option.

-r
--reloc
 Print the relocation entries of the file. If used with '-d' or '-D' options, the relocations are printed interspersed with the disassembly.

-R
--dynamic-reloc
 Print the dynamic relocation entries of the file. This is only meaningful for dynamic objects, such as certain types of shared libraries.

```

-s
--full-contents
    Display the full contents of any sections requested.
-S
--source
    Display source code intermixed with disassembly, if possible.
    Implies '-d'.
--show-raw-insn
    When disassembling instructions, print the instruction in hex as well as in
    symbolic form. Not all targets handle this correctly yet.
--no-show-raw-insn
    When disassembling instructions, do not print the instruction bytes. This is the
    default when using --prefix-addresses.
--stabs
    Display the full contents of any sections requested. Display the contents of the
    .stab and .stab.index and .stab.excl sections from an ELF file. This is only
    useful on systems (such as Solaris 2.0) in which .stab debugging symbol-table
    entries are carried in an ELF section. In most other file formats, debugging
    symbol-table entries are interleaved with linkage symbols, and are visible in the
    '--syms' output. For more information on stabs symbols, see ".stabd, .stabsn,
    and.stabs directives" on page 83 in Using AS.
--start-address=address
    Start displaying data at the specified address. This affects the output of the -d, -r
    and -s options.
--stop-address=address
    Stop displaying data at the specified address. This affects the output of the -d, -r
    and -s options.
-t
--syms
    Print the symbol table entries of the file. This is similar to the information
    provided by the 'nm' program.
-T
--dynamic-syms
    Print the dynamic symbol table entries of the file. This is only meaningful for
    dynamic objects, such as certain types of shared libraries. This is similar to the
    information provided by the 'nm' program when given the '-D' ('--dynamic')
    option.
--version
    Print the version number of objdump and exit.

```

-x

--all-header

Display all available header information, including the symbol table and relocation entries.

Using '-x' is equivalent to specifying all of '-a -f -h -r -t'.

-w

--wide

Format some lines for output devices having more than 80 columns.

ranlib utility

```
ranlib [-vV] archive
```

ranlib generates an index to the contents of an archive and stores it in the archive. The index lists each symbol defined by a member of an archive that is a relocatable object file.

You may use ‘nm -s’ or ‘nm --print-armap’ to list this index.

An archive with such an index speeds up linking to the library and allows routines in the library to call each other without regard to their placement in the archive.

The GNU ranlib program is another form of GNU ar; running ranlib is completely equivalent to executing ‘ar -s’. See “ar utility” on page 257.

-v
-V

Show the version number of ranlib.

size utility

```
size [ -A | -B | --format=compatibility ]  
      [ --help ] [ -d | -o | -x | --radix=number ]  
      [ --target=bfdname ] [ -V | --version ]  
      objfile...
```

The GNU *size* utility lists the section sizes—and the total size—for each of the object or archive files *objfile* in its argument list. By default, one line of output is generated for each object file or each module in an archive.

objfile ... are the object files to be examined.

The command line options have the following meanings.

-A

-B

--format=*compatibility*

Using one of these options, you can choose whether the output from GNU *size* resembles output from System V *size* (using '-A', or '--format=sysv'), or Berkeley *size* (using '-B', or '--format=berkeley'). The default is the one-line format similar to Berkeley's.

The following is an example of the Berkeley (default) format of output from *size*.

```
size --format=Berkeley ranlib size  
text      data      bss      dec      hex      filename  
294880    81920      11592   388392   5ed28     ranlib  
294880    81920      11888   388688   5ee50     size
```

The following example shows the same data, but displayed closer to System V conventions.

```
size --format=SysV ranlib size  
ranlib :  
section      size      addr  
.text        294880     8192  
.data        81920     303104  
.bss         11592     385024  
Total        388392
```

```
size :  
section      size      addr  
.text        294880     8192  
.data        81920     303104  
.bss         11888     385024  
Total        388688
```

--help

Shows a summary of acceptable arguments and options.

-d
-o
-x

--radix=*number*

Using one of these options, you can control whether the size of each section is given: in decimal ('-d', or '--radix=10'); in octal ('-o', or '--radix=8'); or, in hexadecimal ('-x', or '--radix=16').

In '--radix=*number*', only the three values (8, 10, 16) are supported. The total size is always given in two radices; decimal and hexadecimal for '-d' or '-x' output; or octal and hexadecimal if you're using '-o'.

--target=*bfdname*

Specify that the object-code format for *objfile* is *bfdname*. This option may not be necessary; *size* can automatically recognize many formats. See “Target selection” on page 294 for more information.

-V

--version

Display the version number of *size*.

strings utility

```
strings [-afov] [-min-len] [-n min-len] [-t radix] [-]
        [--all] [--print-file-name] [--bytes=min-len]
        [--radix=radix] [--target=bfdname]
        [--help] [--version] file ...
```

For each *file* given, GNU `strings` prints the printable character sequences that are at least 4 characters long (or the number given with the following options) and are followed by an unprintable character. By default, it only prints the strings from the initialized and loaded sections of object files; for other types of files, it prints the strings from the whole file.

`strings` is mainly useful for determining the contents of non-text files.

`-a`

`--all`

-

Do not scan only the initialized and loaded sections of object files; scan the whole files.

`-f`

`--print-file-name`

Print the name of the file before each string.

`--help`

Print a summary of the program usage on the standard output and exit.

`-min-len`

`-n min-len`

`--bytes=min-len`

Print sequences of characters that are at least *min-len* characters long, instead of the default 4.

`-o`

Like `'-t o'`. Some other versions of `strings` have `'-o'` act like `'-t d'`. Since we can not be compatible with both ways, we simply chose one.

`-t radix`

`--radix=radix`

Print the offset within the file before each string. The single character argument specifies the radix of the offset—`'o'` for octal, `'x'` for hexadecimal, or `'d'` for decimal.

`--target=bfdname`

Specify an object code format other than your system's default format. See "Target selection" on page 294 for more information.

`-v`

`--version`

Print the program version number on the standard output and exit.

strip utility

```
strip [ -F bfdname | --target=bfdname ]
      [ -I bfdname | --input-target=bfdname ]
      [ -O bfdname | --output-target=bfdname ]
      [ -s | --strip-all ] [ -S | -g | --strip-debug ]
      [ -K symbolname | --keep-symbol=symbolname ]
      [ -N symbolname | --strip-symbol=symbolname ]
      [ -x | --discard-all ] [ -X | --discard-locals ]
      [ -R sectionname | --remove-section=sectionname ]
      [-o file ] [ -p | --preserve-dates ]
      [ -v | --verbose ] [ -V | --version ] [ --help ]
objfile ...
```

GNU `strip` discards all symbols from object files *objfile*. The list of object files may include archives. At least one object file must be given. `strip` modifies the files named in its argument, rather than writing modified copies under different names.

`-F bfdname`

`--target=bfdname`

Treat the original *objfile* as a file with the object code format *bfdname*, and rewrite it in the same format. See “Target selection” on page 294 for more information.

`--help`

Show a summary of the options to `strip` and exit.

`-I bfdname`

`--input-target=bfdname`

Treat the original *objfile* as a file with the object code format *bfdname*. See “Target selection” on page 294 for more information.

`-O bfdname`

`--output-target=bfdname`

Replace *objfile* with a file in the output format, *bfdname*. See “Target selection” on page 294 for more information.

`-R sectionname`

`--remove-section=sectionname`

Remove any section named *sectionname* from the output file. This option may be given more than once.

NOTE: Using this option inappropriately may make the output file unusable.

`-s`

`--strip-all`

Remove all symbols.

-g
-S
--strip-debug
 Remove debugging symbols only.
--strip-unneeded
 Remove all symbols that are not needed for relocation processing.
-K *symbolname*
--keep-symbol=*symbolname*
 Keep only symbol *symbolname* from the source file. This option may be given more than once.
-N *symbolname*
--strip-symbol=*symbolname*
 Remove symbol *symbolname* from the source file. This option may be given more than once, and may be combined with *strip* options other than -K.
-o *file*
 Put the stripped output in *file*, rather than replacing the existing file. When this argument is used, only one *objfile* argument may be specified.
-p
--preserve-dates
 Preserve the access and modification dates of the file.
-x
--discard-all
 Remove non-global symbols.
-X
--discard-locals
 Remove compiler-generated local symbols. (These usually start with 'L' or '.')-V
--version
 Show the version number for *strip*.
-v
--verbose
 Verbose output: list all object files modified. In the case of archives, '*strip -v*' lists all members of the archive.

c++filt utility

```
c++filt [ -_ | --strip-underscores ]
        [ -n | --no-strip-underscores ]
        [ -sformat | --format=format ]
        [ --help ] [ --version ] [ symbol...]
```

The C++ language provides function overloading, which means that you can write many functions with the same name (providing each takes parameters of different types). All C++ function names are encoded into a low-level assembly label (this process is known as *mangling*). The `c++filt` program does the inverse mapping: it decodes (*demangles*) low-level names into user-level names so that the linker can keep these overloaded functions from clashing.

Every alphanumeric word (consisting of letters, digits, underscores, dollars, or periods) seen in the input is a potential label. If the label decodes into a C++ name, the C++ name replaces the low-level name in the output.

You can use `c++filt` to decipher individual symbols, using a declaration like the following example.

```
c++filt symbol
```

If no *symbol* arguments are given, `c++filt` reads symbol names from the standard input and writes the demangled names to the standard output. All results are printed on the standard output.

-_

--strip-underscores

On some systems, both the C and C++ compilers put an underscore in front of every name. This option removes the initial underscore. Whether `c++filt` removes the underscore by default is target dependent.

-n

--no-strip-underscores

Do not remove the initial underscore.

-s *format*

--format=*format*

GNU `nm` can decode three different methods of mangling, used by different C++ compilers. The argument to this option selects which method it uses:

gnu

The one used by the GNU compiler (the default method).

lucid

The one used by the Lucid compiler.

arm

The one specified by the *C++ Annotated Reference Manual*.

--help

Print a summary of the options to c++filt and exit.

--version

Print the version number of c++filt and exit.

WARNING: c++filt is a developing utility, meaning that the details of its user interface are subject to change as C++ changes. In particular, a command-line option may be required in the future to decode a name passed as an argument on the command line; for example, c++filt *symbol* may in a future release become c++filt *option symbol*.

addr2line utility

```
addr2line [ -b bfdname | --target=bfdname ]
          [ -C | --demangle ]
          [ -e filename | --exe=filename ]
          [ -f | --functions ] [ -s | --basename ]
          [ -H | --help ] [ -V | --version ]
          [ addr addr ... ]
```

`addr2line` translates program addresses into file names and line numbers. Given an address and an executable, it uses the debugging information in the executable to figure out which file name and line number are associated with a given address.

The executable to use is specified with the `-e` option. The default is `a.out`.

`addr2line` has two modes of operation.

In the first, hexadecimal addresses are specified on the command line, and `addr2line` displays the file name and line number for each address.

In the second, `addr2line` reads hexadecimal addresses from standard input, and prints the file name and line number for each address on standard output. In this mode, `addr2line` may be used in a pipe to convert dynamically chosen addresses.

The format of the output is `FILENAME:LINENO`. The file name and line number for each address is printed on a separate line. If the `-f` option is used, then each `FILENAME:LINENO` line is preceded by a `FUNCTIONNAME` line which is the name of the function containing the address.

If the file name or function name can not be determined, `addr2line` will print two question marks in their place. If the line number can not be determined, `addr2line` will print 0.

The long and short forms of options, shown here as alternatives, are equivalent.

`-b bfdname`

`--target=bfdname`

Specify that the object-code format for the object files is *bfdname*.

`-C`

`--demangle`

Decode (*demangle*) low-level symbol names into user-level names. Besides removing any initial underscore prepended by the system, this makes C++ function names readable. See “c++filt utility” on page 283 for more information on demangling.

`-e filename`

`--exe=filename`

Specify the name of the executable for which addresses should be translated. The default file is `a.out`.

addr2line utility

- f
- functions
Display function names as well as file and line number information.
- s
- basenames
Display only the base of each file name.

nlmconv utility

nlmconv converts a relocatable object file into a NetWare Loadable Module.

warning:

nlmconv is not always built as part of the binary utilities, since it is only useful for NLM targets.

```
nlmconv [ -I bfdname | --input-target=bfdname ]
        [ -O bfdname | --output-target=bfdname ]
        [ -T headerfile | --header-file=headerfile ]
        [ -d | --debug ] [ -l linker | --linker=linker ]
        [ -h | --help ] [ -V | --version ]
infile outfile
```

nlmconv converts the relocatable ‘i386’ object file *infile* into the NetWare Loadable Module, *outfile*, optionally reading *headerfile* for NLM header information.

For instructions on writing the NLM command file language used in header files, see “linkers” or “NLMLINK” in particular, in the *NLM Development and Tools Overview*, which is part of the *NLM Software Developer’s Kit* (“*NLM SDK*”), available from Novell, Inc.

nlmconv uses the GNU Binary File Descriptor library (BFD) to read *infile*; see the documentation for “BFD libraries” on page 243 in *Using LD* for more information.

nlmconv can perform a link step. In other words, you can list more than one object file for input if you list them in the definitions file (rather than simply specifying one input file on the command line). In this case, nlmconv calls the linker for you.

-I *bfdname*

--input-target=*bfdname*

Object format of the input file. nlmconv can usually determine the format of a given file (so no default is necessary). See “Target selection” on page 294 for more information.

-O *bfdname*

--output-target=*bfdname*

Object format of the output file. nlmconv infers the output format based on the input format, e.g., for a ‘i386’ input file, the output format is ‘nlm32-i386’. See “Target selection” on page 294 for more information.

-T *headerfile*

--header-file=*headerfile*

Reads *headerfile* for NLM header information. For instructions on writing the NLM command file language used in header files, see “linkers” or “NLMLINK” in particular, of the *NLM Development and Tools Overview*, which is part of the *NLM Software Developer’s Kit* (“*NLM SDK*”), available from Novell, Inc.

-d
--debug
 Displays (on standard error) the linker command line used by nlmconv.
-l *linker*
--linker=*linker*
 Use *linker* for any linking. *linker* can be an absolute or a relative pathname.
-h
--help
 Prints a usage summary.
-V
--version
 Prints the version number for nlmconv.

windres utility

windres may be used to manipulate Windows resources.

WARNING: windres is not always built as part of the binary utilities, since it is only useful for Windows targets.

```
windres [options] [input-file] [output-file]
```

windres reads resources from an input file and copies them into an output file. Either file may be in one of three formats:

`rc`

A text format read by the Resource Compiler.

`res`

A binary format generated by the Resource Compiler.

`coff`

A COFF object or executable.

The exact description of these different formats is available in documentation from Microsoft.

When windres converts from the `rc` format to the `res` format, it is acting like the Windows Resource Compiler. When windres converts from the `res` format to the `coff` format, it is acting like the Windows `CVTRES` program.

When windres generates an `rc` file, the output is similar but not identical to the format expected for the input. When an input `rc` file refers to an external filename, an output `rc` file will instead include the file contents.

If the input or output format is not specified, windres will guess based on the file name, or, for the input file, the file contents. A file with an extension of `‘.rc’` will be treated as an `rc` file, a file with an extension of `‘.res’` will be treated as a `res` file, and a file with an extension of `‘.o’` or `‘.exe’` will be treated as a `coff` format file.

If no output file is specified, windres will print the resources in `rc` format to standard output.

The normal use is for you to write an `rc` file, use windres to convert it to a COFF file, and then link the COFF file into your application.

This will make the resources described in the `rc` file available to Windows.

```
-i filename
```

```
--input filename
```

The name of the input file. If this option is not used, then windres will use the first non-option argument as the input file name. If there are no non-option arguments, then windres will read from standard input. windres can not read a COFF file from standard input.

`-o filename`
`--output filename`
The name of the output file. If this option is not used, then `windres` will use the first non-option argument, after any used for the input file name, as the output file name. If there is no non-option argument, then `windres` will write to standard output. `windres` can not write a COFF file to standard output.

`-I format`
`--input-format format`
The input format to read. *format* may be `res`, `rc`, or `coff`. If no input format is specified, `windres` will guess.

`-O format`
`--output-format format`
The output format to generate. *format* may be `res`, `rc`, or `coff`. If no output format is specified, `windres` will guess.

`-F target`
`--target target`
Specify the BFD format to use for a COFF file as input or output. This is a BFD target name; you can use the `--help` option to see a list of supported targets. Normally `windres` will use the default format, which is the first one listed by the `--help` option.

`--preprocessor program`
When `windres` reads an `rc` file, it runs it through the C preprocessor first. This option may be used to specify the preprocessor to use, including any leading arguments. The default preprocessor argument uses the following commands and arguments:

```
gcc -E -xc-header DRC_INVOKED.
```

`--include-dir directory`
Specify an include directory to use when reading an `rc` file.
`windres` will pass this to the preprocessor as an `-I` option.
`windres` will also search this directory when looking for files named in the `rc` file.

`--define sym[=val]`
Specify a `-D` option to pass to the preprocessor when reading an `rc` file.

`--language val`
Specify the default language to use when reading an `rc` file.
val should be a hexadecimal language code. The low eight bits are the language, and the high eight bits are the sub-language.

`--help`
Prints a usage summary.

`--version`

Prints the version number for windres.

`--yydebug`

If windres is compiled with `YYDEBUG` defined as 1, this will turn on parser debugging.

windres utility

2

Selecting the target system

You can specify three aspects of the target system to the GNU binary file utilities, selecting each in several ways:

- as **target**; see “Target selection” on page 294 for discussions of the lists of ways to specify values in order of decreasing precedence for each specification
- as **architecture**; see “Architecture selection” on page 296 for the ways to manage the binary utilities on different processors
- as **linker emulation**., a *personality* of the linker, giving the linker default values applying only to the linker; see “Linker emulation selection” on page 297

See “Overview of binutils, the GNU binary utilities” on page 255 for more documentation on the individual utilities.

Target selection

A *target* is an object file format. A given target may be supported for multiple architectures (see “Architecture selection” on page 296). A target selection may also have variations for different operating systems or architectures.

The commands to list valid values only list the values for which the programs you are running were configured. If they were configured with `--enable-targets=all`, the commands list most of the available values, but a few are left out; not all targets can be configured in at once because some of them can only be configured *native* (on hosts with the same type as the target system).

The command to list valid target values is `objdump -i` (the first column of output contains the relevant information).

Some sample values are: `'a.out-hp300bsd'`, `'ecoff-littlemips'`, `'a.out-sunos-big'`.

You can also specify a target using a *configuration triplet*. This is the same sort of name that is passed to `configure` to specify a target.

When you use a configuration triplet as an argument, it must be fully *canonicalized*. You can see the canonical version of a triplet by running the shell script, `config.sub`, which is included with the sources.

Some sample configuration triplets are `m68k-hp-bsd`, `mips-dec-ultrix`, and `sparc-sun-sunos`.

objdump target

Ways to specify:

- command line option: `'-b'` or `'--target'`
- environment variable `GNUTARGET`
- deduced from the input file

objcopy and strip input target

Ways to specify:

- command line options, `-I`, `--input-target`, `-F`, or `--target`
- environment variable, `GNUTARGET`
- deduced from the input file

objcopy and strip Output target

Ways to specify:

- command line options, `-O`, `--output-target`, `-F`, or `--target`
- the input target (see “objcopy and strip input target” on page 294)
- environment variable, `GNUTARGET`
- deduced from the input file

nm, size, and strings target

Ways to specify:

- command line option, `--target`
- environment variable `GNUTARGET`
- deduced from the input file

Linker input target

Ways to specify:

- command line option, `-b` or `--format` (see “ld command line options” on page 190 in *Using LD*)
- script command, `TARGET` (see “Commands dealing with object file formats” on page 208 in *Using LD*)
- environment variable, `GNUTARGET` (see “ld environment variables” on page 202 in *Using LD*)
- the default target of the selected linker emulation (see “Linker emulation selection” on page 297).

Linker output target

Ways to specify:

- command line option, `-oformat` (see “ld command line options” on page 190 in *Using LD*)
- script command, `OUTPUT_FORMAT` (see “Commands dealing with object file formats” on page 208 in *Using LD*)
- the linker input target (see “Linker input target”)

Architecture selection

An *architecture* is a type of CPU on which an object file is to run. Its name may contain a colon, separating the name of the processor family from the name of the particular processor.

The command to list valid architecture values is ‘`objdump -i`’ (the second column contains the relevant information).

Sample values: `m68k:68020`, `mips:3000`, `sparc`.

objdump architecture

Ways to specify: `objdump`

- command line option: ‘`-m`’ or ‘`--architecture`’
- deduced from the input file

objcopy, nm, size, strings architecture

Its specification is deduced from the input file.

Linker input architecture

Its specification is deduced from the input file.

Linker output architecture

Ways to specify:

- script command, `OUTPUT_ARCH` (see “Other linker script commands” on page 209 in *Using LD*)
- the default architecture from the linker output target (see “Target selection” on page 294)

Linker emulation selection

A *linker emulation* is a personality of the linker, giving the linker default values for the other aspects of the target system.

In particular, it consists of the linker script, the target and several *hook* functions (which are run at certain stages of the linking process to do special things that some targets require).

The command to list valid linker emulation values is `ld -v`.

Sample values: 'hp300bsd', 'mipsplit', 'sun4'.

Ways to specify:

- command line option, `-m` (see “ld command line options” on page 190 in *Using LD*)
- environment variable, `LDEMULATION`
- compiled-in `DEFAULT_EMULATION` from `Makefile` which comes from `EMUL` in `'config/target.mt'`

Using make

Copyright © 1991-1998 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions, except that the documentation entitled “GNU General Public License,” “Funding for Free Software,” and “Protect Your Freedom; Fight ‘Look And Feel’” and this permission notice, may be included in translations approved by the Free Software Foundation instead of in the original English.

Free Software Foundation

59 Temple Place / Suite 330
Boston, MA 02111-1307 USA

ISBN: 1-882114-66-3

This documentation has been prepared by Cygnus.

Copyright © 1992-1998 Cygnus

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

This documentation has been prepared by Cygnus Technical Publications; contact the Cygnus Technical Publications staff: doc@cygnus.com.

1

Overview of `make`, a program for recompiling

The `make` utility automatically determines which pieces of a large program need to be recompiled, and then issues commands to recompile them.

The following documentation summarizes the `make` utility.

- “Summary of make options” on page 417
- “GNU make quick reference” on page 477

In the following discussions, the first few paragraphs contain introductory or general information while the subsequent paragraphs contain specialized or technical information. The exception is “Introduction to makefiles” on page 303, all of which is overview. If you are familiar with other `make` programs, see “Summary of the features for the GNU make utility” on page 455, “Special built-in target names” on page 338 and “GNU make’s incompatibilities and missing features” on page 459 (which explains the few things GNU `make` lacks that other programs provide).

- “Introduction to makefiles” on page 303
- “Writing makefiles” on page 313
- “Writing rules” on page 323
- “Writing the commands in rules” on page 349
- “How to use variables” on page 367
- “Conditional parts of makefiles” on page 385
- “Functions for transforming text” on page 393

- “How to run the make utility” on page 407
- “Summary of make options” on page 417
- “Implicit rules” on page 423
- “Using make to update archive files” on page 449
- “Summary of the features for the GNU make utility” on page 455
- “GNU make’s incompatibilities and missing features” on page 459
- “Makefile conventions” on page 463
- “GNU make quick reference” on page 477
- “Complex makefile example” on page 483

2

Introduction to makefiles

`make` is a program that was implemented by Richard Stallman and Roland McGrath. GNU `make` conforms to section 6.2 of *IEEE Standard 1003.2-1992* (POSIX.2). The following documentation discusses the fundamentals of `make`.

- “What a rule looks like” on page 305
- “A simple makefile” on page 306
- “How `make` processes a makefile” on page 308
- “Variables make makefiles simpler” on page 309
- “Letting `make` deduce the commands” on page 310
- “Another style of makefile” on page 311
- “Rules for cleaning the directory” on page 312

The examples in the documentation for `make` show C programs, since they are most common, although you can use `make` with any programming language whose compiler can be run with a shell command. Indeed, `make` is not limited to programs. You can use it to describe any task where some files must be updated automatically from others whenever the others change.

To prepare to use `make`, you must write a file called the *makefile*, which describes the relationships among files in your program and provides commands for updating each file. In a program, typically, the executable file is updated from object files, which are in turn made by compiling source files.

Once a suitable makefile exists, each time you change some source files, the following shell command suffices to perform all necessary recompilations.

`make`

The `make` program uses the makefile data base and the last-modification times of the files to decide which of the files need to be updated. For each of those files, it issues the commands recorded in the data base.

You can provide command line arguments to `make` to control how and which files should be recompiled.

If you are new to `make`, or are looking for a general introduction, read the following discussions.

You need a file called a *makefile* to tell `make` what to do. Most often, the makefile tells `make` how to compile and link a program.

In the following discussions, we will describe a simple makefile that tells how to compile and link a text editor which consists of eight C source files and three header files. The makefile can also tell `make` how to run miscellaneous commands when explicitly asked (for example, to remove certain files as a clean-up operation). To see a more complex example of a makefile, see “Complex makefile example” on page 483.

When `make` recompiles the editor, each changed C source file must be recompiled. If a header file has changed, each C source file that includes the header file must be recompiled to be safe. Each compilation produces an object file corresponding to the source file. Finally, if any source file has been recompiled, all the object files, whether newly made or saved from previous compilations, must be linked together to produce the new executable editor.

What a rule looks like

A simple makefile consists of *rules* with the following shape:

```
target ... : dependencies ...  
    command  
    ...  
    ...
```

A *target* is usually the name of a file that is generated by a program; examples of targets are executable or object files. A target can also be the name of an action to carry out, such as `clean` (see “Phony targets” on page 334).

A *dependency* is a file that is used as input to create the target. A target often depends on several files.

A *command* is an action that `make` carries out. A rule may have more than one command, each on its own line.

NOTE: You need to put a tab character at the beginning of every command line! This is an obscurity that catches the unwary.

Usually a command is in a rule with dependencies and serves to create a target file if any of the dependencies change. However, the rule that specifies commands for the target need not have dependencies. For example, the rule containing the delete command associated with the target, `clean`, does not have dependencies.

A *rule*, then, explains how and when to remake certain files which are the targets of the particular rule. `make` carries out the commands on the dependencies to create or update the target. A rule can also explain how and when to carry out an action. See “Writing rules” on page 323.

A makefile may contain other text besides rules, but a simple makefile need only contain rules. Rules may look somewhat more complicated than shown in this template, but all fit the pattern more or less.

A simple makefile

What follows is a straightforward makefile that describes the way an executable file called `edit` depends on eight object files which, in turn, depend on eight C source and three header files. In the following example, all the C files include `defs.h`, but only those defining editing commands include `command.h`, and only low level files that change the editor buffer include `buffer.h`.

```
edit : main.o kbd.o command.o display.o \  
      insert.o search.o files.o utils.o  
      cc -o edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
  
clean :  
      rm edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o
```

NOTE: We split each long line into two lines using backslash-newline (`\`); this is like using one long line, but is easier to read.

To use the previous sample makefile to create the executable file called ‘`edit`’, type: `make`.

To use this makefile to delete the executable file and all the object files from the directory, type: `make clean`.

See the sample makefile with “A simple makefile” on page 306, the targets include the executable file, `edit`, and the object files, `main.o` and `kbd.o`. The dependencies are files such as `main.c` and `defs.h`. In fact, each `.o` file is both a target and a dependency. Commands include `cc -c main.c` and `cc -c kbd.c`.

When a target is a file, it needs to be recompiled or relinked if any of its dependencies change. In addition, any dependencies that are themselves automatically generated should first be updated. In the previous example, `edit` depends on each of the eight object files; the object file, `main.o`, depends on the source file, `main.c`, and on the header file, `defs.h`.

A shell command follows each line that contains a target and dependencies. These shell commands say how to update the target file. A tab character must come at the beginning of every command line to distinguish commands lines from other lines in the makefile. (Bear in mind that `make` does not know anything about how the commands work. It is up to you to supply commands that will update the target file properly. All `make` does is execute the commands in the rule you have specified when the target file needs to be updated.)

The target `clean` is not a file, but merely the name of an action. Since you normally do not want to carry out the actions in this rule, `clean` is not a dependency of any other rule. Consequently, `make` never does anything with it unless you tell it specifically. *This rule not only is not a dependency, it also does not have any dependencies, so the only purpose of the rule is to run the specified commands.* Targets that do not refer to files but are just actions are called *phony targets*. See “Phony targets” on page 334 for information about this kind of target. See “Errors in commands” on page 355 to see how to cause `make` to ignore errors from `rm` or any other command.

How `make` processes a makefile

By default, `make` starts with the first rule (not counting rules whose target names start with `.'`). This is called the default goal. (Goals are the targets that `make` strives ultimately to update. See “Arguments to specify the goals” on page 409.)

See the example shown with “A simple makefile” on page 306; its default goal is to update the executable program `edit`; therefore, we put that rule first.

When you give the command, `make`, `make` reads the makefile in the current directory and begins by processing the first rule. In the example, this rule is for relinking `edit`; but before `make` can fully process this rule, it must process the rules for the files that `edit` depends on; in this case, they are the object files. Each of these files is processed according to its own rule. These rules say to update each `.'o'` file by compiling its source file. The recompilation must be done if the source file, or any of the header files named as dependencies, is more recent than the object file, or if the object file does not exist.

The other rules are processed because their targets appear as dependencies of the goal. If some other rule is not depended on by the goal (or anything it depends on, etc.), that rule is not processed, unless you tell `make` to do so (with a command such as `make clean`).

Before recompiling an object file, `make` considers updating its dependencies, the source file and header files. This makefile does not specify anything to be done for them—the `.'c'` and `.'h'` files are not the targets of any rules—so `make` does nothing for these files. But `make` would update automatically generated C programs, such as those made by Bison or Yacc, by their own rules.

After recompiling whichever object files need it, `make` decides whether to relink `edit`. This must be done if the file `edit` does not exist, or if any of the object files are newer than it. If an object file was just recompiled, it is now newer than `edit`, so `edit` is relinked.

Thus, if we change the file `insert.c` and run `make`, `make` will compile that file to update `insert.o`, and then link `edit`. If we change the file `command.h` and run `make`, `make` will recompile the object files `kbd.o` along with `command.o` and `files.o` and then link the file `edit`.

Variables make makefiles simpler

See the first example with “A simple makefile” on page 306; we had to list all the object files twice in the rule for ‘edit’ and we’ve repeated in this next example.

```
edit : main.o kbd.o command.o display.o \
      insert.o search.o files.o utils.o
      cc -o edit main.o kbd.o command.o display.o \
      insert.o search.o files.o utils.o
```

Such duplication is error-prone; if a new object file is added to the system, we might add it to one list and forget the other. We can eliminate the risk and simplify the makefile by using a variable. *Variables* allow a text string to be defined once and substituted in multiple places later (see “How to use variables” on page 367).

It is standard practice for every makefile to have a variable named `objects`, `OBJECTS`, `objs`, `OBJS`, `obj`, or `OBJ` which is a list of all object file names. We would define such a variable, `objects`, with input like the following in the makefile.

```
objects = main.o kbd.o command.o display.o \
          insert.o search.o files.o utils.o
```

Then, each place we want to put a list of the object file names, we can substitute the variable’s value by writing ‘\$(objects)’. The following example shows how the complete simple makefile looks when you use a variable for the object files.

```
objects = main.o kbd.o command.o display.o \
          insert.o search.o files.o utils.o

edit : $(objects)
      cc -o edit $(objects)
main.o : main.c defs.h
      cc -c main.c
kbd.o : kbd.c defs.h command.h
      cc -c kbd.c
command.o : command.c defs.h command.h
      cc -c command.c
display.o : display.c defs.h buffer.h
      cc -c display.c
insert.o : insert.c defs.h buffer.h
      cc -c insert.c
search.o : search.c defs.h buffer.h
      cc -c search.c
files.o : files.c defs.h buffer.h command.h
      cc -c files.c
utils.o : utils.c defs.h
      cc -c utils.c
clean :
      rm edit $(objects)
```

Letting `make` deduce the commands

It is not necessary to spell out the commands for compiling the individual C source files, because `make` can figure them out: it has an implicit rule for updating a `‘.o’` file from a correspondingly named `‘.c’` file using a `‘cc -c’` command.

For example, it will use the command `‘cc -c main.c -o main.o’` to compile `‘main.c’` into `‘main.o’`. We can therefore omit the commands from the rules for the object files. See “Implicit rules” on page 423.

When a `‘.c’` file is used automatically in this way, it is also automatically added to the list of dependencies. We can therefore omit the `‘.c’` files from the dependencies, provided we omit the commands. The following is the entire example, with both of these changes, and a variable, `objects` (as previously suggested with “Variables make makefiles simpler” on page 309).

```
objects = main.o kbd.o command.o display.o \
          insert.o search.o files.o utils.o

edit : $(objects)
        cc -o edit $(objects)
main.o : defs.h
kbd.o : defs.h command.h
command.o : defs.h command.h
display.o : defs.h buffer.h
insert.o : defs.h buffer.h
search.o : defs.h buffer.h
files.o : defs.h buffer.h command.h
utils.o : defs.h

.PHONY : clean
clean :
        -rm edit $(objects)
```

The example in “Another style of makefile” on page 311 is how we would write the makefile in actual practice. (The complications associated with `‘clean’` are described in “Phony targets” on page 334 and in “Errors in commands” on page 355.)

Because implicit rules are so convenient, they are used frequently.

Another style of makefile

When the objects of a makefile are created only by implicit rules, an alternative style of makefile is possible. In this style of makefile, you group entries by their dependencies instead of by their targets.

The following is what one resembles. ‘defs.h’ is given as a dependency of all the object files; ‘command.h’ and ‘buffer.h’ are dependencies of the specific object files listed for them.

```
objects = main.o kbd.o command.o display.o \  
insert.o search.o files.o utils.o  
  
edit : $(objects)  
cc -o edit $(objects)  
  
$(objects) : defs.h  
kbd.o command.o files.o : command.h  
display.o insert.o search.o files.o : buffer.h
```

Whether this is better is a matter of taste: it is more compact, but some people dislike it because they find it clearer to put all the information about each target in one place.

Rules for cleaning the directory

Compiling a program is not the only thing for which you might want to write rules. Makefiles commonly tell how to do a few other things besides compiling a program; for instance, how to delete all the object files and executables so that the directory is ‘clean’. The following shows how to write a `make` rule for cleaning the example editor.

```
clean:
    rm edit $(objects)
```

In practice, you might want to write the rule in a somewhat more complicated manner to handle unanticipated situations. Use input like the following example.

```
.PHONY : clean
clean :
    -rm edit $(objects)
```

This prevents `make` from using an actual file called ‘`clean`’ allowing it to continue in spite of errors from `rm`. See “Phony targets” on page 334 and in “Errors in commands” on page 355. A rule such as this should not be placed at the beginning of the makefile, since you do not want it to run by default! Thus, in the example makefile, you want the rule for `edit` which recompiles the editor, to remain the default goal. Since `clean` is not a dependency of `edit`, this rule will not run at all if we give the command ‘`make`’ with no arguments. In order to make the rule run, use ‘`make clean`’. See “How to run the make utility” on page 407.

3

Writing makefiles

The information that tells **make** how to recompile a system comes from reading a data base called the *makefile*. The following documentation discusses writing makefiles.

- “What makefiles contain” on page 314
- “What name to give your makefile” on page 315
- “Including other makefiles” on page 316
- “The MAKEFILES variable” on page 318
- “How makefiles are remade” on page 319
- “Overriding part of another makefile” on page 321

What makefiles contain

Makefiles contain five kinds of things: *explicit rules*, *implicit rules*, *variable definitions*, *directives*, and *comments*. Rules, variables, and directives are described in better detail in the corresponding references noted in the following descriptions..

- An *explicit rule* says when and how to remake one or more files called the rule's *targets*. It lists the other files on which the targets depend, and may also give commands to use to create or update the targets. See “Writing rules” on page 323.
- An *implicit rule* says when and how to remake a class of files based on their names. It describes how a target may depend on a file with a name similar to the target and gives commands to create or update such a target. See “Implicit rules” on page 423.
- A *variable definition* is a line that specifies a text string value for a variable that can be substituted into the text later. The simple makefile example shows a variable definition for objects as a list of all object files (see “Variables make makefiles simpler” on page 309).
- A *directive* is a command for make to do something special while reading the makefile. These include:
 - Reading another makefile (see “Including other makefiles” on page 316).
 - Deciding (based on the values of variables) whether to use or ignore a part of the makefile (see “Conditional parts of makefiles” on page 385).
 - Defining a variable from a verbatim string containing multiple lines (see “Defining variables verbatim” on page 382).
- A *comment* in a line of a makefile starts with ‘#’. It and the rest of the line are ignored, except that a trailing backslash not escaped by another backslash will continue the comment across multiple lines. Comments may appear on any of the lines in the makefile, except within a `define` directive, and perhaps within commands (where the shell decides what is a comment). A line containing just a comment (with perhaps spaces before it) is effectively blank, and is ignored.

What name to give your makefile

By default, when **make** looks for the makefile, it tries the following names, in order: ‘GNUmakefile’, ‘makefile’ and ‘Makefile’.

Normally you should call your makefile either ‘makefile’ or ‘Makefile’. (We recommend ‘Makefile’ because it appears prominently near the beginning of a directory listing, right near other important files such as ‘README’.) The first name checked, ‘GNUmakefile’, is not recommended for most makefiles. You should use this name if you have a makefile that is specific to GNU **make**, and will not be understood by other versions of make. Other **make** programs look for ‘makefile’ and ‘Makefile’, but not ‘GNUmakefile’.

If **make** finds none of these names, it does not use any makefile. Then you must specify a goal with a command argument, and **make** will attempt to figure out how to remake it using only its built-in implicit rules. See “Using implicit rules” on page 425.

If you want to use a non-standard name for your makefile, you can specify the makefile name with the ‘-f’ or ‘--file’ option.

The ‘-f *name*’ or ‘--file=*name*’ arguments tell **make** to read the file, *name*, as the makefile. If you use more than one ‘-f’ or ‘--file’ option, you can specify several makefiles.

All the makefiles are effectively concatenated in the order specified.

The default makefile names, ‘GNUmakefile’, ‘makefile’ and ‘Makefile’, are not checked automatically if you specify ‘-f’ or ‘--file’.

Including other makefiles

The `include` directive tells **make** to suspend reading the current make-file and read one or more other makefiles before continuing. The directive is a line in the makefile that looks like `include filenames...`

filenames can contain shell file name patterns. Extra spaces are allowed and ignored at the beginning of the line, but a tab is not allowed. (If the line begins with a tab, it will be considered a command line.) Whitespace is required between `include` and the file names, and between file names; extra whitespace is ignored there and at the end of the directive. A comment starting with `#` is allowed at the end of the line. If the file names contain any variable or function references, they are expanded. See “How to use variables” on page 367.

For example, if you have three `.mk` files, `a.mk`, `b.mk`, and `c.mk`, and `$(bar)` expands to `bish bash`, then the expression, `include foo *.mk $(bar)`, is equivalent to `include foo a.mk b.mk c.mk bish bash`.

When **make** processes an `include` directive, it suspends reading of the containing makefile and reads from each listed file in turn. When that is finished, **make** resumes reading the makefile in which the directive appears. One occasion for using `include` directives is when several programs, handled by individual makefiles in various directories, need to use a common set of variable definitions (see “Setting variables” on page 378) or pattern rules (see “Defining and redefining pattern rules” on page 436).

Another such occasion is when you want to generate dependencies from source files automatically; the dependencies can be put in a file that is included by the main makefile. This practice is generally cleaner than that of somehow appending the dependencies to the end of the main makefile as has been traditionally done with other versions of **make**. See “Generating dependencies automatically” on page 346.

If the specified *name* does not start with a slash, and the file is not found in the current directory, several other directories are searched. First, any directories you have specified with the `-I` or the `--include-dir` options are searched (see “Summary of make options” on page 417). Then the directories (if they exist) are searched, in the following order.

1. `'prefix/include'` (normally, `'/usr/local/include'`)
2. `'/usr/gnu/include'`,
3. `'/usr/local/include'`, `'/usr/include'`.

If an included makefile cannot be found in any of these directories, a warning message is generated, but it is not an immediately fatal error; processing of the makefile containing the `include` continues. Once it has finished reading makefiles, `make` will try to remake any that are out of date or don't exist. See “How makefiles are remade” on page 319. Only after it has tried to find a way to remake a makefile and failed, will `make` diagnose the missing makefile as a fatal error.

If you want `make` to simply ignore a makefile which does not exist and cannot be remade, with no error message, use the `-include` directive instead of `include`, as in `-include filenames....` This acts like `include` in every way except that there is no error (not even a warning) if any of the *filenames* do not exist.

The MAKEFILES variable

If the environment variable, `MAKEFILES`, is defined, **make** considers its value as a list of names (separated by whitespace) of additional makefiles to be read before the others. This works much like the `include` directive in that various directories are searched for those files (see “Including other makefiles” on page 316). In addition, the default goal is never taken from one of these makefiles and it is not an error if the files listed in `MAKEFILES` are not found.

The main use of `MAKEFILES` is in communication between recursive invocations of **make** (see “Recursive use of the make utility” on page 358). It usually is not desirable to set the environment variable before a top-level invocation of **make**, because it is usually better not to mess with a makefile from outside. However, if you are running **make** without a specific makefile, a makefile in `MAKEFILES` can do useful things to help the built-in implicit rules work better, such as defining search paths (see “Directory search and implicit rules” on page 332).

Some users are tempted to set `MAKEFILES` in the environment automatically on login, and program makefiles to expect this to be done. This is a very bad idea, because such makefiles will fail to work if run by anyone else. It is much better to write explicit `include` directives in the makefiles. See “Including other makefiles” on page 316.

How makefiles are remade

Sometimes makefiles can be remade from other files, such as RCS or SCCS files. If a makefile can be remade from other files, you probably want **make** to get an up-to-date version of the makefile to read in.

To this end, after reading in all makefiles, **make** will consider each as a goal target and attempt to update it. If a makefile has a rule which says how to update it (found either in that very makefile or in another one) or if an implicit rule applies to it (see “Using implicit rules” on page 425, it will be updated if necessary. After all makefiles have been checked, if any have actually been changed, **make** starts with a clean slate and reads all the makefiles over again. (It will also attempt to update each of them over again, but normally this will not change them again, since they are already up to date.)

If the makefiles specify a double-colon rule to remake a file with commands but no dependencies, that file will always be remade (see “Double-colon rules” on page 345). In the case of makefiles, a make-file that has a double-colon rule with commands but no dependencies will be remade every time **make** is run, and then again after **make** starts over and reads the makefiles in again. This would cause an infinite loop; **make** would constantly remake the makefile, and never do anything else. So, to avoid this, **make** will not attempt to remake makefiles which are specified as double-colon targets but have no dependencies.

If you do not specify any makefiles to be read with ‘-f’ or ‘--file’ options, **make** will try the default makefile names; see “What name to give your makefile” on page 315. Unlike makefiles explicitly requested with ‘-f’ or ‘--file’ options, **make** is not certain that these makefiles should exist. However, if a default makefile does not exist but can be created by running **make** rules, you probably want the rules to be run so that the makefile can be used.

Therefore, if none of the default makefiles exists, **make** will try to make each of them in the same order in which they are searched for (see “What name to give your makefile” on page 315) until it succeeds in making one, or it runs out of names to try. Note that it is not an error if **make** cannot find or make any makefile; a makefile is not always necessary.

When you use the ‘-t’ or ‘--touch’ option (see “Instead of executing the commands” on page 411), you would not want to use an out-of-date makefile to decide which targets to touch. So the ‘-t’ option has no effect on updating makefiles; they are really updated even if ‘-t’ is specified. Likewise, ‘-q’ (or ‘--question’) and ‘-n’ (or ‘--just-print’) do not prevent updating of makefiles, because an out-of-date makefile would result in the wrong output for other targets.

However, on occasion you might actually wish to prevent updating of even the makefiles. You can do this by specifying the makefiles as goals in the command line as well as specifying them as makefiles. When the makefile name is specified explicitly as a goal, the options ‘-t’ and so on do apply to them.

Thus, ‘make -f mfile -n foo’ will update ‘mfile’, read it in, and then print the commands to update ‘foo’ and its dependencies without running them. The commands printed for ‘foo’ will be those specified in the updated contents of ‘mfile’.

Overriding part of another makefile

Sometimes it is useful to have a makefile that is mostly just like another makefile. You can often use the `'include'` directive to include one in the other, and add more targets or variable definitions. However, if the two makefiles give different commands for the same target, `make` will not let you just do this. But there is another way.

In the containing makefile (the one that wants to include the other), you can use a match-anything pattern rule to say that to remake any target that cannot be made from the information in the containing makefile, `make` should look in another makefile. See “Defining and redefining pattern rules” on page 436 for more information on pattern rules.

For example, if you have a makefile called `'Makefile'` that says how to make the target `'foo'` (and other targets), you can write a makefile called `'GNUmakefile'` that contains the following.

```
foo:
    frobnicate > foo

%: force
    @$(MAKE) -f Makefile $@

force: ;
```

If you say `'make foo'`, `make` will find `'GNUmakefile'`, read it, and see that to make `'foo'`, it needs to run the command `'frobnicate > foo'`. If you say `'make bar'`, `make` will find no way to make `'bar'` in `'GNUmakefile'`, so it will use the commands from the pattern rule: `'make -f Makefile bar'`. If `'Makefile'` provides a rule for updating `'bar'`, `make` will apply the rule. And likewise for any other target that `'GNUmakefile'` does not say how to make.

The way this works is that the pattern rule has a pattern of just `'%'`, so it matches any target whatever. The rule specifies a dependency `'force'`, to guarantee that the commands will be run even if the target file already exists. We give `'force'` target empty commands to prevent `make` from searching for an implicit rule to build it—otherwise it would apply the same match-anything rule to `'force'` itself and create a dependency loop!

4

Writing rules

A *rule* appears in the makefile and says when and how to remake certain files, called the rule's *targets* (most often only one per rule). It lists the other files that are the *dependencies* of the target, and *commands* to use to create or update the target. The following documentation discusses the general rules for makefiles.

- “Rule syntax” on page 325
- “Using wildcard characters in file names” on page 326
- “Wildcard examples” on page 327
- “Pitfalls of using wildcards” on page 328
- “The wildcard function” on page 329
- “Searching directories for dependencies” on page 330
- “Phony targets” on page 334
- “Rules without commands or dependencies” on page 336
- “Empty target files to record events” on page 337
- “Special built-in target names” on page 338
- “Multiple targets in a rule” on page 340
- “Multiple rules for one target” on page 341
- “Static pattern rules” on page 342
- “Double-colon rules” on page 345

■ “Generating dependencies automatically” on page 346

The order of rules is not significant, except for determining the *default goal*: the target for `make` to consider, if you do not otherwise specify one. The default goal is the target of the first rule in the first makefile. If the first rule has multiple targets, only the first target is taken as the default. There are two exceptions: a target starting with a period is not a default unless it contains one or more slashes, ‘/’, as well; and, a target that defines a pattern rule has no effect on the default goal. See “Defining and redefining pattern rules” on page 436.

Therefore, we usually write the makefile so that the first rule is the one for compiling the entire program or all the programs described by the makefile (often with a target called ‘`all`’). See “Arguments to specify the goals” on page 409.

Rule syntax

In general, a rule looks like the following.

```
targets : dependencies
command
...
```

Or like the following.

```
targets : dependencies ; command
command
...
```

The *targets* are file names, separated by spaces. Wildcard characters may be used (see “Using wildcard characters in file names” on page 326) and a name of the form ‘*a(m)*’ represents member *m* in archive file *a* (see “Archive members as targets” on page 450). Usually there is only one target per rule, but occasionally there is a reason to have more (see “Multiple targets in a rule” on page 340). The *command* lines start with a tab character. The first command may appear on the line after the dependencies, with a tab character, or may appear on the same line, with a semicolon. Either way, the effect is the same. See “Writing the commands in rules” on page 349.

Because dollar signs are used to start variable references, if you really want a dollar sign in a rule you must write two of them, ‘*\$\$*’ (see “How to use variables” on page 367). You may split a long line by inserting a backslash followed by a newline, but this is not required, as *make* places no limit on the length of a line in a makefile.

A rule tells *make* two things: when the targets are out of date, and how to update them when necessary.

The criterion for being out of date is specified in terms of the *dependencies*, which consist of file names separated by spaces. Wildcards and archive members (see “Using *make* to update archive files” on page 449) are allowed too. A target is out of date if it does not exist or if it is older than any of the dependencies (by comparison of last-modification times). The idea is that the contents of the target file are computed based on information in the dependencies, so if any of the dependencies changes, the contents of the existing target file are no longer necessarily valid.

How to update is specified by *commands*. These are lines to be executed by the shell (normally, ‘*sh*’), but with some extra features (see “Writing the commands in rules” on page 349).

Using wildcard characters in file names

A single file name can specify many files using *wildcard characters*. The wildcard characters in make are ‘*’, ‘?’ and ‘[. . .]’, the same as in the Bourne shell. For example, ‘*.c’ specifies a list of all the files (in the working directory) whose names end in ‘.c’.

The character ‘~’ at the beginning of a file name also has special significance. If alone, or followed by a slash, it represents your home directory. For example ‘~/bin’ expands to ‘/home/you/bin’. If the ‘~’ is followed by a word, the string represents the home directory of the user named by that word. For example ‘~john/bin’ expands to ‘/home/john/bin’.

Wildcard expansion happens automatically in targets, in dependencies, and in commands (where the shell does the expansion). In other contexts, wildcard expansion happens only if you request it explicitly with the `wildcard` function. The special significance of a wildcard character can be turned off by preceding it with a backslash. Thus, ‘foo*bar’ would refer to a specific file whose name consists of ‘foo’, an asterisk, and ‘bar’.

Wildcard examples

Wildcards can be used in the commands of a rule, where they are expanded by the shell. For example, here is a rule to delete all the object files:

```
clean:
    rm -f *.o
```

Wildcards are also useful in the dependencies of a rule. With the following rule in the makefile, ‘make print’ will print all the ‘.c’ files that have changed since the last time you printed them:

```
print: *.c
    lpr -p $?
    touch print
```

This rule uses ‘print’ as an empty target file; see “Empty target files to record events” on page 337. (The ‘\$?’ automatic variable is used to print only those files that have changed; see “Automatic variables” on page 438.) Wildcard expansion does not happen when you define a variable. Thus, if you write `objects = *.o`, then the value of the variable `objects` is the actual string ‘*.o’. However, if you use the value of `objects` in a target, dependency or command, wildcard expansion will take place at that time. To set `objects` to the expansion, instead use: `objects := $(wildcard *.o)`. See “The wildcard function” on page 329.

Pitfalls of using wildcards

The next is an example of a naive way of using wildcard expansion that does not do what you would intend. Suppose you would like to say that the executable file, `foo`, is made from all the object files in the directory, and you write the following.

```
objects = *.o

foo : $(objects)
    cc -o foo $(CFLAGS) $(objects)
```

The value of `objects` is the actual string `*.o`. Wildcard expansion happens in the rule for `foo`, so that each existing `.o` file becomes a dependency of `foo` and will be recompiled if necessary. But what if you delete all the `.o` files? When a wildcard matches no files, it is left as it is, so then `foo` will depend on the oddly-named file `*.o`. Since no such file is likely to exist, make will give you an error saying it cannot figure out how to make `*.o`. This is not what you want! Actually it is possible to obtain the desired result with wildcard expansion, but you need more sophisticated techniques, including the wildcard function and string substitution. These are described with “The wildcard function” on page 329.

The wildcard function

Wildcard expansion happens automatically in rules. But wildcard expansion does not normally take place when a variable is set, or inside the arguments of a function. If you want to do wildcard expansion in such places, you need to use the `wildcard` function, using: `$(wildcard pattern...)`.

This string, used anywhere in a makefile, is replaced by a space-separated list of names of existing files that match one of the given file name patterns. If no existing file name matches a pattern, then that pattern is omitted from the output of the `wildcard` function. Note that this is different from how unmatched wildcards behave in rules, where they are used verbatim rather than ignored (see “Pitfalls of using wildcards” on page 328).

One use of the `wildcard` function is to get a list of all the C source files in a directory, using: `$(wildcard *.c)`.

We can change the list of C source files into a list of object files by replacing the ‘.o’ suffix with ‘.c’ in the result, using the following.

```
$(patsubst %.c,%.o,$(wildcard *.c))
```

Here we have used another function, `patsubst`. See “Functions for string substitution and analysis” on page 395.

Thus, a makefile to compile all C source files in the directory and then link them together could be written as follows.

```
objects := $(patsubst %.c,%.o,$(wildcard *.c))

foo : $(objects)
    cc -o foo $(objects)
```

This takes advantage of the implicit rule for compiling C programs, so there is no need to write explicit rules for compiling the files. See “The two flavors of variables” on page 370 for an explanation of ‘:=’, which is a variant of ‘=’.)

Searching directories for dependencies

For large systems, it is often desirable to put sources in a separate directory from the binaries. The *directory search* features of `make` facilitate this by searching several directories automatically to find a dependency. When you redistribute the files among directories, you do not need to change the individual rules, just the search paths.

VPATH: search path for all dependencies

The value of the `make` variable, `VPATH`, specifies a list of directories that `make` should search. Most often, the directories are expected to contain dependency files that are not in the current directory; however, `VPATH` specifies a search list that `make` applies for all files, including files which are targets of rules. Thus, if a file that is listed as a target or dependency does not exist in the current directory, `make` searches the directories listed in `VPATH` for a file with that name. If a file is found in one of them, that file becomes the dependency. Rules may then specify the names of source files in the dependencies as if they all existed in the current directory. See “Writing shell commands with directory search” on page 331. In the `VPATH` variable, directory names are separated by colons or blanks. The order in which directories are listed is the order followed by `make` in its search. For example, `VPATH = src:../headers` specifies a path containing two directories, ‘`src`’ and ‘`../headers`’, which `make` searches in that order. With this value of `VPATH`, the rule, `foo.o : foo.c` is interpreted as if it were written: `foo.o : src/foo.c`.

This is assuming the file, ‘`foo.c`’, does not exist in the current directory but is found in the directory, ‘`src`’.

The `vpath` directive

Similar to the `VPATH` variable but more selective is the `vpath` directive (note the use of lower case) which allows you to specify a search path for a particular class of file names, those that match a particular pattern. Thus you can supply certain search directories for one class of file names and other directories (or none) for other file names. There are three forms of the `vpath` directive.

`vpath pattern directories`

Specify the search path directories for file names that match `pattern`.

The search path, *directories*, is a list of directories to be searched, separated by colons or blanks, just like the search path used in the `VPATH` variable.

`vpath pattern`

Clear out the search path associated with *pattern*.

`vpath`

Clear all search paths previously specified with `vpath` directives.

A `vpath` pattern is a string containing a `%` character. The string must match the file name of a dependency that is being searched for, the `%` character matching any sequence of zero or more characters (as in *pattern rules*; see “Defining and redefining pattern rules” on page 436). For example, `%.h` matches files that end in `.h`. (If there is no `%`, the pattern must match the dependency exactly, which is not useful very often.)

`%` characters in a `vpath` directive’s pattern can be quoted with preceding backslashes (`\`). Backslashes that would otherwise quote `%` characters can be quoted with more backslashes. Backslashes that quote `%` characters or other backslashes are removed from the pattern before it is compared to file names. Backslashes that are not in danger of quoting `%` characters go unmolested.

When a dependency fails to exist in the current directory, if the *pattern* in a `vpath` directive matches the name of the dependency file, then the *directories* in that directive are searched just like (and before) the directories in the `VPATH` variable.

For example, `vpath %.h ../headers` tells `make` to look for any dependency whose name ends in `.h` in the directory `../headers` if the file is not found in the current directory.

If several `vpath` patterns match the dependency file’s name, then `make` processes each matching `vpath` directive one by one, searching all the directories mentioned in each directive. `make` handles multiple `vpath` directives in the order in which they appear in the makefile; multiple directives with the same pattern are independent of each other. Thus, the following directive will look for a file ending in `.c` in `foo`, then `blish`, then `bar`.

```
vpath %.c foo
vpath % blish
vpath %.c bar
```

The next example, on the other hand, will look for a file ending in `.c` in `foo`, then `bar`, then `blish`.

```
vpath %.c foo:bar
vpath % blish
```

Writing shell commands with directory search

When a dependency is found in another directory through directory search, this cannot change the commands of the rule; they will execute as written. Therefore, you must write the commands with care so that they will look for the dependency in the directory where `make` finds it.

This is done with the *automatic variables* such as `$(^)` (see “Automatic variables” on page 438).

For instance, the value of ‘`$^`’ is a list of all the dependencies of the rule, including the names of the directories in which they were found, and the value of ‘`$@`’ is the target, as in the following example.

```
foo.o : foo.c
        cc -c $(CFLAGS) $^ -o $@
```

The variable `CFLAGS` exists so you can specify flags for C compilation by implicit rules; we use it here for consistency so it will affect all C compilations uniformly; see “Variables used by implicit rules” on page 432.

Often the dependencies include header files as well, which you do not want to mention in the commands.

The automatic variable ‘`$<`’ is just the first dependency, as in the following declaration.

```
VPATH = src:../headers
foo.o : foo.c defs.h
hack.h cc -c $(CFLAGS) $< -o $@
```

Directory search and implicit rules

The search through the directories specified in `VPATH` or with `vpath` also happens during consideration of implicit rules (see “Using implicit rules” on page 425).

For example, when a file ‘`foo.o`’ has no explicit rule, `make` considers implicit rules, such as the built-in rule to compile ‘`foo.c`’ if that file exists.

If such a file is lacking in the current directory, the appropriate directories are searched for it.

If ‘`foo.c`’ exists (or is mentioned in the makefile) in any of the directories, the implicit rule for C compilation is applied.

The commands of implicit rules normally use automatic variables as a matter of necessity; consequently they will use the file names found by directory search with no extra effort.

Directory search for link libraries

Directory search applies in a special way to libraries used with the linker.

This special feature comes into play when you write a dependency whose name is of the form, ‘`-lname`’. (You can tell something strange is going on here because the dependency is normally the name of a file, and the *file name* of the library looks like ‘`lib name.a`’, not like ‘`-lname`’.)

When a dependency's name has the form `-lname`, `make` handles it specially by searching for the file `libname.a` in the current directory, in directories specified by matching `vpath` search paths and the `VPATH` search path, and then in the directories `/lib`, `/usr/lib`, and `prefix/lib` (normally, `/usr/local/lib`). Use the following example, for instance.

```
foo : foo.c -lcurses
    cc $^ -o $@
```

This would cause the command, `cc foo.c /usr/lib/libcurses.a -o foo`, to execute when `foo` is older than `foo.c` or `/usr/lib/libcurses.a`.

Phony targets

A phony target is one that is not really the name of a file. It is just a name for some commands to be executed when you make an explicit request. There are two reasons to use a phony target: to avoid a conflict with a file of the same name, and to improve performance.

If you write a rule whose commands will not create the target file, the commands will be executed every time the target comes up for remaking. Use the following, for example.

```
clean:
    rm *.o temp
```

Because the `rm` command does not create a file named `'clean'`, probably no such file will ever exist. Therefore, the `rm` command will be executed every time you use `'make clean'`.

The phony target will cease to work if anything ever does create a file named `'clean'` in this directory. Since it has no dependencies, the file `'clean'` would inevitably be considered up to date, and its commands would not be executed. To avoid this problem, you can explicitly declare the target to be phony, using the special target, `.PHONY` (see “Special built-in target names” on page 338), as in: `.PHONY : clean`.

Once this is done, `'make clean'` will run the commands regardless of whether there is a file named `'clean'`. Since it knows that phony targets do not name actual files that could be remade from other files, `make` skips the implicit rule search for phony targets (see “Implicit rules” on page 423). This is why declaring a target phony is good for performance, even if you are not worried about the actual file existing.

Thus, you first write the line that states that `clean` is a phony target, then you write the rule, like the following example.

```
.PHONY: clean
clean:
    rm *.o temp
```

A phony target should not be a dependency of a real target file; if it is, its commands are run every time `make` goes to update that file. As long as a phony target is never a dependency of a real target, the phony target commands will be executed only when the phony target is a specified goal (see “Arguments to specify the goals” on page 409).

Phony targets can have dependencies. When one directory contains multiple programs, it is most convenient to describe all of the programs in one makefile, `./Makefile`. Since the target remade by default will be the first one in the makefile, it is common to make this a phony target named `all` and give it, as dependencies, all the individual programs. Use the following, for example.

```
all : prog1 prog2 prog3
.PHONY : all

prog1 : prog1.o utils.o
       cc -o prog1 prog1.o utils.o

prog2 : prog2.o
       cc -o prog2 prog2.o

prog3 : prog3.o sort.o utils.o
       cc -o prog3 prog3.o sort.o utils.o
```

Now you can use just `make` to remake all three programs, or specify as arguments the ones to remake (as in `make prog1 prog3`).

When one phony target is a dependency of another, it serves as a subroutine of the other. For instance, in the following example, `make cleanall` will delete the object files, the difference files, and the file, `program`.

```
.PHONY: cleanall cleanobj cleandiff

cleanall : cleanobj cleandiff
          rm program

cleanobj :
          rm *.o

cleandiff :
          rm *.diff
```

Rules without commands or dependencies

If a rule has no dependencies or commands, and the target of the rule is a nonexistent file, then `make` imagines this target to have been updated whenever its rule is run. This implies that all targets depending on this one will always have their commands run. The following example will illustrate the rule.

```
clean: FORCE
    rm $(objects)

FORCE:
```

In this case, the target, `'FORCE'`, satisfies the special conditions, so the target, `'clean'`, that depends on it is forced to run its commands. There is nothing special about the name, `'FORCE'`, but that is one name commonly used this way. As you can see, using `'FORCE'` this way has the same results as using `'.PHONY: clean'`. Using `'.PHONY'` is more explicit and more efficient. However, other versions of `make` do not support `'.PHONY'`; thus `'FORCE'` appears in many makefiles. See “Phony targets” on page 334.

Empty target files to record events

The *empty target* is a variant of the phony target; it is used to hold commands for an action that you request explicitly from time to time. Unlike a phony target, this target file can really exist; but the file's contents do not matter, and usually are empty.

The purpose of the empty target file is to record, with its last-modification time, when the rule's commands were last executed. It does so because one of the commands is a `touch` command to update the target file.

The empty target file must have some dependencies. When you ask to remake the empty target, the commands are executed if any dependency is more recent than the target; in other words, if a dependency has changed since the last time you remade the target. Use the following as an example.

```
print: foo.c bar.c
      lpr -p $?
      touch print
```

With this rule, `'make print'` will execute the `lpr` command if either source file has changed since the last `'make print'`. The automatic variable `'$?'` is used to print only those files that have changed (see “Automatic variables” on page 438).

Special built-in target names

Certain names have special meanings if they appear as targets.

`.PHONY`

The dependencies of the special target, `.PHONY`, are considered to be phony targets. When it is time to consider such a target, `make` will run its commands unconditionally, regardless of whether a file with that name exists or what its last-modification time is. See “Phony targets” on page 334.

`.SUFFIXES`

The dependencies of the special target, `.SUFFIXES`, are the list of suffixes to be used in checking for suffix rules. See “Old-fashioned suffix rules” on page 445.

`.DEFAULT`

The commands specified for `.DEFAULT` are used for any target for which no rules are found (either explicit rules or implicit rules). See “Defining last-resort default rules” on page 444. If `.DEFAULT` commands are specified, every file mentioned as a dependency, but not as a target in a rule, will have these commands executed on its behalf. See “Implicit rule search algorithm” on page 447.

`.PRECIOUS`

The targets which `.PRECIOUS` depends on are given the following special treatment: if `make` is killed or interrupted during the execution of their commands, the target is not deleted. See “Interrupting or killing the `make` utility” on page 357. Also, if the target is an intermediate file, it will not be deleted after it is no longer needed, as is normally done. See “Chains of implicit rules” on page 435.

You can also list the target pattern of an implicit rule (such as `% .o`) as a dependency file of the special target, `.PRECIOUS`, to preserve intermediate files created by rules whose target patterns match that file’s name.

`.IGNORE`

If you specify dependencies for `.IGNORE`, then `make` will ignore errors in execution of the commands run for those particular files. The commands for `.IGNORE` are not meaningful.

If mentioned as a target with no dependencies, `.IGNORE` says to ignore errors in execution of commands for all files. This usage of `.IGNORE` is supported only for historical compatibility. Since this affects every command in the makefile, it is not very useful; we recommend you use the more selective ways to ignore errors in specific commands. See “Errors in commands” on page 355.

`.SILENT`

If you specify dependencies for `.SILENT`, then `make` will not print commands to remake those particular files before executing them. The commands for `.SILENT` are not meaningful.

If mentioned as a target with no dependencies, `.SILENT` says not to print any commands before executing them. This usage of `‘.SILENT’` is supported only for historical compatibility. We recommend you use the more selective ways to silence specific commands. See “Command echoing” on page 351.

If you want to silence all commands for a particular run of `make`, use the `‘-s’` or `‘--silent’` options. See “Summary of make options” on page 417.

`.EXPORT_ALL_VARIABLES`

Simply by being mentioned as a target, this tells `make` to export all variables to child processes by default.

See “Communicating variables to a sub-make utility” on page 359.

Any defined implicit rule suffix also counts as a special target if it appears as a target, and so does the concatenation of two suffixes, such as `‘.c.o’`. These targets are suffix rules, an obsolete way of defining implicit rules (but a way still widely used). In principle, any target name could be special in this way if you break it in two and add both pieces to the suffix list. In practice, suffixes normally begin with `‘.’`, so these special target names also begin with `‘.’`. See “Old-fashioned suffix rules” on page 445.

Multiple targets in a rule

A rule with multiple targets is equivalent to writing many rules, each with one target, and all identical aside from that issue. The same commands apply to all the targets, although their effects may vary since you substitute an actual target name into the command using `$(target)`. The rule also contributes the same dependencies to all the targets.

This is useful in two cases.

- You want just dependencies, no commands. Use the following for an example.

```
kbd.o command.o files.o: command.h
```

This input gives an additional dependency to each of the three object files mentioned.

- Similar commands work for all the targets. The commands do not need to be absolutely identical, since the automatic variable `$(target)` can be used to substitute the particular target to be remade into the commands (see “Automatic variables” on page 438). Use the following for an example.

```
bigoutput littleoutput : text.g
generate text.g -$(subst output,,$(target)) > $(target)
```

This input is equivalent to the next example.

```
bigoutput : text.g
generate text.g -big > bigoutput
littleoutput : text.g
generate text.g -little > littleoutput
```

Here we assume the hypothetical program, `generate`, makes two types of output, one if given `-big` and one if given `-little`. See “Functions for string substitution and analysis” on page 395 for an explanation of the `subst` function.

Suppose you would like to vary the dependencies according to the target, much as the variable `$(target)` allows you to vary the commands. You cannot do this with multiple targets in an ordinary rule, but you can do it with a *static pattern rule*. See “Static pattern rules” on page 342.

Multiple rules for one target

One file can be the target of several rules. All the dependencies mentioned in all the rules are merged into one list of dependencies for the target. If the target is older than any dependency from any rule, the commands are executed.

There can only be one set of commands to be executed for a file. If more than one rule gives commands for the same file, `make` uses the last set given and prints an error message. (As a special case, if the file's name begins with a dot, no error message is printed. This odd behavior is only for compatibility with other implementations of `make`.) There is no reason to write your makefiles this way; that is why `make` gives you an error message.

An extra rule with just dependencies can be used to give a few extra dependencies to many files at once. For example, one usually has a variable named `objects` containing a list of all the compiler output files in the system being made. An easy way to say that all of them must be recompiled if `'config.h'` changes is to write the following input.

```
objects = foo.o bar.o
foo.o : defs.h
bar.o : defs.h test.h
$(objects) : config.h
```

This could be inserted or taken out without changing the rules that really specify how to make the object files, making it a convenient form to use if you wish to add the additional dependency intermittently. Another wrinkle is that the additional dependencies could be specified with a variable that you set with a command argument to `make` (see “Overriding variables” on page 414). Use the following for an example.

```
extradeps=
$(objects) : $(extradeps)
```

This input means that the command `'make extradeps=foo.h'` will consider `'foo.h'` as a dependency of each object file, but plain `'make'` will not. If none of the explicit rules for a target has commands, then `make` searches for an applicable implicit rule to find some commands see “Using implicit rules” on page 425).

Static pattern rules

Static pattern rules are rules which specify multiple targets and construct the dependency names for each target based on the target name. They are more general than ordinary rules with multiple targets because the targets do not have to have identical dependencies. Their dependencies must be *analogous* but not necessarily *identical*.

Syntax of static pattern rules

The following shows the syntax of a static pattern rule:

```
targets ...: target-pattern: dep-patterns ...  
      commands  
...
```

The *targets* list specifies the targets to which the rule applies. The targets can contain wildcard characters, just like the targets of ordinary rules (see “Using wildcard characters in file names” on page 326).

The *target-pattern* and *dep-patterns* say how to compute the dependencies of each target.

Each target is matched against the *target-pattern* to extract a part of the target name, called the *stem*.

This stem is substituted into each of the *dep-patterns* to make the dependency names (one from each *dep-pattern*).

Each pattern normally contains the character ‘%’ just once.

When the *target-pattern* matches a target, the ‘%’ can match any part of the target name; this part is called the *stem*. The rest of the pattern must match exactly. For example, the target ‘foo.o’ matches the pattern ‘%.o’, with ‘foo’ as the stem. The targets ‘foo.c’ and ‘foo.out’ do not match that pattern.

The dependency names for each target are made by substituting the stem for the ‘%’ in each dependency pattern. For example, if one dependency pattern is ‘%.c’, then substitution of the stem ‘foo’ gives the dependency name ‘foo.c’. It is legitimate to write a dependency pattern that does not contain ‘%’; then this dependency is the same for all targets.

‘%’ characters in pattern rules can be quoted with preceding backslashes (‘\’). Backslashes that would otherwise quote ‘%’ characters can be quoted with more backslashes. Backslashes that quote ‘%’ characters or other backslashes are removed from the pattern before it is compared to file names or has a stem substituted into it. Backslashes that are not in danger of quoting ‘%’ characters go unmolested. For example, the pattern ‘the\%weird\\%pattern\\’ has ‘the%weird\’ preceding the operative ‘%’ character, and ‘pattern\\’ following it. The final two backslashes are left alone because they cannot affect any ‘%’ character.

The following is an example which compiles each of ‘foo.o’ and ‘bar.o’ from the corresponding ‘.c’ file.

```
objects = foo.o bar.o

$(objects): %.o: %.c
    $(CC) -c $(CFLAGS) $< -o $@
```

In the previous example, ‘\$<’ is the automatic variable that holds the name of the dependency and ‘\$@’ is the automatic variable that holds the name of the target; see “Automatic variables” on page 438.

Each target specified must match the target pattern; a warning is issued for each target that does not. If you have a list of files, only some of which will match the pattern, you can use the filter function to remove nonmatching file names (see “Functions for string substitution and analysis” on page 395), as in the following example.

```
files = foo.elc bar.o lose.o

$(filter %.o,$(files)): %.o: %.c
    $(CC) -c $(CFLAGS) $< -o $@
$(filter %.elc,$(files)): %.elc: %.el
    emacs -f batch-byte-compile $<
```

In this example the result of ‘\$(filter %.o,\$(files))’ is ‘bar.o lose.o’, and the first static pattern rule causes each of these object files to be updated by compiling the corresponding C source file. The result of ‘\$(filter %.elc,\$(files))’ is ‘foo.elc’, so that file is made from ‘foo.el’.

The following example shows how to use \$* in static pattern rules.

```
bigoutput littleoutput : %output : text.g
    generate text.g -$* > $@
```

When the generate command is run, \$* will expand to the stem, either ‘big’ or ‘little’.

Static pattern rules versus implicit rules

A static pattern rule has much in common with an implicit rule defined as a pattern rule (see “Defining and redefining pattern rules” on page 436). Both have a pattern for the target and patterns for constructing the names of dependencies. The difference is in how `make` decides *when* the rule applies.

An implicit rule can apply to any target that matches its pattern, but it *does* apply only when the target has no commands otherwise specified, and only when the dependencies can be found. If more than one implicit rule appears applicable, only one applies; the choice depends on the order of rules.

By contrast, a static pattern rule applies to the precise list of targets that you specify in the rule. It cannot apply to any other target and it invariably does apply to each of the targets specified. If two conflicting rules apply, and both have commands, that’s an error. The static pattern rule can be better than an implicit rule for the following reasons.

- You may wish to override the usual implicit rule for a few files whose names cannot be categorized syntactically but can be given in an explicit list.
- If you cannot be sure of the precise contents of the directories you are using, you may not be sure which other irrelevant files might lead `make` to use the wrong implicit rule. The choice might depend on the order in which the implicit rule search is done. With static pattern rules, there is no uncertainty: each rule applies to precisely the targets specified.

Double-colon rules

Double-colon rules are rules written with ‘: :’ instead of ‘:’ after the target names. They are handled differently from ordinary rules when the same target appears in more than one rule.

When a target appears in multiple rules, all the rules must be the same type: all ordinary, or all double-colon. If they are double-colon, each of them is independent of the others. Each double-colon rule’s commands are executed if the target is older than any dependencies of that rule. This can result in executing none, any, or all of the double-colon rules.

Double-colon rules with the same target are in fact completely separate from one another. Each double-colon rule is processed individually, just as rules with different targets are processed.

The double-colon rules for a target are executed in the order they appear in the makefile. However, the cases where double-colon rules really make sense are those where the order of executing the commands would not matter.

Double-colon rules are somewhat obscure and not often very useful; they provide a mechanism for cases in which the method used to update a target differs depending on which dependency files caused the update, and such cases are rare.

Each double-colon rule should specify commands; if it does not, an implicit rule will be used if one applies. See “Using implicit rules” on page 425.

Generating dependencies automatically

In the makefile for a program, many of the rules you need to write often say only that some object file depends on some header file. For example, if `'main.c'` uses `'defs.h'` via an `#include`, you would write: `main.o: defs.h`.

You need this rule so that `make` knows that it must remake `'main.o'` whenever `'defs.h'` changes. You can see that for a large program you would have to write dozens of such rules in your makefile. And, you must always be very careful to update the makefile every time you add or remove an `#include`.

To avoid this hassle, most modern C compilers can write these rules for you, by looking at the `#include` lines in the source files.

Usually this is done with the `'-M'` option to the compiler. For example, the command, `cc -M main.c`, generates the output: `main.o : main.c defs.h`.

Thus you no longer have to write all those rules yourself. The compiler will do it for you.

NOTE: Such a dependency constitutes mentioning `'main.o'` in a makefile, so it can never be considered an intermediate file by implicit rule search. This means that `make` won't ever remove the file after using it; see “Chains of implicit rules” on page 435.

With old `make` programs, it was common practice to use this compiler feature to generate dependencies on demand with a command like `'make depend'`.

That command would create a file `'depend'` containing all the automatically-generated dependencies; then the makefile could use `include` to read them in (see “Including other makefiles” on page 316).

In GNU `make`, the feature of remaking makefiles makes this practice obsolete—you need never tell `make` explicitly to regenerate the dependencies, because it always regenerates any makefile that is out of date. See “How makefiles are remade” on page 319.

The practice we recommend for automatic dependency generation is to have one makefile corresponding to each source file.

For each source file, `'name.c'`, there is a makefile, `'name.d'`, listing which files on which the object file, `'name.o'`, depends.

That way only the source files that have changed need to be rescanned to produce the new dependencies.

The following is an example of the pattern rule to generate a file of dependencies (i.e., a makefile) called ‘*name.d*’ from a C source file called ‘*name.c*’.

```
%.d: %.c
    $(SHELL) -ec '$(CC) -M $(CPPFLAGS) $< \
    | sed '\''s/$*\\.o[ :]*/& $@/g'\'' > $@'
```

See “Defining and redefining pattern rules” on page 436 for information on defining pattern rules. The ‘-e’ flag to the shell makes it exit immediately if the `$(CC)` command fails (exits with a nonzero status). Normally the shell exits with the status of the last command in the pipeline (`sed` in this case), so `make` would not notice a nonzero status from the compiler.

With the GNU C compiler, you may wish to use the ‘-MM’ flag instead of ‘-M’.

This omits dependencies on system header files. See “Options Controlling the Preprocessor” in *Using GNU CC* in **GNUPro Compiler Tools** for details. For example, the purpose of the `sed` command is to translate `main.o : main.c defs.h` into: `main.o main.d : main.c defs.h`.

This makes each ‘.d’ file depend on all the source and header files on which the corresponding ‘.o’ file depends. `make` then knows it must regenerate the dependencies whenever any of the source or header files changes. Once you’ve defined the rule to remake the ‘.d’ files, you then use the `include` directive to read them all in. See “Including other makefiles” on page 316. Use the following example, for clarification.

```
sources = foo.c bar.c

include $(sources:.c=.d)
```

This example uses a substitution variable reference to translate the list of source files, ‘`foo.c bar.c`’, into a list of dependency makefiles, ‘`foo.d bar.d`’. See “Substitution references” on page 373 for full information on substitution references.) Since the ‘.d’ files are makefiles like any others, `make` will remake them as necessary with no further work from you. See “How makefiles are remade” on page 319.

5

Writing the commands in rules

The commands of a rule consist of shell command lines to be executed one by one. The following documentation discusses this execution of shell commands.

- “Command echoing” on page 351
- “Command execution” on page 352
- “Parallel execution” on page 353
- “Errors in commands” on page 355
- “Interrupting or killing the make utility” on page 357
- “Recursive use of the make utility” on page 358
- “Defining canned command sequences” on page 364
- “Using empty commands” on page 366

Each command line must start with a tab, except that the first command line may be attached to the *target-and-dependencies* line with a semicolon in between. Blank lines and lines of just comments may appear among the command lines; they are ignored. (But beware, an apparently “blank” line that begins with a tab is *not* blank! It is an empty command; see “Using empty commands” on page 366.)

Users use many different shell programs, but commands in makefiles are always interpreted by `/bin/sh` unless the makefile specifies otherwise. See “Command execution” on page 352.

The shell that is in use determines whether comments can be written on command lines, and what syntax they use. When the shell is `/bin/sh`, a `#` starts a comment that extends to the end of the line. The `#` does not have to be at the beginning of a line. Text on a line before a `#` is not part of the comment.

Command echoing

Normally `make` prints each command line before it is executed. We call this *echoing* because it gives the appearance that you are typing the commands yourself.

When a line starts with '@', the echoing of that line is suppressed. The '@' is discarded before the command is passed to the shell. Typically you would use this for a command whose only effect is to print something, such as an `echo` command to indicate progress through the makefile:

```
@echo About to make distribution files
```

When `make` is given the flag '-n' or '--just-print', echoing is all that happens with no execution. See “Summary of make options” on page 417. In this case and only this case, even the commands starting with '@' are printed. This flag is useful for finding out which commands `make` thinks are necessary without actually doing them.

The '-s' or '--silent' flag to `make` prevents all echoing, as if all commands started with '@'. A rule in the makefile for the special target, `.SILENT`, without dependencies has the same effect (see “Special built-in target names” on page 338). `.SILENT` is essentially obsolete since '@' is more flexible.

Command execution

When it is time to execute commands to update a target, they are executed by making a new subshell for each line. (In practice, `make` may take shortcuts that do not affect the results.)

NOTE: The implication that shell commands such as `cd` set variables local to each process will not affect the following command lines. If you want to use `cd` to affect the next command, put the two on a single line with a semicolon between them. Then `make` will consider them a single command and pass them, together, to a shell which will execute them in sequence. Use the following example for clarification.

```
foo : bar/lose
      cd bar; gobble lose > ../foo
```

If you would like to split a single shell command into multiple lines of text, you must use a backslash at the end of all but the last subline. Such a sequence of lines is combined into a single line, by deleting the backslash-newline sequences, before passing it to the shell. Thus, the following is equivalent to the preceding example.

```
foo : bar/lose
      cd bar; \
      gobble lose > ../foo
```

The program used as the shell is taken from the variable, `SHELL`. By default, the program `/bin/sh` is used.

Unlike most variables, `SHELL` is never set from the environment. This is because the `SHELL` environment variable is used to specify your personal choice of shell program for interactive use.

It would be very bad for personal choices like this to affect the functioning of `makefiles`. See “Variables from the environment” on page 383.

Parallel execution

GNU `make` knows how to execute several commands at once. Normally, `make` will execute only one command at a time, waiting for it to finish before executing the next. However, the `-j` or `--jobs` option tells `make` to execute many commands simultaneously.

If the `-j` option is followed by an integer, this is the number of commands to execute at once; this is called the number of *job slots*. If there is nothing looking like an integer after the `-j` option, there is no limit on the number of job slots. The default number of job slots is one which means serial execution (one thing at a time).

One unpleasant consequence of running several commands simultaneously is that output from all of the commands comes when the commands send it, so messages from different commands may be interspersed.

Another problem is that two processes cannot both take input from the same device; so to make sure that only one command tries to take input from the terminal at once, `make` will invalidate the standard input streams of all but one running command. This means that attempting to read from standard input will usually be a fatal error (a ‘Broken pipe’ signal) for most child processes if there are several.

It is unpredictable which command will have a valid standard input stream (which will come from the terminal, or wherever you redirect the standard input of `make`). The first command run will always get it first, and the first command started after that one finishes will get it next, and so on.

We will change how this aspect of `make` works if we find a better alternative. In the mean time, you should *not* rely on any command using standard input at all if you are using the parallel execution feature; but if you are *not* using this feature, then standard input works normally in all commands.

If a command fails (i.e., it is killed by a signal or exits with a nonzero status), and errors are not ignored for that command (see “Errors in commands” on page 355), the remaining command lines to remake the same target will not be run. If a command fails and the `-k` or `--keep-going` option was not given (see “Summary of make options” on page 417), `make` aborts execution. If `make` terminates for any reason (including a signal) with child processes running, it waits for them to finish before actually exiting.

When the system is heavily loaded, you will probably want to run fewer jobs than when it is lightly loaded. You can use the `-l` option to tell `make` to limit the number of jobs to run at once, based on the load average. The `-l` or `--max-load` option is followed by a floating-point number.

For example, `-l 2.5` will not let `make` start more than one job if the load average is above 2.5. The `'-l'` option with no following number removes the load limit, if one was given with a previous `'-l'` option.

More precisely, when `make` goes to start up a job, and it already has at least one job running, it checks the current load average; if it is not lower than the limit given with `'-l'`, `make` waits until the load average goes below that limit, or until all the other jobs finish. By default, there is no load limit.

Errors in commands

After each shell command returns, `make` looks at its exit status. If the command completed successfully, the next command line is executed in a new shell; after the last command line is finished, the rule is finished.

If there is an error (the exit status is nonzero), `make` gives up on the current rule, and perhaps on all rules.

Sometimes the failure of a certain command does not indicate a problem. For example, you may use the `mkdir` command to ensure that a directory exists. If the directory already exists, `mkdir` will report an error, but you probably want `make` to continue regardless.

To ignore errors in a command line, write a ‘-’ at the beginning of the line’s text (after the initial tab). The ‘-’ is discarded before the command is passed to the shell for execution, as in the following example.

```
clean:
    -rm -f *.o
```

This causes `rm` to continue even if it is unable to remove a file.

When you run `make` with the ‘-i’ or ‘--ignore-errors’ flag, errors are ignored in all commands of all rules. A rule in the makefile for the special target, `.IGNORE`, has the same effect, if there are no dependencies. These ways of ignoring errors are obsolete because ‘-’ is more flexible.

When errors are to be ignored, because of either a ‘-’ or the ‘-i’ flag, `make` treats an error return just like success, except that it prints out a message that tells you the status code the command exited with, and says that the error has been ignored.

When an error happens that `make` has not been told to ignore, it implies that the current target cannot be correctly remade, and neither can any other that depends on it either directly or indirectly. No further commands will be executed for these targets, since their preconditions have not been achieved.

Normally `make` gives up immediately in this circumstance, returning a nonzero status. However, if the ‘-k’ or ‘--keep-going’ flag is specified, `make` continues to consider the other dependencies of the pending targets, remaking them if necessary, before it gives up and returns nonzero status. For example, after an error in compiling one object file, ‘`make -k`’ will continue compiling other object files even though it already knows that linking them will be impossible. See “Summary of make options” on page 417.

The usual behavior assumes that your purpose is to get the specified targets up to date; once `make` learns that this is impossible, it might as well report the failure immediately. The `-k` option says that the real purpose is to test as many of the changes made in the program as possible, perhaps to find several independent problems so that you can correct them all before the next attempt to compile.

This is why Emacs' `compile` command passes the `-k` flag by default.

Usually when a command fails, if it has changed the target file at all, the file is corrupted and cannot be used—or at least it is not completely updated. Yet the file's timestamp says that it is now up to date, so the next time `make` runs, it will not try to update that file. The situation is just the same as when the command is killed by a signal; see “Interrupting or killing the make utility” on page 357. So generally the right thing to do is to delete the target file if the command fails after beginning to change the file. `make` will do this if `.DELETE_ON_ERROR` appears as a target. This is almost always what you want `make` to do, but it is not historical practice; so for compatibility, you must explicitly request it.

Interrupting or killing the `make` utility

If `make` gets a fatal signal while a command is executing, it may delete the target file that the command was supposed to update. This is done if the target file's last-modification time has changed since `make` first checked it.

The purpose of deleting the target is to make sure that it is remade from scratch when `make` is next run. Why is this? Suppose you use `Ctrl-c` while a compiler is running, and it has begun to write an object file `'foo.o'`. The `Ctrl-c` kills the compiler, resulting in an incomplete file whose last-modification time is newer than the source file `'foo.c'`. But `make` also receives the `Ctrl-c` signal and deletes this incomplete file. If `make` did not do this, the next invocation of `make` would think that `'foo.o'` did not require updating—resulting in a strange error message from the linker when it tries to link an object file half of which is missing.

You can prevent the deletion of a target file in this way by making the special target, `.PRECIOUS`, depend on it. Before remaking a target, `make` checks to see whether it appears on the dependencies of `.PRECIOUS`, and thereby decides whether the target should be deleted if a signal happens. Some reasons why you might do this are that the target is updated in some atomic fashion, or exists only to record a modification-time (its contents do not matter), or must exist at all times to prevent other sorts of trouble.

Recursive use of the `make` utility

Recursive use of `make` means using `make` as a command in a makefile. This technique is useful when you want separate makefiles for various subsystems that compose a larger system. For example, suppose you have a subdirectory, `'subdir'`, which has its own makefile, and you would like the containing directory's makefile to run `make` on the subdirectory. You can do it by writing the following

```
subsystem:
    cd subdir; $(MAKE)
```

Or, equivalently (see “Summary of make options” on page 417), use the following input.

```
subsystem:
    $(MAKE) -C subdir
```

You can write recursive `make` commands just by copying this example, but there are many things to know about how they work and why, and about how the sub-`make` relates to the top-level `make`.

How the `MAKE` variable works

Recursive `make` commands should always use the variable, `MAKE`, not the explicit command name, `make`, as the following example shows.

```
subsystem:
    cd subdir; $(MAKE)
```

The value of this variable is the file name with which `make` was invoked. If this file name was `'/bin/make'`, then the command executed is `'cd subdir; /bin/make'`. If you use a special version of `make` to run the top-level makefile, the same special version will be executed for recursive invocations. As a special feature, using the variable, `MAKE`, in the commands of a rule alters the effects of the `'-t'` (`'--touch'`), `'-n'` (`'--just-print'`), or `'-q'` (`'--question'`) options. Using the `MAKE` variable has the same effect as using a `'+'` character at the beginning of the command line. See “Instead of executing the commands” on page 411.

Consider the command `'make -t'` for example. The `'-t'` option marks targets as up to date without actually running any commands; see “Instead of executing the commands” on page 411. Following the usual definition of `'-t'`, a `'make -t'` command would create a file named `'subsystem'` and do nothing else. What you really want it to do is run `'cd subdir; make -t'` although that would require executing the command, and `'-t'` says not to execute commands.

The special feature makes this do what you want: whenever a command line of a rule contains the variable, `MAKE`, the `'-t'`, `'-n'` and `'-q'` flags do not apply to that line. Command lines containing `MAKE` are executed normally despite the presence of a flag that causes most commands not to be run. The usual `MAKEFLAGS` mechanism passes the flags to the sub-`make` (see “Communicating options to a sub-make utility” on page 361), so your request to touch the files, or print the commands, is propagated to the subsystem.

Communicating variables to a sub-make utility

Variable values of the top-level `make` can be passed to the sub-`make` through the environment by explicit request. These variables are defined in the sub-`make` as defaults, but do not override what is specified in the makefile used by the sub-`make` makefile unless you use the `'-e'` switch (see “Summary of make options” on page 417).

To pass down, or *export*, a variable, `make` adds the variable and its value to the environment for running each command. The sub-`make`, in turn, uses the environment to initialize its table of variable values. See “Variables from the environment” on page 383.

Except by explicit request, `make` exports a variable only if it is either defined in the environment initially or set on the command line, and if its name consists only of letters, numbers, and underscores.

Some shells cannot cope with environment variable names consisting of characters other than letters, numbers, and underscores. The special variables, `SHELL` and `MAKEFLAGS`, are always exported (unless you `unexport` them). `MAKEFILES` is exported if you set it to anything.

`make` automatically passes down variable values that were defined on the command line, by putting them in the `MAKEFLAGS` variable. See “Communicating options to a sub-make utility” on page 361.

Variables are *not* normally passed down if they were created by default by `make` (see “Variables used by implicit rules” on page 432). The sub-`make` will define these for itself.

If you want to export specific variables to a sub-`make`, use the `export` directive like:
`export variable`

If you want to prevent a variable from being exported, use the `unexport` directive, like:
`unexport variable`

As a convenience, you can define a variable and export it at the same time by using `export variable = value` (which has the same result as `variable = value export variable`) and `export variable := value` (which has the same result as `variable := value export variable`).

Likewise, `export variable += value` is just like `variable += value export variable`.

See “Appending more text to variables” on page 379. You may notice that the `export` and `unexport` directives work in `make` in the same way they work in the shell, `sh`.

If you want all variables to be exported by default, you can use `export` by itself: `export`. This tells `make` that variables which are not explicitly mentioned in an `export` or `unexport` directive should be exported. Any variable given in an `unexport` directive will still not be exported. If you use `export` by itself to export variables by default, variables whose names contain characters other than alphanumeric characters and underscores will not be exported unless specifically mentioned in an `export` directive.

The behavior elicited by an `export` directive by itself was the default in older versions of GNU `make`. If your makefiles depend on this behavior and you want to be compatible with old versions of `make`, you can write a rule for the special target, `.EXPORT_ALL_VARIABLES`, instead of using the `export` directive. This will be ignored by old makes, while the `export` directive will cause a syntax error.

Likewise, you can use `unexport` by itself to tell `make` not to export variables by default. Since this is the default behavior, you would only need to do this if `export` had been used by itself earlier (in an included makefile, perhaps). You *cannot* use `export` and `unexport` by themselves to have variables exported for some commands and not for others. The last `export` or `unexport` directive that appears by itself determines the behavior for the entire run of `make`.

As a special feature, the variable, `MAKELEVEL`, is changed when it is passed down from level to level. This variable’s value is a string which is the depth of the level as a decimal number. The value is ‘0’ for the top-level `make`; ‘1’ for a sub-`make`; ‘2’ for a sub-sub-`make`, and so on. The incrementation happens when `make` sets up the environment for a command.

The main use of `MAKELEVEL` is to test it in a conditional directive (see “Conditional parts of makefiles” on page 385); this way you can write a makefile that behaves one way if run recursively and another way if run directly by you.

You can use the variable, `MAKEFILES`, to cause all sub-`make` commands to use additional makefiles. The value of `MAKEFILES` is a whitespace-separated list of file names. This variable, if defined in the outer-level makefile, is passed down through the environment; then it serves as a list of extra makefiles for the sub-`make` to read before the usual or specified ones. See “The `MAKEFILES` variable” on page 318.

Communicating options to a sub-make utility

Flags such as `-s` and `-k` are passed automatically to the sub-make through the variable, `MAKEFLAGS`. This variable is set up automatically by `make` to contain the flag letters that `make` received. Thus, if you do `'make -ks'` then `MAKEFLAGS` gets the value `'ks'`.

As a consequence, every sub-make gets a value for `MAKEFLAGS` in its environment. In response, it takes the flags from that value and processes them as if they had been given as arguments. See “Summary of make options” on page 417.

Likewise variables defined on the command line are passed to the sub-make through `MAKEFLAGS`. Words in the value of `MAKEFLAGS` that contain `'='`, `make` treats as variable definitions just as if they appeared on the command line. See “Overriding variables” on page 414.

The options `-C`, `-f`, `-o`, and `-w` are not put into `MAKEFLAGS`; these options are not passed down.

The `-j` option is a special case (see “Parallel execution” on page 353). If you set it to some numeric value, `-j 1` is always put into `MAKEFLAGS` instead of the value you specified. This is because if the `-j` option was passed down to sub-makes, you would get many more jobs running in parallel than you asked for. If you give `-j` with no numeric argument, meaning to run as many jobs as possible in parallel, this is passed down, since multiple infinities are no more than one. If you do not want to pass the other flags down, you must change the value of `MAKEFLAGS`, like the following example shows.

```
MAKEFLAGS=
subsystem:
    cd subdir; $(MAKE)
```

Alternately, use the following example’s input.

```
subsystem:
    cd subdir; $(MAKE) MAKEFLAGS=
```

The command line variable definitions really appear in the variable, `MAKEOVERRIDES`, and `MAKEFLAGS` contains a reference to this variable. If you do want to pass flags down normally, but don’t want to pass down the command line variable definitions, you can reset `MAKEOVERRIDES` to empty, like `MAKEOVERRIDES =`.

This is not usually useful to do. However, some systems have a small fixed limit on the size of the environment, and putting so much information in into the value of `MAKEFLAGS` can exceed it. If you see the error message `'Arg list too long'`, this may be the problem. (For strict compliance with POSIX.2, changing `MAKEOVERRIDES` does not affect `MAKEFLAGS` if the special target `_.POSIX` appears in the makefile. You probably do not care about this.)

A similar variable `MFLAGS` exists also, for historical compatibility. It has the same value as `MAKEFLAGS` except that it does not contain the command line variable definitions, and it always begins with a hyphen unless it is empty (`MAKEFLAGS` begins with a hyphen only when it begins with an option that has no single-letter version, such as `--warn-undefined-variables`). `MFLAGS` was traditionally used explicitly in the recursive `make` command, like the following.

```
subsystem:
    cd subdir; $(MAKE) $(MFLAGS)
```

Now `MAKEFLAGS` makes this usage redundant. If you want your makefiles to be compatible with old `make` programs, use this technique; it will work fine with more modern `make` versions too.

The `MAKEFLAGS` variable can also be useful if you want to have certain options, such as `-k` (see “Summary of make options” on page 417), set each time you run `make`. You simply put a value for `MAKEFLAGS` in your environment. You can also set `MAKEFLAGS` in a makefile, to specify additional flags that should also be in effect for that makefile.

NOTE: You cannot use `MFLAGS` this way. That variable is set only for compatibility; `make` does not interpret a value you set for it in any way.)

When `make` interprets the value of `MAKEFLAGS` (either from the environment or from a makefile), it first prepends a hyphen if the value does not already begin with one. Then it chops the value into words separated by blanks, and parses these words as if they were options given on the command line (except that `-C`, `-f`, `-h`, `-o`, `-w`, and their long-named versions are ignored; and there is no error for an invalid option).

If you do put `MAKEFLAGS` in your environment, you should be sure not to include any options that will drastically affect the actions of `make` and undermine the purpose of makefiles and of `make` itself. For instance, the `-t`, `-n`, and `-q` options, if put in one of these variables, could have disastrous consequences and would certainly have at least surprising and probably annoying effects.

The `--print-directory` option

If you use several levels of recursive `make` invocations, the options, `-w` or `--print-directory` can make the output a lot easier to understand by showing each directory as `make` starts processing it and as `make` finishes processing it. For example, if `make -w` is run in the directory `/u/gnu/make`, `make` will print a line like the following before doing anything else.

```
make: Entering directory `/u/gnu/make'.
```

Then, a line of the following form when processing is completed.

```
make: Leaving directory `/u/gnu/make'.
```

Normally, you do not need to specify this option because `make` does it for you: `-w` is turned on automatically when you use the `-C` option, and in sub-makes. `make` will not automatically turn on `-w` if you also use `-s`, which says to be silent, or if you use `--no-print-directory` to explicitly disable it.

Defining canned command sequences

When the same sequence of commands is useful in making various targets, you can define it as a canned sequence with the `define` directive, and refer to the canned sequence from the rules for those targets. The canned sequence is actually a variable, so the name must not conflict with other variable names.

The following is an example of defining a canned sequence of commands.

```
define run-yacc
yacc $(firstword $^ )
mv y.tab.c $@
endif
```

`run-yacc` is the name of the variable being defined; `endif` marks the end of the definition; the lines in between are the commands. The `define` directive does not expand variable references and function calls in the canned sequence; the ‘\$’ characters, parentheses, variable names, and so on, all become part of the value of the variable you are defining. See “Defining variables verbatim” on page 382 for a complete explanation of `define`.

The first command in this example runs Yacc on the first dependency of whichever rule uses the canned sequence. The output file from Yacc is always named ‘`y.tab.c`’. The second command moves the output to the rule’s target file name.

To use the canned sequence, substitute the variable into the commands of a rule. You can substitute it like any other variable (see “Basics of variable references” on page 369). Because variables defined by `define` are recursively expanded variables, all the variable references you wrote inside the `define` are expanded now. Use the following for example.

```
foo.c : foo.y
      $(run-yacc)
```

‘`foo.y`’ will be substituted for the variable ‘`$^`’ when it occurs in `run-yacc`’s value, and ‘`foo.c`’ for ‘`$@`’. This is a realistic example, but this particular one is not needed in practice because `make` has an implicit rule to figure out these commands based on the file names involved (see “Using implicit rules” on page 425).

In command execution, each line of a canned sequence is treated just as if the line appeared on its own in the rule, preceded by a tab. In particular, `make` invokes a separate subshell for each line.

You can use the special prefix characters that affect command lines (‘@’, ‘-’, and ‘+’) on each line of a canned sequence. See “Summary of make options” on page 417.

For example, use the following example of a canned sequence.

```
define frobnicate
```



```
@echo "frobnicating target $@"  
frob-step-1 $< -o $@-step-1  
frob-step-2 $@-step-1 -o $@  
endif
```

make will not echo the first line, the `echo` command. But it will echo the following two command lines. On the other hand, prefix characters on the command line that refers to a canned sequence apply to every line in the sequence. So the following rule statement does not echo any commands. (See “Command Echoing” on page Command echoing for a full explanation of ‘@’.)

```
frob.out: frob.in  
@$(frobnicate)
```

Using empty commands

It is sometimes useful to define commands which do nothing. This is done simply by giving a command that consists of nothing but whitespace. For example, `target: ;` defines an empty command string for ‘`target`’. You could also use a line beginning with a tab character to define an empty command string, but this would be confusing because such a line looks empty.

You may be wondering why you would want to define a command string that does nothing. The only reason this is useful is to prevent a target from getting implicit commands (from implicit rules or the `.DEFAULT` special target; see “Implicit rules” on page 423 and “Defining last-resort default rules” on page 444).

You may be inclined to define empty command strings for targets that are not actual files, but only exist so that their dependencies can be remade. However, this is not the best way to do that, because the dependencies may not be remade properly if the target file actually does exist. See “Phony targets” on page 334 for a better way to execute this requirement.

6

How to use variables

A *variable* is a name defined in a makefile to represent a string of text, called the variable's *value*. The following documentation discusses using variables.

- “Basics of variable references” on page 369
- “The two flavors of variables” on page 370
- “Advanced features for reference to variables” on page 373
- “How variables get their values” on page 377
- “Setting variables” on page 378
- “Appending more text to variables” on page 379
- “The override directive” on page 381
- “Defining variables verbatim” on page 382
- “Variables from the environment” on page 383

Values are substituted by explicit request into targets, dependencies, commands, and other parts of the makefile. In some other versions of `make`, variables are called *macros*.

Variables and functions in all parts of a makefile are expanded when read, except for the shell commands in rules, the right-hand sides of variable definitions using ‘=’, and the bodies of variable definitions using the `define` directive.

Variables can represent lists of file names, options to pass to compilers, programs to run, directories to look in for source files, directories to write output in, or anything else you can imagine.

A variable name may be any sequence of characters not containing ‘:’, ‘#’, ‘=’, or leading or trailing whitespace. However, variable names containing characters other than letters, numbers, and underscores should be avoided because they may be given special meanings in the future, and with some shells they cannot be passed through the environment to a sub-**make** (see “Communicating variables to a sub-make utility” on page 359).

Variable names are case-sensitive. The names ‘foo’, ‘FOO’, and ‘Foo’ all refer to different variables.

It is traditional to use uppercase letters in variable names, but we recommend using lowercase letters for variable names that serve internal purposes in the makefile, and reserving uppercase for parameters that control implicit rules or for parameters that the user should override with command options (see “Overriding variables” on page 414).

A few variables have names that are a single punctuation character or just a few characters. These are the automatic variables, and they have particular specialized uses. See “Automatic variables” on page 438.

Basics of variable references

To substitute a variable's value, write a dollar sign followed by the name of the variable in parentheses or braces: either `$(foo)` or `${foo}` is a valid reference to the variable `foo`. This special significance of `$` is why you must write `$$` to have the effect of a single dollar sign in a file name or command. Variable references can be used in any context: targets, dependencies, commands, most directives, and new variable values. The following is an example of a common case, where a variable holds the names of all the object files in a program.

```
objects = program.o foo.o utils.o
program : $(objects)
        cc -o program $(objects)
```

```
$(objects) : defs.h
```

Variable references work by strict textual substitution. Thus, the following rule could be used to compile a C program `prog.c`.

```
foo = c
prog.o : prog.$(foo)
        $(foo)$(foo) -$(foo) prog.$(foo)
```

Since spaces before the variable value are ignored in variable assignments, the value of `foo` is precisely `c`. (Don't actually write your makefiles this way!) A dollar sign followed by a character other than a dollar sign, open-parenthesis or open-brace treats that single character as the variable name. Thus, you could reference the variable `x` with `$x`. However, this practice is strongly discouraged, except in the case of the automatic variables (see "Automatic variables" on page 438).

The two flavors of variables

There are two ways that a variable in GNU make can have a value; we call them the two *flavors* of variables. The two flavors are distinguished in how they are defined and in what they do when expanded.

The first flavor of variable is a *recursively expanded* variable. Variables of this sort are defined by lines using ‘=’ (see “Setting variables” on page 378) or by the define directive (see “Defining variables verbatim” on page 382). The value you specify is installed verbatim; if it contains references to other variables, these references are expanded whenever this variable is substituted (in the course of expanding some other string). When this happens, it is called *recursive expansion*. Consider the following example.

```
foo = $(bar)
bar = $(ugh)
ugh = Huh?

all:;echo $(foo)
```

This input will echo ‘Huh?’; ‘\$(foo)’ expands to ‘\$(bar)’ which expands to ‘\$(ugh)’ which finally expands to ‘Huh?’.

This flavor of variable is the only sort supported by other versions of **make**. It has its advantages and its disadvantages. An advantage (most would say) is that the following statement will do what was intended: when ‘CFLAGS’ is expanded in a command, it will expand to ‘-Ifoo -Ibar -O’.

```
CFLAGS = $(include_dirs) -O
include_dirs = -Ifoo -Ibar
```

A major disadvantage is that you cannot append something on the end of a variable, as in `CFLAGS = $(CFLAGS) -O`, because it will cause an infinite loop in the variable expansion. (Actually **make** detects the infinite loop and reports an error.)

Another disadvantage is that any functions referenced in the definition will be executed every time the variable is expanded (see “Functions for transforming text” on page 393).

This makes **make** run slower; worse, it causes the wildcard and shell functions to give unpredictable results because you cannot easily control when they are called, or even how many times.

To avoid all the problems and inconveniences of recursively expanded variables, there is another flavor: *simply expanded variables*.

Simply expanded variables are defined by lines using ‘:=’ (see “Setting variables” on page 378). The value of a simply expanded variable is scanned once and for all, expanding any references to other variables and functions, when the variable is defined.

The actual value of the simply expanded variable is the result of expanding the text that you write. It does not contain any references to other variables; it contains their values *as of the time this variable was defined*.

Therefore, consider the following statement.

```
x :=foo
y := $(x) bar
x := later
```

This is equivalent to the next statement.

```
y := foo bar
x := later
```

When a simply expanded variable is referenced, its value is substituted verbatim.

The following is a somewhat more complicated example, illustrating the use of ‘:=’ in conjunction with the shell function. See “The shell function” on page 405. This example also shows use of the variable, `MAKELEVEL`, which is changed when it is passed down from level to level. See “Communicating variables to a sub-make utility” on page 359 for information about `MAKELEVEL`.)

```
ifeq (0,${MAKELEVEL})
cur-dir := $(shell pwd)
whoami := $(shell whoami)
host-type := $(shell arch)
MAKE := ${MAKE} host-type=${host-type} whoami=${whoami}
endif
```

An advantage of this use of ‘:=’ is that a typical ‘descend into a directory’ command then looks like this:

```
${subdirs}:
    ${MAKE} cur-dir=${cur-dir}/$@ -C $@ all
```

Simply expanded variables generally make complicated makefile programming more predictable because they work like variables in most programming languages. They allow you to redefine a variable using its own value (or its value processed in some way by one of the expansion functions) and to use the expansion functions much more efficiently (see “Functions for transforming text” on page 393).

You can also use them to introduce controlled leading whitespace into variable values. Leading whitespace characters are discarded from your input before substitution of variable references and function calls; this means you can include leading spaces in a variable value by protecting them with variable references like the following example’s input.

```
nullstring :=  
space := $(nullstring) # end of the line
```

With this statement, the value of the variable `space` is precisely one space.

The comment ‘# end of the line’ is included here just for clarity. Since trailing space characters are not stripped from variable values, just a space at the end of the line would have the same effect (but be rather hard to read). If you put whitespace at the end of a variable value, it is a good idea to put a comment like that at the end of the line to make your intent clear.

Conversely, if you do *not* want any whitespace characters at the end of your variable value, you must remember *not* to put a random comment on the end of the line after some whitespace, such as the following.

```
dir := /foo/bar          # directory to put the frobs in
```

With this statement, the value of the variable, `dir`, is ‘/foo/bar ’ (with four trailing spaces), which was probably not the intention. (Imagine something like ‘\$(dir)/file’ with this definition!)

Advanced features for reference to variables

The following discussion describes some advanced features that you can use to reference variables in more flexible ways.

Substitution references

A substitution reference substitutes the value of a variable with alterations that you specify. It has the form `$(var: a= b)` or `{var:a=b}`, and its meaning is to take the value of the variable, `var`, replace every `a` at the end of a word with `b` in that value, and substitute the resulting string. When we say “at the end of a word”, we mean that `a` must appear either followed by whitespace or at the end of the value in order to be replaced; other occurrences of `a` in the value are unaltered. Consider the following statement.

```
foo := a.o b.o c.o
bar := $(foo:.o=.c)
```

This input sets `'bar'` to `'a.c b.c c.c'`. See “Setting variables” on page 378. A substitution reference is actually an abbreviation for use of the `patsubst` expansion function; see “Functions for string substitution and analysis” on page 395. We provide substitution references as well as `patsubst` for compatibility with other implementations of `make`.

Another type of substitution reference lets you use the full power of the `patsubst` function. It has the same form `$(var:a=b)` described above, except that now `a` must contain a single `'%'` character. This case is equivalent to `$(patsubst a, b, $(var))`. See “Functions for string substitution and analysis” on page 395 for a description of the `patsubst` function. Consider the following example.

```
foo := a.o b.o c.o
bar := $(foo:%.o=%.c)
```

This statement sets `'bar'` to `'a.c b.c c.c'`.

Computed variable names

Computed variable names are a complicated concept needed only for sophisticated makefile programming. For most purposes you need not consider them, except to know that making a variable with a dollar sign in its name might have strange results. However, if you are the type that wants to understand everything, or you are actually interested in what they do, read on.

Variables may be referenced *inside* the name of a variable. This is called a *computed variable name* or a *nested variable reference*. Consider the next example.

```
x =y
y =z
a := $( $(x) )
```

The previous example defines `a` as `'z'`; the `'$(x)'` inside `'$($(x))'` expands to `'y'`, so `'$($(x))'` expands to `'$(y)'` which in turn expands to `'z'`. Here the name of the variable to reference is not stated explicitly; it is computed by expansion of `'$(x)'`. The reference, `'$(x)'`, is nested within the outer variable reference.

The previous example shows two levels of nesting; however, any number of levels is possible. For instance, the following example shows three levels.

```
x =y
y =z
z =u
a := $( $( $(x) ) )
```

The previous example shows the innermost `'$(x)'` expands to `'y'`, so `'$($(x))'` expands to `'$(y)'` which in turn expands to `'z'`; now we have `'$(z)'`, which becomes `'u'`.

References to recursively-expanded variables within a variable name are reexpanded in the usual fashion. Consider the following example.

```
x = $(y)
y =z
z = Hello
a := $( $(x) )
```

The previous example shows `a` defined as `'Hello'`; `'$($(x))'` becomes `'$($(y))'` which becomes `'$(z)'` which becomes `'Hello'`.

Nested variable references can also contain modified references and function invocations (see “Functions for transforming text” on page 393), just like any other reference. For instance, the following example uses the `subst` function (see “Functions for string substitution and analysis” on page 395):

```
x = variable1
variable2 := Hello
y = $(subst 1,2,$(x))
z =y a := $( $( $(z) ) )
```

The previous example eventually defines `a` as `'Hello'`. It is doubtful that anyone would ever want to write a nested reference as convoluted as this one, but it works. `'$($($(z)))'` expands to `'$($(y))'` which becomes `'$($(subst 1,2,$(x)))'`. This gets the value `'variable1'` from `x` and changes it by substitution to `'variable2'`, so that the entire string becomes `'$(variable2)'`, a simple variable reference whose value is `'Hello'`.

A computed variable name need not consist entirely of a single variable reference. It can contain several variable references as well as some invariant text. Consider the following example.

```
a_dirs := dira dirb
l_dirs := dir1 dir2

a_files := filea fileb
l_files := file1 file2

ifeq "$(use_a)" "yes"
a1 := a
else
a1 := 1
endif

ifeq "$(use_dirs)" "yes"
df := dirs
else
df := files
endif

dirs := $($a1)_$(df))
```

The previous example will give `dirs` the same value as `a_dirs`, `l_dirs`, `a_files` or `l_files` depending on the settings of `use_a` and `use_dirs`.

Computed variable names can also be used in substitution references:

```
a_objects := a.o b.o c.o
l_objects := 1.o 2.o 3.o

sources := $($a1)_objects:.o=.c)
```

The previous example defines `sources` as `'a.c b.c c.c'` or `'1.c 2.c 3.c'`, depending on the value of `a1`.

The only restriction on this sort of use of nested variable references is that they cannot specify part of the name of a function to be called. This is because the test for a recognized function name is done before the expansion of nested references as in the following example.

```
ifdef do_sort
func := sort
else
func := strip
endif

bar:=a d bg q c

foo := $($func) $(bar))
```

The previous example attempts to give `foo` the value of the variable `sort a d b g q c` or `strip a db gq c`, rather than giving `ad b gq c` as the argument to either the `sort` or the `strip` function. This restriction could be removed in the future if that change is shown to be a good idea.

You can also use computed variable names in the left-hand side of a variable assignment, or in a `define` directive as in the following example.

```
dir = foo
$(dir)_sources := $(wildcard $(dir)/*.c)
define $(dir)_print
lpr $($dir)_sources
endef
```

This example defines the variables `dir`, `foo_sources`, and `foo_print`.

NOTE: *Nested variable references* are quite different from *recursively expanded variables* (see “The two flavors of variables” on page 370), though both are used together in complex ways when doing makefile programming.

How variables get their values

Variables can get values in several different ways:

- You can specify an overriding value when you run `make`. See “Overriding variables” on page 414.
- You can specify a value in the makefile, either with an assignment (see “Setting variables” on page 378) or with a verbatim definition (see “Defining variables verbatim” on page 382).
- Variables in the environment become `make` variables. See “Variables from the environment” on page 383.
- Several automatic variables are given new values for each rule. Each of these has a single conventional use. See “Automatic variables” on page 438.
- Several variables have constant initial values. See “Variables used by implicit rules” on page 432.

Setting variables

To set a variable from the makefile, write a line starting with the variable name followed by `=` or `:=`. Whatever follows the `=` or `:=` on the line becomes the value. For example, `objects = main.o foo.o bar.o utils.o` defines a variable named `objects`. Whitespace around the variable name and immediately after the `=` is ignored.

Variables defined with `=` are *recursively expanded variables*. Variables defined with `:=` are *simply expanded variables*; these definitions can contain variable references which will be expanded before the definition is made. See “The two flavors of variables” on page 370.

The variable name may contain function and variable references, which are expanded when the line is read to find the actual variable name to use. There is no limit on the length of the value of a variable except the amount of swapping space on the computer.

When a variable definition is long, it is a good idea to break it into several lines by inserting backslash-newline at convenient places in the definition. This will not affect the functioning of `make`, but it will make the makefile easier to read.

Most variable names are considered to have the empty string as a value if you have never set them. Several variables have built-in initial values that are not empty, but you can set them in the usual ways (see “Variables used by implicit rules” on page 432). Several special variables are set automatically to a new value for each rule; these are called the *automatic variables* (see “Automatic variables” on page 438).

Appending more text to variables

Often it is useful to add more text to the value of a variable already defined. You do this with a line containing `+=`, as in `objects += another.o`.

This takes the value of the variable `objects`, and adds the text `'another.o'` to it (preceded by a single space), as in the following example.

```
objects = main.o foo.o bar.o utils.o
objects += another.o
```

The last line sets `objects` to `'main.o foo.o bar.o.o an utils other.o'`.

Using `+=` is similar to the following example.

```
objects = main.o foo.o bar.o utils.o
objects := $(objects) another.o
```

This last example's statement differs in ways that become important when you use more complex values. When the variable in question has not been defined before, `+=` acts just like normal `:=`: it defines a recursively-expanded variable. However, when there is a previous definition, exactly what `+=` does depends on what flavor of variable you defined originally. See "The two flavors of variables" on page 370 for an explanation of the two flavors of variables.

When you add to a variable's value with `+=`, **make** acts essentially as if you had included the extra text in the initial definition of the variable. If you defined it first with `:=`, making it a simply-expanded variable, `+=` adds to that simply-expanded definition, and expands the new text before appending it to the old value just as `:=` does (see "Setting variables" on page 378 for a full explanation of `:=`). Consider the following definition.

```
variable := value
variable += more
```

This last statement is exactly equivalent to this next definition.

```
variable := value
variable := $(variable) more
```

On the other hand, when you use `+=` with a variable that you defined first to be recursively-expanded using plain `=`, **make** does something a bit different. Recall that when you define a recursively-expanded variable, **make** does not expand the value you set for variable and function references immediately. Instead it stores the text verbatim, and saves these variable and function references to be expanded later, when you refer to the new variable (see "The two flavors of variables" on page 370). When you use `+=` on a recursively-expanded variable, it is this unexpanded text to which **make** appends the new text you specify.

```
variable = value
variable += more
```

This last statement is roughly equivalent to this next definition.

```
temp = value
variable = $(temp) more
```

Of course it never defines a variable called `temp`. The importance of this comes when the variable's old value contains variable references. Take this common example.

```
CFLAGS = $(includes) -O
```

```
...
CFLAGS += -pg # enable profiling
```

The first line defines the `CFLAGS` variable with a reference to another variable, `includes`. (`CFLAGS` is used by the rules for C compilation; see “Catalogue of implicit rules” on page 427.) Using `=` for the definition makes `CFLAGS` a recursively-expanded variable, meaning `$(includes) -O` is not expanded when `make` processes the definition of `CFLAGS`. Thus, `includes` need not be defined yet for its value to take effect. It only has to be defined before any reference to `CFLAGS`. If we tried to append to the value of `CFLAGS` without using `+=`, we might do it with this following definition.

```
CFLAGS := $(CFLAGS) -pg # enable profiling
```

This is pretty close, but not quite what we want. Using `:=` redefines `CFLAGS` as a simply-expanded variable; this means `make` expands the text, `$(CFLAGS) -pg` before setting the variable. If `includes` is not yet defined, we get `-O -pg`, and a later definition of `includes` will have no effect. Conversely, by using `+=` we set `CFLAGS` to the unexpanded value `$(includes) -O -pg`. Thus we preserve the reference to `includes` so that, if that variable gets defined at any later point, a reference like `$(CFLAGS)` still uses its value.

The override directive

If a variable has been set with a command argument (see “Overriding variables” on page 414), then ordinary assignments in the makefile are ignored. If you want to set the variable in the makefile even though it was set with a command argument, you can use an `override` directive which is a line looking like `override variable = value`, or `override variable := value`.

To append more text to a variable defined on the command line, use the statement, `override variable += more text`.

See “Appending more text to variables” on page 379. The `override` directive was not invented for escalation in the war between makefiles and command arguments. It was invented so you can alter and add to values that the user specifies with command arguments.

For example, suppose you always want the ‘-g’ switch when you run the C compiler, but you would like to allow the user to specify the other switches with a command argument just as usual. You could use this `override` directive: `override CFLAGS += -g`.

You can also use `override` directives with `define` directives. This is done as you might expect, as in the following.

```
override define foo
bar
endef
```

See “Defining variables verbatim” on page 382 for information about `define`.

Defining variables verbatim

Another way to set the value of a variable is to use the `define` directive. This directive has an unusual syntax which allows newline characters to be included in the value, which is convenient for defining canned sequences of commands (see “Defining canned command sequences” on page 364).

The `define` directive is followed on the same line by the name of the variable and nothing more. The value to give the variable appears on the following lines. The end of the value is marked by a line containing just the word, `endef`. Aside from this difference in syntax, `define` works just like `=`; it creates a recursively-expanded variable (see “The two flavors of variables” on page 370). The variable name may contain function and variable references which are expanded when the directive is read to find the actual variable name to use.

```
define two-lines
echo foo
echo $(bar)
endef
```

The value in an ordinary assignment cannot contain a newline; the newlines that separate the lines of the value in a `define` become part of the variable’s value (except for the final newline which precedes the `endef` and is not considered part of the value). The previous example is functionally equivalent to `two-lines = echo foo; echo $(bar)` since two commands separated by semicolon behave much like two separate shell commands.

However, using two separate lines means `make` will invoke the shell twice, running an independent sub-shell for each line. See “Command execution” on page 352.

If you want variable definitions made with `define` to take precedence over command-line variable definitions, you can use the `override` directive together with `define` as in the following example.

```
override define two-lines
foo
$(bar)
endef
```

See “The override directive” on page 381.

Variables from the environment

Variables in `make` can come from the environment in which `make` is run. Every environment variable that `make` sees when it starts up is transformed into a `make` variable with the same name and value. But an explicit assignment in the makefile, or with a command argument, overrides the environment. If the `-e` flag is specified, then values from the environment override assignments in the makefile (see “Summary of make options” on page 417), although this is not recommended practice.)

Thus, by setting the `CFLAGS` variable in your environment, you can cause all C compilations in most makefiles to use the compiler switches you prefer. This is safe for variables with standard or conventional meanings because you know that no makefile will use them for other things. However, this is not totally reliable; some makefiles set `CFLAGS` explicitly and therefore are not affected by the value in the environment.

When `make` is invoked recursively, variables defined in the outer invocation can be passed to inner invocations through the environment (see “Recursive use of the make utility” on page 358). By default, only variables that came from the environment or the command line are passed to recursive invocations. You can use the `export` directive to pass other variables. See “Communicating variables to a sub-make utility” on page 359 for full details.

Other use of variables from the environment is not recommended. It is not wise for makefiles to depend for their functioning on environment variables set up outside their control, since this would cause different users to get different results from the same makefile. This is against the whole purpose of most makefiles.

Such problems would be especially likely with the variable, `SHELL`, which is normally present in the environment to specify the user’s choice of interactive shell. It would be very undesirable for this choice to affect `make`. So `make` ignores the environment value of `SHELL`.

7

Conditional parts of makefiles

A *conditional* causes part of a makefile to be obeyed or ignored depending on the values of variables. Conditionals can compare the value of one variable to another, or the value of a variable to a constant string. Conditionals control what **make** actually “sees” in the makefile, so they cannot be used to control shell commands at the time of execution. The following documentation describes conditionals in more detail.

- “Example of a conditional” on page 386
- “Syntax of conditionals” on page 388
- “Conditionals that test flags” on page 391

Example of a conditional

The following example of a conditional tells **make** to use one set of libraries if the `cc` variable is `'gcc'`, and a different set of libraries otherwise. It works by controlling which of two command lines will be used as the command for a rule. The result is that `'CC=gcc'` as an argument to **make** changes not only which compiler is used but also which libraries are linked.

```
libs_for_gcc = -lgnu
normal_libs =

foo: $(objects)
ifeq ($(CC),gcc)
    $(CC) -o foo $(objects) $(libs_for_gcc)
else
    $(CC) -o foo $(objects) $(normal_libs)
endif
```

This conditional uses three directives: one `ifeq`, one `else` and one `endif`.

The `ifeq` directive begins the conditional, specifying the condition. It contains two arguments, separated by a comma and surrounded by parentheses. Variable substitution is performed on both arguments and then they are compared. The lines of the makefile following the `ifeq` are obeyed if the two arguments match; otherwise they are ignored.

The `else` directive causes the following lines to be obeyed if the previous conditional failed. In the example above, this means that the second alternative linking command is used whenever the first alternative is not used. It is optional to have an `else` in a conditional.

The `endif` directive ends the conditional. Every conditional must end with an `endif`. Unconditional makefile text follows.

As the following example illustrates, conditionals work at the textual level; the lines of the conditional are treated as part of the makefile, or ignored, according to the condition. This is why the larger syntactic units of the makefile, such as rules, may cross the beginning or the end of the conditional.

When the variable, `cc`, has the value `'gcc'`, the previous example has this effect.

```
foo: $(objects)
    $(CC) -o foo $(objects) $(libs_for_gcc)
```

When the variable, `CC`, has any other value, it takes the following effect.

```
foo: $(objects)
    $(CC) -o foo $(objects) $(normal_libs)
```

Equivalent results can be obtained in another way by conditionalizing a variable assignment and then using the variable unconditionally as in the following example.

```
libs_for_gcc = -lgnu
normal_libs =

ifeq ($(CC),gcc)
    libs=$(libs_for_gcc)
else
    libs=$(normal_libs)
endif

foo: $(objects)
    $(CC) -o foo $(objects) $(libs)
```

Syntax of conditionals

The syntax of a simple conditional with no `else` is as follows.

```
conditional-directive
text-if-true
endif
```

The `text-if-true` may be any lines of text, to be considered as part of the makefile if the condition is true. If the condition is false, no text is used instead. The syntax of a complex conditional is as follows.

```
conditional-directive
text-if-true
else
text-if-false
endif
```

If the condition is true, `text-if-true` is used; otherwise, `text-if-false` is used instead.

The `text-if-false` can be any number of lines of text.

The syntax of the `conditional-directive` is the same whether the conditional is simple or complex. There are four different directives that test different conditions. The following is a description of them.

```
ifeq (arg1, arg2)
ifeq 'arg1' 'arg2'
ifeq "arg1" "arg2"
ifeq "arg1" 'arg2'
ifeq 'arg1' "arg2"
```

Expand all variable references in `arg1` and `arg2` and compare them. If they are identical, the `text-if-true` is effective; otherwise, the `text-if-false`, if any, is effective.

Often you want to test if a variable has a non-empty value.

When the value results from complex expansions of variables and functions, expansions you would consider empty may actually contain whitespace characters and thus are not seen as empty.

However, you can use the `strip` function (see “Functions for string substitution and analysis” on page 395) to avoid interpreting whitespace as a non-empty value.

For example, the following will evaluate `text-if-empty` even if the expansion of `$(foo)` contains whitespace characters.

```
ifeq ($(strip $(foo)),)
text-if-empty
endif
```



```
ifneq ( arg1, arg2)
ifneq ` arg1' `arg2'
ifneq " arg1" "arg2"
ifneq " arg1" `arg2'
ifneq ` arg1' "arg2"
```

Expand all variable references in *arg1* and *arg2* and compare them.

If they differ, the *text-if-true* is effective; otherwise, the *text-if-false*, if any, is effective.

```
ifdef variable-name
```

If the variable *variable-name* has a non-empty value, the *text-if-true* is effective; otherwise, the *text-if-false*, if any, is effective. Variables that have never been defined have an empty value.

NOTE: *ifdef* only tests whether a variable has a value. It does not expand the variable to see if that value is nonempty. Consequently, tests using *ifdef* return true for all definitions except those like `foo =`.

To test for an empty value, use *ifeq* (`$(foo)`), as in the following example.

```
bar =
foo = $(bar)
ifdef foo
frobozz = yes
else
frobozz = no
endif
```

The previous definition example sets ‘frobozz’ to ‘yes’ while the following definition sets ‘frobozz’ to ‘no’.

```
foo =
ifdef foo
frobozz = yes
else
frobozz = no
endif
```

```
ifndef variable-name
```

If the variable *variable-name* has an empty value, the *text-if-true* is effective; otherwise, the *text-if-false*, if any, is effective.

Extra spaces are allowed and ignored at the beginning of the conditional directive line, but a tab is not allowed. (If the line begins with a tab, it will be considered a command for a rule.) Aside from this, extra spaces or tabs may be inserted with no effect anywhere except within the directive name or within an argument. A comment starting with ‘#’ may appear at the end of the line.

The other two directives that play a part in a conditional are `else` and `endif`. Each of these directives is written as one word, with no arguments. Extra spaces are allowed and ignored at the beginning of the line, and spaces or tabs at the end. A comment starting with ‘#’ may appear at the end of the line.

Conditionals affect which lines of the makefile **make** uses. If the condition is true, **make** reads the lines of the *text-if-true* as part of the makefile; if the condition is false, **make** ignores those lines completely. It follows that syntactic units of the makefile, such as rules, may safely be split across the beginning or the end of the conditional.

make evaluates conditionals when it reads a makefile. Consequently, you cannot use automatic variables in the tests of conditionals because they are not defined until commands are run (see “Automatic variables” on page 438). To prevent intolerable confusion, it is not permitted to start a conditional in one makefile and end it in another. However, you may write an `include` directive within a conditional, provided you do not attempt to terminate the conditional inside the included file.

Conditionals that test flags

You can write a conditional that tests **make** command flags such as ‘-t’ by using the variable `MAKEFLAGS` together with the `findstring` function (see “Functions for string substitution and analysis” on page 395). This is useful when `touch` is not enough to make a file appear up to date.

The `findstring` function determines whether one string appears as a substring of another. If you want to test for the ‘-t’ flag, use ‘t’ as the first string and the value of `MAKEFLAGS` as the other.

For example, the following shows how to arrange to use ‘`ranlib -t`’ to finish marking an archive file up to date.

```
archive.a... :
ifneq (, $(findstring t, $(MAKEFLAGS)))
    +touch archive.a
    +ranlib -t archive.a
else
    ranlib archive.a
endif
```

The ‘+’ prefix marks those command lines as “recursive” so that they will be executed despite use of the ‘-t’ flag. See “Recursive use of the make utility” on page 358.

8

Functions for transforming text

Functions allow you to do text processing in the makefile to compute the files to operate on or the commands to use. You use a function in a *function call*, where you give the name of the function and some text (the *arguments*) on which the function operates. The result of the function's processing is substituted into the makefile at the point of the call, just as a variable might be substituted.

The following documentation discusses functions in more detail.

- “Function call syntax” on page 394
- “Functions for string substitution and analysis” on page 395
- “Functions for file names” on page 398
- “The foreach function” on page 401
- “The origin function” on page 403
- “The shell function” on page 405

Function call syntax

A function call resembles a variable reference.

It looks like: `$(function arguments)`; or like: `${function arguments}`.

Here, *function* is a function name; one of a short list of names that are part of **make**. There is no provision for defining new functions.

The *arguments* are the arguments of the function. They are separated from the function name by one or more spaces or tabs, and if there is more than one argument, then they are separated by commas. Such whitespace and commas are not part of an argument's value. The delimiters which you use to surround the function call, whether paren-theses or braces, can appear in an argument only in matching pairs; the other kind of delimiters may appear singly. If the arguments themselves contain other function calls or variable references, it is wisest to use the same kind of delimiters for all the references; write `$(subst a,b,$(x))`, not `$(subst a,b,${x})`. This is because it is clearer, and because only one type of delimiter is matched to find the end of the reference.

The text written for each argument is processed by substitution of variables and function calls to produce the argument value, which is the text on which the function acts. The substitution is done in the order in which the arguments appear.

Commas and unmatched parentheses or braces cannot appear in the text of an argument as written; leading spaces cannot appear in the text of the first argument as written. These characters can be put into the argument value by variable substitution. First define variables `comma` and `space` whose values are isolated comma and space characters; then, substitute these variables where such characters are wanted, like the following.

```
comma:= ,
empty:=
space:= $(empty) $(empty)
foo:= a b c
bar:= $(subst $(space),$(comma),$(foo))
# bar is now 'a,b,c'.
```

Here the `subst` function replaces each space with a comma, through the value of `foo`, and substitutes the result.

Functions for string substitution and analysis

The following functions operate on strings.

`$(subst from, to, text)`

Performs a textual replacement on the text *text*: each occurrence of *from* is replaced by *to*. The result is substituted for the function call.

For example, `$(subst ee,EE,feet on the street)` substitutes the string ‘fEEt on the strEEt’.

`$(patsubst pattern, replacement, text)`

Finds whitespace-separated words in *text* that match *pattern* and replaces them with *replacement*. In this string, *pattern* may contain a ‘%’ which acts as a wildcard, matching any number of any characters within a word. If *replacement* also contains a ‘%’, the ‘%’ is replaced by the text that matched the ‘%’ in *pattern*. ‘%’ characters in `patsubst` function invocations can be quoted with preceding backslashes (‘\’).

Backslashes that would otherwise quote ‘%’ characters can be quoted with more backslashes. Backslashes that quote ‘%’ characters or other backslashes are removed from the pattern before it compares file names or has a stem substituted into it. Backslashes that are not in danger of quoting ‘%’ characters go unmolested.

For example, the pattern ‘the\%weird\%\%pattern\%’ has ‘the%weird\’ preceding the operative ‘%’ character, and ‘pattern\%’ following it. The final two backslashes are left alone because they cannot affect any ‘%’ character.

Whitespace between words is folded into single space characters; leading and trailing whitespace is discarded.

For example, `$(patsubst %.c,%.o,x.c.c bar.c)` produces the value, ‘x.c.o bar.o’. Substitution references are a simpler way to get the effect of the `patsubst` function; see “Substitution references” on page 373.

`$(var: pattern=replacement)`

The previous example of a substitution reference is equivalent to the following example’s input.

`$(patsubst pattern,replacement,$(var))`

The second shorthand simplifies one of the most common uses of `patsubst:`, replacing the suffix at the end of file names.

`$(var:suffix=replacement)`

The previous example’s shorthand is equivalent to the following example’s input.

`$(patsubst %suffix,%replacement,$(var))`

For example, you might have a list of object files: `objects = foo.o bar.o baz.o`

To get the list of corresponding source files, you could simply write:

```
$(objects:.o=.c)
```

instead of using the general form:

```
$(patsubst %.o,%.c,$(objects))
```

`$(strip string)`

Removes leading and trailing whitespace from *string* and replaces each internal sequence of one or more whitespace characters with a single space. Thus, `$(strip a b c)` results in `'a b c'`.

The function, `strip`, can be very useful when used in conjunction with conditionals.

When comparing something with an empty string using `ifeq` or `ifneq`, you usually want a string of just whitespace to match the empty string (see “Conditional parts of makefiles” on page 385).

Thus, the following may fail to have the desired results.

```
.PHONY: all
ifneq "$$(needs_made)" ""
all: $(needs_made)
else
all:;@echo 'Nothing to make!'
endif
```

Replacing the variable reference, `$(needs_made)`, with the function call `$(strip $(needs_made))` in the `ifneq` directive would make it more robust.

`$(findstring find,in)`

Searches *in* for an occurrence of *find*. If it occurs, the value is *find*; otherwise, the value is empty. You can use this function in a conditional to test for the presence of a specific substring in a given string.

Thus, the two following examples produce, respectively, the values `'a'` and an empty string.

```
$(findstring a,a b c)
```

```
$(findstring a,b c)
```

See “Conditionals that test flags” on page 391 for a practical application of `findstring`.

`$(filter pattern ...,text)`

Removes all whitespace-separated words in *text* that *do not* match any of the *pattern* words, returning only words that *do* match. The patterns are written using `'%'`, just like the patterns used in the `patsubst` function.

The `filter` function can be used to separate out different types of strings (such as

file names) in a variable. Consider the following, for example.

```
sources := foo.c bar.c baz.s ugh.h
foo: $(sources)
    cc $(filter %.c %.s,$(sources)) -o foo
```

With this statement, ‘foo’ depends on ‘foo.c’, ‘bar.c’, ‘baz.s’ and ‘ugh.h’ but only ‘foo.c’, ‘bar.c’ and ‘baz.s’ should be specified in the command to the compiler.

```
$(filter-out pattern ...,text)
```

Removes all whitespace-separated words in *text* that *do* match the *pattern* words, returning only the words that *do not* match. This is the exact opposite of the *filter* function. Consider the following for example.

```
objects=main1.o foo.o main2.o bar.o
mains=main1.o main2.o
```

Given the previous lines, the following then generates a list which contains all the object files not in ‘mains’.

```
$(filter-out $(mains),$(objects))
```

```
$(sort list)
```

Sorts the words of *list* in lexical order, removing duplicate words. The output is a list of words separated by single spaces.

Thus, `$(sort foo bar lose)` returns the value, ‘bar foo lose’.

Incidentally, since *sort* removes duplicate words, you can use it for this purpose even if you don’t care about the *sort* order.

The following is a realistic example of the use of *subst* and *patsubst*.

Suppose that a makefile uses the *VPATH* variable to specify a list of directories that make should search for dependency files (see “*VPATH*: search path for all dependencies” on page 330). The following example shows how to tell the C compiler to search for header files in the same list of directories. The value of *VPATH* is a list of directories separated by colons, such as ‘src:../headers’. First, the *subst* function is used to change the colons to spaces:

```
$(subst :, ,$(VPATH))
```

This produces ‘src ../headers’. Then, *patsubst* is used to turn each directory name into a ‘-I’ flag. These can be added to the value of the variable *CFLAGS* which is passed automatically to the C compiler, as in the following.

```
override CFLAGS += $(patsubst %,-I%,$(subst :, ,$(VPATH)))
```

The effect is to append the text, ‘-Isrc -I../headers’, to the previously given value of *CFLAGS*. The *override* directive is used so that the new value is assigned even if the previous value of *CFLAGS* was specified with a command argument (see “The *override* directive” on page 381).

Functions for file names

Several of the built-in expansion functions relate specifically to taking apart file names or lists of file names.

Each of the following functions performs a specific transformation on a file name. The argument of the function is regarded as a series of file names, separated by whitespace. (Leading and trailing whitespace is ignored.) Each file name in the series is transformed in the same way and the results are concatenated with single spaces between them.

`$(dir names...)`

Extracts the directory-part of each file name in *names*. The directory-part of the file name is everything up through (and including) the last slash in it. If the file name contains no slash, the directory part is the string `./`.

For example, `$(dir src/foo.c hacks)` produces the result, `'src/ ./'`.

`$(notdir names ...)`

Extracts all but the directory-part of each file name in *names*. If the file name contains no slash, it is left unchanged. Otherwise, everything through the last slash is removed from it.

A file name that ends with a slash becomes an empty string. This is unfortunate because it means that the result does not always have the same number of whitespace-separated file names as the argument had; but we do not see any other valid alternative.

For example, `$(notdir src/foo.c hacks)` produces the resulting file name, `'foo.c hacks'`.

`$(suffix names ...)`

Extracts the suffix of each file name in *names*. If the file name contains a period, the suffix is everything starting with the last period. Otherwise, the suffix is the empty string. This frequently means that the result will be empty when *names* is not, and if *names* contains multiple file names, the result may contain fewer file names.

For example, `$(suffix src/foo.c hacks)` produces the result, `'.c'`.

`$(basename names...)`

Extracts all but the suffix of each file name in *names*. If the file name contains a period, the basename is everything starting up to (and not including) the last period. Otherwise, the basename is the entire file name. For example, `$(basename src/foo.c hacks)` produces the result, `'src/foo hacks'`.

`$(addsuffix suffix, names...)`

The argument, *names*, is regarded as a series of names, sep-arated by whitespace; *suffix* is used as a unit. The value of *suffix* is appended to the end of each individual name and the resulting larger names are concatenated with single spaces between them.

For example, `$(addsuffix .c,foo bar)` results in `'foo.c bar.c'`.

`$(addprefix prefix, names...)`

The argument, *names*, is regarded as a series of names, separated by whitespace; *prefix* is used as a unit. The value of *prefix* is prepended to the front of each individual name and the resulting larger names are concatenated with single spaces between them.

For example, `$(addprefix src/,foo bar)` results in `'src/foo src/bar'`.

`$(join list1, list2)`

Concatenates the two arguments word by word; the two first words (one from each argument), concatenated, form the first word of the result; the two second words form the second word of the result, and so on. So the *n*th word of the result comes from the *n*th word of each argument. If one argument has more words than the other, the extra words are copied unchanged into the result.

For example, `'$(join a b,.c .o)'` produces `'a.c b.o'`.

Whitespace between the words in the lists is not preserved; it is replaced with a single space. This function can merge the results of the `dir` and `notdir` functions to produce the original list of files which was given to those two functions.

`$(word n, text)`

Returns the *n*th word of *text*. The legitimate values of *n* start from 1. If *n* is bigger than the number of words in *text*, the value is empty.

For example, `$(word 2, foo bar baz)` returns `'bar'`.

`$(words text)`

Returns the number of words in *text*. Thus, the last word of *text* is `$(word $(words text),text)`.

`$(firstword names...)`

The argument, *names*, is regarded as a series of names, separated by whitespace. The value is the first name in the series. The rest of the names are ignored.

For example, `$(firstword foo bar)` produces the result `'foo'`.

Although `$(firstword text)` is the same as `$(word 1, text)`, the `firstword` function is retained for its simplicity.

`$(wildcard pattern)`

The argument *pattern* is a file name pattern, typically containing wildcard characters (as in shell file name patterns). The result of `wildcard` is a space-separated list of the names of existing files that match the pattern. See “Using wildcard characters in file names” on page 326.

The `foreach` function

The `foreach` function is very different from other functions. It causes one piece of text to be used repeatedly, each time with a different substitution performed on it. It resembles the `for` command in the shell `sh` and the `foreach` command in the C-shell, `csh`.

The syntax of the `foreach` function is: `$(foreach var,list,text)`

The first two arguments, `var` and `list`, are expanded before anything else is done; the last argument, `text`, is not expanded at the same time. Then for each word of the expanded value of `list`, the variable named by the expanded value of `var` is set to that word, and `text` is expanded.

Presumably `text` contains references to that variable, so its expansion will be different each time.

The result is that `text` is expanded as many times as there are whitespace-separated words in `list`. The multiple expansions of `text` are concatenated, with spaces between them, to make the result of `foreach`.

The following example sets the variable, 'files', to the list of all files in the directories in the list, 'dirs'.

```
dirs := a b c d
files := $(foreach dir,$(dirs),$(wildcard $(dir)/*))
```

With the previous example, `text` is '`$(wildcard $(dir)/*)`'. The first repetition finds the value 'a' for `dir`, so it produces the same result as '`$(wildcard a/*)`'; the second repetition produces the result of '`$(wildcard b/*)`'; and the third, that of '`$(wildcard c/*)`'. The previous example has the same result (except for setting 'dirs') as `files:= $(wildcard a/* b/* c/* d/*)`.

When `text` is complicated, you can improve readability by giving it a name, with an additional variable, as in the following.

```
find_files = $(wildcard $(dir)/*)
dirs := a b c d
files := $(foreach dir,$(dirs),$(find_files))
```

Here we use the variable, `find_files`, this way. We use plain '=' to define a recursively-expanding variable, so that its value contains an actual function call to be re-expanded under the control of `foreach`; a simply-expanded variable would not do, since `wildcard` would be called only once at the time of defining `find_files`.

The `foreach` function has no permanent effect on the variable, `var`; its value and flavor after the `foreach` function call are the same as they were beforehand. The other values which are taken from `list` are in effect only temporarily, during the execution of `foreach`. The variable, `var`, is a simply-expanded variable during the execution of `foreach`. If `var` was undefined before the `foreach` function call, it is undefined after the call. See “The two flavors of variables” on page 370.

You must take care when using complex variable expressions that result in variable names because many strange things are valid variable names, and are probably not what you intended. Consider the following, for example.

```
files := $(foreach Esta escrito en espanol!,b c ch,$(find_files))
```

This expression might be useful if the value of `find_files` references the variable whose name is ‘`Esta escrito en espanol!`’ (es un nombre bastante largo, no?), but it is more likely to be a mistake.

The origin function

The `origin` function is unlike most other functions in that it does not operate on the values of variables; it tells you something about a variable.

Specifically, it tells you its origin. Its syntax is:

```
$(origin variable)
```

NOTE: *variable* is the name of a variable to inquire about; it is *not* a *reference* to that variable. Therefore you would not normally use a ‘\$’ or parentheses when writing it. (You can, however, use a variable reference in the name if you want the name not to be a constant.)

The result of this function is a string telling you how the variable, *variable*, was defined as the following descriptions discuss.

‘undefined’

Used if *variable* was never defined.

‘default’

Used if *variable* has a default definition as is usual with `CC` and so on. See “Variables used by implicit rules” on page 432.

NOTE: If you have redefined a default variable, the `origin` function will return the origin of the later definition.

‘environment’

Used if *variable* was defined as an environment variable and the ‘-e’ option is not turned on (see “Summary of make options” on page 417).

‘environment override’

Used if *variable* was defined as an environment variable and the ‘-e’ option is turned on (see “Summary of make options” on page 417).

‘file’

Used if *variable* was defined in a makefile.

‘command line’

Used if *variable* was defined on the command line.

‘override’

Used if *variable* was defined with an `override` directive in a makefile (see “The override directive” on page 381).

‘automatic’

Used if *variable* is an automatic variable defined for the execution of the commands for each rule (see “Automatic variables” on page 438).

This information is primarily useful (other than for your curiosity) to determine if you want to believe the value of a variable. For example, suppose you have a makefile, `'foo'`, that includes another makefile, `'bar'`.

You want a variable, `bletch`, to be defined in `'bar'` if you run the command, `'make -f bar'`, even if the environment contains a definition of `bletch`. However, if `'foo'` defined `bletch` before including `'bar'`, you do not want to override that definition. This could be done by using an `override` directive in `'foo'`, giving that definition precedence over the later definition in `'bar'`; unfortunately, the `override` directive would also override any command line definitions. So, `'bar'` could include the following.

```
ifdef bletch
ifeq "$(origin bletch)" "environment"
bletch = barf, gag, etc.
endif
endif
```

If `bletch` has been defined from the environment, this will redefine it. If you want to override a previous definition of `bletch` if it came from the environment, even under `'-e'`, you could instead write the following.

```
ifneq "$(findstring environment,$(origin bletch))" ""
bletch = barf, gag, etc.
endif
```

Here the redefinition takes place if `'$(origin bletch)'` returns either `'environment'` or `'environment override'`. See “Functions for string substitution and analysis” on page 395.

The `shell` function

The `shell` function is unlike any other function except the `wildcard` function (see “The wildcard function” on page 329) in that it communicates with the world outside of `make`.

The `shell` function performs the same function that backquotes (``) perform in most shells: it does *command expansion*. This means that it takes an argument that is a shell command and returns the output of the command. The only processing `make` does on the result, before substituting it into the surrounding text, is to convert newlines to spaces.

The commands run by calls to the `shell` function are run when the function calls are expanded. In most cases, this is when the makefile is read in. The exception is that function calls in the commands of the rules are expanded when the commands are run, and this applies to `shell` function calls like all others. The following is an example of the use of the `shell` function which sets `contents` to the contents of the file, `foo`, with a space (rather than a newline) separating each line.

```
contents := $(shell cat foo)
```

The following is an example of the use of the `shell` function which sets `files` to the expansion of `*.c`. Unless `make` is using a very strange shell, this has the same result as `$(wildcard *.c)`.

```
files := $(shell echo *.c)
```

The `shell` function

9

How to run the `make` utility

A makefile that says how to recompile a program can be used in more than one way. The simplest use is to recompile every file that is out of date.

Usually, makefiles are written so that if you run `make` with no arguments, it does just that. However, you might want to update only some of the files; you might want to use a different compiler or different compiler options; you might want just to find out which files are out of date without changing them.

The following documentation provides more details with running `make` for recompilation.

- “Arguments to specify the makefile” on page 408
- “Arguments to specify the goals” on page 409
- “Instead of executing the commands” on page 411
- “Avoiding recompilation of some files” on page 413
- “Overriding variables” on page 414
- “Testing the compilation of a program” on page 415

Arguments to specify the makefile

By giving arguments when you run `make`, you can do many things.

The exit status of `make` is always one of three values.

- 0
The exit status is zero if `make` is successful.
- 2
The exit status is two if `make` encounters any errors. It will print messages describing the particular errors.
- 1
The exit status is one if you use the `-q` flag and `make` determines that some target is not already up to date. See “Instead of executing the commands” on page 411.

The way to specify the name of the makefile is with the `-f` or `--file` options (`--makefile` also works). For example, `-f altmake` says to use the file `altmake` as the makefile.

If you use the `-f` flag several times and follow each `-f` with an argument, all the specified files are used jointly as makefiles.

If you do not use the `-f` or `--file` flag, the default is to try `gnumakefile`, `makefile`, and `Makefile`, in that order, and use the first of these three which exists or can be made (see “Writing makefiles” on page 313).

Arguments to specify the goals

The *goals* are the targets that `make` should strive ultimately to update. Other targets are updated as well if they appear as dependencies of goals, or dependencies of dependencies of goals, etc.

By default, the goal is the first target in the makefile (not counting targets that start with a period). Therefore, makefiles are usually written so that the first target is for compiling the entire program or programs they describe. If the first rule in the makefile has several targets, only the first target in the rule becomes the default goal, not the whole list.

You can specify a different goal or goals with arguments to `make`. Use the name of the goal as an argument. If you specify several goals, `make` processes each of them in turn, in the order you name them. Any target in the makefile may be specified as a goal (unless it starts with ‘-’ or contains an ‘=’, in which case it will be parsed as a switch or variable definition, respectively). Even targets not in the makefile may be specified, if `make` can find implicit rules that say how to make them.

One use of specifying a goal is if you want to compile only a part of the program, or only one of several programs. Specify as a goal each file that you wish to remake. For example, consider a directory containing several programs, with a makefile that starts like the following.

```
.PHONY: all all: size nm ld ar as
```

If you are working on the program `size`, you might want to say ‘`make size`’ so that only the files of that program are recompiled.

Another use of specifying a goal is to make files that are not normally made. For example, there may be a file of debugging output, or a version of the program that is compiled specially for testing, which has a rule in the makefile but is not a dependency of the default goal.

Another use of specifying a goal is to run the commands associated with a phony target (see “Phony targets” on page 334) or empty target (see “Empty target files to record events” on page 337). Many makefiles contain a phony target named ‘`clean`’ which deletes everything except source files. Naturally, this is done only if you request it explicitly with ‘`make clean`’.

Following is a list of typical phony and empty target names. See “Standard targets for users” on page 466 for a detailed list of all the standard target names which GNU software packages use.

‘`all`’

Makes all the top-level targets about which the makefile knows.

`'clean'`

Deletes all files that are normally created by running `make`.

`'mostlyclean'`

Like `'clean'`, but may refrain from deleting a few files that people normally don't want to recompile. For example, the `'mostlyclean'` target for GCC does not delete `'libgcc.a'`, because recompiling it is rarely necessary and takes a lot of time.

`'distclean'`

`'realclean'`

`'clobber'`

Any of these targets might be defined to delete more files than `'clean'` does. For example, this would delete configuration files or links that you would normally create as preparation for compilation, even if the makefile itself cannot create these files.

`'install'`

Copies the executable file into a directory that users typically search for commands; copy any auxiliary files that the executable uses into the directories where it will look for them.

`'print'`

Prints listings of the source files that have changed.

`'tar'`

Creates a `tar` file of the source files.

`'shar'`

Creates a shell archive (`shar` file) of the source files.

`'dist'`

Creates a distribution file of the source files. This might be a `tar` file, or a `shar` file, or a compressed version of one of the previous targets, or even more than one of the previous targets.

`'TAGS'`

Updates a tags table for this program.

`'check'`

`'test'`

Performs self tests on the program this makefile builds.

Instead of executing the commands

The makefile tells `make` how to tell whether a target is up to date, and how to update each target. But updating the targets is not always what you want.

The following options specify other activities for `make`.

`'-n'`

`'--just-print'`

`'--dry-run'`

`'--recon'`

“No-op.” The activity is to print what commands would be used to make the targets up to date, but not actually execute them.

`'-t'`

`'--touch'`

“Touch.” The activity is to mark the targets as up to date without actually changing them. In other words, `make` pretends to compile the targets but does not really change their contents.

`'-q'`

`'--question'`

“Question.” The activity is to find out silently whether the targets are up to date already; but execute no commands in either case. In other words, neither compilation nor output will occur.

`'-W file'`

`'--what-if=file'`

`'--assume-new=file'`

`'--new-file=file'`

“What if.” Each `'-w'` flag is followed by a file name. The given files' modification times are recorded by `make` as being the present time, although the actual modification times remain the same. You can use the `'-w'` flag in conjunction with the `'-n'` flag to see what would happen if you were to modify specific files.

With the `'-n'` flag, `make` prints the commands that it would normally execute but does not execute them.

With the `'-t'` flag, `make` ignores the commands in the rules and uses (in effect) the command, `touch`, for each target that needs to be remade. The `touch` command is also printed, unless `'-s'` or `.SILENT` is used. For speed, `make` does not actually invoke the program, `touch`. It does the work directly.

With the `-q` flag, `make` prints nothing and executes no commands, but the exit status code it returns is zero if and only if the targets to be considered are already up to date. If the exit status is one, then some updating needs to be done. If `make` encounters an error, the exit status is two, so you can distinguish an error from a target that is not up to date.

It is an error to use more than one of the three flags, `-n`, `-t`, and `-q`, in the same invocation of `make`.

The `-n`, `-t`, and `-q` options do not affect command lines that begin with `+` characters or contain the strings `$(MAKE)` or `${MAKE}`. Only the line containing the `+` character or the strings `$(MAKE)` or `${MAKE}` is run, regardless of these options. Other lines in the same rule are not run unless they too begin with `+` or contain `$(MAKE)` or `${MAKE}`. See “How the MAKE variable works” on page 358.

The `-w` flag provides two features:

- If you also use the `-n` or `-q` flag, you can see what `make` would do if you were to modify some files.
- Without the `-n` or `-q` flag, when `make` is actually executing commands, the `-w` flag can direct `make` to act as if some files had been modified, without actually modifying the files.

The options, `-p` and `-v`, allow you to obtain other information about `make` or about the makefiles in use (see “Summary of make options” on page 417).

Avoiding recompilation of some files

Sometimes you may have changed a source file but you do not want to recompile all the files that depend on it. For example, suppose you add a macro or a declaration to a header file that many other files depend on. Being conservative, `make` assumes that any change in the header file requires recompilation of all dependent files, but you know that they do not need to be recompiled and you would rather not waste the time waiting for them to compile.

If you anticipate the problem before changing the header file, you can use the `'-t'` flag. This flag tells `make` not to run the commands in the rules, but rather to mark the target up to date by changing its last-modification date.

Use the following procedure.

1. Use the command, `'make'`, to recompile the source files that really need recompilation
2. Make the changes in the header files.
3. Use the command, `'make -t'`, to mark all the object files as up to date. The next time you run `make`, the changes in the header files will not cause any recompilation.

If you have already changed the header file at a time when some files do need recompilation, it is too late to do this. Instead, you can use the `'-o file'` flag which marks a specified file as “old” (see “Summary of make options” on page 417), meaning that the file itself will not be remade, and nothing else will be remade on its account.

Use the following procedure.

1. Recompile the source files that need compilation for reasons independent of the particular header file, with `'make -o headerfile'`. If several header files are involved, use a separate `'-o'` option for each header file.
2. Touch all the object files with `'make -t'`.

Overriding variables

An argument that contains ‘=’ specifies the value of a variable: ‘*v*=*x*’ sets the value of the variable, *v*, to *x*. If you specify a value in this way, all ordinary assignments of the same variable in the makefile are ignored; we say they have been *overridden* by the command line argument.

The most common way to use this facility is to pass extra flags to compilers. For example, in a properly written makefile, the variable, `CFLAGS`, is included in each command that runs the C compiler. So, a file, ‘`foo.c`’, would be compiled using something like: `cc -c $(CFLAGS) foo.c`.

Thus, whatever value you set for `CFLAGS` affects each compilation that occurs.

The makefile probably specifies the usual value for `CFLAGS`, like: `CFLAGS=-g`.

Each time you run `make`, you can override this value if you wish. For example, if you say ‘`make CFLAGS='-g -O'`’, each C compilation will be done with ‘`cc -c -g -O`’. (This illustrates how you can use quoting in the shell to enclose spaces and other special characters in the value of a variable when you override it.)

The variable, `CFLAGS`, is only one of many standard variables that exist just so that you can change them this way. See “Variables used by implicit rules” on page 432 for a complete list.

You can also program the makefile to look at additional variables of your own, giving the user the ability to control other aspects of how the makefile works by changing the variables.

When you override a variable with a command argument, you can define either a recursively-expanded variable or a simply-expanded variable. The examples shown previously make a recursively-expanded variable; to make a simply-expanded variable, write ‘:=’ instead of ‘=’. Unless you want to include a variable reference or function call in the *value* that you specify, it makes no difference which kind of variable you create.

There is one way that the makefile can change a variable that you have overridden. This is to use the `override` directive, which is a line using something like ‘`override variable=value`’ (see “The `override` directive” on page 381).

Testing the compilation of a program

Normally, when an error happens in executing a shell command, `make` gives up immediately, returning a nonzero status. No further commands are executed for any target. The error implies that the goal cannot be correctly remade, and `make` reports this as soon as it knows.

When you are compiling a program that you have just changed, this is not what you want. Instead, you would rather that `make` try compiling every file that can be tried, to show you as many compilation errors as possible.

On these occasions, you should use the `-k` or `--keep-going` flag. This tells `make` to continue to consider the other dependencies of the pending targets, remaking them if necessary, before it gives up and returns nonzero status. For example, after an error in compiling one object file, `'make -k'` will continue compiling other object files even though it already knows that linking them will be impossible. In addition to continuing after failed shell commands, `'make -k'` will continue as much as possible after discovering that it does not know how to make a target or dependency file. This will always cause an error message, but without `-k`, it is a fatal error (see “Summary of make options” on page 417).

The usual behavior of `make` assumes that your purpose is to get the goals up to date; once `make` learns that this is impossible, it might as well report the failure immediately. The `-k` flag says that the real purpose is to test as much as possible of the changes made in the program, perhaps to find several independent problems so that you can correct them all before the next attempt to compile. This is why Emacs’ `Meta-x compile` command passes the `-k` flag by default.

10

Summary of make options

The following documentation discusses the options for `make`.

`'-b'`

`'-m'`

These options are ignored for compatibility with other versions of `make`.

`'-C dir'`

`'--directory=dir'`

Change to directory, *dir*, before reading the makefiles.

If multiple `'-C'` options are specified, each is interpreted relative to the previous one. `'-C / -C etc'` is equivalent to `'-C /etc'`.

This is typically used with recursive invocations of `make` (see “Recursive use of the make utility” on page 358).

`'-d'`

`'--debug'`

Print debugging information in addition to normal processing.

The debugging information says which files are being considered for remaking, which file-times are being compared and with what results, which files actually need to be remade, which implicit rules are considered and which are applied—everything interesting about how `make` decides what to do.

`-e`
`--environment-overrides`
Give variables taken from the environment precedence over variables from makefiles. See “Variables from the environment” on page 383.

`-f file`
`--file=file`
`--makefile=file`
Read the file named *file* as a makefile. See “Writing makefiles” on page 313.

`-h`
`--help`
Remind you of the options that make understands and then exit.

`-i`
`--ignore-errors`
Ignore all errors in commands executed to remake files. See “Errors in commands” on page 355.

`-I dir`
`--include-dir=dir`
Specifies a directory, *dir*, to search for included makefiles. See “Including other makefiles” on page 316. If several `-I` options are used to specify several directories, the directories are searched in the order specified.

`-j [jobs]`
`--jobs=[jobs]`
Specifies the number of jobs (commands) to run simultaneously. With no argument, **make** runs as many jobs simultaneously as possible. If there is more than one `-j` option, the last one is effective. See “Parallel execution” on page 353 for more information on how commands are run.

`-k`
`--keep-going`
Continue as much as possible after an error. While the target that failed, and those that depend on it, cannot be remade, the other dependencies of these targets can be processed all the same. See “Testing the compilation of a program” on page 415.

`-l [load]`
`--load-average[=load]`
`--max-load[=load]`
Specifies that no new jobs (commands) should be started if there are other jobs running and the load average is at least *load* (a floating-point number). With no argument, removes a previous load limit. See “Parallel execution” on page 353.

```

'-n'
'--just-print'
'--dry-run'
'--recon'

```

Print the commands that would be executed, but do not execute them. See “Instead of executing the commands” on page 411.

```

'-o file'
'--old-file=file'
'--assume-old=file'

```

Do not remake the file, *file*, even if it is older than its dependencies, and do not remake anything on account of changes in *file*. Essentially the file is treated as very old and its rules are ignored. See “Avoiding recompilation of some files” on page 413.

```

'-p'
'--print-data-base'

```

Print the data base (rules and variable values) that results from reading the makefiles; then execute as usual or as other-wise specified. This also prints the version information given by the ‘-v’ switch (see “-v” on page 420). To print the data base without trying to remake any files, use ‘make -p -f /dev/null’.

```

'-q'
'--question'

```

“Question mode”. Do not run any commands, or print anything; just return an exit status that is zero if the specified targets are already up to date, one if any remaking is required, or two if an error is encountered. See “Instead of executing the commands” on page 411.

```

'-r'
'--no-builtin-rules'

```

Eliminate use of the built-in implicit rules (see “Using implicit rules” on page 425). You can still define your own by writing pattern rules (see “Defining and redefining pattern rules” on page 436). The ‘-r’ option also clears out the default list of suffixes for suffix rules (see “Old-fashioned suffix rules” on page 445). But you can still define your own suffixes with a rule for `.SUFFIXES`, and then define your own suffix rules.

```

'-s'
'--silent'
'--quiet'

```

Silent operation; do not print the commands as they are executed. See “Command echoing” on page 351.

`'-s'`

`'--no-keep-going'`

`'--stop'`

Cancel the effect of the `'-k'` option. This is never necessary except in a recursive `make` where `'-k'` might be inherited from the top-level `make` via `MAKEFLAGS` or if you set `'-k'` in `MAKEFLAGS` in your environment (see “Recursive use of the `make` utility” on page 358).

`'-t'`

`'--touch'`

Touch files (mark them up to date without really changing them) instead of running their commands. This is used to pretend that the commands were done, in order to fool future invocations of `make`. See “Instead of executing the commands” on page 411.

`'-v'`

`'--version'`

Print the version of the `make` program plus a copyright, a list of authors, and a notice that there is no warranty; then exit.

`'-w'`

`'--print-directory'`

Print a message containing the working directory both before and after executing the makefile. This may be useful for tracking down errors from complicated nests of recursive `make` commands. See “Recursive use of the `make` utility” on page 358. (In practice, you rarely need to specify this option since `'make'` does it for you; see “The `'--print-directory'` option” on page 362.)

`'--no-print-directory'`

Disable printing of the working directory under `-w`. This option is useful when `-w` is turned on automatically, but you do not want to see the extra messages. See “The `'--print-directory'` option” on page 362.

`'-W file'`

`'--what-if=file'`

`'--new-file=file'`

`'--assume-new=file'`

Pretend that the target file has just been modified.

When used with the `'-n'` flag, this shows you what would happen if you were to modify that file. Without `'-n'`, it is almost the same as running a `touch` command on the given file before running `make`, except that the modification time is changed only in the imagination of `make`. See “Instead of executing the commands” on page 411.

`--warn-undefined-variables`

Issue a warning message whenever `make` sees a reference to an undefined variable. This can be helpful when you are trying to debug makefiles which use variables in complex ways.

Implicit rules

Certain standard ways of remaking target files are used very often. For example, one customary way to make an object file is from a C source file using the GNU C compiler, `gcc`. The following documentation describes in more detail the rules of remaking target files.

- “Using implicit rules” on page 425
- “Catalogue of implicit rules” on page 427
- “Variables used by implicit rules” on page 432
- “Chains of implicit rules” on page 435
- “Defining and redefining pattern rules” on page 436
- “Defining last-resort default rules” on page 444
- “Old-fashioned suffix rules” on page 445
- “Implicit rule search algorithm” on page 447

Implicit rules tell `make` how to use customary techniques so that you do not have to specify them in detail when you want to use them. For example, there is an implicit rule for C compilation. File names determine which implicit rules are run. For example, C compilation typically takes a `.c` file and makes a `.o` file. So `make` applies the implicit rule for C compilation when it sees this combination of file name endings.

A *chain of implicit rules* can apply in sequence; for example, **make** will remake a ‘.o’ file from a ‘.y’ file by way of a ‘.c’ file. See “Chains of implicit rules” on page 435.

The built-in implicit rules use several variables in their commands so that, by changing the values of the variables, you can change the way the implicit rule works. For example, the variable, `CFLAGS`, controls the flags given to the C compiler by the implicit rule for C compilation. See “Variables used by implicit rules” on page 432.

You can define your own implicit rules by writing *pattern rules*. See “Defining and redefining pattern rules” on page 436.

Suffix rules are a more limited way to define implicit rules. Pattern rules are more general and clearer, but suffix rules are retained for compatibility. See “Old-fashioned suffix rules” on page 445.

Using implicit rules

To allow **make** to find a customary method for updating a target file, all you have to do is refrain from specifying commands yourself. Either write a rule with no command lines, or don't write a rule at all. Then **make** will figure out which implicit rule to use based on which kind of source file exists or can be made. For example, suppose the makefile looks like the following specification.

```
foo : foo.o bar.o
    cc -o foo foo.o bar.o $(CFLAGS) $(LDFLAGS)
```

Because you mention 'foo.o' but do not give a rule for it, **make** will automatically look for an implicit rule that tells how to update it. This happens whether or not the file 'foo.o' currently exists.

If an implicit rule is found, it can supply both commands and one or more dependencies (the source files). You would want to write a rule for 'foo.o' with no command lines if you need to specify additional dependencies (such as header files) which the implicit rule cannot supply.

Each implicit rule has a target pattern and dependency patterns. There may be many implicit rules with the same target pattern. For example, numerous rules make '.o' files: one, from a '.c' file with the C compiler; another, from a '.p' file with the Pascal compiler; and so on. The rule that actually applies is the one whose dependencies exist or can be made. So, if you have a file 'foo.c', **make** will run the C compiler; otherwise, if you have a file 'foo.p', **make** will run the Pascal compiler; and so on. Of course, when you write the makefile, you know which implicit rule you want **make** to use, and you know it will choose that one because you know which possible dependency files are supposed to exist. See "Catalogue of implicit rules" on page 427 for a catalogue of all the predefined implicit rules.

Previously, we said an implicit rule applies if the required dependencies "exist or can be made". A file "can be made" if it is mentioned explicitly in the makefile as a target or a dependency, or if an implicit rule can be recursively found for how to make it. When an implicit dependency is the result of another implicit rule, we say that chaining is occurring. See "Chains of implicit rules" on page 435.

In general, **make** searches for an implicit rule for each target, and for each double-colon rule, that has no commands. A file that is mentioned only as a dependency is considered a target whose rule specifies nothing, so implicit rule search happens for it. See "Implicit rule search algorithm" on page 447 for the details of how the search is done.

NOTE: Explicit dependencies do not influence implicit rule search. For example, consider the explicit rule: `foo.o: foo.p`.

The dependency on `'foo.p'` does not necessarily mean that **make** will remake `'foo.o'` according to the implicit rule to make an object file, a `'o'` file, from a Pascal source file, a `'p'` file. For example, if `'foo.c'` also exists, the implicit rule to make an object file from a C source file is used instead, because it appears before the Pascal rule in the list of predefined implicit rules (see “Catalogue of implicit rules” on page 427).

If you do not want an implicit rule to be used for a target that has no commands, you can give that target empty commands by writing a semicolon (see “Using empty commands” on page 366).

Catalogue of implicit rules

The following is a catalogue of predefined implicit rules which are always available unless the makefile explicitly overrides or cancels them. See “Canceling implicit rules” on page 443 for information on canceling or overriding an implicit rule.

The ‘-r’ or ‘--no-builtin-rules’ option cancels all predefined rules.

Not all of these rules will always be defined, even when the ‘-r’ option is not given. Many of the predefined implicit rules are implemented in **make** as suffix rules, so which ones will be defined depends on the *suffix list* (the list of dependencies of the special target, `.SUFFIXES`).

The default suffix list is: `.out, .a, .ln, .o, .c, .cc, .C, .p, .f, .F, .r, .y, .l, .s, .S, .mod, .sym, .def, .h, .info, .dvi, .tex, .texinfo, .texi, .txinfo, .w, .ch, .web, .sh, .elc, .el`.

All of the implicit rules (described in the following) whose dependencies have one of these suffixes are actually suffix rules. If you modify the suffix list, the only predefined suffix rules in effect will be those named by one or two of the suffixes that are on the list you specify; rules whose suffixes fail to be on the list are disabled.

See “Old-fashioned suffix rules” on page 445 for full details on suffix rules.

Compiling C programs

‘`n.o`’ is made automatically from ‘`n.c`’ with a command of the form, ‘`$(CC) -c $(CPPFLAGS) $(CFLAGS)`’.

Compiling C++ programs

‘`n.o`’ is made automatically from ‘`n.cc`’ or ‘`n.C`’ with a command of the form, ‘`$(CXX) -c $(CPPFLAGS) $(CXXFLAGS)`’. We encourage you to use the suffix ‘`.cc`’ for C++ source files instead of ‘`.C`’.

Compiling Pascal programs

‘`n.o`’ is made automatically from ‘`n.p`’ with the command of the form, ‘`$(PC) -c $(PFLAGS)`’.

Compiling Fortran and Ratfor programs

‘`n.o`’ is made automatically from ‘`n.r`’, ‘`n.F`’ or ‘`n.f`’ by running the Fortran compiler. The precise command used is as follows:

```
‘.f’
‘$(FC) -c $(FFLAGS)’.
```

```
‘.F’
‘$(FC) -c $(FFLAGS) $(CPPFLAGS)’.
```

```
‘.r’
‘$(FC) -c $(FFLAGS) $(RFLAGS)’.
```

Preprocessing Fortran and Ratfor programs

'*n.f*' is made automatically from '*n.r*' or '*n.F*'. This rule runs just the preprocessor to convert a Ratfor or preprocessable Fortran program into a strict Fortran program. The precise command used is as follows:

```
'.F'  
$(FC) -F $(CPPFLAGS) $(FFLAGS)  
  
'.r'  
$(FC) -F $(FFLAGS) $(RFLAGS)
```

Compiling Modula-2 programs

'*n.sym*' is made from '*n.def*' with a command of the form:

```
$(M2C) $(M2FLAGS) $(DEFFLAGS)
```

'*n.o*' is made from '*n.mod*'; the form is:

```
$(M2C) $(M2FLAGS) $(MODFLAGS)
```

Assembling and preprocessing assembler programs

'*n.o*' is made automatically from '*n.s*' by running the assembler, **as**. The precise command is:

```
$(AS) $(ASFLAGS)
```

'*n.s*' is made automatically from '*n.S*' by running the C preprocessor, **cpp**. The precise command is:

```
$(CPP) $(CPPFLAGS)
```

Linking a single object file

'*n*' is made automatically from '*n.o*' by running the linker (usually called **ld**) via the C compiler. The precise command used is:

```
$(CC) $(LDFLAGS) n.o $(LOADLIBES)
```

This rule does the right thing for a simple program with only one source file. It will also do the right thing if there are multiple object files (presumably coming from various other source files), one of which has a name matching that of the executable file.

Thus, *x: y.o z.o*, when '*x.c*', '*y.c*' and '*z.c*' all exist will execute the following.

```
cc -c x.c -o x.o  
cc -c y.c -o y.o  
cc -c z.c -o z.o  
cc x.o y.o z.o -o x  
rm -f x.o  
rm -f y.o  
rm -f z.o
```

In more complicated cases, such as when there is no object file whose name derives from the executable file name, you must write an explicit command for linking. Each kind of file automatically made into '*.o*' object files will be

automatically linked by using the compiler (`$(CC)`, `$(FC)` or `$(PC)`; the C compiler, `$(CC)`, is used to assemble `.s` files) without the `-c` option. This could be done by using the `.o` object files as intermediates, but it is faster to do the compiling and linking in one step, so that's how it's done.

Yacc for C programs

`'n.c'` is made automatically from `'n.y'` by running Yacc with the command:

```
$(YACC) $(YFLAGS)
```

Lex for C programs

`'n.c'` is made automatically from `'n.l'` by by running Lex. The actual command is:

```
$(LEX) $(LFLAGS)
```

Lex for Ratfor programs

`'n.r'` is made automatically from `'n.l'` by by running Lex. The actual command is:

```
$(LEX) $(LFLAGS)
```

The convention of using the same suffix `.l` for all Lex files regardless of whether they produce Ccode or Ratfor code makes it impossible for **make** to determine automatically which of the two languages you are using in any particular case. If **make** is called upon to remake an object file from a `.l` file, it must guess which compiler to use. It will guess the C compiler, because that is more common. If you are using Ratfor, make sure make knows this by mentioning `'n.r'` in the makefile. Or, if you are using Ratfor exclusively, with no C files, remove `.c` from the list of implicit rule suffixes with the following:

```
.SUFFIXES:
.SUFFIXES: .o .r .f .l ...
```

Making Lint Libraries from C, Yacc, or Lex programs

`'n.ln'` is made from `'n.c'` by running `lint`. The precise command is shown in the following example's input.

```
$(LINT) $(LINTFLAGS) $(CPPFLAGS) -i
```

The same command is used on the C code produced from `'n.y'` or `'n.l'`.

TEX and Web

`'n.dvi'` is made from `'n.tex'` with the command `$(TEX)`. `'n.tex'` is made from `'n.web'` with `$(WEAVE)`, or from `'n.w'` (and from `'n.ch'` if it exists or can be made) with `$(CWEAVE)`.

`'n.p'` is made from `'n.web'` with `$(TANGLE)` and `'n.c'` is made from `'n.w'` (and from `'n.ch'` if it exists or can be made) with `$(CTANGLE)`.

Texinfo and Info

To make `'n.dvi'` from either `'n.texinfo'`, `'n.texi'`, or `'n.txinfo'`, use the command:

```
$(TEXI2DVI) $(TEXI2DVI_FLAGS)
```

To make `'n.info'` from either `'n.texinfo'`, `'n.texi'`, or `'n.txinfo'`, use the command in the form:

```
$(MAKEINFO) $(MAKEINFO_FLAGS)
```

RCS

Any file `'n'` is extracted if necessary from an RCS file named either `'n,v'` or `'RCS/n,v'`. The precise command used is the following.

```
$(CO) $(COFLAGS)
```

`'n'` will not be extracted from RCS if it already exists, even if the RCS file is newer. The rules for RCS are terminal (see “Match-anything pattern rules” on page 441), so RCS files cannot be generated from another source; they must actually exist.

SCCS

Any file `'n'` is extracted if necessary from an SCCS file named either `'s.n'` or `'SCCS/s.n'`. The precise command used is the following.

```
$(GET) $(GFLAGS)
```

The rules for SCCS are terminal (see “Match-anything pattern rules” on page 441), so SCCS files cannot be generated from another source; they must actually exist.

For the benefit of SCCS, a file `'n'` is copied from `'n.sh'` and made executable (by everyone). This is for shell scripts that are checked into SCCS. Since RCS preserves the execution permission of a file, you do not need to use this feature with RCS.

We recommend that you avoid using of SCCS. RCS is widely held to be superior, and is also free. By choosing free software in place of comparable (or inferior) proprietary software, you support the free software movement.

Usually, you want to change only the variables listed in the catalogue of implicit rules; for documentation on variables, see “Variables used by implicit rules” on page 432.

However, the commands in built-in implicit rules actually use variables such as `COMPILE.c`, `LINK.p`, and `PREPROCESS.s`, whose values contain the commands listed in the catalogue of implicit rules.

make follows the convention that the rule to compile a `'x'` source file uses the variable `COMPILE.x`. Similarly, the rule to produce an executable from a `'x'` file uses `LINK.x`; and the rule to preprocess a `'x'` file uses `PREPROCESS.x`.

Every rule that produces an object file uses the variable, `OUTPUT_OPTION`. **make** defines this variable either to contain `'-o $@'` or to be empty, depending on a compile-time option. You need the `'-o'` option to ensure that the output goes into the right file when the source file is in a different directory, as when using `VPATH` (see “Searching directories for dependencies” on page 330). However, compilers on some systems do not accept a `'-o'` switch for object files. If you use such a system, and use `VPATH`, some compilations will put their output in the wrong place. A possible workaround for this problem is to give `OUTPUT_OPTION` the value:

```
; mv $*.o $@
```

Variables used by implicit rules

The commands in built-in implicit rules make liberal use of certain predefined variables. You can alter these variables in the makefile, with arguments to make, or in the environment to alter how the implicit rules work without redefining the rules themselves.

For example, the command used to compile a C source file actually says `$(CC) -c $(CFLAGS) $(CPPFLAGS)`. The default values of the variables used are `cc` and nothing, resulting in the command `cc -c`. By redefining `CC` to `ncc`, you could cause `ncc` to be used for all C compilations performed by the implicit rule. By redefining `CFLAGS` to be `-g`, you could pass the `-g` option to each compilation. All implicit rules that do C compilation use `$(CC)` to get the program name for the compiler and all include `$(CFLAGS)` among the arguments given to the compiler.

The variables used in implicit rules fall into two classes:

- Those being names of programs (like `CC`).
- Those containing arguments for the programs (like `CFLAGS`). (The “name of a program” may also contain some command arguments, but it must start with an actual executable program name.) If a variable value contains more than one argument, separate them with spaces.

The following are variables used as names of programs in built-in rules.

AR

Archive-maintaining program; default: `ar`.

AS

Program for doing assembly; default: `as`.

CC

Program for compiling C programs; default: `cc`.

CXX

Program for compiling C++ programs; default: `g++`.

CO

Program for extracting a file from RCS; default: `co`.

CPP

Program for running the C preprocessor, with results to standard output; default: `$(CC) -E`.

FC

Program for compiling or preprocessing Fortran and Ratfor programs; default: `f77`.

GET

Program for extracting a file from SCCS; default: `get`.

LEX
 Program to use to turn Lex grammars into C programs or Ratfor programs;
 default: `'lex'`.

PC
 Program for compiling Pascal programs; default: `'pc'`.

YACC
 Program to use to turn Yacc grammars into C programs; default: `'yacc'`.

YACCR
 Program to use to turn Yacc grammars into Ratfor programs; default: `'yacc -r'`.

MAKEINFO
 Program to convert a Texinfo source file into an Info file; default: `'makeinfo'`.

TEX
 Program to make TEX DVI files from TEX source; default: `'tex'`.

TEXI2DVI
 Program to make TEX DVI files from Texinfo source; default: `'texi2dvi'`.

WEAVE
 Program to translate Web into TEX; default: `'weave'`.

CWEAVE
 Program to translate C Web into TEX; default: `'cweave'`.

TANGLE
 Program to translate Web into Pascal; default: `'tangle'`.

CTANGLE
 Program to translate C Web into C; default: `'ctangle'`.

RM
 Command to remove a file; default: `'rm -f'`.

The following are variables whose values are additional arguments for the previous list of programs associated with variables. The default values for all of these is the empty string, unless otherwise noted.

ARFLAGS
 Flags to give the archive-maintaining program; default: `'rv'`.

ASFLAGS
 Extra flags to give to the assembler (when explicitly invoked on a `'s'` or `'.s'` file).

CFLAGS
 Extra flags to give to the C compiler.

CXXFLAGS
 Extra flags to give to the C++ compiler.

COFLAGS
 Extra flags to give to the RCS `co` program.

CPPFLAGS

Extra flags to give to the C preprocessor and programs that use it (the C and Fortran compilers).

FFLAGS

Extra flags to give to the Fortran compiler.

GFLAGS

Extra flags to give to the SCCS `get` program.

LDFLAGS

Extra flags to give to compilers when they are supposed to invoke the linker, `'ld'`.

LFLAGS

Extra flags to give to Lex.

PFLAGS

Extra flags to give to the Pascal compiler.

RFLAGS

Extra flags to give to the Fortran compiler for Ratfor programs.

YFLAGS

Extra flags to give to Yacc.

Chains of implicit rules

Sometimes a file can be made by a sequence of implicit rules. For example, a file ‘*n.o*’ could be made from ‘*n.y*’ by running first Yacc and then *cc*. Such a sequence is called a *chain*.

If the file ‘*n.c*’ exists, or is mentioned in the makefile, no special searching is required: *make* finds that the object file can be made by C compilation from ‘*n.c*’; later on, when considering how to make ‘*n.c*’, the rule for running Yacc is used. Ultimately both ‘*n.c*’ and ‘*n.o*’ are updated.

However, even if ‘*n.c*’ does not exist and is not mentioned, *make* knows how to envision it as the missing link between ‘*n.o*’ and ‘*n.y*’! In this case, ‘*n.c*’ is called an intermediate file. Once *make* has decided to use the *intermediate file*, it is entered in the data base as if it had been mentioned in the makefile, along with the implicit rule that says how to create it.

Intermediate files are remade using their rules just like all other files. The difference is that the intermediate file is deleted when *make* is finished. Therefore, the intermediate file which did not exist before *make* also does not exist after *make*. The deletion is reported to you by printing a ‘*rm -f*’ command that shows what *make* is doing. You can list the target pattern of an implicit rule (such as ‘%.o’) as a dependency of the special target, *.PRECIOUS*, to preserve intermediate files made by implicit rules whose target patterns match that file’s name; see “Interrupting or killing the make utility” on page 357.

A chain can involve more than two implicit rules. For example, it is possible to make a file ‘*foo*’ from ‘*RCS/foo.y,v*’ by running RCS, Yacc and *cc*. Then both ‘*foo.y*’ and ‘*foo.c*’ are intermediate files that are deleted at the end.

No single implicit rule can appear more than once in a chain. This means that *make* will not even consider such a ridiculous thing as making ‘*foo*’ from ‘*foo.o.o*’ by running the linker twice. This constraint has the added benefit of preventing any infinite loop in the search for an implicit rule chain.

There are some special implicit rules to optimize certain cases that would otherwise be handled by rule chains. For example, making ‘*foo*’ from ‘*foo.c*’ could be handled by compiling and linking with separate chained rules, using ‘*foo.o*’ as an intermediate file. But what actually happens is that a special rule for this case does the compilation and linking with a single *cc* command. The optimized rule is used in preference to the step-by-step chain because it comes earlier in the ordering of rules.

Defining and redefining pattern rules

You define an implicit rule by writing a *pattern rule*. A pattern rule looks like an ordinary rule, except that its target contains the character ‘%’ (exactly one of them). The target is considered a pattern for matching file names; the ‘%’ can match any non-empty substring, while other characters match only themselves. The dependencies likewise use ‘%’ to show how their names relate to the target name. Thus, a pattern rule ‘%.o : %.c’ says how to make any file ‘*stem.o*’ from another file ‘*stem.c*’.

NOTE: Expansion using ‘%’ in pattern rules occurs after any variable or function expansions, which take place when the makefile is read. See “How to use variables” on page 367 and “Functions for transforming text” on page 393.

Fundamentals of pattern rules

A pattern rule contains the character ‘%’ (exactly one of them) in the target; otherwise, it looks exactly like an ordinary rule. The target is a pattern for matching file names; the ‘%’ matches any nonempty substring, while other characters match only themselves.

For example, ‘%.c’ as a pattern matches any file name that ends in ‘.c’. ‘%.%.c’ as a pattern matches any file name that starts with ‘%.’, ends in ‘.c’ and is at least five characters long. (There must be at least one character to match the ‘%’.) The substring that the ‘%’ matches is called the stem.

‘%’ in a dependency of a pattern rule stands for the same stem that was matched by the ‘%’ in the target. In order for the pattern rule to apply, its target pattern must match the file name under consideration, and its dependency patterns must name files that exist or can be made. These files become dependencies of the target.

Thus, a rule of the following form specifies how to make a file ‘*n.o*’, with another file ‘*n.c*’ as its dependency, provided that ‘*n.c*’ exists or can be made.

```
%.o : %.c ; command...
```

There may also be dependencies that do not use ‘%’; such a dependency attaches to every file made by this pattern rule. These unvarying dependencies are useful occasionally.

A pattern rule need not have any dependencies that contain ‘%’, or in fact any dependencies at all. Such a rule is effectively a general wildcard. It provides a way to make any file that matches the target pattern. See “Defining last-resort default rules” on page 444.

Pattern rules may have more than one target. Unlike normal rules, this does not act as many different rules with the same dependencies and commands. If a pattern rule has multiple targets, **make** knows that the rule's commands are responsible for making all of the targets. The commands are executed only once to make all the targets. When searching for a pattern rule to match a target, the target patterns of a rule other than the one that matches the target in need of a rule are incidental; **make** worries only about giving commands and dependencies to the file presently in question. However, when this file's commands are run, the other targets are marked as having been updated themselves.

The order in which pattern rules appear in the makefile is important since this is the order in which they are considered. Of equally applicable rules, only the first one found is used. The rules you write take precedence over those that are built in. However, a rule whose dependencies actually exist or are mentioned always takes priority over a rule with dependencies that must be made by chaining other implicit rules.

Pattern rule examples

The following are some examples of pattern rules actually predefined in **make**. First, the rule that compiles `‘.c’` files into `‘.o’` files:

```
% .o : % .c
    $(CC) -c $(CFLAGS) $(CPPFLAGS) $< -o $@
```

defines a rule that can make any file `‘x.o’` from `‘x.c’`. The command uses the automatic variables `‘$@’` and `‘$<’` to substitute the names of the target file and the source file in each case where the rule applies (see “Automatic variables” on page 438).

The following is a second built-in rule:

```
% :: RCS/% ,v
    $(CO) $(COFLAGS) $<
```

This statement defines a rule that can make any file `‘x’` whatsoever from a corresponding file `‘x,v’` in the subdirectory `‘RCS’`. Since the target is `‘%’`, this rule will apply to any file whatever, provided the appropriate dependency file exists. The double colon makes the rule *terminal*, meaning that its dependency may not be an intermediate file (see “Match-anything pattern rules” on page 441).

The following pattern rule has two targets:

```
%.tab.c %.tab.h: %.y
    bison -d $<
```

This tells make that the command `bison -dx.y` will make both `x.tab.c` and `x.tab.h`. If the file `foo` depends on the files `parse.tab.o` and `scan.o` and the file `scan.o` depends on the file `parse.tab.h`, when `parse.y` is changed, the command `bison -d parse.y` will be executed only once, and the dependencies of both `parse.tab.o` and `scan.o` will be satisfied. (Presumably the file `parse.tab.o` will be recompiled from `parse.tab.c` and the file `scan.o` from `scan.c`, while `foo` is linked from `parse.tab.o`, `scan.o`, and its other dependencies, and it will execute happily ever after.)

Automatic variables

Suppose you are writing a pattern rule to compile a `.c` file into a `.o` file: how do you write the `cc` command so that it operates on the right source file name? You cannot write the name in the command, because the name is different each time the implicit rule is applied. What you do is use a special feature of **make**, the *automatic variables*. These variables have values computed afresh for each rule that is executed, based on the target and dependencies of the rule. By example, you would use `$(@)` for the object file name and `$(<)` for the source file name.

The following is a list of automatic variables.

`$(@)`

The file name of the target of the rule. If the target is an archive member, then `$(@)` is the name of the archive file. In a pattern rule that has multiple targets (see “Fundamentals of pattern rules” on page 436), `$(@)` is the name of whichever target caused the rule’s commands to be run.

`$(%)`

The target member name, when the target is an archive member. See “Using make to update archive files” on page 449. For example, if the target is `foo.a(bar.o)` then `$(%)` is `bar.o` and `$(@)` is `foo.a`. `$(%)` is empty when the target is not an archive member.

`$(<)`

The name of the first dependency. If the target got its commands from an implicit rule, this will be the first dependency added by the implicit rule (see “Using implicit rules” on page 425).

`$(?)`

The names of all the dependencies that are newer than the target, with spaces between them. For dependencies which are archive members, only the member named is used (see “Using make to update archive files” on page 449).

\$^

The names of all the dependencies, with spaces between them. For dependencies which are archive members, only the member named is used (see “Using make to update archive files” on page 449). A target has only one dependency on each other file it depends on, no matter how many times each file is listed as a dependency. So if you list a dependency more than once for a target, the value of ‘\$^’ contains just one copy of the name.

\$+

This is like ‘\$^’, but dependencies listed more than once are duplicated in the order they were listed in the makefile. This is primarily useful for use in linking commands where it is meaningful to repeat library file names in a particular order.

\$*

The stem with which an implicit rule matches (see “How patterns match” on page 441). If the target is ‘dir/a.foo.b’ and the target pattern is ‘a.%.b’ then the stem is ‘dir/foo’. The stem is useful for constructing names of related files. In a static pattern rule, the stem is part of the file name that matched the ‘%’ in the target pattern.

In an explicit rule, there is no stem; so ‘\$*’ cannot be determined in that way. Instead, if the target name ends with a recognized suffix (see “Old-fashioned suffix rules” on page 445), ‘\$*’ is set to the target name minus the suffix. For example, if the target name is ‘foo.c’, then ‘\$*’ is set to ‘foo’, since ‘.c’ is a suffix. gnu make does this bizarre thing only for compatibility with other implementations of make. You should generally avoid using ‘\$*’ except in implicit rules or static pattern rules.

If the target name in an explicit rule does not end with a recognized suffix, ‘\$*’ is set to the empty string for that rule.

‘\$?’ is useful even in explicit rules when you wish to operate on only the dependencies that have changed. For example, suppose that an archive named ‘lib’ is supposed to contain copies of several object files. This rule copies just the changed object files into the archive:

```
lib: foo.o bar.o lose.o win.o
    ar r lib $?
```

Of the variables previously listed, four have values that are single file names, and two have values that are lists of file names. These six have variants that get just the file’s directory name or just the file name within the directory.

The variant variables’ names are formed by appending ‘D’ or ‘F’, respectively. These variants are semi-obsolete in GNU **make** since the functions `dir` and `notdir` can be used to get a similar effect (see “Functions for file names” on page 398).

NOTE: The ‘F’ variants all omit the trailing slash that always appears in the output of

the `dir` function.

The following is a list of the variants.

`$(@D)`

The directory part of the file name of the target, with the trailing slash removed. If the value of `$_` is `dir/foo.o` then `$(@D)` is `dir`. This value is `.` if `$_` does not contain a slash.

`$(@F)`

The file-within-directory part of the file name of the target. If the value of `$_` is `dir/foo.o` then `$(@F)` is `foo.o`. `$(@F)` is equivalent to `$(notdir $_)`.

`$(*D)`

`$(*F)`

The directory part and the file-within-directory part of the stem; `dir` and `foo` in this instance.

`$(%D)`

`$(%F)`

The directory part and the file-within-directory part of the target archive member name. This makes sense only for archive member targets of the form `archive(member)` and is useful only when member may contain a directory name. See “Archive members as targets” on page 450.

`$(<D)`

`$(<F)`

The directory part and the file-within-directory part of the first dependency.

`$(^D)`

`$(^F)`

Lists of the directory parts and the file-within-directory parts of all dependencies.

`$(?D)`

`$(?F)`

Lists of the directory parts and the file-within-directory parts of all dependencies that are newer than the target.

We use a special stylistic convention when we discuss these automatic variables; we write “the value of `$_`”, rather than “the variable, `$_`” as we would write for ordinary variables such as `objects` and `CFLAGS`. We think this convention looks more natural in this special case. Do not assume it has a deep significance; `$_` refers to the variable named `$_` just as `$(CFLAGS)` refers to the variable named `CFLAGS`. You could just as well use `$($_)` in place of `$_`.

How patterns match

A target pattern is composed of a ‘%’ between a prefix and a suffix, either or both of which may be empty. The pattern matches a file name only if the file name starts with the prefix and ends with the suffix, without overlap. The text between the prefix and the suffix is called the *stem*. Thus, when the pattern ‘%.o’ matches the file name ‘test.o’, the stem is ‘test’. The pattern rule dependencies are turned into actual file names by substituting the stem for the character ‘%’. Thus, if in the same example one of the dependencies is written as ‘%.c’, it expands to ‘test.c’.

When the target pattern does not contain a slash (and it usually does not), directory names in the file names are removed from the file name before it is compared with the target prefix and suffix. After the comparison of the file name to the target pattern, the directory names, along with the slash that ends them, are added on to the dependency file names generated from the pattern rule’s dependency patterns and the file name. The directories are ignored only for the purpose of finding an implicit rule to use, not in the application of that rule. Thus, ‘%t’ matches the file name ‘src/eat’, with ‘src/a’ as the stem. When dependencies are turned into file names, the directories from the stem are added at the front, while the rest of the stem is substituted for the ‘%’. The stem ‘src/a’ with a dependency pattern ‘c%r’ gives the file name ‘src/car’.

Match-anything pattern rules

When a pattern rule’s target is just ‘%’, it matches any file name whatever. We call these rules *match-anything* rules. They are very useful, but it can take a lot of time for **make** to think about them, because it must consider every such rule for each file name listed either as a target or as a dependency.

Suppose the makefile mentions ‘foo.c’. For this target, make would have to consider making it by linking an object file ‘foo.c.o’, or by C compilation-and-linking in one step from ‘foo.c.c’, or by Pascal compilation-and-linking from ‘foo.c.p’, and many other possibilities.

We know these possibilities are ridiculous since ‘foo.c’ is a C source file, not an executable. If **make** did consider these possibilities, it would ultimately reject them, because files such as ‘foo.c.o’ and ‘foo.c.p’ would not exist. But these possibilities are so numerous that make would run very slowly if it had to consider them.

To gain speed, we have put various constraints on the way make considers match-anything rules. There are two different constraints that can be applied, and each time you define a match-anything rule you must choose one or the other for that rule.

One choice is to mark the match-anything rule as *terminal* by defining it with a double colon. When a rule is terminal, it does not apply unless its dependencies actually exist. Dependencies that could be made with other implicit rules are not good enough. In other words, no further chaining is allowed beyond a terminal rule.

For example, the built-in implicit rules for extracting sources from RCS and SCCS files are terminal; as a result, if the file `'foo.c,v'` does not exist, **make** will not even consider trying to make it as an intermediate file from `'foo.c,v.o'` or from `'RCS/SCCS/s.foo.c,v'`. RCS and SCCS files are generally ultimate source files, which should not be remade from any other files; therefore, **make** can save time by not looking for ways to remake them.

If you do not mark the match-anything rule as terminal, then it is nonterminal. A non-terminal match-anything rule cannot apply to a file name that indicates a specific type of data. A file name indicates a specific type of data if some non-match-anything implicit rule target matches it.

For example, the file name `'foo.c'` matches the target for the pattern rule `'%.c : %.Y'` (the rule to run Yacc). Regardless of whether this rule is actually applicable (which happens only if there is a file `'foo.Y'`), the fact that its target matches is enough to prevent consideration of any non-terminal match-anything rules for the file `'foo.c'`. Thus, **make** will not even consider trying to make `'foo.c'` as an executable file from `'foo.c.o'`, `'foo.c.c'`, `'foo.c.p'`, etc.

The motivation for this constraint is that nonterminal match-anything rules are used for making files containing specific types of data (such as executable files) and a file name with a recognized suffix indicates some other specific type of data (such as a C source file).

Special built-in dummy pattern rules are provided solely to recognize certain file names so that nonterminal match-anything rules will not be considered. These dummy rules have no dependencies and no commands, and they are ignored for all other purposes. For example, the built-in implicit rule, `%.p :`, exists to make sure that Pascal source files such as `'foo.p'` match a specific target pattern and thereby prevent time from being wasted looking for `'foo.p.o'` or `'foo.p.c'`.

Dummy pattern rules such as the one for `'%.p'` are made for every suffix listed as valid for use in suffix rules (see “Old-fashioned suffix rules” on page 445).

Canceling implicit rules

You can override a built-in implicit rule (or one you have defined yourself) by defining a new pattern rule with the same target and dependencies, but different commands. When the new rule is defined, the built-in one is replaced. The new rule's position in the sequence of implicit rules is determined by where you write the new rule. You can cancel a built-in implicit rule by defining a pattern rule with the same target and dependencies, but no commands. For example, the following would cancel the rule that runs the assembler:

```
%.o : %.s
```

Defining last-resort default rules

You can define a last-resort implicit rule by writing a terminal match-anything pattern rule with no dependencies (see “Match-anything pattern rules” on page 441). This is just like any other pattern rule; the only thing special about it is that it will match any target. So such a rule’s commands are used for all targets and dependencies that have no commands of their own and for which no other implicit rule applies.

For example, when testing a makefile, you might not care if the source files contain real data, only that they exist. Then you might do the following.

```
% ::  
    touch $@
```

This causes all the source files needed (as dependencies) to be created automatically.

You can instead define commands to be used for targets for which there are no rules at all, even ones which don’t specify commands. You do this by writing a rule for the target, `.DEFAULT`. Such a rule’s commands are used for all dependencies which do not appear as targets in any explicit rule, and for which no implicit rule applies. Naturally, there is no `.DEFAULT` rule unless you write one.

If you use `.DEFAULT` with no commands or dependencies like: `.DEFAULT:`, the commands previously stored for `.DEFAULT` are cleared. Then **make** acts as if you had never defined `.DEFAULT` at all.

If you do not want a target to get the commands from a match-anything pattern rule or `.DEFAULT`, but you also do not want any commands to be run for the target, you can give it empty commands (see “Using empty commands” on page 366).

You can use a last-resort rule to override part of another makefile (see “Overriding part of another makefile” on page 321).

Old-fashioned suffix rules

Suffix rules are the old-fashioned way of defining implicit rules for **make**. Suffix rules are obsolete because pattern rules are more general and clearer. They are supported in GNU **make** for compatibility with old makefiles. They come in two kinds: *double-suffix* and *single-suffix*.

A double-suffix rule is defined by a pair of suffixes: the target suffix and the source suffix. It matches any file whose name ends with the target suffix. The corresponding implicit dependency is made by replacing the target suffix with the source suffix in the file name.

A two-suffix rule (whose target and source suffixes are `‘.o’` and `‘.c’`) is equivalent to the pattern rule, `%.o : %.c`.

A single-suffix rule is defined by a single suffix, which is the source suffix. It matches any file name, and the corresponding implicit dependency name is made by appending the source suffix. A single-suffix rule whose source suffix is `‘.c’` is equivalent to the pattern rule `% : %.c`.

Suffix rule definitions are recognized by comparing each rule’s target against a defined list of known suffixes. When **make** sees a rule whose target is a known suffix, this rule is considered a single-suffix rule. When **make** sees a rule whose target is two known suffixes concatenated, this rule is taken as a double-suffix rule. For example, `‘.c’` and `‘.o’` are both on the default list of known suffixes. Therefore, if you define a rule whose target is `‘.c.o’`, **make** takes it to be a double-suffix rule with source suffix, `‘.c’` and target suffix, `‘.o’`. The following is the old-fashioned way to define the rule for compiling a C source file.

```
.c.o:
    $(CC) -c $(CFLAGS) $(CPPFLAGS) -o $@ $<
```

Suffix rules cannot have any dependencies of their own. If they have any, they are treated as normal files with funny names, not as suffix rules. Thus, use the following rule.

```
.c.o: foo.h
    $(CC) -c $(CFLAGS) $(CPPFLAGS) -o $@ $<
```

This rule tells how to make the file, `‘.c.o’`, from the dependency file, `‘foo.h’`, and is not at all like the following pattern rule.

```
%.o: %.c foo.h
    $(CC) -c $(CFLAGS) $(CPPFLAGS) -o $@ $<
```

This rule tells how to make `‘.o’` files from `‘.c’` files, and makes all `‘.o’` files using this pattern rule also depend on `‘foo.h’`.

Suffix rules with no commands are also meaningless. They do not remove previous rules as do pattern rules with no commands (see “Canceling implicit rules” on page 443). They simply enter the suffix or pair of suffixes concatenated as a target in the data base.

The known suffixes are simply the names of the dependencies of the special target, `.SUFFIXES`. You can add your own suffixes by writing a rule for `.SUFFIXES` that adds more dependencies, as in: `.SUFFIXES: .hack .win`, which adds `‘.hack’` and `‘.win’` to the end of the list of suffixes.

If you wish to eliminate the default known suffixes instead of just adding to them, write a rule for `.SUFFIXES` with no dependencies. By special dispensation, this eliminates all existing dependencies of `.SUFFIXES`.

You can then write another rule to add the suffixes you want. For example, use the following.

```
.SUFFIXES: # Delete the default suffixes
.SUFFIXES: .c .o .h # Define our suffix list
```

The `‘-r’` or `‘--no-builtin-rules’` flag causes the default list of suffixes to be empty. The variable, `SUFFIXES`, is defined to the default list of suffixes before `make` reads any makefiles. You can change the list of suffixes with a rule for the special target, `.SUFFIXES`, but that does not alter this variable.

Implicit rule search algorithm

The following is the procedure **make** uses for searching for an implicit rule for a target, t . This procedure is followed for each double-colon rule with no commands, for each target of ordinary rules none of which have commands, and for each dependency that is not the target of any rule. It is also followed recursively for dependencies that come from implicit rules, in the search for a chain of rules.

Suffix rules are not mentioned in this algorithm because suffix rules are converted to equivalent pattern rules once the makefiles have been read in. For an archive member target of the form, `archive(member)`, the following algorithm is run twice, first using the entire target name, t , and, second, using `(member)` as the target, t , if the first run found no rule.

1. Split t into a directory part, called d , and the rest, called n . For example, if t is `src/foo.o`, then d is `src/` and n is `foo.o`.
2. Make a list of all the pattern rules one of whose targets matches t or n . If the target pattern contains a slash, it is matched against t ; otherwise, against n .
3. If any rule in that list is not a match-anything rule, then remove all non-terminal match-anything rules from the list.
4. Remove from the list all rules with no commands.
5. For each pattern rule in the list:
 - a. Find the stem s , which is the non-empty part of t or n matched by the `%` in the target pattern.
 - b. Compute the dependency names by substituting s for `%`; if the target pattern does not contain a slash, append d to the front of each dependency name.
 - c. Test whether all the dependencies exist or ought to exist. (If a file name is mentioned in the makefile as a target or as an explicit dependency, then we say it ought to exist.)
If all dependencies exist or ought to exist, or there are no dependencies, then this rule applies.
4. If no pattern rule has been found so far, try harder. For each pattern rule in the list:
 - a. If the rule is terminal, ignore it and go on to the next rule.
 - b. Compute the dependency names as before.
 - c. Test whether all the dependencies exist or ought to exist.
 - d. For each dependency that does not exist, follow this algorithm recursively to see if the dependency can be made by an implicit rule.
 - e. If all dependencies exist, ought to exist, or can be made by implicit rules, then

this rule applies.

6. If no implicit rule applies, the rule for `.DEFAULT`, if any, applies. In that case, give t the same commands that `.DEFAULT` has. Otherwise, there are no commands for t .

Once a rule that applies has been found, for each target pattern of the rule other than the one that matched t or n , the ‘%’ in the pattern is replaced with s and the resultant file name is stored until the commands to remake the target file, t , are executed. After these commands are executed, each of these stored file names are entered into the database and marked as having been updated and having the same update status as the file, t .

When the commands of a pattern rule are executed for t , the automatic variables are set corresponding to the target and dependencies. See “Automatic variables” on page 438.

12

Using `make` to update archive files

Archive files are files containing named subfiles called *members*; they are maintained with the program, `ar`, and their main use is as subroutine libraries for linking.

The following documentation discusses `make`'s updating of your archive files.

- “Archive members as targets” on page 450
- “Implicit rule for archive member targets” on page 451
- “Updating archive symbol directories” on page 451
- “Dangers when using archives” on page 453
- “Suffix rules for archive files” on page 454

Archive members as targets

An individual member of an archive file can be used as a target or dependency in **make**. You specify the member named *member* in archive file, *archive*, as follows:

```
archive(member)
```

This construct is available only in targets and dependencies, not in commands! Most programs that you might use in commands do not support this syntax and cannot act directly on archive members. Only **ar** and other programs specifically designed to operate on archives can do so. Therefore, valid commands to update an archive member target probably must use **ar**. For instance, this rule says to create a member, 'hack.o', in archive, 'foolib', by copying the file, 'hack.o' as in the following.

```
foolib(hack.o) : hack.o
    ar cr foolib hack.o
```

In fact, nearly all archive member targets are updated in just this way and there is an implicit rule to do it for you.

NOTE: The 'c' flag to **ar** is required if the archive file does not already exist.

To specify several members in the same archive, write all the member names together between the parentheses, as in the following example.

```
foolib(hack.o kludge.o)
```

The previous statement is equivalent to the following statement.

```
foolib(hack.o) foolib(kludge.o)
```

You can also use shell-style wildcards in an archive member reference. See “Using wildcard characters in file names” on page 326. For example, 'foolib(*.o)' expands to all existing members of the 'foolib' archive whose names end in '.o'; perhaps 'foolib(hack.o) foolib(kludge.o)'.

Implicit rule for archive member targets

Recall that a target that looks like ‘*a(m)*’ stands for the member named *m* in the archive file, *a*.

When **make** looks for an implicit rule for such a target, as a special feature, it considers implicit rules that match ‘*(m)*’ as well as those that match the actual target, ‘*a(m)*’.

This causes one special rule whose target is ‘*(%)*’ to match. This rule updates the target, ‘*a(m)*’, by copying the file, *m*, into the archive. For example, it will update the archive member target, ‘*foo.a(bar.o)*’ by copying the *file*, ‘*bar.o*’, into the archive, ‘*foo.a*’, as a *member* named ‘*bar.o*’.

When this rule is chained with others, the result is very powerful. Thus, ‘**make** “*foo.a(bar.o)*”’ (the quotes are needed to protect the ‘*(*’ and ‘*)*’ from being interpreted specially by the shell) in the presence of a file ‘*bar.c*’ is enough to cause the following commands to be run, even without a makefile:

```
cc -c bar.c -o bar.o
ar r foo.a bar.o
rm -f bar.o
```

Here **make** has envisioned the ‘*bar.o*’ as an intermediate file. See “Chains of implicit rules” on page 435.

Implicit rules such as this one are written using the automatic variable ‘*\$\$*’. See “Automatic variables” on page 438.

An archive member name in an archive cannot contain a directory name, but it may be useful in a makefile to pretend that it does. If you write an archive member target ‘*foo.a(dir/file.o)*’, **make** will perform automatic updating with the command, **ar r foo.a dir/file.o**, having the effect of copying the file, ‘*dir/file.o*’, into a member named ‘*file.o*’. In connection with such usage, the automatic variables, *%D* and *%F*, may be useful.

Updating archive symbol directories

An archive file that is used as a library usually contains a special member named ‘*__SYMDEF*’ which contains a directory of the external symbol names defined by all the other members.

After you update any other members, you need to update ‘*__SYMDEF*’ so that it will summarize the other members properly. This is done by running the **ranlib** program: **ranlib archivefile**.

Normally you would put this command in the rule for the archive file, and make all the members of the archive file dependencies of that rule. Use the following example, for instance.

```
libfoo.a: libfoo.a(x.o) libfoo.a(y.o) ...  
    ranlib libfoo.a
```

The effect of this is to update archive members ‘x.o’, ‘y.o’, etc., and then update the symbol directory member, ‘__SYMDEF’, by running `ranlib`. The rules for updating the members are not shown here; most likely you can omit them and use the implicit rule which copies files into the archive, as described in the preceding section (see “Implicit rule for archive member targets” on page 451 for more information).

This is not necessary when using the GNU **ar** program which automatically updates the ‘__SYMDEF’ member.

Dangers when using archives

It is important to be careful when using parallel execution (the `-j` switch; see “Parallel execution” on page 353) and archives. If multiple `ar` commands run at the same time on the same archive file, they will not know about each other and can corrupt the file.

Possibly a future version of `make` will provide a mechanism to circumvent this problem by serializing all commands that operate on the same archive file. But for the time being, you must either write your makefiles to avoid this problem in some other way, or not use `-j`.

Suffix rules for archive files

You can write a special kind of suffix rule for with archive files. See “Old-fashioned suffix rules” on page 445 for a full explanation of suffix rules.

Archive suffix rules are obsolete in GNU **make** because pattern rules for archives are a more general mechanism (see “Implicit rule for archive member targets” on page 451 for more information). But they are retained for compatibility with other **makes**.

To write a suffix rule for archives, you simply write a suffix rule using the target suffix, **‘.a’** (the usual suffix for archive files).

For example, the following is the old-fashioned suffix rule to update a library archive from C source files:

```
.c.a:
    $(CC) $(CFLAGS) $(CPPFLAGS) -c $< -o $*.o
    $(AR) r $@ $*.o
    $(RM) $*.o
```

This works just as if you had written the following pattern rule.

```
(%.o): %.c
    $(CC) $(CFLAGS) $(CPPFLAGS) -c $< -o $*.o
    $(AR) r $@ $*.o
    $(RM) $*.o
```

In fact, this is just what **make** does when it sees a suffix rule with **‘.a’** as the target suffix. Any double-suffix rule **‘.x.a’** is converted to a pattern rule with the target pattern, **‘(%.o)’** and a dependency pattern of **‘%.x’**. Since you might want to use **‘.a’** as the suffix for some other kind of file, **make** also converts archive suffix rules to pattern rules in the normal way (see “Old-fashioned suffix rules” on page 445). Thus a double-suffix rule, **‘.x.a’**, produces two pattern rules: **‘(%.o): %.x’** and **‘%.a: %.x’**.

13

Summary of the features for the GNU `make` utility

The following summary compares the features of GNU `make` with other versions of `make`.

We consider the features of `make` in 4.2 BSD systems as a baseline. If you are concerned with writing portable makefiles, consider the following features of `make`, using them with caution; see also “GNU `make`’s incompatibilities and missing features” on page 459. Many features come from the System V version of `make`.

- The `VPATH` variable and its special meaning. This feature exists in System V `make`, but is undocumented. It is documented in 4.3 BSD `make` (which says it mimics System V’s `VPATH` feature). See “Searching directories for dependencies” on page 330.
- Included makefiles. See “Including other makefiles” on page 316. Allowing multiple files to be included with a single directive is a GNU extension.
- Variables are read from and communicated, using the environment. See “Variables from the environment” on page 383.
- Options passed through the variable `MAKEFLAGS` to recursive invocations of `make`. See “Communicating options to a sub-make utility” on page 361.
- The automatic variable, `%`, is set to the member name in an archive reference. See “Automatic variables” on page 438.
- The automatic variables, `$(F)`, `$(*)`, `$(<)`, `$(%)`, and `$(?)`, have corresponding forms like `$(@F)` and `$(@D)`. We have generalized this to `$(^)` as an obvious extension. See

“Automatic variables” on page 438.

- Substitution variable references. See “Basics of variable references” on page 369.
- The command-line options, ‘-b’ and ‘-m’, are accepted and ignored. In System V `make`, these options actually do something.
- Execution of recursive commands to run `make`, using the variable `MAKE` even if ‘-n’, ‘-q’ or ‘-t’ is specified. See “Recursive use of the make utility” on page 358.
- Support for suffix ‘.a’ in suffix rules. See “Suffix Rules for Archive Files” on page . This feature is obsolete in GNU `make` because the general feature of rule chaining (see “Chains of Implicit Rules” on page) allows one pattern rule for installing members in an archive (see “Implicit Rule for Archive Member Targets” on page) to be sufficient.
- The arrangement of lines and backslash-newline combinations in commands is retained when the commands are printed, so they appear as they do in the makefile, except for the stripping of initial whitespace.

The following features were inspired by various other versions of `make`. In some cases it is unclear exactly which versions inspired others.

- Pattern rules using ‘%’. This has been implemented in several versions of `make`. We’re not sure who invented it first, but it’s been spread around a bit. See “Defining and Redefining Pattern Rules” on page .
- Rule chaining and implicit intermediate files. This was implemented by Stu Feldman in his version of `make` for AT&T Eighth Edition Research Unix, and later by Andrew Hume of AT&T Bell Labs in his `mk` program (where he terms it “transitive closure”). We do not really know if we got this from either of them or thought it up ourselves at the same time. See “Chains of Implicit Rules” on page .
- The automatic variable, `$$`, containing a list of all dependencies of the current target. We did not invent this, but we have no idea who did. See “Automatic variables” on page 438. The automatic variable, `$$+`, is a simple extension of `$$`.
- The “what if” flag (‘-w’ in GNU `make`) was (as far as we know) invented by Andrew Hume in `mk`. See “Instead of Executing the Commands” on page .
- The concept of doing several things at once (parallelism) exists in many incarnations of `make` and similar programs, though not in the System V or BSD implementations. See “Command Execution” on page .
- Modified variable references using pattern substitution come from SunOS 4. See “Basics of variable references” on page 369. This functionality was provided in GNU `make` by the `patsubst` function before the alternate syntax was implemented for compatibility with SunOS 4. It is not altogether clear who inspired whom, since GNU `make` had `patsubst` before SunOS 4 was released.

- The special significance of ‘+’ characters preceding command lines (see “Instead of Executing the Commands” on page) is mandated by IEEE Standard 1003.2-1992 (POSIX.2).
- The ‘+=’ syntax to append to the value of a variable comes from SunOS 4 `make`. See “Appending more text to variables” on page 379.
- The syntax ‘*archive(mem1 mem2 ...)*’ to list multiple members in a single archive file comes from SunOS 4 `make`. See “Implicit rule for archive member targets” on page 451.
- The `-include` directive to include makefiles with no error for a nonexistent file comes from SunOS 4 `make`. (But note that SunOS 4 `make` does not allow multiple makefiles to be specified in one `-include` directive.)

The remaining features are inventions new in GNU `make`:

- Use the ‘-v’ or ‘--version’ option to print version and copyright information.
- Use the ‘-h’ or ‘--help’ option to summarize the options to `make`.
- Simply-expanded variables. See “The two flavors of variables” on page 370.
- Pass command-line variable assignments automatically through the variable, `MAKE`, to recursive `make` invocations. See “Recursive use of the `make` utility” on page 358.
- Use the ‘-C’ or ‘--directory’ command option to change directory. See “Summary of `make` options” on page 417.
- Make verbatim variable definitions with `define`. See “Defining variables verbatim” on page 382.
- Declare phony targets with the special target, `.PHONY`.
Andrew Hume of AT&T Bell Labs implemented a similar feature with a different syntax in his `mk` program. This seems to be a case of parallel discovery. See “Phony targets” on page 334.
- Manipulate text by calling functions. See “Functions for transforming text” on page 393.
- Use the ‘-o’ or ‘--old-file’ option to pretend a file’s modification-time is old. See “Avoiding recompilation of some files” on page 413.
- Conditional execution has been implemented numerous times in various versions of `make`; it seems a natural extension derived from the features of the C preprocessor and similar macro languages and is not a revolutionary concept. See “Conditional parts of makefiles” on page 385.
- Specify a search path for included makefiles. See “Including other makefiles” on page 316.

- Specify extra makefiles to read with an environment variable. See “The `MAKEFILES` variable” on page 318.
- Strip leading sequences of `./` from file names, so that `./file` and `file` are considered to be the same file.
- Use a special search method for library dependencies written in the form, `-l name`. See “Directory search for link libraries” on page 332.
- Allow suffixes for suffix rules (see “Old-fashioned suffix rules” on page 445) to contain any characters. In other versions of `make`, they must begin with `.` and not contain any `/` characters.
- Keep track of the current level of `make` recursion using the variable, `MAKELEVEL`. See “Recursive use of the `make` utility” on page 358.
- Specify static pattern rules. See “Static pattern rules” on page 342.
- Provide selective `vpath` search. See “Searching directories for dependencies” on page 330.
- Provide computed variable references. See “Basics of variable references” on page 369.
- Update makefiles. See “How makefiles are remade” on page 319. System V `make` has a very, very limited form of this functionality in that it will check out SCCS files for makefiles.
- Various new built-in implicit rules. See “Catalogue of implicit rules” on page 427.
- The built-in variable, `MAKE_VERSION`, gives the version number of `make`.

GNU `make`'s incompatibilities and missing features

The following documentation describes some incompatibilities and missing features in GNU `make`. See also “Problems and bugs with make tools” on page 461. The `make` programs in various other systems support a few features that are not implemented in GNU `make`. The POSIX.2 standard (*IEEE Standard 1003.2-1992*) that specifies `make` does not require any of these features.

- A target of the form `'file((entry))'` stands for a member of archive file `file`. The member is chosen, not by name, but by being an object file which defines the linker symbol, `entry`.

This feature was not put into GNU `make` because of the non-modularity of putting knowledge into `make` of the internal format of archive file symbol tables. See “Updating archive symbol directories” on page 451.

- Suffixes (used in suffix rules) that end with the character `'~'` have a special meaning to System V `make`; they refer to the SCCS file that corresponds to the file one would get without the `'~'`. For example, the suffix rule, `'.c~.o'`, would make the file, `'n.o'`, from the SCCS file, `'s.n.c'`. For complete coverage, a whole series of such suffix rules is required. See “Old-fashioned suffix rules” on page 445.

In GNU `make`, this entire series of cases is handled by two pattern rules for extraction from SCCS, in combination with the general feature of rule chaining. See “Chains of implicit rules” on page 435.

- In System V `make`, the string, `'$$@'`, has the strange meaning that, in the

dependencies of a rule with multiple targets, it stands for the particular target that is being processed.

This is not defined in GNU `make` because '\$\$' should always stand for an ordinary '\$'. It is possible to get this functionality through the use of static pattern rules (see "Static pattern rules" on page 342).

The System V `make` rule: `$(targets): $$@.o lib.a` can be replaced with the GNU `make` static pattern rule: `$(targets): %: %.o lib.a`.

- In System V and 4.3 BSD `make`, files found by `VPATH` search (see "Searching directories for dependencies" on page 330) have their names changed inside command strings. We feel it is much cleaner to always use automatic variables and thus make this feature obsolete.
- In some Unix `makes`, the automatic variable, `$*`, appearing in the dependencies of a rule, has the amazingly strange "feature" of expanding to the full name of the target of that rule. We cannot imagine what went on in the minds of Unix `make` developers to do this; it is utterly inconsistent with the normal definition of `$*`.
- In some Unix `makes`, implicit rule search is apparently done for all targets, not just those without commands (see "Using implicit rules" on page 425). This means you can use: `foo.o: cc -c foo.c`, and Unix `make` will intuit that '`foo.o`' depends on '`foo.c`'.

We feel that such usage is broken. The dependency properties of `make` are well-defined (for GNU `make`, at least), and doing such a thing simply does not fit the model.

- GNU `make` does not include any built-in implicit rules for compiling or preprocessing EFL programs. If we hear of anyone who is using EFL, we will gladly add them.
- It appears that in SVR4 `make`, a suffix rule can be specified with no commands, and it is treated as if it had empty commands (see "Using empty commands" on page 366).

For example, `.c.a:` will override the built-in '`.c.a`' suffix rule. We feel that it is cleaner for a rule without commands to always simply add to the dependency list for the target. `.c.a:` can be easily rewritten to get the desired behavior in GNU `make`: `.c.a: ;`

- Some versions of `make` invoke the shell with the '`-e`' flag, except under '`-k`' (see "Testing the compilation of a program" on page 415). The '`-e`' flag tells the shell to exit as soon as any program it runs returns a nonzero status.

We feel it is cleaner to write each shell command line to stand on its own without requiring this special treatment.

Problems and bugs with `make` tools

If you have problems with GNU `make` or think you've found a bug, please report it to the developers; we cannot promise to do anything but we might well want to fix it.

Before reporting a bug, make sure you've actually found a real bug. Carefully reread the documentation and see if it really says you can do what you're trying to do. If it's not clear whether you should be able to do something or not, report that too; it's a bug in the documentation!

Before reporting a bug or trying to fix it yourself, try to isolate it to the smallest possible makefile that reproduces the problem. Then send us the makefile and the exact results `make` gave you. Also say what you expected to occur; this will help us decide whether the problem was really in the documentation.

Once you've got a precise problem, please send electronic mail either through the Internet or via UUCP.

Internet address:

`bug-gnu-utils@prep.ai.mit.edu`

UUCP path:

`mit-eddie!prep.ai.mit.edu!bug-gnu-utils`

Please include the version number of `make` you are using. You can get this information with the command, `make --version`. Be sure also to include the type of machine and operating system you are using. If possible, include the contents of the file, `config.h`, generated by the configuration process.

Non-bug suggestions are always welcome as well. If you have questions about things that are unclear in the documentation or that are just obscure features, send a message to the bug reporting address. We cannot guarantee that you'll get help with your problem, but many seasoned `make` users read the mailing list and they will probably try to help you out. The maintainers sometimes answer such questions as well, when time permits.

15

Makefile conventions

The following discusses conventions for writing the Makefiles for GNU programs.

- “General conventions for makefiles” on page 464
- “Utilities in makefiles” on page 465
- “Standard targets for users” on page 466
- “Variables for specifying commands” on page 470
- “Variables for installation directories” on page 471

General conventions for makefiles

Every Makefile should contain the following line to avoid trouble on systems where the `SHELL` variable might be inherited from the environment. (This is never a problem with GNU `make`.)

```
SHELL = /bin/sh
```

Different `make` programs have incompatible suffix lists and implicit rules, and this sometimes creates confusion or misbehavior.

So it is a good idea to set the suffix list explicitly using only the suffixes you need in the particular Makefile, using something like the following.

```
.SUFFIXES:
.SUFFIXES: .c .o
```

The first line clears out the suffix list, the second introduces all suffixes which may be subject to implicit rules in this Makefile.

Don't assume that `.` is in the path for command execution. When you need to run programs that are a part of your package during the make, make sure to use `./` if the program is built as part of the make or `$(srcdir)/` if the file is an unchanging part of the source code. Without one of these prefixes, the current search path is used.

The distinction between `./` and `$(srcdir)/` is important when using the `--srcdir` option to `configure`. A rule of the following form will fail when the current directory is not the source directory, because `foo.man` and `sedscrip`t are not in the current directory.

```
foo.1 : foo.man sedscrip
sed -e sedscrip foo.man > foo.1
```

When using GNU `make`, relying on `VPATH` to find the source file will work in the case where there is a single dependency file, since the `make` automatic variable, `$<`, will represent the source file wherever it is. (Many versions of `make` set `$<` only in implicit rules.) A makefile target like the following should instead be re-written in order to allow `VPATH` to work correctly.

```
foo.o : bar.c
$(CC) -I. -I$(srcdir) $(CFLAGS) -c bar.c -o foo.o
```

The makefile should be written like the following.

```
foo.o : bar.c
$(CC) -I. -I$(srcdir) $(CFLAGS) -c $< -o $@
```

When the target has multiple dependencies, using an explicit `\$(srcdir)` is the easiest way to make the rule work well. For instance, the previous target for `foo.1` is best written as the following.

```
foo.1 : foo.man sedscrip
sed -e $(srcdir)/sedscrip $(srcdir)/foo.man > $@
```

Utilities in makefiles

Write the Makefile commands (and any shell scripts, such as `configure`) to run in `sh`, not in `csh`. Don't use any special features of `ksh` or `bash`.

The `configure` script and the Makefile rules for building and installation should not use any utilities directly except the following:

```
cat cmp cp echo egrep expr grep
ln mkdir mv pwd rm rmdir sed test touch
```

Stick to the generally supported options for these programs. For example, don't use `'mkdir -p'`, convenient as it may be, because most systems don't support it. The Makefile rules for building and installation can also use compilers and related programs, but should do so using `make` variables so that the user can substitute alternatives.

The following are some of the programs.

```
ar bison cc flex install ld lex
make makeinfo ranlib texi2dvi yacc
```

Use the following `make` variables:

```
$(AR) $(BISON) $(CC) $(FLEX) $(INSTALL) $(LD) $(LEX)
$(MAKE) $(MAKEINFO) $(RANLIB) $(TEXI2DVI) $(YACC)
```

When you use `ranlib`, you should make sure nothing bad happens if the system does not have `ranlib`. Arrange to ignore an error from that command, and print a message before the command to tell the user that failure of the `ranlib` command does not mean a problem.

If you use symbolic links, you should implement a fallback for systems that don't have symbolic links.

It is acceptable to use other utilities in Makefile portions (or scripts) intended only for particular systems where you know those utilities to exist.

Standard targets for users

All GNU programs should have the following targets in their Makefiles:

‘all’

Compile the entire program. This should be the default target. This target need not rebuild any documentation files; Info files should normally be included in the distribution, and DVI files should be made only when explicitly asked for.

‘install’

Compile the program and copy the executables, libraries, and so on to the file names where they should reside for actual use. If there is a simple test to verify that a program is properly installed, this target should run that test.

If possible, write the `install` target rule so that it does not modify anything in the directory where the program was built, provided `‘make all’` has just been done. This is convenient for building the program under one user name and installing it under another. The commands should create all the directories in which files are to be installed, if they don’t already exist. This includes the directories specified as the values of the variables, `prefix` and `exec_prefix`, as well as all sub-directories that are needed. One way to do this is by means of an `installdirs` target.

Use `‘-’` before any command for installing a man page, so that `make` will ignore any errors. This is in case there are systems that don’t have the Unix man page documentation system installed. The way to install info files is to copy them into `‘$(infodir)’` with `$(INSTALL_DATA)` (see “Variables for specifying commands” on page 470), and then run the `install-info` program if it is present. `install-info` is a script that edits the Info `‘dir’` file to add or update the menu entry for the given info file; it will be part of the Texinfo package. The following is a sample rule to install an Info file:

```
$(infodir)/foo.info: foo.info
# There may be a newer info file in . than in srcdir.
  -if test -f foo.info; then d=.; \
    else d=$(srcdir); fi; \
  $(INSTALL_DATA) $$d/foo.info $@; \
# Run install-info only if it exists.
# Use 'if' instead of just prepending '-' to the
# line so we notice real errors from install-info.
# We use '$(SHELL) -c' because some shells do not
# fail gracefully when there is an unknown command.
  if $(SHELL) -c 'install-info --version' \
    >/dev/null 2>&1; then \
    install-info --infodir=$(infodir) $$d/foo.info; \
  else true; fi
```

`'uninstall'`

Delete all the installed files that the `'install'` target would create (but not the non-installed files such as `'make all'` would create).

This rule should not modify the directories where compilation is done, only the directories where files are installed.

`'clean'`

Delete all files from the current directory that are normally created by building the program. Don't delete the files that record the configuration. Also preserve files that could be made by building, but normally aren't because the distribution comes with them. Delete `'.dvi'` files here if they are not part of the distribution.

`'distclean'`

Delete all files from the current directory that are created by configuring or building the program. If you have unpacked the source and built the program without creating any other files, `'make distclean'` should leave only the files that were in the distribution.

`'mostlyclean'`

Like `'clean'`, but may refrain from deleting a few files that people normally don't want to recompile. For example, the `'mostlyclean'` target for GCC does not delete `'libgcc.a'`, because recompiling it is rarely necessary and takes a lot of time.

`'maintainer-clean'`

Delete almost everything from the current directory that can be reconstructed with this Makefile. This typically includes everything deleted by `distclean`, plus more: C source files produced by Bison, tags tables, info files, and so on.

The reason we say “almost everything” is that `'make maintainer-clean'` should not delete `'configure'` even if `'configure'` can be remade using a rule in the Makefile. More generally, `'make maintainer-clean'` should not delete anything that needs to exist in order to run `'configure'` and then begin to build the program. This is the only exception; `maintainer-clean` should delete everything else that can be rebuilt. The `'maintainer-clean'` is intended to be used by a maintainer of the package, not by ordinary users. You may need special tools to reconstruct some of the files that `'make maintainer-clean'` deletes. Since these files are normally included in the distribution, we don't take care to make them easy to reconstruct. If you find you need to unpack the full distribution again, don't blame us. To help make users aware of this, `maintainer-clean` should start with the following two commands.

```
@echo "This command is intended for maintainers \
to use;"
@echo "it deletes files that may require special \
tools to rebuild."
```

‘TAGS’

Update a tags table for this program.

‘info’

Generate any Info files needed. The best way to write the rules is as follows.

```
info: foo.info

foo.info: foo.texi chap1.texi chap2.texi $(MAKEINFO)
        $(srcdir)/foo.texi
```

You must define the variable `MAKEINFO` in the Makefile. It should run the `makeinfo` program which is part of the Texinfo distribution.

‘dvi’

Generate DVI files for all Texinfo documentation. For example:

```
dvi: foo.dvi

foo.dvi: foo.texi chap1.texi chap2.texi $(TEXI2DVI)
        $(srcdir)/foo.texi
```

You must define the variable, `TEXI2DVI`, in the Makefile. It should run the program, `texi2dvi`, which is part of the Texinfo distribution. Alternatively, write just the dependencies, and allow GNU `make` to provide the command.

‘dist’

Create a distribution tar file for this program. The `tar` file should be set up so that the file names in the `tar` file start with a subdirectory name which is the name of the package it is a distribution for. This name can include the version number. For example, the distribution `tar` file of GCC version 1.40 unpacks into a subdirectory named `‘gcc-1.40’`.

The easiest way to do this is to create a subdirectory appropriately named, use `ln` or `cp` to install the proper files in it, and then `tar` that subdirectory. The `dist` target should explicitly depend on all non-source files that are in the distribution, to make sure they are up to date in the distribution. See section “Making Releases” in GNU Coding Standards.

‘check’

Perform self-tests (if any). The user must build the program before running the tests, but need not install the program; you should write the self-tests so that they work when the program is built but not installed.

The following targets are suggested as conventional names, for programs in which they are useful.

installcheck

Perform installation tests (if any). The user must build and install the program before running the tests. You should not assume that `‘\$(bindir)’` is in the search path.

`installdirs`

It's useful to add a target named 'installdirs' to create the directories where files are installed, and their parent directories. There is a script called 'mkinstalldirs' which is convenient for this; find it in the Texinfo package. You can use a rule like this:

```
# Make sure all installation directories
# (e.g. $(bindir)) actually exist by
# making them if necessary.
installdirs: mkinstalldirs
              $(srcdir)/mkinstalldirs      $(bindir) $(datadir) \
                                              $(libdir) $(infodir) \
                                              $(mandir)
```

This rule should not modify the directories where compilation is done. It should do nothing but create installation directories.

Variables for specifying commands

Makefiles should provide variables for overriding certain commands, options, and so on.

In particular, you should run most utility programs via variables. Thus, if you use Bison, have a variable named `BISON` whose default value is set with `'BISON = bison'`, and refer to it with `\$(BISON)` whenever you need to use Bison.

File management utilities such as `ln`, `rm`, `mv`, and so on, need not be referred to through variables in this way, since users don't need to replace them with other programs.

Each program-name variable should come with an options variable that is used to supply options to the program. Append `'FLAGS'` to the program-name variable name to get the options variable name—for example, `BISONFLAGS`. (The name `CFLAGS` is an exception to this rule, but we keep it because it is standard.) Use `CPPFLAGS` in any compilation command that runs the preprocessor, and use `LDFLAGS` in any compilation command that does linking as well as in any direct use of `ld`.

If there are C compiler options that *must* be used for proper compilation of certain files, do not include them in `CFLAGS`. Users expect to be able to specify `CFLAGS` freely themselves. Instead, arrange to pass the necessary options to the C compiler independently of `CFLAGS`, by writing them explicitly in the compilation commands or by defining an implicit rule, like the following.

```
CFLAGS = -g
ALL_CFLAGS = -I. $(CFLAGS)
.c.o:
    $(CC) -c $(CPPFLAGS) $(ALL_CFLAGS) $<
```

Do include the `'-g'` option in `CFLAGS` because that is not required for proper compilation. You can consider it a default that is only recommended. If the package is set up so that it is compiled with GCC by default, then you might as well include `'-O'` in the default value of `CFLAGS` as well. Put `CFLAGS` last in the compilation command, after other variables containing compiler options, so the user can use `CFLAGS` to override the others. Every Makefile should define the variable, `INSTALL`, which is the basic command for installing a file into the system. Every Makefile should also define the variables `INSTALL_PROGRAM` and `INSTALL_DATA`. (The default for each of these should be `\$(INSTALL)`.) Then it should use those variables as the commands for actual installation, for executables and nonexecutables respectively. Use these variables as follows:

```
$(INSTALL_PROGRAM) foo $(bindir)/foo
$(INSTALL_DATA) libfoo.a $(libdir)/libfoo.a
```

Always use a file name, not a directory name, as the second argument of the installation commands. Use a separate command for each file to be installed.

Variables for installation directories

Installation directories should always be named by variables, so it is easy to install in a nonstandard place. The standard names for these variables are described in the following documentation. They are based on a standard filesystem layout; variants of it are used in SVR4, 4.4BSD, Linux, Ultrix v4, and other modern operating systems. These two variables set the root for the installation. All the other installation directories should be subdirectories of one of these two, and nothing should be directly installed into these two directories.

`'prefix'`

A prefix used in constructing the default values of the variables listed in the following discussions for installing directories. The default value of `prefix` should be `'/usr/local'`

When building the complete GNU system, the prefix will be empty and `'/usr'` will be a symbolic link to `'/'`.

`'exec_prefix'`

A prefix used in constructing the default values of some of the variables listed in the following discussions for installing directories. The default value of `exec_prefix` should be `\$(prefix)`.

Generally, `\$(exec_prefix)` is used for directories that contain machine-specific files (such as executables and subroutine libraries), while `\$(prefix)` is used directly for other directories.

Executable programs are installed in one of the following directories.

`'bindir'`

The directory for installing executable programs users run. This should normally be `'/usr/local/bin'` but write it as `\$(exec_prefix)/bin`.

`'sbindir'`

The directory for installing executable programs that can be run from the shell but are only generally useful to system administrators. This should normally be `'/usr/local/sbin'` but write it as `\$(exec_prefix)/sbin`.

`'libexecdir'`

The directory for installing executable programs to be run by other programs rather than by users. This directory should normally be `'/usr/local/libexec'`, but write it as `\$(exec_prefix)/libexec`.

Data files used by `make` during its execution are divided into categories in the following two ways.

- Some files are normally modified by programs; others are never normally modified (though users may edit some of these).

- Some files are architecture-independent and can be shared by all machines at a site; some are architecture-dependent and can be shared only by machines of the same kind and operating system; others may never be shared between two machines.

This makes for six different possibilities. However, we want to discourage the use of architecture-dependent files, aside from object files and libraries. It is much cleaner to make other data files architecture-independent, and it is generally not hard.

Therefore, the following are the variables makefiles should use to specify directories.

`'datadir'`

The directory for installing read-only architecture independent data files. This should normally be `'/usr/local/share'`, but write it as `'\$(prefix)/share'`. As a special exception, see `'\$(infodir)'` and `'\$(includedir)'` in the following discussions for them.

`'sysconfdir'`

The directory for installing read-only data files that pertain to a single machine—that is to say, files for configuring a host. Mailer and network configuration files, `'/etc/passwd'`, and so forth belong here. All the files in this directory should be ordinary ASCII text files. This directory should normally be `'/usr/local/etc'`, but write it as `'\$(prefix)/etc'`.

Do not install executables in this directory (they probably belong in `'\$(libexecdir)'` or `'\$(sbindir)'`). Also do not install files that are modified in the normal course of their use (programs whose purpose is to change the configuration of the system excluded). Those probably belong in `'\$(localstatedir)'`.

`'sharedstatedir'`

The directory for installing architecture-independent data files which the programs modify while they run. This should normally be `'/usr/local/com'`, but write it as `'\$(prefix)/com'`.

`'localstatedir'`

The directory for installing data files which the programs modify while they run, and that pertain to one specific machine. Users should never need to modify files in this directory to configure the package's operation; put such configuration information in separate files that go in `'datadir'` or `'\$(sysconfdir)'`.

`'\$(localstatedir)'` should normally be `'/usr/local/var'`, but write it as `'\$(prefix)/var'`.

‘libdir’

The directory for object files and libraries of object code. Do not install executables here, they probably belong in ‘`\$(libexecdir)`’ instead. The value of `libdir` should normally be ‘`/usr/local/lib`’, but write it as ‘`\$(exec_prefix)/lib`’.

‘infodir’

The directory for installing the Info files for this package. By default, it should be ‘`/usr/local/info`’, but it should be written as ‘`\$(prefix)/info`’.

‘includedir’

The directory for installing header files to be included by user programs with the C ‘`#include`’ preprocessor directive. This should normally be ‘`/usr/local/include`’, but write it as ‘`\$(prefix)/include`’.

Most compilers other than GCC do not look for header files in ‘`/usr/local/include`’. So installing the header files this way is only useful with GCC. Sometimes this is not a problem because some libraries are only really intended to work with GCC. But some libraries are intended to work with other compilers. They should install their header files in two places, one specified by `includedir` and one specified by `oldincludedir`.

‘oldincludedir’

The directory for installing ‘`#include`’ header files for use with compilers other than GCC. This should normally be ‘`/usr/include`’.

The Makefile commands should check whether the value of `oldincludedir` is empty. If it is, they should not try to use it; they should cancel the second installation of the header files. A package should not replace an existing header in this directory unless the header came from the same package. Thus, if your Foo package provides a header file ‘`foo.h`’, then it should install the header file in the `oldincludedir` directory if either (1) there is no ‘`foo.h`’ there, or, (2), the ‘`foo.h`’ that exists came from the Foo package. To tell whether ‘`foo.h`’ came from the Foo package, put a magic string in the file—part of a comment—and `grep` for that string.

Unix-style `man` pages are installed in one of the following:

‘mandir’

The directory for installing the man pages (if any) for this package. It should include the suffix for the proper section of the documentation—usually ‘`1`’ for a utility. It will normally be ‘`/usr/local/man/man1`’ but you should write it as ‘`\$(prefix)/man/man1`’.

‘man1dir’

The directory for installing section 1 `man` pages.

`'man2dir'`

The directory for installing section 2 `man` pages.

`'...'`

Use these names instead of `'mandir'` if the package needs to install `man` pages in more than one section of the documentation.

WARNING: Don't make the primary documentation for any GNU software be a `man` page.

Write a manual in Texinfo instead. `man` pages are just for the sake of people running GNU software on Unix which is only a secondary application.

`'manext'`

The file name extension for the installed `man` page. This should contain a period followed by the appropriate digit; it should normally be `'.1'`.

`'man1ext'`

The file name extension for installed section 1 `man` pages.

`'man2ext'`

The file name extension for installed section 2 `man` pages.

`'...'`

Use these names instead of `'manext'` if the package needs to install `man` pages in more than one section of the documentation.

And finally, you should set the following variable:

`'srcdir'`

The directory for the sources being compiled. The value of this variable is normally inserted by the `configure` shell script.

Use the following for example.

```
# Common prefix for installation directories.
# NOTE: This directory must exist when you start the install.
prefix = /usr/local
exec_prefix = $(prefix)
# Where to put the executable for the command 'gcc'.
bindir = $(exec_prefix)/bin
# Where to put the directories used by the compiler.
libexecdir = $(exec_prefix)/libexec
# Where to put the Info files.
infodir = $(prefix)/info
```

If your program installs a large number of files into one of the standard user-specified directories, it might be useful to group them into a subdirectory particular to that program. If you do this, you should write the `install` rule to create these subdirectories.

Do not expect the user to include the subdirectory name in the value of any of the variables previously discussed. The idea of having a uniform set of variable names for

installation directories is to enable the user to specify the exact same values for several different GNU packages. In order for this to be useful, all the packages must be designed so that they will work sensibly when the user does so.

GNU `make` quick reference

The following summary describes the directives, text manipulation functions, and special variables that GNU `make` understands and recognizes. See “Special built-in target names” on page 338, “Catalogue of implicit rules” on page 427, and “Summary of make options” on page 417 for other discussions.

```
define variable
endef
```

Define a multi-line, recursively-expanded variable. See “Defining canned command sequences” on page 364.

```
ifndef variable
ifndef variable
ifeq (a,b)
ifeq "a" "b"
ifeq 'a' 'b'
ifneq (a,b)
ifneq "a" "b"
ifneq 'a' 'b'
else
endif
```

Conditionally evaluate part of the makefile. See “Conditional parts of makefiles” on page 385.

`include file`

Include another makefile. See “Including other makefiles” on page 316.

`override variable= value`

`override variable:= value`

`override variable+= value`

`override define variable`

`endif`

Define a variable, overriding any previous definition, even one from the command line. See “The override directive” on page 381.

`export`

Tell make to export all variables to child processes by default. See “Communicating variables to a sub-make utility” on page 359.

`export variable`

`export variable= value`

`export variable:= value`

`export variable+= value`

`unexport variable`

Tell **make** whether or not to export a particular variable to child processes. See “Communicating variables to a sub-make utility” on page 359.

`vpath pattern path`

Specify a search path for files matching a ‘%’ pattern. See “The vpath directive” on page 330.

`vpath pattern`

Remove all search paths previously specified for *pattern*.

`vpath`

Remove all search paths previously specified in any `vpath` directive.

The following is a summary of the text manipulation functions (see “Functions for transforming text” on page 393):

`$(subst from, to, text)`

Replace *from* with *to* in *text*. See “Functions for string substitution and analysis” on page 395.

`$(patsubst pattern, replacement, text)`

Replace words matching *pattern* with *replacement* in *text*. See “Functions for string substitution and analysis” on page 395.

`$(strip string)`

Remove excess whitespace characters from *string*. See “Functions for string substitution and analysis” on page 395.

`$(findstring find, text)`
 Locate *find* in *text*. See “Functions for string substitution and analysis” on page 395.

`$(filter pattern..., text)`
 Select words in *text* that match one of the *pattern* words. See “Functions for string substitution and analysis” on page 395.

`$(filter-out pattern..., text)`
 Select words in *text* that do not match any of the *pattern* words. See “Functions for string substitution and analysis” on page 395.

`$(sort list)`
 Sort the words in *list* lexicographically, removing duplicates. See “Functions for string substitution and analysis” on page 395.

`$(dir names...)`
 Extract the directory part of each file name. See “Functions for file names” on page 398.

`$(notdir names...)`
 Extract the non-directory part of each file name. See “Functions for file names” on page 398.

`$(suffix names...)`
 Extract the suffix (the last ‘.’ and following characters) of each file name. See “Functions for file names” on page 398.

`$(basename names...)`
 Extract the base name (name without suffix) of each file name. See “Functions for file names” on page 398.

`$(addsuffix suffix, names...)`
 Append *suffix* to each word in *names*. See “Functions for file names” on page 398.

`$(addprefix prefix, names...)`
 Prepend *prefix* to each word in *names*. See “Functions for file names” on page 398.

`$(join list1, list2)`
 Join two parallel lists of words. See “Functions for file names” on page 398.

`$(word n, text)`
 Extract the *n*th word (one-origin) of *text*. See “Functions for file names” on page 398.

`$(words text)`
 Count the number of words in *text*. See “Functions for file names” on page 398.

`$(firstword names...)`
 Extract the first word of *names*. See “Functions for file names” on page 398.

`$(wildcard pattern ...)`

Find file names matching a shell file name, *pattern* (not a ‘%’ pattern). See “The wildcard function” on page 329.

`$(shell command)`

Execute a shell command and return its output. See “The shell function” on page 405.

`$(origin variable)`

Return a string describing how the make variable, *variable*, was defined. See “The origin function” on page 403.

`$(foreach var, words, text)`

Evaluate *text* with *var* bound to each word in *words*, and concatenate the results. See “The foreach function” on page 401.

The following is a summary of the automatic variables. See “Automatic variables” on page 438 for full information.

`$@`

The file name of the target.

`$%`

The target member name, when the target is an archive member.

`$<`

The name of the first dependency.

`$?`

The names of all the dependencies that are newer than the target, with spaces between them. For dependencies which are archive members, only the member named is used (see “Using make to update archive files” on page 449).

`$^`

`$+`

The names of all the dependencies with spaces between them. For dependencies which are archive members, only the member named is used (see “Using make to update archive files” on page 449). The value of `$^` omits duplicate dependencies while `$+` retains them and preserves their order.

`$*`

The stem with which an implicit rule matches (see “How patterns match” on page 441).

`$(@D)`

`$(@F)`

The directory part and the file-within-directory part of `$@`.

`$(*D)`

`$(*F)`

The directory part and the file-within-directory part of `$*`.

`$(%D)`

`$(%F)`

The directory part and the file-within-directory part of `$%`.

`$(<D)`

`$(<F)`

The directory part and the file-within-directory part of `$<`.

`$(^D)`

`$(^F)`

The directory part and the file-within-directory part of `$^`.

`$(+D)`

`$(+F)`

The directory part and the file-within-directory part of `$+`.

`$(?D)`

`$(?F)`

The directory part and the file-within-directory part of `$?`.

The following variables are used specially by GNU `make`.

MAKEFILES

Makefiles to be read on every invocation of `make`. See “The **MAKEFILES** variable” on page 318.

VPATH

Directory search path for files not found in the current directory.

See “**VPATH**: search path for all dependencies” on page 330.

SHELL

The name of the system default command interpreter, usually `/bin/sh`. You can set **SHELL** in the makefile to change the shell used to run commands. See “Command execution” on page 352.

MAKE

The name with which `make` was invoked. Using this variable in commands has special meaning. See “How the **MAKE** variable works” on page 358.

MAKELEVEL

The number of levels of recursion (sub-`makes`). See “Communicating variables to a sub-make utility” on page 359.

MAKEFLAGS

The flags given to `make`. You can set this in the environment or a makefile to set flags. See “Communicating variables to a sub-make utility” on page 359.

SUFFIXES

The default list of suffixes before `make` reads any makefiles.

Complex makefile example

The following is the makefile for the GNU **tar** program. This is a moderately complex makefile.

Because it is the first target, the default goal is `'all'`. An interesting feature of this makefile is that `'testpad.h'` is a source file automatically created by the testpad program, itself compiled from `'testpad.c'`.

If you type `'make'` or `'make all'`, then **make** creates the `'tar'` executable, the `'rmt'` daemon that provides remote tape access, and the `'tar.info'` Info file.

If you type `'make install'`, then **make** not only creates `'tar'`, `'rmt'`, and `'tar.info'`, but also installs them.

If you type `'make clean'`, then **make** removes the `'.'o'` files, and the `'tar'`, `'rmt'`, `'testpad'`, `'testpad.h'`, and `'core'` files.

If you type `'make distclean'`, then **make** not only removes the same files as does `'make clean'` but also the `'TAGS'`, `'Makefile'`, and `'config.status'` files. (Although it is not evident, this makefile (and `'config.status'`) is generated by the user with the `configure` program which is provided in the **tar** distribution; not shown in this documentation.)

If you type `'make realclean'`, then **make** removes the same files as does `'make distclean'` and also removes the Info files generated from `'tar.texinfo'`.

In addition, there are targets `shar` and `dist` that create distribution kits.

```
# Generated automatically from Makefile.in by configure.
# Un*x Makefile for GNU tar program.
# Copyright (C) 1991 Free Software Foundation, Inc.

# This program is free software; you can redistribute
# it and/or modify it under the terms of the GNU
# General Public License...
...
...
SHELL = /bin/sh

#### Start of system configuration section. ####

srcdir = .
# If you use gcc, you should either run the
# fixincludes script that comes with it or else use
# gcc with the -traditional option. Otherwise ioctl
# calls will be compiled incorrectly on some systems.
CC = gcc -O
YACC = bison -y
INSTALL = /usr/local/bin/install -c
INSTALLDATA = /usr/local/bin/install -c -m 644

# Things you might add to DEFS:

# -DSTDC_HEADERS      If you have ANSI C headers and
#                    libraries.
# -DPOSIX             If you have POSIX.1 headers and
#                    libraries.
# -DBSD42             If you have sys/dir.h (unless
#                    you use -DPOSIX), sys/file.h,
#                    and st_blocks in 'struct stat'.
# -DUSG               If you have System V/ANSI C
#                    string and memory functions
#                    and headers, sys/sysmacros.h,
#                    fcntl.h, getcwd, no valloc,
#                    and ndir.h (unless
#                    you use -DDIRENT).
```



```

# -DNO_MEMORY_H      If USG or STDC_HEADERS but do not
#                    include memory.h.
# -DDIRENT            If USG and you have dirent.h
#                    instead of ndir.h.
# -DSIGTYPE=int       If your signal handlers
#                    return int, not void.
# -DNO_MTIO           If you lack sys/mtio.h
#                    (magtape ioctls).
# -DNO_REMOTE         If you do not have a remote shell
#                    or rexec.
# -DUSE_REXEC         To use rexec for remote tape
#                    operations instead of
#                    forking rsh or remsh.
# -DVPRINTF_MISSING   If you lack vprintf function
#                    (but have _doprnt).
# -DDOPRNT_MISSING    If you lack _doprnt function.
#                    Also need to define
#                    -DVPRINTF_MISSING.
# -DFTIME_MISSING     If you lack ftime system call.
# -DSTRSTR_MISSING    If you lack strstr function.
# -DVALLOC_MISSING    If you lack valloc function.
# -DMKDIR_MISSING     If you lack mkdir and
#                    rmdir system calls.
# -DRENAME_MISSING    If you lack rename system call.
# -DFTRUNCATE_MISSING If you lack ftruncate
#                    system call.
# -DV7                On Version 7 Unix (not
#                    tested in a long time).
# -DEMUL_OPEN3        If you lack a 3-argument version
#                    of open, and want to emulate it
#                    with system calls you do have.
# -DNO_OPEN3          If you lack the 3-argument open
#                    and want to disable the tar -k
#                    option instead of emulating open.
# -DXENIX             If you have sys/inode.h
#                    and need it 94 to be included.
DEFS = -DSIGTYPE=int -DDIRENT -DSTRSTR_MISSING \
      -DVPRINTF_MISSING -DBSD42
# Set this to rtapelib.o unless you defined NO_REMOTE,
# in which case make it empty.

```

```
RTAPELIB = rtapelib.o
LIBS =
DEF_AR_FILE = /dev/rmt8
DEFBLOCKING = 20

CDEBUG = -g
CFLAGS = $(CDEBUG) -I. -I$(srcdir) $(DEFS) \
        -DDEF_AR_FILE=\ "$(DEF_AR_FILE)\ " \
        -DDEFBLOCKING=$(DEFBLOCKING)
LDFLAGS = -g

prefix = /usr/local
# Prefix for each installed program,
# normally empty or 'g'.
binprefix =

# The directory to install tar in.
bindir = $(prefix)/bin

# The directory to install the info files in.
infodir = $(prefix)/info

#### End of system configuration section. ####

SRC1 = tar.c create.c extract.c buffer.c \
      getoldopt.c update.c gnu.c mangle.c
SRC2 = version.c list.c names.c diffarch.c \
      port.c wildmat.c getopt.c
SRC3 = getopt1.c regex.c getdate.y
SRCS = $(SRC1) $(SRC2) $(SRC3)
OBJ1 = tar.o create.o extract.o buffer.o \
      getoldopt.o update.o gnu.o mangle.o
OBJ2 = version.o list.o names.o diffarch.o \
      port.o wildmat.o getopt.o
OBJ3 = getopt1.o regex.o getdate.o $(RTAPELIB)
OBJS = $(OBJ1) $(OBJ2) $(OBJ3)
AUX  = README COPYING ChangeLog Makefile.in \
      makefile.pc configure configure.in \
      tar.texinfo tar.info* texinfo.tex \
      tar.h port.h open3.h getopt.h regex.h \
```

```

rmt.h rmt.c rtapelib.c alloca.c \
msd_dir.h msd_dir.c tcexparg.c \
level-0 level-1 backup-specs testpad.c

all: tar rmt tar.info

tar: $(OBJS)
$(CC) $(LDFLAGS) -o $@ $(OBJS) $(LIBS)
rmt: rmt.c
$(CC) $(CFLAGS) $(LDFLAGS) -o $@ rmt.c
tar.info: tar.texinfo
makeinfo tar.texinfo

install: all
$(INSTALL) tar $(bindir)/$(binprefix)tar
-test ! -f rmt || $(INSTALL) rmt /etc/rmt
$(INSTALLDATA) $(srcdir)/tar.info* $(infodir)

$(OBJS): tar.h port.h testpad.h
regex.o buffer.o tar.o: regex.h
# getdate.y has 8 shift/reduce conflicts.

testpad.h: testpad
./testpad

testpad: testpad.o
$(CC) -o $@ testpad.o

TAGS: $(SRCS)
etags $(SRCS)

clean:
rm -f *.o tar rmt testpad testpad.h core

distclean: clean
rm -f TAGS Makefile config.status
realclean: distclean
rm -f tar.info*
shar: $(SRCS) $(AUX)
shar $(SRCS) $(AUX) | compress \

```

```

> tar-\sed -e '/version_string/!d' \
            -e 's/[^0-9.]*\([0-9.]*\).*/\1/' \
            -e q
            version.c`.shar.Z
dist: $(SRCS) $(AUX)
    echo tar-\sed \
            -e '/version_string/!d' \
            -e 's/[^0-9.]*\([0-9.]*\).*/\1/' \
            -e q
            version.c` > .fname
    -rm -rf `cat .fname`
    mkdir `cat .fname`
    ln $(SRCS) $(AUX) `cat .fname`
    -rm -rf `cat .fname` .fname
    tar chZf `cat .fname`.tar.Z `cat .fname`
tar.zoo: $(SRCS) $(AUX)
    -rm -rf tmp.dir
    -mkdir tmp.dir
    -rm tar.zoo
    for X in $(SRCS) $(AUX) ; do \
        echo $$X ; \
        sed 's/$$/^M/' $$X \
        > tmp.dir/$$X ; done
    cd tmp.dir ; zoo aM ../tar.zoo *
    -rm -rf tmp.dirIndex

```

Using diff & patch

Copyright © 1988-1998 Free Software Foundation

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions.

GNU `diff` was written by Mike Haertel, David Hayes, Richard Stallman, Len Tower, and Paul Eggert. Wayne Davison designed and implemented the unified output format.

The basic algorithm is described in “An O(ND) Difference Algorithm and its Variations” by Eugene W. Myers, in *Algorithmica*; Vol. 1, No. 2, 1986; pp. 251–266; and in “A File Comparison Program” by Webb Miller and Eugene W. Myers, in *Software—Practice and Experience*; Vol. 15, No. 11, 1985; pp. 1025–1040.

The algorithm was independently discovered as described in “Algorithms for Approximate String Matching” by E. Ukkonen, in *Information and Control*; Vol. 64, 1985, pp. 100–118.

GNU `diff3` was written by Randy Smith.

GNU `sdiff` was written by Thomas Lord.

GNU `cmp` was written by Torbjorn Granlund and David MacKenzie.

`patch` was written mainly by Larry Wall; the GNU enhancements were written mainly by Wayne Davison and David MacKenzie. Parts of the documentation are adapted from a material written by Larry Wall, with his permission.

This documentation has been prepared by Cygnus.

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

1

Overview of `diff` & `patch`, the compare & merge tools

Computer users often find occasion to ask how two files differ. Perhaps one file is a newer version of the other file. Or maybe the two files initially were identical copies that were then changed by different people.

You can use the `diff` command to show differences between two files, or each corresponding file in two directories. `diff` outputs differences between files line by line in any of several formats, selectable by command line options. This set of differences is often called a `diff` or `patch`.

The following documentation discusses using the commands and other related commands.

- “What comparison means” on page 493
- “diff output formats” on page 503
- “Comparing directories” on page 525
- “Making diff output prettier” on page 527
- “diff performance tradeoffs” on page 531
- “Comparing three files” on page 533
- “Merging from a common ancestor” on page 539
- “sdiff interactive merging” on page 547
- “Merging with the patch utility” on page 551

- “Tips for making distributions with patches” on page 561
- “Invoking the `cmp` utility” on page 563
- “Invoking the `diff` utility” on page 565
- “Invoking the `diff3` utility” on page 573
- “Invoking the `patch` utility” on page 577
- “Invoking the `sdiff` utility” on page 587
- “Incomplete lines” on page 591
- “Future projects for `diff` and `patch` utilities” on page 593

For files that are identical, `diff` normally produces no output; for binary (non-text) files, `diff` normally reports only that they are different.

You can use the `cmp` command to show the offsets and line numbers where two files differ. `cmp` can also show all the characters that differ between the two files, side by side. Another way to compare two files character by character is the Emacs command, `M-x compare-windows`. See “Comparing Files” in *The GNU Emacs Manual*¹ for more information, particularly on that command.

You can use the `diff3` command to show differences among three files. When two people have made independent changes to a common original, `diff3` can report the differences between the original and the two changed versions, and can produce a merged file that contains both persons’ changes together with warnings about conflicts.

You can use the `sdiff` command to merge two files interactively.

You can use the set of differences produced by `diff` to distribute updates to text files (such as program source code) to other people. This method is especially useful when the differences are small compared to the complete files. Given `diff` output, you can use the `patch` program to update, or `patch`, a copy of the file. If you think of `diff` as subtracting one file from another to produce their difference, you can think of `patch` as adding the difference to one file to reproduce the other.

This documentation first concentrates on making `diffs`, and later shows how to use `diffs` to update files.

¹ *The GNU Emacs Manual* is published by the Free Software Foundation (ISBN 1-882114-03-5).

2

What comparison means

There are several ways to think about the differences between two files. One way to think of the differences is as a series of lines that were deleted from, inserted in, or changed in one file to produce another file. `diff` compares two files line by line, finds groups of lines that differ, and reports each group of differing lines. It can report the differing lines in several formats, which have different purposes. See the following documentation for more information.

- “Hunks” on page 495
- “Suppressing differences in blank and tab spacing” on page 496
- “Suppressing differences in blank lines” on page 497
- “Suppressing case differences” on page 498
- “Suppressing lines matching a regular expression” on page 499
- “Summarizing which files differ” on page 500
- “Binary files and forcing text comparisons” on page 501

GNU `diff` can show whether files are different without detailing the differences. It also provides ways to suppress certain kinds of differences that are not important to you. Most commonly, such differences are changes in the amount of white space between words or lines. `diff` also provides ways to suppress differences in alphabetic case or in lines that match a regular expression that you provide. These options can accumulate; for example, you can ignore changes in both white space and alphabetic case.

Another way to think of the differences between two files is as a sequence of pairs of characters that can be either identical or different. `cmp` reports the differences between two files character by character, instead of line by line. As a result, it is more useful than `diff` for comparing binary files. For text files, `cmp` is useful mainly when you want to know only whether two files are identical.

To illustrate the effect that considering changes character by character can have compared with considering them line by line, think of what happens if a single newline character is added to the beginning of a file. If that file is then compared with an otherwise identical file that lacks the newline at the beginning, `diff` will report that a blank line has been added to the file, while `cmp` will report that almost every character of the two files differs.

`diff3` normally compares three input files line by line, finds groups of lines that differ, and reports each group of differing lines. Its output is designed to make it easy to inspect two different sets of changes to the same file.

Hunks

When comparing two files, `diff` finds sequences of lines common to both files, interspersed with groups of differing lines called *hunks*. Comparing two identical files yields one sequence of common lines and no hunks, because no lines differ. Comparing two entirely different files yields no common lines and one large hunk that contains all lines of both files. In general, there are many ways to match up lines between two given files. `diff` tries to minimize the total hunk size by finding large sequences of common lines interspersed with small hunks of differing lines. For example, suppose the file ‘F’ contains the three lines ‘a’, ‘b’, ‘c’, and the file ‘G’ contains the same three lines in reverse order ‘c’, ‘b’, ‘a’. If `diff` finds the line ‘c’ as common, then the command ‘`diff F G`’ produces the following output:

```
1,2d0
< a
< b
3a2,3
> b
> a
```

But if `diff` notices the common line ‘b’ instead, it produces the following output:

```
1c1
< a
---
> c
3c3
< c
---
> a
```

It is also possible to find ‘a’ as the common line. `diff` does not always find an optimal matching between the files; it takes shortcuts to run faster. But its output is usually close to the shortest possible. You can adjust this tradeoff with the ‘`--minimal`’ option (see “diff performance tradeoffs” on page 531).

Suppressing differences in blank and tab spacing

The `-b` and `--ignore-space-change` options ignore white space at line end, and considers all other sequences of one or more white space characters to be equivalent. With these options, `diff` considers the following two lines to be equivalent, where `$` denotes the line end:

```
Here lyeth  mucho rychnesse in lytell space. -- John Heywood$
Here lyeth mucho rychnesse in lytell space. -- John Heywood $
```

The `-w` and `--ignore-all-space` options are stronger than `-b`. They ignore difference even if one file has white space where the other file has none. *White space* characters include tab, newline, vertical tab, form feed, carriage return, and space; some locales may define additional characters to be white space. With these options, `diff` considers the following two lines to be equivalent, where `$` denotes the line end and `^M` denotes a carriage return:

```
Here lyeth mucho rychnesse in lytell space. --   John Heywood$
He relyeth much erychnes seinly tells pace. --John Heywood ^M$
```

Suppressing differences in blank lines

The `'-B'` and `'--ignore-blank-lines'` options ignore insertions or deletions of blank lines. These options normally affect only lines that are completely empty; they do not affect lines that look empty but contain space or tab characters. With these options, for instance, consider a file containing only the following.

1. A point is that which has no part.
 2. A line is breadthless length.
- Euclid, The Elements, I

The last example is considered identical to another file containing only the following.

1. A point is that which has no part.
 2. A line is breadthless length.
- Euclid, The Elements, I

Suppressing case differences

GNU `diff` can treat lowercase letters as equivalent to their uppercase counterparts, so that, for example, it considers ‘Funky Stuff’, ‘funky STUFF’, and ‘fUNKy stuFf’ to all be the same. To request this, use the ‘`-i`’ or ‘`--ignore-case`’ option.

Suppressing lines matching a regular expression

To ignore insertions and deletions of lines that match a regular expression, use the `-I regex` or `--ignore-matching-lines= regex` option. You should escape regular expressions that contain shell meta-characters to prevent the shell from expanding them.

For example, `diff -I `^[0-9]`` ignores all changes to lines beginning with a digit.

However, `-I` only ignores the insertion or deletion of lines that contain the regular expression if every changed line in the hunk—every insertion and every deletion—matches the regular expression.

In other words, for each non-ignorable change, `diff` prints the complete set of changes in its vicinity, including the ignorable ones.

You can specify more than one regular expression for lines to ignore by using more than one `-I` option. `diff` tries to match each line against each regular expression, starting with the last one given.

Summarizing which files differ

When you only want to find out whether files are different, and you don't care what the differences are, you can use the summary output format. In this format, instead of showing the differences between the files, `diff` simply reports whether files differ. The `'-q'` and `'--brief'` options select this output format.

This format is especially useful when comparing the contents of two directories. It is also much faster than doing the normal line by line comparisons, because `diff` can stop analyzing the files as soon as it knows that there are any differences.

You can also get a brief indication of whether two files differ by using `cmp`. For files that are identical, `cmp` produces no output. When the files differ, by default, `cmp` outputs the byte offset and line number where the first difference occurs. You can use the `'-s'` option to suppress that information, so that `cmp` produces no output and reports whether the files differ using only its exit status (see “Invoking the `cmp` utility” on page 563).

Unlike `diff`, `cmp` cannot compare directories; it can only compare two files.

Binary files and forcing text comparisons

If `diff` thinks that either of the two files it is comparing is binary (a non-text file), it normally treats that pair of files much as if the summary output format had been selected (see “Summarizing which files differ” on page 500), and reports only that the binary files are different. This is because line by line comparisons are usually not meaningful for binary files.

`diff` determines whether a file is text or binary by checking the first few bytes in the file; the exact number of bytes is system dependent, but it is typically several thousand. If every character in that part of the file is non-null, `diff` considers the file to be text; otherwise it considers the file to be binary.

Sometimes you might want to force `diff` to consider files to be text. For example, you might be comparing text files that contain null characters; `diff` would erroneously decide that those are non-text files. Or you might be comparing documents that are in a format used by a word processing system that uses null characters to indicate special formatting. You can force `diff` to consider all files to be text files, and compare them line by line, by using the ‘`-a`’ or ‘`--text`’ option. If the files you compare using this option do not in fact contain text, they will probably contain few newline characters, and the `diff` output will consist of hunks showing differences between long lines of whatever characters the files contain.

You can also force `diff` to consider all files to be binary files, and report only whether they differ (but not how). Use the ‘`--brief`’ option for this.

In operating systems that distinguish between text and binary files, `diff` normally reads and writes all data as text. Use the ‘`--binary`’ option to force `diff` to read and write binary data instead. This option has no effect on a Posix-compliant system like GNU or traditional Unix. However, many personal computer operating systems represent the end of a line with a carriage return followed by a newline. On such systems, `diff` normally ignores these carriage returns on input and generates them at the end of each output line, but with the ‘`--binary`’ option `diff` treats each carriage return as just another input character, and does not generate a carriage return at the end of each output line. This can be useful when dealing with non-text files that are meant to be interchanged with Posix-compliant systems.

If you want to compare two files byte by byte, you can use the `cmp` program with the ‘`-l`’ option to show the values of each differing byte in the two files. With GNU `cmp`, you can also use the ‘`-c`’ option to show the ASCII representation of those bytes. See “Invoking the `cmp` utility” on page 563 for more information.

If `diff3` thinks that any of the files it is comparing is binary (a non-text file), it normally reports an error, because such comparisons are usually not useful. `diff3` uses the same test as `diff` to decide whether a file is binary. As with `diff`, if the input files contain a few non-text characters but otherwise are like text files, you can force `diff3` to consider all files to be text files and compare them line by line by using the `'-a'` or `'--text'` options.

3

diff output formats

`diff` has several mutually exclusive options for output format. The following documentation discusses the output and formats.

- “Two sample input files” on page 504
- “Showing differences without context” on page 505
- “Showing differences in their context” on page 507
- “Showing differences side by side” on page 513
- “Controlling side by side format” on page 514
- “Merging files with if-then-else” on page 518

Two sample input files

The following are two sample files that we will use in numerous examples to illustrate the output of `diff` and how various options can change it. The following is the file, 'lao'.

```
The Way that can be told of is not the eternal Way;
The name that can be named is not the eternal name.
The Nameless is the origin of Heaven and Earth;
The Named is the mother of all things.
Therefore let there always be non-being,
    so we may see their subtlety,
And let there always be being,
    so we may see their outcome.
The two are the same,
But after they are produced,
    they have different names.
```

The following is the file, 'tzu'.

```
The Nameless is the origin of Heaven and Earth;
The named is the mother of all things.

Therefore let there always be non-being, so we may see their
    subtlety,
And let there always be being,
    so we may see their outcome.
The two are the same,
But after they are produced,
    they have different names.
They both may be called deep and profound.
Deeper and more profound,
The door of all subtleties!
```

In this example, the first hunk contains just the first two lines of 'lao', the second hunk contains the fourth line of 'lao' opposing the second and third lines of 'tzu', and the last hunk contains just the last three lines of 'tzu'.

Showing differences without context

The *normal* `diff` output format shows each hunk of differences without any surrounding context. Sometimes such output is the clearest way to see how lines have changed, without the clutter of nearby unchanged lines (although you can get similar results with the context or unified formats by using 0 lines of context). However, this format is no longer widely used for sending out patches; for that purpose, the context format (see “Context format” on page 507) and the unified format (see “Unified format” on page 509) are superior. Normal format is the default for compatibility with older versions of `diff` and the Posix standard.

Detailed description of normal format

The normal output format consists of one or more hunks of differences; each hunk shows one area where the files differ. Normal format hunks look like the following:

```
change-command
< from-file-line
< from-file-line ...
---
> to-file-line
> to-file-line ...
```

There are three types of change commands. Each consists of a line number or comma-separated range of lines in the first file, a single character indicating the kind of change to make, and a line number or comma-separated range of lines in the second file. All line numbers are the original line numbers in each file. The types of change commands are the following.

lar

Add the lines in range *r* of the second file after line *l* of the first file. For example, ‘8a12,15’ means append lines 12–15 of file 2 after line 8 of file 1; or, if changing file 2 into file 1, delete lines 12–15 of file 2.

fmt

Replace the lines in range *f* of the first file with lines in range *t* of the second file. This is like a combined add and delete, but more compact. For example, ‘5,7c8,10’ means change lines 5–7 of file 1 to read as lines 8–10 of file 2; or, if changing file 2 into file 1, change lines 8–10 of file 2 to read as lines 5–7 of file 1.

rdl

Delete the lines in range *r* from the first file; line *l* is where they would have appeared in the second file had they not been deleted. For example, ‘5,7d3’ means delete lines 5–7 of file 1; or, if changing file 2 into file 1, append lines 5–7 of file 1 after line 3 of file 2.

An example of normal format

The following is the output of the command `'diff lao tzu'` (see “Two sample input files” on page 504 for the complete contents of the two files).

Notice that the following example shows only the lines that are different between the two files.

```
1,2d0
< The Way that can be told of is not the eternal Way;
< The name that can be named is not the eternal name.
4c2,3
< The Named is the mother of all things.
---
> The named is the mother of all things.
>
11a11,13
> They both may be called deep and profound.
> Deeper and more profound,
> The door of all subtleties!
```

Showing differences in their context

Usually, when you are looking at the differences between files, you will also want to see the parts of the files near the lines that differ, to help you understand exactly what has changed. These nearby parts of the files are called the *context*.

GNU `diff` provides two output formats that show context around the differing lines: *context format* and *unified format*. It can optionally show in which function or section of the file the differing lines are found.

If you are distributing new versions of files to other people in the form of `diff` output, you should use one of the output formats that show context so that they can apply the diffs even if they have made small changes of their own to the files. `patch` can apply the diffs in this case by searching in the files for the lines of context around the differing lines; if those lines are actually a few lines away from where the `diff` says they are, `patch` can adjust the line numbers accordingly and still apply the `diff` correctly. See “Applying imperfect patches” on page 554 for more information on using `patch` to apply imperfect diffs.

Context format

The context output format shows several lines of context around the lines that differ. It is the standard format for distributing updates to source code.

To select this output format, use the ‘`-C lines`’, ‘`--context[=lines]`’, or ‘`-c`’ option. The argument *lines* that some of these options take is the number of lines of context to show. If you do not specify lines, it defaults to three. For proper operation, `patch` typically needs at least two lines of context.

Detailed description of context format

The context output format starts with a two-line header, which looks like the following.

```
*** from-file from-file-modification-time
--- to-file to-file-modification time
```

You can change the header’s content with the ‘`-L label`’ or ‘`--label=label`’ option; see “Showing alternate file names” on page 512. Next come one or more hunks of differences; each hunk shows one area where the files differ. Context format hunks look like the following.

```
*****
*** from-file-line-range ***
    from-file-line
    from-file-line ...
--- to-file-line-range ----
```

```
to-file-line
to-file-line...
```

The lines of context around the lines that differ start with two space characters. The lines that differ between the two files start with one of the following indicator characters, followed by a space character:

‘!’

A line that is part of a group of one or more lines that changed between the two files. There is a corresponding group of lines marked with ‘!’ in the part of this hunk for the other file.

‘+’

An “inserted” line in the second file that corresponds to nothing in the first file.

‘-’

A “deleted” line in the first file that corresponds to nothing in the second file.

If all of the changes in a hunk are insertions, the lines of *from-file* are omitted. If all of the changes are deletions, the lines of *to-file* are omitted.

An example of context format

Here is the output of ‘diff -c lao tzu’ (see “Two sample input files” on page 504 for the complete contents of the two files). Notice that up to three lines that are not different are shown around each line that is different; they are the context lines. Also notice that the first two hunks have run together, because their contents overlap.

```
*** lao Sat Jan 26 23:30:39 1991
--- tzu Sat Jan 26 23:30:50 1991
*****
*** 1,7 ****
- The Way that can be told of is not the eternal Way;
- The name that can be named is not the eternal name.
  The Nameless is the origin of Heaven and Earth;
! The Named is the mother of all things.
  Therefore let there always be non-being,
    so we may see their subtlety,
  And let there always be being,
--- 1,6 ----
  The Nameless is the origin of Heaven and Earth;
! The named is the mother of all things.
!
  Therefore let there always be non-being,
    so we may see their subtlety,
  And let there always be being,
*****
*** 9,11 ****
--- 8,13 ----
  The two are the same,
```



```

    But after they are produced,
        they have different names.
+ They both may be called deep and profound.
+ Deeper and more profound,
+ The door of all subtleties!

```

An example of context format with less context

The following example shows the output of `diff --context=1 lao tzu` (see “Two sample input files” on page 504 for the complete contents of the two files). Notice that at most one context line is reported here.

```

*** lao Sat Jan 26 23:30:39 1991
--- tzu Sat Jan 26 23:30:50 1991
*****
*** 1,5 ****
- The Way that can be told of is not the eternal Way;
- The name that can be named is not the eternal name.
  The Nameless is the origin of Heaven and Earth;
! The Named is the mother of all things.
  Therefore let there always be non-being,
--- 1,4 ----
  The Nameless is the origin of Heaven and Earth;
! The named is the mother of all things.
!
  Therefore let there always be non-being,
*****
*** 11 ****
--- 10,13 ----
    they have different names.
+ They both may be called deep and profound.
+ Deeper and more profound,
+ The door of all subtleties!

```

Unified format

The unified output format is a variation on the context format that is more compact because it omits redundant context lines. To select this output format, use the `‘-u lines’`, `‘--unified[=lines]’`, or `‘-u’` option. The argument *lines* is the number of lines of context to show. When it is not given, it defaults to three. At present, only GNU diff can produce this format and only GNU patch can automatically apply diffs in this format. For proper operation, patch typically needs at least two lines of context.

Detailed description of unified format

The unified output format starts with a two-line header, which looks like this:

```

--- from-file from-file-modification-time
+++ to-file to-file-modification-time

```

You can change the header's content with the `'-L label'` or `'--label=label'` option; see See "Showing alternate file names" on page 512. Next come one or more hunks of differences; each hunk shows one area where the files differ. Unified format hunks look like the following

```
@@ from-file-range to-file-range @@
line-from-either-file
line-from-either-file...
```

The lines common to both files begin with a space character. The lines that actually differ between the two files have one of the following indicator characters in the left column:

'+' A line was added here to the first file.

'-' A line was removed here from the first file.

An example of unified format

Here is the output of the command, `'diff -u lao tzu'` (see "Two sample input files" on page 504 for the complete contents of the two files):

```
--- lao Sat Jan 26 23:30:39 1991
+++ tzu Sat Jan 26 23:30:50 1991
@@ -1,7 +1,6 @@
-The Way that can be told of is not the eternal Way;
-The name that can be named is not the eternal name.
  The Nameless is the origin of Heaven and Earth;
-The Named is the mother of all things.
+The named is the mother of all things.
+
  Therefore let there always be non-being,
    so we may see their subtlety,
  And let there always be being,
@@ -9,3 +8,6 @@
  The two are the same,
  But after they are produced,
    they have different names.
+They both may be called deep and profound.
+Deeper and more profound,
+The door of all subtleties!
```

Showing which sections differences are in

Sometimes you might want to know which part of the files each change falls in. If the files are source code, this could mean which function was changed. If the files are documents, it could mean which chapter or appendix was changed. GNU `diff` can show this by displaying the nearest section heading line that precedes the differing lines. Which lines are "section headings" is determined by a regular expression.

Showing lines that match regular expressions

To show in which sections differences occur for files that are not source code for C or similar languages, use the `-F regexp` or `--show-function-line=regexp` option. `diff` considers lines that match the argument, *regexp*, to be the beginning of a section of the file. Here are suggested regular expressions for some common languages:

C, C++, Prolog

`^[A-Za-z_]`

Lisp

`^(`

Texinfo

`^@\(chapter\|appendix\|unnumbered\|chapheading\)`

This option does not automatically select an output format; in order to use it, you must select the context format (see “Context format” on page 507) or unified format (see “Unified format” on page 509). In other output formats it has no effect.

The `-F` and `--show-function-line` options find the nearest unchanged line that precedes each hunk of differences and matches the given regular expression. Then they add that line to the end of the line of asterisks in the context format, or to the `@@` line in unified format. If no matching line exists, they leave the output for that hunk unchanged. If that line is more than 40 characters long, they output only the first 40 characters. You can specify more than one regular expression for such lines; `diff` tries to match each line against each regular expression, starting with the last one given. This means that you can use `-p` and `-F` together, if you wish.

Showing C function headings

To show in which functions differences occur for C and similar languages, you can use the `-p` or `--show-c-function` option. This option automatically defaults to the context output format (see “Context format” on page 507), with the default number of lines of context. You can override that number with `-C lines` elsewhere in the command line. You can override both the format and the number with `-U lines` elsewhere in the command line.

The `-p` and `--show-c-function` options are equivalent to `-F'^[_a-zA-Z$]'` if the unified format is specified, otherwise `-C -F'^[_a-zA-Z$]'` (see “Showing lines that match regular expressions” on page 511).

GNU `diff` provides them for the sake of convenience.

Showing alternate file names

If you are comparing two files that have meaningless or uninformative names, you might want `diff` to show alternate names in the header of the context and unified output formats.

To do this, use the `'-L label'` or `'--label=label'` option. The first time you give this option, its argument replaces the name and date of the first file in the header; the second time, its argument replaces the name and date of the second file. If you give this option more than twice, `diff` reports an error. The `'-L'` option does not affect the file names in the `pr` header when the `'-l'` or `'--paginate'` option is used (see “Paginating diff output” on page 529). The following are the first two lines of the output from `'diff -C2 -Loriginal -Lmodified lao tzu'`:

```
*** original
--- modified
```

Showing differences side by side

`diff` can produce a side by side difference listing of two files. The files are listed in two columns with a gutter between them. The gutter contains one of the following markers:

white space

The corresponding lines are in common. That is, either the lines are identical, or the difference is ignored because of one of the `--ignore` options (see “Suppressing differences in blank and tab spacing” on page 496).

‘|’

The corresponding lines differ, and they are either both complete or both incomplete.

‘<’

The files differ and only the first file contains the line.

‘>’

The files differ and only the second file contains the line.

‘(’

Only the first file contains the line, but the difference is ignored.

‘)’

Only the second file contains the line, but the difference is ignored.

‘\’

The corresponding lines differ, and only the first line is incomplete.

‘/’

The corresponding lines differ, and only the second line is incomplete.

Normally, an output line is incomplete if and only if the lines that it contains are incomplete; see “Incomplete lines” on page 591. However, when an output line represents two differing lines, one might be incomplete while the other is not. In this case, the output line is complete, but its the gutter is marked ‘\’ if the first line is incomplete, ‘/’ if the second line is.

Side by side format is sometimes easiest to read, but it has limitations. It generates much wider output than usual, and truncates lines that are too long to fit. Also, it relies on lining up output more heavily than usual, so its output looks particularly bad if you use varying width fonts, nonstandard tab stops, or nonprinting characters.

You can use the `sdiff` command to interactively merge side by side differences. See “`sdiff` interactive merging” on page 547 for more information on merging files.

Controlling side by side format

The `-y` or `--side-by-side` option selects side by side format. Because side by side output lines contain two input lines, they are wider than usual. They are normally 130 columns, which can fit onto a traditional printer line. You can set the length of output lines with the `-W columns` or `--width=columns` option. The output line is split into two halves of equal length, separated by a small gutter to mark differences; the right half is aligned to a tab stop so that tabs line up. Input lines that are too long to fit in half of an output line are truncated for output. The `--left-column` option prints only the left column of two common lines. The `--suppress-common-lines` option suppresses common lines entirely.

An example of side by side format

The following is the output of the command `diff -y -W 72 lao tzu` (see “Two sample input files” on page 504 for the complete contents of the two files).

```
The Way that can be told of is <
The name that can be named is <
The Nameless is the origin of   The Nameless is the origin of
The Named is the mother of all  | The named is the mother of all
                                >

Therefore let there always be   Therefore let there always be
    so we may see their subtlet    so we may see their subtlet
And let there always be being   And let there always be being
    so we may see their outcome    so we may see their outcome
The two are the same,           The two are the same,
But after they are produced,     But after they are produced,
they have different names.       they have different names.
                                >
                                > They both may be called deep
                                > Deeper and more profound,
                                > The door of all subtleties!
```

Making edit scripts

Several output modes produce command scripts for editing *from-file* to produce *to-file*.

ed scripts

`diff` can produce commands that direct the `ed` text editor to change the first file into the second file. Long ago, this was the only output mode that was suitable for editing one file into another automatically; today, with `patch`, it is almost obsolete. Use the `-e` or `--ed` option to select this output format. Like the normal format (see “Showing differences without context” on page 505), this output format does not show any context; unlike the normal format, it does not include the information necessary to apply the diff in reverse (to produce the first file if all you have is the second file and the diff). If the file `d` contains the output of `diff -e old new`, then the command, `(cat d && echo w) | ed - old`, edits `old` to make it a copy of `new`.

More generally, if `d1`, `d2`, ..., `dN` contain the outputs of `diff -e old new1`, `diff -e new1 new2`, ..., `diff -e newN-1 newN`, respectively, then the command, `(cat d1 d2 ...dN && echo w) | ed - old`, edits `old` to make it a copy of `newN`.

Detailed description of ed format

The `ed` output format consists of one or more hunks of differences. The changes closest to the ends of the files come first so that commands that change the number of lines do not affect how `ed` interprets line numbers in succeeding commands. `ed` format hunks look like the following:

```
change-command
to-file-line
to-file-line...
```

Because `ed` uses a single period on a line to indicate the end of input,

GNU `diff` protects lines of changes that contain a single period on a line by writing two periods instead, then writing a subsequent `ed` command to change the two periods into one. The `ed` format cannot represent an incomplete line, so if the second file ends in a changed incomplete line, `diff` reports an error and then pretends that a newline was appended.

There are three types of change commands. Each consists of a line number or comma-separated range of lines in the first file and a single character indicating the kind of change to make. All line numbers are the original line numbers in the file. The types of change commands are:

`1a`

Add text from the second file after line 1 in the first file. For example, `8a` means to add the following lines after line 8 of file 1.

`'rc'`

Replace the lines in range *r* in the first file with the following lines. Like a combined add and delete, but more compact. For example, `'5,7c'` means change lines 5–7 of file 1 to read as the text file 2.

`'rd'`

Delete the lines in range *r* from the first file. For example, `'5,7d'` means delete lines 5–7 of file 1.

Example ed Script

The following is the output of `'diff -e lao tzu'` (see “Two sample input files” on page 504 for the complete contents of the two files):

```
11a
They both may be called deep and profound.
Deeper and more profound,
The door of all subtleties!
.
4c
The named is the mother of all things.
.
1,2d
```

Forward ed scripts

`diff` can produce output that is like an `ed` script, but with hunks in forward (front to back) order. The format of the commands is also changed slightly: command characters precede the lines they modify, spaces separate line numbers in ranges, and no attempt is made to disambiguate hunk lines consisting of a single period. Like `ed` format, forward `ed` format cannot represent incomplete lines. Forward `ed` format is not very useful, because neither `ed` nor `patch` can apply diffs in this format. It exists mainly for compatibility with older versions of `diff`. Use the `'-f'` or `'--forward-ed'` option to select it.

RCS scripts

The RCS output format is designed specifically for use by the Revision Control System, which is a set of free programs used for organizing different versions and systems of files. Use the `'-n'` or `'--rcs'` option to select this output format. It is like the forward `ed` format (see “Forward ed scripts” on page 516), but it can represent arbitrary changes to the contents of a file because it avoids the forward `ed` format’s problems with lines consisting of a single period and with incomplete lines. Instead of

ending text sections with a line consisting of a single period, each command specifies the number of lines it affects; a combination of the ‘a’ and ‘d’ commands are used instead of ‘c’. Also, if the second file ends in a changed incomplete line, then the output also ends in an incomplete line. The following is the output of ‘diff -n lao tzu’ (see “Two sample input files” on page 504 for the complete contents of the two files):

```
d1 2
d4 1
a4 2
The named is the mother of all things.
a11 3
They both may be called deep and profound.
Deeper and more profound,
The door of all subtleties!
```

Merging files with if-then-else

You can use `diff` to merge two files of C source code. The output of `diff` in this format contains all the lines of both files. Lines common to both files are output just once; the differing parts are separated by the C preprocessor directives, `#ifdef name` or `#ifndef name`, `#else`, and `#endif`. When compiling the output, you select which version to use by either defining or leaving undefined the macro name.

To merge two files, use `diff` with the `'-D name'` or `'--ifdef=name'` option. The argument, *name*, is the C preprocessor identifier to use in the `#ifdef` and `#ifndef` directives. For example, if you change an instance of `wait (&s)` to `waitpid (-1, &s, 0)` and then merge the old and new files with the `'--ifdef=HAVE_WAITPID'` option, then the affected part of your code might look like the following declaration.

```
do {
#ifndef HAVE_WAITPID
    if ((w = wait (&s)) < 0 && errno != EINTR)
#else /* HAVE_WAITPID */
    if ((w = waitpid (-1, &s, 0)) < 0 && errno != EINTR)
#endif /* HAVE_WAITPID */
        return w; }
while (w != child);
```

You can specify formats for languages other than C by using line group formats and line formats.

Line group formats

Line group formats let you specify formats suitable for many applications that allow if-then-else input, including programming languages and text formatting languages. A line group format specifies the output format for a contiguous group of similar lines. For example, the following command compares the TeX files `'old'` and `'new'`, and outputs a merged file in which old regions are surrounded by `'\begin{em}'`-`'\end{em}'` lines, and new regions are surrounded by `'\begin{bf}'`-`'\end{bf}'` lines.

```
diff \
    --old-group-format='\begin{em}
%<\end{em}
' \
    --new-group-format='\begin{bf}
%>\end{bf}
' \
    old new
```

The following command is equivalent to the previous example, but it is a little more verbose, because it spells out the default line group formats.

```
diff \
  --old-group-format='\begin{em}
%<\end{em}
' \
  --new-group-format='\begin{bf}
%>\end{bf} '
\
  --unchanged-group-format='%=' \
  --changed-group-format='\begin{em}
%<\end{em}
\begin{bf}
%>\end{bf}
' \
  old new
```

Here is a more advanced example, which outputs a diff listing with headers containing line numbers in a “plain English” style.

```
diff \
  --unchanged-group-format=' ' \
  --old-group-format='----- %dn line%(n=1?:s) deleted at %df:
%<' \
  --new-group-format='----- %dN line%(N=1?:s) added after %de:
%>' \
  --changed-group-format='----- %dn line%(n=1?:s) changed at %df:
%<----- to:
%>' \
  old new
```

To specify a line group format, use `diff` with one of the options listed below. You can specify up to four line group formats, one for each kind of line group. You should quote *format*, because it typically contains shell metacharacters.

`--old-group-format=format`

These line groups are hunks containing only lines from the first file. The default old group format is the same as the changed group format if it is specified; otherwise it is a format that outputs the line group as-is.

`--new-group-format=format`

These line groups are hunks containing only lines from the second file. The default new group format is same as the the changed group format if it is specified; otherwise it is a format that outputs the line group as-is.

`--changed-group-format=format`

These line groups are hunks containing lines from both files. The default changed group format is the concatenation of the old and new group formats.

`--unchanged-group-format=format`

These line groups contain lines common to both files. The default unchanged group format is a format that outputs the line group as-is.

In a line group format, ordinary characters represent themselves; conversion specifications start with `'%'` and have one of the following forms.

`'%<'`

Stands for the lines from the first file, including the trailing newline. Each line is formatted according to the old line format (see “Line formats” on page 521).

`'%>'`

Stands for the lines from the second file, including the trailing newline. Each line is formatted according to the new line format.

`'%= '`

Stands for the lines common to both files, including the trailing newline. Each line is formatted according to the unchanged line format.

`'%%'`

Stands for `'%'`.

`'%c' c'`

Where *c* is a single character, stands for *c*. *c* may not be a backslash or an apostrophe. For example, `'%c' : '` stands for a colon, even inside the ‘then-’ part of an if-then-else format, which a colon would normally terminate.

`'%c' \ o'`

Stands for the character with octal code *o*, where *o* is a string of 1, 2, or 3 octal digits. For example, `'%c' \0'` stands for a null character.

`'Fn'`

Stands for *n*’s value formatted with *F* where *F* is a `printf` conversion specification and *n* is one of the following letters.

`'e'`

The line number of the line just before the group in the old file.

`'f'`

The line number of the first line in the group in the old file; equals *e* + 1.

`'l'`

The line number of the last line in the group in the old file.

`'m'`

The line number of the line just after the group in the old file; equals *l* + 1.

`'n'`

The number of lines in the group in the old file; equals *l* - *f* + 1.

`'E, F, L, M, N'`

Likewise, for lines in the new file.

The `printf` conversion specification can be `%d`, `%o`, `%x`, or `%X`, specifying decimal, octal, lower case hexadecimal, or upper case hexadecimal output respectively. After the `%` the following options can appear in sequence: a `-` specifying left-justification; an integer specifying the minimum field width; and a period followed by an optional integer specifying the minimum number of digits. For example, `%5dN` prints the number of new lines in the group in a field of width 5 characters, using the `printf` format, `"%5d"`.

`'(A=B?T:E)'`

If *A* equals *B*, then *T*, else, *E*. *A* and *B* are each either a decimal constant or a single letter interpreted as above. This format spec is equivalent to *T* if *A*'s value equals *B*'s; otherwise it is equivalent to *E*. For example, `%(N=0?no:%dN)line%(N=1?:s)'` is equivalent to `'no lines'` if *N* (the number of lines in the group in the the new file) is 0, to `'1 line'` if *N* is 1, and to `'%dN lines'` otherwise.

Line formats

Line formats control how each line taken from an input file is output as part of a line group in if-then-else format. For example, the following command outputs text with a one-column change indicator to the left of the text. The first column of output is `-` for deleted lines, `|` for added lines, and a space for unchanged lines. The formats contain newline characters where newlines are desired on output.

```
diff \
  --old-line-format='-%l
' \
  --new-line-format='| %l
' \
  --unchanged-line-format=' %l
' \
  old new
```

To specify a line format, use one of the following options. You should quote *format*, since it often contains shell metacharacters.

`--old-line-format=format`

Formats lines just from the first file.

`--new-line-format=format`

Formats lines just from the second file.

`--unchanged-line-format=format`

Formats lines common to both files.

`--line-format=format`

Formats all lines; in effect, it simultaneously sets all three of the previous options.

In a line format, ordinary characters represent themselves; conversion specifications start with `%` and have one of the following forms.

`'%l'`

Stands for the contents of the line, not counting its trailing newline (if any). This format ignores whether the line is incomplete; see “Incomplete lines” on page 591.

`'%L'`

Stands for the contents of the line, including its trailing newline (if any). If a line is incomplete, this format preserves its incompleteness.

`'%%'`

Stands for `'%'`.

`'%C' C'`

Stands for `C`, where `C` is a single character. `C` may not be a backslash or an apostrophe. For example, `'%C' ':'` stands for a colon.

`'%C' \ o'`

Stands for the character with octal code `o` where `o` is a string of 1, 2, or 3 octal digits. For example, `'%C' \0'` stands for a null character.

`'Fn'`

Stands for the line number formatted with `F` where `F` is a `printf` conversion specification. For example, `'%.5dn'` prints the line number using the `printf` format, `"%.5d"`. See “Line group formats” on page 518 for more about `printf` conversion specifications.

The default line format is `'%l'` followed by a newline character.

If the input contains tab characters and it is important that they line up on output, you should ensure that `'%l'` or `'%L'` in a line format is just after a tab stop (e.g., by preceding `'%l'` or `'%L'` with a tab character), or you should use the `'-t'` or `'--expand-tabs'` option.

Taken together, the line and line group formats let you specify many different formats. For example, the following command uses a format similar to `diff`’s normal format. You can tailor this command to get fine control over `diff`’s output.

```
diff \
  --old-line-format='< %l
' \
  --new-line-format='> %l

' \
  --old-group-format='%df%(f=l?:,%dl)d%dE
%<' \
  --new-group-format='%dea%dF%(F=L?:,%dL)
%>' \
  --changed-group-format='%df%(f=l?:,%dl)c%dF%(F=L?:,%dL)
%<---
%>' \
```

```
--unchanged-group-format=' ' \
old new
```

Detailed description of if-then-else format

For lines common to both files, `diff` uses the unchanged line group format. For each hunk of differences in the merged output format, if the hunk contains only lines from the first file, `diff` uses the old line group format; if the hunk contains only lines from the second file, `diff` uses the new group format; otherwise, `diff` uses the changed group format.

The old, new, and unchanged line formats specify the output format of lines from the first file, lines from the second file, and lines common to both files, respectively.

The option ‘`--ifdef=name`’ is equivalent to the following sequence of options using shell syntax:

```
--old-group-format='#ifndef name %<#endif /* not name */ ' \
--new-group-format='#ifdef name %>#endif /* name */ ' \
--unchanged-group-format='%=' \ --changed-group-format='#ifndef name
%<#else /* name */ %>#endif /* name */ '
```

You should carefully check the `diff` output for proper nesting. For example, when using the the ‘`-D name`’ or ‘`--ifdef=name`’ option, you should check that if the differing lines contain any of the C preprocessor directives ‘`#ifdef`’, ‘`#ifndef`’, ‘`#else`’, ‘`#elif`’, or ‘`#endif`’, they are nested properly and match. If they don’t, you must make corrections manually. It is a good idea to carefully check the resulting code any-way to make sure that it really does what you want it to; depending on how the input files were produced, the output might contain duplicate or otherwise incorrect code.

The `patch` ‘`-D name`’ option behaves just like the `diff` ‘`-D name`’ option, except it operates on a file and a diff to produce a merged file; see “patch options” on page 583.

An example of if-then-else format

The following is the output of ‘`diff -DTWO lao tzu`’ (see “Two sample input files” on page 504 for the complete contents of the two files):

```
#ifndef TWO
The Way that can be told of is not the eternal Way;
The name that can be named is not the eternal name.
#endif /* not TWO */
The Nameless is the origin of Heaven and Earth;
#ifdef TWO
The Named is the mother of all things.
#else /* TWO */
The named is the mother of all things.
```

```
#endif /* TWO */
Therefore let there always be non-being,
so we may see their subtlety,
And let there always be being,
so we may see their outcome.
The two are the same,
But after they are produced,
they have different names.
#ifdef TWO
They both may be called deep and profound.
Deeper and more profound,
The door of all subtleties!
#endif /* TWO */
```


4

Comparing directories

You can use `diff` to compare some or all of the files in two directory trees. When both file name arguments to `diff` are directories, it compares each file that is contained in both directories, examining file names in alphabetical order. Normally `diff` is silent about pairs of files that contain no differences, but if you use the `-s` or `--report-identical-files` option, it reports pairs of identical files. Normally `diff` reports subdirectories common to both directories without comparing subdirectories' files, but if you use the `-r` or `--recursive` option, it compares every corresponding pair of files in the directory trees, as many levels deep as they go.

For file names that are in only one of the directories, `diff` normally does not show the contents of the file that exists; it reports only that the file exists in that directory and not in the other. You can make `diff` act as though the file existed but was empty in the other directory, so that it outputs the entire contents of the file that actually exists. (It is output as either an insertion or a deletion, depending on whether it is in the first or the second directory given.) To do this, use the `-N` or `--new-file` option.

If the older directory contains one or more large files that are not in the newer directory, you can make the patch smaller by using the `-P` or `--unidirectional-new-file` options instead of `-N`. This option is like `-N` except that it only inserts the contents of files that appear in the second directory but not the first (that is, files that were added). At the top of the patch, write instructions for the user applying the patch to remove the files that were deleted before applying the patch. See “Tips for making distributions with patches” on page 561 for more discussion of making patches for distribution.

To ignore some files while comparing directories, use the `-x pattern` or `--exclude=pattern` option. This option ignores any files or subdirectories whose base names match the shell pattern pattern. Unlike in the shell, a period at the start of the base of a file name matches a wildcard at the start of a pattern. You should enclose pattern in quotes so that the shell does not expand it. For example, the option `-x '*.ao'` ignores any file whose name ends with `.a` or `.o`.

This option accumulates if you specify it more than once. For example, using the options `-x 'RCS' -x '*,v'` ignores any file or subdirectory whose base name is `'RCS'` or ends with `,v`.

If you need to give this option many times, you can instead put the patterns in a file, one pattern per line, using the `-X file` or `--exclude-from=file` option.

If you have been comparing two directories and stopped partway through, later you might want to continue where you left off. You can do this by using the `-S file` or `--starting-file=file` option. This compares only the file, *file*, and all alphabetically subsequent files in the topmost directory level.

5

Making `diff` output prettier

`diff` provides several ways to adjust the appearance of its output. These adjustments can be applied to any output format. For more information, see “Preserving tabstop alignment” on page 528 and “Paginating `diff` output” on page 529.

Preserving tabstop alignment

The lines of text in some of the `diff` output formats are preceded by one or two characters that indicate whether the text is inserted, deleted, or changed. The addition of those characters can cause tabs to move to the next tabstop, throwing off the alignment of columns in the line. GNU `diff` provides two ways to make tab-aligned columns line up correctly.

The first way is to have `diff` convert all tabs into the correct number of spaces before outputting them.

Select this method with the `-t` or `--expand-tabs` option. `diff` assumes that tabstops are set every 8 columns. To use this form of output with `patch`, use the `-l` or `--ignore-white-space` option (see “Applying patches in other directories” on page 579 for more information).

The other method for making tabs line up correctly is to add a tab character instead of a space after the indicator character at the beginning of the line. This ensures that all following tab characters are in the same position relative to tabstops that they were in the original files, so that the output is aligned correctly. Its disadvantage is that it can make long lines too long to fit on one line of the screen or the paper. It also does not work with the unified output format which does not have a space character after the change type indicator character.

Select this method with the `-T` or `--initial-tab` option.

Paginating `diff` output

It can be convenient to have long output page-numbered and time-stamped. The `'-l'` and `'--paginate'` options do this by sending the `diff` output through the `pr` program. The following is what the page header might look like for `'diff -lc lao tzu'`:

```
Mar 11 13:37 1991 diff -lc lao tzu Page 1
```


6

diff performance tradeoffs

GNU `diff` runs quite efficiently; however, in some circumstances you can cause it to run faster or produce a more compact set of changes. There are two ways that you can affect the performance of GNU `diff` by changing the way it compares files.

Performance has more than one dimension. These options improve one aspect of performance at the cost of another, or they improve performance in some cases while hurting it in others.

The way that GNU `diff` determines which lines have changed always comes up with a near-minimal set of differences. Usually it is good enough for practical purposes. If the `diff` output is large, you might want `diff` to use a modified algorithm that sometimes produces a smaller set of differences. The `-d` or `--minimal` option does this; however, it can also cause `diff` to run more slowly than usual, so it is not the default behavior.

When the files you are comparing are large and have small groups of changes scattered throughout them, use the `-H` or `--speed-large-files` option to make a different modification to the algorithm that `diff` uses. If the input files have a constant small density of changes, this option speeds up the comparisons without changing the output. If not, `diff` might produce a larger set of differences; however, the output will still be correct.

Normally `diff` discards the prefix and suffix that is common to both files before it attempts to find a minimal set of differences. This makes `diff` run faster, but occasionally it may produce non-minimal output.

The ‘`--horizon-lines= lines`’ option prevents `diff` from discarding the last *lines* lines of the prefix and the first *lines* lines of the suffix. This gives `diff` further opportunities to find a minimal output.

7

Comparing three files

Use the program `diff3` to compare three files and show any differences among them. (`diff3` can also merge files; see “Merging from a common ancestor” on page 539).

The “normal” `diff3` output format shows each hunk of differences without surrounding context. Hunks are labeled depending on whether they are two-way or three-way, and lines are annotated by their location in the input files.

See “Invoking the `diff3` utility” on page 573 for more information on how to run `diff3`.

The following documentation discusses comparing files.

- “A third sample input file” on page 534
- “Detailed description of `diff3` normal format” on page 535
- “`diff3` hunks” on page 536
- “An example of `diff3` normal format” on page 537

A third sample input file

The following is a third sample file that will be used in examples to illustrate the output of `diff3` and how various options can change it. The first two files are the same that we used for `diff` (see “Two sample input files” on page 504). This is the third sample file, called ‘`tao`’:

```
The Way that can be told of is not the eternal Way;
The name that can be named is not the eternal name.
The Nameless is the origin of Heaven and Earth;
The named is the mother of all things.
Therefore let there always be non-being,
    so we may see their subtlety,
And let there always be being,
    so we may see their result.
The two are the same,
But after they are produced,
    they have different names.
```

```
-- The Way of Lao-Tzu, tr. Wing-tsit Chan
```

Detailed description of `diff3` normal format

Each hunk begins with a line marked `'===='`. Three-way hunks have plain `'===='` lines, and two-way hunks have `'1'`, `'2'`, or `'3'` appended to specify which of the three input files differ in that hunk. The hunks contain copies of two or three sets of input lines each preceded by one or two commands identifying the origin of the lines.

Normally, two spaces precede each copy of an input line to distinguish it from the commands. But with the `'-T'` or `'--initial-tab'` option, `diff3` uses a tab instead of two spaces; this lines up tabs correctly.

See “Preserving tabstop alignment” on page 528 for more information.

Commands take the following forms.

`'file: la'`

This hunk appears after line `l` of file, `file`, and contains no lines in that file. To edit this file to yield the other files, one must append hunk lines taken from the other files. For example, `'1:11a'` means that the hunk follows line 11 in the first file and contains no lines from that file.

`'file: rc'`

This hunk contains the lines in the range `r` of file `file`. The range `r` is a comma-separated pair of line numbers, or just one number if the range is a singleton. To edit this file to yield the other files, one must change the specified lines to be the lines taken from the other files. For example, `'2:11,13c'` means that the hunk contains lines 11 through 13 from the second file.

If the last line in a set of input lines is incomplete, it is distinguished on output from a full line by a following line that starts with `'\'` (see “Incomplete lines” on page 591).

diff3 hunks

Groups of lines that differ in two or three of the input files are called diff3 hunks, by analogy with diff hunks (see “Hunks” on page 495). If all three input files differ in a diff3 hunk, the hunk is called a three-way hunk; if just two input files differ, it is a two-way hunk. As with diff, several solutions are possible. When comparing the files ‘A’, ‘B’, and ‘C’, diff3 normally finds diff3 hunks by merging the two-way hunks output by the two commands ‘diff A B’ and ‘diff A C’. This does not necessarily minimize the size of the output, but exceptions should be rare. For example, suppose ‘F’ contains the three lines ‘a’, ‘b’, ‘f’, ‘G’ contains the lines ‘g’, ‘b’, ‘g’, and ‘H’ contains the lines ‘a’, ‘b’, ‘h’. ‘diff3 F G H’ might output the following:

```
====2
1:1c
3:1c
a
2:1c
g
====
1:3c
f
2:3c
g
3:3c
h
```

because it found a two-way hunk containing ‘a’ in the first and third files and ‘g’ in the second file, then the single line ‘b’ common to all three files, then a three-way hunk containing the last line of each file.

An example of `diff3` normal format

Here is the output of the command `'diff3 lao tzu tao'` (see “A third sample input file” on page 534 for the complete contents of the files). Notice that it shows only the lines that are different among the three files.

```
====2
1:1,2c
3:1,2c
    The Way that can be told of is not the eternal Way;
    The name that can be named is not the eternal name.
2:0a
====
1 1:4c
    The Named is the mother of all things.
2:2,3c
3:4,5c
    The named is the mother of all things.

====3
1:8c
2:7c
    so we may see their outcome.
3:9c
    so we may see their result.
====
1:11a
2:11,13c
    They both may be called deep and profound.
    Deeper and more profound,
    The door of all subtleties!
3:13,14c

    -- The Way of Lao-Tzu, tr. Wing-tsit Chan
```

An example of `diff3` normal format

8

Merging from a common ancestor

When two people have made changes to copies of the same file, `diff3` can produce a merged output that contains both sets of changes together with warnings about conflicts. One might imagine programs with names like `diff4` and `diff5` to compare more than three files simultaneously, but in practice the need rarely arises. You can use `diff3` to merge three or more sets of changes to a file by merging two change sets at a time.

`diff3` can incorporate changes from two modified versions into a common preceding version. This lets you merge the sets of changes represented by the two newer files. Specify the common ancestor version as the second argument and the two newer versions as the first and third arguments, like this: `diff3 mine older yours`.

You can remember the order of the arguments by noting that they are in alphabetical order.

You can think of this as subtracting *older* from yours and adding the result to *mine*, or as merging into *mine* the changes that would turn *older* into *yours*. This merging is well-defined as long as *mine* and *older* match in the neighborhood of each such change. This fails to be true when all three input files differ or when only *older* differs; we call this a *conflict*. When all three input files differ, we call the conflict an *overlap*.

`diff3` gives you several ways to handle overlaps and conflicts. You can omit overlaps or conflicts, or select only overlaps, or mark conflicts with special ‘<<<<<<’ and ‘>>>>>>’ lines.

`diff3` can output the merge results as an `ed` script that that can be applied to the first file to yield the merged output. However, it is usually better to have `diff3` generate the merged output directly; this bypasses some problems with `ed`.

Selecting which changes to incorporate

You can select all unmerged changes from *old* to *yours* for merging into *mine* with the ‘-e’ or ‘--ed’ option. You can select only the nonoverlapping unmerged changes with ‘-3’ or ‘--easy-only’, and you can select only the overlapping changes with ‘-x’ or ‘--overlap-only’.

The ‘-e’, ‘-3’ and ‘-x’ options select only unmerged changes, i.e., changes where *mine* and *yours* differ; they ignore changes from *older* to *yours* where *mine* and *yours* are identical, because they assume that such changes have already been merged. If this assumption is not a safe one, you can use the options, ‘-A’ or ‘--show-all’ (see “Marking conflicts” on page 542). The following is the output of the command, `diff3`, with each of these three options (see “A third sample input file” on page 534 for the complete contents of the files). Notice that ‘-e’ outputs the union of the disjoint sets of changes output by ‘-3’ and ‘-x’.

Output of ‘`diff3 -e lao tzu tao`’:

```
11a
.
-- The Way of Lao-Tzu, tr. Wing-tsit Chan
.
8c
    so we may see their result.
.
```

Output of ‘`diff3 -3 lao tzu tao`’:

```
8c
    so we may see their result.
.
```

Output of ‘`diff3 -x lao tzu tao`’:

```
11a
.
-- The Way of Lao-Tzu, tr. Wing-tsit Chan
.
```

Marking conflicts

`diff3` can mark conflicts in the merged output by bracketing them with special marker lines. A conflict that comes from two files *A* and *B* is marked as follows:

```
<<<<<< A
lines from A
=====
lines from B
>>>>>> B
```

A conflict that comes from three files, *A*, *B* and *C*, is marked as follows:

```
<<<<<< A
lines from A
||||||| B
lines from B
=====
lines from C
>>>>>> C
```

The ‘`-A`’ or ‘`--show-all`’ option acts like the ‘`-e`’ option, except that it brackets conflicts, and it outputs all changes from *older* to *yours*, not just the unmerged changes. Thus, given the sample input files (see “A third sample input file” on page 534), ‘`diff3 -A lao tzu tao`’ puts brackets around the conflict where only ‘*tzu*’ differs:

```
<<<<<< tzu
=====
The Way that can be told of is not the eternal Way;
The name that can be named is not the eternal name.
>>>>>> tao
```

And it outputs the three-way conflict as follows:

```
<<<<<< lao
||||||| tzu
They both may be called deep and profound.
Deeper and more profound,
The door of all subtleties!
=====

-- The Way of Lao-Tzu, tr. Wing-tsit Chan
>>>>>> tao
```

The `-E` or `--show-overlap` option outputs less information than the `-A` or `--show-all` option, because it outputs only unmerged changes, and it never outputs the contents of the second file. Thus the `-E` option acts like the `-e` option, except that it brackets the first and third files from three-way overlapping changes. Similarly, `-X` acts like `-x`, except it brackets all its (necessarily overlapping) changes. For example, for the three-way overlapping change above, the `-E` and `-X` options output the following:

```
<<<<<< lao
=====

-- The Way of Lao-Tzu, tr. Wing-tsit Chan
>>>>>> tao
```

If you are comparing files that have meaningless or uninformative names, you can use the `-L label` or `--label= label` option to show alternate names in the `<<<<<<`, `|||||` and `>>>>>>` brackets. This option can be given up to three times, once for each input file. Thus `diff3 -A -L X -L Y -L Z A B C` acts like `diff3 -A A B C`, except that the output looks like it came from files named `'X'`, `'Y'` and `'Z'` rather than from files named `'A'`, `'B'` and `'C'`.

Generating the merged output directly

With the ‘-m’ or ‘--merge’ option, `diff3` outputs the merged file directly. This is more efficient than using `ed` to generate it, and works even with non-text files that `ed` would reject. If you specify ‘-m’ without an `ed` script option, ‘-A’ (‘--show-all’) is assumed.

For example, the command ‘`diff3 -m lao tzu tao`’ (see “A third sample input file” on page 534 for a copy of the input files) would output the following:

```
<<<<<< tzu
=====
The Way that can be told of is not the eternal Way;
The name that can be named is not the eternal name.
>>>>>> tao
The Nameless is the origin of Heaven and Earth;
The Named is the mother of all things.
Therefore let there always be non-being,
    so we may see their subtlety,
And let there always be being,
    so we may see their result.
The two are the same,
But after they are produced,
    they have different names.
<<<<<< lao
||||||| tzu
They both may be called deep and profound.
Deeper and more profound,
The door of all subtleties!
=====

-- The Way of Lao-Tzu, tr. Wing-tsit Chan
>>>>>> tao
```

How `diff3` merges incomplete lines

With `-m`, incomplete lines (see “Incomplete lines” on page 591) are simply copied to the output as they are found; if the merged output ends in an conflict and one of the input files ends in an incomplete line, succeeding `|||||`, `=====` or `>>>>>>` brackets appear somewhere other than the start of a line because they are appended to the incomplete line.

Without `-m`, if an `ed` script option is specified and an incomplete line is found, `diff3` generates a warning and acts as if a newline had been present.

Saving the changed file

Traditional Unix `diff3` generates an `ed` script without the trailing `'w'` and `'q'` commands that save the changes. System V `diff3` generates these extra commands. GNU `diff3` normally behaves like traditional Unix `diff3`, but with the `'-i'` option it behaves like System V `diff3` and appends the `'w'` and `'q'` commands.

The `'-i'` option requires one of the `ed` script options `'-AeExX3'`, and is incompatible with the merged output option `'-m'`.

9

`sdiff` interactive merging

With `sdiff`, you can merge two files interactively based on a side-by-side format comparison with ‘-y’ (see “Showing differences side by side” on page 513). Use ‘-o *file*’ or ‘--output=*file*’ to specify where to put the merged text. See “Invoking the `sdiff` utility” on page 587 for more details on the options to `sdiff`.

Another way to merge files interactively is to use the Emacs Lisp package, `emerge`. See “Merging Files with Emerge” in *The GNU Emacs Manual*[†] for more information.

[†] The GNU Emacs Manual is published by the Free Software Foundation (ISBN 1-882114-03-5).

Specifying `diff` options to the `sdiff` utility

The following `sdiff` options have the same meaning as for `diff`. See “diff options” on page 566 for the use of these options.

```
-a -b -d -i -t -v
-B -H -I regexp

--ignore-blank-lines --ignore-case
--ignore-matching-lines= regexp --ignore-space-change
--left-column --minimal --speed-large-files
--suppress-common-lines --expand-tabs
--text --version --width= columns
```

For historical reasons, `sdiff` has alternate names for some options. The ‘-l’ option is equivalent to the ‘--left-column’ option, and similarly ‘-s’ is equivalent to ‘--suppress-common-lines’. The meaning of the `sdiff` ‘-w’ and ‘-W’ options is interchanged from that of `diff`: with `sdiff`, ‘-w *columns*’ is equivalent to ‘--width=*columns*’, and ‘-W’ is equivalent to ‘--ignore-all-space’. `sdiff` without the ‘-o’ option is equivalent to `diff` with the ‘-y’ or ‘--side-by-side’ option (see “Showing differences side by side” on page 513).

Merge commands

Groups of common lines, with a blank gutter, are copied from the first file to the output. After each group of differing lines, `sdiff` prompts with ‘%’ and pauses, waiting for one of the following commands. Follow each command using the Enter key.

‘e’

Discard both versions. Invoke a text editor on an empty temporary file, then copy the resulting file to the output.

‘eb’

Concatenate the two versions, edit the result in a temporary file, then copy the edited result to the output.

‘el’

Edit a copy of the left version, then copy the result to the output.

‘er’

Edit a copy of the right version, then copy the result to the output.

‘l’

Copy the left version to the output.

‘q’

Quit.

‘r’

Copy the right version to the output.

‘s’

Silently copy common lines.

‘v’

Verbosely copy common lines. This is the default.

The text editor invoked is specified by the `EDITOR` environment variable if it is set. The default is system-dependent.

10

Merging with the `patch` utility

`patch` takes comparison output produced by `diff` and applies the differences to a copy of the original file, producing a patched version. With `patch`, you can distribute just the changes to a set of files instead of distributing the entire file set; your correspondents can apply `patch` to update their copy of the files with your changes. `patch` automatically determines the diff format, skips any leading or trailing headers, and uses the headers to determine which file to patch. This lets your correspondents feed an article or message containing a difference listing directly to `patch`.

`patch` detects and warns about common problems like forward patches. It saves the original version of the files it patches, and saves any patches that it could not apply. It can also maintain a `patchlevel.h` file to ensure that your correspondents apply diffs in the proper order.

`patch` accepts a series of diffs in its standard input, usually separated by headers that specify which file to patch. It applies `diff` hunks (see “Hunks” on page 495) one by one. If a hunk does not exactly match the original file, `patch` uses heuristics to try to patch the file as well as it can. If no approximate match can be found, `patch` rejects the hunk and skips to the next hunk. `patch` normally replaces each file, `f`, with its new version, saving the original file in `f.orig`, and putting reject hunks (if any) into `f.rej`.

See “Invoking the patch utility” on page 577 for detailed information on the options to patch. See “Backup file names” on page 580 for more information on how patch names backup files. See “Naming reject files” on page 582 for more information on where patch puts reject hunks.

Selecting the `patch` input format

`patch` normally determines which `diff` format the patch file uses by examining its contents. For patch files that contain particularly confusing leading text, you might need to use one of the following options to force `patch` to interpret the `patch` file as a certain format of `diff`. The output formats shown in the following discussion are the only ones that `patch` can understand.

```
'-c'
'--context'
    context diff.
'-e'
'--ed'
    ed script.
'-n'
'--normal'
    normal diff.
'-u'
'--unified'
    unified diff.
```

Applying imperfect patches

`patch` tries to skip any leading text in the `patch` file, apply the diff, and then skip any trailing text. Thus you can feed a news article or mail message directly to `patch`, and it should work. If the entire diff is indented by a constant amount of white space, `patch` automatically ignores the indentation. However, certain other types of imperfect input require user intervention.

Applying patches with changed white space

Sometimes mailers, editors, or other programs change spaces into tabs, or vice versa. If this happens to a patch file or an input file, the files might look the same, but `patch` will not be able to match them properly. If this problem occurs, use the `‘-l’` or `‘--ignore-white-space’` option, which makes `patch` compare white space loosely so that any sequence of white space in the patch file matches any sequence of whitespace in the input files. Non-whitespace characters must still match exactly. Each line of the context must still match a line in the input file.

Applying reversed patches

Sometimes people run `diff` with the new file first instead of second. This creates a diff that is “reversed”. To apply such patches, give `patch` the `‘-R’` or `‘--reverse’` option. `patch` then attempts to swap each hunk around before applying it. Rejects come out in the swapped format. The `‘-R’` option does not work with `ed` scripts because there is too little information in them to reconstruct the reverse operation. Often `patch` can guess that the patch is reversed. If the first hunk of a patch fails, `patch` reverses the hunk to see if it can apply it that way. If it can, `patch` asks you if you want to have the `‘-R’` option set; if it can’t, `patch` continues to apply the patch normally. This method cannot detect a reversed patch if it is a normal diff and the first command is an append (which should have been a delete) since appends always succeed, because a null context matches anywhere. But most patches add or change lines rather than delete them, so most reversed normal diffs begin with a delete, which fails, and `patch` notices.

If you apply a patch that you have already applied, `patch` thinks it is a reversed patch and offers to un-apply the patch. This could be construed as a feature. If you did this inadvertently and you don’t want to un-apply the patch, just answer `‘n’` to this offer and to the subsequent “apply anyway” question—or use the keystroke sequence, **C-C**, to kill the `patch` process.

Helping `patch` find inexact matches

For context diffs, and to a lesser extent normal diffs, `patch` can detect when the line numbers mentioned in the patch are incorrect, and it attempts to find the correct place to apply each hunk of the patch. As a first guess, it takes the line number mentioned in the hunk, plus or minus any offset used in applying the previous hunk. If that is not the correct place, `patch` scans both forward and backward for a set of lines matching the context given in the hunk.

First `patch` looks for a place where all lines of the context match. If it cannot find such a place, and it is reading a context or unified diff, and the maximum fuzz factor is set to 1 or more, then `patch` makes another scan, ignoring the first and last line of context. If that fails, and the maximum fuzz factor is set to 2 or more, it makes another scan, ignoring the first two and last two lines of context are ignored. It continues similarly if the maximum fuzz factor is larger.

The `'-F lines'` or `'--fuzz=lines'` option sets the maximum fuzz factor to *lines*. This option only applies to context and unified diffs; it ignores up to *lines* lines while looking for the place to install a hunk. Note that a larger fuzz factor increases the odds of making a faulty patch. The default fuzz factor is 2; it may not be set to more than the number of lines of context in the diff, ordinarily 3.

If `patch` cannot find a place to install a hunk of the patch, it writes the hunk out to a reject file (see “Naming reject files” on page 582 for information on how reject files are named). It writes out rejected hunks in context format no matter what form the input patch is in. If the input is a normal or `ed` diff, many of the contexts are simply null. The line numbers on the hunks in the reject file may be different from those in the patch file; they show the approximate location where `patch` thinks the failed hunks belong in the new file rather than in the old one.

As it completes each hunk, `patch` tells you whether the hunk succeeded or failed, and if it failed, on which line (in the new file) `patch` thinks the hunk should go. If this is different from the line number specified in the diff, it tells you the offset. A single large offset may indicate that `patch` installed a hunk in the wrong place. `patch` also tells you if it used a fuzz factor to make the match, in which case you should also be slightly suspicious.

`patch` cannot tell if the line numbers are off in an `ed` script, and can only detect wrong line numbers in a normal `diff` when it finds a change or delete command. It may have the same problem with a context diff using a fuzz factor equal to or greater than the number of lines of context shown in the diff (typically 3). In these cases, you should probably look at a context diff between your original and patched input files to see if the changes make sense. Compiling without errors is a pretty good indication that the patch worked, but not a guarantee.

`patch` usually produces the correct results, even when it must make many guesses. However, the results are guaranteed only when the patch is applied to an exact copy of the file that the patch was generated from.

Removing empty files

Sometimes when comparing two directories, the first directory contains a file that the second directory does not. If you give `diff` the `-N` or `--new-file` option, it outputs a diff that deletes the contents of this file. By default, `patch` leaves an empty file after applying such a diff. The `-E` or `--remove-empty-files` option to `patch` deletes output files that are empty after applying the diff.

Multiple patches in a file

If the patch file contains more than one patch, patch tries to apply each of them as if they came from separate patch files. This means that it determines the name of the file to patch for each patch, and that it examines the leading text before each patch for file names and prerequisite revision level (see “Tips for making distributions with patches” on page 561 for more on that topic). For the second and subsequent patches in the patch file, you can give options and another original file name by separating their argument lists with a ‘+’. However, the argument list for a second or subsequent patch may not specify a new patch file, since that does not make sense. For example, to tell `patch` to strip the first three slashes from the name of the first patch in the patch file and none from subsequent patches, and to use ‘`code.c`’ as the first input file, you can use: `patch -p3 code.c + -p0 < patchfile`.

The ‘`-s`’ or ‘`--skip`’ option ignores the current patch from the patch file, but continue looking for the next patch in the file. Thus, to ignore the first and third patches in the patch file, you can use: `patch -S + + -S + < patch file`.

Messages and questions from the `patch` utility

`patch` can produce a variety of messages, especially if it has trouble decoding its input. In a few situations where it's not sure how to proceed, `patch` normally prompts you for more information from the keyboard. There are options to suppress printing non-fatal messages and stopping for keyboard input. The message 'Hmm...' indicates that `patch` is reading text in the `patch` file, attempting to determine whether there is a patch in that text, and if so, what kind of patch it is. You can inhibit all terminal output from `patch`, unless an error occurs, by using the `-s`, `--quiet`, or `--silent` option. There are two ways you can prevent `patch` from asking you any questions. The `-f` or `--force` option assumes that you know what you are doing. It assumes the following:

- skip patches that do not contain file names in their headers;
- patch files even though they have the wrong version for the 'Prereq:' line in the patch;
- assume that patches are not reversed even if they look like they are.

The `-t` or `--batch` option is similar to `-f`, in that it suppresses questions, but it makes somewhat different assumptions:

- skip patches that do not contain file names in their headers (the same as `-f`);
- skip patches for which the file has the wrong version for the 'Prereq:' line in the patch;
- assume that patches are reversed if they look like they are.

`patch` exits with a non-zero status if it creates any reject files. When applying a set of patches in a loop, you should check the exit status, so you don't apply a later patch to a partially patched file.

11

Tips for making distributions with patches

The following discussions detail some things you should keep in mind if you are going to distribute patches for updating a software package.

Make sure you have specified the file names correctly, either in a context `diff` header or with an `'Index:'` line. If you are patching files in a subdirectory, be sure to tell the patch user to specify a `'-p'` or `'--strip'` option as needed. Take care to not send out reversed patches, since these make people wonder whether they have already applied the patch.

To save people from partially applying a patch before other patches that should have gone before it, you can make the first patch in the patch file update a file with a name like `'patchlevel.h'` or `'version.c'`, which contains a patch level or version number. If the input file contains the wrong version number, patch will complain immediately.

An even clearer way to prevent this problem is to put a `'Prereq:'` line before the patch. If the leading text in the patch file contains a line that starts with `'Prereq:'`, patch takes the next word from that line (normally a version number) and checks whether the next input file contains that word, preceded and followed by either white space or a newline. If not, patch prompts you for confirmation before proceeding. This makes it difficult to accidentally apply patches in the wrong order.

Since patch does not handle incomplete lines properly, make sure that all the source files in your program end with a newline whenever you release a version.

To create a patch that changes an older version of a package into a newer version, first make a copy of the older version in a scratch directory. Typically you do that by unpacking a `tar` or `shar` archive of the older version.

You might be able to reduce the size of the patch by renaming or removing some files before making the patch. If the older version of the package contains any files that the newer version does not, or if any files have been renamed between the two versions, make a list of `rm` and `mv` commands for the user to execute in the old version directory before applying the patch. Then run those commands yourself in the scratch directory.

If there are any files that you don't need to include in the patch because they can easily be rebuilt from other files (for example, `TAGS` and output from `yacc` and `makeinfo`), replace the versions in the scratch directory with the newer versions, using `rm` and `ln` or `cp`.

Now you can create the patch. The de-facto standard `diff` format for patch distributions is context format with two lines of context, produced by giving `diff` the `'-c 2'` option. Do not use less than two lines of context, because patch typically needs at least two lines for proper operation.

Give `diff` the `'-p'` option in case the newer version of the package contains any files that the older one does not. Make sure to specify the scratch directory first and the newer directory second.

Add to the top of the patch a note telling the user any `rm` and `mv` commands to run before applying the patch. Then you can remove the scratch directory.

12

Invoking the `cmp` utility

The `cmp` command compares two files, and if they differ, tells the first byte and line number where they differ.

Its arguments are: `cmp options ... from-file [to-file]`.

The file name `-` is always the standard input. `cmp` also uses the standard input if one file name is omitted.

An exit status of 0 means no differences were found, 1 means some differences were found, and 2 means trouble.

cmp options

The following is a summary of all of the options that GNU `cmp` accepts. Most options have two equivalent names, one of which is a single letter preceded by `'-'`, and the other of which is a long name preceded by `--`. Multiple single letter options (unless they take an argument) can be combined into a single command line word: `'-cl'` is equivalent to `'-c -l'`.

`'-c'`

Print the differing characters. Display control characters as a `'^'` followed by a letter of the alphabet and precede characters that have the high bit set with `'M-'` (which stands for “meta”).

`--ignore-initial=bytes`

Ignore any differences in the the first *bytes* bytes of the input files. Treat files with fewer than *bytes* bytes as if they are empty.

`'-l'`

Print the (decimal) offsets and (octal) values of all differing bytes.

`--print-chars`

Print the differing characters. Display control characters as a `'^'` followed by a letter of the alphabet and precede characters that have the high bit set with `'M-'` (which stands for “meta”).

`--quiet`

`'-s'`

`--silent`

Do not print anything; return exit status indicating whether files differ.

`--verbose`

Print the (decimal) offsets and (octal) values of all differing bytes.

`'-v'`

`--version`

Output the version number of `cmp`.

13

Invoking the `diff` utility

`diff options ...from-file to-file` is the format for running the `diff` command.

In the simplest case, `diff` compares the contents of the two files *from-file* and *to-file*. A file name of `-` stands for text read from the standard input. As a special case, `diff - -` compares a copy of standard input to itself. If *from-file* is a directory and *to-file* is not, `diff` compares the file in *from-file* whose file name is that of *to-file*, and vice versa. The non-directory file must not be `-`. If both *from-file* and *to-file* are directories, `diff` compares corresponding files in both directories, in alphabetical order; this comparison is not recursive unless the `-r` or `--recursive` option is given. `diff` never compares the actual contents of a directory as if it were a file. The file that is fully specified may not be standard input, because standard input is nameless and the notion of “file with the same name” does not apply.

`diff` options begin with `-`, so normally *from-file* and *to-file* may not begin with `-`. However, `--` as an argument by itself treats the remaining arguments as file names even if they begin with `-`.

An exit status of 0 means no differences were found, 1 means some differences were found, and 2 means trouble.

diff options

The following is a summary of all of the options that GNU `diff` accepts. Most options have two equivalent names, one of which is a single letter preceded by ‘-’, and the other of which is a long name preceded by ‘--’.

Multiple single letter options (unless they take an argument) can be combined into a single command line word: ‘-ac’ is equivalent to ‘-a -c’.

Long named options can be abbreviated to any unique prefix of their name.

Brackets ([and]) indicate that an option takes an optional argument.

‘-lines’

Show *lines* (an integer) lines of context. This option does not specify an output format by itself; it has no effect unless it is combined with ‘-c’ (see “Context format” on page 507) or ‘-u’ (see “Unified format” on page 509). This option is obsolete. For proper operation, `patch` typically needs at least two lines of context.

‘-a’

Treat all files as text and compare them line-by-line, even if they do not seem to be text. See “Binary files and forcing text comparisons” on page 501.

‘-b’

Ignore changes in amount of white space. See “Suppressing differences in blank and tab spacing” on page 496.

‘-B’

Ignore changes that just insert or delete blank lines. See “Suppressing differences in blank lines” on page 497.

‘--binary’

Read and write data in binary mode. See “Binary files and forcing text comparisons” on page 501.

‘--brief’

Report only whether the files differ, not the details of the differences. See “Summarizing which files differ” on page 500.

‘-c’

Use the context output format. See “Context format” on page 507.

‘-C lines’

‘--context[=*lines*]

Use the context output format, showing *lines* (an integer) lines of context, or three if *lines* is not given. See “Context format” on page 507. For proper operation, `patch` typically needs at least two lines of context.

-
- ‘--changed-group-format=*format*’
Use *format* to output a line group containing differing lines from both files in if-then-else format. See “Line group formats” on page 518.
 - ‘-d’
Change the algorithm perhaps find a smaller set of changes. This makes `diff` slower (sometimes much slower). See “diff performance tradeoffs” on page 531.
 - ‘-D *name*’
Make merged ‘`#ifdef`’ format output, conditional on the pre-processor macro *name*. See “Merging files with if-then-else” on page 518.
 - ‘-e’
 - ‘--ed’
Make output that is a valid `ed` script. See “ed scripts” on page 515.
 - ‘--exclude=*pattern*’
When comparing directories, ignore files and subdirectories whose basenames match *pattern*. See “Comparing directories” on page 525.
 - ‘--exclude-from=*file*’
When comparing directories, ignore files and subdirectories whose basenames match any pattern contained in *file*. See “Comparing directories” on page 525.
 - ‘--expand-tabs’
Expand tabs to spaces in the output, to preserve the alignment of tabs in the input files. See “Preserving tabstop alignment” on page 528.
 - ‘-f’
Make output that looks vaguely like an `ed` script but has changes in the order they appear in the file. See “Forward ed scripts” on page 516.
 - ‘-F *regexp*’
In context and unified format, for each hunk of differences, show some of the last preceding line that matches *regexp*. See “Suppressing lines matching a regular expression” on page 499.
 - ‘--forward-ed’
Make output that looks vaguely like an `ed` script but has changes in the order they appear in the file. See “Forward ed scripts” on page 516.
 - ‘-h’
This option currently has no effect; it is present for Unix compatibility.
 - ‘-H’
Use heuristics to speed handling of large files that have numerous scattered small changes. See “diff performance tradeoffs” on page 531.
 - ‘--horizon-lines=*lines*’
Do not discard the last *lines* lines of the common prefix and the first *lines* lines of the common suffix. See “diff performance tradeoffs” on page 531.

- `'-i'`
Ignore changes in case; consider uppercase and lowercase letters equivalent. See “Suppressing case differences” on page 498.
- `'-I regexp'`
Ignore changes that just insert or delete lines that match *regexp*. See “Suppressing lines matching a regular expression” on page 499.
- `'--ifdef=name'`
Make merged if-then-else output using *name*. See “Merging files with if-then-else” on page 518.
- `'--ignore-all-space'`
Ignore white space when comparing lines. See “Suppressing differences in blank and tab spacing” on page 496.
- `'--ignore-blank-lines'`
Ignore changes that just insert or delete blank lines. See “Suppressing differences in blank lines” on page 497.
- `'--ignore-case'`
Ignore changes in case; consider upper- and lower-case to be the same. See “Suppressing case differences” on page 498.
- `'--ignore-matching-lines=regexp'`
Ignore changes that just insert or delete lines that match *regexp*. See “Suppressing lines matching a regular expression” on page 499.
- `'--ignore-space-change'`
Ignore changes in amount of white space. See “Suppressing differences in blank and tab spacing” on page 496.
- `'--initial-tab'`
Output a tab rather than a space before the text of a line in normal or context format. This causes the alignment of tabs in the line to look normal. See “Preserving tabstop alignment” on page 528
- `'-l'`
Pass the output through `pr` to paginate it. See “Paginating diff output” on page 529.
- `'-L label'`
Use *label* instead of the file name in the context format (see “Detailed description of context format” on page 507) and unified format (see “Detailed description of unified format” on page 509) headers. See “RCS scripts” on page 516.

-
- ‘--label=*label*’
Use *label* instead of the file name in the context format (see “Detailed description of context format” on page 507) and unified format (see “Detailed description of unified format” on page 509) headers.
 - ‘--left-column’
Print only the left column of two common lines in side by side format. See “Controlling side by side format” on page 514.
 - ‘--line-format=*format*’
Use *format* to output all input lines in if-then-else format. See “Line formats” on page 521.
 - ‘--minimal’
Change the algorithm to perhaps find a smaller set of changes. This makes `diff` slower (sometimes much slower). See “diff performance tradeoffs” on page 531.
 - ‘-n’
Output RCS-format diffs; like ‘-f’ except that each command specifies the number of lines affected. See “RCS scripts” on page 516.
 - ‘-N’
 - ‘--new-file’
In directory comparison, if a file is found in only one directory, treat it as present but empty in the other directory. See “Comparing directories” on page 525.
 - ‘--new-group-format=*format*’
Use *format* to output a group of lines taken from just the second file in if-then-else format. See “Line group formats” on page 518.
 - ‘--new-line-format=*format*’
Use *format* to output a line taken from just the second file in if-then-else format. See “Line formats” on page 521.
 - ‘--old-group-format=*format*’
Use *format* to output a group of lines taken from just the first file in if-then-else format. See “Line group formats” on page 518.
 - ‘--old-line-format=*format*’
Use *format* to output a line taken from just the first file in if-then-else format. See “Line formats” on page 521.
 - ‘-p’
Show which C function each change is in. See “Showing C function headings” on page 511.
 - ‘-P’
When comparing directories, if a file appears only in the second directory of the two, treat it as present but empty in the other. See “Comparing directories” on page 525.

- `--paginate`
Pass the output through `pr` to paginate it. See “Paginating diff output” on page 529.
- `-q`
Report only whether the files differ, not the details of the differences. See “Summarizing which files differ” on page 500.
- `-r`
When comparing directories, recursively compare any sub-directories found. See “Comparing directories” on page 525.
- `--rcs`
Output RCS-format diffs; like `-f` except that each command specifies the number of lines affected. See “RCS scripts” on page 516.
- `--recursive`
When comparing directories, recursively compare any sub-directories found. See “Comparing directories” on page 525.
- `--report-identical-files`
Report when two files are the same. See “Comparing directories” on page 525.
- `-s`
Report when two files are the same. See “Comparing directories” on page 525.
- `-S file`
When comparing directories, start with the file, *file*. This is used for resuming an aborted comparison. See “Comparing directories” on page 525.
- `--sdiff-merge-assist`
Print extra information to help `sdiff`. `sdiff` uses this option when it runs `diff`. This option is not intended for users to use directly.
- `--show-c-function`
Show which C function each change is in. See “Showing C function headings” on page 511.
- `--show-function-line=regex`
In context and unified format, for each hunk of differences, show some of the last preceding line that matches *regex*. See “Suppressing lines matching a regular expression” on page 499.
- `--side-by-side`
Use the side by side output format. See “Controlling side by side format” on page 514.
- `--speed-large-files`
Use heuristics to speed handling of large files that have numerous scattered small changes. See “diff performance tradeoffs” on page 531.

-
- ‘--starting-file=*file*’
When comparing directories, start with the file, *file*. This is used for resuming an aborted comparison. See “Comparing directories” on page 525.
 - ‘--suppress-common-lines’
Do not print common lines in side by side format. See “Controlling side by side format” on page 514.
 - ‘-t’
Expand tabs to spaces in the output, to preserve the alignment of tabs in the input files. See “Preserving tabstop alignment” on page 528.
 - ‘-T’
Output a tab rather than a space before the text of a line in normal or context format. This causes the alignment of tabs in the line to look normal. See “Preserving tabstop alignment” on page 528.
 - ‘--text’
Treat all files as text and compare them line-by-line, even if they do not appear to be text. See “Binary files and forcing text comparisons” on page 501.
 - ‘-u’
Use the unified output format. See “Unified format” on page 509.
 - ‘--unchanged-group-format=*format*’
Use *format* to output a group of common lines taken from both files in if-then-else format. See “Line group formats” on page 518.
 - ‘--unchanged-line-format=*format*’
Use *format* to output a line common to both files in if-then-else format. See “Line formats” on page 521.
 - ‘--unidirectional-new-file’
When comparing directories, if a file appears only in the second directory of the two, treat it as present but empty in the other. See “Comparing directories” on page 525.
 - ‘-U *lines*’
‘--unified[=*lines*]’
Use the unified output format, showing *lines* (an integer) lines of context, or three if *lines* is not given. See “Unified format” on page 509. For proper operation, patch typically needs at least two lines of context.
 - ‘-v’
‘--version’
Output the version number of diff.
 - ‘-w’
Ignore white space when comparing lines. See “Suppressing differences in blank and tab spacing” on page 496.

`'-W columns'`

`'--width=columns'`

Use an output width of *columns* in side by side format. See “Controlling side by side format” on page 514.

`'-x pattern'`

When comparing directories, ignore files and subdirectories whose basenames match *pattern*. See “Comparing directories” on page 525.

`'-X file'`

When comparing directories, ignore files and subdirectories whose basenames match any pattern contained in *file*. See “Comparing directories” on page 525.

`'-y'`

Use the side by side output format. See “Controlling side by side format” on page 514.

14

Invoking the `diff3` utility

The `diff3` command compares three files and outputs descriptions of their differences.

Its arguments are as follows: `diff3 options ...mine older yours`.

The files to compare are *mine*, *older*, and *yours*. At most one of these three file names may be '-', which tells `diff3` to read the standard input for that file. An exit status of 0 means `diff3` was successful, 1 means some conflicts were found, and 2 means trouble.

diff3 options

The following is a summary of all of the options that GNU `diff3` accepts. Multiple single letter options (unless they take an argument) can be combined into a single command line argument.

‘-a’

Treat all files as text and compare them line-by-line, even if they do not appear to be text. See “Binary files and forcing text comparisons” on page 501.

‘-A’

Incorporate all changes from older to yours into mine, sur-rounding all conflicts with bracket lines. See “Marking conflicts” on page 542.

‘-e’

Generate an `ed` script that incorporates all the changes from older to yours into mine. See “Selecting which changes to incorporate” on page 541.

‘-E’

Like ‘-e’, except bracket lines from overlapping changes’ first and third files. See “Marking conflicts” on page 542. With ‘-e’, an overlapping change looks like this:

```
<<<<<<< mine
lines from mine
=====
lines from yours
>>>>>>> yours
```

‘--ed’

Generate an `ed` script that incorporates all the changes from older to yours into mine. See “Selecting which changes to incorporate” on page 541.

‘--easy-only’

Like ‘-e’, except output only the non-overlapping changes. See “Selecting which changes to incorporate” on page 541.

‘-i’

Generate ‘w’ and ‘q’ commands at the end of the `ed` script for System V compatibility. This option must be combined with one of the ‘-AeExX3’ options, and may not be combined with ‘-m’. See “Saving the changed file” on page 546.

‘--initial-tab’

Output a tab rather than two spaces before the text of a line in normal format. This causes the alignment of tabs in the line to look normal. See “Preserving tabstop alignment” on page 528.

`'-L label'`

`'--label=label'`

Use the label, *label*, for the brackets output by the `'-A'`, `'-E'` and `'-X'` options. This option may be given up to three times, one for each input file. The default labels are the names of the input files. Thus `'diff3 -L X -L Y -L Z -m A B C'` acts like `'diff3 -m A B C'`, except that the output looks like it came from files named `'x'`, `'y'` and `'z'` rather than from files named `'A'`, `'B'` and `'C'`. See “Marking conflicts” on page 542.

`'-m'`

`'--merge'`

Apply the edit script to the first file and send the result to standard output. Unlike piping the output from `diff3` to `ed`, this works even for binary files and incomplete lines. `'-A'` is assumed if no edit script option is specified. See “Generating the merged output directly” on page 544.

`'--overlap-only'`

Like `'-e'`, except output only the overlapping changes. See “Selecting which changes to incorporate” on page 541.

`'--show-all'`

Incorporate all unmerged changes from *older* to *yours* into *mine*, surrounding all overlapping changes with bracket lines. See “Marking conflicts” on page 542.

`'--show-overlap'`

Like `'-e'`, except bracket lines from overlapping changes' first and third files. See “Marking conflicts” on page 542.

`'-T'`

Output a tab rather than two spaces before the text of a line in normal format. This causes the alignment of tabs in the line to look normal. See “Preserving tabstop alignment” on page 528.

`'--text'`

Treat all files as text and compare them line-by-line, even if they do not appear to be text. See “Binary files and forcing text comparisons” on page 501.

`'-v'`

`'--version'`

Output the version number of `diff3`.

`'-x'`

Like `'-e'`, except output only the overlapping changes. See “Selecting which changes to incorporate” on page 541.

`'-X'`

Like `'-E'`, except output only the overlapping changes. In other words, like `'-x'`, except bracket changes as in `'-E'`. See “Marking conflicts” on page 542.

‘-3’

Like ‘-e’, except output only the nonoverlapping changes. See “Selecting which changes to incorporate” on page 541.

15

Invoking the `patch` utility

Normally `patch` is invoked using: `patch < patchfile`.

The full format for invoking `patch` is a declaration like the following example.

```
patch options...[origfile [patchfile]] [+ options ...[origfile]]...
```

If you do not specify `patchfile`, or if `patchfile` is '-', `patch` reads the patch (that is, the `diff` output) from the standard input.

You can specify one or more of the original files as `orig` arguments; each one and options for interpreting it is separated from the others with a '+'. See "Multiple patches in a file" on page 558 for more information.

If you do not specify an input file on the command line, `patch` tries to figure out from the leading text (any text in the patch that comes before the `diff` output) which file to edit. In the header of a context or unified `diff`, `patch` looks in lines beginning with '***', '---', or '+++'; among those, it chooses the shortest name of an existing file. Otherwise, if there is an 'Index:' line in the leading text, `patch` tries to use the file name from that line. If `patch` cannot figure out the name of an existing file from the leading text, it prompts you for the name of the file to patch.

If the input file does not exist or is read-only, and a suitable RCS or SCCS file exists, `patch` attempts to check out or get the file before proceeding. By default, `patch` replaces the original input file with the patched version, after renaming the original file into a backup file (see “Backup File Names” on page Backup file names for a description of how `patch` names backup files). You can also specify where to put the output with the `-o output-file` or `--output= output-file` option.

Applying patches in other directories

The `'-d directory'` or `'--directory=directory'` option to `patch` makes `directory` the current directory for interpreting both file names in the patch file, and file names given as arguments to other options (such as `'-B'` and `'-o'`). For example, while in a news reading program, you can patch a file in the `'/usr/src/emacs'` directory directly from the article containing the patch like the following example:

```
| patch -d /usr/src/emacs
```

Sometimes the file names given in a patch contain leading directories, but you keep your files in a directory different from the one given in the patch. In those cases, you can use the `'-p[number]'` or `'--strip=[number]'` option to set the file name strip count to *number*. The strip count tells `patch` how many slashes, along with the directory names between them, to strip from the front of file names. `'-p'` with no number given is equivalent to `'-p0'`. By default, `patch` strips off all leading directories, leaving just the base file names, except that when a file name given in the patch is a relative file name and all of its leading directories already exist, `patch` does not strip off the leading directory. (A relative file name is one that does not start with a slash.)

`patch` looks for each file (after any slashes have been stripped) in the current directory, or if you used the `'-d directory'` option, in that directory. For example, suppose the file name in the patch file is `'/gnu/src/emacs/etc/new'`. Using `'-p'` or `'-p0'` gives the entire file name unmodified, `'-p1'` gives `'gnu/src/emacs/etc/new'` (no leading slash), `'-p4'` gives `'etc/news'`, and not specifying `'-p'` at all gives `'news'`.

Backup file names

Normally, patch renames an original input file into a backup file by appending to its name the extension `.orig`, or `~` on systems that do not support long file names. The `-b backup-suffix` or `--suffix=backup-suffix` option uses *backup-suffix* as the backup extension instead. Alternately, you can specify the extension for backup files with the `SIMPLE_BACKUP_SUFFIX` environment variable, which the options over-ride.

patch can also create numbered backup files the way GNU Emacs does. With this method, instead of having a single backup of each file, patch makes a new backup file name each time it patches a file. For example, the backups of a file named `'sink'` would be called, successively, `'sink.~1~'`, `'sink.~2~'`, `'sink.~3~'`, etc. The `-v backup-style` or `--version-control=backup-style` option takes as an argument a method for creating backup file names. You can alternately control the type of backups that patch makes with the `VERSION_CONTROL` environment variable, which the `-v` option overrides. The value of the `VERSION_CONTROL` environment variable and the argument to the `-v` option are like the GNU Emacs version-control variable (see “Transposing Text” in *The GNU Emacs Manual*[†], for more information on backup versions in Emacs). They also recognize synonyms that are more descriptive. The valid values are listed in the following; unique abbreviations are acceptable.

`'t'`

`'numbered'`

Always make numbered backups.

`'nil'`

`'existing'`

Make numbered backups of files that already have them, simple backups of the others. This is the default.

`'never'`

`'simple'`

Always make simple backups.

Alternately, you can tell patch to prepend a prefix, such as a directory name, to produce backup file names.

The `-B backup-prefix` or `--prefix=backup-prefix` option makes backup files by prepending *backup-prefix* to them. If you use this option, patch ignores any `-b` option that you give.

[†] The Free Software Foundation publishes *The GNU Emacs Manual* (ISBN 1-882114-03-5).

If the backup file already exists, `patch` creates a new backup file name by changing the first lowercase letter in the last component of the file name into uppercase. If there are no more lowercase letters in the name, it removes the first character from the name. It repeats this process until it comes up with a backup file name that does not already exist.

If you specify the output file with the `‘-o’` option, that file is the one that is backed up, not the input file.

Naming reject files

The names for reject files (files containing patches that `patch` could not find a place to apply) are normally the name of the output file with `.rej` appended (or `#` on systems that do not support long file names).

Alternatively, you can tell `patch` to place all of the rejected patches in a single file. The `-r reject-file` or `--reject-file= reject-file` option uses *reject-file* as the reject file name.

patch options

The following summarizes the options that `patch` accepts. Older versions of `patch` do not accept long-named options or the `-t`, `-E`, or `-V` options.

Multiple single-letter options that do not take an argument can be combined into a single command line argument (with only one dash). Brackets (⌈ and ⌋) indicate that an option takes an optional argument.

`-b backup-suffix`

Use *backup-suffix* as the backup extension instead of `.orig` or `~`. See “Backup file names” on page 580.

`-B backup-prefix`

Use *backup-prefix* as a prefix to the backup file name. If this option is specified, any `-b` option is ignored. See “Backup file names” on page 580.

`--batch`

Do not ask any questions. See “Messages and questions from the patch utility” on page 559.

`-c`

`--context`

Interpret the `patch` file as a context diff. See “Selecting the patch input format” on page 553.

`-d directory`

`--directory= directory`

Makes directory *directory* the current directory for interpreting both file names in the patch file, and file names given as arguments to other options. See “Applying patches in other directories” on page 579.

`-D name`

Make merged if-then-else output using format. See “Merging files with if-then-else” on page 518.

`--debug=number`

Set internal debugging flags. Of interest only to `patch` patchers.

`-e`

`--ed`

Interpret the patch file as an `ed` script. See “Selecting the patch input format” on page 553.

`-E`

Remove output files that are empty after the patches have been applied. See “Removing empty files” on page 557.

- `'-f'`
Assume that the user knows exactly what he or she is doing, and ask no questions. See “Messages and questions from the patch utility” on page 559.
- `'-F lines'`
Set the maximum fuzz factor to *lines*. See “Helping patch find inexact matches” on page 555.
- `'--force'`
Assume that the user knows exactly what he or she is doing, and ask no questions. See “Messages and questions from the patch utility” on page 559.
- `'--forward'`
Ignore patches that `patch` thinks are reversed or already applied. See also `'-R'`. See “Applying reversed patches” on page 554.
- `'--fuzz=lines'`
Set the maximum fuzz factor to *lines*. See “Helping patch find inexact matches” on page 555.
- `'--help'`
Print a summary of the options that `patch` recognizes, then exit.
- `'--ifdef= name'`
Make merged if-then-else output using format. See “Merging files with if-then-else” on page 518.
- `'--ignore-white-space'`
- `'-l'`
Let any sequence of white space in the patch file match any sequence of white space in the input file. See “Applying patches with changed white space” on page 554.
- `'-n'`
- `'--normal'`
Interpret the `patch` file as a normal diff. See “Selecting the patch input format” on page 553.
- `'-N'`
Ignore patches that `patch` thinks are reversed or already applied. See also `'-R'`. See “Applying reversed patches” on page 554.
- `'-o output-file'`
- `'--output=output-file'`
Use *output-file* as the output file name.
- `'-p[number]'`
Set the file name strip count to *number*. See “Applying patches in other directories” on page 579.

`--prefix=backup-prefix`

Use *backup-prefix* as a prefix to the backup file name. If this option is specified, any `-b` option is ignored. See “Backup file names” on page 580.

`--quiet`

Work silently unless an error occurs. See “Messages and questions from the patch utility” on page 559.

`-r reject-file`

Use *reject-file* as the reject file name. See “Naming reject files” on page 582.

`-R`

Assume that this patch was created with the old and new files swapped. See “Applying reversed patches” on page 554.

`--reject-file=reject-file`

Use *reject-file* as the reject file name. See “Naming reject files” on page 582.

`--remove-empty-files`

Remove output files that are empty after the patches have been applied. See “Removing empty files” on page 557.

`--reverse`

Assume that this patch was created with the old and new files swapped. See “Applying reversed patches” on page 554.

`-s`

Work silently unless an error occurs. See “Messages and questions from the patch utility” on page 559.

`-S`

Ignore this patch from the patch file, but continue looking for the next patch in the file. See “Multiple patches in a file” on page 558.

`--silent`

Work silently unless an error occurs. See “Messages and questions from the patch utility” on page 559.

`--skip`

Ignore this patch from the *patch* file, but continue looking for the next patch in the file. See “Multiple patches in a file” on page 558.

`--strip[=number]`

Set the file name strip count to *number*. See “Applying patches in other directories” on page 579.

`--suffix=backup-suffix`

Use *backup-suffix* as the backup extension instead of `.orig` or `~`. “Backup file names” on page 580.

- '-t'
Do not ask any questions. See “Messages and questions from the patch utility” on page 559.
- '-u'
- '--unified'
Interpret the patch file as a unified diff. See “Selecting the patch input format” on page 553.
- '-v'
Output the revision header and patch level of patch.
- '-v *backup-style*'
Select the kind of backups to make. See “Backup file names” on page 580.
- '--version'
Output the revision header and patch level of patch, then exit.
- '--version=control=*backup-style*'
Select the kind of backups to make. See “Backup file names” on page 580.
- '-x *number*'
Set internal debugging flags. Of interest only to patch patchers.

16

Invoking the `sdiff` utility

The `sdiff` command merges two files and interactively outputs the results.

Its arguments are: `sdiff -o outfile options ...from-file to-file`.

This merges *from-file* with *to-file*, with output to *outfile*. If *from-file* is a directory and *to-file* is not, `sdiff` compares the file in *from-file* whose file name is that of *to-file*, and vice versa. *from-file* and *to-file* may not both be directories.

`sdiff` options begin with ‘-’, so normally *from-file* and *to-file* may not begin with ‘-’. However, ‘--’ as an argument by itself treats the remaining arguments as file names even if they begin with ‘-’. You may not use ‘-’ as an input file. An exit status of 0 means no differences were found, 1 means some differences were found, and 2 means trouble.

`sdiff` without ‘-o’ (or ‘--output’) produces a side-by-side difference. This usage is obsolete; use ‘diff --side-by-side’ instead.

sdiff options

The following is a summary of all of the options that GNU sdiff accepts. Each option has two equivalent names, one of which is a single letter preceded by '-', and the other of which is a long name preceded by '--'. Multiple single letter options (unless they take an argument) can be combined into a single command line argument. Long named options can be abbreviated to any unique prefix of their name.

'-a'

Treat all files as text and compare them line-by-line, even if they do not appear to be text. See "Binary files and forcing text comparisons" on page 501.

'-b'

Ignore changes in amount of whitespace. See "Suppressing differences in blank and tab spacing" on page 496.

'-B'

Ignore changes that just insert or delete blank lines. See "Suppressing differences in blank lines" on page 497.

'-d'

Change the algorithm to perhaps find a smaller set of changes. This makes sdiff slower (sometimes much slower). See "diff performance tradeoffs" on page 531.

'-H'

Use heuristics to speed handling of large files that have numerous scattered small changes. See "diff performance tradeoffs" on page 531.

'--expand-tabs'

Expand tabs to spaces in the output, to preserve the alignment of tabs in the input files. See "Preserving tabstop alignment" on page 528.

'-i'

Ignore changes in case; consider uppercase and lowercase to be the same. See "Suppressing case differences" on page 498.

'-I *regexp*'

Ignore changes that just insert or delete lines that match *regexp*. See "Suppressing lines matching a regular expression" on page 499.

'--ignore-all-space'
Ignore white space when comparing lines. See "Suppressing differences in blank and tab spacing" on page 496.

'--ignore-blank-lines'

Ignore changes that just insert or delete blank lines. See "Suppressing differences in blank lines" on page 497.

`--ignore-case`
 Ignore changes in case; consider uppercase and lowercase to be the same. See “Suppressing case differences” on page 498.

`--ignore-matching-lines=regex`
 Ignore changes that just insert or delete lines that match *regex*. See “Suppressing lines matching a regular expression” on page 499.

`--ignore-space-change`
 Ignore changes in amount of whitespace. See “Suppressing differences in blank and tab spacing” on page 496.

`-l`
`--left-column`
 Print only the left column of two common lines. See “Controlling side by side format” on page 514.

`--minimal`
 Change the algorithm to perhaps find a smaller set of changes. This makes *sdiff* slower (sometimes much slower). See “diff performance tradeoffs” on page 531.

`-o file`
`--output=file`
 Put merged output into *file*. This option is required for merging.

`-s`
`--suppress-common-lines`
 Do not print common lines. See “Controlling side by side format” on page 514.

`--speed-large-files`
 Use heuristics to speed handling of large files that have numerous scattered small changes. See “diff performance tradeoffs” on page 531.

`-t`
 Expand tabs to spaces in the output, to preserve the alignment of tabs in the input files. See “Preserving tabstop alignment” on page 528.

`--text`
 Treat all files as text and compare them line-by-line, even if they do not appear to be text. See “Binary files and forcing text comparisons” on page 501.

`-v`
`--version`
 Output the version number of *sdiff*.

`-w columns`
`--width=columns`
 Use an output width of *columns*. See “Controlling side by side format” on page 514. For historical reasons, this option is `-w` in *indiff*, `-w` in *sdiff*.

‘-w’

Ignore horizontal white space when comparing lines. See “Suppressing differences in blank and tab spacing” on page 496. For historical reasons, this option is ‘-w’ in `diff`, ‘-w’ in `sdiff`.

Incomplete lines

When an input file ends in a non-newline character, its last line is called an incomplete line because its last character is not a newline. All other lines are called full lines and end in a newline character. Incomplete lines do not match full lines unless differences in white space are ignored (see “Suppressing differences in blank and tab spacing” on page 496).

An incomplete line is normally distinguished on output from a full line by a following line that starts with ‘\’. However, the RCS format (see “RCS scripts” on page 516) outputs the incomplete line as-is, without any trailing newline or following line. The side by side format normally represents incomplete lines as-is, but in some cases uses a ‘\’ or ‘/’ gutter marker; See “Controlling side by side format” on page 514. The if-then-else line format preserves a line’s incompleteness with ‘%L’, and discards the new-line with ‘%1’; see “Line formats” on page 521. Finally, with the ed and forward ed output formats (see “diff output formats” on page 503) diff cannot represent an incomplete line, so it pretends there was a newline and reports an error. For example, suppose ‘f’ and ‘g’ are one-byte files that contain just ‘f’ and ‘g’, respectively. Then ‘diff f g’ outputs

```
1c1
< f
\ No newline at end of file
---
> g
\ No newline at end of file
```

Incomplete lines

(The exact message may differ in non-English locales.) `'diff -n f g'` outputs the following without a trailing newline:

```
d1 1
a1 1
g
```

`'diff -e f g'` reports two errors and outputs the following:

```
1c
g
.
```

18

Future projects for `diff` and `patch` utilities

The following discussions have some ideas for improving GNU `diff` and `patch`. The GNU project has identified some improvements as potential programming projects for volunteers. You can also help by reporting any bugs that you find. If you are a programmer and would like to contribute something to the GNU project, please consider volunteering for one of these projects. If you are seriously contemplating work, please write to 'gnu@prep.ai.mit.edu' to coordinate with other volunteers.

18: Future projects for `diff` and
`patch` utilities

Suggested projects for improving GNU `diff` and `patch` utilities

One should be able to use GNU `diff` to generate a patch from any pair of directory trees, and given the patch and a copy of one such tree, use `patch` to generate a faithful copy of the other. Unfortunately, some changes to directory trees cannot be expressed using current patch formats; also, `patch` does not handle some of the existing formats. These shortcomings motivate the following suggested projects.

Handling changes to the directory structure

`diff` and `patch` do not handle some changes to directory structure. For example, suppose one directory tree contains a directory named ‘`D`’ with some subsidiary files, and another contains a file with the same name ‘`D`’. ‘`diff -r`’ does not output enough information for `patch` to transform the the directory subtree into the file. There should be a way to specify that a file has been deleted without having to include its entire contents in the patch file. There should also be a way to tell `patch` that a file was renamed, even if there is no way for `diff` to generate such information. These problems can be fixed by extending the `diff` output format to represent changes in directory structure, and extending `patch` to understand these extensions.

Files that are neither directories nor regular files

Some files are neither directories nor regular files: they are unusual files like symbolic links, device special files, named pipes, and sockets. Currently, `diff` treats symbolic links like regular files; it treats other special files like regular files if they are specified at the top level, but simply reports their presence when comparing directories. This means that `patch` cannot represent changes to such files. For example, if you change which file a symbolic link points to, `diff` outputs the difference between the two files, instead of the change to the symbolic link.

`diff` should optionally report changes to special files specially, and `patch` should be extended to understand these extensions.

File names that contain unusual characters

When a file name contains an unusual character like a newline or white space, ‘`diff -r`’ generates a patch that `patch` cannot parse. The problem is with format of `diff` output, not just with `patch`, because with odd enough file names one can cause `diff` to generate a patch that is syntactically correct but patches the wrong files. The format of `diff` output should be extended to handle all possible file names.

Arbitrary limits

GNU `diff` can analyze files with arbitrarily long lines and files that end in incomplete lines. However, `patch` cannot patch such files. The `patch` internal limits on line lengths should be removed, and `patch` should be extended to parse `diff` reports of incomplete lines.

Handling files that do not fit in memory

`diff` operates by reading both files into memory. This method fails if the files are too large, and `diff` should have a fallback.

One way to do this is to scan the files sequentially to compute hash codes of the lines and put the lines in equivalence classes based only on hash code. Then compare the files normally. This does produce some false matches.

Then scan the two files sequentially again, checking each match to see whether it is real. When a match is not real, mark both the “matching” lines as changed. Then build an edit script as usual.

The output routines would have to be changed to scan the files sequentially looking for the text to print.

Ignoring certain changes

It would be nice to have a feature for specifying two strings, one in from-file and one in to-file, which should be considered to match. Thus, if the two strings are ‘foo’ and ‘bar’, then if two lines differ only in that ‘foo’ in file 1 corresponds to ‘bar’ in file 2, the lines are treated as identical. It is not clear how general this feature can or should be, or what syntax should be used for it.

Reporting bugs

If you think you have found a bug in GNU `cmp`, `diff`, `diff3`, `sdiff`, or `patch`, report it by electronic mail to ‘bug-gnu-utils@prep.ai.mit.edu’. Send as precise a description of the problem as you can, including sample input files that produce the bug, if applicable.

Using info

Copyright © 1988-1998 Free Software Foundation

Permission is granted to make and distribute verbatim copies of this documentation provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this documentation under the conditions for verbatim copying, provided also that the “GNU General Public License” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this documentation into another language, under the above conditions for modified versions.

This document was written by Brian J. Fox: bfox@ai.mit.edu

This documentation has been prepared by Cygnus.

All rights reserved.

GNUPro™, the GNUPro™ logo, and the Cygnus logo are trademarks of Cygnus. All other brand and product names are trademarks of their respective owners.

To contact the Cygnus Technical Publications staff, email: doc@cygnus.com.

1

Overview of `info`, the GNU online documentation

The GNU `info` program is online documentation used to view help files on an ASCII terminal. `info` files are the result of processing Texinfo files with the program, `Makeinfo`, using the Emacs editing command: `M-x texinfo-format-buffer`.

Texinfo is a documentation language allowing printed and online documentation (an `info` file) to be produced from a single source file. The following documentation discusses `info` in more detail.

- “Using the `info` program” on page 600
- “Reading `info` files” on page 601
- “Making `info` files from Texinfo files” on page 623

Using the `info` program

`info` is organized into *nodes*, corresponding to the chapters and sections of printed books. You can follow them in sequence just like in the printed book or, using menus, go quickly to the node having the information you need.

`info` has *hot* references; if one section refers to another, you can tell `info` to take you immediately to that other section. You can get back again easily to take up your reading where you left off. Naturally, you can also search for particular words and phrases.

The best way to get started with the online documentation system is to use a programmed tutorial by running `info` itself. You run `info` by just typing its name—no options or arguments are necessary—at your shell prompt (shown here as ‘#’), as in the following example.

```
# info
```

`info` displays its first screen, a menu of available documentation, and waits. To request a tutorial for learning `info`, type `h`. Or, from Emacs document browsing mode, in the window’s command buffer, type `C-h, i`.

You can exit Info any time by typing `q`.

`info` also displays a summary of all its commands when you type `?`.

2

Reading `info` files

You can read the documentation for GNU software either on paper, as with any other instruction manual, or as online `info` files, using an ordinary ASCII terminal.

You can browse through the online documentation with GNU Emacs. The program, `info`, is a small program intended just for the purpose of viewing help files.

`info` files are generated by the program, `makeinfo`, from a Texinfo source file.

Texinfo is a documentation markup language designed to allow the same source file to generate either printed or online documentation. The Texinfo language is described in *Texinfo (The GNU Documentation Format)*[†].

[†] *Texinfo (The GNU Documentation Format)* is published by the Free Software Foundation (ISBN 1-882114-12-4).

GNU info command line options

GNU `info` accepts several options to control the initial node being viewed, and to specify which directories to search for `info` files. The following is a template showing an invocation of GNU `info` from the shell.

```
info [-- option-name option-value] menu-item ...
```

The following *option-names* are available when invoking `info` from the shell.

`--directory directory-path`

`-d directory-path`

Adds *directory-path* to the list of directory paths searched when Info needs to find a file. You may issue `--directory` multiple times; once for each directory which contains `info` files.

Alternatively, you may specify a value for the environment variable, `INFOPATH`; if `--directory` is not given, the value of `INFOPATH` is used. The value of `INFOPATH` is a colon separated list of directory names. If you do not supply `INFOPATH` or `--directory-path`, a default path is used.

`--file filename`

`-f filename`

Specifies a particular `info` file to visit. Instead of visiting the file, `dir`, `info` will start with (*filename*) `Top` as the first file and node.

`--node nodename`

`-n nodename`

Specifies a particular node to visit in the initial file loaded. This is especially useful in conjunction with `--file`[†].

You can specify `--node` multiple times. For an interactive Info session, each *nodename* is visited in its own window. For a non-interactive Info (such as when `--output` is given), each *node-name* is processed sequentially.

`--output filename`

`-o filename`

Specify *filename* as the name of a file to output to. Each node that Info visits will be output to *filename* instead of interactively viewed. A value of `'-'` for *filename* specifies the standard output.

[†] Of course, you can specify both the file and node in a `--node` command; but don't forget to escape the open and close parentheses from the shell as in: `info --node '(emacs)Buffers'`

--subnodes

This option only has meaning given in conjunction with `--output`. It means to recursively output the nodes appearing in the menus of each node being output. Menu items which resolve to external info files are not output, and neither are menu items which are members of an index. Each node is only output once.

--help**-h**

Produces a relatively brief description of the available Info options.

--version

Prints the version information of `info` and exits.

menu-item

Remaining arguments to `info` are treated as the names of menu items. The first argument would be a menu item in the initial node visited, while the second argument would be a menu item in the first argument's node. You can easily move to the node of your choice by specifying the menu names which describe the path to that node, as in the following example.

info emacs buffers

This input example selects the menu item 'Emacs' in the node '(dir)Top', and then selects the menu item 'Buffers' in the node '(emacs)Top'.

Moving the cursor

Many people find that reading screens of text page by page is made easier when one is able to indicate particular pieces of text with some kind of pointing device. Since this is the case, GNU `info` (both the Emacs and standalone versions) have several commands which let you move the cursor about the screen. The notation in this documentation to describe keystrokes is identical to the notation used within the Emacs manual, and the GNU Readline manual. See “Characters, Keys and Commands” in *GNU Emacs Manual*^d, if you are unfamiliar with the notation.

The following table lists the basic `info` cursor movement commands.

Each entry consists of the key sequence to type to perform the cursor movement: a command like `M-x\`, and a short description of what it does.

Cursor motion also uses a *numeric* argument; for further discussion, see “Miscellaneous commands” on page 617. A numeric argument executes a command that many times; for example, `4` given to `next-line` moves the cursor *down* 4 lines.

A negative numeric argument reverses the motion; thus, an argument of `-4` for the `next-line` command moves the cursor *up* 4 lines.

`C-n` (`next-line`)

Moves the cursor down to the next line.

`C-p` (`prev-line`)

Move the cursor up to the previous line.

`C-a` (`beginning-of-line`)

Move the cursor to the start of the current line.

`C-e` (`end-of-line`)

Moves the cursor to the end of the current line.

`C-f` (`forward-char`)

Move the cursor forward a character.

`C-b` (`backward-char`)

Move the cursor backward a character.

`M-f` (`forward-word`)

Moves the cursor forward a word.

`M-b` (`backward-word`)

Moves the cursor backward a word.

`M-<` (`beginning-of-node`)

^d *GNU Emacs Manual* is published by the Free Software Foundation (ISBN 1-882114-03-5).

[\] `M-x` is also a command; it invokes `execute-extended-command`. See “Keyboard Input” in the *GNU Emacs Manual* for more detailed information.

b

Moves the cursor to the start of the current node.

M-> (*end-of-node*)

Moves the cursor to the end of the current node.

M-r (*move-to-window-line*)

Moves the cursor to a specific line of the window. Without a numeric argument, **M-r** moves the cursor to the start of the line in the center of the window. With a numeric argument of *n*, **M-r** moves the cursor to the start of the *n*th line in the window.

Moving text within a window

Sometimes you are looking at a full screen of text, and only part of the current paragraph you are reading is visible on the screen. The commands detailed in the following section are used to shift which part of the current node is visible on the screen.

SPC / Spacebar (`scroll-forward`)

C-v

Shifts the text in this window up to show more of the node which is currently below the bottom of the window. A numeric argument shows that many more lines at the bottom of the window; a numeric argument of 4 shifts all text in the window up 4 lines (discarding the top 4 lines), and shows four new lines at the bottom of the window. Without a numeric argument, Spacebar takes the bottom two lines of the window and places them at the top of the window, redisplaying almost a completely new screenful of lines.

Del /Delete (`scroll-backward`)

M-v

Shifts the text in this window down, the inverse of scroll-forward.

`scroll-forward` and `scroll-backward` moves forward and backward through the node structure of the file. If you press Spacebar while viewing the end of a node, or Del while viewing the beginning of a node, what happens is controlled by the variable `scroll-behavior`. See “Manipulating Variables,” page Manipulating variables for more information.

C-l (`redraw-display`)

Redraws the display from scratch, or shifts the line with the cursor to a specified location. With no numeric argument, ‘C-l’ clears the screen, and then redraws its entire contents. With a numeric argument of *n*, the line with the cursor is shifted to the *n*th line of the window.

C-x w (`toggle-wrap`)

Toggles the state of line wrapping in the current window. Normally, when lines *wrap* when they are longer than the screen width,; i.e., they continue on the next line. Lines which wrap display a ‘\’ in the rightmost column of the screen. You can cause such lines to be terminated at the rightmost column by changing the state of line wrapping in the window with C-x, w. When a line contains more than one screen width, ‘\$’ appears in the rightmost column of the line, and the remainder is invisible.

Selecting a new node

The following documentation details the numerous `info` commands which select a new node to view in the current window.

The most basic node commands are ‘n’, ‘p’, ‘u’, and ‘l’.

When you are viewing a node, the top line of the node contains some *info pointers* which describe where the *next*, *previous*, and *up* nodes are. `info` uses this line to move about the node structure of the file when you type the following commands.

`n` (`next-node`)

Selects the `Next` node.

`p` (`prev-node`)

Selects the `Prev` (previous) node.

`u` (`up-node`)

Selects the `Up` node.

You can easily select a node that you have already viewed in this window by using the ‘l’ command (this name stands for *last*), to actually move through the list of already visited nodes for this window. ‘l’ with a negative numeric argument moves forward through the history of nodes for this window, so you can quickly step between two adjacent (in viewing history) nodes.

`l` (`history-node`)

Selects the most recently selected node in this window.

Two additional commands make it easy to select the most commonly selected nodes; they are ‘t’ and ‘d’.

`t` (`top-node`)

Selects the node `Top` in the current `info` file.

`d` (`dir-node`)

Selects the directory node (i.e., the node, `(dir)`).

The following are some other commands which immediately result in the selection of a different node in the current window.

`<` (`first-node`)

Selects the first node which appears in this file. This node is most often ‘Top’, but it doesn’t have to be.

`>` (`last-node`)

Selects the last node which appears in this file.

] (`global-next-node`)

Moves forward or down through node structure. If the node that you are currently viewing has a ‘Next’ pointer, that node is selected. Otherwise, if this node has a menu, the first menu item is selected. If there is no ‘Next’ and no menu, the same process is tried with the ‘Up’ node of this node.

[(`global-prev-node`)

Moves backward or up through node structure. If the node that you are currently viewing has a ‘Prev’ pointer, that node is selected. Otherwise, if the node has an ‘Up’ pointer, that node is selected, and if it has a menu, the last item in the menu is selected.

`global-next-node` and `global-prev-node` behave the same as simply scrolling through the file with **Spacebar** and **Del**; see `scroll-behavior` in “Manipulating variables” on page 619 for more information.

g (`goto-node`)

Reads the name of a node and selects it. No completion is done while reading the node name, since the desired node may reside in a separate file. The node must be typed exactly as it appears in the `info` file. A file name may be included as with any node specification, as in the following example.

```
g(emacs)Buffers
```

This input finds the ‘Buffers’ node in the ‘emacs’ `info` file.

C-x, k (`kill-node`)

Kills a node. The node name is prompted for in the echo area, with a default of the current node. *Killing* a node means that `info` tries hard to forget about it, removing it from the list of history nodes kept for the window where that node is found. Another node is selected in the window which contained the killed node.

C-x, C-f (`view-file`)

Reads the name of a file and selects the entire file.

C-x, C-f *filename*, is equivalent to typing `g(filename)*`

C-x, C-b (`list-visited-nodes`)

Makes a window containing a menu of all of the currently visited nodes. This window becomes the selected window, and you may use the standard `info` commands within it.

C-x, b (`select-visited-node`)

Selects a node which has been previously visited in a visible window. This is similar to C-x, C-b followed by ‘m’, but no window is created.

Searching an `info` file

GNU `info` allows you to search for a sequence of characters throughout an entire `info` file, search through the indices of an `info` file, or find areas within an `info` file which discuss a particular topic.

s (`search`)

Reads a string in the echo area and searches for it.

C-s (`isearch-forward`)

Interactively searches forward through the `info` file for a string you type.

C-r (`isearch-backward`)

Interactively searches backward through the `info` file for a string as you type it.

i (`index-search`)

Looks up a string in the indices for this `info` file, and selects the node that the found index entry points to.

, (`next-index-match`)

Moves to the node containing the next matching index item from the last ‘i’ command.

The most basic searching command is ‘s’ (`search`). The ‘s’ command prompts you for a string in the echo area, and then searches the remainder of the `info` file for an occurrence of that string. If the string is found, the node containing it is selected, and the cursor is left positioned at the start of the found string. Subsequent ‘s’ commands show you the default search string within ‘[’ and ‘]’; pressing Enter, instead of typing a new string will use the default search string.

Incremental searching is similar to basic searching, but the string is looked up while you are typing it, instead of waiting until the entire search string has been specified.

Selecting cross references

We have already discussed the ‘Next’, ‘Prev’, and ‘Up’ pointers which appear at the top of a node. In addition to these pointers, a node may contain other pointers which refer you to a different node, perhaps in another `info` file. Such pointers are called *cross references*, or *xrefs* for short.

Parts of a cross reference

Cross references have two major parts: the first part is called the *label*; it is the name that you can use to refer to the cross reference, and the second is the *target*; it is the full name of the node to which the cross reference points.

The target is separated from the label by a colon ‘:’; first, the label appears, and then the target. For instance, the following example’s input shows a cross reference menu, where the single colon separates the label from the target.

```
* Foo Label: Foo Target. More information about Foo.
```

The ‘.’ is not part of the target; it serves only to let `info` know where the target name ends.

A shorthand way of specifying references allows two adjacent colons to stand for a target name, as in the following example.

```
* Foo Commands:: Commands pertaining to Foo.
```

In the previous example, the name of the target is the same as the name of the label, in this case `Foo Commands`.

You will normally see two types of cross references while viewing nodes: *menu* references, and *note* references. Menu references appear within a node’s menu; they begin with a ‘*’ at the beginning of a line, and continue with a label, a target, and a comment which describes what the contents of the node pointed to contains.

NOTE: References appear within the body of the node text; they begin with `*Note`, and continue with a label and a target.

Like ‘Next’, ‘Prev’ and ‘Up’ pointers, cross references can point to any valid node. They are used to refer you to a place where more detailed information can be found on a particular subject.

See “Cross References” in *Texinfo (The GNU Documentation Format)*^{*}, for more information on creating your own Texinfo cross references.

Selecting xrefs

The following lists the `info` commands that operate on menu items.

^{*} *Texinfo (The GNU Documentation Format)* is published by the Free Software Foundation (ISBN 1-882114-12-4).

1 (menu-digit)

2 ...9

Within an `info` window, pressing a single digit, (such as '1'), selects that menu item, and places its node in the current window. For convenience, there is one exception; pressing '0' selects the *last* item in the node's menu.

0 (last-menu-item)

Select the last item in the current node's menu.

m (menu-item)

Reads the name of a menu item in the echo area and selects its node. Completion is available while reading the menu label.

M-x find-menu

Moves the cursor to the start of this node's menu.

The following lists the `info` commands which operate on note cross references.

f (xref-item)

r

Reads the name of a note cross reference in the echo area and selects its node. Completion is available while reading the cross reference label.

Finally, the next few commands operate on both menu or note references.

Tab (move-to-next-xref)

Moves the cursor to the start of the next nearest menu item or note reference in the current node. You can also then use the following command, Return (`select-reference- this-line`), to select the menu or note reference.

M-Tab (move-to-prev-xref)

Moves the cursor to the start of the nearest previous menu item or note reference in the current node.

Enter / Return(`select-reference-this-line`)

Selects the menu item or note reference appearing on the line where the cursor currently is.

Manipulating multiple windows

A *window* is a place to show the text of a node. Windows have a view area where the text of the node is displayed, and an associated *mode* line, which briefly describes the node being viewed.

GNU *info* supports multiple windows appearing in a single screen; each window is separated from the next by its modeline. At any time, there is only one *active* window, that is, the window in which the cursor appears. There are commands available for creating windows, changing the size of windows, selecting which window is active, and for deleting windows.

The mode line

A *mode* line is a line of inverse video which appears at the bottom of an *info* window. It describes the contents of the previously displayed window; this information includes the name of the file and node appearing in that window, the number of screen lines it takes to display the node, and the percentage of text that is above the top of the window. It can also tell you if the indirect tags table for this *info* file needs to be updated, and whether or not the *info* file was compressed when stored on disk. The following is a sample mode line for a window containing an uncompressed file named ‘*dir*’, showing the node ‘*Top*’.

```
-----Info: (dir)Top, 40 lines --Top-----
             ^^^  ^^^             ^^
             (file)Node #lines where
```

When a node comes from a file which is compressed on disk, this is indicated in the mode line with two small ‘z’s. In addition, if the *info* file containing the node has been split into subfiles, the name of the subfile containing the node appears in the modeline as well.

```
--zz-Info: (emacs)Top, 291 lines --Top-- Subfile: emacs-1.Z-
```

When *info* makes a node internally, such that there is no corresponding *info* file on disk, the name of the node is surrounded by asterisks (*). The name itself tells you what the contents of the window are; the following sample mode line shows an internally constructed node showing possible one possible completion.

```
-----Info: *Completions*, 7 lines --All-----
```

Window commands

To view more than one node at a time, *info* can display more than one window. Each window has its own mode line (see “The mode line” on page 612) and history of nodes viewed in that window (for information on *history-node*, see “Selecting a new node” on page 607).

C-x, o (*next-window*)

Selects the next window on the screen. The echo area can *only* be selected if it is

already in use, and you have left it temporarily. Normally, `C-x, o` simply moves the cursor into the next window on the screen, or if you are already within the last window, into the first window on the screen. Given a numeric argument, `C-x, o` moves over that many windows. A negative argument causes `C-x, o` to select the previous window on the screen.

`M-x` (`prev-window`)

Selects the previous window on the screen. This is identical to `C-x, o` with a negative argument.

`C-x, 2` (`split-window`)

Splits the current window into two windows, both showing the same node. Each window is one half the size of the original window, and the cursor remains in the original window. The variable, `automatic-tiling`, can cause all of the windows on the screen to be resized for you automatically; for more information on `automatic-tiling`, see “Manipulating variables” on page 619.

`C-x, 0` (`delete-window`)

Deletes the current window from the screen. If you have made too many windows and your screen appears cluttered, this is the way to get rid of some of them.

`C-x, 1` (`keep-one-window`)

Deletes all of the windows excepting the current one.

`Esc C-v` (`scroll-other-window`)

Scrolls the other window, in the same fashion that ‘`C-v`’ might scroll the current window. Given a negative argument, the *other* window is scrolled backward.

`C-x, ^` (`grow-window`) Grows (or shrinks) the current window. Given a numeric argument, grows the current window that many lines; with a negative numeric argument, the window is shrunk instead.

`C-x, t` (`tile-windows`)

Divides the available screen space among all of the visible windows. Each window is given an equal portion of the screen in which to display its contents. The variable `automatic-tiling` can cause `tile-windows` to be called when a window is created or deleted. For more information on `automatic-tiling`, see “Manipulating variables” on page 619.

The Echo Area

The *echo area* is a one line window which appears at the bottom of the screen. It is used to display informative or error messages, and to read lines of input from you when that is necessary. Almost all of the commands available in the echo area are identical to their Emacs counterparts, so please refer to GNU Emacs documentation for greater depth of discussion on the concepts of editing a line of text.

The following briefly details the commands that are available while input is being

read in the echo area.

C-f (echo-area-forward)

Moves forward a character.

C-b (echo-area-backward)

Moves backward a character.

C-a (echo-area-beg-of-line)

Moves to the start of the input line.

C-e (echo-area-end-of-line)

Moves to the end of the input line.

M-f (echo-area-forward-word)

Moves forward a word.

M-b (echo-area-backward-word)

Moves backward a word.

Cd (echo-area-delete)

Deletes the character under the cursor.

Del (echo-area-rubout)

Deletes the character behind the cursor.

C-g (echo-area-abort)

Cancels or quits the current operation. If completion is being read, **C-g** discards the text of the input line which does not match any completion. If the input line is empty, **C-g** aborts the calling function.

RET (echo-area-newline)

Accepts (or forces completion of) the current input line.

C-q (echo-area-quoted-insert)

Inserts the next character verbatim; for example, so you can insert control characters into a search string.

printing character (echo-area-insert)

Inserts the character.

M-Tab (echo-area-tab-insert)

Inserts a Tab character.

Ctrl-t (echo-area-transpose-chars)

Transposes the characters at the cursor.

The next group of commands deal with killing and yanking text. For an in depth discussion of killing and yanking, see “Killing and Moving Text” in the *GNU Emacs Manual*^{*}.

^{*} *GNU Emacs Manual* is published by the Free Software Foundation (ISBN 1-882114-03-5).

M-d (echo-area-kill-word)

Kills the word following the cursor.

M-Del (echo-area-backward-kill-word)

Kills the word preceding the cursor.

C-k (echo-area-kill-line)

Kills the text from the cursor to the end of the line.

C-x, Del (echo-area-backward-kill-line)

Kills the text from the cursor to the beginning of the line.

C-y (echo-area-yank)

Yanks back the contents of the last kill.

M-y (echo-area-yank-pop)

Yanks back a previous kill, removing the last yanked text first.

Sometimes when reading input in the echo area, the command that needed input will only accept one of a list of several choices. The choices represent the *possible completions*, and you must respond with one of them. Since there are a limited number of responses you can make, `info` allows you to abbreviate what you type, only typing as much of the response as is necessary to uniquely identify it. In addition, you can request `info` to fill in as much of the response as is possible; this is called *completion*.

The following commands are available when completing in the echo area.

Tab (echo-area-complete)

SPACEBAR

Inserts as much of a completion as is possible.

? (echo-area-possible-completions)

Displays a window containing a list of the possible completions of what you have typed so far. For example, say the available choices are the following if you typed an 'f', followed by '?'.

```
bar foliate
food forget
```

bar foliate
food forget

Possible completions would contain the choices which begin with 'f'.

foliate food forget

Pressing Spacebar or Tab would result in 'fo' appearing in the echo area, since all of the choices which begin with 'f' continue with 'o'. Now, typing 'l' followed by pressing Tab results in 'foliate' appearing in the echo area, since that is the only choice which begins with 'fol'.

Esc Ctrl-v (echo-area-scroll-completions-window)

Scrolls the completions window, if that is visible, or, if not, the *other* window.

Printing out nodes

You may wish to print out the contents of a node as a quick reference document for later use. `info` provides you with a command for printing.

In general, we recommend that you use the C program utility, `Makeinfo`, to create an `info` file from a Texinfo source file and then, by using the command, `texify`, format the document and print the DVI (*Device Independent*) file.

See *Texinfo (The GNU Documentation Format)*^{††} manual for more details.

`info` also provides you with a command for printing.

M-x `print-node`

Pipes the contents of the current node through the command in the environment variable, `INFO_PRINT_COMMAND`. If the variable doesn't exist, the node is simply piped to `lpr`.

^{††} *Texinfo (The GNU Documentation Format)* is published by the Free Software Foundation (ISBN 1-882114-12-4).

Miscellaneous commands

GNU `info` contains several commands which self-document GNU `info` as the following discussions help to clarify.

M-x `describe-command`

Reads the name of an `info` command in the echo area and then displays a brief description of what that command does.

M-x `describe-key`

Reads a key sequence in the echo area, and then displays the name and documentation of the `info` command which a given key sequence invokes.

M-x `describe-variable`

Reads the name of a variable in the echo area and then displays a brief description of what the variable affects.

M-x `where-is`

Reads the name of an `info` command in the echo area, and then displays a key sequence which can be typed in order to invoke that command.

C-h (`get-help-window`)

?

Creates (or moves into) the window displaying ***Help***, and places a node containing a quick reference card into it. This window displays the most concise information about GNU `info` available.

h (`get-info-help-node`)

Tries hard to visit the node `(info)Help`. The `info` file, '`info.texi`,' distributed with GNU `info`, contains this node. Of course, the file must first be processed with `makeinfo`, and then placed into the location of your `info` directory.

The following are the commands for creating a numeric argument.

C-u (`universal-argument`)

Starts (or multiplies by 4) the current numeric argument. '`C-u`' is a good way to give a small numeric argument to cursor movement or scrolling commands.

`C-u`, `C-v` scrolls the screen 4 lines, while '`C-u`, `C-u`, `C-n`' moves the cursor down 16 lines.

M-1 (`add-digit-to-numeric-arg`)

M-2... M-9

Adds the digit value of the invoking key to the current numeric argument. Once `info` is reading a numeric argument, you may just type the digits of the argument, without the `M` prefix. For example, you might give `C-1` a numeric argument of 32 by using the keystroke sequence, `C-u`, 3, 2, `C-1` or `M-3`, 2, `C-1`.

`C-g` is used to abort the reading of a multi-character key sequence, to cancel

lengthy operations (such as multi-file searches) and to cancel reading input in the echo area.

C-g (`abort-key`)

 Cancels current operation.

q (`quit`)

 Exits `info`.

If the operating system tells `info` that the screen is 60 lines tall, and it is actually only 40 lines tall, the following is a way to tell `info` that the operating system is correct.

M-x `set-screen-height`

 Reads a height value in the echo area and sets the height of the displayed screen to that value.

Finally, `info` provides a convenient way to display footnotes which might be associated with the current node that you are viewing:

Esc C-f (`show-footnotes`)

 Shows the footnotes (if any) associated with the current node in another window.

 You can have `info` automatically display the footnotes associated with a node when the node is selected by setting the variable, `automatic-footnotes`; for more information on `automatic-footnotes`, see “Manipulating variables” on page 619.

Manipulating variables

GNU **info** contains several variables whose values are looked at by various **info** commands. You can change the values of these variables, and thus change the behavior of **info** to more closely match your environment and **info** file reading manner.

M-x set-variable

Reads the name of a variable, and the value for it, in the echo area and then sets the variable to that value. Completion is available when reading the variable name; often, completion is available when reading the value to give to the variable, but that depends on the variable itself. If a variable does not supply multiple choices to complete over, it expects a numeric value.

M-x describe-variable

Reads the name of a variable in the echo area and then displays a brief description of what the variable affects.

What follows is a list of the variables that you can set in **info**.

automatic-footnotes

When set to **on**, footnotes appear and disappear automatically. This variable is **on** by default. When a node is selected, a window containing the footnotes which appear in that node is created, and the footnotes are displayed within the new window. The window that **info** creates to contain the footnotes is called ‘*Footnotes*’. If a node is selected which contains no footnotes, and a ‘*Footnotes*’ window is on the screen, the ‘*Footnotes*’ window is deleted. Footnote windows created in this fashion are not automatically tiled so that they can use as little of the display as is possible.

automatic-tiling

When set to **on**, creating or deleting a window *resizes* other windows. This variable is **off** by default. Normally, typing ‘Ctrl-x, 2’ divides the current window into two equal parts. When **automatic-tiling** is set to **on**, all of the windows are resized automatically, keeping an equal number of lines visible in each window. There are exceptions to the automatic tiling; specifically, the windows ‘*Completions*’ and ‘*Footnotes*’ are *not* resized through automatic tiling; they remain their original size.

visible-bell

When set to **on**, GNU **info** attempts to flash the screen instead of ringing the bell. This variable is **off** by default.

Of course, **info** can only flash the screen if the terminal allows it; in the case that the terminal does not allow it, the setting of this variable has no effect.

However, you can set the **errors-ring-bell** variable to **off** to make Info

perform quietly.

`errors-ring-bell`

When set to `On`, errors cause the bell to ring. The default setting of this variable is `On`.

`gc-compressed-files`

When set to `On`, `info` garbage collects files which had to be uncompressed. The default value of this variable is `Off`. Whenever a node is visited in `info`, the `info` file containing that node is read into core, and `info` reads information about the tags and nodes contained in that file. Once the tags information is read by `info`, it is never forgotten. However, the actual text of the nodes does not need to remain in core unless a particular `info` window needs it. For non-compressed files, the text of the nodes does not remain in core when it is no longer in use. But decompressing a file can be a time consuming operation, and so `info` tries hard not to do it twice. `gc-compressed-files` tells `info` it is okay to garbage collect the text of the nodes of a file which was compressed on disk.

`show-index-match`

When set to `On`, the portion of the matched search string is highlighted in the message which explains where the matched search string was found. The default value of this variable is `On`. When `info` displays the location where an index match was found, (for more information on `next-index-match`, see “Searching an info file” on page 609), the portion of the string that you had typed is highlighted by displaying it in the inverse case from its surrounding characters.

`scroll-behaviour`

Controls what happens when forward scrolling is requested at the end of a node, or when backward scrolling is requested at the beginning of a node. The default value for this variable is `Continuous`. There are three possible values for this variable:

`Continuous`

Tries to get the first item in this node’s menu, or failing that, the ‘Next’ node, or failing that, the ‘Next’ of the ‘Up’. This behavior is identical to using the ‘]’ (`global-next-node`) and ‘[’ (`global-prev-node`) commands.

`Next Only`

Only tries to get the ‘Next’ node.

`Page Only`

Simply gives up, changing nothing. If `scroll-behaviour` is `Page Only`, no scrolling command can change the node that is being viewed.

`scroll-step`

The number of lines to scroll when the cursor moves out of the window. Scrolling happens automatically if the cursor has moved out of the visible portion of the node text when it is time to display. Usually the scrolling is done so as to put the cursor on the center line of the current window. However, if the variable `scroll-step` has a nonzero value, **info** attempts to scroll the node text by that many lines; if that is enough to bring the cursor back into the window, that is what is done. The default value of this variable is 0, thus placing the cursor (and the text it is attached to) in the center of the window. Setting this variable to 1 causes a kind of *smooth scrolling* which some people prefer.

`ISO-Latin`

When set to `on`, **info** accepts and displays ISO Latin characters. By default, Info assumes an ASCII character set. `ISO-Latin` tells **info** that it is running in an environment where the European standard character set is in use, and allows you to input such characters to **info**, as well as display them.

3

Making `info` files from Texinfo files

`makeinfo` is the program that builds `info` files from Texinfo files. Before reading this documentation, you should be familiar with reading `info` files. If you want to run `makeinfo` on a Texinfo file prepared by someone else, this documentation contains most of what you need to know. However, to write your own Texinfo files, you should also read *Texinfo (The GNU Documentation)*[†].

[†] *Texinfo (The GNU Documentation Format)* is published by the Free Software Foundation (ISBN 1-882114-12-4).

Controlling paragraph formats

In general, `makeinfo` *fills* the paragraphs that it outputs to the `info` file. Filling is the process of breaking up and connecting lines such that the output is nearly justified. With `makeinfo`, you can control the following.

- The width of each paragraph (the *fill-column*).
- The amount of indentation that the first line of the paragraph receives (the *paragraph-indentation*).

makeinfo command line options

The following command line options are available for `makeinfo`.

- I **dir**
Adds **dir** to the directory search list for finding files which are included with the `@include` command. By default, only the current directory is searched.
- D **var**
Defines the `texinfo` flag, **var**. This is equivalent to `@set var` in the `Texinfo` file.
- U **var**
Makes the `Texinfo` flag, **var**, undefined. This is equivalent to `@clear var` in the `Texinfo` file.
- error-limit num
Sets the maximum number of errors that `makeinfo` will print before exiting (on the assumption that continuing would be useless). The default number of errors printed before `makeinfo` gives up on processing the input file is 100.
- fill-column num
Specifies the maximum right-hand edge of a line. Paragraphs that are filled will be filled to this width. The default value for `fill-column` is 72.
- footnote-style style
Sets the footnote style to `style`. `style` should either be `'separate'` to have `makeinfo` create a separate node containing the footnotes which appear in the current node, or `'end'` to have `makeinfo` place the footnotes at the end of the current node.
- no-headers
Suppress the generation of menus and node headers. This option is useful together with the `'--output file'` and `'--no-split'` options (see following options) to produce a simple formatted file (suitable for printing on a dumb printer) from `Texinfo` source. If you do not have `TEX`, these two options may allow you to get readable hard copy.
- no-split
Suppress the splitting stage of `makeinfo`. In general, large output files (where the size is greater than 70k bytes) are split into smaller subfiles, each one approximately 50k bytes.
If you specify `'--no-split'`, `makeinfo` will not split up the output file.
- no-pointer-validate
- no-validate
Suppress the validation phase of `makeinfo`. Normally, after the file is processed, some consistency checks are made to ensure that cross references can be resolved, and so forth. See “What makes a valid info file?” on page 627.

`--no-warn`

Suppress the output of warning messages. This does not suppress the output of error messages, simply warnings. You might want this if the file you are creating has texinfo examples in it, and the nodes that are referenced don't actually exist.

`--no-number-footnotes`

Suppress the automatic numbering of footnotes. The default is to number each footnote sequentially in a single node, resetting the current footnote number to 1 at the start of each node.

`--output file`

`-o file`

Specify that the output should be directed to **file** instead of the file name specified in the `@setfilename` command found in the Texinfo source. **file** can be the special token '-', which specifies standard output.

`--paragraph-indent num`

Sets the paragraph indentation to **num**. The value of **num** is interpreted as follows:

- A value of 0 (or 'none') means not to change the existing indentation (in the source file) at the start of paragraphs.
- A value less than zero means to indent paragraph starts to column zero by deleting any existing indentation.
- A value greater than zero is the number of spaces to leave at the front of each paragraph start.

`--reference-limit num`

When a node has many references in a single Texinfo file, this may indicate an error in the structure of the file. **num** is the number of times a given node may be referenced before `makeinfo` prints a warning message about it (with `@prev`, `@next`, or `@note` appearing in an `@menu`, for example).

`--verbose`

Causes `makeinfo` to inform you as to what it is doing. Normally `makeinfo` only outputs text if there are errors or warnings.

`--version`

Displays the `makeinfo` version number.

What makes a valid `info` file?

If you have not used ‘`--no-pointer-validate`’ to suppress validation, `makeinfo` will check the validity of the final `info` file. Mostly, this means ensuring that nodes you have referenced really exist. What follows is a complete list of what is checked.

- If a node reference such as `Prev`, `Next` or `Up` is a reference to a node in this file (meaning that it is not an external reference such as ‘`(DIR)`’), then the referenced node must exist.
- In a given node, if the node referenced by the `Prev` is different than the node referenced by the `Up`, then the node referenced by the `Prev` must have a `Next` which references this node.
- Every node except `Top` must have an `Up` field.
- The node referenced by `Up` must contain a reference to this node, other than a `Next` reference. Obviously, this includes menu items and followed references.
- If the `Next` reference is not the same as the `Next` reference of the `Up` reference, then the node referenced by `Next` must have a `Prev` reference pointing back at this node. This rule still allows the last node in a section to point to the first node of the next chapter.

Defaulting the Prev, Next, and Up pointers

If you write the `@node` commands in your Texinfo source file without `Next`, `Prev`, and `Up` pointers, `makeinfo` will fill in the pointers from context (by reference to the menus in your source file). Although the definition of an info file allows a great deal of flexibility, there are some conventions that you are urged to follow. By letting `makeinfo` default the `Next`, `Prev`, and `Up` pointers you can follow these conventions with a minimum of effort.

A common error occurs when adding a new node to a menu; often the nodes which are referenced in the menu do not point to each other in the same order as they appear in the menu.

`makeinfo` node defaulting helps with this particular problem by not requiring any explicit information beyond adding the new node (so long as you do include it in a menu). The node to receive the defaulted pointers must be followed immediately by a sectioning command, such as `@chapter` or `@section`, and must appear in a menu that is one sectioning level or more above the sectioning level that this node is to have.

What follows is an example of how to use this feature.

```
@setfilename default-nodes.info
@node Top
@chapter Introduction
@menu
* foo:: the foo node
* bar:: the bar node
@end menu
@node foo
@section foo
this is the foo node.
@node bar
@section Bar
This is the Bar node.
@bye
```

The previous input produces the following output.

```
Info file default-nodes.info, produced by makeinfo, -*- Text -*-

from input file default-nodes.texinfo.

File: default-nodes.info, Node: Top

Introduction *****
* Menu:
```



```
* foo:: the foo node
* bar:: the bar node

File: default-nodes.info, Node: foo, Next: bar, Up: Top
foo
===

this is the foo node.

File: default-nodes.info, Node: bar, Prev: foo, Up: Top

Bar
===

This is the Bar node.
```


Index

Symbols

- - file, command line option 602
 - help, command line option 603
 - node, command line option 602
 - output, command line option 602
 - subnodes, command line option 603
 - version, command line option 603
- - directory, command line option 602
- as prefix character 364
- to add or subtract hexadecimal constants or symbols 193
- 513, 540, 607
- !, for assembler 35
- !, indicator character 508
- ", for assembler 35
- #, for assembler 35
- #, for comments 316
- #else directive 518
- #endif directive 518
- #ifdef directive 518
- #ifndef directive 518
- \$ 438, 480
- \$ in symbol names 36
- \$\$ variables 480
- \$\$, automatic variable 438
- \$(440, 440, 481, 481
- \$(%D) variables 481
- \$(%D) variants 440
- \$(%F) variables 481
- \$(%F) variants 440
- \$(%D) variables 480
- \$(%D) variants 440
- \$(%F) variables 480
- \$(%F) variants 440
- \$(+D) variables 481
- \$(+F) variables 481
- \$(?D) variables 481
- \$(?D) variants 440

- \$(?F) variables 481
- \$(?F) variants 440
- \$(@D) variables 480
- \$(@D) variants 440
- \$(@F) variables 480
- \$(@F) variants 440
- \$(^D) variables 481
- \$(^D) variants 440
- \$(^F) variables 481
- \$(^F) variants 440
- \$(addprefix prefix, names...) 399
- \$(addsuffix suffix, names...) 399
- \$(basename names...) 398
- \$(dir names...) 398
- \$(filter pattern ...,text) 396
- \$(filter-out pattern ...,text) 397
- \$(findstring find,in) 396
- \$(firstword names...) 399
- \$(join list1, list2) 399
- \$(notdir names ...) 398
- \$(patsubst pattern, replacement, text) 395
- \$(sort list) 397
- \$(strip string) 396
- \$(subst from, to, text) 395
- \$(suffix names ...) 398
- \$(wildcard pattern) 400
- \$(word n, text) 399
- \$(words text) 399
- \$* variables 480
- *, automatic variable 439
- \$+ variables 480
- +, automatic variable 439
- \$? variables 480
- \$?, automatic variable 438
- \$@ variables 480
- \$@, automatic variable 438
- \$@, for varying commands 340
- \$^ variables 480
- \$^, automatic variable 439
- % 520
- % character 331, 342
- % in pattern rules 436
- %, for line group formatting 520, 522
- %, for conversion specifications 520
- %=, for line group formatting 520
- %>, for line group formatting 520
- %c' C', for line group formatting 520, 522
- %c', for line group formatting 522
- %c':' 522
- %c'' 522
- %c'O', for line group formatting 520
- %d, for line group formatting 521
- %L, for line group formatting 522
- %l, for line group formatting 522
- %o, for line group formatting 521
- %X, for line group formatting 521
- %x, for line group formatting 521
- (191
- (, or its synonym, --start-group 191
- (markers 513
- (info)Help 617
-) markers 513
- *, comments 249
- *, in comments 261
- + as prefix character 364
- + or %, as operators 63
- + prefix 391
- +, adding a textline 510
- +, indicator character 508
- +, line continuation character 261
- +, to add or subtract hexadecimal constants or symbols 193
- += syntax 457
- +=, appending values to variables 380
- , 609
- , deleting a text line 510
- , for specifying left-justification 521
- , in comments 261
- , in macro definitions 77
- , indicator character 508
- , option 10
- , special token, standard output 626

--, standard input file 10
., as current address 56
., in names 65
., indicator of end of output in ed 515
.ABORT 68
.abort 67
.app-file 67, 68, 72
.ascii 68
.asciz 69
.balign 69
.block 83
.bss output section 206
.byte 69
.c file 423
.code32 directive 111
.comm 69
.data 28, 70
.data output section 206
.def 70
.def and .endef directives 58
.def/.endef 85
.DEFAULT 338, 444, 448
.desc 70
.dfloat, for Vax 176
.dim 58
.double 70, 110, 175
.eject 70
.else 71
.endef 70, 71
.endef directives 58
.endif 71, 73
.endm 76, 77
.endr 74, 79
.equ 71
.exitm 77
.EXPORT 73
.EXPORT_ALL_VARIABLES 339, 360
.extern 72
.ffloat, for Vax 176
.file 72
.fill 72
.float 72, 110, 175
.global 73
.globl 73
.hword 73
.ident 73
.if 73
.ifdef 73
.ifndef 73
.IGNORE 338
.include 74
.include directive 34
.include directives 20
.int 74, 110
.irp 74
.irpc 74
.lcomm 75
.lcomm on HPPA 75
.lflags 75
.line 58, 67, 75
.list 76
.ln 76
.long 76, 110
.macro 76
.mri 77
.nolist 76, 77
.octa 77
.org 78
.PHONY 338, 457
.POSIX 361
.PRECIOUS 338, 357
.psize 79
.quad 79, 110
.rept 79
.sbttl 79
.scl 58, 80
.section 80
.set 81
.set on HPPA 81
.short 82

- .SILENT 338
- .single 72, 82, 110
- .size 58, 82
- .space 82
 - warnings 82
- .space and .subspace directives 44
- .space on AMD 29K 83
- .stab 275
- .stab directives 164
- .stab.excl 275
- .stab.index 275
- .stabd 83
- .stabs 83
- .stabs 83
- .string 84
- .subspace directives 44
- .SUFFIXES 338, 419, 427, 446
- .tag 58, 85
- .text 28, 85, 206
- .text section 206
- .tfloat 110
- .title 85
- .type 85
- .val 85
- .word 110, 176
- / markers 513
- /* to */ 35
- /lib 196
- /usr/lib 196
- ., for labels 52
- :=, for simply expanded variables 378
- =, equals sign 191
- =, for setting variable values 378
- =, in expressions 53
- load-average 418
- max-load 418
- > 607
 - >, last-node 607
- > markers 513
- >>>>>>>, marker 540

- ? 54, 615
- ?, object file format specific 264
- @ as prefix character 364
- @@, line in unified format 511
- @menu, for referencing nodes 626
- @next, for referencing nodes 626
- @node 628
- @note, for referencing nodes 626
- @prev, for referencing nodes 626
- _ 283
- __SYMDEF 451
- __main 218
- | markers 513
- |, for assembler 35

Numerics

- 0, in Info windows 611
- 1...9, in Info windows 611
- 80386
 - jump instructions 109
 - machine dependent options 105
 - opcode prefixes 108
 - register naming 107
 - registers 109
 - syntax for Intel or AT&T 109
- 80387
 - 16-, 32-, and 64-bit integer formats 110

A

- A 191, 241, 264, 278
- a 264, 272, 280, 502
- a, adding new archive files 259
- a, listing 17
- a.out 187, 195, 246
- a.out file format 57
- a.out object file format 249, 263
- a.out or b.out 27, 75
- a.out, or b.out, for Intel 80960 12
- a.out-hp300bsd 294
- a.out-sunos-big 294

- Aarchitecture 191, 241
- aborting execution of make 353
- abort-key 618
- ABSOLUTE 250
- absolute address 192, 240
- absolute addresses 250
- absolute expression 73
- absolute expressions
 - result, size, value 72
- absolute locations of memory 195
- absolute number 59
- abstraction payback 243
- ACA 241
- ad, to omit debugging 17
- add text lines 505
- add-digit-to-numeric-arg 617
- ADDLIB 261
- ADDMOD 261
- addprefix 479
- addr2line 285
- address 222, 269
- address displacements 28
- addresses 43, 204
- addressing modes 6
- add-section 270
- addsuffix 479
- adjust-section-vma 270
- adjust-start 270
- adjust-vma 270
- adjust-vma= 272
- adjust-warnings 270
- advanced features of variables 373
- aggregates 248
- ah, high-level language listing 17
- al, assembly listing 17
- ALIAS 250
- align 68
- alignment 213, 221, 246
- alignment constraints 197
- alignment information
 - propagating 246
- alignment of input section 213
- alignment, specified 68
- all 280
- all 409
- all target 466
- all-header 276
- allocatable 204
- alphanumerics 283
- alternate file names 512
- AMD 29K
 - command-line options 92
 - floating-point numbers 94
 - machine directives 95
 - macros 93
 - opcodes 95
 - register names 93
 - special characters 93
- an, forms preprocessing 17
- ar 257
 - command line control 257
 - configuring 257
 - controlling with library scripts 260
 - creating indexes to symbols 257
- ar (creates, modifies, extracts archives) 257
- ar program 452
- ar -s 277
- AR variable 432
- ar, brief description 255
- arbitrary sections, supported 197
- ARC
 - command line options 98
 - floating point 98
 - machine directives 99
- architecture selection 294
- architecture values 296
- architecture, defined 296
- architecture= 274
- architecture-independent files 472
- archive
 - contents 257
 - defined 257

Index

- extraction 257
- formats 257
- generating index of symbols 277
- members 257
- archive file symbol tables 459
- archive files 187, 200, 449
- archive header information 272
- archive libraries 241
- archive member 438
- archive member target 447
- archive member targets 440
- archive suffix rules 454
- archive-header 272
- archive-maintaining program 433
- archives 255
- archives, specifying 201
- ARFLAGS variable 433
- argument, for assembly 59
- arguments 62, 602
- arguments, delimited 61
- arithmetic expressions 38
- arithmetic operands 62
- ARM
 - local labels with 22
- arm 283
- AS variable 432
- as, symbol table listing 17
- ASFLAGS variable 433
- ASM68K, ASM960 23
- assembler 433
- assembler & linker 43
- assembler directives 65
 - periods 65
- assembler programs 428
- assembly
 - 35
 - ! 35
 - " 37
 - # 35
 - \$ 36
 - + or -, as constants 40
 - , escape characters 39
 - . 36, 56
 - _ 36
 - | 35
 - ' 39
 - ' , escape characters 39
 - a.out file format 57
 - absolute expressions 35
 - addresses 43
 - argument 59
 - asm statements 34
 - base-10 flonums 41
 - bignums 40
 - characters 39
 - COFF format 58
 - comments 35
 - constants 38
 - converting character constants 34
 - directives and . 65
 - directives for 80386
 - jump instructions 109
 - double-quotes 38
 - escapes 38
 - exponents as number constants 41
 - expression 59
 - expressions 40
 - f option 34
 - file processing 34
 - floating point number 40
 - flonum 40
 - flonums, how they work 41
 - include file handling 34
 - logical line 35
 - machine-independent syntax 33
 - macro processing 34
 - newline character () 37
 - newlines 34, 37
 - operator 59
 - prefix operator (-) 40
 - preprocessing 34
 - processing, using integers 41

- removing comments 34
- section 43
- separators 37
- SOM format 58
- statement 37
- string 35, 38
- symbol 36
- symbol attributes 57
- syntax 33
- type 57
- value 57
- whitespace 34
- assignment operators 210
- assume-new= option 411
- assume-new=file option 420
- assume-old=file option 419
- AT&T
 - Link Editor Command Language syntax 187
 - opcodes 106
 - register operands 106
- AT&T System V/386 assembler syntax 106
- attributes 58
- authors, listing 420
- automatic 403
- automatic variable 403, 455, 456
- automatic variables 331, 368, 377, 378, 438, 480
- automatic-footnotes 619
- automatic-tiling 619
- auxiliary symbol table 58

B

- B 264, 278
- b 191, 269, 273, 294, 496
- b (or its synonym, -format) 191
- b , or -format 202
- b option 417
- b, adding new archive files 259
- b, as a command in Info 604
- b, as a meta-command in Info 605

- b, backwards 54
- b.out 68, 246
- backslash (37
- backslashes 395
- backslashes (331
- backslash-newline 378
- backslash-newline (306
- backup file names 580
- backward-char 604
- backward-word 604
- bal 241
- balx 241
- BASE 250
- basename 479
- bash script 465
- beginning-of-line 604
- beginning-of-node 604
- Berkeley
 - size 278
- BFD 202, 243, 246
 - abstraction payback 243
 - archive files 243
 - back end conversions 245
 - back end, symbol table 247
 - canonical relocation record 248
 - coff-m68k variant 250
 - conversion and output, information loss 244
 - converting between formats 243
 - files 247
 - internal canonical form of external formats 246
 - object and archive files 243
 - object file format 243
 - optimizing algorithms 243
 - section 247
 - subroutines 245
- BFD libraries 191, 195, 243
- bfdname 268
- big endian 246
- big endian COFF 246
- bignum 40, 62
- binary and text file comparison 501

- binary file utilities
 - target, architecture, linker emulation 293
- binary files 267
- binary format for output object file 195
- binary input files 202
- binary input files, specifying 191
- binary integer 39
- binary integers 39
- binary utility, defined 257
- bindir 471
- Bison 308, 467, 470
- BISONFLAGS 470
- blank and tab difference suppression 496
- blank lines 497
- blocks 43
- bra 240
- braces, unmatched 394
- branch optimization, special instructions 89
- brief difference reports 500
 - brief option 500
- Broken pipe signal 353
- BSD 456, 471
- bsd 265
- BSD 4.2 assembler 33, 39
- BSD make 460
- bsr 240
- bss section 50
- bss subsections 50
- Bstatic 192
- Bsymbolic 192
- buffer.h 306
- bugs, reporting 461
- built-in expansion functions 398
- built-in implicit rule, overriding 443
- built-in implicit rules 318, 424, 432, 458
- built-in rules 432
- byte= 269
- bytes=min-len 280

C

- C 264
- c 249
- C compilation 441
- C function headings, showing 511
- C if-then-else output format 518
- C option 417
- C option, multiple specifications 417
- C preprocessor 428, 457
- C programs 427
- C++ Annotated Reference Manual 283
- C++ function names 265, 283
- C++ function overloading 283
- C++ programs 427
- c++filt 283
 - inverse mapping for linkers 283
- c++filt option symbol 284
- c++filt symbol 284
- c++filt, brief description 255
- C, C++, Prolog regular expressions 511
- c, creating new archive files 259
- cal 241
- callj 149
- calx 241
- canned sequence 364
- canned sequence of commands 364
- canned sequence with the define directive 364
- canonical format, destination format 247
- canonicalization 246
- canonicalized 294
- case 10
- case difference suppression 498
- case... instruction 176
- catalogue of predefined implicit rules 427
- CC variable 432
- CFLAGS 370, 380, 383, 414, 424, 440, 470
- CFLAGS variable 433
- chain, defined 435
- change commands 505, 515
- change compared files 515

- changed-group-format= 519
- change-leading-char 270
- character constants 37, 38
- check 410
- check target 468
- child processes 353
- CHIP 250
- clean 305, 410
- clean target 467
- CLEAR 261
- clobber 410
- cmp 501
- cmp command options 563
- CO variable 432
- COFF 187
- coff 289
- COFF debuggers 247
- COFF files, sections 246
- COFF format
 - auxiliary symbol attributes 58
- COFF format output 85
 - with .def and .endef 71
- COFF object file format 197
- COFF output 28, 44
- COFF section names 246
- coff-m68k 250
- COFLAGS variable 433
- columnar output 514
- command 305
 - next-line 604
- command line 192, 316, 403
- command line options 602
- command line override 511
- command sequence, defining 364
- command.h 306
- command.o 308
- command-line options 10, 15, 456
- commands
 - key sequence 604
- commands, echoing 351
- commands, empty 366
- commands, errors 355
- commas 70, 72, 394
- commas, in commands 261
- comments 205, 314, 372
- common symbol 199
- common symbols 216, 263
- common variables 247
- compare-and-branch instructions 150
- comparing text files with null characters 501
- compilation errors 415
- compilation-and-linking 441
- compile 356
- COMPILE.c 430
- COMPILE.x 430
- compiler 218
- compiler switches 383
- compiling 187, 427
- compiling a program, testing 415
- compiling C programs 329
- complex makefile, example 483
- complex values 379
- computed variable names 373
- concatenation 11
- conditional 385
- conditional directive 360
- conditional example 386
- conditional execution 457
- conditional syntax 388
- conditionals, omitting 17
- config/target.mt 297
- configuration triplet 294
- configure 474
- configuring 294
- conflict 539
- constraints 441
- constructor 218
- CONSTRUCTORS 218
- constructors 197
- context 505
- context format and unified format 507

- context output format 511
- Continuous 620
- controlling a linker 189
- conventions for writing the Makefiles 463
- COPY 220
- copying and translating object files 255
- copyright 420
- counter relative instructions 240
- CPP variable 432
- CPPFLAGS 470
- CPPFLAGS variable 434
- c-r 609
- CREATE 261
- create, modify, and extract from archives 255
- CREATE_OBJECT_SYMBOLS 218
- cross references 610
 - followed by a colon 610
 - label 610
 - menu 610
 - menu, followed by * 610
 - note 610
 - periods within a cross-reference 610
 - pointers 610
 - Next 610
 - Prev 610
 - Up 610
 - specifying with adjacent colons 610
 - target 610
- cross-references
 - (DIR) 627
 - errors in Texinfo files 626
 - Next 627
 - pointers
 - Next 628
 - Prev 628
 - Up 628
 - Prev 627
 - Up 627
- c-s 609
- csh script 465
- CTANGLE variable 433
- Ctrl-a 614
- Ctrl-a, in Info windows 604
- Ctrl-b 614
- Ctrl-b, in Info windows 604
- Ctrl-d 614
- Ctrl-e 614
- Ctrl-e, in Info windows 604
- Ctrl-f 614
- Ctrl-f, in Info windows 604
- Ctrl-g 614, 618
- Ctrl-h 617
- Ctrl-k 615
- Ctrl-l 606
- Ctrl-n 604
- Ctrl-p 604
- Ctrl-q 614
- Ctrl-t 614
- Ctrl-u 617
- Ctrl-v 606
- Ctrl-w 606
- Ctrl-x, ^ 613
- Ctrl-x, 0 613
- Ctrl-x, 1 613
- Ctrl-x, 2 613
- Ctrl-x, Delete 615
- Ctrl-x, t 613
- Ctrl-y 615
- cursor, moving in an Info file 604
- CVTRES 289
- CWEAVE variable 433
- CXX variable 432
- CXXFLAGS variable 433

D

- D 265, 273
- d 273
- d 607
- d option 417
- D var 625
- d, -dc, -dp 192

- d, deleting modules in archive 258
- dashes in arguments 257
- dashes in option names 191
- data objects 263
- data section 204
- data segments 194
- datadir 472
- data-section data 28
- debug 288
- debug option 417
- debugger symbol information 197
- debugging 269, 273
- debugging 70, 264
 - addresses in output files 248
 - mapping between symbols 248
 - relocated addresses of output files 248
 - source line numbers 248
 - symbols 282
- debugging directives 17
- debugging information 265, 273
- debugging mips 164
- debugging symbols 9, 268
- debugging symbol-table entries 275
- debug-syms 264
- DEC mnemonics 176
- decimal integer 40
- decimal integers 40
- default 403
- DEFAULT_EMULATION 297
- defaulting to Makeinfo 628
- define 477
- define directive 376, 382
- define directives 381
- define option 457
- defined-only 266
- definitions of variables 316
- defs.h 306, 307
- defsym 191, 192
- DELETE 261
- Delete 614
- delete text lines 505
- Delete, in Info windows 606
- delete-window 613
- demand pageable bit 247
- demand pageable bit, write protected text bit set 247
- demangle 264, 273
- demangle 265
- demangling 265, 283
- demangling low-level symbol names 264
- dependencies 307, 415, 419, 439
- dependencies, analogous 342
- dependencies, duplicate 480
- dependencies, explicit 425
- dependencies, from source files 316
- dependencies, generating automatically 346
- dependency 305
- dependency loop 321
- dependency patterns 425
- dependency tracking (--MD) 26
- dep-patterns 342
- destination format, canonical format 247
- destructors 218
- diagnostics 188
- diff and patch, overview 491
- diff output, example 504
- diff sample files 504
- diff, older versions 505
- diff3 502
- diff3 comparison 533
- difference tables 21
- dir 372, 479
- directives 314, 477
- directories 531
- directories, implicit rules 332
- DIRECTORY 261
- directory option 417, 457
- directory search features 330
- directory search with linker libraries 332
- directory, rules for cleaning 312
- directory, specifying 418

- directory-path 602
- dir-node 607
- disassemble 273
- disassemble-all 273
- disassembling instructions 275
- discard-all 269, 282
- discarded symbols 255
- discard-locals 269, 282
- displaying files to open 198
- displaying information from object files 255
- dist 410, 484
- dist target 468
- distclean 410
- distclean target 467
- documentation 474
- dollar signs 325, 369
- dot (.), in symbol names 54
- double-colon rule 319, 425
- double-colon rules 345
- double-suffix 445
- double-suffix rule 445
- dry-run option 411, 419
- DSECT 220
- DT SONAME field 197
- dummy pattern rules 442
- duplication 309
- dvi target 468
- dynamic linker 193
- dynamic objects 275
- dynamic relocation 274
- dynamic, for symbols 265
- dynamic-linker 193
- dynamic-reloc 274
- dynamic-syms 275

E

- e 193
- e flag, for exiting 347
- e option 418
- echo area 612, 613

- completion 615
- echo command 365
- echo-area-abort 614
- echo-area-backward 614
- echo-area-backward-kill-line 615
- echo-area-backward-kill-word 615
- echo-area-backward-word 614
- echo-area-beg-of-line 614
- echo-area-complete 615
- echo-area-delete 614
- echo-area-end-of-line 614
- echo-area-forward 614
- echo-area-forward-word 614
- echo-area-insert 614
- echo-area-kill-line 615
- echo-area-kill-word 615
- echo-area-newline 614
- echo-area-possible-completions 615
- echo-area-quoted-insert 614
- echo-area-rubout 614
- echo-area-scroll-completions-window 615
- echo-area-tab-insert 614
- echo-area-transpose-chars 614
- echo-area-yank 615
- echo-area-yank-pop 615
- echoing 351
- ecoff-littlemips 294
- ed output format 515
- ed script example 516
- ed script output format 515
- edit 306
- editor 304
- editor, recompiling 304
- EFL programs 460
- ELF executables 193, 196
- ELF output
 - for HPPA 44
- ELF platforms 192
- ELF systems 196
- else 386, 477
- Emacs' compile command 356

Emacs' M-x compile command 415
-embedded-relocs 193
empty command 349
empty commands 321
empty files 557
empty lines of command language 260
empty strings 398
empty target 337
EMUL 297
emulation linker 194
--enable-targets=all 294
encoded C++ symbols 255
END 250, 262
endif 364, 382, 477, 478
endif 386, 477
end-of-file 37
end-of-line 604
end-of-node 605
entry command-line option 207
entry point 207
environment 403
environment override 403
environment variable 202, 318
environment variables 294, 383
environment, redefining 404
--environment-overrides option 418
environments 195
environments that override assignments 383
error messages 11, 13
--error-limit num 625
errors 188
errors, continuing make after 418
errors, ignoring, in commands 418
errors, tracking 420
errors-ring-bell 620
Esc Ctrl-f 618
Esc Ctrl-v 613, 615
exec_prefix variable 471
executable 204, 441
executables 304

execution of programs 193
execution time 197
exit status 355
exit status of make 408
--expand-tabs option 522
explicit dependencies 425
explicit rules 314
export 478
export directive 360, 383
exporting variables 359
expression 71, 213
 absolute numbers 59
 empty 60
expression, for assembly 59
extern int i 199
--extern-only 265
EXTRACT 262

F

-F 193, 268, 294
-f 265, 274
f 611
-F bfdname 281
-F option 511
-f option 418
-F regexp option 511
F variants 439
f, forwards 54
-f, or --file for naming makefile 315
f, truncating names in archive 259
failed shell commands 415
false conditionals, omitting 17
fatal error 317
fatal messages 559
fatal signal 357
FC variable 432
FFLAGS variable 434
file 403
file differences, summarizing 500
file name endings 423

- file names 398
- file names, showing alternate 512
- file section sizes and total size 255
- file, reading 418
- file=file option 418
- file-header 274
- filename 195
- files 247
- files, marking them up to date 420
- fill-column num 625
- filter 479
- find-menu 611
- findstring 479
- first-node 607
- firstword 479
- FLAGS 470
- fldt, load temporary real to stack top 110
- floating point numbers 70, 72
- flonum 40, 62
- flonums 70, 72
 - on ARC (D, F, R, or S) 40
 - on H8/300, H8/500, Hitachi SH, and AMD 29K (D, F, P, R, S, or X) 40
 - on HPPA (E) 40
 - on Intel 960 (D, F, or T) 40
- FN, for line group formatting 522
- footnote-style style 625
- for 32-, 64-, and 80-bit formats 110
- force 321
- FORCE_COMMON_ALLOCATION 192
- foreach 480
- foreach function 401
- FORMAT 250
- format 193
- format= 265, 278
- format=compatibility 278
- format=format 283
- Fortran programs 427
- forward ed script output format 516
- forward-char 604
- forward-word 604

- from-file 508
- fstpt, store temporary real and pop stack 110
- full lines 591
- full-contents 275
- function 204
- function call 393
- function call syntax 394
- function invocations 374
- function overloading 283
- function references 378
- functions 393
- functions, transforming text with 457

G

- G 193
- g 193, 265, 268, 282
- g 608
- g, compiler debugging 17
- g, debugging 17
- gap-fill 269
- gc-compressed-files 620
- generating index to archive 255
- generating merged output directly 544
- GET variable 432
- get-help-window 617
- get-info-help-node 617
- getting started 600
- GFLAGS variable 434
- global array 263
- global constructors
 - warnings 200
- global int variable 263
- global optimizations 195
- global or static variable 204
- global symbol 210
- global symbols 192, 199
- global, static, and common variables 247
- global-next-node 608
- global-prev-node 608
- gnu 283

compiler 193
 GNU assembler
 16-bit mode 111
 gnu assembler 5
 .o 12
 contributors 181
 -D, for compatibility 18
 -f, faster preprocessing 19
 -I, adding path 20
 introduction 5
 invoking summary, example 8
 -K, difference tables 21
 listing output 17
 Microtec 23
 -o option 12
 object file formats 7
 output file 12
 pseudo-ops 7
 syntax 7
 gnu compiler 283
 GNU Emacs 580
 GNUTARGET 202, 294
 goals 308, 409
 goto-node 608
 GP register 193
 GROUP, for loading objects 191
 grow-window 613

H

-h 274
 h 617
 -h option 418
 hash table 83
 --header 274
 headerfile 287
 --header-file=headerfile 287
 headings 511
 hello.o 190
 --help 266, 271, 274, 278, 280, 281, 284, 288
 -help 193

--help option 418, 457
 hexadecimal integer 40
 hexadecimal integers 40
 history-node 607
 Hitachi h8/300 195
 addressing modes 123
 command-line options 122
 floating point numbers 124
 machine directives 125
 opcodes 126
 register names 123
 special characters 123
 Hitachi h8/500
 addressing modes 129
 command-line options 128
 opcodes 131
 register names 129
 special characters 129
 Hitachi SH
 addressing modes 141
 command-line options 140
 floating point numbers 142
 machine-dependent directives 143
 opcodes 143
 register names 140
 special characters 140
 hook functions 297
 how to use Info 600
 HP9000 Series 800 Assembly Language Reference
 Manual 44
 HPPA 22, 133
 assembler directives 135, 136
 assembler syntax 135
 command-line options 134
 floating point numbers 40
 floating-point numbers 135
 local labels with L\$ 22
 SOM and ELF object file formats 133
 HPPA targets with -R 28
 HPPA targets, labels 37
 hunks 495, 515

I

- I 193, 294
- i 195, 197, 274
- i 609
- I bfdname 281
- I dir 625
- I dir option 418
- i interleave 269
- i option 418
- I, for searches 74
- i, inserting new archive files 259
- i386 code, 32-bit 111
- i960 191, 195, 241
 - command-line options 146
 - compare & branch instructions 150
 - opcodes 148
- IEEE 250
- IEEE Standard 1003.2-1992 459
- IEEE Standard 1003.2-1992 (POSIX.2) 457
- ifdef 389, 477
- ifdef=HAVE_WAITPID option 518
- ifeq 386, 388, 477
- ifndef 389, 477
- ifneq 389, 477
- ifnotdef 73
- if-then-else example 523
- if-then-else format 520, 523
- if-then-else output format 518
- ignore option 513
- ignore-case option 498
- ignore-errors 355
- ignore-errors option 418
- ignore-space-change 496
- imperfect patch 554
- implicit intermediate files 456
- implicit rule 310, 329, 344
- implicit rule for archive member targets 451
- implicit rule search algorithm 447
- implicit rule, canceling 443
- implicit rule, cancelling or overriding 427
- implicit rule, using 425
- implicit rules 314, 423
- implicit rules, chains of 435
- implicit rules, predefined 425
- implicit rules, special 435
- include 478
- include directive 457
- include directive 316
- include-dir 316
- includedir 473
- include-dir=dir option 418
- includes 380
- incomplete lines 591
- incomplete lines, merging with diff3 545
- incr 270
- incremental links 193
- incremental searching 609
- index generating 277
- index table, listing 257
- index, listing archives 277
- index-search 609
- indicator characters 508
- infile 268
- infix operators 63
- INFO 220
- info 274
- info 599
- Info commands 600
- Info files, building Texinfo files from 623
- info files, description of 601
- info target 468
- Info, learning to use 600
- info.texi 617
- INFO_PRINT_COMMAND, environment variable 616
- infodir 466, 473
- information loss 246
- information loss with BFD 244
- initial values of variables 377
- initialized data 206
- input files 241

input section 204
input section description 214, 215
input section wildcard patterns 215
INPUT, for loading objects 191
input, standard 257
--input-target 268, 294
--input-target=bfdname 281
insert.c 308
insert.o 308
insertions 508
INSTALL 470
install 410
install target 466
INSTALL_DATA 466, 470
INSTALL_PROGRAM 470
installcheck target 468
installdirs target 469
instruction operands 62
instruction set 6
instructions in hex 275
instructions in symbolic form 275
int i 199
int i = 1 199
integer expression 61
integer formats for 80387 110
integers 39
integers for assembly use 39
Intel syntax 106
interleave 269
--interleave=interleave 269
interleave-1 269
intermediate file 435
internal canonicals 246
internal counters 76
internal fields 197
invariant text 375
isearch-backward 609
isearch-forward 609
ISO-Latin 621

J

-j 274
jmp 240
job slots 353
--jobs= 418
-j 418
--jobs, for simultaneous commands 353
jobs, running 418
jobs, running simultaneously 418
jobs, specifying number 418
join 479
jsr 240
--just-print option 411, 419
--just-print, or -n 319

K

-K 268
-k option 418
-K symbolname 282
kbd.o 306, 308
--keep-going 355
--keep-going option 415, 418
keep-one-window 613
--keep-symbol= 268
--keep-symbol=symbolname 282
key symbol 37
keywords 218
keywords, unrecognized 249
killing and yanking text 614
killing the compiler 357
kill-node 608
ksh script 465

L

-L 191, 194, 241
L 200
-l 191, 241, 265, 272
l 607
-L option 507, 512

- l option 512
- l, modifier for archive files 259
- la, add text command 515
- label 37
- label option 507
- labels 52
 - local 22
- last-menu-item 611
- last-node 607
- last-resort implicit rule 444
- ld
 - BFD libraries for operating on object files 187
 - command-line options 190
 - compiling 187
 - configuring 194
 - configuring with default input format 191
 - controlling a linker 189
 - dynamic libraries 192
 - Link Editor Command Language syntax 187
 - Linker Command Language files 187
 - machine-dependent features 239
 - MRI compatible linker scripts 249
 - MRI linker 249
 - object and archive files 187
 - optimizing 194
 - shared libraries 192
 - shared link 196
 - symbol references 187
 - symbols 247
 - warning messages 249
- ld sections 46
 - absolute 46
 - bss 46
 - data 46
 - named 46
 - text 46
 - undefined 46
- LD_LIBRARY_PATH 196
- LD_RUN_PATH 196
- LD_RUN_PATH 196
- LDEMULATION 297
- LDFLAGS 470
- LDFLAGS variable 434
- leading character 270
- leaf routines 241
- learning Info 600
- left-column option 514
- Lex for C programs 429
- Lex for Ratfor programs 429
- LEX variable 433
- LFLAGS variable 434
- libc.a 190
- libdir 473
- libexecdir 471
- libraries 332
- libraries, shared 274
- libraries, subroutines 257
- library archive, updating 454
- library dependencies 458
- line comment character
 - 680x0 35
 - AMD 29K 35
 - ARC 35
 - h8/300, h8/500 35
 - Hitachi SH 35
 - HPPA 35
 - i960 35
 - SPARC 35
 - Vax 35
 - z8000 35
- line comment characters 35
- line formats 516, 521
- line group formats 518, 519
- line number list 248
- line numbers 248
- linebreaks in Makeinfo 624
- line-numbers 265, 274
- link map 250
- linkage symbols 275
- linker 332, 357, 428
 - addressing 195
 - canonical form 245

- default values 297
 - emulation 297
 - invoking 191
 - m 297
 - object file format 243
 - scripts 191
 - specifying for input target 295
 - specifying for target 295
 - linker & assembler 43
 - linker commands 249
 - linker emulation 293
 - linker input architecture 296
 - linker input target 295
 - linker ld 22
 - linker output architecture 296
 - linker script example 206
 - linker script, defined 203
 - linker scripts commands 209
 - linker, basic concepts 204
 - linker=linker 288
 - linking 187, 470
 - absolute or a relative pathnames 288
 - faster links to library 277
 - linking C++ 197
 - linking C++ programs 198
 - linking libraries 190
 - linking, partial 196
 - Lint Libraries from C, Yacc, or Lex programs 429
 - Linux 471
 - Lisp regular expressions 511
 - LIST 250, 262
 - listing output directives
 - .list, .nolist, .psize, .eject, .title, and .sbtbl 17
 - listing printable strings from files 255
 - listing symbols from object files 255
 - lists 293
 - list-visited-nodes 608
 - literal characters 39
 - little endian 246
 - little endian COFF 246
 - LMA 204, 220
 - ln utilities 470
 - LOAD 250
 - l 418
 - load memory address 204
 - loadable 204
 - loading shared objects 197
 - local labels 22
 - local symbol names 54
 - local symbols 54, 200, 269
 - localstatedir 472
 - locating a line of input 11
 - locating shared objects 196
 - location counter 49, 78, 210
 - active 49
 - logical file 35
 - logical files 11
 - loop names, suffixed 241
 - ls 272
 - ltry 241
 - lucid 283
 - Lucid compiler 283
- M**
- M 194, 250, 257
 - m 194, 274
 - m option 417
 - m, moving members in archive 258
 - M, option for controlling ar script 260
 - m68k
 - options 24
 - machine 274
 - machine instruction sets 89
 - macros, defining 76
 - magic numbers 194, 197, 247
 - main.c 306, 307
 - main.o 306, 307, 346
 - maintainer-clean target 467
 - MAKE 412
 - make 26
 - commands 323

- default goal of rules 324
- targets 323
- make commands, recursive 358
- make options 411
- MAKE variable 358
- MAKE variables 481
- make with no arguments 407
- make, invoking recursively 383
- make, modifying 420
- make, recursive invocations 318
- make, running 407
- make, version number 458
- MAKE_VERSION 458
- Makefile 297
- makefile 418
- Makefile commands 465
- makefile, defined 303
- makefile, example 483
- makefile, naming 315, 408
- makefile, overriding another makefile 321
- makefile, using variables for object files 309
- makefile=file option 418
- MAKEFILES 318, 359
- MAKEFILES variables 481
- makefiles, debugging 421
- makefiles, portable 455
- makefiles, what they contain 314
- MAKEFLAGS 359, 361, 391, 420
- MAKEFLAGS variables 481
- Makeinfo command line options 625
- Makeinfo output, controlling 624
- MAKEINFO variable 433
- MAKELEVEL 360, 371, 458
- MAKELEVEL variables 481
- MAKEOVERRIDES 361
- man1dir 473
- man1ext 474
- man2dir 474
- man2ext 474
- mandir 473
- manext 474
- mangling 283
- manipulating variables 619
- Map 194
- mark conflicts 542
- markers 513
- match-anything pattern rule 444
- match-anything rules 441
- matching 428
- max-load option 353
- membedded-pic option 193
- member name 439
- memory descriptor 245
- memory, symbolic 195
- menu items as options 603
- menu-digit 611
- menu-item 611
- merge commands 549
- merged output format 518
- merged output with diff3 545
- merging two files 518
- messages 559
- Meta- 604
- Meta-> 605
- Meta-1 617
- Meta-2... Meta-9 617
- Meta-b 614
- Meta-d 615
- Meta-Delete 615
- Meta-f 614
- Meta-f, in Info windows 604
- Meta-Tab 614
- Meta-Tab, in Info windows 611
- Meta-v 606
- Meta-x 613
- Meta-x describe-command 617
- Meta-x describe-key 617
- Meta-x describe-variable 617, 619
- Meta-x set-screen-height 618
- Meta-x set-variable 619

Meta-x where-is 617
 Meta-y 615
 Microtec 23
 --minimal 495
 --minimal option 531
 -min-len 280
 mips 161
 (ISA) Instruction Set Architecture 165
 assembling in 32-bit mode 165
 debugging 164
 ECOFF targets 163
 options 162
 r2000 161
 r3000 161
 r4000 161
 r6000 161
 mk program 456, 457
 mod 257
 mode line 612
 modified references 374
 modifiers 257
 modifiers for p 259
 Modula-2 programs 428
 mostlyclean 410
 mostlyclean target 467
 Motorola
 syntax 155
 Motorola 680x0
 addressing modes 155
 compatibility with Sun assembler 157
 floating point formats 156
 machine dependent options 152
 opcodes 157
 options 152
 pseudo opcodes 157
 special characters 159
 syntax 153
 mov.b instructions 240
 move-to-next-xref 611
 move-to-prev-xref 611
 MRI 192

a.out object file format 249
 -c 192
 linker scripts 249
 script files 192
 -T 192
 MRI assembler
 complex relocations 23
 global symbols in common section 23
 specifying pseudo-ops 23
 MRI compatibility 262
 MRI compatibility mode 23
 MRI librarian program 257
 MRI linker 249
 multiple -C options 417
 multiple members in a archive file 457
 multiple patches 558
 multiple targets 340
 mv utilities 470

N

-N 194, 269
 -n 194, 265, 283
 n 607
 -n flag 420
 -n option 411, 419, 516
 -N symbolname 282
 NAME 251
 name 274
 name patterns 316
 names of files, alternates 512
 nested references 375
 nested variable reference 374
 nests of recursive make commands 420
 NetWare Loadable Module 287
 Netware Loadable Module (NLM) 256
 new jobs commands 418
 --new-file= option 411
 --new-file=file option 420
 --new-group-format= 519
 new-lc 78

- Next Only 620
 - next-index-match 609
 - next-line 604
 - next-node 607
 - NLM Development and Tools Overview 287
 - NLM Software Developer's Kit 287
 - NLM targets 287
 - nlmconv 256, 287
 - converting relocatable object files to NetWare Loadable Modules 287
 - version 288
 - NLMLINK 287
 - nm 263, 283
 - common symbols as uninitialized data 263
 - linking 263
 - specifying for architecture 296
 - specifying for target 295
 - symbol type 263
 - symbol value 263
 - uninitialized data section 263
 - nm --print-armap 257, 277
 - nm -s 257, 277
 - nm, brief description 255
 - nm, listing symbols from object files 263
 - NMAGIC 194
 - no-adjust-warnings 270
 - no-builtin-rules option 419, 427
 - NOCROSSREFS 222
 - node
 - Top 627
 - Up 627
 - node, selection of 606
 - no-demangle 265
 - nodes, description of 600
 - no-headers 625
 - noinhibit-exec 194
 - no-keep-going option 420
 - no-keep-memory 194
 - NOLOAD 220
 - non-fatal messages 559
 - non-global symbols in source files 269
 - non-text files 501
 - no-number-footnotes 626
 - no-pointer-validate 625
 - no-print-directory option 420
 - no-print-directory, disabling 363
 - normal diff output format 505
 - no-sort 265
 - no-split 625
 - no-strip-underscores 283
 - notdir 479
 - no-validate 625
 - Novell 287
 - no-warn 626
 - number constants 39
 - bignums 39
 - flonums 39
 - integers 39
 - numeric expressions 38
 - numeric-sort 265
- O**
- O 268
 - o 195, 251, 264, 280
 - O bfdname 281, 287
 - o file 626
 - o file option 419
 - o, preserving member dates of archive files 259
 - Oasys 246
 - Oasys compilers 273
 - objcopy 267
 - specifying for architecture 296
 - specifying for input target 294
 - specifying for output target 295
 - S-records 267
 - temporary files 267
 - objcopy, brief description 255
 - objcopy, for copying contents of object files 267
 - objcopy, leading char 270
 - objdump 195, 204, 243, 272
 - displaying data for object files 272

- specifying for architecture 296
 - objdump -h 274
 - objdump -i 296
 - objdump, brief description 255
 - object and archive files 243
 - object file 204, 304, 426
 - object file (-o) 27
 - object file format 204, 243
 - object file formats 208, 218
 - object file names 309
 - object file output 32
 - object files 187, 191
 - object files, required 200
 - object formats, alternatives, supported 191
 - object-code format 273
 - object-file format 193
 - objects 309
 - objfile 278
 - objfile ... 278
 - objfile..., linking 191
 - octa- (16 bytes), defined 78
 - octal integer 39
 - octal integers 39
 - offsets in an address 64
 - oformat 195, 295
 - oformat srec, --oformat=sre 191
 - oformat, -oformat 191
 - old files 419
 - old-fashioned suffix rule 454
 - old-file option 457
 - old-file=file option 419
 - old-group-format= 519
 - oldincludedir 473
 - OMAGIC 194, 197
 - OPEN 261, 262
 - opened input files 198
 - operand 62
 - operation 257
 - operator, for assembly 59
 - option 418
 - options
 - exceptions to arguments 190
 - repeating 190
 - options to objcopy 271
 - options variable 470
 - options with machine dependent effects 195
 - ORDER 251
 - origin 403, 480
 - origin function 403
 - outfile 268
 - OUTPUT 195
 - output file 626
 - output file 12
 - output files 262
 - output section 204, 213, 216
 - output section address 213
 - output to a file 602
 - output, understanding easier 362
 - OUTPUT_ARCH 296
 - OUTPUT_FORMAT 195, 250, 295
 - OUTPUT_OPTION 431
 - output-name 251
 - output-target= 268
 - output-target=bfdname 281, 287
 - overlap 539
 - overlapping contents comparison 508
 - OVERLAY 220, 222
 - overlay description 222
 - overloaded functions 283
 - overridden by command line argument 414
 - override 403, 478
 - override define 478
 - override directive 381
 - override directives with define directives 381
 - overview 491
- P**
- P 265
 - p 265
 - p 257, 607
 - p option 419

- p, specifying archive members to output file 258
- pad-to 269
- Page Only 620
- page-numbered and time-stamped output 529
- paginate option 512
- paginating diff output 529
- paragraph-indent num 626
- paragraphs, controlling formats with
 - Makeinfo 624
- parallel execution 353, 453
- parallelism 456
- parentheses 450
- parentheses, unmatched 394
- Pascal compilation 441
- Pascal programs 427
- passing down variables 359
- patch, merging with 551
- patches in another directory 579
- patches with whitespace 554
- patsubst 373, 456, 478
- pattern matches 441
- pattern rule 344, 436, 447
- pattern rules 316, 456
- pattern rules, examples 437
- pattern rules, writing 419
- PC variable 433
- performance of diff 531
- periods, in assembler directives 65
- PFLAGS variable 434
- phdr 221
- phony target 334
- phony targets 457
- physical files 11
- PIC 193
- PIC code 193
- portability 265
- posix 265
- POSIX.2 standard 459
- POSIX.2 standard output format 265
- Posix-compliant systems, comparing 501
- predefined implicit rules 425

- predefined rules 427
- prefix operators 63
- prefix variable 471
- PREPROCESS.S 430
- PREPROCESS.x 430
- preprocessor 470
- prev-line 604
- prev-node 607
- prev-window 613
- print 327, 410
- print the commands 419
- print-armap 265
- print-data-base option 419
- print-directory 362
- print-directory option 420
- printf conversion specification 521
- print-file-name 264, 280
- printing 200, 616
- printing character 614
- printing input filenames 198
- processing symbol tables 245
- program version 280
- program-counter relative instructions 241
- PROVIDE 211
- pseudo-ops 7
- PUBLIC 251
- punctuation 368

Q

- q 618
- q option 411, 419, 500
- q, quick appends to archive 258
- q, update operation 257
- question mode 419
- question option 411, 419
- question, or -q 319
- questions 559
- quiet option 419
- quit 618
- quoted numeric string 213

quotes 72
 quotes, escape characters 39

R

-R 191, 195, 268, 274
 -r 192, 196, 265, 274
 r 611
 -r option 419, 427
 -R sectionname 281
 -R warnings 28
 -R, for data-section data 28
 r, replacing member... into archive 258
 radix 266
 --radix= 266, 279
 --radix=number 279
 --radix=radix 280
 ranlib 257, 277, 452, 465
 ranlib, brief description 255
 Ratfor programs 427
 rc 289
 rc, replace text command 516
 RCS 430, 435
 --rcs option 516
 RCS or SCCS files 319
 RCS output format 516
 RCS rules, terminal 430
 rd, delete text command 516
 realclean 410
 recompilation 308
 recompilation of dependent files 413
 recompilation, avoiding 413
 recompiling 304
 --recon option 411, 419
 recursion 481
 recursive commands 456
 recursive expansion 370
 recursive invocations of make 318, 417
 recursive make 362, 420
 recursive make commands 420
 recursive make invocations 457

recursive use of make 358, 391
 recursively expanded variable 370
 recursively expanded variables 376, 378
 redefinition 404
 redistribution of input sections 197
 redraw-display 606
 redundant context 509
 reference 199
 --reference-limit num 626
 references, binding 192
 region 221
 registers 6
 regular expression suppression 499
 regular expressions, matching comparisons 511
 -relax 195, 240, 241
 relaxing address modes 240
 relinking 307
 --reloc 274
 relocatable object modules 257
 relocatable output 198
 relocatable output file 192
 relocation 43, 248, 274
 relocation and symbol information 268
 relocation entry 265
 relocation level 248
 relocations 195, 197
 relpos argument, before archive specification 259
 remaking target files 423
 --remove-leading-char 271
 --remove-section= 268
 --remove-section=sectionname 281
 replace text lines 505
 res 289
 resolving references to constructors 198
 Resource Compiler 289
 result 72
 -retain-symbols-file 195
 --reverse-sort 266
 Revision Control System 516
 RFLAGS variable 434

- rm utilities 470
- RM variable 433
- rm, errors 307
- rpath 195
- rpath-link 196
- rpath-link option 196
- rule 305
- rule chaining 456
- rule syntax 325
- rule without dependencies or commands 336
- rules 314
- rules for make 305
- run-time addresses 43
- runtime linker 196
- run-yacc 364

S

- S 195, 197, 268, 275, 282
- S 250
- s 195, 197, 265, 275, 281
- s 257, 609
- s format 283
- S option 420
- s option 419
- s, writing object-file index to archive 260
- SAVE 261, 262
- saving a changed file 546
- sbindir 471
- SCCS file 459
- SCCS files 430, 458
- SCCS or RCS files 319
- scripts, specifying 191
- scroll-backward 606
- scroll-behavior 606, 608
- scroll-behaviour 620
- scroll-forward 606
- scrolling through node structure 606
- scroll-other-window 613
- scroll-step 621
- sdiff options 588

- search 609
- search directories 190
- search paths 330
- search through indices of info files 609
- search, implicit rule 425
- SEARCH_DIR 194
- searches 20
- searching 609
 - incremental 609
- SECT 251
- section 43
 - location counter 49
 - offsets in address 64
 - undefined, undefined U 45
- section addresses 272
- section attributes 219
- section contents 204
- section fill 221
- section headers 274
- section headings 510
- section sizes 278
- section= 274
- section=flags 270
- section-header 274
- sectionname 268, 281
- sectionname=filename 270
- SECTIONS 222
- sections 247
 - bss 44
 - data 44
 - internal use 47
 - text 44
 - text and data 48
- sections differences 510
- sections, assembly addresses 43
- selecting unmerged changes 541
- select-reference-this-line 611
- select-visited-node 608
- semicolon 352, 426
- semicolons 205

- serial execution 353
- set-section-flags 270
- set-start 270
- sh script 465
- shar 410, 484
- shared 197
- shared libraries 197
- shared objects, locating 196
- sharedstatedir 472
- SHELL 352, 359, 383
- shell 480
- shell commands 352
- shell commands, failed 415
- shell file 316
- shell function 405
- shell metacharacters 519
- SHELL variable 464
- SHELL variables 481
- show-c-function option 511
- show-footnotes 618
- show-function-line option 511
- show-function-line=regex option 511
- show-index-match 620
- show-raw-insn 275
- side by side comparison of files 513
- side by side format 514
- side-by-side option 514
- silent 351
- silent operation 419
- silent option 419
- SIMPLE_BACKUP_SUFFIX environment variable 580
- simply expanded variables 370, 378
- simultaneous commands 353
- single-suffix 445
- single-suffix rule 445
- size 72, 278
 - command line options 278
 - formats 279
 - specifying for architecture 296
 - specifying for target 295
- size, brief description 255
- size-sort 266
- Solaris 2.0 275
- SOM format for HPPA 58
- SOM or ELF output
 - for HPPA 44
- SONAME 197
- soname 197
- sort 479
- sort-common 197
- sorting 197
- source 275
- source 11
- source file 304, 426
- source file information 268
- source files 316
 - non-global symbols 269
- source, source program 11
- space character 510
- space or tab characters 497
- SPACEBAR 615
- SPACEBAR, in Info windows 606
- spaces 316
- SPARC
 - floating point numbers 170
 - machine directives 171
 - options 170
 - v6, v7, v8, v8+, v9, v9a 170
- sparclite 170
- special built-in target names 338
- special prefix characters 364
- special target 427
- specifying a goal 409
- specifying output file names 195
- speed-large-files option 531
- split-by-file 197
- split-by-reloc 197
- split-window 613
- sreclir 474
- S-record 267
- S-records 217, 273

Index

- stabs 275
- standard mnemonics 6
- standard output streams 262
- start address 251, 270
- start-address= 275
- start-group archives--end-group 201
- static pattern rule 340, 439
- static pattern rule, syntax 342
- static pattern rules 342, 458
- static variable 204
- static variables 247
- statistics 29
- statistics 197
- stats 197
- stem 439, 441
- stem of a target name 342
- stop option 420
- stop-address= 275
- storage boundary 78
- string 72
- string constants, literals 38
- strings 241, 280
 - determining contents of non-text files 280
 - printing character sequences 280
 - specifying for architecture 296
 - specifying for target 295
- strings, brief description 255
- strip 281, 478
 - modifying files named in arguments 281
 - specifying for input target 294
 - specifying for output target 295
 - version 282
- strip -v 282
- strip, brief description 255
- strip-all 268, 281
- strip-debug 268, 282
- strip-symbol= 269
- strip-symbol=symbolname 282
- strip-underscores 283
- strip-unneeded 268
- sub-make 359
- sub-make options 361
- subroutine libraries for linking 449
- subroutines 241
- subsections 48
- subst 394, 478
- subst function 340
- substitution references 373
- substitution variable reference 347
- suffix 479
- suffix list 427
- suffix rules 439, 445, 456
- SUFFIXES variables 481
- summary output format 500
- Sun
 - compatibility with Motorola 680x0 153
- SunOS 196
- suppress-common-lines option 514
- SVR4 471
- switching formats 192
- symbol
 - debugging 264
 - descriptor 70
 - descriptors with .desc statements 57
 - global (external) 263
 - in the text (code) section 264
 - line number records 248
 - name 264
 - name setting 58
 - read only 264
 - small objects 264
 - type 57, 263
 - undefined 264
 - weak 264
 - with ld 73
- symbol characters, using 36
- symbol information 197, 247, 265
- symbol names 247, 273, 451
- symbol pointer 248
- symbol references 187
- symbol string table, duplications 198

symbol table 204
 symbol table and relocation entries 276
 symbol tables
 processing 245
 symbol tables, caching 194
 symbolic debuggers 83
 symbolname 268, 282
 symbols 62, 204, 247
 local 200
 members 277
 non-global, removing 282
 processing 282
 relocatable object files 277
 removing 281
 removing local, compiler-generated 282
 value 57
 warnings 199
 symbols as uninitialized data 263
 symbols, auxiliary attributes 57
 symbols, gaps 197
 --syms 275
 syntax and pseudo-op handling 23
 syntax error 360
 synthesizing instructions 240
 sysconfdir 472
 System V 456
 size 278
 System V make 459
 sysv 265

T

-T 275
 -t 198, 275
 t 607
 -T headerfile 287
 -t option 411, 420, 522
 -t radix 266, 280
 t, table, listing contents of archive 259
 t, temporary real 110
 Tab 615

tab and blank difference suppression 496
 tab character 305
 Tab, in Info windows 611
 tabs, unallowed 316
 tabstop alignment 528
 TAGS 410
 tags 467
 TAGS atrget 468
 TANGLE variable 433
 tar 410
 TARGET 192
 --target 294
 target
 object file format 294
 supporting multiple architectures 294
 target empty commands 426
 target file, modified 420
 target pattern 343, 425, 439, 442
 target system, specifying 293
 --target= 268, 273, 279
 --target=bfdname 279, 280, 281
 targets, required 466
 -Tbss 198, 274
 -Tdata 198, 274
 temp 380
 terminal match-anything pattern rule 444
 terminal rules 437
 test 410
 testing compilation 415
 TEX 429
 TEX and Web 429
 TEX variable 433
 TEXI2DVI 468
 TEXI2DVI variable 433
 Texinfo 429, 468
 Texinfo and Info 429
 Texinfo regular expressions 511
 text 206
 text and binary file comparison 501
 text and data sections, setting 194
 text files, comparing 501

- text manipulation functions 477
- text option 502
- text section 28
- then- part of if-then-else format 520
- tiling, automatic 613
- time-stamped output 529
- to-file 508
- toggle-wrap 606
- top-node 607
- touch command 420
- touch option 411, 420
- traditional formats 198
- traditional-format 198
- transitive closure 456
- triplet 294
- Ttext 198, 274
- two column output comparison 514
- two-suffix rule 445
- type descriptor 248
- type information 247

U

- u 191, 198, 266
- u 607
- U var 625
- u, replaces new archive files 260
- Ultrix 471
- unchanged-group-format= 520
- undefined 403
- undefined section 45
- undefined symbol 265
- undefined symbols 198
- undefined-only 266
- underscore 271, 273
- underscore (_), in symbol names 54
- underscores 283
- unexport 478
- unified format 507
- unified output format 509
- uninitialized data 206

- uninstall target 467
- universal-argument 617
- unmatched braces 394
- unmatched parentheses 394
- unmerged changes, selecting 541
- unprintable character 280
- updates 507
- updating MAKEFILES 319
- updating software 561
- up-node 607
- uppercase usage in variable names 368
- Ur 197, 198

V

- v 199, 266, 271, 277
- v 30, 31, 199, 265, 271, 277, 280
- v option 420
- v, for verbose version 260
- V, version numbering 260
- val 270
- valid Info file 627
- value 72
- variable assignments 457
- variable definitions 314
- variable names 368
- variable names, computing 373
- variable reference 394
- variable references 374, 458
- variable values 359
- variable, undefined 421
- variable, using 367
- variable's value 367
- variables 309, 314, 316, 477
- variables and functions 367
- variables for overriding commands 470
- variables for overriding options 470
- variables used in implicit rules 432
- variables used in implicit rules, classes 432
- variables, adding more text to definitions 379
- variables, manipulating 619

variables, setting 378
 variables, specifying 377
 variables, values 377
 VAX 273
 Vax
 command line options 174
 conversion of floating point numbers 175
 debugging 174
 displacement sizing character 178
 floating point constants 176
 JUMPify, for longer branches 174
 JUMPs 174
 opcodes, pseudo opcodes 176
 options for VMS 174
 register names 178
 symbol table 174
 temporary file generation 174
 Vax bit-fields 33
 verbatim variable definitions 457
 VERBOSE 261, 262
 --verbose 194, 198, 271, 282, 626
 --version 266, 271, 275, 280, 282, 284, 288, 626
 -version 30, 199
 version 277
 version information 603
 version of make, printing 420
 --version option 420, 457
 VERSION_CONTROL environment variable 580
 view-file 608
 virtual memory address 204
 visible-bell 619
 VMA 204, 213
 VPATH 330, 397, 431, 455, 464
 vpath 478
 vpath directive 330
 vpath search 458
 VPATH variable 330
 VPATH variables 481
 VxWorks 195

W

-W columns option 514
 -W file option 420
 -W option 411
 -w option 420
 -Wa, passing arguments 16
 -warn-common 199
 -warn-constructors 200
 warning message 421
 warning messages 317
 warnings 13, 199
 -warn-once 200
 --warn-undefined-variables 362
 --warn-undefined-variables option 421
 warranty 420
 WEAVE variable 433
 --what-if= option 411
 --what-if=file option 420
 white space 193
 white space characters 496
 white space markers 513
 whitespace 60, 205, 316
 whitespace, in commands 261
 whitespace, using 366
 --wide 276
 --width=columns option 514
 wildcard 480
 wildcard characters 325
 wildcard expansion 326
 wildcard function 329
 wildcard patterns 215
 wildcard pitfalls 328
 wildcards 450
 windows
 automatic-tiling 613
 in Info, manipulating 612
 manipulating 613
 multiple 612
 Windows Resource Compiler 289
 Windows utilities 289

Index

windres 289
.p2align 78
word 479
words 479
write protected text bit set 247
writing a semicolon 426

X

-X 200, 269, 282
-x 200, 269, 282
x 259
x, extracting members from archive 259
xrefs 610

Y

-y option 514
Yacc 308, 364, 435
Yacc for C programs 429
YACC variable 433
YACCR variable 433
yanking text 614
YFLAGS variable 434

Z

-Z, object file output warning 32
ZMAGIC 247